

Jienluv2bake - change history

Every version of the backoffice app, from v54 to v265 - newest first. The green 'Engine v###' pill on the More screen tells you which build a phone is running.

About this history: it honestly begins at version 54, when the two-phone shared-cloud era started (2 September 2026). The app grew through earlier versions before that, but they were never written down version by version, so they are left out rather than invented.

What changed in each version of the backoffice app, newest first. Each version number is the "Engine" you can see on the app's **More** screen, so you can always tell which build a phone is running.

01 Oct 2026 - engine v265, THE CONSIGNMENT BOX IS A PARCEL'S, SO IT IS DRAWN WITH THE PARCEL (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "after i get price from lalamove, click use this fee, it closes the window, but user might confuse as at that page there is stick fields that user have not fill in, maybe that field or button should not be there? this create confuse"

What was happening. On the **+ New order** card the courier half drew the **Courier tracking number (optional)** box on every courier order, whatever kind of courier it was. A consignment number is the thing a carrier gives you for a **parcel**, and a Lalamove trip is not a parcel - so on a trip that box was empty, it was not about anything on the card, and it stood beside a delivery price that was already finished. That is what read as an order with something still to fill in.

What changed on screen. The consignment box now belongs to the parcel, so it is drawn with the parcel and nowhere else. Take a Lalamove price and the card under it is the address, the pin, the courier charge, the price and the delivery note - no empty box waiting to be filled in. Name a parcel carrier and the consignment box appears **under that carrier**, which is where a consignment number comes from in the first place.

The order of the courier half on the card is now the address with its pin, then the parcel carrier, then its consignment box, then the courier charge, then the price. Only the last two positions are new: the parcel used to sit under the tracking box, and the tracking box used to sit on its own, above everything to do with a parcel.

Nothing became unreachable. The parcel carrier picker is still on the card, in the same place, saying **Not a parcel - nothing recorded** until you pick one - it is the press that brings the consignment box out. A number you had already typed is never hidden from you either: a box that holds something is always drawn, so the card can never keep a consignment number you cannot see or edit.

A trip's own number is untouched. When you book a real Lalamove trip from the order, the live share link is written into that same slot by the booking - the box is a parcel's, and the slot is still one slot, exactly as it was.

Swapping carrier does not move you. Because the consignment box comes and goes with the carrier list, the card now has to decide, each time you pick a carrier, whether the box should be standing. It only touches it when the answer actually changes - so if you are halfway through typing a consignment number and you change your mind about who is carrying it, the box you are typing in stays put and stays focused, rather than being thrown away and drawn again under your finger.

The Edit window and the Note / tracking box keep the layout they had. Their consignment box still stands above their parcel carrier, and the hint there still says so ("the consignment number goes in the tracking box above"). The two screens were allowed to differ before and they are allowed to differ now; only the **+ New order** card changed.

What did NOT change. The delivery note, the courier charge and who bore it, the delivery address with its pin, **Get a delivery price** and **Use this fee**, the parcel carrier list, the hand-over press, and the rule that nothing at all is written until you press **Place Order**.

Nothing to run in Supabase, and one push. This is app code only, exactly like v264.

01 Oct 2026 - engine v264, A TAP ON THE MAP DOES NOT MOVE THE PIN (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "click on the map should not move the pin, only dragging the pin will"

What changed on screen. A tap on the map no longer moves the pin. Only dragging the pin itself moves it. This is true of both maps you use to place a pin: the pin map in the order form (the one behind **Move this pin** and **Done moving**), and the pin window that opens when you set a doorstep or the bakery's pickup pin.

Why it matters on a phone. That map is a 200-pixel strip inside a card you scroll. Before this, any accidental touch on it moved the pin to wherever your finger landed **and** pulled the map onto that spot - two things at once, from one touch you did not mean. The pin window even said so in its own words all along: *"Look the address up, then drag the pin to the exact door."* Dragging is what it promised, and dragging is now the only thing that moves it.

The one place a tap still works, and why it is not an exception. A map with **nothing on it yet** has no pin to move and nothing to drag - so there, a tap places the first pin. That is the state the pin window opens in when an order has no point at all, and taking that ability away would leave you with a map that does nothing. Once a pin exists, a tap does nothing at all.

What did NOT change. Looking the address up, the list of other matches, pasting a Google Maps link or a pair of numbers, pressing **Use this spot**, and the **Move this pin** button all work exactly as before. So does zooming - the plus and minus, the two-finger pinch and the double-tap.

What was deliberately left alone. The shop's own map, the one your customer uses at checkout, is untouched. For a customer the map is often the only way to place their pin at all (they have no coordinate box and no list of matches), and a map a customer can only aim at by dragging a small marker with a thumb is the one thing that would make them give up. If you want that map changed too, say so - it is a one-line change.

01 Oct 2026 - engine v263, A DAY YOU DO NOT DELIVER IS DARKER STILL (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "can make the quite day more solid?" - asked straight after seeing v262.

What changed on screen. A day you do not deliver is now a solid, dark day. v262 had already taken the fade off it, but it was still faint - **3.50:1** against the card. It is now **6.71:1**, very nearly twice as strong again. The colour is exactly **halfway between the app's quiet grey and its ink**, so it reads as a real day rather than a whisper, while still looking nothing like a day you deliver.

The three steps are further apart now, in the same order. The past is still the faintest thing on the calendar, a day you do not deliver is still the middle step, and a day you deliver is still the darkest. The middle step has simply moved a long way down towards the bottom. A day you deliver is still **twice as dark** as a day you do not - **13.45:1** against **6.71:1** - so the two can never be mistaken for each other.

What did NOT change. The past grey itself is untouched: still the shop's own faint grey, exactly as you asked for it to stay. A day you deliver is exactly as dark as it was. So are the **green pill** on a delivery day, the **green tint** on a day a product sells, the **brown** day you are looking at, and today's halo.

One thing worth knowing about, deliberately left alone. On a product's own availability calendar there is a second, rarer kind of quiet: a day you have **marked as selling** but which the bakery does not deliver. That one is kept faded on purpose - a green tint turned right down - because it means "this mark cannot become an order", which is something to notice rather than something to read easily. It is untouched by this version. Say the word if you would rather it read stronger too.

Why 6.71 and not darker. The same rule draws the product availability calendar, where it is the only thing separating **a day you can sell on** from **a day you cannot**. Pushed much further, those two would stop being different pictures. Measured on that calendar: a day you can sell still reads **13.45:1** against a day you cannot at **6.71:1**, which is a clear step and not a blur.

Nothing else moved. No number, price, booking count or order changed. The database was not touched, and there is nothing to upload - this is the app's own picture of a calendar.

01 Oct 2026 - engine v262, A DAY YOU DO NOT DELIVER IS NOW A SOLID DAY, AND THE PAST IS REALLY ONE GREY THIS TIME (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "I still feel that the greyed and the non grey contrast is not big?" - and then, when offered the choice, **make the non-grey more solid.**

You were right, and v261 had not finished its own job. v261 put the shop's grey in the right place, but it left the older **half-transparent** rule standing on a day you do not deliver. That half-transparency is applied to whatever colour the day ends up with, so the past was still being drawn in **two shades** - and worse, a past day you do not deliver came out **fainter than a future day you do not deliver**, which is backwards. Measured on the real Orders calendar: a quiet day sat at **1.98:1** against the card and a past quiet day at **1.22:1**, so the two were almost the same picture.

What changed on screen. A day you do not deliver is no longer faded - it is a solid, plainly readable day that is simply quieter than a day you do deliver. On the real grid that took it from **1.98:1 to 3.50:1**, very nearly twice as strong. And because nothing fades the calendar cell any more, **the past is now genuinely one shade:** a past day you deliver and a past day you do not are the identical colour on screen, which is what v261 was meant to give you. Three clear steps now, everywhere: the past faintest, a day you do not deliver solid in the middle, a day you deliver darkest of all.

What did NOT change. The shop's past grey itself is untouched - you said you liked it, so it stays exactly as faint as the shop's is. Anything on a past day that is telling you something still keeps its own colour, so on a day that has gone the **green pill** on a delivery day stays green and so does the **green tint** on a day a product sells. The day you are looking at still turns brown, on every calendar. Today's halo is untouched.

Why this is the right lever and not the other one. Between making the past darker and making everything else stronger, the second is the one that keeps the shop's calm look you asked for while still giving the past something solid to stand against.

Nothing else moved. No number, price, booking count or order changed. The database was not touched, and there is nothing to upload - this is the app's own pictures of its calendars.

01 Oct 2026 - engine v261, EVERY CALENDAR NOW GREYS OUT THE PAST THE SAME WAY YOUR SHOP DOES (no database step, nothing to upload - pushing this one is the whole of it).

What you asked. "Is all calendar passed day grey out same thru all calendar? I like the one at store."

It was not. On the Orders calendar a day that had gone by came out in **two different greys side by side in the same row** - one shade for the days you deliver, and a slightly different one for the days you do not. On the other calendars there was only one grey, but it was the app's own faded grey, not the shop's. The shop's is the one you picked.

What changed on screen. A day already gone is now drawn in the shop's own quiet grey, **flat and one shade, whether or not you deliver that day.** That covers the Orders calendar at the top of the Orders screen, the calendar inside **Delivery Dates**, the small calendar you open when you pick a delivery day, and the days you mark on a product's availability - they all now use the one colour, so a past day looks the same wherever you meet it. Nothing ahead of today is touched - a day still to come looks exactly as it did.

What deliberately did NOT go grey. Anything on a past day that is telling you something keeps its own colour, because greying it out would hide a fact you are looking at. So on a day that has gone by, the **green pill** on one of your delivery days stays green, and so does the **green tint** on a day a product sells. Greying those would have made a past delivery day look like a day you do not deliver, which is the opposite of useful.

The day you are looking at still turns brown. That has not changed anywhere - so on the Orders calendar you can still tell at a glance which day is open in front of you, and it stays plainly different from the days that have simply gone by.

Why it is safe to be this quiet. The store calendar's grey is deliberately faint but legible, so a day that has gone still reads as a day rather than a hole in the week. That reasoning is the shop's, and it now holds here too.

Nothing else moved. No number, price, booking count or order changed. The database was not touched, and there is nothing to upload - this is the app's own pictures of its calendars.

01 Oct 2026 - engine v260, YOUR MAILS NOW COME FROM YOUR BAKERY, NOT FROM THE APP (two functions must be uploaded - pushing alone does not change this one).

What a customer saw. The two mails this system sends - the wish-list mail you send out, and a customer's suggestion coming back to you from the shop page - each arrived with a person's name at the top of the inbox list. That name read "**BakeAdmin wishes**". **BakeAdmin is the app's own internal name.** It is a word that appears nowhere on your homepage, nowhere in your shop, and on no label you print. It was the first thing a customer read from you, before they opened anything.

What they see now. Jien Luv 2 Bake - your own name, taken from your homepage title, which is where your brand is written down. The address the mail is actually sent from has not moved: it is the same verified address as before, and only the name in front of it changed.

One thing to check, because it can silently win. If the project has a saved setting called RESEND_FROM, that setting is used **instead of** the name in the code, and this change would make no difference to what a customer sees. Run supabase secrets list in **Terminal** to see whether it is set. If it is there, tell me, and I will give you the one command to set it to your name. If it is not listed, nothing further is needed.

Why pushing this one is not enough. These two mails are sent by the two functions that live in Supabase, not by the app on your phone. Uploading the code - a push - leaves Supabase running the old copy. **Both functions must be uploaded separately**, from the repository folder, in **Terminal**: wish-mail and shop-feedback. The changelog cannot do that part for you; the commands are with the version notes.

No database step. Nothing was added or changed in the database, and no existing setting moved.

01 Oct 2026 - engine v259, THE EDIT CARD STOPS CLOSING ITSELF WHEN YOU LOOK AN ADDRESS UP AGAIN (no database step - pushing this one is the whole of it).

Your words were "when i edit order, go in and press look this address up again, it exit the page and go into the order list and without allowing me to get a price from lalamove". That is exactly what it did - and it was not the fault the four reports before it sent me chasing.

The whole of it, in one sentence. The Edit card holds two things that can save: the **Save changes** button at the bottom, and the pin work, which saves by itself the moment it finds a door - because a real vehicle on a real road must never be thrown away by closing a card. The line doing the pin's saving was written to reach for the app's own save - and both of them were called save, so it reached for the **Save changes** button instead. Pressing **Look this address up again** therefore found the address, wrote it, and then pressed Save for you: your order was saved, the card closed, and you were dropped back on the Orders list.

You can still see it in what the app said at the time - "Order updated". That is the Save button's own message. The pin's press never had a message of its own to give you, which is how one word gave the whole thing away.

What happens now. The press finds the address, writes it onto the order, says what it found on the card's own line, and **leaves the card open and exactly where it was, with the pin moved.** **Save changes** is still yours to press when you are ready - and **Get a delivery price** is still on the card to press, which is the thing you could not get to before.

Nothing else about the Edit card changes. The delivery day, the customer, the items, the prices, the charge and the parcel are all exactly as they were.

There is no database step and nothing to upload - pushing this one is the whole of it.

01 Oct 2026 - engine v258, THE APP NOW TELLS YOU WHEN YOUR PHONE IS RUNNING AN OLD BUILD (no database step - pushing this one is the whole of it).

Your words were "still the same problem engine no 257", which was the fifth time you reported the same press - and this time it sent me looking somewhere else. The v257 fix is on your website; I checked that from outside, and the file your phone downloads does contain it. So the fix was not missing. Something was stopping it reaching your phone.

What was happening. Your website hands out every file of the app with a ten-minute cache, and there is no build number written into any of the addresses. So a phone that had the app open - or that opened it within those ten minutes - was given the copy it already had: the old one. Nothing on the screen said so. And the **Engine v257** pill on the More screen could not say so either, because that pill reads a small file of its own rather than the one carrying the fix. It would have read v257 perfectly happily on a phone still running v254.

Two changes, so that this cannot be quiet again.

The app's own files are now confirmed with the website every time it opens. The part of the app that runs in the background used to ask the network for each file - but "the network" was allowed to answer out of those ten minutes of cache. It now asks the website to confirm that each file is still the current one before using its own copy. An unchanged file costs one very small check rather than a re-download, so opening the app does not get slower; a changed file is fetched whole. There is nothing here for you to press.

And a strip now says so when your phone is behind. Where the site is serving a newer build than the one on your phone, an amber strip appears just above the bottom menu - **New version ready** - this phone has v257, the site has v258 - with an **Update now** button beside it. Press it when you are ready and the app reloads on the new build. Its words say "save what you are writing, then update" for a reason: a half-written order does not survive a reload.

If pressing Update now does not take - which can happen inside those first ten minutes - the strip does not go on offering a button that did nothing. It changes its words, says the update still has not taken, and tells you to close the app completely and open it again.

What this means for you, from now on. When you report something, look at the bottom of the screen first. If that strip is not there, your phone is on the newest build and the fault is in the code. If it is there, press Update now and try again before you tell me. Either way you are no longer guessing, and neither am I - which is worth more than the fix itself.

Also put right while here: the change history had v256 printed above v257, so the two newest entries read in the wrong order. They are back the right way round.

30 Sep 2026 - engine v257, TAPPING "RESET THE PIN" NO LONGER FOLDS THE ORDER YOU ARE WRITING (no database step - pushing this one is the whole of it).

This is the one you have been reporting. You press **Look this address up again** on the order's own door block, the app asks you to confirm because there is already a pin to replace, you tap **Reset the pin** - and you land back on the Orders list with the half-written order gone. The look-up itself then carries on and writes the order.

The cause was the rule that folds the + New order card when you press anywhere else. That rule is right for a press on the page. It was wrong here, because the confirmation the card itself opened is drawn on a layer that sits **beside** the page rather than inside it - so the app counted your tap on **Reset the pin** as a tap on the page behind, folded the card you were writing, and then did the look-up anyway. A press on the confirmation, on a pop-up, or on the lock screen no longer folds the card. A press on the page still does, exactly as before.

Why the last two versions missed it, in plain words. v255 and v256 both changed the pin window's **Look it up**; your words named the button on the order's own door block, which is a different button on a different window. And the check that cleared that door block measured where each row sat on the screen - it never asked whether the card was still open, and it never pressed the case where the confirmation appears, which only happens when there is already a pin to replace.

30 Sep 2026 - engine v256, THE MATCHES IN THE PIN WINDOW NO LONGER PUSH THE CARD AT ALL (no database step - pushing this one is the whole of it).

Your words were "it is still the same, press look up this address again make it exit the page we are working in, why?", and the honest answer is that v255 made it worse. v255 tried to hold the map still, and while the list of matches takes up room on the card, holding the map still can only be done by scrolling the whole window by the height of that list - 269 pixels on a phone's width. That scroll carried the address box and the **Look it up** button you were pressing clean off the top of the card: the button you were holding went from 248 pixels down the screen to 21 pixels **above** the card's own top edge. Instead of the map jumping, the window jumped.

The list of matches now floats over the map rather than taking up room on the card. It drops from under the button you pressed, sits on top of the map, and is capped at about four rows with its own scroll. Because it takes up no space, nothing above it can be pushed anywhere - not the button, not the address box, not the window's own scroll.

And the list puts itself away the moment you choose from it. That is the one thing that had to change for it to float: a list left sitting on top of the map would hide the very thing you want next, which is the map, to check the pin landed on the right door. Pick a match and the list closes, the pin moves to it, and the map is there.

Measured on a phone's width with four matches: the button moved 0 pixels and the window scrolled 0 pixels. The map slides down by the height of the one line that appears under the button (38 pixels) and nothing else moves - including, now, the button itself. It holds the same at two matches as at four, when the look-up finds only one place, and when the look-up fails.

Separately measured, because your words name that button too: "Look this address up again" never moved anything. That is the same press on the order's own door block, and it was measured on both kinds of window - the + New order card and the Edit order pop-up - and on both kinds of answer, a look-up that finds your house and one that fails. In every case the row stayed on the same line and the card stayed open. The window that misbehaved was the pin window's **Look it up**, and that is the one this version fixes.

What did not change: the first match still lands on the map by itself, the list is still only how you say "not that one", the line under the button still says when the answer is the road rather than your house, and no wording, price or order of questions is touched.

30 Sep 2026 - engine v255, LOOKING AN ADDRESS UP NO LONGER THROWS THE PIN CARD AROUND (no database step - pushing this one is the whole of it).

Your question was "when i say look this address up, why the interface jump out of the page?", and the answer was in the shape of that card. Pressing **Look it up** in the pin window asks the address service for other matches, and the list of them is drawn *under the button*, above the map. So the list appearing shoved everything beneath it down the window: on a phone's width the map was thrown 269 pixels down the card, and the **Use this spot** button went with it. The button you had just pressed did not move at all, which is exactly why it read as the page jumping rather than the button.

The map now stays where it is while the matches appear. The window is held by the map - the thing you are checking the pin against - so when the list arrives the window scrolls to keep the map on the same line of your screen, and the matches are revealed above it, where you are already looking. Measured on a phone's width, the map and the **Use this spot** button both stayed still to the pixel, and three presses in a row left them still. The same holds when the look-up finds only one place, and when you tap one of the matches.

This is the same rule that stopped the + New order card jumping at v254, and it now lives in one place. Both screens use one copy of it, so any later screen that needs the same protection inherits something already tested rather than a second hand-written version. Nothing else on either screen moved: the wording, the prices, the order things are asked in and the address look-up itself are all unchanged.

30 Sep 2026 - engine v254, THE + NEW ORDER CARD HOLDS STILL UNDER YOUR THUMB, AND ITS BIG BUTTON NOW SAYS PLACE ORDER (no database step - pushing this one is the whole of it).

The button inside the card says + Place Order. It used to say + Add order, and that is the only thing that changed: the card's own title is still + **New order**, the + New order button on the Home screen is untouched, and the Edit order pop-up's **Save changes** is untouched. The two sentences that send you to the button now say

"then press Place Order", so they no longer point at a button that is not there.

The card no longer jumps out from under your finger when you reset the pin. Your words were "once i click reset pin the screen jump". Resetting the pin makes the card change its own height in up to eleven ways at once - the answer line appears, the words re-wrap, the map is shown or hidden, the price list is rebuilt - and nothing was putting the page back where you were reading. The card is now held by the row of buttons under your thumb: whatever the block does to its own height, that row stays on the same line of your screen. Measured at a phone's width, the row sat still to the pixel while the page scrolled 15 pixels to absorb the change. The same press used to push everything below it by 210 pixels when the press was the one that made a map appear.

A map no longer keeps running behind a card you have moved on from. Every time the card is rebuilt - you change Fulfilment, you tap a day in the card's calendar, a pop-up refreshes, a sync pull arrives - the old map was left alive, holding its tiles and its screen-size listener until some later resize happened to notice. On a phone that notice is the keyboard opening, which is exactly what happens while you are typing in this form, so the maps piled up. Each new map now clears away the ones whose box has left the page, at the moment it is built. Nothing on screen changes.

A map can no longer paint over a question you are being asked. You put this as "sometimes pop up like half at back layer", and it was real. With a live map in the card, the map was drawn over the bottom half of the confirm dialog: the sentence cut off mid-sentence, both buttons hidden, and a tap in that band landing on the map's + / - zoom instead of the button. The map's own layers climb to 1000, and the box it sat in was not containing them, so they were competing with every layer this app draws. The map is now sealed inside its own box, in all four places the app draws one.

Nothing else moved. The door block's wording, the address look-up, the prices and the order the card asks things in are all unchanged.

30 Sep 2026 - engine v253, THE SHOP'S FEEDBACK MAIL NOW READS IN THE SAME ORDER AS YOUR WISH-LIST MAIL (no database step, but the shop-feedback function must be redeployed).

The heading block moved to the top, and the customer's words now come last. Engine v252 put the words first and the heading underneath a short rule. You asked for the wish-list mail to be the reference, and the wish-list mail does the opposite: it opens with what the mail is, then who it is about and when, and puts the content underneath. So the feedback mail now does exactly that, with no rule in between and no closing line:

New feedback for the shop page (Engine v253).

Project: jienluv2bake.com.my/store/

Sent: 2026-09-30 17:31

Written in English.

the words the customer typed

The language is written out as a word. Engine v252 printed the short code it was handed, so the line read Written in en. The shop asks its questions in three languages, so the line is worth having, but "en" is not a word anybody reads. It now says **Written in English.**, **Written in Malay.** or **Written in Chinese.**, and it is left out entirely when the page did not say which language the customer was reading.

The subject line and the Project line are unchanged from v252: the subject is still Shop feedback · Engine v<n> · <date>, the project is still the live address the customer was reading, and the time is still on your own clock in Penang. Only the order of the body and the wording of the language line have moved.

One step on your machine. The part of this that builds the email lives on Supabase, not on GitHub, so pushing is not enough on its own. Engine v252's version of it was never actually running there - the mail you received after v252 was still the build from before, which is why it carried none of the new heading. After pushing this version, redeploy it once:

supabase functions deploy shop-feedback --project-ref hzpyblqygnntixkijeem

30 Sep 2026 - engine v252, THE SHOP'S FEEDBACK MAIL SAYS WHICH SHOP AND WHICH BUILD (no

database step, but the shop-feedback function must be redeployed).

A customer's message now arrives with the same kind of heading your wish-list mail has. Before this version the email was the customer's words and nothing else, so a message saying "the cart is confusing" could have come from either of your shops and from any build of the page. Now every one of these emails carries, under the words and a short rule:

- **Project** - the address the customer was actually reading, for example jienluv2bake.com.my/store. It is read off the live page rather than typed in anywhere, so your two shops can never be confused for one another, and the second half names which page of that shop the words were written on. **Sent** - the date and time on **your** clock, Penang time, so "16:52" means the same thing on the email, on your phone and in the backoffice. It follows the same shape as the wish-mail mail: 2026-09-30 16:52. **Engine** - the build the customer was looking at, both in the subject line and in the line above the project, so you can tell whether a complaint about the shop came from a phone that is still running yesterday's copy. **Written in** - the language, when the customer was reading the shop in Malay or Chinese.

The subject line names the build too, so the mailbox listing alone tells you what you need: Shop feedback · Engine v252 · 2026-09-30. Everything else about the mail is unchanged - the customer's words still come first, the heading still sits below them, and the mail still goes to the developer addresses you set in Settings.

And the shop itself now says which build it is running. At the very foot of the shop page, under "Website by" and the developer's WhatsApp and email links, there is a new small line:

Engine v252

It is the same number your app shows on More, and the two are read from one file rather than kept as two copies that could drift apart. That is the point of it: when a customer tells you something looks wrong, you can compare the number at the foot of their shop page with the number on your own More screen and know in one glance whether they are seeing the same page you are. It is the smallest type on the page and the same muted grey as the credit above it, in the developer's own corner, where it is there when you go looking for it and out of the way of everything a customer came for.

One step on your machine. The part of this that builds the email lives on Supabase, so it must be redeployed once, exactly as before:

```
supabase functions deploy shop-feedback --project-ref hzpyblqygnntixkijeem
```

Nothing else to set up - no new secret, no new database table, no new key. The email reuses the same Resend account your wish list already sends with, and the same published settings row the shop page reads for its "Website by" line.

30 Sep 2026 - engine v251, THE SHOP'S QUESTION IN YOUR OWN WORDS (no database step, no redeploy, one push).

The line at the foot of the shop now reads exactly as you wrote it. "Webmaster: Like the User Interface? Tell me & I will improve it!" - your words, your punctuation, with "Tell me" back where you put it.

It is a measurement, not a rewrite, and this version takes the measurement in both directions. A phone gives that box about 323px of room, and "the User Interface" spelled out costs about 150px of that on its own. Your sentence runs to about 377px, so on a phone it takes **two lines** - and the guard built for exactly this opens the box to show the whole question rather than cutting the end off it. Nothing is clipped on any phone: at 320px the box simply opens a little further, and the page never scrolls sideways.

The Malay is your own sentence and holds one line - as long as it keeps the shape you wrote it in. "Suka User Interface ini? Komen & Saya akan perbaiki!" measures about 322px against a phone's 323px, which is a fit to within a pixel. It carries no "Webmaster:" prefix for that reason: the prefix would add another 76px and push it onto a second line. The English and Chinese still say who is asking.

The Chinese needed no change at all - it already used the everyday word for a screen's interface rather than a shorthand, which is why it reads the same as before.

30 Sep 2026 - engine v250, THE SHOP'S QUESTION SPELLS OUT "USER INTERFACE" (no database step, no redeploy, one push).

Shorthand is out. You looked at the built page and said it plainly - "UI should be user interface, ui is not a laymen term" - so the line at the foot of the shop read **"Webmaster: Like the user interface? I'll improve it"**, with the Malay and Chinese equivalents.

And it stayed one line on a phone, which is what made this a measurement rather than a rewrite. A phone gives that box about 323px of room, and "user interface" spelled out takes about 150px of it on its own. The full sentence you first wrote runs to about 373px - wider than the box - so its tail would have been cut off. Everything you asked for was kept: who is asking, the question in full words, and the promise that it gets improved. Only "Let me know" went, because the question mark and a box waiting to be typed in already say it. The Chinese needed no change - it already used the everyday word for a screen's interface rather than a shorthand. The Malay used "UI" too and spelled the term out as "antara muka", for exactly the same reason the English does.

30 Sep 2026 - engine v249, THE BOX AT THE FOOT OF THE SHOP SAYS WHO IS ASKING (no database step, no redeploy, one push).

The one quiet line at the foot of the shop page now introduces itself. It reads **"Webmaster: Like this UI? Tell me, I'll make it better"** - the customer types over that sentence exactly as before, and the whole of it still goes out as their own words when they send, with nothing to press but Enter.

It is one line on a phone, and it is never a line and a half. The sentence was measured against the width of the box at phone size rather than guessed at, which is why it says "this UI" and "Tell me" rather than spelling everything out: the longer wording you first wrote runs wider than the box on a phone, and its tail would have been **cut off mid-sentence** - the question would have read as far as "...Let me know, and I" and stopped. What you asked for is kept in full, in the shorter form, and all three languages carry it.

And if a phone is narrow enough that the question does wrap anyway, the box now opens itself to show the whole question instead of hiding the rest. That is the one thing this version can do that a shorter sentence alone could not: on any phone, in any of the three languages, the customer always sees the entire question rather than its first line.

The Chinese and the Malay say the same thing as the English, in the same few words: the Chinese names the person asking as the site's webmaster and asks whether the customer likes this page, and the Malay uses "Webmaster" as it is, since that is the word a Malaysian website uses too.

30 Sep 2026 - engine v248, NOBODY HAS TO PRESS ANYTHING, AND NOBODY LOSES WHAT THEY WROTE (no database step, no redeploy - the same one function, one push).

Two more things you asked for on the little box at the foot of the shop.

Closing the page is now a send. A customer who writes a sentence and then closes the shop, or taps a link out of it, has had their say - the words go as they leave, without them pressing Enter at all. **Nothing goes out while they are still on the page**, so a sentence somebody is midway through writing is never mailed off behind their back. Pressing Enter is still a send, and still shows your reply straight away.

And nothing they wrote is lost on the way. If they step away from the box with words in it, or the page is closed before anything could be sent, the sentence is **kept on their own device** and is waiting in the box the next time they open the shop - it goes only once it has been sent. If a send does not go through, their words are put back, so a failure nobody was around to read is never a sentence lost. And **with no internet, nothing is sent at all**: the words simply wait on their device for the next visit instead of going down with the page.

The box itself is unchanged and stays as quiet as it was: one line, no Send button, and one small grey line that says how to send.

And the Bahasa Malaysia has been rewritten to read the way Malaysians actually write. Several lines on the shop were correct textbook Malay but nobody writes that way on a Malaysian website - "pembangun" for your developer, "telah" where anyone here would say "sudah", "kemas kini baharu", "Menghantar..." where a

shop would say "Sedang dihantar...", and a reply that read like a letter from an office rather than from you. The words are the same promise, in the voice a Malaysian customer expects: **"Idea anda sudah kami terima! Update terbaru akan datang tak lama lagi."**

The Chinese is now Malaysian Chinese, which is what your customers read. Malaysian Chinese read simplified characters exactly as mainland China does, so the characters stay as they are - what changes is the words. A few lines had been written the way a shopping app from China writes them: the items count used the mainland measure word rather than the one a Malaysian uses, the basket was called a shopping bag, two lines said "coming soon" and "enjoy" in the flat official way, and the two arrows on the delivery calendar were in **traditional** characters, which is simply the wrong script for your customers. The reply now says, in Chinese, that their idea has been received and will be taken good care of, with new updates coming very soon - not that it "will be processed". Taiwanese Chinese was not used: it would mean changing every character to traditional, and your customers do not read that.

30 Sep 2026 - engine v247, THE SHOP FRONT NOW ASKS THE CUSTOMER WHAT THEY WOULD CHANGE (no database step, one new function to deploy once, one push).

You asked for a way for a customer to tell your developer what they would improve about the shop page, in their own words, with a reply straight away.

A single line now sits at the foot of the shop, under the "Website by" line. It is printed with your own question - **"Like this UI? Comment, and I will make it better."** The customer types over that sentence and presses **Enter**. There is no Send button, because you asked for none: Enter is the send. The line is one line tall and opens further only when the words no longer fit, so at rest the whole thing is a single quiet row at the foot of the page. Once there are words in the box a small grey line under it says **"Press Enter to send"**, so the one key that does something is never something they have to discover by accident.

The moment the send lands, the box is replaced by your reply - **"Your idea is well taken care of. New updates soon!"** - so they see the answer without doing anything else. They are never asked for an email address, and nothing is created on your side for them.

The words come to you as an ordinary email, at the developer address you already set in **Settings -> Website & developer** - the same address the shop's "Website by" line shows. It runs on the **same email service your wish list already uses**, so there is no new account, no new key and no new DNS record. Each message says which page it was written on, in which of the three languages, and when it arrived.

There is no box until you have set that address, because the box has nowhere to send to without it. That is on purpose: a box that quietly collects sentences nobody will ever read is worse than no box at all.

A send that did not go through says so, plainly. If the email could not be sent the customer is told, in words, to try again or to use the WhatsApp link that is already drawn just above the box - and **their words stay in the box**, exactly as typed, so nothing they wrote is lost. A thank-you drawn over a message that never left would stop them trying again and nobody would ever find out, so that is the one thing this never does.

It speaks all three languages. The question, the line that says how to send and the reply are written in English, Mandarin and Bahasa Malaysia, and a customer who switches language keeps whatever they had half-typed and whatever reply they had already been given.

What you have to do, once. This needs one small new function deployed in Supabase the same way the wish-list function was - the file is supabase/functions/shop-feedback. Until it is deployed the box still appears and still says plainly that the send did not go through. Nothing else is required: no database step, no new key, no new setting - it uses the developer email you have already saved.

30 Sep 2026 - engine v246, THE NOTES ON THE ITEMS ARE NOW UNDERLINED ON THE COMPACT LABEL TOO (no database step, no redeploy, one push).

You spotted this one: on the Compact label a note like **no nuts** was printing as ordinary words, while the same note stood out on the Full and Mailing labels. It was the last place a customer's own words went unmarked, and it was easy to miss - Compact puts every item on ONE line, so the note sat in the middle of that line with nothing

to pick it out.

Each noted item's words are underlined on that joined line now, on the item they belong to, so a note on the second of four items is underlined on the second item and not somewhere else on the line. **If two items carry the same note, each one gets its own underline** - the line is not allowed to mark one item twice and leave the other plain.

The line itself has not changed at all. It reads exactly as it did - the items in order, the notes in brackets beside the item they belong to - and a label for an order with no notes on it is byte-for-byte what it was. Nothing else about any label moved: the Full, Mailing and Name-only styles print exactly as they did, and the delivery note on the Compact label is untouched by this.

30 Sep 2026 - engine v245, THE DELIVERY NOTE IS UNDERLINED TOO, AND IT NOW PRINTS ON THE COMPACT LABEL (no database step, no redeploy, one push).

You made the point that the **delivery note** is not a courier thing - it applies to a self-collect order just as much, so it should stand out and be printed wherever it belongs.

It is underlined now wherever it is drawn. The order row carries it as Auntie Bee · 012-345 6789 · deliver after 3pm, with the note itself drawn under, exactly as a note on a single item already was. The same on the printed label: the Note: line keeps its label plain and underlines the customer's words.

The Compact label prints it too. Compact used to join every item onto one line and leave the delivery note off entirely - so a note about the doorstep, the gate code or a collect time disappeared the moment you picked the denser label, on a self-collect order as much as a courier one. It now prints in the same place the Full label puts it: after the items, before the courier address.

Everything else is unchanged. The Mailing and Full labels read exactly as they did, the note box is still offered on both self-collect and courier orders, and an order with no note on it looks and prints exactly as before. Name-only stays a bag tag - code and name, no fields to read.

30 Sep 2026 - engine v244, A NOTE A CUSTOMER LEAVES ON ONE ITEM IS NOW UNDERLINED, SO IT IS HARD TO MISS (no database step, no redeploy, one push).

You asked where the words a customer leaves actually reach you. The answer turned out to be uneven: the **delivery note** on the whole order pings your phone the moment the order lands, but the **note on a single item** - "no nuts", "write Happy Birthday" - only ever lived inside the app, sitting quietly in grey beside its item where it is easy to skim past. You decided that is enough for it to be **underlined**, so it reads as something you were meant to notice.

On the Orders list, an item that carries the customer's own words now has those words underlined - Focaccia 800g ×2 (no nuts) with the brackets drawn under. It is the same line in the same place, just marked. Nothing moves, nothing is added, and an order with no notes on it looks exactly as it did.

On the printed label, the note beside its item is underlined the same way, on both the Full sheet and the Mailing sheet - so what you hold in your hand while you kit an order is as easy to read as the screen.

The line itself is unchanged. Only the way it is drawn changed, so the row, the label and the search still read the very same words - the underline cannot quietly disagree with what the line says.

30 Sep 2026 - engine v243, THE DELIVERY-DAY CALENDAR IN THE ORDER SCREENS NOW FOLLOWS TODAY, THE SAME AS YOUR SHOP (no database step, no redeploy, one push).

Your words: you wanted the calendar in the + New order card to behave like the one on your shop.

It does now, and it is the **same calendar in all three places it appears** - the + New order card, the strip at the top of Orders, and the day picker in the Edit-order pop-up. One behaviour in three places, so a day sits in the same square wherever you meet it.

What was wrong with it. It was a **month** - the whole of September, then the whole of October - exactly the calendar v231 replaced on your shop, for the same reason. On the 30th, most of what it drew was days already gone; a delivery day on the 29th could not be seen at all; and both arrows were drawn even at the ends, so a

greyed button sat there inviting a press that would do nothing.

From this version it is five whole weeks anchored on today. The first row is the week just gone, the second row is this week with today ringed so you can see where "now" is, and the three rows under it are the weeks ahead, which is where your delivery days actually live. Today is always in the second row. **Every square is a real date now** - no blanks, no padding, so a delivery day at the start of next month is visible and openable from the moment you publish it.

Your delivery days are unchanged on it. Still circled, still carrying how booked each one is under the number, still FULL in red where a day is at capacity, and a day already gone still dimmed but **still open** - you backfill and review old days. A day you do not deliver still answers a tap, and still names where it gets added, so a holiday falling on a day you do not deliver is never the one square with nothing to say.

The arrows move one whole week, and are gone where there is nowhere to go. A dead control reads as a bug, so at the ends of the days you have set there is simply **no arrow** - nothing greyed out, nothing to press that would do nothing. The window still stops where your shop's stops for the same list, so the two calendars agree about where the far end is. A new week arrives **sliding up** going forward and **coming down** going back - the direction you are travelling in - and only when you pressed an arrow: an ordinary redraw never replays it.

The marks you made on Delivery Dates are on it too - the same single-day boxes and the same see-through bands, in their own colours and their own depths.

It opens on the day you are working on, not on today. The Edit-order pop-up opens on the day that order is actually on, so moving an order a season out does not leave you looking at a week with nothing to do with it. And picking a day that is already on screen **does not move the window**, so the grid never slides out from under your finger.

One thing I deliberately left alone. v231 ended with "nothing in the backoffice calendars changed; this is the shop's calendar only" - that was your instruction then, and this version is you asking for the opposite, for this one calendar. The **look** of the past days in the backoffice is untouched, and so are the other admin calendars: the free order-date field, Delivery Dates, Profit and Deliveries. What changed here is this one calendar's behaviour.

30 Sep 2026 - engine v242, A CUSTOMER WHO ALREADY HAS A COURIER BOOKING IS NEVER QUIETLY PUT ON A SECOND VAN (no database step, no redeploy, one push).

Your report: if a customer's order already has a courier booked, the delivery run should not tick that order, and it needs to say so, so that a delivery is not booked twice.

Here is what was happening. A booked trip is recorded on the order's own row, and the delivery run did not look at it when it built the day. Every customer on the day opened already ticked, the Tick them all press ticked them all, and a booked customer looked exactly like any other. The only thing that refused was the app's own sentence at the Book press - so you could tick a customer a driver was already on the way to, price the run, and press Book, and the refusal arrived afterwards as a message rather than as anything you could see while you were choosing.

From this version, a customer whose trip is already running opens OFF the run, with a line under their name saying so. The line names the courier and its own word for where the trip has got to, and it says plainly that ticking them and booking the run would send a second vehicle to the same door. The head of the list counts them too - it will read something like "Who is on the run - 3 of 5, 1 already booked", so the number you tick and the number of rows in front of you never disagree.

It is not a gate. The tick is still live and still yours. If you tick a booked customer by hand the app honours it, and the refusal comes at the price instead, in the same words as before. That is deliberate: the tick is a decision you make, and the app's job is to make sure you can see what you are deciding.

And the row offers the way out, rather than going grey. Under the warning there is one press: Call off the trip and add to this run. It asks you first, in the open, because calling a trip off is the half that cannot be undone - the driver stops being sent, and the customer's tracking box keeps its link but nothing will update it any more. If you

say yes, the original booking is cancelled with the courier, the order's own row records that it was you who called it off and when, the warning disappears, and that customer is put on this run so their delivery can be consolidated with the rest of the day. If the courier refuses to call the trip off, its own reason is shown and nothing is written.

A trip that has already finished - delivered, or cancelled earlier - is not a barrier at all. Those customers open ticked like anyone else, with no warning, because a delivery that has to be done again still has to go on a run.

Orders going out by parcel carrier are untouched by this version. They were never on the run list in the first place, so there is no second van for them and nothing here to warn you about.

No database step, no redeploy, no new key, and no new setting to switch on.

30 Sep 2026 - engine v241, THE RESET NOW ANSWERS ON EVERY CARD, NOT ONLY A FRESH ONE (no database step, no redeploy, one push).

You pushed v240 and you were right again: the press worked, but the card could still stay silent. This version closes the last way it could.

Here is what was happening. The answer your press writes - the pin moved, or it found the same spot, or the address could not be reached - has two possible places to be said. One is the line under the buttons on the door card itself. The other is the line inside the delivery-price section. That section folds away, and its line folds away with it. v240 made the press work before the price section existed at all, which was the right fix - but the moment you opened **Get a delivery price** even once, the card handed every later answer to the folded-away section and stopped using the door card's own line.

So the sequence that caught you was this: open an order, press **Get a delivery price**, close it again, then reset the pin. The look-up really ran and the pin really moved, but the sentence saying so was written inside the section you had folded away, and the door card said nothing at all. And where the look-up answers with the same spot as before, nothing moved on the map either - so the whole press looked exactly like a press that had never happened.

From this version the answer is always said on the line you can actually see. While the price section is open it is said there, exactly as it always was. While it is folded away - or if you have never opened it at all - it is said on the door card, right under the buttons. The two lines are never both used at once, so they cannot come to disagree about what your press did.

The same fault was on the pin you drag. Moving the pin by hand ends through the same sentence about the price section's numbers, and that sentence had the same problem: with the price section folded away, a drag you had just made reported into a section you could not see. It now says what it did on the door card in exactly the same way.

Nothing else changed. No database step, no redeploy, no new key, and no new setting to switch on.

30 Sep 2026 - engine v240, THE RESET BUTTON NOW ACTUALLY PRESSES (no database step, no redeploy, one push).

You pushed v239, saw the reset button at last, and pressed it - and nothing happened. That was exactly what your report said, and it was exactly right. The button was on the card, but its press was not connected to anything yet.

Here is what was happening. The reset press was wired up at the same moment as the delivery-price half of the card, and that half is only built the first time you press **Get a delivery price**. So on a card you had just opened, the button was drawn and waiting but its press led nowhere: no look-up, no message, no movement. The instant you happened to open the price fold once, the button started working - which is why it looked like an ordinary working button that simply ignored you.

From this version the reset press is connected the moment the card opens. It never waits on the delivery price again. You can open an order, see the pin, and reset it without touching the price button at all - which is the order you would naturally do it in anyway, because you reset a stale pin when you call the customer, not when you ask for a price.

And the press now always tells you what it did, on the card itself, right under the buttons. It says it is looking the address up, then it says one of the three things that can happen: the pin moved, or the look-up found the same spot and there was nothing to move, or the address could not be reached and the pin has been left exactly as it was. Those words used to appear only up in the price section, which is another reason a press on a fresh card looked like it had done nothing at all.

If a reset is about to replace a pin a person chose - the customer's own pin, or one you placed by hand - it still asks you first, exactly as v239 described.

Nothing else changed. No database step, no redeploy, no new key, and no new setting to switch on.

30 Sep 2026 - engine v239, THE RESET NOW SHOWS UP WHERE YOU ACTUALLY WORK (no database step, no redeploy, one push).

This is the fix for v238. You pushed it, opened an order, and the reset was still not there. You were right, and the reason is this: v238 offered the reset when the pin was the customer's own, and hid it when the pin was one **you** had placed by hand.

But placing the pin by hand is the only thing that card ever offered you. So a drag is what you did, and a drag is exactly what then hid the reset. Every customer whose pin you had ever corrected by hand still showed **Move this pin** and nothing else. v238 changed nothing at all for the doors you had touched yourself, which is why it looked like it had not arrived.

From this version the reset is offered wherever there is a pin to replace. If a pin exists on the order and an address is typed, the press is on the card - the customer's own pin, a pin you dragged by hand, a pin an older look-up wrote, or one a reset already wrote. The only case with no press is an order that has no pin yet, where there is simply nothing to replace; the price button looks one up on its own.

Where it would replace something a person chose, it still asks first. That is the customer's own pin, and it is also now your own hand-placed pin. The wording names which of the two it is about to replace. Over your own pin it says this is the door you placed on the map by hand, that a look-up may only find the road and can be a step back from a door you already had right, and that you can drag the pin again afterwards. Over the customer's pin it says, as v238 did, that this is their pin and you can switch back to it. Either way the confirming button says **Reset the pin**, and Cancel changes nothing at all.

Where it only replaces the app's own guess, it does not ask - a look-up's answer, or a reset of one. That is what this press has always done, so nothing new is put in your way there. The button says **Look this address up again** on those, and **Reset the pin from the address** on the two that will ask.

Nothing else changed. No database step, no redeploy, no new key, and no new setting to switch on.

30 Sep 2026 - engine v238, RESET THE PIN, AND LET THE MAP ZOOM (no database step, no redeploy, one push).

This one is yours, word for word: when you call a customer, their address may already have changed, and the only thing on offer was to move the pin. The map would not zoom out, and dragging a pin onto the right rooftop by thumb was slow and easy to get wrong. You asked why there was no way to reset the pin. There wasn't one. There is now, and the map zooms.

There is now a reset, on the pin card, where you can see it. When the pin the driver is sent to is the customer's own - the one they dropped on your shop page - the button under the map reads **Reset the pin from the address**. One press asks your address service for the address on the order again and puts the pin where that answer lands. That is the case you described: the pin was right when they dropped it, and they have moved since.

It asks before it replaces their pin. The customer's own pin is a fact they gave you, so a button that quietly overwrote it would be the pin moving on its own again - the fault several versions of this card have been spent ending. So the press asks first, in plain words: this is their pin, resetting replaces it with a fresh look-up of the address, and a look-up may only find the road. The confirming button says **Reset the pin**. Cancel changes nothing at all, not even a look-up.

And you can put their pin back. Once a reset has replaced their pin, the same card offers **Use the customer's pin instead**, one press, back to their exact point. The reset is not a one-way door.

A reset yields to a pin they drop afterwards. If that customer pins a new spot on your shop page after you reset, their new pin wins again straight away - the same rule as before. A reset replaces one specific stale pin, not every pin that customer will ever drop.

The map zooms the whole time now. The plus and minus buttons are drawn on the pin card, and pinch, double-tap and a desktop box zoom all work - before you press anything. Only moving the pin is still behind **Move this pin**: the map will not pan under a pin you are reading, and the pin will not take a drag, until you say so. That split is the point of this version. Looking and correcting used to be one action; they are two again.

Nothing else changed. Your prices, your delivery calendar, your messages, the order card, the packing slip and the label sheet are all exactly as they were. Adding an order is still the one pass v237 made it, and a pin you placed by your own hand is still not offered up for replacement.

On two phones. If your second phone has not been reloaded yet and is still running v237, it reads a reset as the older kind of pin and will price to the pin you replaced until it is reloaded. Reloading it fixes it, and there is nothing to run in Supabase for any of this.

30 Sep 2026 - engine v237, A COURIER ORDER IN ONE PASS, AND THE DAY ON ONE LINE (no database step, no redeploy, one push).

This one is yours, and it came from a complaint about your own way of working. Adding an order used to be two steps: fill in the New order card, press Add order, then find the order again and open the Edit window to finish it. You said that was redundant, that the flow was not smooth, and that it started in the wrong place.

A courier order is now one pass. The reason it was ever two steps is that four things lived only in the Edit window and not on the card: the courier charge, the tracking number, the parcel carrier, and the Get a delivery price block. They are all on the card now. Choose Courier delivery and they unfold underneath it, in the order you need them: the address with the map pin, the charge, the tracking number, the parcel, and the price. Take the order, price the trip and record what it cost, all before you press Add order once.

The price is on the card. The booking is not. This is the way you asked for it. You can ask for a price and take the fee onto the order from the card, but Book this trip stays on the order itself, where it has always been. Booking a real vehicle against an order you have not finished taking would be the wrong thing to put under your thumb one press away.

The day is one line now. The card used to open on a whole month of calendar, even though you had already picked the day on the screen behind it. It now reads as one line - Delivering Fri, 2 Oct - with the calendar unfolding underneath only when you tap it. Nothing about how a day is chosen has changed: tapping one still switches the screen, so the products and the day's limits are right.

The items come first. After the day comes what the customer wants, then the customer. It is the order things are actually said to you on the phone: they tell you what they want, then who they are. Everything else on the card is unchanged and in the place it has always been.

Your own note box is now the delivery note. Your app and your shop both used to call the order's note box simply Note. Now that every item carries its own note, that box was left doing a different job, so it says what it does: it is the delivery note, for the gate code, the landmark and the time you should arrive. The box on your shop and the box in your app now show a grey hint saying exactly that, so a customer knows what belongs in it without you having to tell them.

Nothing is written until you press Add order. The card was, and still is, a draft. The charge, the tracking number and the parcel are written onto the order at the same moment the order itself is created, so an order you decide against leaves nothing behind it. And if you type a charge but do not say who paid it, the card refuses in words and changes nothing, exactly as the Edit window does.

Nothing else changed. The order's own field order in the Edit window, your prices, the delivery calendar, your WhatsApp messages, the packing slip and the label sheet are all exactly as they were.

30 Sep 2026 - engine v236, A NOTE ON EACH ITEM, SWITCHED ON PER PRODUCT (no database step, no redeploy, one push).

This one is yours, word for word: you wanted to switch a note box on or off on a product card, so a customer can leave a note on each thing they order - and switch it off on the products where a note makes no sense. And you asked for the same box in your own app, so a note you are told over the phone can be written down in the same place.

The switch lives on the product card. In your app, open a product and you will find a tick that reads **Ask the customer for a note on this item**, sitting beside the other selling switches. Tick it and the customer is offered a note on that item; leave it alone and they are not. It is off until you tick it, so nothing on your shop changes until you decide it should.

On the shop, the note is quiet. A customer who puts an item in the basket sees a small **Add a note** link under it. Nothing is opened on their screen until they tap it, so your menu still looks exactly as it looks today. Tap it and a box opens on that item, with a short example in grey: no nuts, or a name to write on the loaf. Whatever they type stays with that one item and rides on that item's line when the order comes in.

The note belongs to the item, not to the order. This is the whole point of it. If someone orders a focaccia and a sandwich and writes "no nuts" on the focaccia, that note comes through on the focaccia line only. The sandwich line is untouched, and the order's own note box at the bottom still works exactly as it did for anything that is about the whole order.

A note in your own app is offered on every line. Your New order card, and the Edit order window, both give every line its own small note box - whichever product it is, and whether or not you have switched the customer's box on. The switch decides what you ASK a customer for; it must never decide what you are allowed to write down. A phone order of "no nuts on the focaccia" has to have somewhere to go.

You see the note where you work. On the orders list it appears in brackets beside the item it belongs to, so it sits next to the thing it is about rather than in a note line at the foot of the order. The same is true of the packing slip and the label sheet: it prints beside that item and never as the order's own note.

Nothing about a note blocks an order. An empty box and no box at all place exactly the same order - there is no such thing as a half-filled note to chase. Switching a product's note box off after the fact never hides or loses a note you have already collected, and clearing a note you wrote puts that line back exactly as it was.

Your WhatsApp messages do not change. The note stays in your app, which is what you asked for. Every message you send reads exactly as it reads today, and your customers' tracking page is untouched by this version.

Nothing to set up. There is no database step and nothing to redeploy. The switch travels to your shop with the rest of the product details, and a note travels with the order from one of your phones to the other on its own, the same way the rest of the order already does.

Nothing else changed. The order's own note box, the prices, the delivery calendar, the basket and the place-order bar are all exactly as they were.

29 Sep 2026 - engine v235, A BOOKED TRIP NOW SHOWS WHAT IT COST AGAINST WHAT YOU CHARGED (no database step, no redeploy, one push).

This one is yours, and it started as a worry rather than a request. A courier price you pick when you take an order can turn out to be higher or lower by the time the trip is actually booked, and nothing was watching that gap. You asked to see it rather than be protected from it: a real cost makes you aware, leaves room for a promotion later, and lets you decide not to collect the delivery at all.

So the booked-trip card now does the sum. Under the line that tells you when calling the trip off stops being free, the card compares what the trip actually cost against the courier charge on the order, and says which way the difference fell. If the trip cost more than you charged, it says so and tells you that much came out of your own pocket. If it cost less, it says that too, and that the difference stayed with you. Neither direction is treated as the bad one.

It shows both directions on purpose. A gap that only spoke up when you were short would be an alarm; you asked for a reading, so a surplus and a shortfall are worded the same way, in the same plain sentences.

A trip you decided not to charge for says so. If you have recorded no courier charge at all, the card tells you the whole cost of the trip is your own, rather than staying silent. That is the free-delivery case you named, and the one number worth seeing is the one you chose to give away.

A charge that matches the trip says nothing at all. When the two agree there is no line, because a difference of nothing is not a difference, and an "RM 0.00" row on every card would be noise you would learn to skip.

A delivered trip keeps the reading. The calling-off deadline disappears when a trip is over, because there is nothing left to call off. This line stays, because a cost is money that has already been spent and the whole point is to learn from it afterwards.

This is on your screen only. Nothing here goes near a customer. The price a customer was quoted does not move, their messages do not change, and their tracking page is untouched by this version.

Nothing else changed. The van you pick, the prices, the five-minute life of a price, the booking buttons, the charge box and the calling-off deadline are all exactly as they were.

29 Sep 2026 - engine v234, A BOOKED TRIP NOW TELLS YOU WHEN CALLING IT OFF STOPS BEING FREE (no database step, no redeploy, one push).

This one started as your own question: if you book a trip and then need to call it off, when does it start costing you money? The answer is a rule rather than a feeling. Lalamove lets you call off a scheduled pickup at no charge up to 45 minutes before it, and it may charge a fee after that. That rule was living in your head, so this version puts it on the page. Under the status line on a booked trip, the card now says **Free to call off until 30 Sep, 10:15 am - Lalamove may charge a fee after that.**

It turns over by itself when the window shuts. Leave the card open and the sentence changes on its own to **Lalamove's free calling-off window shut at 10:15 am, so a fee may apply from here.** You never have to work out which side of the deadline you are on, because the card says it for you.

A trip booked for as soon as possible is told apart honestly. On an immediate booking the free window runs from the moment a driver takes the job, and Lalamove never sends that moment back to the app. So rather than inventing a time that could be wrong, the card states the rule instead: **Booked for collection as soon as possible, so there is no pickup time to count back from. An immediate trip is free to call off only for a short while after a driver takes it - check with Lalamove before you count on it.**

Nothing is claimed that the app cannot stand behind. Every sentence says the fee may apply rather than that it will, because it is Lalamove's own terms that decide it and not this app. A trip that has finished, or one you have already called off, shows no deadline line at all, since there is nothing left to call off.

Nothing else about booking changed. The vans, the prices, the five-minute life of a price, the customer's tracking box, the charge box and the buttons on this card are all exactly as they were. This is one line added to a card you already had.

29 Sep 2026 - engine v233, YOUR ORDER PAGE NOW HAS A WAY BACK TO YOUR HOMEPAGE (no database step, no redeploy, one push).

Your order page had no route to your homepage at all. Someone who lands on it straight from an Instagram bio or a shared link sees the bakery name and the menu, and if they want the reviews, the gallery or the story, there was nowhere to tap. From this version there is a small **Our homepage** link in the orange banner at the top, marked with a little house, on its own line under your tagline - and it takes them to **jienluv2bake.com.my**.

It sits on its own line on purpose. The banner's top row already holds the "Made to order" line and the EN / Mandarin / BM buttons, and on a phone the "Made to order" line already wraps onto two or three lines - a third thing squeezed in there would have crowded all three. Under the tagline it belongs to the banner, reads clearly in white on the orange, and is a full-size tap target rather than a thin line of text.

It is translated like everything else on the page. A customer reading in Chinese or in Bahasa Malaysia sees it

in their own language, and it changes with the rest of the page the moment they tap a language, with nothing re-loaded.

Nothing else moved. The menu, the delivery calendar, the basket, the order form, the track-your-order box and the place-order bar are all exactly as they were. Your homepage still opens your order page in a new tab, so a customer who came that way keeps both pages open.

28 Sep 2026 - engine v232, A TRANSLATED LINE YOU DELETE STAYS DELETED, AND THE GREY HINT IS SHORT ENOUGH TO READ (no database step, no redeploy, one push).

This one comes straight from you, and you described it exactly: you delete the Chinese or the Bahasa Malaysia wording on a product, and **the machine translation comes back** - greyed, sitting in the box you had just emptied, with the little sentence about filling in the English, and the right-pointing arrow over it. Then the arrow went away and left the greyed words behind with no way to accept them.

What was actually happening. The greyed words and the arrow are supposed to be **one offer**, made together: the arrow is drawn on the words, and the words are only there because the arrow can take them. They were being written in two different places, so they could disagree - and when you emptied a line, the words stayed and the arrow did not.

From this version a line you have made yours is yours. If you type your own wording into a translated line, or you **delete** what was there, that line goes **quiet**: nothing greyed comes back into it, and no arrow is put over it. It stays as you left it when you save.

The circular-arrow button is still your way back. A line you have emptied shows the **circular-arrow button** on its right instead of the arrow, so a translation is always one tap away if you change your mind - the line is never a dead end. The same circular-arrow button re-translates a line the machine filled for you.

A line you have not touched is unchanged. Empty, never typed into, with English above it - it still shows the suggested words greyed with the -> to take them, exactly as before. What is on offer, and how to take it, are now decided in the one place, so they can never disagree again.

And the grey hint is short enough to read now. It used to say **"e.g.if blank, it will be filled with English"** after the suggested words, and that is too long for a one-line box: on your phone it ran past the right edge and was **cut off mid-word**, right where the arrow sits. It now shows just the suggested words with the plain **"e.g."** in front, the same shape as every other suggested box in the app. The promise it was making is still made, in full, in the card's own line above: leave a line blank and the English shows instead.

A line you emptied says what blank means. Instead of the machine's words lingering in it, an emptied line now carries a small grey note - **"Left blank - English shows."** - so it is clear that leaving it empty is a real choice and not a mistake, and that your customers will simply see the English.

Nothing else moved. The translations still come from the same free service, still only ever appear as grey suggestions, are still never saved over your own words, and the shop still falls back to the English wherever a translation is missing. Nothing here touches your shop, your orders or your prices.

28 Sep 2026 - engine v231, THE PICK-A-DELIVERY-DAY CALENDAR IS NOW FIVE WEEKS THAT FOLLOW TODAY (no database step, no redeploy, one push).

The calendar on your shop was a **month** - the whole of September, then the whole of October. That sounds tidy and it is not: on the 28th, almost everything on it is days you can no longer sell, and the days just before the 1st and just after the last day of the month were **blank squares**, so a delivery day at the start of next month was invisible until the month turned over.

From this version it is **five whole weeks anchored on today**, and it moves with the date.

The first row is the week just gone, faded the way it has been faded, and faded **harder than before** now that there is a week of it - it is there as a sense of where you are in the month, not as something to read. **The second row is this week**, with today ringed so the customer can see where "now" is. The three rows under that are the weeks ahead, which is where the delivery days actually live.

Every square is a real date now. No blanks, no padding - the days before the 1st and after the last of the month are simply part of the row they fall in, so a delivery day at the very start of next month is visible and bookable from the moment you publish it.

Your delivery days are unchanged on it - still circled green with the day you have open, and still struck out where a day is full. The line under the calendar still names the day they are getting.

Paging, and the motion you asked about. If you publish a delivery day further out than the window reaches, **arrows appear either side of the title**. Tapping the one on the right moves the window **one whole week later** and the dates arrive **sliding up**; tapping the one on the left moves it back and they arrive **coming down** - the direction you are travelling in, so a tap reads as a tap. The arrows are **not there today**, and that is on purpose: all five of your published dates sit inside the one window, so there is nowhere to go and **nothing dead is drawn on your shop**. They will appear the first time you publish a date roughly three weeks out.

Your shop's rules are untouched. The 18:00-the-day-before closing, the sold-out days, the day the customer has chosen, the product counts, the refresh that keeps a page open in someone's hand up to date - all exactly as they were. **Nothing in the backoffice calendars changed**; this is the shop's calendar only.

28 Sep 2026 - engine v230, THE DELIVERY ADDRESS BOX IS NOW THE FULL WIDTH OF THE FORM (no database step, no redeploy, one push).

Last version made that box **tall**; you told me it was still cramped, and you were right - it was **narrow**, and I had not seen it. From this version the delivery address box is **as wide as the form**, so a whole address sits on the line it belongs on instead of wrapping three words at a time.

Why it was narrow. The order form lays its fields out in **two columns**, and the address was sitting in **one** of them - with the space beside it standing **empty**. The address is the longest thing you type into that form, so it now takes **both columns**, and that wasted gap is gone. On a phone it went from about **137 pixels wide to about 315** - the same width as the shop's own address box, which was never the problem.

It is in both order forms - the + **New order** card and the **Edit** pop-up - so the two behave alike, and **no other field on the form moved**.

The map card's address box is unchanged by this one. The box labelled "**The address you have**" was already full width - v229 gave it the height it needed, and nothing about its width was wrong.

Nothing else changed. The suggestions still appear as you type, the address still fills itself in when you pick a customer, and saving an order works exactly as it did. There is **no database step and no redeploy** - it is one push.

28 Sep 2026 - engine v229, THE DELIVERY ADDRESS BOX HAS ROOM FOR A WHOLE ADDRESS (no database step, no redeploy, one push).

A delivery address is four or five lines - the name, the street, the area, the postcode and the state. The box you typed it into was **one line tall**, so most of what you had written was hidden and you were reading your own address back through a letterbox. From this version that box is **tall enough to show the whole address at once**.

It is in both order forms - the + **New order** card and the **Edit** pop-up - so the two behave alike. The box is also **draggable at its bottom edge** if you want it taller still.

The map card's address box got the same treatment. The one labelled "**The address you have**", where you type the address before looking it up, is now multi-line as well - it holds a whole address too, and it had the same problem.

Nothing else changed. The suggestions still appear as you type, the address still fills itself in when you pick a customer, and saving an order works exactly as it did. There is **no database step and no redeploy** - it is one push.

28 Sep 2026 - engine v228, THE ADDRESS BOX SUGGESTS REAL ADDRESSES AS YOU TYPE (one Google setting to switch on, one redeploy - the app itself is one push).

You said last version's address helper only helped people who had ordered before, and that what you actually wanted was help **keying in an address for someone new**. That is what this version is.

From now, in the **+ New order** card and in the **Edit** pop-up, a **short list of real addresses appears under the delivery address box as you type**. Type the first few words of the street, and the list offers the addresses that match - each one a **whole address with its postcode and area**, not the half-remembered line you would have pecked in yourself. Tap the one you want and it fills the box. Then the app's own pin lookup can find the door from there, as it already does.

It never blocks you and it never changes what you have typed. Your words go into the box the moment you type them, whether or not the list appears, and nothing is saved until you press **Add order** or **Save**. If the list does not appear - no signal in the kitchen, or the lookup simply found nothing - you type the address by hand exactly as you do today, and the order saves just the same. **No suggestion can ever hold up a sale.** The list is the only thing that is ever missing.

It is only the delivery address, and only for you. The shop's address box is public and stays exactly as it is; this is for the orders you take by hand.

What you have to switch on, once. This one genuinely needs a Google setting: **Places API (New)** has to be turned on for your Google project, and ticked onto the key you already have, and the **courier** function has to be redeployed. **Until you do that, nothing is broken** - the list simply never appears and the address box behaves exactly as it does today. The step-by-step is in the Google Cloud setup guide, and I will walk you through it on your Mac.

What it costs: nothing at your size. The first **10,000 lookups a month are free**, and an order uses a handful. There is deliberately **no daily limit** on this one - you asked for it to keep working rather than stop partway through a day, so a runaway month would be a few ringgit rather than nothing. Only your signed-in phone can reach it, too, so no one else can spend it.

Worth knowing: it sends your customer's half-typed address to Google as you type it. That is a step further than the door lookup, which sends a complete address once. It is the price of a suggestion list, and it is the same Google account the door lookup already uses - but you should know it rather than find out later.

28 Sep 2026 - engine v227, THE CUSTOMER'S ADDRESS FILLS ITSELF IN WITH THEIR NAME (no database step, no redeploy, one push).

When you type a customer's name into an order form, the app already offers you the people you have served before, each with their number and what they usually buy. Tap the right one and their **name and number fill in**. From this version their **delivery address fills in with them** - the address on the last order you actually delivered to them.

It is a default, never an overwrite. The address fills only a box that is still empty; it does not touch one that already says something. So an address you typed yourself, or the one already on an order you are editing, stays exactly as it is. A pick can never quietly move a delivery to a different door. And if a customer has only ever collected, or has no address on any of their orders, the box is simply left alone - nothing is written into it.

Both order forms do it - the **+ New order** card and the **Edit** pop-up - so the two behave alike rather than one of them surprising you.

It never asks the internet for anything. The addresses are the ones already in your app, from your own orders, so there is nothing to pay, no key to switch on and no Google account involved. It draws on the orders you have saved, so a customer who has never ordered before has nothing to offer yet, and it gets more useful as your customer list grows.

28 Sep 2026 - engine v226, THE SECOND KIND OF COURIER: PARCELS YOU POST YOURSELF, RECORDED IN THE APP (no database step, no redeploy, no key to set - one push).

You asked whether Lalamove is one kind of courier and SPX, Ninja Van, Line Clear and EasyParcel are another kind, and whether we could just have **two kinds**. Yes, and this is the second one.

There are exactly two kinds. Lalamove is *kind one* - on demand, point to point: it prices a vehicle for a journey,

a driver accepts, and you get a name, a plate and a live share link. J&T, Ninja Van, Line Clear, Pos Laju and SPX are *kind two* - parcel carriers: they price a box by weight and by where it is going, you hand the box over, and it arrives next day or later. **EasyParcel is not a third kind** - it is one account and one API standing in front of most of kind two. So there are two, exactly as you said, and kind two is what this version adds.

What is new: More -> Parcel couriers. A list of your own, the same shape as your suppliers. You add a carrier by name, and there is **one press that adds the usual five** - J&T Express, Ninja Van, Line Clear, Pos Laju and SPX Express - so a fresh phone fills itself in. You can rename or delete any of them at any time.

On a courier order you now record the parcel. Open an order, and under the fulfilment picker you choose **which carrier took it**, type the **consignment number**, and press **Handed to the carrier** when you have actually handed it over. Press **Undo - not handed over yet** and that moment is taken back, while the carrier stays on the order. The cost goes into the **courier charge box the order already had**, with the same question of who paid it, so nothing about your books changes shape.

The order row shows one quiet line - for example *J&T Express · JT998877665 · handed over* - and only once a parcel is recorded. An order that is not a parcel looks exactly as it did before.

The customer is told. Their track card gains one line naming the carrier, and once you mark the parcel handed over it reads **Collected** - the same word a Lalamove pickup publishes. That carrier line is drawn **only on an order with no driver**, so every existing Lalamove card is left exactly as it was.

A product can now be ticked "Can travel as a parcel", and that tick blocks nothing. A courier order always offers the carrier picker. If a line is not marked, the app simply **names it** - "Not marked as able to travel as a parcel: Fresh Focaccia" - and you can still send it that way. Nothing is hidden, removed or turned off, because no rule in this app may block a sale you take by hand. The tick only feeds that one warning, and it is where a parcel's weight and box size will sit when the time comes.

Two things you should know about how it behaves. A parcel and a booked Lalamove trip are **mutually exclusive on one order** - whichever exists, the other's controls are replaced by a sentence naming it, and neither is deleted. And an order recorded as a parcel **can no longer be swept into a delivery run**, because it is not going on your van.

No price is shown anywhere, and that is deliberate. A parcel's cost moves with the zone and the box, and the way a box is measured into a price is not the same at every carrier - J&T and Line Clear divide by 6000, Ninja Van by 5000, which is about 20 per cent more on the identical box. One stored price per carrier would quietly lie to you, so there is no rate field at all. **There is also no key, no wallet, no account and no database step in this version** - you keep booking the parcel exactly as you do now, on the carrier's own site or over the counter, and the app remembers the rest.

Be prepared for this to look like almost nothing in your own shop, and that is correct. Every product here is fresh focaccia, and a parcel network is not where fresh bread goes. This is an engine feature and it is for **Munchies**, which is the shop whose products are the kind that survive a night in a van. You told me the plan: the parcel work is done here, in Bakeadmin, and Munchies syncs to it. When the booking side does come, it comes through EasyParcel with **Munchies' own account and its own wallet**, kept separate from this shop's - which is what you decided, and it is the right call.

One more thing, carried over from last version. A product that is on the shop can always be *hidden* - see v225 below - so nothing here changes how you take a product down.

Two faults were found and repaired on the way, and neither had ever reached you. In the app's own picker, the code that was supposed to run when you choose a new value had been quietly dead for about a hundred versions. **Nothing you could see was wrong**, because nothing in the app had yet depended on it - the new parcel card is the first thing that does, and a control that arrives one reopen late would have read as missing rather than as conditional. It is repaired once, in the shared picker, so it is now correct for the parcel card and for anything added later. The second was in the new card itself: what you had typed was being thrown away whenever the form redrew, and that is fixed before it ever shipped.

28 Sep 2026 - engine v225, A PRODUCT THAT IS ON THE SHOP CAN ALWAYS BE HIDDEN (no database

step, no redeploy - one push).

You said *after publish a product, only delete is allow, it should not be*. You were right, and this is fixed. A product that is **on the shop** now always carries a *Hide* button. Before, the row decided for you: a product that already had orders, or that another product uses to make a set, offered *Hide*, but a **freshly published product with no orders yet offered only Delete**. So the one product you had just put up was the one you could take down only by destroying it.

Hiding never touches your recipe. Hide simply takes the product off the customer's menu and leaves everything of yours intact - the recipe, the ingredients and amounts, the photo, the price, the category, and any order history it has. *Unhide* puts it straight back. Nothing is deleted, and no database step is needed for any of this.

Delete is still there, beside Hide, on a product that is clean and unused - one with no orders and not used inside a set. That is the same rule as before, just no longer the only button on the row. So on a live product you now read *Edit · Hide · Delete*, and on one with history you read *Edit · Hide* (Delete never appears where it would break an order or a set).

Why this helps: the mistake this removes is a real one - publish a product, look at the card, and the only way off the shop is the button that also deletes the recipe. Now taking a product down is one tap and completely safe, and Delete stays for the case where you really do want it gone.

28 Sep 2026 - engine v224, THE PICTURE WINDOW IS SQUARE AND YOUR PHOTO IS CROPPED TO FILL IT, AND A PRODUCT ROW SHOWS ONLY WHAT A CUSTOMER SEES (no database step, no redeploy - one push).

You asked me for the window to be *square*, and *photo upload can be auto crop to fit the window*. Both are done.

The window is a square, and **your photo is cut to a square as it is saved**, so it fills the window edge to edge with no cream strips down the sides. The middle of your picture is the part that is kept, and the app shrinks it for you at the same time, so you never have to resize or measure anything. What you pick is exactly what shows.

This changes one thing from last version (v223), where your photo was shown whole. In v223 the window kept its shape and your photo kept its own, with the window's soft cream colour filling the spare space. You have now asked for the other thing, and this is it. Everything else about v223 stands: the window is still **one standard size in all three places** - your shop's card, your app's Products list, and the editor where you choose the photo.

What that means: whatever the square cuts off is gone. A very tall or a very wide photo loses its outer edges. If a crop has taken off a part you wanted, choosing the photo again is how you change it - the app keeps only the small square copy, so a picture cannot be un-cropped afterwards.

The pictures you already saved still work, and are now trimmed to the square. They were saved in whatever shape they were, so on your shop and in your list they show trimmed to the middle of the square window. Choosing the photo again gives you one cut for this window properly - that is the step you told me you would take.

A product with no picture is unchanged. It simply has no window, and its name and price read across the whole card exactly as before.

One thing to know, carried over from v223: a wider window leaves the words beside it a little less room.

On a product that has both a picture and a long name, the name and price wrap onto one or two more lines than they used to. If that crowds the words more than you like once you see it, say so and I will make the window a little smaller - it is one number and a quick change.

The other half of this version: your product row is now only what a customer sees. You said *we dont need ingredient and cost price for products in app that shown on. On the shop*. So both came off **every** row in your app's Products list - the rows on the shop, the drafts, and the hidden ones alike.

What a row shows now is the name, the unit, an optional daily limit, the price you sell at, and the category. The two things that went are **the cost per unit** - the small figure that used to print beside your sell price, worked out from the recipe's ingredients - and **the ingredient line**, the list of what goes into the product

with the amounts. Both are gone from the row.

Nothing was deleted, and nothing about your recipes changed. Both are still on each product's **Edit** screen exactly where they were, where you build the recipe and where the cost is worked out. The row was only ever meant to be a view of what you sell; the recipe is yours, and it stays yours on the Edit screen.

Why this helps: the row is shorter, so more products fit on the screen at once, and the one figure left on it - the price a customer pays - is the one you are actually checking when you scan the list. If you want to see a cost, tap **Edit**.

Nothing to run. No new table, no Supabase change, no redeploy of anything. This reaches your shop and your app the next time you push, and you will see **Engine v224** on your **More** screen.

28 Sep 2026 - engine v223, ONE STANDARD PICTURE WINDOW, BIGGER AND THE SAME IN BOTH PLACES (no database step, no redeploy - one push).

You asked me for a *standard photo window that is big enough for effectively shows a photo of product for APP side and Store side*. That is exactly what this is: one picture window, one size, used in every place you meet a product picture.

The window is the same size in all three places: your shop's card, your app's Products list, and the editor where you choose the photo. Before this version each place had its own size - 96 wide on the shop and following the card's own height, and only 64 in the app - so the same picture showed three different sizes, and none of them big enough to make out the product. Now it is one window everywhere, and it is bigger than before.

Your photo still goes in whole. Nothing about last version changes. A tall photo and a wide photo both still show in full, with the window's soft cream colour in the spare space beside them. The window is the fixed thing; your photo keeps its own shape.

A product with no picture is unchanged. It simply has no window, and its name and price read across the whole card exactly as before.

One thing to know: a wider window leaves the words beside it a little less room. On a product that has both a picture and a long name, the name and price now wrap onto one or two more lines than they did. If that crowds the words more than you like once you see it, say so and I will make the window a little smaller - it is one number and a quick change.

To use the new window at its best, choose the photo again. The pictures already saved were cut to the old shape back when they were saved, so re-choosing the photo is the step that gives you one made for this window. You told me you would reload the product photo, and this is the version it belongs to.

Nothing to run. No new table, no Supabase change, no redeploy of anything. This reaches your shop and your app the next time you push, and you will see **Engine v223** on your **More** screen.

28 Sep 2026 - engine v222, YOUR PHOTO KEEPS ITS OWN SHAPE AND NOTHING IS CUT OFF (no database step, no redeploy - one push).

You said the photo should be *restored to its original wide tall ratio*, and then you told me what to do about it: *fix tall and wide*. So this version fixes both. Your photo is no longer cut to fit a shape. It now goes on the card exactly as the shape you took it in, and the picture panel stays the same size and holds the whole of it.

A tall photo and a wide photo both work now. Take a photo of a tall stack of focaccia and the whole stack shows, top to bottom. Take a wide photo across a tray and the whole tray shows, side to side. Nothing is trimmed off the top, the bottom or the sides, whichever way round your photo is.

The picture panel is the thing that stays the same, not the photo. The panel beside a product is still the same size it was, so every product's card on your shop still lines up with the next. If your photo is not exactly the panel's shape, the panel's own soft cream colour shows in the spare space beside it. That is the panel, not a mistake, and it is the trade for never cutting your photo again.

This is the same in your own app. The picture beside a product in your Products list, and the preview in the

product's editor when you choose a photo, both show the whole photo the same way.

A photo you have already chosen is still the old cut. I have to be straight with you about this one: the cutting happened when the photo was saved, so the pictures already in the app are the old cropped version and I cannot put back what was trimmed off. They will show whole on the card now, which is an improvement, but they are still the old shape. To get a picture that uses this properly, open that product in the app, choose its photo again, and the new one is kept in its own shape.

Nothing to run. No new table, no Supabase change, no redeploy of anything. This reaches your shop and your app the next time you push, and you will see **Engine v222** on your **More** screen.

28 Sep 2026 - engine v221, THE PICTURE IS NOW BIGGER AND RUNS THE WHOLE HEIGHT OF THE CARD (no database step, no redeploy - one push).

Last version I made the picture a tall shape. You told me you meant *bigger - take the fullest height as possible*. You are right, and last version only changed the shape; this one changes the size.

The picture now runs the whole height of the product's card. It starts at the top of the card and carries on down to the bottom, past the price and the description and alongside the + and - buttons, so the strip beside it is never left empty. On your shop the picture is half again as wide as it was, and the full height of the card rather than a small square in the corner. It is the biggest the picture can be before your product names would start breaking onto a third line, which is the point you asked me to stop at.

Your own Products list is the same. In the app, the small picture beside a product is now bigger and stretches from the product's name down past everything written about it. A product that has a description therefore gets a taller picture than a bare one - the picture follows the row rather than being a fixed stamp, so nothing is left empty and nothing is cut off.

A new photo you choose is now stored bigger, so it stays sharp at this size. The file the app keeps is about 10 KB instead of about 6 KB. That is still tiny - your whole app still travels in one small package to the cloud and to every customer's phone - but it means the picture will not look soft now that it is drawn so much larger. Choose a product's photo again in its editor and you get the bigger version automatically.

The pictures you already have are fine, and nothing is stretched. An old square picture is fitted into the new tall box by taking a slice from the middle, so it fills the height without being pulled out of shape. It will look right straight away. If you want a picture that uses the new size at its best, choose that product's photo again - a tall photo of the tray or the plate is what suits this shape.

That middle slice is gone as of v222 above: the picture is no longer cut to shape at all, and the panel holds the whole of it instead.

Nothing to run. No new table, no Supabase change, no redeploy of anything. This reaches your shop and your app the next time you push, and you will see **Engine v221** on your **More** screen.

28 Sep 2026 - engine v220, THE PICTURE IS NOW TALLER THAN IT IS WIDE, AND YOUR OWN LIST ARRANGES ITSELF LIKE YOUR SHOP (no database step, no redeploy - one push).

Your request: *the thumbnail photo should be taller* - and, when I asked how far it should reach, *taller for store and app*. Both places, then: the picture a customer sees on your shop, and the small picture beside the product in your own Products list.

The picture is now a tall shape rather than a square - four parts wide to five parts tall. Last version you chose a square, and a square crop of a tray or a plate cuts off the top and bottom of the food, which is usually the part worth seeing. The app crops your photo to the tall shape by itself, taken from the middle, exactly as it did before - you still choose a photo and the app does the measuring. The same crop now sits at the left of the customer's row, in your own Products list, and in the little preview in the product editor, so the picture you pick is the picture you get in all three places, and no part of it is squashed out of shape.

Your existing product pictures are not lost, and they are not stretched. The photos you have already chosen are squares, and the new box is taller; the shop shows a square photo inside the tall box by taking a slice from the middle and filling the rest of the height, so it looks right and nothing is distorted. If you want a

picture that uses the full height, simply choose that product's photo again in its editor and the app stores a genuinely taller one. Nothing to do if you are happy with what is there.

Your Products screen now arranges itself the way your shop does. Your request: *in app the product should arrange like it should at shop*. This is the change you will notice most. The three lists you already have - **On the shop, Draft and Hidden** - all stay, because those are about whether a product is on sale, not about what it is. Inside *each* of the three, your products now group under your own headings, in your own order, in exactly the order a customer scrolling your shop would meet them. A product with no heading yet sits last under a plain heading called **More items**, the same word your shop uses, so what you see in the app is what a customer sees.

And a product can be moved inside its heading, by a handle. Your request: *and it should be able to swap position by a handle in same category*. Beside each product's row there is a small handle; press it and drag the row up or down, and the products in that heading re-arrange behind your finger. You asked for a handle rather than up-and-down buttons or number boxes, so a handle is what it is. The handle appears only where there is something to move - a heading holding a single product shows none, because there is nowhere for it to go - and a drag can never pull a product out of its heading into another one, which is what you asked for. The position you set is stored, not just drawn: it travels to your other phone and it is the order your shop lists in, so dragging here is the same thing as arranging your shop.

One thing that was quietly broken, found while building this, and fixed. Adding a heading, renaming one, re-ordering one or deleting one did not reach your shop on its own. The shop only caught up the next time something else published - so a heading you renamed could sit under its old name on the customer page until you happened to edit a product. Every one of those four actions now publishes on the spot. It was there in v219 and you would have met it as a heading that "would not change"; it does now, at once and on its own.

Nothing to run. No database step and no redeploy: this is the app's own screens and the shop page, and one push is the whole of it. **Engine v220** on your **More** screen once this is pushed.

28 Sep 2026 - engine v219, YOUR SHOP NOW FOLLOWS YOUR OWN HEADINGS, AND EVERY PRODUCT CAN WEAR A PICTURE (no database step, no redeploy - one push). Product categories and thumbnails.

Your request: *I need a category to be setup for each products, and category be manage like how we manage ingredient. Now our product listed in store are random, when there is category existed, the list will follow category, and in a category sort my name, or other criteria. Need a thumbnail photo for each product as well, with photo size adjusted accordingly.* - and, a moment later, *one product can be with more than one category*.

There is a new screen, and it starts empty on purpose. More -> Categories is where your headings live, built the same way your ingredients are: press **New category**, give it a name, and it appears in the list. Nothing is in it until you put it there, because your shop should say what *you* sell, not what a program guessed for you.

Your shop now lists by your headings instead of in whatever order the products happened to be stored in. That is the whole point of this version. Under each heading sit the products you filed there, and the order they appear in is the order *you* set - not alphabetical, not by price, not by anything the app decided. Your own arrangement is the arrangement.

A product with no heading yet is never hidden - it is shown last. A customer scrolling to the bottom of your shop finds a plain heading called **More items** with those products under it. So if you add a loaf this evening and file it tomorrow, it is on sale tonight; it simply sits at the end until you say where it belongs.

A product can be in more than one heading, and you choose which one it is listed under - by the order you tick. In the product editor your categories appear as a list of tick boxes, and the box you tick *first* is the heading the product is drawn under. Tick a second one and the product does not vanish from the first, and it is not drawn twice either: **one product appears in exactly one place on your shop**, under its first tick. The other ticks are kept for you, as the note that says "this one is also a snack", and the product's own row in the app spells them out so you can see them at a glance. You asked for this to be made simple rather than switchable,

and that is what it is: *1st category, 2nd category, 3rd category; the first one groups it.*

A heading can hold headings under it, as deep as you like. Your bakery's headings are flat - Food, Drink, Snack, Appetiser - and they can stay that way. The engine underneath is shared with the Munchies shop you are planning, where you said the shape is **For Dog, For Cat** at the top and then **Pork, Duck, Fruits & Vegetables** underneath, so that is built in now and will need no second version later.

Every product can have a picture, and it is a square. Choose a photo from your phone in the product editor and the app crops it to a tidy square by itself, taken from the middle of the picture - you do not have to resize or measure anything. Your shop shows it at the left of the product's row, and your own Products list shows it too, so you can see at a glance which products have a picture and which are still blank.

Two things to know about pictures, said plainly rather than found out later. First, **the homepage is untouched** - you asked to leave it alone, and it still carries the photographs already written into it, so a picture set here reaches your shop and your own product list but not the front page. Second, **a picture costs weight twice**: it travels inside the copy of your records that every backup and every export carries, and it is fetched again by every customer who opens your shop. That is why it is kept small - a square about 160 pixels across, in the compact photo format - which looks right on a phone and keeps your shop quick to open. A picture that is far too large, or in a format the shop cannot show, is quietly left out rather than published broken.

Moving a heading, or a product inside one, is a drag. Each row carries a small handle at its left; press that handle and drag the row up or down, and the list re-arranges behind your finger. You asked for exactly this - a drag rather than up-and-down arrows or number boxes - so that is what it is. It is the first drag in the app and there was nothing here to copy from, so it is the part most worth a second look on your own phone.

Deleting a heading is guarded, the way your ingredients and suppliers are. A heading that still has headings under it will not delete. A heading with products filed in it will not delete, and it tells you how many. An empty one simply goes. There is no undo, and the question it asks is the last thing before it happens.

Two things you would have spotted, and they are fixed rather than left to puzzle over. The row for a heading shows how many products your *shop actually draws* under it. If you have ticked a product into a later heading, the shop has moved it there, so the row now also says *"1 more ticked here, listed elsewhere"* instead of quietly counting it and leaving you to wonder why the number does not match your customer page. And a product's own row in the app now *names* the other headings you ticked instead of saying "1 more ticked", because what you want to know is *which*, not how many.

Nothing to run. No database step and no redeploy: this is the app's own screens and the shop page, and one push is the whole of it. **Engine v219** on your **More** screen once this is pushed.

27 Sep 2026 - engine v218, ONE CUSTOMER'S DOOR IS NOT ANOTHER CUSTOMER'S BUSINESS (no database step, no redeploy - one push). The delivery run's tracking link.

Your report: *Can the courier link be not shared as one, as when the customer track it, they become aware of the other drop off point.*

You were right, and there is no way to split it. The courier hands back *one* link for the *whole trip* - its own documents describe that link as a thing for sharing delivery information with a third party, and there is no per-drop-off link at all. So a customer cannot be given live tracking of their own box without also being handed every other doorstep on the van. The only way to stop them seeing each other is to stop giving that link to customers on a run, and that is what this version does.

A run of two or more doorsteps no longer puts the courier's link on the customers' orders. It does not reach their track page and it does not go in their WhatsApp message, so there is nothing for them to open that would show them anybody else. **Everything else they had, they still have:** the app's own tracking card keeps saying how far the delivery has got - finding a driver, on the way, collected, delivered - and it still carries the driver's name, the plate to look for at the gate, and the press that rings him. All of that comes from your own app rather than from the courier's page, so none of it can leak a second address.

What changes for them is the live map. Without the courier's link there is no moving dot for a customer to watch, and their page is as fresh as the last time you press **Check the trip** - which is exactly how it already works for any order that has no link in its tracking box. That is the price of not showing them each other, and it is worth saying out loud rather than discovering.

A single customer's own courier order is untouched - you chose this, and it is deliberate. When you book one order on its own, that trip is your bakery to that one door and nobody else is in it, so its link gives nothing away and the customer keeps the live tracking exactly as before. The same is true if you ever run a day with one doorstep on it. **The rule is: the courier's link reaches a customer only when the trip carries nothing but that customer's own doorstep.**

The card tells you which one you are about to do, before the money is spent. The booking question on the Delivery screen now says which of the two it is: on a run it says the courier's link is deliberately kept off the customers' orders and why; on a single doorstep it says the customer's tracking box takes the link. And the note afterwards agrees with it instead of claiming everyone now shares one link. **Your own card still shows you the link** on every trip, run or not, because that is the one you use to check the trip and to call it off.

One more line promised the old thing, and it was found by driving a real run rather than by reading the code. Under the prices, the Delivery run screen's own note still read that the trip had *one share link that goes on every customer's own track card and message* - true until this version, and the exact opposite of what this version does. Nothing was asserting that line, so it had survived the first pass of this work; pricing a two-customer run in a browser is what caught it saying one thing while the booking did another. It now says what actually happens, and a test names that line by its own words, so the screen and the booking cannot drift apart again.

Nothing to run. No database step and no redeploy: this is the app's own screens, and one push is the whole of it. **Engine v218** on your **More** screen once this is pushed.

27 Sep 2026 - engine v217, A BUTTON THAT COULD GO QUIETLY DEAD CANNOT ANY MORE (no database step, no redeploy - one push). The courier price and booking presses.

Your report: *the get price from lalamove not responding*.

What I found, and what I could not. I drove both screens here and could not make the press go dead - the **Delivery** run's **[Price this run with Lalamove]** answered with a real price, and the order card's press answered too (it said the courier had refused, which is what this machine's test login gets). So this is not a version I can tell you was your fault - it is the one shape of fault that leaves no trace on the screen, and I found it by reading the code rather than by watching it happen.

The shape of it. Every press in those two screens carried its own little flag - it raised the flag before it asked Lalamove anything, and lowered it again on each way out that it knew about. Every throw that happened in between left the flag up. And a raised flag makes the button return **every** later press instantly, with **nothing said at all**, on a button still drawn grey. That is a dead control: it looks pressable, it is pressable, and it will never do anything again for the life of that card. It is also exactly what *not responding* looks like from your side. The app's own rule for this is written in the courier code: an error thrown at a button in a card, while you are standing in a kitchen, is *a dead screen with no words on it* - and the flag was the one place that rule was not kept.

What happens now, every time, on all eight presses. The flag is raised and lowered by one shared piece of code, and the lowering happens in its **last** step rather than on a list of happy endings - so **there is no way out of a press that leaves it up**. The button always comes back. And a throw is **said**, in a sentence, in the same place the price would have been: *The price could not be asked for, and nothing has been priced - [what went wrong]*. Nothing is ever swallowed, and no press is ever swallowed again. This covers **Get a price from Lalamove, Get a price / Price this run, Book this trip, Book this run, Check, Call it off**, and the one-press address lookup - on the order card **and** on the Delivery run.

The one sentence that is deliberately not a promise. A booking that throws after the request left but before the reply came back cannot tell whether the trip was booked. That one says *The booking could not be finished* -

check the trip in Lalamove before pressing again, in case it went through. It does **not** say nothing was booked, because it cannot know that, and telling you the money did not move when it might have is worse than telling you to look. An **unbooked** trip still says so plainly, as before.

Nothing to run. No database step and no redeploy: this is the app's own screens, and one push is the whole of it. **Engine v217** on your **More** screen once this is pushed.

27 Sep 2026 - engine v216, THE CHARGE BOX OPENS ON THE CUSTOMER, AND A GREYED BUTTON FINALLY SAYS WHY (no database step, no redeploy - one push). Two days after v215.

Two things you asked for, in your own words: *Can you default the who paid courier to The customer paid it?* and *now the greyed out book button.*

1. Who paid the courier now opens on The customer paid it. Open an order's **Note / tracking** card - or its **Edit** form - and the question **Who paid the courier** is already answered: **The customer paid it**. Type the courier's fee and it is theirs: the line under the box moves straight away, *The customer owes RM 38.00 - items total RM 30.00 + courier charge RM 8.00*, and **Save** writes it in one press. Before this, that same amount went in with **nobody named as the payer**, which is what v215 had to refuse - so this version removes the step you were being asked to take every time.

It is a default, and never an overwrite. An order that *already* records a payer opens showing the answer that order carries - press **Save** without touching it and nothing about it changes. **Not recorded** is still in the list, still yours to choose, and still how a charge is **deleted** (clearing the payer takes the amount with it, the rule you set on 19 September 2026). And the v215 refusal is still there for the one press that needs it: take the payer back to **Not recorded** with an amount in the box and press **Save**, and the app refuses in words rather than dropping it.

The delivery run is deliberately left alone. On the **Delivery** screen the same question looks identical, but there it is read **before** the amounts and decides **which amounts they are**: settling on *The customer paid it* makes the app price every doorstep on its own and write a charge onto **every** customer's bill. That is a decision about money, so on that screen it stays a choice **you** make - the run still opens on **Not recorded**, and says so.

2. A greyed-out [Book this trip] now says why. You reported the button drawn grey with nothing to explain it. Each price row used to work out its own, shorter list of reasons - a price that had run out of time, no price id, fewer than two doors - so a price the courier's own reader refused for any **other** reason was drawn inert and **said nothing at all**. The reason is now asked of the one place that really refuses the booking, and printed once under the prices, in the courier's own words. The commonest one, and the one you are most likely to meet: *That price did not come back with the courier's own handle for each door, so it cannot be booked. Ask for a fresh price - and if this keeps happening, the app is reading the courier's reply wrongly, which is worth fixing rather than working around.* The same fix is on the **Delivery** run's **[Book this run]**, and the reason is drawn at the moment the row appears rather than a second later.

Nothing to run. No database step and no redeploy: this is the app's own screens, and one push is the whole of it. **Engine v216** on your **More** screen once this is pushed.

27 Sep 2026 - engine v215, THE COURIER CHARGE SAVES (no database step, no redeploy - one push). The order's charge box.

Your report: *the selected courier charges cannot save*. It was true, and this version fixes it.

What was happening. You pressed **Get a delivery price**, then **Use this fee**, and the courier's price appeared in the **Courier charge** box on the order. You pressed **Save**. The app said *Order updated* - and the charge was gone. The amount had been written into the box with **nobody named as the payer**, and a charge is the amount *and* who bore it. With no payer, the app read the whole thing as no charge at all, quietly dropped it, and still told you it had saved. A save that says it wrote something it did not write is the worst kind of fault, because there is nothing on screen to tell you to look again.

What happens now instead of that. Press **Save** with an amount in the box and no payer chosen, and the app **refuses, in words**: it tells you the charge is RM 12.50 but nobody is down as the payer, so it would not be

saved, and it gives you the two ways out - choose who paid the courier under the box, or clear the amount.

Nothing is written, and your card is left exactly as you left it - the amount still in the box, the note and the tracking number untouched. Choose a payer and press **Save** again, and it saves.

One thing this version is careful not to break. Clearing the payer is **your own way of deleting a charge** - the rule you set on 19 September 2026. So the refusal only applies to an amount that arrived while the card was open (from the courier's **Use this fee**, or typed by you). An order that already has a charge, whose payer you set back to *Not recorded* and then save, is still **deleted**, exactly as before. Two older tests pin that, and this version had to be narrowed until they passed again.

Both doors, one answer. The same refusal guards the **Edit** form's **Save changes**, in the same words - so a charge cannot be lost through one door after being protected at the other.

And the question you asked, answered. Booking a trip is not what saves the charge. The courier's charge saves on its own, with no booking at all - that was the whole of the bug. **[Get a delivery price]** asks the courier what it would charge and costs nothing. **[Use this fee]** copies that price into the charge box so you can save it. **[Book this trip]** is the one press that **sends a real driver to the bakery and spends real money** - press it when you are ready for the driver to come, and it is also what puts the trip's share link into the customer's tracking box. The charge is yours to record whenever you like; the booking is the dispatch.

Booking in advance: the app cannot do it yet. The time box on the card changes **the price only** - it does not schedule the trip, and the card now says so on itself. A booking is a driver sent now, so do not press **Book this trip** until you want the driver to arrive. When a booked trip can be scheduled for a later time, that becomes its own setting, and the card will say so then.

Nothing to run. No database step and no redeploy: this is the app's own screens, and one push is the whole of it. **Engine v215** on your **More** screen once this is pushed.

27 Sep 2026 - engine v214, THE ADDRESS IT FOUND IS WRITTEN IN FOR THEM (no database step, no redeploy - one push). The shop page.

Your words, once the lookup started landing on the right house: *the address is very accurate, it can go into the delivery address instead of customer type full*. So now it does.

What a customer sees. They type their address on the shop page. The list of doors appears underneath, as it has since v202. They tap the one that is their house - and the **Delivery address** box is filled in with the complete address the lookup found, in place of the part they had typed. Where they wrote *23 Jalan Seang Tek, Penang*, the box now reads the map service's own answer, house number, town, postcode and all. The pin still moves to that door, the map still comes to it, and everything else on the page is exactly as it was.

The house number is what decides it, and nothing else. The box is only written when the row they tapped **holds every number they typed** - so the address being written in is provably an answer about the door they meant, rather than a nearby one. If the row only reached the road, or a different town, or has one of their numbers but not the other, **their own words are left exactly as they are and only the pin moves**. That is deliberate, and it is the whole of the safety here.

Why this was refused for so long, and why that was right. Back at v205 you saw a tapped suggestion put its own words on your screen, and your words for it were that it *will contaminate the customer keyin address*. You were right: the free map services answer a Malaysian address with the **road**, and they throw the house number away. Writing a street and a town into the box would have replaced the customer's one complete address with a worse one. **What changed is not the rule but the service** - since v212 the lookup asks Google first, and Google's answer is a complete address, often with details the customer left out. So the rule you first refused is now safe, and the test for it is exact: the house number.

One consequence worth knowing. A customer who types an address with **no house number in it** - *Taman Sri Nibong, Penang* - can never have their words written over, because there is no number for the row to have found and so no way to tell a row that matched from one that guessed. Their words stand and the pin moves.

And the one thing that must not change, does not. The address on the order and the pin's name are still **one string** - the v205 promise. Whatever ends up in the box is what names the pin, so your screen still cannot show you two place names and have them disagree.

Nothing else moved. A customer who never taps a suggestion is unaffected. Editing the address afterwards still takes the pin, and still says so. **Use my location** and **Pin on the map** are untouched. The bakery side is untouched. No address the lookup could not match is written anywhere.

Nothing to run. No database step and no redeploy: this is the shop page's own code, and one push is the whole of it. **Engine v214** on your **More** screen once this is pushed.

27 Sep 2026 - engine v213, ASK THAT ADDRESS UP AGAIN (no database step, no redeploy - one push). The door card.

This is the missing half of v212, and without it you would have set up your Google key and seen nothing change.

The thing that was wrong. Once the app has worked out where a customer's door is, it keeps that answer and never asks again. Ask for a price a month later and it reads the saved point back and prices that. So every customer you pinned *before* v212 is holding a road-level point that nothing will ever revisit - and the better lookup v212 switched on only ever runs for an address the app has never looked up at all. Your key would have been set, working, and invisible.

So there is now one more press on the door card: "Look this address up again". It sits next to **Move this pin**, under the map. Press it and the app asks that address up a fresh, exactly the way it did the first time, and moves the dot to the answer it gets now.

It is a press, and never something the app does on its own. A pin that moves under your finger while you are reading the card is the exact fault v209 was written to stop, and this must not bring it back. Nothing here runs by itself.

It appears only where the app itself put the door. One press is offered where a lookup decided that point, and where the door was saved before the app began keeping a note of how it was made - because before that, a lookup was the only thing that ever wrote one. **It is not offered over a pin you placed yourself**, by dragging or by picking on the map: that is your correction, not a guess, and the app will not offer to hand your work back to a service. And **it is not offered over the customer's own pin** from the shop page: they were standing at their door when they dropped it, and the app never looked that one up in the first place.

Every press tells you what it found, and there are three answers. It moved: the dot is on the new point, and the line says *The door moved - the lookup answers this address with a different point now*. It did not move: the line says it found the same spot, and names the spot it is standing on. It could not be asked: the reason is printed in the service's own words, followed by *The door has been left as it was*. A press that says nothing at all would read as a broken button, so there is no such press.

Prices follow the same rule they always have. If the press moves the door, the prices on screen were quoted for the old point and they go, with the line telling you to ask again. If the press finds the same spot, nothing moved, so the prices stand - a re-lookup that finds the same road does not make you pay for a quote twice.

And the road warning from v211 keeps its own rule. If the fresh lookup still cannot find the house number, the door keeps the warning line and the press says so in the same breath - *The door moved, and it is still only the road, not number 23*. A point that has been asked about twice is not on that account a better point.

Nothing to run. No database step and no redeploy: this is all inside the app, and one push is the whole of it. **Engine v213** on your **More** screen once this is pushed.

27 Sep 2026 - engine v212, THE PIN CAN NOW LAND ON THE HOUSE (no database step - but this one needs two commands from you before it does anything).

Your words: *the pin dont pin to the right house*. You were right, and v211 only put a warning on it. This version fixes the cause. It also comes with a small job for you, so please read the last part of this note.

Why the pin could not reach the house, and why no amount of work in the app would have fixed it. The two map services the app asks are free and need no account. Both of them are built from the same worldwide map, and that map holds Malaysian **roads**. It holds very few house numbers. So when you typed **23 Jalan Seang Tek**, the service was not being difficult - the number 23 is simply not written down in the file it reads. The best honest answer it could give was the middle of Jalan Seang Tek. Nothing in the app could turn that into your customer's gate.

So the app now asks a service that does have the house numbers. Google keeps house-level positions in Malaysia. When a key for it is set up, the app asks Google *first*, and the two free services stay behind it as the fallback. That is the whole of the change: the answer you get is now the house rather than the road.

The map you see and drag does not change at all. It is the same map, on the same screen, with the same pin, exactly as it is today. What is being replaced is not the map - it is the quiet lookup, the thing that decides where the dot lands when you type an address. You chose to leave the map alone, and it is left alone.

This is asleep until you do your part. With no key set, the app behaves exactly as it did at v211 - nothing better, nothing worse, and nothing broken. Setting the key is what switches the better lookup on. You can do it today or in a month; nothing in between is half-on.

Your part, in a terminal, about ten minutes, once. Never type a key into our chat, and never put one in a file. There are four small steps.

Make the key. At console.cloud.google.com, make a Google Cloud account and turn on the **Geocoding API**.

Restrict it. Create an API key, and then restrict it so a leaked key can do nothing else: under **API restrictions**, choose **Restrict key** and tick only **Geocoding API**. Google will ask for a card on file. The first 10,000 lookups each month are free, and a bakery this size will not come close to that.

Cap it. Set a **quota cap** on the key - a daily limit you are comfortable with. This is the important one. Past the cap Google simply stops answering, and the app falls back to the free map services, so the pin goes to the road instead of the house. It cannot turn into a bill you did not agree to.

Set it and build it. In the terminal, in this folder, run the three lines at the bottom of this note - one to give the app the key, and two to build the two functions that use it.

What is behind the key, so the "card on file" sentence is not doing all the work. The key lives on the server, in the same place your Lalamove key lives. It is never put on the shop page, never in a file the browser downloads, and no phone ever holds it. A customer's phone only ever receives the finished answer. And the shop's own lookup has a rate limit in front of it already, so one person cannot spend your allowance in a loop.

If Google ever refuses - a spent cap, a mistyped key - the lookup does not stop and does not break. It falls back to the free services and you get a road-level pin with v211's warning line under it, which is exactly the behaviour you have today. Failing quietly back to what already worked was the whole design.

One genuine bug fixed on the way. One of the free services does sometimes send a house number, and the app was throwing it away - labelling a result that HAD found number 23 as **Jalan Seang Tek**, the road's name. So a correct answer looked like the road, and v211's own warning fired on it, telling you the pin was on the road when it was on the house. The number is kept now. This one is fixed whether or not you ever set up the key.

The commands (the = in the first one is required; without it the tool refuses):

```
supabase secrets set GOOGLE_GEOCODING_KEY=<the key> --project-ref hzpyblqygnntixkijeem supabase
functions deploy shop-geocode --project-ref hzpyblqygnntixkijeem supabase functions deploy courier
--project-ref hzpyblqygnntixkijeem
```

That is three lines, and the last two are two separate deploys on purpose: the lookup lives in both the shop's function and your app's function, and they are built separately, so doing only one would leave the other still asking the free services.

No database step. Engine v212 on your **More** screen once this is pushed.

27 Sep 2026 - engine v211, A LOOKUP THAT CAN ONLY FIND THE ROAD NOW SAYS SO (no database step, no redeploy - one push). The doorstep lookup, and the lines that report it.

Your words: *make it say this is the road not the house*. It does now.

Here is the thing you were seeing. A Malaysian address leads with its house number. The free map services the app asks - the ones that need no key and no bill - very often can only answer down to the street. So when you typed **23 Jalan Seang Tek**, the answer came back as **Seang Tek Road, George Town, 10400**: your street, named as your street. The pin landed on the road. And the app reported that answer in exactly the same voice it uses when it has found the actual house, so there was no way for you to tell the two apart.

From now on the line says which one it is. Where you type an address with a number and the geocoder's own answer does not contain that number, the line adds, in plain words: *The lookup found the road, not number 23 - drag the pin to the door*. Where the answer *does* contain your number, the house was found and the line stays quiet - no needless worrying.

It shows in all three places that report a lookup. The line under the pin window's **Look it up** press, whether it found one match or several. The card's own line under **The door the driver is sent to**, right where you are already reading where the dot is. And each row of the **Price this run** list, so a run you price in one press tells you the same thing about the same door.

And it is remembered, so it keeps being said. The number the lookup missed is written down with the door itself, not just printed once - so the caveat is still there on the next repaint, and on the next phone that syncs, and it survives the one-press switch from one door to the other. A door you place by hand or drag yourself carries no caveat at all, because no lookup wrote it.

It stays quiet where it should. An address you type with no number in it gets no such line - there is no house to have missed. And a postcode is not mistaken for a house number, so typing 10400 does not make the app look for a house called 10400. The comparison is by whole number, never by a piece of one: the postcode 10400 can never stand in for number 23, and 123 is a different house from 23.

Nothing is blocked, and nothing was gated. The point is still placed, the price is still asked for, the trip can still be booked - only the words changed, exactly as a warning should work. And when the lookup can only reach the road, the box for **Coordinates, if you have them** is right there: paste a pair of numbers or a Google Maps link and the pin goes exactly where you put it.

Nothing to run, nothing to set up. No database step and no deploy command. **Engine v211** on your **More** screen once this is pushed.

27 Sep 2026 - engine v210, THE WINDOW THAT OPENS WHEN THE MAP CANNOT (no database step, no redeploy - one push). The pin window on the courier card.

Small, but it is the one screen you fall back to when a map will not load - so it mattered.

On a courier order there is a press under **The door the driver is sent to**. When a map is on the card, that press unlocks the pin so you can drag it. When there is no map at all - its tiles never arrived, or there is no point to drag yet - the press opens a small window instead, where you type an address, look it up, and place the pin by its numbers. That window is the only way left to pin a door on a phone that cannot hold a map, and it had two things wrong.

It opened standing on **the door you keep for that customer**, even where the door in force was the customer's own pin from the shop page. And its sentence promised to look the address up - the one thing a customer who pinned their own door has already answered. So on exactly the orders it was most likely to be used on, the window stood on the wrong point under the wrong words.

There was one more wrong word on it, found while measuring this in a browser rather than in a test. Opened standing on a door, the line above the Keep still read *No spot chosen yet* until the window's own map happened to arrive - so on a phone whose map never comes, the line and the button beside it disagreed about the same

point.

It now opens on the door that is actually in force, and says which one that is - from the first moment, map or no map. Where the door is the customer's own pin, the window says so; where their written address is all you have, it offers the lookup exactly as before. Keeping what you see writes whichever of the two you were standing on, and - like every press you make - it is recorded as your own hand, so it stays put.

The delivery run had the same shape of mistake, and it is gone. The window reached from a run's **Put it on the map** carried two sentences, but only one could ever be shown: that press is drawn only when an order has no door at all, and a customer who left a pin has a door. The unreachable sentence is removed; what you read there is unchanged.

Nothing to run, nothing to set up. No database step and no deploy command. **Engine v210** on your **More** screen once this is pushed.

27 Sep 2026 - engine v209, THE PIN THE CUSTOMER DROPPED IS NOW THE DOOR - ALWAYS (no database step, no redeploy - one push). The dot on an order, and the door a driver is sent to.

You told me where to look, and then what was wrong: *my app, the order; the dot is in the wrong place; a different town or area.* And when I asked you which door to trust when both exist, you said: *their own pin - always.* That is what this version is.

Here is what was happening. Every order could carry two doors: the pin the customer tapped on the map when they placed the order, and a door your app had kept for that customer - filled in, on an earlier version, by looking their written address up on a map.

The door your app had kept was the one in charge, so the customer's own pin was ignored. And a map lookup for a Malaysian house number can only answer down to the street - or, worse, the wrong town altogether. That is why your dot sat somewhere else.

From now on, when a customer tapped their own door, that pin IS the door. It is the point a delivery price is asked for and the point a driver is booked to. You were standing nowhere near their house; they were standing at their own door when they tapped. Your app's own lookup is what no longer wins.

One exception, so you are never trapped. A pin you placed or dragged yourself still wins, and stays exactly where you put it. If you had already corrected a doorstep by hand, this version will not move it back.

And the card tells you which door it is using. Under **The door the driver is sent to**, the line now reads either "*<the address> - the door you keep for <name>*" or "*<the address> - <name>'s own pin from the shop page. This is the door the driver is sent to.*" When the two doors disagree, one press underneath switches to the other one - whichever of the two is not in use - so you can always overrule this with one tap.

This also applies to the delivery run and to the price. Pressing **Get a delivery price**, or **Price this run with ...**, prices the door on the screen, and a booked trip sends the driver there. Nothing to set up: an order whose customer left no pin simply behaves as it always has.

Nothing to run, nothing to set up. No database step and no deploy command. The engine on your **More** screen reads **Engine v209** once this is pushed.

27 Sep 2026 - engine v208, THE DOT NO LONGER SLIDES OFF THE CUSTOMER'S OWN DOOR WHEN YOU ASK FOR A PRICE (no database step, no redeploy - one push). The pin on the courier card.

Your report, the fourth time: the pin still wrong. When you told me where, it was **the order in your app**, and what was wrong was **the dot is in the wrong place** - not its name. This is that.

The courier card opens with the dot on the customer's *own* pin - the spot they tapped on the map when they placed the order. Pressing **Get a delivery price** then looked the written address up and moved the dot onto whatever the map's lookup found for it. But a lookup, for a Malaysian house number, can only ever answer down to the street - it has no house numbers. So the dot slid off the customer's own door and out onto the road, and the price that came back was a price for the road.

A customer who pinned their own door has already answered where the door is. Their pin is now the door, and asking for a price prices the door on screen and spends no lookup at all. Only a customer who left no pin of their own has their address looked up - and then, as at v207, the point is the map's and the words are the address on the order.

The same is fixed on the delivery run. Pressing **Price this run with ...** no longer looks up every customer in the list: each customer who left a pin of their own is priced at their own pin, and only the others are looked up. On that screen too, a looked-up door is now named with the order's address instead of the map's row - a fix v207 had missed on this one screen.

Nothing to run, nothing to set up. No database step and no deploy command. The engine on your **More** screen reads **Engine v208** once this is pushed.

27 Sep 2026 - engine v207, THE DOOR IS NOW CALLED BY ITS ADDRESS, NOT BY HALF AN ADDRESS FROM THE MAP (no database step, no redeploy - one push). The pin on the courier card.

Your report, three times: the pin still wrong. Read on the card, the one door was getting *two names*, and they did not agree:

12 Jalan Bunga, 10450 Penang - the door you keep for Mei Ling: Taman Sri Nibong, George Town.

The first is the address you and the customer both use. The second was the name the map's lookup had left on that door the first time it was looked up - and a lookup's answer is only a row off the map: a street and a town, no house number. So one door was being called by its address and by half of a different one.

A door is now named with the address it was looked up for. When your app looks the address up on its way to a delivery price, it keeps the *point* the map found and the *words* from the order - your address - instead of the map's row. The same is true when you press "Look it up" in the pin window and keep the spot: the pin moves to the map's point, and the door is named with the address you have. Where an order carries no address at all, the words the door already has still stand - it is never left as two bare numbers.

Old orders are fixed too, without touching a single record of yours: because the card now says the door with the address on the order, a door that was saved with a map's row for a name reads correctly from the moment this version is on your phone.

Nothing to run, nothing to set up. No database step and no deploy command. The engine on your **More** screen reads **Engine v207** once this is pushed.

27 Sep 2026 - engine v206, THE MAILING ADDRESS NO LONGER ASKS TO BE TYPED ON EVERY PHONE (no database step, no redeploy - one push). One sentence on More > Settings was telling you something that stopped being true at v200.

The card under **Mailing labels (courier)** used to end with "Type it on each phone you print labels from." Since v200 the mailing address is one of the settings that travel, so it is typed once and your other phones have it the next time they open - which is what the manual has said all along. The card now says so: *Saved with your settings, so it travels to your other phones - type it once.*

Nothing else changed, and nothing you have typed is touched. This is a sentence, not a behaviour: the address has been travelling since v200 and still does.

Nothing to run, nothing to set up. No database step and no deploy command. The engine on your app's **More** screen will read 206 once the phone has the new build.

27 Sep 2026 - engine v205, THE PIN IS NAMED WITH THE CUSTOMER'S OWN ADDRESS (no database step, no redeploy - one push). Version 204 gave a dropped pin the wording of the suggestion row the customer tapped. That was the wrong half: those rows are pieces of an address, not addresses. From now the pin is named with the address the customer *typed themselves*, and the row's wording is used for nothing except the row on their own screen.

What you said, and it settled it. "the customer know their address well, when i tap the address the address is

not a complete one, if it is plaste into the address line, it will contaminate the customer keyin address." You are right, and it is worth saying exactly why: a suggestion row reads like "Taman Sri Nibong, George Town" - *a street and a town, and no house number*. It is a good enough answer to "roughly where is this" and a hopeless answer to "which door". Putting it on the order was giving a fragment the standing of a complete address.

What now happens.

- **The pin travels named with the customer's own typed address**, word for word as they wrote it. So the address on the order and the name of the pin are the same string, and there is nothing left that can say a different place. **The suggestion row's wording never leaves the shop page**. It is printed in the list they tap and nowhere else - not in the pin, not on the order, not on your screen. **A tap on a suggestion still never writes into their address box**. Their words stay exactly as they typed them. A suggestion moves the pin; it does not edit anybody's address. **On your card the address is printed once, not twice**. Because the pin now carries the same words as the address, the card's line about their own pin adds nothing when the two match - and it still says the thing that matters, that this is *their* pin and not yet the door the driver is sent to. A pin from an older order, which does carry a name of its own, still shows that name, so you never lose the ability to see where a point was meant to be. **A pin with no address to name it goes as the bare point it is**. A customer who pins without writing anything has nothing to name the spot with, and the order stays exactly what it has always been.

One consequence worth knowing. The pin's words are the customer's, taken at the moment they ordered, so they do not follow a later edit you make to that order's address. The pin stays where they drew it and keeps saying what they called it, which is the honest thing for it to do.

Nothing to run, nothing to set up. No database step and no deploy command: the change is in the shop page and the app itself, and it arrives with the push. The engine on your app's More screen will read 205 once the phone has the new build.

27 Sep 2026 - engine v204, THE PIN AND THE ADDRESS CAN NO LONGER DISAGREE (no database step, no redeploy - one push). A customer's dropped pin used to arrive as two bare numbers with nothing tying it to the address written beside it, so the two could say different places and your screen had no way to show you. That is fixed, and the one sequence that actually produced it - taking a door from the suggestion list and then editing the address - now takes the pin back with the address instead of leaving it behind.

What you were seeing. "When the pin arrive at backoffice, it did not tally": the order carried an address that said one place and a pin that sat somewhere else entirely. Both were shown, and they contradicted each other.

Two things were wrong, and only one of them was visible.

The first: **the words the customer tapped were being thrown away**. When they tap one of the address suggestions, the row they tapped says which door it is - "Taman Sri Nibong, George Town" - and that wording was dropped one line before the order was sent. What reached your app was "5.33250, 100.30204" and nothing else. A pin of two numbers cannot be checked against an address by eye, so the app had no way to tell you the pin and the address were in different places. A pin taken from the list now travels with the words it was found for, so the two can be read side by side.

The second, and this is the one that produced the order you saw: **the pin was never told that the address had changed**. A customer types "Taman Sri Nibong, Penang", taps the suggestion, and then edits the box to "Bayan Lepas, Penang". Measured on the live page, that gave an order with Bayan Lepas written on it and a pin **5.8 km away** at Taman Sri Nibong. Two answers to one question, written at two different moments, with nothing joining them. Nothing on your screen could show it, because the pin had no words to compare.

What now happens.

- **Changing the address takes back a pin that came from the suggestion list**, and says so in a line above the list in the customer's own language: the words they tapped no longer go with what is in the box, so tap a suggestion above or set the pin again. The list they were reading stays where it is until its replacement arrives - it is not yanked out from under their thumb. **A pin they placed with their own hand is left alone**. Dragged on the map, or taken from "Use my location": that is a door they chose themselves and not an answer to what they

typed, and no edit to the address box touches it. **A suggestion row tapped a moment too late is refused out loud.** The list is deliberately left up while they keep typing, so for the length of the lookup's own pause a row drawn for the old address is still tappable. That tap is now refused with its own sentence, and the rows it came from are taken away with it. A tap that does nothing at all is the dead-control fault this shop has a standing rule about, so it always says why.

One more thing, on the customer's confirmation card. When no product has a change/cancel window, the card has one line fewer - and that missing line was printing as the word **null** between two sentences. It prints as nothing now. This was reachable on your live storefront: none of your products sets a window, so it was on every order's receipt. Your cancellation wording in More -> Settings is a different thing and is untouched.

Nothing to run, nothing to set up. No database step and no deploy command: the change is in the shop page and the app itself, and it arrives with the push. The engine on your app's More screen will read 204 once the phone has the new build.

26 Sep 2026 - engine v203, AN ADDRESS WITH A UNIT NUMBER IS NOW FOUND (no database step, no redeploy - one push). Some shop addresses the lookup used to give up on now bring the map to the door. There is nothing for you to run and nothing to set up: this is a change to the shop's own page, and it is live the moment you push.

What was going wrong, exactly. The lookup matches the words a customer typed against the words the map service holds, and a Malaysian address carries two kinds of word that service does not have.

The first is the part that says WHICH door: "No 5", "Blk A", "Lot 1234", "Tkt 3", "Mukim 12" - unit numbers, block letters, lot numbers, floors and sub-districts. The map service holds STREETS and BUILDINGS, and those are neither. Worse than being no help, they crowd out the street that IS in the map, so a perfectly good address came back empty. This was measured on real Penang addresses rather than reasoned about, and every one of these four pairs repeated every single time:

Taman Sri Nibong, 11900 Bayan Lepas, Penang found 1 door 12-3-4 Blk A, Taman Sri Nibong, 11900 Bayan Lepas found nothing

Pangsapuri Sri Indah, Penang found 5 doors Blk 12-3-4, Pangsapuri Sri Indah, Penang found nothing

Lorong Seri Nibong 3, 11900 Bayan Lepas found 4 doors No 5 Lorong Seri Nibong 3, 11900 Bayan Lepas found nothing

Jalan Teluk Kumbar, Penang found 2 doors Lot 1234, Mukim 12, Jalan Teluk Kumbar, Penang found nothing

The second is shorthand. "Rd" for Road, "Jln" for Jalan, "Tmn" for Taman. The map service holds one spelling and not the other, so the same street written the short way came back empty too.

What now happens. When the lookup finds nothing, it asks once more - with the unit number and the block letter taken out and the shorthand spelled out - and only when it found nothing. Checked against the live map services, those four addresses now come back with one, five, four and two doors instead of an empty list.

The customer's own words are always asked first, and this is the part that matters. The exact address a customer types is the first thing the page sends, every time. So an address that already worked is not slowed by so much as a moment and cannot come back with a different door. The forgiving wording is only ever spent on a question that has already come back with nothing, which is also why taking a word out can never do harm: it is only ever tried where the full address has already failed.

What it costs. A lookup that finds nothing now takes about a second longer before it settles, because it is asking a second question. An address that works is not affected at all. It also means a customer who types an address the map truly does not hold waits a moment longer to be told so - which is the honest price of the four addresses above now working.

No deploy and no database step. Unlike v202, nothing here lives on your Supabase. The change is entirely in the shop's page, so if you pushed v202 and ran its deploy command, you are already done and this simply arrives with the push.

26 Sep 2026 - engine v202, **THE SHOP NOW COMES TO THE ADDRESS THE CUSTOMER TYPES** (no database step - but **ONE** deploy command, and the feature does nothing until you run it - see the bottom of this entry). A customer on the shop page used to type their address into a plain box and then press "Pin on the map", which opened the map on the whole of Penang. They then had to drag and zoom around the island looking for their own street. You asked whether the map could bring them to their address instead, the way a ride app does. It now does.

What they see. As they type their address, a short list appears under the box with the doors they might mean - the street on one line, the town and postcode with it. They tap the right one, and the map opens already standing on that door, close enough to pick out the house, **with the pin already on it**. Then they drag or tap to fine-tune it, exactly as before, and press Keep this spot. The list holds at most five doors, only in Malaysia, and only doors it has a real point for. Once they have tapped one, the line above the list steps aside - the screen must not go on telling them to do something they have just done. The doors themselves stay where they are, so changing their mind is one more tap rather than typing the street again.

Nothing is asked until they have actually typed an address, and nothing is asked mid-word. The page waits until they pause, and it only bothers with a question worth asking - a word or two, or a postcode on its own, gets no list. It also never blocks anything: if nothing matches, or if the service having a bad day, one short sentence appears in their own language, and the map button is exactly where it was. A customer who ignores the list completely can still pin by hand the way they always could. That has not changed at all.

Their typed address now leaves the phone one more time - as your bakery, not as them. This was your own decision and it is worth stating. The address goes to your Supabase, which asks the map service. So the map service sees your bakery's name and your server, not the customer and not their phone. That is the same posture your own app already takes when it looks an address up, and it is the reason the shop could not simply ask the map service from the customer's own browser. Nothing is stored anywhere: the service is asked a question and gives an answer.

One thing to do before this works, and it is the only step. The shop page is a plain website, so the lookup has to live somewhere that can hold your bakery's name - a small function on your Supabase. Until you put it there, the list says the lookup is unavailable and the customer pins by hand, which is exactly what they did before this version. Push alone changes the shop's screen; it does not make the list work. The command is repeated at the bottom of this entry, and it is one line.

Nothing about the pin's standing has changed. A pin the customer drops - from the list or by hand - is still a SUGGESTION and never a fact. Nothing is priced from it, nothing is booked from it, and no driver is ever told it. It is still offered to you on the order's delivery-price card and on the delivery run, and it still only becomes their doorstep when you press **Use the customer's pin**. A suggestion is a suggestion however the customer arrived at it.

And the door you keep for them still wins. A lookup the shop did has no say in what you have saved against a customer. Where you already keep a door for somebody, that is the door the driver is sent to, and their new suggestion is offered beside it as it always was.

The one step, written out. In the Terminal, from the folder that holds the app:

```
supabase functions deploy shop-geocode --project-ref hzpyblqygnntixkijeem
```

The --project-ref part matters - without it the tool asks you to pick a project from a list, and that list also holds an unrelated one. There is nothing to configure and no key to paste: this function holds no secrets and can reach nothing but the two public map services. Run the line once and the list works from then on.

One thing was wrong in this version, and it is fixed. No second deploy. The first build of this feature still said "the address lookup isn't available right now" after the function was deployed and working. The reason was small and entirely in the shop's own page: the lookup was sending its key twice, and one of those two ways is a header the function was not allowed to accept. A browser refuses that before the request ever leaves, which is why the screen said the service was down when it was in fact answering perfectly. The fix is one line in the shop's page and nothing on the server, so the deploy above is still the only command - if you have already run

it, you are done, and the lookup now works. It was found by asking the function directly, not by reading the code: the same call with the extra header failed and without it returned four real Penang streets.

26 Sep 2026 - engine v201, THE COURIER BOX NOW SHOWS THE DOOR (no database step, no redeploy - one push). Open an order's Note / tracking box and, for an order that goes by courier, the top of it now tells you which door the driver would be sent to: the customer, the address, and a small map with the pin on it. You can look at it without pressing anything, and you can move it if it is not the door. That is the whole of it - and it is the box you asked about, the one that used to send you to Edit for the address, which meant the one place you book a trip was the one place you could not check where the van was going.

Why this was worth doing, in one line. A courier is not given an address, it is given a point. The pin is the stop the van drives to. A pin that landed on the wrong estate is a delivery you pay for and then chase, and until now the only place you could see it was a map screen that threw away whatever you had half-typed in this box. So the door moved to the top of the box you were already in.

It is a picture until you say otherwise. The map opens read-only, which is what you asked for: a drag does nothing, a tap does nothing, and pinching does nothing. Press **Move this pin** and that same map wakes up - drag the pin, and the moment you let go it is saved against that customer, with nothing you had typed beside it thrown away. A second press (**Done moving**) puts it back to a map you cannot nudge by accident while reaching past it. If you do not want to drag a small pin with a thumb, a tap anywhere on the unlocked map drops it there instead.

If the customer pinned their own door on the shop page, that is what you see, and the box says so. It reads as their own pin, not as the door the driver is sent to - because it is not one until you make it one. Press Move and it becomes the door you keep for them, exactly as it always worked, and from then on the line reads as the door you keep for them rather than as their suggestion. Nothing they dropped reaches a driver unaccepted.

Moving the pin clears the prices that were on the screen, and tells you why. A price belongs to the door it was asked for, so a price quoted for the old one would have been a lie in the box you book from. The box says, in the line where prices appear, that the door moved and to ask again for this spot. It does not quietly re-ask on its own - asking is eight requests, and that has always stayed your tap.

And it still works when the map does not. On a phone with one bar of signal, or a day when the map service itself is having trouble, you get the words instead of the picture: the door named as before, the pin's own coordinates so you can still check them, and a line saying the map is not available right now. There is always a way to put a doorstep on the map even then, because the pin screen can take coordinates as well as an address.

An order with nothing pinned shows its address and offers to pin one. No map, because a map with no pin on it is a picture of nothing, and it would spend a fifth of the box saying so. And an order that is collected rather than delivered gets no door block at all - there is no door to check.

One thing to know about the Edit form. It gains the same block, so the customer end has one kind of control rather than two. There it sits directly under the address row rather than at the very top, because the address on that screen is a draft you are part-way through typing and is not the address the driver is sent to - so the check you can trust sits beside the real one rather than floating above a box you have not saved yet.

Nothing here can stop you taking a delivery, and that is deliberate. An order with no pin, or a pin you have never looked at, prices and books exactly as it did before. This is a way to look and a way to fix; it is not a step you have to pass.

The map stays where you leave it. If you drag the map across to look at the next street, nothing on the card will pull it back. And once you move the pin, it lands where you let go of it and stays there, rather than sliding back to the middle of the box. This was wrong in the first build of this version and was caught by dragging the map on a real render and watching the box snap back the moment anything redrew. It is the same rule the planner already follows: a redraw must never move you. The one thing that *does* move the map is the address

lookup finding a door the map has not shown yet, because that is a door you would otherwise not be looking at.

Seven new tests, and ten faults put back to prove they work. The faults included the block being built inside the price fold instead of at the top of the box (where it would be invisible until you asked for a price - the exact fault you reported), a dragged pin never being saved, a dragged pin leaving the old door's prices on the screen, the map not following the pin the panel looks up on its way to a price, the block being drawn for collect orders, a missing map being left as a blank space with no word about why, the map opening already draggable, a locked map swallowing the card's scroll on a phone, the map never being measured once the card had settled, and every redraw pulling the map back to the pin. All ten were caught by the tests that name them, and each was injected, seen to fail, and removed again with the files proved byte-identical. The whole suite is 1711 tests.

26 Sep 2026 - engine v200, THE SETTINGS THAT STAYED BEHIND NOW TRAVEL (no database step, no redeploy - one push). Five things you set on one phone never reached the other: the two lists on Money - what an expense was for, and the ways you get paid - your mailing address, and the planner's people, meaning the names you give the workers and which of them the day calls. Set them on your phone, open the app on the other, and the built-in names were back, the address box was empty, and your workers were numbers again. Nothing you set was ever lost; it saved properly on the phone that set it and had no way across. All five now travel, so whichever phone you pick up has them.

Here is what was actually wrong, and it is not what it looks like. Your settings go up to the cloud as one block, and the app only puts a thing into that block if it knows to look for it. Five of them were never on that list. So they were saved on the phone, left out of everything sent up, and the other phone - never told - kept whatever it already had. It was never a failed upload and never anything you did; the app simply was not carrying them.

Set any of the five on either phone and the other phone has it the next time it opens. The two Money lists travel as the whole list, not blended together - editing them replaces them, exactly the way the rest of Settings behaves, so the two phones cannot end up with a half-and-half version of your chart. Your mailing address reaches the other phone as the FROM block on your labels. The planner's names are the same names on every plan you make, because the app keeps them by the worker, not by the plan - name somebody once and every plan reads as people.

Emptying one is an answer and now travels as one. If you clear your chart back to nothing, that means "back to the built-in names", and that meaning now reaches the other phone instead of the old list coming straight back at you. The same goes for an address you clear, and for a worker's name you delete. Before, a list you had cleared and a list you had never touched looked identical on the way up, so the phone that still had your old list would have handed it back - the app could not tell "she cleared this" from "she has no opinion here". It can now.

And it is safe to open the app on a phone you have not set up. A phone that has never touched one of these can never delete it from the other phone: staying quiet about a setting is not an instruction to remove it. If the phone that has never seen your chart opens first, it receives yours rather than pushing its own built-in one over you. This is the same rule that stopped a newly set-up phone from emptying your saved days, and it now covers these five as well.

Still on the phone it is on, on purpose, and worth knowing. Your sign-in details and your app lock never travel - those belong to the phone in your hand. The short list of occasion names Home remembers stays put too, since it is a convenience for the phone you name them on; if you would rather it followed you as well, say so and it is a small change.

Twenty-one new tests, and ten faults put back to prove they work. The faults included each list being left out of what goes up, each one being made deletable by a phone that never set it, a cleared address going out as silence instead of as an empty one, and a name list arriving as something that would have stopped the planner drawing. All ten were caught by the tests that name them; a test that stays green while the thing it tests is broken is worth nothing, so each one was injected, seen to fail, and removed. The whole suite is 1704 tests.

25 Sep 2026 - engine v199, THE TOTAL NOW SHOWS ITS WORKING (no database step, no redeploy - one push). On the WhatsApp message at every stage of an order, and on the Track page your customer

opens, the money is set out as a small addition rather than one figure on its own: what they ordered comes to this, plus the courier's charge if there is one, and it all reaches the total - with the total set apart by a blank line and shown bold, so the eye lands on the number that matters. Every message that names money now names it the same way, and every track card says the same figures in the same order as the message.

There were two different things wrong, and only one of them was the emphasis you asked about. On a message, the total was simply the next line down from the working, in the same weight as everything else, so it read as one more line of the paragraph rather than the answer. And on an order that carried no courier charge at all there was no working to read in the first place: the "Items:" line is followed by what the customer ordered - two loaves, one sandwich - and not by what it costs. So a customer with a plain self-collect order had a total arrive from nowhere and nothing to check it against. That is the one you spotted, and it was the more serious of the two.

What every message now reads, from the confirmation onwards. The goods, then what they come to, then the courier's charge if the customer is bearing one, then the total. The charge keeps the two wordings it has always had - the plain charge, and the charge the courier collects at the door, which says so and stays out of the total because the customer hands it over in person. A self-collect order with no charge skips straight from what the goods come to down to the total, so the two figures sit one above the other and agree.

One of the three messages was different from the other two, and it is now the same. The message you send when a courier order is on its way used to leave the total out completely unless the order carried a charge - so the one message a customer gets when their order is in a van was the one message that never said what they had paid, while the payment reminder for the same order always quoted it. It now carries the same lines as the rest. Nothing about where the order is, or the tracking number, has changed.

The Track page draws the same three lines, in the same order. Before, it named the courier's charge and then ran the goods and the total together onto one line joined by a dash, so the charge appeared with no clue how it related to the figure after it. It now lists what they ordered, what that comes to, the charge if there is one, and the total underneath - the same reading order as the message, so the two cannot be read as two different sums. It also used to write the charge as "RM8.00" with no space, unlike every other figure on the site; it now matches.

Bold is done the way each place can do it. WhatsApp has no bold setting in a typed message, so the total is wrapped in a pair of asterisks, which is the one way to make a line stand out there. The Track page is a web page, so it simply styles the line instead - you will never see an asterisk on the shop side, because that would be the WhatsApp message leaking onto the page. Both were worth saying, because they are the same intention carried out two different ways.

The figures are worked out from what the app already knows, so no new column was needed. The total is the one figure the app publishes to the customer's page; the courier's charge is published beside it, and it is inside that total unless the courier is collecting it at the door. So the Track page takes the charge back off to show what the goods came to - which is why a doorstep charge does not move that figure at all, since it was never added to it. The page also reads the currency off the published total rather than assuming ringgit, so a price written in another currency cannot end up with its working-out written in RM.

Sixteen faults put back, all sixteen caught. The faults included the total quietly losing its asterisks, the blank line going missing above it, the working disappearing altogether so the old "total from nowhere" came back, and - the one most worth having - the Track page taking the courier's charge off an order the courier is collecting at the door, which would have shown a customer a subtotal of RM22 on an order they are being asked RM30 for. Also among them: the shipped message going back to naming a total only when there was a charge, and one of the translations losing the placeholder that carries its figure, which would print "Items total:" with no amount beside it. Every file was restored byte-identically and proved by sha256 rather than by counting lines. The suite is 1670 passing with none failing.

No database step, and no redeploy. This one is all in the app itself, so the push is the whole of it - nothing to

run in Supabase, and the courier function is not touched. The v196 and v198 lookup changes still need the redeploy they were waiting for; this one does not add to that. After pushing, close and reopen the app on your phone.

25 Sep 2026 - engine v198, THE ADDRESS LOOKUP OFFERS ITS OTHER MATCHES (no database step; the courier function must be redeployed). When you press "Look it up" on the pin card, the card now lists the other places the lookup found, underneath the button, so the right door is a press rather than a drag. The first match still lands on the map by itself and nothing else about the flow changes - the list is there for one thing only, saying "not that one".

Every answer but the first was being thrown away, on every lookup. The address services answer with a handful of places, nearest-sounding first, and the app used to keep only the first one and drop the rest on the floor. So when the top answer was the wrong end of the street, the pin landed there, and the only way to fix it was to notice and drag it. The other answers were arriving all along.

What you do now, and it is one press more than nothing. Type or paste the address, press "Look it up", and the pin lands on the best match exactly as it always did. If the lookup found more than one place, the line under the button says so - "Found: 12 Jalan Bunga - and 3 more below" - and the other places are listed under it, each showing the street on one line and the town and postcode underneath. Press the one you want and the pin goes there.

The row the pin is on wears a tick, and the tick follows the pin rather than your press. It is not a note that you tapped a row; it is read back from where the pin actually is. So if you tap a row and then drag the pin, or paste a pair of coordinates, the tick moves with it and stops claiming a street the pin is not on. The list stays open after you press a row, so you can try one, watch the map move, and try another before you settle.

One match draws no list. A single answer is not a choice, and the line under the button already names it, so nothing is drawn under it. This is also what keeps the card behaving exactly as it does today until the courier function is redeployed.

A new lookup clears the old rows first, and a lookup that finds nothing takes them away with it. The previous answer is never left sitting under a new question.

Why the same free map service, and not Google. You asked about Google's address service, and I read its terms before building anything. It requires the results to be shown on a Google map and it forbids keeping them - and this app keeps your customers' doors for good, which is exactly what the pin card is for. The two cannot be made to agree by any amount of code, so the answer was to ask the service already in use for all five of its answers instead of one. It is built from the same open map data, it covers Malaysia, and it needs no account, no key and no card.

The suggestions arrive when you press, and not while you type, and that is deliberate. This is not the typing-ahead suggestion box you see on shop websites. A suggestion per keystroke would put several requests per word onto a free service that asks to be used lightly, and its own terms say heavy use may be throttled or refused. One press is one request - the same number of requests the app made before this change.

A quiet detail, put where a test can see it. Long addresses are split into the two lines of a row, and how to split them is decided in the app's own rules file rather than inside the card, so the split that cannot be got wrong is the one that gets tested. The honest limit: an address from the second service can read as a bare house number on the top line, with the street underneath. The coordinates are unaffected, and the list is a correction rather than the thing that places the pin.

Ten faults put back, all ten caught. The faults included pressing a row and the pin landing on the first match anyway, a single match drawing a list of one, and the old rows being left behind under a new lookup. Every file was restored byte-identically and proved by sha256 rather than by counting lines, because a one-letter difference has hidden behind a matching length here before. The suite is 1658 passing with none failing.

This one needs the courier function redeployed. The list is assembled on your server, so pushing the app's files is not enough by itself - the deploy step is what puts it live, and until then the card behaves exactly as it does today. One redeploy settles this and the v196 lookup fix together. After that, close and reopen the app on

your phone.

No database step. Nothing stored is added, changed or moved.

25 Sep 2026 - engine v197, THE CUSTOMER CAN PIN THEIR OWN DOOR (no database step). Your customer, on the shop page, can now drop a pin on the door the courier should stop at, and you are offered it where you already look at a doorstep: on the order's delivery-price section, and on the delivery run under that customer's row. Their pin is a SUGGESTION and never a fact - nothing is priced, booked or sent to a driver from a pin you have not taken up yourself. An order from a customer who pins nothing is posted exactly as it always was, byte for byte.

The two ways a customer pins, because one is not enough. A customer standing at their own door presses "Use my location" and the phone hands the pin over. A customer ordering from somewhere else - the office, the bus, their mother's house - presses "Pin on the map" and drags a pin to their door, which is how everybody already drops a pin in a ride app. Where they order from has nothing to do with where the focaccia goes, which is the whole reason both exist.

Your three answers are the feature. Offered every time, even when you already keep a door for that customer: the offer is not silenced by your own door, only the words change - "pinned a different spot this time, the doorstep you keep for them is untouched until you take this one", and the press says "Use the customer's pin instead". A courier customer who pins nothing sends as it always has. And it appears in both places, the order card and the delivery run, drawn the same way in both so the same press means the same thing.

Taking it up is the only thing that makes it a door. Pressing "Use the customer's pin" writes it with the same single writer your own map pin already uses, so a pin you accepted and a pin you placed by hand are the same record on the same customer. That is also what makes it stop being offered - no "dismissed" flag is stored anywhere, because accepting is what ends it. And because a delivery run's row is one big label you tap to tick a customer, the offer sits beside the row rather than inside it: a press in there would have ticked the customer off the run instead of pinning their door.

Where to pin, for a block or a condo. The shop now says it in words, in all three languages: drop the pin where the courier can stop - the guard house or your block's entrance - and put the block and unit number in the address box above. A pin is two numbers and cannot hold a unit number, and a pin dropped inside a tower is exactly where a phone's own position goes vague. Lalamove's own Malaysian guidance says the same thing: pin the lobby, the main entrance or a landmark a driver can actually stop at, and sort a wrong pin out by talking to the driver. That is also why the typed address still travels with every order, with or without a pin.

A vague position is kept and said, never refused. If a phone reports a position that is only accurate to more than about 150 metres, the customer is told the number and pointed at the map to move it - rather than being blocked from ordering. Your own rule, applied here: the shop never stands between you and a sale, and you reconfirm every pin anyway.

Every way the location press can fail has its own sentence - a browser with no location at all, a permission refused, a position unavailable, a wait that timed out, and a map that could not load. None of them blocks the order.

The map is only fetched when somebody presses for it. A customer standing at their own front door never loads it and never pays for it, and the location press works with no map at all.

The suite is 1649 passing with none failing, and nothing of yours was written. The new coverage drives the real customer controls: a stubbed phone position, a press, and the posted order read back. One finding worth naming was caught by these tests rather than by you - a half-written pin would otherwise have become a real point on the equator, because a missing latitude reads as zero to the computer. No database step and no redeploy: the pin rides inside the order the shop already posts.

One honest limit. Your phone's own permission sheet is something only a real phone can show, so "Use my location" is proved by tests rather than by me driving it. Please try it once on your own phone - stand somewhere you know, press it, and check the words that come back.

25 Sep 2026 - engine v196, THE ADDRESS LOOKUP ASKS A SECOND SERVICE (no database step; the courier function must be redeployed). You pressed "Look it up" on the pickup-pin card and read "The address lookup service did not answer. Put the pin on the map instead." - and then read it again, a second time. The lookup now asks a second address service first, and that one answers: typing an address fills the pin in again. Nothing about your day, your chart, your orders or your saved data was touched, and putting the pin on the map by hand still works exactly as it did - it is still the thing the pin is really for.

What you read, and what it proved. That sentence is written in one place, and it appears only when the address service has answered and refused. So the news was not that your phone could not reach it, and not that the service was down: it was that the service looked at the question and turned it away. Three other sentences cover the unreachable, the too-slow and the not-found cases, and you would have read one of those instead of this one.

The service itself is healthy, and I checked rather than assumed. Asked for the same kind of request your app makes, it answered with a correct Penang result - the right street, the right town, the right postcode, first time. So the address was not the trouble, and neither was the question or the way the question is written.

Why it turned the app away. That service asks every caller to say who it is, and it turns away callers it cannot account for. The app does say who it is, and says it the way that service asks - but the request leaves through a shared machine, and whether the name survives the trip is not something the app can see. What settled it was not the wording but the repetition: an ordinary address, refused identically twice, while the same service happily answers for a different caller (me). That is a house rule being applied rather than a fault to retry, so the answer was not to ask harder but to ask somebody else as well.

The fix: a second service is asked first. The new one is built from the same open map data, covers Malaysia, needs no account and no key, and asks nothing of the caller beyond the question itself. The original service is kept and is now the second ask rather than the only one. Whichever answers first wins, so in ordinary use only one service ever sees a customer's address - and an address still never leaves your own server for the browser, which is how this was built and how it stays.

The old service was kept rather than thrown out. It is the better-resourced index of the two, and keeping it costs nothing: if whatever turned it away is ever lifted, it is already in the queue. Losing its first place is the only thing it lost.

The four ways a lookup can fail are now told apart. Not found, refused, too slow, and unreachable each get their own words, so the sentence you read describes the problem you actually have - a house in a new estate that the map simply does not carry is not the same news as a service that refused, and neither should wear the other's words.

A quiet trap in the new service, caught by a test rather than by you. The new service writes a point's two numbers in the opposite order to every other point in this app. Read in the order it writes them, a Penang street becomes a latitude of 100 degrees, which is not a place at all - and a pin that silently never appears is a fault nobody notices until a driver is sent somewhere wrong. A test now names that order explicitly.

Eleven new tests, six faults put back, all six caught. The faults included the original fault itself - a refusal from the first service ending the whole lookup - and every one of them failed a test that named it. Every file was restored byte-identically and proved by sha256 rather than by counting lines, because a one-letter difference has hidden behind a matching length here before. The suite is 1618 passing with none failing.

This one needs the courier function redeployed. Unlike the app's own screens, this fix lives on your server, so pushing the files is not enough by itself - the deploy step is what puts it live. After that, close and reopen the app on your phone before trying the address again.

No database step. Nothing stored is added, changed or moved.

25 Sep 2026 - engine v195, THE PIN CARD DRAWS ITSELF (no database step). The screen that puts your bakery's door on the map was not drawing a card. It was printing the words "[object HTMLParagraphElement],[object HTMLDivElement],[object HTMLDivElement]..." - which is what you

saw on your phone, and it is what the card had become. Every field, button and map it was supposed to hold had been built and then thrown away, so there was nothing on it to press and no way to save the pin. It draws now. Nothing about your day, your chart, your orders or your saved data was touched, and this was the same day the courier card above it began reading "Lalamove: sandbox", which is that connection working.

What you saw, and what it was. You opened Settings, pressed "Put the pickup pin on the map", and instead of a card you got a line of bracketed text. Those brackets are how JavaScript says "here is an object" when something asks it to write an object out as words. Each one was one part of the card that never reached the screen: a line of explanation, the box for your address, the map itself, the note under the map, the box for coordinates, the line reading where the pin sits, and the row with the press that keeps it. Seven parts, all of them lost to one line of code.

Why it happened, in one sentence. The app builds a pop-up's contents as a list, and then hands that list to a browser call that only accepts bare items, not a list - so the browser turned the whole list into a single piece of text instead of drawing the seven things inside it. This is a subtle trap rather than carelessness: the helper used everywhere else in the app happily takes a list, so the two look as though they should behave the same, and only one of them does.

This was the only card of the thirty-eight that did it. Every other pop-up in the app - the order edit, the product editor, the money journals, the day's availability - hands over one piece at a time, which is why none of them has ever shown you this. The pin card is the one that hands over a list, and it is the one that broke.

And it got that far because that screen had no test of any kind. Not a weak test, not an outdated test - none. It is the newest screen in the app and it was the only one nothing checked. The file that keeps this class of mistake from ever reaching a screen already existed, and it already covered three other screens that had printed stray words at you; it simply did not know about this one. It does now.

Where the fix sits, and it is deliberately not the pin card. The pin card could have been rewritten to hand its parts over one at a time, which would have cured the symptom and left the trap in place for the next card. Instead the pop-up helper itself was taught to accept a list as well as a single piece, so the two builders in the app now behave the same way - and a test asserts both halves: that a card handing over a list gets the list drawn, and that the thirty-seven cards handing over one piece each still draw exactly as they did.

Three faults put back, all three caught, every file restored byte-identically. The original fault was reinstated, and the test that names it failed; the list-check was removed, and cards across the app failed, proving that guard is what keeps the other thirty-seven working; and the map's own failure message was silenced, and the pin card's test failed. Restoration was proved by sha256 of each file, not by counting lines, because a one-letter difference has hidden behind a matching length before. The suite is 1607 passing with none failing.

No database step. Nothing stored is added or changed, and there is nothing to run in Supabase. This is a screen-drawing fault, so it is cured by the app itself being updated - push it, then close and reopen the app on your phone before you try the pin again.

25 Sep 2026 - engine v194, THE COURIER STOPS BLAMING THE BUILD (no database step; the courier function was redeployed). The sentence that appears on the Settings card when the courier cannot be asked was false, and it pointed you at the wrong place. It said the build has no courier called Lalamove, when the truth was that no Lalamove key had been added to the server yet. Those are two different problems that until now shared one sentence, and they no longer do. Nothing about your day, your chart, your orders or your saved data was touched.

What you read, and why it was wrong. Having just deployed the courier function, you opened More -> Settings and read this: "Lalamove could not be asked: This build has no courier called lalamove." Read plainly, that says the app is missing the courier - so the hunt begins in the build, in the deploy, in the code. There is no fault in any of them, and there never was. The server knew the real reason and had already written it to its own log, and then handed you a sentence that hid it.

Where the sentence came from. One small function answers for the courier, and it builds its configuration

through a single door. Two different things could shut that door: a request naming a courier this build does not carry, which is genuinely a fault in the app - and a Lalamove key that has not been put on the server, which is one secret's worth of setup. Both produced the same nothing, and one sentence was written for both. It was the sentence for the fault, so a missing key was reported to you as a broken build.

What it says now, and it is the exact line you will read. "Lalamove could not be asked: Lalamove's api key and secret have not been added to the server yet." The courier's name is still there, the thing that is missing is named, and the word build is gone from it, because the build is not the problem.

Two problems, two sentences, and a test that keeps them apart. The sentence about the key comes from the courier's own file, beside the other two sentences that courier can be refused by, so a second courier one day brings its own words rather than inheriting Lalamove's. And because the function that chooses between them cannot be loaded by the test suite at all - it begins with a line only Deno understands - a test now reads those two branches out of the file and requires each to keep its own words. Put the old single sentence back into either one and the suite goes red.

A guard of your own caught this change while it was being made. The first attempt wrote the missing key's name into a log line inside the dispatcher, and the app's existing rule - only the registry and the courier's own file may name the courier in live code - refused it. The guard was right and the line was changed to suit it, rather than the guard being widened to suit the line. That guard exists because a courier's name written by hand into a screen is the defect that stops a second courier ever being added cleanly.

One more thing corrected in the same file, and it cost real time today. The setup note at the top of the courier function told us to run `supabase secrets set LALAMOVE_KEY <key>`. That command is refused: the tool requires the equals sign, `LALAMOVE_KEY=<key>`, and answers "Invalid secret pair ... Must be NAME=VALUE." The note also now says to name the project when deploying, because the tool otherwise offers a list in which a wrong project can be chosen by a stray keystroke, and it records that the sandbox switch lives in the top right corner of Lalamove's own page.

No database step. Nothing stored is added or changed, and there is nothing to run in Supabase. The courier function itself was redeployed, which is the only thing that had to change for this to take effect. Suite is 1604 passing with none failing, and five faults were put back in one at a time - the collapse of the two sentences among them - each watched failing a test that names it and then restored byte-identically.

25 Sep 2026 - engine v193, THE WORKERS' LINE BECOMES A TRUE TIMELINE (no database step). The seven changes you asked for to the person's window on the Production line, and one question you asked about its countdown. A coach is now as long as the minutes your hand is needed on it, the three people's lines share one clock so they can be read against each other, the line starts at the left and the clock sweeps in to the middle before the line itself moves, a coach turns from green to light red when its moment has passed and eases rather than snapping, and the line is dragged with the right mouse button or with two fingers on a phone. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "1. The should start left hand side, not centred, this request I did in previous chat, but never implemented. 2. Make the coach width relative to its duration, and others batch duration, their labour requirement mark by hand needed, not batch arrival. 3. So the 3 person's train head, should be timed and position relatively to each other, when time start, 3 train started together. 4. The wait time a person interested is not the batch arriving, but their hand needed. 5. The coach should have 2 color only. Green and light red, the coach change to red only when their hand needed time is up, the change of color should be fade in and out, not sudden so that the user dont feel the screen is jumping here and there." And then: "and the drag, should be by right mouse button hold down". One more, asked after the first pass: "can the time show under their names, accurate to 5m 55s?"

2. A coach is as long as your hand is needed, and nothing else. Until this release every coach was the same width, and each person's line had its own private map from minutes to pixels, so a twelve-minute job and a forty-minute one looked identical. The width is now the minutes your hand is needed, multiplied by the day's own

scale - the same figure the modules' window above it is already drawn from. Measured live at the Detail scale: a hand-needed stretch of 6 minutes drew 43.2 pixels, one of 2 minutes drew 14.4 pixels, one of 18 minutes drew 129.6, and every gap between them was an exact multiple of the same figure. A one-minute fold is therefore a hairline, which is your own answer for it.

3. Three people, one clock. Because every line is now drawn from that one scale, a minute of the morning is the same point on all three, and the three trains start together. Measured live through one whole walk: all three people's rulers stood at the identical pixel at every moment - 0 at the start of the day, then 121 pixels, then 452 - and all three strips had slid by the identical amount. Before this release each line carried its own map, so the three could not be compared at all, which is the fault you named.

4. The line starts at the left, and the clock sweeps to the middle. This is your point 1, the one you had asked for in an earlier conversation and which had never been built. At the start of the day the day's own left edge is glued to the left edge of the window and the clock stands at the left, so the whole queue ahead of you can be seen without dragging anything. Measured live on a walk: with the clock 0.7 minutes in, the ruler stood at 5 pixels and the line had not moved at all. Once the clock reaches the middle it stops there and the line begins to move instead. Measured live at the same scale: the ruler pinned at 121 pixels - the middle of the 241-pixel window - while the line slid on to 181 pixels. Your own words for the reason were that at the start a person would otherwise see only one train head, with the rest hidden until the line was dragged.

5. Two colours, and they ease rather than snap. Green means nothing is outstanding: the moment has not come yet, or you have already taken it. Light red means the moment your hand is needed has passed and nobody has taken it - the same minute the bell rings on. There are exactly two, read off the drawn coaches: green as rgb(231, 242, 232) and light red as rgb(251, 233, 228). The coach carries a 0.45-second fade on its colour, so it eases from one to the other instead of jumping, which was your worry about the screen moving under you. A tap on a coach takes it and fades it back to green, and where a coach is wide enough to hold its face it also carries a tick.

6. The coach is the hand-needed section, labelled. Your points 2 and 4 are one thing seen twice, and it was already true in the model underneath: the section a coach is drawn from IS the stretch your hand is needed, taken from the same hands-window the day's own plan is built on, and never from the moment a batch arrives. What this release changes is that you can now see it, because the length of the coach is that stretch.

7. A coach too small for its words shows no words, not smaller words. Your answer was that if there is no room it should simply not be shown, because the same job is already named under the person's name. Measured live: coaches of 14.4 and 28.8 pixels showed no face at all, while one of 43.2 pixels showed its icon, its name, its clock and its batch. The line under each person's name still names what is coming, so nothing is lost. And a hairline coach is still a 36-pixel tap target, so even a very short job can be taken by hand.

8. The drag is the right mouse button, or two fingers on a phone. Your correction, taken literally. On a computer the line is dragged by holding the right mouse button, and a left press no longer drags at all - which is the point of it, because a left press is also the tap that says "I am on it". On a phone the gesture is two fingers, as you asked. Measured live on the drawn board: a left press with a hundred-pixel sideways move left the line exactly where it was, to the pixel; the same move with the right button carried it a hundred pixels; two fingers carried it forty. A drag also leaves behind a click, and that click is not read as a tap on the work - measured live, the press that ended a two-finger drag took nothing, while a still press on the same coach took it.

9. Scale, on the board's own row. You asked for it and there it is: the control row is three groups now, with the Scale step as the third, laid out like the planner's own so that the two screens behave alike. It moves through the app's six scales and stops at both ends rather than running past them. Measured live: stepping from the Standard scale to the Close one drew every coach half again as wide and the line from 2304 pixels to 3456, and pressing on at the top of the range left the button looking switched off and changed nothing at all.

10. The countdown reads to the second. Your question. The line under each person's name now gives seconds as well as minutes - measured live as "Next 22m 7s", "Next 160m 7s" and "Next 918m 7s" - and because the same instant is also printed on the link beside each coach, both are written by one shared piece of

arithmetic, so the screen cannot state one moment two ways.

11. Both windows are kept. You chose to keep the train AND the planner's own people's row underneath it, and both are on the screen. Measured live, the three windows stand in the order train, modules, people, with all three present at the phone width and at the desk width.

12. What is measured live, and at what widths. Everything above was read off the drawn screen at 375 by 812 and at 1280 by 900, including the whole sweep of a walk: the ruler's own pixel, the line's own slide, a coach's drawn width against its minutes, the three rulers standing at the same pixel, the two colours and the fade read off the drawn element, the tap target of a hairline, and the Scale step's ends. The desk width was reached by widening the window with no reload at all, and the board re-measured itself within its own second - which is how it will behave on your own desk.

13. Nothing of yours is rewritten, and one honest note about that claim. Your stored data was read before anything was measured, and every write the measurement made was captured and never allowed to land - seven of them across the whole pass - so your phone's storage holds exactly the three keys it held before and not one byte of it changed. The note: the first copy I made of your data to compare against carried a typo that I typed, a single letter in the name of a stored list, so the file-by-file comparison could not honestly be made from that first copy. The copy was taken again from the screen itself and verified, and it does match your stored data byte for byte. But this is stated plainly rather than overclaimed: the strongest form of the claim here rests on the reading taken before the measurement began and on every write having been intercepted.

14. Every rule that stands on it was proved load-bearing. Seven faults were put back into the real files, one at a time; each was watched failing a test that names it, and every file was put back byte-identically afterwards. Among them: the moment the ruler is meant to stop at the middle, the one scale that every line must share, the minute at which a coach is allowed to turn red, and the left press being allowed to drag again. The suite is 1602 passing with none failing.

15. No database step. Not one stored field is added or changed. The train is drawn from the day's own plan and the day's own scale, both of which the app already saves, so there is nothing to run in Supabase.

CHANGELOG and the guide both carry all of this, both PDFs rebuilt and read back with PyMuPDF, and the app's own More screen reads Engine v193.

25 Sep 2026 - engine v192, THE SAVING STAYS WITH YOU, AND A CUSTOMER IS CHARGED THE ORIGINAL PRICE (no database step). A correction to v191, in your own words and taken as the rule. When a customer bears the courier charge, they are charged what their own doorstep costs on its own - never a share of the one-trip fee - so the money that consolidating saves stays with you rather than quietly becoming their discount. When you bear it, the run's real fee is your own cost and is apportioned across the orders as before. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "the benefit of consolidated charges, should go to merchant, not the customer. And if the courier charges were reveal to them, it will shown as the original cost." Two sentences and one rule. v191 split the run's one fee evenly across the orders, which was right for your own books and wrong for a customer's: a customer who was quoted RM 16.25 to their own house was being asked for RM 11.00 - the saving from the shared vehicle had been handed to them without you deciding to. The rule now: a customer who pays carries their own doorstep's standalone price; the difference between that and the trip's real fee is yours. Nothing else about v191 changed.

2. Where the money lands, and which half of the screen changed. One function decides the amounts, and it now has three answers rather than two. If the customer bears the charge, each order carries its OWN doorstep's standalone cost, exactly as Lalamove quoted it - and if the number of costs does not match the number of doorsteps, the answer is no charge at all rather than a guess, because a mismatched list means the list is about a different run. If YOU bear it, the trip's one fee is apportioned across the orders as v191 did, and the customer is charged none of it - the amount then lands on your books as a "Delivery & fuel" cost and never reaches a customer's total. If nobody has answered who paid, the answer is no charge, which is the rule the charge model

has carried since v127.

3. Where those standalone costs come from, and what they cost you. They are real quotations from Lalamove, one per doorstep, asked when you press Book. There is no bulk pricing call, so two doorsteps is two requests - and Lalamove allows two requests a second, so the app spaces them and waits for every one before anything is booked. Measured live on the wire: the requests went out 602 milliseconds apart, and the confirmation appeared 633 milliseconds after the press - so the question genuinely waits rather than showing you numbers it does not have yet. If you have ALREADY pressed Compare with sending them separately, those costs are reused rather than asked for and paid for a second time; that is the same set of prices, so it is the same answer, and the screen says where they came from.

4. What the screen now says, measured on the drawn page rather than reasoned about. The line under the price list no longer talks about an even split. With the customer bearing it, it reads: "Each customer is charged what their own doorstep costs on its own, never a share of the one-trip fee - so the saving from going together stays with you. Those 2 costs come from Lalamove when you book, and the confirmation shows each of them before anything is asked for." The confirmation then does the arithmetic in front of you: "Each customer is charged what their own doorstep costs on its own, never a share of the trip: RM 13.50 - RM 16.25 - RM 29.75 in all, which is RM 7.75 more than the RM 22.00 the trip costs you - that difference stays with you." The RM 7.75 is the consolidation saving, said as a number and said as whose it is. When the same run LOSES you money - which happens when the doorsteps are close together and the trip fee is high - the sentence says so in those words rather than polishing the arithmetic into a saving that is not there: on a run built for the purpose whose standalone costs came to RM 12.00 against a RM 22.00 trip, it read "RM 12.00 in all ... RM 10.00 SHORT of the RM 22.00 ... loses you money".

5. Your own half of it, unchanged in the arithmetic and clearer in the words. With "I paid it" chosen, the line under the prices reads: "You are bearing it, so the run's own fee is your cost and the customer is charged none of it: RM 22.00 over 2 orders - RM 11.00 - RM 11.00." The confirmation says the same thing as a fact about your books: "The run's fee goes on your books as one cost of RM 22.00, over 2 orders - RM 11.00 - RM 11.00 - and the customer is charged none of it." Measured live after confirming that run: two "Delivery & fuel" rows of RM 11.00 each, the orders carrying RM 11.00 each and paid-by-you, and the customers' own track cards carrying no courier charge at all - their totals read RM 45.00, which is their items and nothing else. That is the half where the even split is correct, and it is the only half where it is.

6. The odd cents, and why the last order carries them. When a fee does not divide evenly, the apportioning puts the remainder on the LAST order rather than spreading a fraction of a cent. Measured on a fee of RM 22.00 across three orders: RM 7.33, RM 7.33 and RM 7.34, which add up to the RM 22.00 you actually paid. This matters only on your own books - it is the arithmetic that stops the three rows summing to RM 21.99 or RM 22.02 - and it is unchanged from how the charge model has always rounded.

7. What is measured live, and at what widths. Every number in sections 4 and 5 was read off the drawn screen at 375 by 812, with the world built for the purpose and taken back out afterwards. The Book press measures 110 by 36 pixels, the row 315 wide with no sideways scrolling, and the line under the price list 90 pixels tall, which fits without clipping at that width. At 1280 by 900 the two presses measure 110 by 36 and 270 by 36.

8. Nothing of yours is rewritten. The only press this release changes is the one that decides what each order's charge box is to hold, and it writes the same three fields it has written since v188 - the amount, who paid it, and whether it is collected on delivery. Measured across a whole pass - a price asked, a run priced separately, a run booked with the customer bearing it, and the same run booked again with you bearing it - the only fields that moved were the charge on those three orders and the trip itself. Your saved days and your settings were not touched.

One limit stated plainly rather than glossed: the copy of the app's stored data taken before the measurement was lost when the browser reloaded partway through, so the usual byte-for-byte comparison could not be made this time. The bakery the measurement ran in was an empty one built for the purpose and it is empty again, with nothing of yours in it at any point - but the comparison itself was not performed, and that is a weaker claim than

the one every other release in this history makes.

9. One thing this release deliberately did NOT change, and you should know it. If you book a run with the payer question left unanswered, the orders on it lose any charge you had set on them by hand. That is not new - it is how v191 behaved, because the charge is written as a set of three keys and an unanswered payer is an empty answer. It is left exactly as it was rather than quietly changed under this release's heading, and it is named here so you can decide whether you want it to behave differently.

10. Every rule that stands on it was proved load-bearing. Ten faults were put back in the real source, each watched failing a test that NAMES it and then restored byte-identically. Among them: a wrong-length list of costs accepted as a charge; a customer given the even split instead of their own doorstep; your own share repeated whole on every order instead of apportioned; the first customer's cost written onto every order; the costs re-asked for even when a comparison was already on screen; and the losing run's shortfall smoothed into a saving. One of the ten is worth singling out because only ONE test catches it: taking the spacing between the two standalone requests away fails the timing assertion and nothing else, so the rate limit is guarded by the only thing that can see it. Two tests were added to the pure charge function for the answers that had no test - a list of the wrong length, and the no-payer case. The suite is 1596 passing with none failing.

11. No database step. Not one stored field is added or changed. The charge lands in the fields the order already carried, the trip in the fields v189 already writes, and an order record syncs whole - so all of it reaches your other phone with no sync change and no migration.

CHANGELOG and the guide both carry all of this, both PDFs rebuilt and read back with PyMuPDF, and the app's own More screen reads Engine v192.

25 Sep 2026 - engine v191, ONE TRIP, SEVERAL DOORSTEPS, AND WHAT IT SAVES (no database step). The fourth and last step of the courier work you asked for, and the one that saves you money. The app now groups a day's courier orders into ONE trip, prices it against the same doorsteps sent one at a time, and shows you the difference as two numbers rather than as a claim. Booking it spends one vehicle on one journey and gives every customer on it the same share link. What your shop promises those customers becomes a WINDOW that you type, and only when you type one. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "i want to do Lalamove, API to manage courier, make it ready for other courier as well. Focus now for laalamove. And check our backoffice readiness. From my understanding lalamove support last day deliver consolidation to save on cost." v188 was the price, v189 was the booking, v190 was the customer watching the trip. This is the fourth and last, and it is the one your own hunch was about. Your consolidation belief was confirmed in v188: one trip carries two to sixteen drop-offs, charged as one base fare plus a fee for every stop after the first, which outside the Klang Valley is RM 1 on a motorcycle, RM 2 on a car, RM 5 on a van or four-by-four and RM 10 on a lorry. This release is that mechanism with a screen around it.

2. One trip, and the arithmetic you can read rather than take on trust. The run screen prices the whole run twice. Once as the one multi-stop trip it is about to become, and once as the same doorsteps sent one at a time, and then it writes the difference out as a sentence. Measured live on a two-customer run on a car: "One trip RM 22.00 - one at a time RM 44.00 - this one trip saves RM 22.00 against sending them one at a time. Each of those was asked now, for the same collection time, so they are this journey split up rather than a guess at it." That last sentence matters more than it looks: the comparison prices are REAL quotations asked at the same moment for the same collection time, not an estimate the app did on its own. The saving is therefore a number you can act on rather than a claim off a leaflet. It also costs you waiting time rather than money - there is no bulk pricing call, so a run of eight drops is still eight price requests - and the screen says so rather than making the second pass look free.

3. Your four decisions, and where each one landed. You were asked four things before any of this was built, because a single trip cannot let each customer choose their own hour, and the answers are built rather than described. **The window** - "A window I type when I book the run" - so the run screen has two boxes, opening and closing, that you fill in yourself, and nothing is promised to any customer until you do. **The fee** - "Split evenly into

each order's box" - so the trip's one fee is divided across the customers on it and each order's own charge box carries its share, which is what lets the existing charge model, the books and the customer's own total keep running exactly as they run. **What fits** - "Show me the load and let me judge" - so the screen counts the doorsteps, counts the items and names them, and leaves the judgement to you rather than refusing the run on your behalf. **The trip link** - "The same trip link for everyone" - because there is one vehicle on one journey, so every customer on it gets the one share link and there is nothing to keep apart.

4. A customer's own doorstep, and the rule that makes consolidation honest. One customer who ordered three things is ONE doorstep, not three. The screen counts stops per customer, and so does the price it asks for: measured live, a two-line order for one customer was priced as one drop, and a two-customer run was priced as two. Without that rule a customer ordering twice would have been charged as two deliveries to the same house, which is exactly the overcharge consolidation exists to remove. The row for each customer says what is being sent and where it is going, and a customer whose doorstep has never been pinned carries a press to put it on the map instead of a price it could not honestly be given.

5. What your shop promises now, and only when you have said it. A delivery window you type reaches the customer through the slots that were already there rather than through anything new: it is folded into the delivery line on their track card, and into the shipped WhatsApp message and the reminder. Measured live: the card read "Sat, 26 Sep - Courier - 1 Jalan A, 2-5 pm", and both messages read "Delivery: Sat, 26 Sep, 2-5 pm - Courier delivery" with "Track your delivery:" above the share link. Leave the two boxes empty and every customer keeps the promise they already had - the day, and no hour - so a run booked in a hurry is not a run that promises something nobody chose. And the window is checked before it is priced: a window whose end is before its start is refused in words on the spot rather than being priced and then written onto customers' cards, which it would otherwise have been, because the formatter that draws a window renders an end-before-start one perfectly happily. That is a fault this release's own test sweep found and the test now holds.

6. Nothing is booked until you press the press that books it. The run screen prices freely and books once, and the two are separate presses with a question between them. Measured live: the question read "Book the Car with Lalamove for RM 22.00? It carries ONE trip with 2 doorsteps and 6 items. Customers will be told 2-5 pm... (RM 11.00 - RM 11.00), marked paid by you." - so the number being spent, the vehicle being spent on, the load it has to carry, the promise about to be made and what each customer is about to be charged are all on the screen before anything is sent. Behind it there is exactly ONE booking call for the whole run, and the toast afterwards read "Run booked with Lalamove - 2 customers now share one trip and one link."

7. The money, and why a share and not a fee. One trip is one fee, so charging each customer the whole fee would take the cost of one van several times over. The trip's fee is divided across the customers on the run and the remainder lands on the last order rather than being lost, so the shares add up to the fee exactly. Measured live on a RM 13.50 motorcycle price for two customers, the split line read "Split evenly over the 2 orders ticked above, in that order: RM 6.75 - RM 6.75" - and on the booked car run the two orders were stamped RM 11.00 each against the RM 22.00 fee, with two RM 11.00 "Delivery & fuel" expenses on your books, one per order, paid by you. An odd cent is not invented and not dropped: on three customers the same RM 22.00 fee lands RM 7.33, RM 7.33 and RM 7.34, and there is a test that fails if the parts ever stop adding up to the whole. Answering the charge questions is optional and stays optional: a run booked with nothing said about the money records no charge at all rather than a charge of nothing, which is the difference between a customer who owes you nothing and a customer whose box says they owe you zero.

8. Three places where the honest answer is not the convenient one. A price describes one list of customers on one journey, so anything that moves either throws it away rather than leaving a stale number standing beside a new question: untick a customer after pricing and every Book press on the screen goes inert with the sentence "The list of customers has changed since this price was asked, so it describes a run you are no longer taking. Price the run again for the list above." Change the collection day and the prices go entirely. And a live trip on any order in the run stops the whole run being booked, with one sentence saying which order is already out - the same one-trip-at-a-time guard the single order card carries, at the level of the run.

9. Two faults found by reading the drawn screen rather than by reasoning. The first: a customer's tick row

was drawn as a whole pressable line, but the live area was only the words inside it - the top and bottom nine pixels of a 49-pixel line landed on the row rather than the tick and ticked nobody, which is the exact opposite of what that row's own comment promises. Measured with probes before and after, and fixed by moving the padding from the row onto the label so the whole line is the tick's. The second: ticking ONE customer left the head line reading "Who is on the run - 2 of 2" over a list with one customer on it, while the tick-everyone press corrected it. Two controls that look alike behaving differently, and the count she reads is the one that went stale. The cause was that the single row's handler repainted the load, the prices and the charge boxes but not the list the head lives in, and only the bulk press rebuilt it. Fixed so a single tick corrects the count and the bulk press WITHOUT rebuilding the rows - measured, the row she pressed is still the same node afterwards and the list does not move under her finger. Both are named rather than glossed because no assertion had ever watched either one.

10. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Your stored data was copied before the presses that write, put back afterwards and compared byte-identically, with no backup key left behind. The one thing this release adds to an order is the window you typed and the trip and share of the fee you booked - the same fields the single-order booking already writes - and an order with no run on it is drawn exactly as it was.

11. Every rule that stands on it was proved load-bearing. Thirteen faults were put back in the real source, each watched failing a test that NAMES it, and each restored byte-identically. Among them: pricing one drop per order LINE rather than per customer, charging each customer the whole fee instead of a share, writing the window onto only the first line of a customer's group, stamping the trip on only the first line, a stale price still booking, an end-before-start window slipping past its refusal, the odd cent dropped so the shares no longer add up, a zero fee recorded as a charge, and the stale head from section 9. Two of the thirteen are worth singling out because the sweep is what found the holes rather than what confirmed them: the publishing gate that decides whether a window reaches a customer was asking "can this be read" rather than "is this a window at all", so an end-before-start window published happily - the test now compares the customer's card against the card of an order with no window, which is the only comparison that catches it. And the zero fee was initially written as a literal zero rather than left absent, and nothing failed, because no test had ever booked a run with the money questions left unanswered. A test was written for it and the fault fails by name. The suite is 1587 passing with none failing.

12. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. A run's window is folded into the delivery line the customer's card already carried, the trip and the share of the fee land in fields the single-order booking already writes, and an order record syncs whole - so all of it reaches your other phone with no sync change and no migration.

13. Still to come, and only if you want it. The public callback is the one piece left unbuilt from the courier work. It would let a customer's page update by itself within seconds of a parcel being picked up, without anybody pressing anything. It is also the one part of this work that would be reachable by anybody on the internet rather than only by your own signed-in app, so it stays your decision and nothing about it is promised here.

CHANGELOG and the guide both carry all of this, both PDFs rebuilt and read back with PyMuPDF, and the app's own More screen reads Engine v191.

25 Sep 2026 - engine v190, THE DELIVERY'S PROGRESS AND ITS DRIVER REACH THE CUSTOMER, AND AN ORDER MOVES ITSELF (one database step). The third step of the courier work you asked for. A booked trip's own progress now shows on the customer's track card in the customer's own language, with the driver's name, the number plate and a press to ring them. And when the courier says the parcel is on the vehicle, the order moves itself to Collected / Shipped - a note appears on the row naming who moved it and when, and one press puts it back. Your order row is the first thing in this app that has ever changed itself, so it is built to be SEEN and UNDONE rather than to be trusted. There is one database step, and it must be run BEFORE this build is deployed or every customer's tracking page stops updating silently. Not one module, batch, start time, cycle or saved day of yours was rewritten.

1. What you asked, in your own words. "i want to do Lalamove, API to manage courier, make it ready for other

courier as well. Focus now for laalamove. And check our backoffice readiness. From my understanding lalamove support last day deliver consolidation to save on cost." v188 was the price, v189 was the booking, and this is the third of the four stages you chose. The seam and your consolidation belief are unchanged, and the fourth stage - the delivery run that saves you money - is still to come.

2. What the customer sees, and it is three lines rather than a tracking number. A customer who opens their track page on a courier order now reads where the parcel has got to ("Delivery: The driver is on the way"), who is bringing it ("Driver: Ah Meng · PMM 1234"), and a press that rings them ("Call the driver"). Each line draws only when there is something to put on it: an order with no trip on it is drawn exactly as it was before any of this existed, and a trip booked a moment ago carries no driver yet because the courier has not matched one - so the card has two lines and not a line reading "Driver: -". That is deliberate. A card that named an empty driver would be telling your customer to wait for somebody who does not exist yet.

3. The customer's page is never taught a courier's vocabulary. This is the seam, said as a rule a customer can be caught by. What the backoffice publishes about a trip is one of a handful of NEUTRAL words - finding, on_the_way, collected, delivered, stopped, nodriver - and the customer's own words for those live in the shop's own language file, in all three languages. The courier's status string (ASSIGNING_DRIVER, PICKED_UP) never leaves the backoffice. A phase the shop's page has not been taught draws NOTHING rather than the raw word, so a second courier arriving with a status list nobody has heard of cannot make your customers' page read wrongly. Six new words were added to all three languages, and a test holds every one of them and refuses a label that has lost its placeholder - the broken-label class this shop has been caught by before.

4. The driver is a SECOND call, and for most of the day it refuses - which is not an error. The trip record itself carries only a driver id, and the driver's own record - the name, the plate, the number - sits behind its own endpoint. That endpoint answers nothing until about an hour before the pickup, so a check you make in the morning comes back with the trip and no driver at all. That refusal says nothing about the trip, which is perfectly healthy, so it is folded in as an absence and never reaches your screen as a fault. Failing a working check over a name that is not available yet would be breaking a feature over a detail you had not asked for. The same rule runs in the other direction: a check that learns nothing this time keeps the driver it already had, because treating an empty answer as an erasure would take the driver's name and number off the customer's card at the exact moment they became useful. The card in the backoffice says all of this under the booking button, in words.

5. The order moves itself, and the row says so where you will actually see it. When a check reads the trip as collected, the order moves to Collected / Shipped and a note appears on the row in your own list: "Lalamove says it collected at 10:35 am - this row moved itself. Put it back if that is not right." Written on the row rather than said in a message, because a message is gone in seconds and this is a change to an order's own money that you may not look at until the evening. The note names WHO moved it and WHEN, so the row can never read as something you did and cannot remember doing. The message at the moment it happens carries both facts in one line - what the courier said and what this app did about it - because two messages would let you read the first and miss the second, and the second is the one that changed an order.

6. A check made at the end of the day still moves it. A parcel cannot have arrived without being collected first, so a trip that already reads delivered moves the order exactly as a trip reading collected does. Without that, an order whose collection you were never told about would sit on Baked for ever and the app would look like it had missed the trip entirely - and the check you actually make is often the evening one, after the run.

7. The Undo, and the two things it deliberately does not do. The row gains an Undo beside its other presses. It puts the stage back and it puts the money flag back, and it is careful about both. An order that had no payment flag at all goes back to having none, rather than to a "false" this Undo invented - because a flag that was never there and a flag saying she has not been paid are two different things on this app's own rules. And the stock rule is deliberately NOT run on the way back. That is not an oversight: the stock rule reads a row's current stage, which at that moment is the last one, so it cannot tell "put this back" from "step this out of Baked" - and stepping OUT of Baked puts ingredients back on your shelf. Charging one order's ingredients twice because an API spoke is exactly the kind of thing this release must not do. The move itself changed no stock (only stepping into or out of Baked does), so there is nothing to undo. This was measured, not reasoned: the whole shelf count was

read before and after.

8. An Undo has to survive you looking again. The guard that stops the rule moving a row twice is the stamp it writes when it moves it - NOT the trip's phase. That difference is the whole thing. If the guard were the phase, then looking at the trip again would move the row straight back, and an Undo that a look undoes has not undone anything. So after an Undo the row is once again a row the rule has not moved, and the trip being collected does not move it a second time. Measured live: after an Undo, a fresh check left the row exactly where she had put it back.

9. A row you have already marked Collected / Shipped is left alone. There is nothing to move, so nothing moves. A trip booked on an order you had already marked as gone needs no help from anybody, and the row stays yours.

10. The callback you chose to leave until after this release, and what that costs. The plan for this stage included a public callback - a small function the courier could call the moment a parcel is picked up, so the customer's card would update by itself within seconds. You chose to leave that until after this release, and the reason is on the record: Lalamove's webhook specification could not be read from here, and a public endpoint built against a guessed shape is exactly how a security hole gets shipped - it is the one part of this work that is reachable by anybody on the internet rather than only by your own signed-in app. So the honest description of where things stand is this: the customer's page is as fresh as the last time a phone asked the courier, and pressing Check is what asks. Her own order row catches up the same way. Nothing is lost, and nothing is claimed that is not true - the callback needs no database change and can be added whenever you want it.

11. One database step, and the order it must be run in. supabase/courier_job.sql adds five columns to the order_tracking table: the courier's name, the neutral phase, the driver's name, the plate and the number. It is safe to re-run and each column is added only if missing. RUN IT BEFORE THIS BUILD IS DEPLOYED. The backoffice publishes the whole tracking row in one call and PostgREST rejects the entire call if it names a column that is not there - and the publish is deliberately silent about its own failures, so it would not look like a failure: every customer's tracking page would simply stop updating, for every order, with nothing on any screen saying why. This is the same trap the courier charge migration documented, and the SQL file carries the warning in its first paragraph.

12. The shop's own page had to learn five new columns, and forgetting one fails silently. The customer's card asks PostgREST for a named list of columns, and it returns ONLY what is named there. A column the backoffice publishes and this list forgets is a line the card can never draw - and it fails with no error anywhere: the row arrives complete, the card is simply missing a section, and nothing says why. So the five names were added to the request, and a test now reads the request the page actually sends rather than trusting that it was edited.

13. A fault found by measuring the drawn card, which no test in this repository could have seen. The two things on the customer's card that matter most at the door - ringing the driver and opening the courier's own tracking page - were built from padding alone and came out 30 pixels tall, while every other thing a customer taps on that page (the delivery day, the collect / courier pair, Track, Place order) is 37 to 48. The two smallest targets on the page were the two the customer most needs at the door. Found by measuring the drawn card at 375 pixels, not by reading the rule. They now take the page's own floor and centre their word in it, re-measured at 38 pixels. And the reason no assertion caught it is worth the sentence: every stand-in screen in this repository answers a box with the number it was handed, so a link 30 pixels tall and a link 48 pixels tall measure exactly the same. The guard therefore reads the stylesheet's own declared numbers instead, and it holds both halves at once - the floor and the box that makes the floor mean anything.

14. Every rule that stands on this was proved load-bearing. Sixteen faults were put back in all, each watched failing a test that NAMES it and then restored byte-identically. Among them: a check that moves an order every time it looks; an Undo that restores an invented "false" where there was no flag at all; a stage change that forgets to tell the customer's own card; an order already marked paid being un-paid by a move past it; a late check that skips the collection it was never told about; a courier's raw status word published to the customer; a plate labelled as a driver; a number that cannot be dialled turned into a button anyway; a driver read as a bare

string; and the customer's card dropping one of the five columns from its request. Two of the sixteen are worth naming because a test could not have caught them before a person read the screen: the dispatcher's own import list, where a name called but not imported is a 500 that only appears when that action is used, and the two links on the customer's card being 30 pixels tall, which is section 13. One guard turned out to be unreachable by pressing anything at all - the transition rule in the courier panel and a second guard inside the order rule both answer a repeated check, so the inner one can never be reached from the screen. Rather than leave a guard nothing can prove, the rule was exported and a test written for the one state that isolates it: the state where YOU have moved the stage back by hand after the rule moved it, which is a state you really can reach. The suite is 1545 passing with none failing.

15. Where the secret lives, and where it does not. Unchanged and worth restating, because this release adds a second call that spends nothing but reaches the courier's own records: the signing of every request happens on Supabase, in a function only your own signed-in app can call. The key and the secret are Supabase secrets you set yourself; they never enter this repository, this app's screen, or any message. The driver's number is the one thing this release publishes on purpose, and it does so because you asked for it and because the person at the door is a stranger your customer has to meet - but it is published only for a trip that has one, and only as a number that can actually be dialled.

16. Nothing of yours is rewritten. Measured on the running app, on a phone-sized screen, with the whole state read out of storage before the pass and read out again afterwards and compared field by field. The pass was the full one: a trip checked, the order moving itself, the Undo pressed, and then a SECOND check made afterwards. The result, in full. After the Undo, the order's stage and its payment flag were byte-for-byte what they had been - the Undo really does put the row back and not merely draw it back. The only fields that changed at all, across the entire pass, were the trip's own: what the courier last said, when it said it, and the driver it named. Nothing else in the whole stored state moved - not one ingredient, not one order's items, not one setting. The three saved days still read 24, 4 and 24 pans, exactly as they did before. And the phone's storage holds exactly the same three keys it held before the pass, with no backup key and no scratch key left behind.

17. What is still to come. The delivery run - one trip carrying several drop-offs at once, which is the stage that saves you the money your own hunch was about. It is v191, and it opens with a question to you rather than with code.

25 Sep 2026 - engine v189, BOOKING THE TRIP (no database step). The second step of the courier work you asked for: the price you chose becomes a real van. One press books the trip with Lalamove, the share link it sends back goes into the tracking box your orders already have, and one press checks where the driver has got to or calls the trip off. Two real defects were found by reading the drawn screen rather than the code - the job card had been printing the courier's own word `ASSIGNING_DRIVER` instead of "Finding a driver", and the price panel was printing the word "null" under its last row - and both are fixed, guarded and proved. The customer-facing half had no test at all and has three now. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "i want to do Lalamove, API to manage courier, make it ready for other courier as well. Focus now for laalamove. And check our backoffice readiness. From my understanding lalamove support last day deliver consolidation to save on cost." Three separate things. v188 built the quote, which was the first. This is the second: booking. The seam and your consolidation belief were both answered in v188 and nothing about them changed here.

2. What a booking does, and what it deliberately does not. The trip comes back with its own share link, and that link is written into the tracking box this app has carried since v97 - the box that is already printed on the customer's card and in the shipped WhatsApp message. That reuse is why a booking needs no database step at all: an order row syncs whole, so a trip booked on one phone reaches your other phone with no migration and no new column. A share link IS the courier's reference for this delivery, so putting it where a tracking number used to go is honest rather than a shortcut. One copy touch goes with it: when that box holds a web address the message says "Track your delivery" instead of "Tracking number", because it is a page the customer opens rather than digits they read out.

- 3. The driver's name and number plate cannot be shown at booking - and the plan promised one.** This is the one place this release does not deliver what was planned, and it is said here rather than glided over. Lalamove has no tracking number and no driver at the moment a trip is booked; the driver is assigned afterwards and Lalamove does not release the name and plate until about an hour before pickup. So booking shows the price, the vehicle and the share link, and the driver arrives later, if you want that screen. There is nothing to fix; the promise was simply wrong when it was written.
- 4. The two defects, both found by reading the drawn screen rather than the code.** The first: the job card printed "ASSIGNING_DRIVER" where the app has a translation table with a test on it. The table was correct and tested; what was missing was the wire from the card to it, because the card asks the courier "what is this status in words" and the courier object did not answer that question. The fallback printed the courier's own enum instead of complaining, which is what a quiet fallback does. It reads "Finding a driver" now, and the whole interface has a guard so a second courier cannot ship without the answer. The second: the price panel printed the word "null" under its last price row, because two optional lines were handed to a browser function that turns a null into the word. It was invisible to every test the app had, because the ordinary stand-in screen quietly drops nulls - only the strict one keeps them, and it now draws this panel too.
- 5. Your three decisions, carried forward and unchanged.** "Locate it, let me fix it" for addresses, so a courier is handed a point and never words on a card. "All four stages", so this is the second of four and the remaining two follow as v190 and v191. "Build now, keys later" - the app still carries no key and no secret, and says so in plain words rather than failing in a way that looks like a bug.
- 6. Measured live on the drawn screen, at 375 by 812 and at 1280 by 900.** The job card reads "Lalamove is on this order", then "Car · RM 14.00", then "Booked 1:51 am", then the two presses, then "Status: Finding a driver - read just now", then the customer's own link. Pressing Check the trip against a reply of PICKED_UP put up "Trip status: Collected" and the card then read "Status: Collected - read just now". Pressing Cancel trip opened the app's own confirmation, which names the consequence and the money: "Call off this Lalamove trip? The driver stops being sent, and the customer's tracking box keeps the link but nothing will update it. This cannot be undone from here - you would have to book again, at a fresh price." Confirming it put up "Trip called off" and the card then read "Status: Called off from here 9:28 am". The presses measure 115 by 36, 99 by 36 and 110 by 36, this app's own tap-target floor, and the confirmation's two presses 87 by 46 and 100 by 46.
- 7. Where the honest answer is not the convenient one.** Calling a trip off sends the courier a request that answers with nothing at all - so the card does NOT claim the courier said "Cancelled". It records that YOU called it off, with the moment, in the app's own words: "Called off from here". The same reason runs the other way on a finished trip: the heading changes from "Lalamove is on this order" to "A Lalamove trip on this order", the Cancel press is withdrawn rather than offered and then refused, and the Book presses come back to life so you can rebook a trip you cancelled. And a checked status is merged into the trip's record rather than replacing it, so looking at a trip never forgets the price it was booked at.
- 8. The money guard, measured.** A trip that is still running blocks a second one. Measured live with a live trip on the order: all three Book presses drawn disabled, and the panel's own sentence under them - "This order is already on a trip." That is the guard that stops you paying for a second van to the same door.
- 9. The customer-facing half had no test, and that gap was found by writing the faults rather than by reading the code.** v189 puts a trip's share link into the tracking box, which makes the "is this a link or a number" branch customer-facing on every courier order. The number half was covered; the link half, which is exactly what a booking newly produces, was not - on either side. Three tests were written: the storefront now proves a link is drawn tappable and a javascript: value is not; and the shipped message now proves it routes the value through the one label-maker rather than labelling a URL a tracking number.
- 10. Nothing of yours is rewritten.** No module, batch, start time, cycle or saved day. Nothing was written to Supabase, and no real network call left the browser for the whole measurement: the courier was answered by a stand-in. The pane's own stored data was compared field by field afterwards and put back - the only fields the measurement had changed were the trip's own state and the app's two customer edit-stamps, and the fake trip, the fake share link and the fake session token are all removed. The pane holds exactly the three storage keys it

held before, with no backup key left behind.

11. Every rule that stands on it was proved load-bearing. Twenty-four faults were put back in all, each watched failing a test that NAMES it and then restored byte-identically. Among them: the request envelope the API needs; a reply not being unwrapped; the stop ceiling checked AFTER the doors; a doorstep matched by position instead of by where it is; a finished trip still counting as live; a phone number losing its plus; a trip number dropped into a URL unescaped; the vehicle and time being re-sent on a booking, which would be a second and quieter way to choose a vehicle; a GET carrying an empty body; and a cancel sending one. Four of them are named rather than counted, because a test could not have caught them and one was found only by a fault: the printed "null" and the printed enum, both read off the drawn panel; the storefront's link branch; and the missing "done" flag on a job written before that flag existed, which every test passed until a fault exposed it. The suite is 1519 passing with none failing.

12. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. A booked trip is kept on the order row, which syncs whole.

25 Sep 2026 - engine v188, A DELIVERY PRICE BEFORE YOU PROMISE ONE (no database step). The first step of the courier work you asked for: on an order you are delivering, one press asks Lalamove what the trip would cost on every vehicle it runs, shows the prices side by side with the distance and the quotation's own five-minute life counting down, and the one you choose fills the charge box that was already there. Nothing is booked, nothing customer-facing changes, and not one module, batch, start time, cycle or saved day of yours was rewritten. There is no database step.

1. What you asked, in your own words. "i want to do Lalamove, API to manage courier, make it ready for other courier as well. Focus now for laalamove. And check our backoffice readiness. From my understanding lalamove support last day deliver consolidation to save on cost." Those are three separate things, and this release is the first of them. The other two are answered below and built on top of this.

2. Your consolidation belief is correct, and here is the mechanism. Lalamove does not consolidate by holding parcels back for a later trip; it consolidates by letting ONE trip carry several drop-offs, two to sixteen of them, priced as one base fare plus a fee for each extra stop. Outside the Klang Valley, which is Penang, that fee is RM 1 on a motorcycle, RM 2 on a car, RM 5 on a four-by-four or van and RM 10 on a lorry, so eight cakes to eight houses on one motorcycle is one base fare plus RM 7 against eight base fares. Three things go with it and all three are said rather than hidden. The vehicle has to physically hold the load, so one motorcycle is not eight focaccia and the run screen will say so. There is no bulk pricing call, so a run of eight drops is still eight price requests, which costs you waiting time and not money. And route optimisation is described two different ways in Lalamove's own documents, so the stop order is yours to arrange and the screen will never claim it found you an optimum. Building that screen is the last stage of this work and the only one that changes what your shop promises, so you will be asked before it is built.

3. The seam, which is what "ready for other courier as well" had to mean - and a hole in it that the audit found. A claim in a comment is not readiness, so readiness here is a file that holds the one interface every courier must answer, a file that is the only place any courier's own name and rules live, and one channel that both use and that never throws at you. Behind it all is a **fake courier written for the tests, which quotes without the name Lalamove anywhere near it** - and that fake is the proof that a second courier is a new file and a line in the registry rather than a rewrite.

And the audit you asked for, walked against the built thing rather than against the code meant to satisfy it, found the claim was not yet true. It found the same defect twice, once on each side of the seam. On the app's side, two screens and the one shared file named the courier in their own words: the Settings card printed "Courier (Lalamove)" and sent the literal word to the server itself, and the shared channel defaulted the courier to Lalamove when nobody said otherwise. A default is that file quietly choosing a courier, which is the one thing the seam exists to prevent - the day a second courier arrives, a default left behind would send its calls to the first one and nothing on screen would say so. On the server's side, the dispatcher wrote the courier's name into a refusal you would read: "Lalamove did not name any vehicles for this market". That is the identical fault, and it is the one that matters more, because a refusal is exactly the sentence a second courier would have to have found

and edited by hand. Both are fixed with no visible change: the card asks the registry for the name, the channel refuses to name one, and the refusal now comes from the provider's own label. The words on your screen are character for character what they were. And the rule has a test it never had, which is the real lesson here and the same one v187 taught. It now covers both halves: inside the app the word is allowed in two files only, the provider and the registry that imports it, and the walk refuses any third file that names it; on the server the walk does the same, and the dispatcher gets a tighter rule of its own because its environment variables ARE named after the courier and it therefore had to be on the allowed list - it may name the courier only in those variables and in the provider's own file path, and in nothing that could be shown to you.

4. Your three decisions, and where each one landed. You chose "Locate it, let me fix it" for addresses, so a courier is never handed words on a card: it is handed a point. The address already on the order is looked up once, the answer is kept against that customer so a second order from the same number costs no lookup at all, and where the lookup misses a map opens and you tap the exact door. The bakery's own door is pinned once, from the same place you asked for it. You chose "All four stages", so this is the first of four and the other three follow as v189, v190 and v191. You chose "Build now, keys later", and that is exactly where this stands - see paragraph 9.

5. What the screen does, on your own order, measured live at two widths. The price panel opens inside the card that already holds the charge box, folded away until you ask for it, and that is deliberate: this app has one pop-up layer, so a second card would destroy the charge box, the note you were part-way through and the Save button, and the fee would land in a box nothing could save. Measured at 375 by 812 with the test courier answering, the panel drew three price rows - Motorcycle RM 8.50 counting its own life down, Car RM 14.00, Van RM 0.00 - then the line that says which vehicle could not be priced and why, then "3 prices from Lalamove, just now, for collection 2026-09-30 at 10:00. Each one dies on its own clock." The four presses at the foot of the rows measure 99 by 36 pixels, which is this app's own tap-target floor, and the line naming the two ends of the trip reads "From Jienluv2bake, George Town to 12 Jalan Bunga, George Town". At 1280 by 900 the same card drew all three rows, the failure line and the live status line with nothing clipped.

6. Taking a price is one press, and it goes through the money path that already worked. Pressing Use this fee on the RM 14.00 car put 14 in the charge box and said so, and the panel folded itself away so the number it had just written is what you are looking at with the payer question under it. Setting the payer to "The customer paid it" then read back "The customer owes RM 58.00 - items total RM 44.00 + courier charge RM 14.00". Nothing new was invented for the money: the fee lands through the same model that has carried the courier charge since it was built, so the amount, who paid it and COD all keep running the way they already ran.

7. Three places where the honest answer is not the convenient one. First, a quotation DIES - Lalamove's last five minutes - so every price carries its own clock, a dead price greys itself out and says "ask again" rather than sitting there looking live. Second, a courier can reply with a total of zero, and a zero is not a charge: measured live on the RM 0.00 van, the press left the charge box exactly as it was and said "Lalamove priced this trip at RM 0.00 - that is not a charge, so nothing was put in the box. Ask again, or type the amount," and the panel stayed open with all three rows still there. A toast claiming a number had landed in a box the box had refused is the screen lying to you, and this is the one place in the app where a lie about money is easiest to ship by accident. Third, a price that arrives after you have closed the card is not written into it: measured live with a deliberately slow reply, closing the card the moment after asking left the pop-up closed, no charge box and your stored data unchanged, character for character the same length before and after.

8. Where the secret lives, and where it does not. The signing of every request happens on Supabase, in a function only your own signed-in app can call, and the key and secret are Supabase secrets that you set yourself and that never enter this repository, this app's screen or this conversation. The browser half is built to the same shape as the mail function that already works: it never throws, and every refusal arrives as a sentence you can act on rather than a code. Two you will meet: a missing pin reads as a pin to be placed rather than a failure, and an address Lalamove cannot find reads as an address to put on the map. Until you open the partner account and set those secrets the app carries the courier and no keys, and it says so in plain words instead of failing in a way that looks like a bug.

9. What is measured and what is not, said plainly. The panel, the two ends of the trip, the five minute clock, the money, the refusal and the closed-card case were all measured live, in a real browser, on the drawn screen, at 375 by 812 and at 1280 by 900 - against a stand-in courier that returns Lalamove's own documented replies. What is NOT measured is Lalamove itself, because there is no account yet: the signing and the reading of its real replies are tested against its published sample answers in the suite, and three details are flagged at the head of the provider file where the first real call will show them - whether a stop's point is sent under one field name or another, the exact language tag this market expects, and whether the distance comes back with a unit or as a bare number. Each is one line in one file, and each will be settled in the first minute of the sandbox, not guessed at now.

10. Nothing of yours is rewritten, and every rule that stands on this was proved load-bearing. Measured live: your stored data was backed up before the single press that writes anything, put back afterwards and compared, and your phone's own storage holds exactly the two keys it held before, with no backup key left behind. Eleven faults were put back in all, each watched failing a test that names it and then restored byte-identically - among them the zero being accepted as a charge, the closed card being written into anyway, the panel's clock failing to clear itself when the card it belonged to was gone, and all four seam leaks named in paragraph 3. Two of the eleven are worth naming because a test could not have caught them, and both are said rather than glossed. The clock one: the panel stops its own clock when a fee is accepted and folds away, so only the case where the fee is REFUSED and the panel stays open exercises the rule at all, and the test that catches it is the one written for the refused zero. And the server's own hand-written refusal: when that fault was put back, the first version of the new rule did NOT catch it, because the rule read only ordinary quoted text and the sentence was written as a template literal - so a real leak sat in a blind spot of the very guard written to catch it. The rule was widened to read those too, and the fault then fails by name. The suite is 1460 passing with none failing. **No database step** - a price is transient and an accepted one lands in fields that already exist, so there is nothing to run in Supabase.

24 Sep 2026 - engine v187, EVERY PERSON'S WINDOW IS THE HEIGHT OF ITS OWN PEOPLE (no database step). One fault, found by measuring the drawn board against your own eleven points and then fixed: the workers' windows on the Production page carried a ceiling, so on a phone only the first person stood on screen and the other three sat inside a box that scrolled. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked. You told me the Production page deferred from your instruction, and then: "fix it". So this release is your eleven points read back against the board that was actually drawn, not against the board that was intended, and it carries the one place they did not agree.

2. The fault, in one sentence. Your point 8 says "The person's window be taller" - the *window*, not the row. v186 made the row taller and left the window's own ceiling exactly where it was.

3. Why a ceiling was there at all, and why it was wrong here. The window a person's rows sit in had a ceiling of min(20vh, 190px) - 190 pixels on your phone. That number was written for the Scenario planner's own people window, where the rows are 35 pixels tall and five of them is 190. The Production page's rows are 72 pixels, because a taller row is what you asked for. Four of them need 406. The ceiling did not know the rows had grown.

4. What that looked like on your phone, measured. At 375 by 812 the window was 162 pixels tall while holding 406 pixels of people: one person on screen, and persons 2, 3 and 4 behind a scrollbar that was not visible and that you had no reason to look for. Read off the drawn board: 0 of 4 rows fully visible. Your point 2, "Create windows for each person", was answered with one window and three hidden behind it.

5. The fix. One line of the stylesheet. The board's own person window is now as tall as its people and carries no ceiling of its own; the ceiling stays exactly where it was written, on the planner's window, which still has the five 35-pixel rows it was written for. Nothing else about either window changed, and nothing inside either of them was restyled.

6. Measured after, at the same two widths. At 375 by 812: the window 407 pixels tall, its inside 406 and its content 406 - equal, so nothing scrolls - all 4 of 4 rows fully visible, the coaches still 37 pixels wide with four

filling the line, four red rulers still standing at 107 and the one clock face still at 215 reading 11:27 pm. At 1280 by 900: the window 401 pixels, no scroll, 4 of 4 visible, the coaches 155 and the line 738. The modules' window above and the planner's borrowed window below are the heights they were.

7. Nothing of yours is rewritten. The whole of this release is one stylesheet line and one test, and neither of them writes anything: a ceiling is a drawing rule, and no module, batch, start time, cycle or saved day is read by it, let alone changed. Measured live on a four-person day built for the purpose, the app's own stored data was put back from a copy taken beforehand and compared byte for byte the same, with no backup key left behind.

8. And the rule has a test now, which it never did. This is the part worth reading, because the reason the fault shipped is not the fault itself. No test could ever have caught a window that was too short. The stand-in screen the tests draw on answers every box with the same measurement - nothing high and nothing of content - so a window with three persons hidden inside it looks exactly like one that fits. A stand-in that says every box is zero high cannot see a box that is too short, and a guard that measured that screen would have passed just as happily with the bug in place. So the guard reads the stylesheet's own declared numbers instead, and it holds both halves of the rule at once: the board's window declares no ceiling, and the planner's still declares its own. Two faults were put back and each was watched failing that test by name - the ceiling returned to the board's window, and the ceiling taken away from the planner's - so the test cannot be passed by deleting the cap everywhere and leaving the screen these windows were borrowed from unbounded. Each was restored byte-identically. The suite is 1343 passing with none failing.

9. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. A ceiling lives in the stylesheet and in the test, and your plan, your days and your own settings are exactly as they were.

24 Sep 2026 - engine v186, THE MORNING CAN BE WALKED FROM THE WORKERS' OWN WINDOW (no database step). Eleven things you asked for on the Production page, built as one: the lines drag left and right and everything on them moves together, Start the day now / Stop / the bell / Back to now sit where the workers read them, and the walk's own pace is a drop-down of Slow, Mid and Fast. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "let me tell your what i want for production page. 1. Keep the Next Up. 2. Create windows for each person, the detail able to be drag left or right. 3. Put the person windows parallel, close to each other, the 2nd person and subsequent person have the centered red ruler, but the 2nd and other don have to show the clock face. 4. Replicated for person 2, 3, etc. 5. I need walk the day, Button 'START THE DAY NOW', 'STOP' , 'The N Called', these button are for the person's windows. 6. The scenario dual window dont join to the person's window. 7. Name of the scenario should be on Top of page. 8. The person's window be taller, so that coach size better fit wordings. 9. Each of the person line has a clock line at centre of line, when we drag to the right, the train move to right and the clock and red ruler move relatively. 10. When i click outside the person window, the clock back to center. 11. For simulation, i want some button to do this, left buttton make the clock back to current, Right button, when clicked, dropdown a choice list, slow, Mid, fast. This button is to make the chart move faster for simulation purpose."

2. The Next Up stays. Your point 1, and nothing about it moved. The card that says what is next is still there, still first, still telling you the job and its clock.

3. The scenario's name, on top of the page. Your point 7: the heading card now carries the name of the day you are looking at, so a page of numbers cannot be mistaken for another day's.

4. A taller person's line, so the words fit. Your point 8. Each of the workers' lines is now 72 pixels tall, against the 56 it was and the planner's own 35 - measured at both 375 pixels wide and 1280, and every coach measures 52 of those 72. Nothing on a coach face is clipped at either width: the face reads its whole derived word - at 375 pixels the words on your own day read *mixer, fold, Wash, Load, Unload, Dimple, oven, Cool* - and not one word is cut off.

5. The two windows stay your borrowed pair, and they do not join the lines. Your point 6. The modules window and the planner's people window are drawn below the lines by the very code that draws them in Scenario planning, and they are **not** tied to the train: the trains do not scroll sideways at all, so no drag on a

worker's line can move the two windows above it. What v185 put back - the two borrowed windows panning together on both screens - is still true and still theirs; it is the train that stands apart, and that is your point 6 said as a rule.

6. A red ruler on every line, but only one clock face. Your points 3 and 4. Every person's line carries its own red ruler standing at the middle of that line - at 375 pixels wide the ruler sits at 107 and at 1280 the line's own centre is 694 - and the face is drawn **once**, on the first line, reading the clock at the line's own centre in the pane. The second person and every person after carry the ruler and no face, exactly as you asked. Measured at 375 pixels: the face stands at 215 in the pane and each row's ruler at 107 in its own track, and the two differ by exactly the 108-pixel name column between them - so they are one line, drawn in two coordinate systems.

7. The drag, and everything on the line moves together. Your point 9, in your own words: "when we drag to the right, the train move to right and the clock and red ruler move relatively." A sideways drag on a person's line moves the strip, that line's red ruler and the clock face by the same number of pixels - measured live at 375 pixels wide, an 80-pixel drag took the strip from -437 to -357, the ruler from 107 to 187 and the face from 215 to 295, *all three by exactly 80*. Nothing about it is saved, so it dies with the screen.

8. Letting go of the line, or pressing off it, puts the clock back at the centre. Your point 10, and your point 11's left button. A press anywhere off the line re-centres it, and the *Back to now* press above the lines does the same thing on purpose - both put the ruler, the face and the trains back where the day puts them. It is never saved, so it cannot move a module, a batch or a start time.

9. Walking the day, without touching the day. Your point 5, in your own words: "I need walk the day, Button 'START THE DAY NOW', 'STOP', 'The N Called', these button are for the person's windows." So the presses sit above the lines, inside the workers' own window:

- **Start the day now** begins the walk at today's own clock time and takes the screen to the day's first minute. **Stop** ends the walk and puts the lines straight back on the real clock. **The bell** tells you how many jobs would be called and throws the calling on and off. **Back to now** re-centres the line whether or not a walk is running.

Measured live on a 375-pixel screen with the walk started: the line went to the day's first minute and the row read *Now: mixer*, the link's count was 30m, and *12 coaches stood red*, falling away as the clock ate into them. Nothing about the walk is written down anywhere, so a day you saved is exactly the day you saved no matter how long the walk runs.

10. The walk's own pace, and it is only the walk's. Your point 11, in your own words: "Right button, when clicked, dropdown a choice list, slow, Mid, fast. This button is to make the chart move faster for simulation purpose." The drop-down carries *Slow 1x*, *Mid 10x* and *Fast 60x*. Measured live at Fast: three seconds of walking took the clock face from 10:58 pm to 11:01 pm and brought a link's countdown down from 30m to 27m with it. Choosing a speed changes nothing in your plan - it is the walk's speed and it lives nowhere but the walk. **Stop** then put the face and the line straight back to 10:58 pm, turned the press off and brought Start back, saying "Stopped - the line is back on the real clock."

11. The countdown comes down by itself. Nothing has to be tapped for a worker to watch their next job get nearer: measured on the board, one minute of the clock brought a row's own count from *Next 20m* to *Next 19m* with nobody touching the screen. And that minute rebuilt nothing - the app's own rule stands, because a repaint throws away a worker's scroll, so a beat writes the numbers on the line and never rebuilds the line around them.

12. One more thing was mended underneath. A tap on a coach had to be proved to reach your other phone and the store, not just the screen. Reading the tests showed that the tick's write - and the Clear press's write - were being read off the screen's own copy of the day rather than off the saved day, so a tap that moved the picture and skipped its save would have looked exactly like a working one. Both are read off the saved day now, and both fail by name if the write ever goes missing again. This is the app's own standing rule: a stand-in must be as unforgiving as the real thing, because a forgiving one hides a bug instead of catching it.

13. Where a tick lives, and why there. A tick is written in a key of its own beside your plan, never inside it - so **clearing the board can never touch a saved day, and deleting a day can never wipe the board.** *Clear the board* asks first, in the app's own words, before it takes a single green coach back.

14. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Measured live, across a whole pass of this page - the board drawn, the day started and stopped, the speed changed three times, the line dragged and put back, a coach ticked and then the board cleared - the one and only thing that changed in the app's stored data was that new key of the tick's own, sitting beside the plan: an empty board written as a spoken empty so your other phone cannot read a clear as ignorance, and the plan itself untouched. The app's own version reads *Engine v186*.

15. Every rule that stands on it was proved load-bearing. Eleven faults were put back in all, each watched failing a test that names it, and each restored byte-identically. One of them escaped on the first pass and is named rather than glossed: with the beat no longer repainting a row's next-job line, nothing failed - because no test had ever watched a countdown come down on a beat alone, with nobody touching the screen, which is the half a worker at a bench actually reads. A test was written for it, and the fault then fails by name, reading that a minute of the clock moved a row's own countdown from 20 minutes to "Next 20m". The suite is *1343 passing* with none failing.

16. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The tick is written under a settings key the app already guards, the walk's speed is a number in the walk and nowhere else, and nothing new is written anywhere.

24 Sep 2026 - engine v185, THE COACH KEEPS ITS WORD, AND THE BOARD'S PAGE IS PUT IN YOUR ORDER (no database step). Two things: every coach on the Production line's train now carries a word at every width, and the board's page reads top to bottom the way you asked - the person's train first, then the two windows borrowed from the Scenario planner, kept original. Nothing of yours is rewritten, and there is no database step.

1. What you asked, in your own words. On the coach face: "Phone coach" reads as "Two lines: icon + name, then time", and on the broken label, "Yes, fix it." Then, on the page itself: "let me emphasise, the production page start with the person's windows the train inside, after the N person, we have the 2 windows that we bring in from scenario planning, we keep it original." Asked what should sit below the train rows, you chose both planner windows - the module window AND the people window exactly as Scenario planning draws them.

2. The coach keeps its word. v184's refinement - "if the coach box is too small to house the full words, then just show meaningful, hover mouse on show tooltip, with more details" - had been built as *drop the name and keep the icon*, and reading the drawn board on a phone showed what that cost: at 375 pixels every coach was nameless. So a narrow coach now gives up its type size and its side padding instead - 10 pixels to 9, and 3 to 2 - and never the word. The stylesheet rule that hid the name outright is gone, and a test now refuses to let it come back.

3. And the word keeps its punctuation off. The word on a coach is derived from the module's own name, and it was keeping whatever punctuation sat against it: "Wash, oil and fill" was showing as *Wash*, with the comma welded on. It reads *Wash* now. That one was found by reading the drawn board rather than by reasoning - the same way the last one was found - and it is exactly the kind of broken label your refinement refuses.

4. The board's page, in your order. The Production page now reads top to bottom: the person's train rows first, then the modules window and the people window, drawn by the same code that draws them in Scenario planning. They are borrowed, not re-drawn - a second renderer of one day is the thing this app refuses - so the bars, the rows and the counts in both windows are the planner's own.

5. What that order costs, said plainly rather than hidden. Two things follow, and neither is dressed up. The people now appear twice on that one page: once as train rows at the top, and once inside the planner's own people window below. And because that borrowed people window sits on the planner's minute axis, it pans together with the modules window again on both screens, while only the train stays fixed - v184 had taken the pan off it because the train had replaced it, and you have put it back. The train itself still never scrolls sideways, which is what keeps its centre the centre.

6. A seam, so the train is not mistaken for the map. A hairline rule now runs above the modules window on the board, so a worker can see where the coaches end and the day's own map begins. It is a seam on the

wrapper only - nothing inside either borrowed window is restyled, which is what "keep it original" means.

7. One more thing had to be mended underneath. The app remembers where you had scrolled in each window and puts it back after it redraws - your standing rule that a redraw must not move what you are looking at. It was finding the FIRST window of each kind; with one page now holding two people-shaped windows, the second one was being left behind. It keeps every one of them now, each by its own numbers.

8. Measured, on the drawn board at 375 x 812, signed out of the cloud. The three windows are in your order on the page - the train 57 pixels tall at the top, the modules window 493 under it, the planner's people window 35 at the foot - with 31 coaches 37 pixels wide, all four to the visible line, the clock standing at the line's own centre, and every coach face reading a whole word: *mixer, fold, Wash, Load, Unload, Dimple, oven, Cool*. Not one clipped word and not one nameless coach anywhere on the board.

9. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Your three saved days still read 24, 4 and 24 pans, and there is nothing to run in Supabase - not one stored field is added or changed.

10. Every rule that stands on it was proved load-bearing. The coach's word, the punctuation it keeps off, and the seam above the modules window were each put back as a fault, each watched failing a test that names it, and each restored byte-identically. Two of the three were caught by name; the seam was caught by the page-order test, which now refuses a board whose train and day's map run together. The suite is 1341 passing with none failing.

24 Sep 2026 - engine v184, THE TRAIN ON THE PRODUCTION LINE, AND TWO CLOSER STOPS ON THE DIAL (no database step). The board's workers' window is no longer a strip of time - it is a train of work. One coach is one job, joined in the order they happen, with the clock standing at the centre of the line, a countdown on every link between two coaches, and one tap on a coach turning it green. Your own words for it: "For the production, the effectiveness need to be improve, for each of person line, make his series of work join up like a train, coaches represent work, the current time at the center sharing with all person, showing next work in minutes on top, something like next station countdown for subway. Person should click on their work to turn it green indicating acknowledgement." Nothing of yours is rewritten, and there is no database step.

1. What you asked, in your own words. Six points, and then one refinement while it was being built: (1) "i want to have 2 step more closer scale"; (2) the Production line should join each person's work up "like a train, coaches represent work, the current time at the center sharing with all person, showing next work in minutes on top, something like next station countdown for subway", and "Person should click on their work to turn it green indicating acknowledgement"; (3) "The coach can pass the current timeline, but stay red, click it turn green"; (4) "Each title of the line show their name, make each line distinctively seperated"; (5) its purpose - "to tell person, what they have done, and what is coming dynamiccally, on the work that have not done, every person can see it, and can land. helping hand if their process can tolerate an interruption"; and (6) "Ask me question if not sure" - which is what settled the six design answers below. The refinement: "if the coach box is too smalll to house the full words, then just show meaningful, hoover mouse on show tooltip, with more details."

2. Two closer stops on the dial. The Scenario planner's Scale control went from four stops to six. The two that were there stay where they were and only one name changed: the old Closest is now called Closer, and the new top stop is Detail. Three lists have to be the same length or the ruler loses its step and the chart breaks, so a test now says so out loud rather than trusting it.

3. The answers you gave, which the build follows. Asked five rounds of questions, you chose: the clock is "Pinned to the centre"; one coach is "One job they do"; "Anyone, on any line" may tick; clearing is "A Clear the board button only"; green means "I'm on it (seen)"; a coach goes red "When the job is due"; coaches are "All the same width" with numbers "Only on the links"; ticks "Travels to your other phone"; the two windows are kept with the person's line taller; and the architecture is "The train is its own strip". And on the clock itself: "lock the clock relative to the train, when drag, clock line move, when click outside, the clock centred on the line, make the coach not too big, in a full shown line should be able to visualize 4 coach, we can also show the coming time countdown on the coach link."

4. What the person's line is now. A taller 56-pixel row per person, separated from the next by a rule and a rail in that person's own colour - your point 4, "make each line distinctively seperated". The name cell carries who it is and what is next ("Next 26m", "Now: rests"). Alongside it, one coach for every job that person has, all the same width, in the order the jobs happen. Each coach carries its icon, a word that means it, and the clock time. That word is DERIVED and never cut: "Into the proofer" shows as "proofer", "Dimple and top" as "Dimple", "The oven swap and the bake" as "oven" - because a face showing the word "The" would be showing a word and saying nothing. The full module name, the two clocks, the length, the batch and the state in words are all on the coach's tooltip, which is your refinement read at the only moment it applies.

5. The clock, the countdown, and the colours. The clock is one line at the centre of the line, shared by every person, and its label is the clock time - never the word "now", so it can never be mistaken for the planner's own now-line. Every wait between two jobs carries its countdown in minutes on the link, "like next station countdown for subway", and it counts down live. A coach turns red the moment its job is due and STAYS red after the clock has gone past it - your point 3 exactly. One tap turns it green and names the job back to you; the same tap takes it off again. A green coach travels to your other phone, and a Clear the board press appears beside the calling switch, asks first, and clears every coach at once.

6. Dragging reads the day without moving it. A sideways drag on the workers' window moves the clock line only and leaves every train exactly where it was, because a drag is somebody asking "how long have I got" and not a change to the day. The label then reads the minute the line is standing on and the line is drawn in the reading colour. Letting go, or pressing anywhere off a row, puts it back at the centre. It is never saved, so it dies with the screen.

7. One wording fault found by measuring, and fixed. Live at 4:30 am on your own day, Person 3's one job had run 4:00 to 4:12, so their row read "All done" - beside a coach that was standing RED and untaken. That is the board telling a worker they have finished something nobody has touched, and talking them out of the very tap your point 3 asks for. A row with work still untaken now counts it instead ("1 not taken"), and gives the count up the moment the last coach goes green.

8. One real bug, found only by measuring in the browser. The widths of a train's coaches were being assigned to the strip in a way that looks right and writes nothing at all, so every coach fell through to the stylesheet's own fallback size. Measured at 375 pixels wide: 34-pixel coaches with SIX of them filling the visible line, where the plan promises four. It is fixed, and the test that watches it was made stricter FIRST so the fault was caught by a named failing assertion before the fix landed - the same lesson as every earlier version where a forgiving stand-in let a real fault through.

9. Measured on your own day, at 375 x 812 and at 1280 x 900, signed out of the cloud. At 375: the line the coaches stand on is 213 pixels, and a coach is 37, a link 18, a side stub 4 - four coaches and three links and two stubs come to 210, so exactly four coaches fill the line, which is the figure you gave. The clock stands at 215 pixels, the centre of the line. At 1280 x 900: a 723-pixel line, 151-pixel coaches, 30-pixel links, and again exactly four to the line, with the clock at the centre at 470. Wei's row read "Now: rests" with her five coaches at 4:00, 4:13, 4:30, 5:01 and 5:32, the first three red and the one at 4:30 ringed as the job running now; the links counted 31m and 62m, and five minutes later the same links read 26m and 57m with the clock label moving to 4:35 am. One tap on Person 4's coach turned it from red to green, said "Oil the pans and weigh the dough out at 4:00 am - marked as taken.", wrote the tick, and turned that row's line from "1 not taken" to "All done". A drag from the centre to 150 pixels read 4:13 am in the reading colour with every train's position unchanged to the pixel; letting go put the clock back at 215 and the label back to 4:30 am. The Clear press asked first, wrote the empty map rather than deleting it, turned every coach back to red, and made itself quiet again.

10. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. After all of the above - the taps, the clear, the drag, the clock moving, the window resized - your stored data was put back from a copy taken beforehand and compared: byte for byte the same 13,888 characters, your three saved days still reading 24, 4 and 24 pans, and only the app's own two keys left in the phone's storage.

11. Every rule that stands on it was proved load-bearing, not assumed. Sixteen faults were put back in all, each watched failing a test that names it, then restored byte-identically. Among them: with a coach sized as a

plain quarter of the line instead of out of the budget it fails reading that four coaches no longer fill the visible line; with a wait between two jobs given no width it fails the round trip that reads every minute of a row back off its own strip; with the due test given an upper bound it fails reading that a coach which has passed the clock is no longer red; with the tap no longer deleting the key it fails reading that the last tick off left something behind; with the centre fixed instead of measured it fails the placement of every row; with the pan guard removed it fails the test that says the board's window cannot pan while the planner's still can; and - the two worth naming, because they escaped the first sweep - with the tap's save or the clear's save skipped, nothing failed at all until the tests read the STORE and not just the object the app holds in memory. A tick that moved on screen and never reached the phone's saved file looks perfect until the phone is opened again, and no test could see it. Both now fail by name. The suite is 1339 passing with none failing.

12. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The ticks live in the settings row your phones already sync, beside the calling switch, and that row learned one thing in this version: an emptied board is now SPOKEN rather than passed over in silence, so the other phone takes your clear instead of reading it as ignorance and putting every tick back.

CHANGELOG and the guide both carry all of this, both PDFs rebuilt and read back with PyMuPDF, and the app's own More screen reads Engine v184.

24 Sep 2026 - engine v183, REASSIGNMENT HAPPENS BY ITSELF, AND THE PERSON'S CARD SAYS SO (no database step). One sentence on one card, and nothing else in the app changed. You went looking for a reassign button on a person's card and there was not one, so the card now says why there is nothing to press: the day hands out its own work, so changing that person's hours or their ticks changes their job list on the spot, and a job that no longer fits them is given to somebody else. Your own words that settled it: "if you know that reassignment is automatic, then forget about my request for reassignment button, just need to mention work reassignment is automatic with worker card detail change." Nothing of yours is rewritten, and there is no database step.

1. What you asked, in your own words. "where is the reassign button?" - and then, having answered it yourself: "you can put the reassign work button inside person's card", and then the correction that ended it: "if you know that reassignment is automatic, then forget about my request for reassignment button, just need to mention work reassignment is automatic with worker card detail change".

2. Why there was nothing to find. Reassignment has been automatic since v177, and it is automatic on that very card. The day gives out every job nobody has been named for, and the only two things it reads about a person are the two facts you type on their card: the hours they are here, and the modules they are ticked for. Change either and the day re-picks hands immediately - which is why the card never needed a button, and why a button there would have been a second way to say the same thing.

3. What the card says now. One sentence, in place of a control: the day gives out this work by itself, so changing the hours above or the ticks changes these lines on the spot, and a job that no longer fits is handed to somebody else. And what you placed by hand stays yours - a person you put on a module yourself, or a stretch you moved off somebody's row, is not moved by this card. That last half matters as much as the first: the card only ever moves the work the day was already choosing between.

4. The rule itself, written down the way you stated it. "in a module finding his person, module will look at person's with the least skill and fit its module. if fit then assign him first." So a module asks the person who FITS it - trained for that module, and inside the hours they are here - and between two people who both fit it asks the one trained for the fewest modules first, with the person who can work anything at all asked last of all, exactly as you confirmed it. Ties go to the lower person number. It is a preference and never a refusal: a job nobody trained is free for still goes to somebody, and the row says why.

5. Measured, on your own day, at 375 x 812, signed out of the cloud. Your card for Wei opened reading "11 min of work, in 3 places, and nothing collides", "Here all day", and her five jobs: 4:00 am Dimple and top, 4:13 am The oven swap and the bake, and the three stretch and folds at 4:30, 5:01 and 5:32. The new sentence was on that card, and the card carried no press about the work at all - the only button on it was its own close. Then

one tick, "Mixing the dough in the tub", and four of Wei's five jobs left her on the spot: her list fell from five lines to one, and Person 3's rose from one to five. The day still stood at the same four people, so nothing invented a hand and nothing was lost - the work was handed on, which is the whole claim. Taking the tick off put all five lines back exactly as they were.

6. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day, and no word of your plan. Measured live: after working the card, ticking a module and taking the tick off, your stored data was put back from a copy taken beforehand and compared - byte for byte the same 13,888 characters, with your three saved days still reading 24, 4 and 24 pans, and only the app's own two keys left in the phone's storage. One trace is named rather than glossed, because the same shape was named in v180: taking the last tick off leaves twelve characters - an empty skills entry - where there was nothing before, exactly as your stored data already carries for merges. It reads as no training at all, it changes nothing about the day, and a tick you never press never writes it.

7. The sentence was proved true, not just printed. A card that said reassignment was automatic and did not do it would be worse than saying nothing, so the test reads the job list off the card itself after one tick. Three faults were put back and each was watched failing a test that names it: with the sentence deleted, the test fails reading that the card says nothing about the reassign button; with the training made unable to keep a job off anybody, it fails at the day's own end, reading that the day went on giving the free job to somebody trained for something else; and with the card no longer re-reading itself on a tick, it fails at the card's end, reading that the card went on drawing the day it was built from. The claim can break at either end and the two are told apart on purpose. The suite is 1313 passing with none failing.

8. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The card's sentence reads the same two facts v177 already stores in the scenario, and nothing new is written anywhere.

24 Sep 2026 - engine v182, THE SHADE IS ONLY WHAT YOU SET, ONE PRESS MOVES ONE BATCH, AND THE PRODUCTION-LINE NUMBERS ARE PROTECTED (no database step). Three things in one version, all of them inside the Scenario planner except the last. First: the shaded band the app used to work out for you is gone for good. Her rule, in her own words: "Wei should not have any shade because there is no work time set for him, shade should just follow what i set, not other consideration." So the only shade a person's row ever carries is the two hours you typed on their card, and a person whose hours nobody has typed carries none - no band, no "Working ..." line, nothing. Second: "delta t on one batch of the module dont change the module batches, it should not be." It did - one press on the first batch of a module moved every batch of it - and it does not now. One press moves the batch under your finger and nothing else, and Back onto the line takes off that one batch's hold and leaves any other batch's hold standing. Third, carried over from what v181 measured and left unprotected: the numbers on More -> Production line are now guarded exactly like your saved days, so a newly set-up phone can no longer push its own defaults over the numbers you typed.

1. What you reported, in your own words. "build v182. And the working time shade is wrong", and then, while this was being built: "delta t on one batch of the module dont change the module batches, it should not be." Both were measured before anything was changed, and both were real.

2. The shade, and why it goes rather than gets quieter. v179 drew a band on each person's row worked out from their own day - their first job to their last, the gaps between included - and v180 made that band stand down on a row where you had typed hours. Both were answers to a question you had already answered yourself, and v180 left the app holding two opinions about one working day. The measurement that settled it was on your own day: the computed band was drawn 93 minutes wide, 4:00 am to 5:33 am, on a row whose five jobs come to eleven minutes, beside a card that already read "11 min of work". It was measuring the stretch and calling it the work. So it is deleted - from the day, from the row, from the card, and from the drawing itself - and what is left is the one thing you actually said: the hours you type. Typed hours draw a band; no typed hours draw nothing. The card still says "Here all day" for a person with no hours set, because that is the real answer for every person on every day you have ever built, and a blank where a fact should be reads as a fault.

3. The delta t, and the batch it belongs to. A press on a batch is a hold - the batch runs later than the line puts it, and it comes with the module above it when that module moves. Since v154 the first batch of a module had been read as the module's own offset, so holding it moved every batch of that module with it: on a module running four batches you pressed one bar and four moved. Batch three, pressed the same way, moved only itself - which is the answer you asked for, so every batch now behaves like batch three. Moving a whole module is still one press away and it is the honest one: the module's own start time on the module card, from which every other batch is spaced by the module's own pace. The words changed with the rule, so nothing on the screen claims a move the app no longer makes: the toast now reads "Batch 3 held back 5 minutes from where the line puts it", the tip on a held bar "held back 5 min", and the way back "Batch 1 back at 4:00 am - exactly where the line puts it."

4. The Production-line numbers, which v181 measured and left open. v181 closed the same family of fault for your saved days - silence must never delete - and named this one as not fixed: the numbers on More -> Production line have no "empty" of their own, so a brand-new phone could still push its defaults over what the cloud had. It is fixed here, and it needed no new stored field. The signature of a phone with no opinion about the line is a plan that still reads exactly the stock numbers; that signature is now asked of both sides, so a phone that has never touched the line takes the cloud's numbers on its first pull and carries them in its first push, a phone whose numbers are your own keeps them against a stock cloud and queues one publish to put them back, and a plan you deliberately set to the stock numbers is the one case that cannot be told apart from silence - by design, and said plainly rather than hidden.

5. Measured on your own phone, at 375 x 812, signed out of the cloud. Your own day first: four person rows, not one computed band anywhere in the app, not one typed band, and the word "Working" appearing nowhere on any screen; Wei's card read "17 min of work, in 3 places", "Here all day", and nothing else about her hours. Then hours typed onto Wei's card, 5:00 am and 9:00 am: stored as 60 and 300 minutes from your 4:00 am start, and drawn as one band 96 pixels along the row and 384 wide at your own 1.6 pixels a minute - 60 minutes and 240 minutes - with the computed band still absent and the card reading "Wei is set to be here 5:00 am -> 9:00 am, but Dimple and top runs to 4:06 am - a job that cannot fit inside those hours." Then the delta, measured by giving the oven four batches for the length of the test: its four bars sat at 0, 139, 278 and 418 pixels. Pressing + 5 min on batch 1 moved that bar 0 to 8 pixels - five minutes at 1.6 - its card read 4:00 am -> 4:15 am becoming 4:05 am -> 4:20 am, its tag read B1 delta t=+5, and batches 2, 3 and 4 did not move a pixel, at 139, 278 and 418 exactly as they were. Pressing + 5 min on batch 3 moved batch 3 alone, 278 to 286. Pressing Back onto the line on batch 1 put it back to 4:00 am with a plain B1 tag and said "Batch 1 back at 4:00 am - exactly where the line puts it", and batch 3's hold was still standing.

6. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. The copy of your stored data taken before the first press was put back afterwards and compared: byte for byte the same 13,888 characters, with the four batches used for the measurement gone and the oven back to the one batch you had it at, and with the two hours typed onto Wei's card taken back off again. Your three saved days are exactly the ones you had - 24, 4 and 24 pans. The only keys left in the phone's storage are the app's own two, and nothing here blocks a sale or reads an order.

7. Every rule that stands on it was proved load-bearing, not assumed. Eleven faults were put back in all, each watched failing by a test that names it, then restored byte-identically. Among them: with the computed band put back into the day's own model it fails reading that a person's row carries nothing but the hours she typed; with it put back into the row's drawing it fails that same test from the other end; with the rule put back into the styling it fails the third time; and with it put back onto the worker's board it fails reading that the board shows the hours she typed and nothing the app worked out for her. With the first batch's press carrying the rest of the module again, three named tests fail - including the one on the last module's own bar. Of the production guard: dropping the key from the guarded list fails four tests; making every plan count as somebody's opinion fails five; dropping the cloud-side guard fails two; and reading a plan that has no field at all as an opinion fails the test that says a cloud with no plan and a phone that has one is put back. One fault escaped on the first pass and is worth naming: restoring the old "Back onto the line" spread was NOT caught, because every test pressed the way

back on a batch that happened to be the only one held - the tests had pinned the press and left the undo unpinned. A test was written that holds two batches of one module and takes one back, and the fault now fails by name. The suite is 1312 passing with none failing.

8. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The production guard reads a signature off the numbers your phone already syncs, and the batch hold has always lived in the delta your plan already carries.

24 Sep 2026 - engine v181, A NEW PHONE NO LONGER EMPTIES THE OTHER ONE (no database step). You reported that signing in on a new phone showed none of your saved scenarios. It was real, and it was the app rather than the phone. The cause is one word: silence. When the new phone pulled from the cloud it had nothing to say about your saved days, and the app read that silence as a decision to delete them - so the next thing it wrote back to the cloud left them out, and the cloud lost them. Three rules now stand on one sentence: silence means this phone has no opinion, and silence must never delete. A newly set-up phone now receives your saved days and the day you built, and its first push carries them back up instead of dropping them; and a phone that still holds something the cloud has lost puts it back by itself, with no press from you. A deletion you actually made is spoken out loud, so it is honoured and never refilled.

1. What you reported, in your own words. "i try log in with the new phone, the save scenario is not loaded? why". The "why" is answered below. This version is the answer to it happening again.

2. Why it happened. Everything you have set up in this app - your saved scenarios, the day you built, your to-do list, your wish list, the developer line - travels to your other phone in one single row called "settings". There is only ever one of those rows, on every phone, from the first day you install the app. A phone that has just signed in therefore already has a settings row of its own to publish, even though it has never seen the cloud copy, and the app decided which side to keep by the clock: whichever phone noticed last won. A new phone wins that comparison simply by being newer, and when it won, it wrote its own empty settings up to the cloud. The saved days were not emptied on purpose and were not refused - they were simply not there to write. A push replaces the whole row, so a key the payload leaves out is a key the cloud loses. That is the fault: silencing a thing is not the same as deleting it, and the app could not tell the two apart.

3. The three rules, and they are one sentence between them. Silence means this phone has no opinion, and silence must never delete. First: if you empty a shelf on one phone, that deletion is now spoken out loud - the app says "the saved days are now an empty list" instead of saying nothing at all, so the other phone cannot mistake your decision for ignorance. Second: when a phone's own settings win, every one of those five keys it has no opinion about is taken from the cloud instead of being left out - taken into what it is about to publish, so its push cannot delete them, and taken into its own app, so a phone that has never had your days receives them on that same pull. Third: if the cloud has lost one of those keys and this phone still holds it, this phone queues one publish to put it back - no press from you, and never for a key you emptied on purpose.

4. What a newly set-up phone does now. Measured. Against a stand-in cloud holding your own saved days, the new phone receives all three of them and the day you had built, and the settings it then sends up still carries those same days - so the push that used to empty the shelf now leaves it whole. Measured again from the other end: with the cloud's settings row missing your saved days and your own phone still holding them, one pull from that phone queues exactly one publish carrying them, and a second pull queues nothing, because there is nothing left missing. And a shelf you really did empty is not put back - the empty answer is spoken, read as an answer, and left alone.

5. What this version protects, and what it does not yet. Read this part. This version protects your saved scenarios, the day you have built, your to-do list, your wish list and the developer line. It does not protect the numbers on More -> Production line. Those are a different kind of value, with no "empty" of their own - a brand-new phone's numbers can still replace what the cloud has, and that was measured, not guessed: the cloud's production plan was pushed over as a brand-new phone's defaults. It is the same family of fault one key over, it is not fixed here, and it is not dressed up as fixed. Until it is, check More -> Production line on both phones and set the numbers on the phone you trust; and if a number has already been replaced, the daily cloud

backups are the way back, not a re-typing.

6. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day was changed by this version, and nothing here blocks a sale or reads an order. Measured live on your own phone: after unlocking the app, signing out of the cloud and opening the Scenario planner, your stored data was put back from a copy taken beforehand and compared - byte for byte the same 13,888 characters. Your three saved days read 24, 4 and 24 pans. That last figure is worth one line, because two earlier notes in this history said "24, 12 and 24": 12 is what the third day is set to aim at, 4 is what that day as built actually makes, and 4 is what the app has always computed and still computes. The number in those two notes was the mistake, not the app, and the earlier entries are left standing rather than quietly tidied.

7. Every rule that stands on it was proved load-bearing, not assumed. Eleven faults were put back in all, and each was watched failing by a test that names it, then restored byte-identically. Taking the shelf the phone is silent about out of what it publishes, so the push deletes it again, fails reading that a phone that never had her saved days receives them. Taking the same value out of her own app instead, so the phone stays empty, fails the same test from the other end. Letting the phone overwrite its own days with the cloud's, and refusing the cloud's empty answer when she has deleted a shelf on purpose, each fail the test that says a phone with its own days keeps them and a shelf she emptied stays emptied. Dropping the repair of a short cloud row fails reading that a phone that still has her days puts them back. Making that repair fire on every pull fails reading that a cloud row with nothing missing queues no publish. Dropping the built day from the guarded keys fails reading that a phone that never built a day receives the one she built. Removing the spoken deletion fails reading that emptying her shelf is said out loud, and making every phone speak - including one that never held the key - fails reading that a missing field is not announced as an empty shelf. Two of the eleven were faults in the first draft of the fix and were found by this sweep rather than by reasoning: one change was unfalsifiable as written, because the cloud's value was copied into her own app before the comparison the fault was supposed to break, and one guard turned out to be protecting a real case - a hand-edited or corrupted file - which had no test until the sweep asked for one. The suite is 1309 passing with none failing.

8. No database step. Not one stored field is added or changed on the server, so there is nothing to run in Supabase. The three rules are read by the sync code your app already runs, and they act on the same five keys that already travel between your phones.

23 Sep 2026 - engine v180, THE HOURS YOU TYPED TAKE THE SHADE OFF, AND THE LAST BAR OF THE DAY CAN BE HELD (no database step). On a person whose working hours you have filled in on their card, the shaded band the app worked out for you is no longer drawn. Your two clock boxes are the answer for that person, and a window you typed with a second, computed window drawn inside it was two answers to one question - which is what you saw and reported. Nothing else about the shade changes: on every person whose hours you have NOT filled in, it is drawn exactly as it was, and on every person it is still named in words - on their row, in the tip and on their card. The bar at the very end of your last module now carries the same readings and the same pair of presses as every other bar, so it can be held like any other; the card it opens is still the day's own card, with the whole day read off it underneath.

1. What you reported, in your own words. "dont shade if the person have indicated work time", and then, so there was no doubt which screen you meant, "and those person with work time specify, the shade is wrongly indicated". Both true, and neither was your phone.

2. What was wrong. The shade added in v179 is the stretch of the day a person's own jobs actually occupy, worked out from the day. It was drawn on every person, including the ones you had already filled the two hour boxes in for. So on those rows two bands lay on top of each other: the hours you typed, and a second, computed band sitting inside it. Measured on a row with 5:00 am to 9:00 am typed in, at 1.6 pixels a minute: the typed band ran 96 pixels along the row and 384 wide, and the computed one was drawn on top of it, starting later and ending earlier wherever the work did not fill the hours. Two different lengths for one person's working day, and no way to tell which one the row meant.

3. What it does now, on your rule. The moment you have said the hours for somebody, that person is not

shaded at all. Your typed hours are the only band on their row, and they are read as the answer. Where you have NOT filled the boxes in - which is every person on every day you have ever built, including all of today's - nothing changes at all: the shade is drawn from their own first job to their own last exactly as it was in v179, because there is no typed answer for it to argue with. The rule is one line, held in the model rather than on the screen, so the row, the worker's board and the person's card cannot come to different conclusions.

4. And the words stay, deliberately. Taking the drawing away must not take the fact away. The working stretch is still named in words everywhere it was: on the person's row ("Working 4:00 am -> 5:33 am"), in the tip, and on their card, where it sits beside the hours you typed so you can still read the two against each other. That is the half of this release you may not have asked for, and it is deliberate: you said do not SHADE it, not do not SAY it. If you would rather the words went too on a person whose hours you have typed, that is a one-line change - say the word.

5. And the other thing you reported, which I had explained away. "the last module batch pop up, still dont mark his delta?" - the very last bar of your day. Tapping it opened the day's own card and nothing else: it carried no batch number, no delta t, no cycle reading and no way back, while every other bar on the chart carried all four. In v179 I answered that this was by design, because since v152 that bar is where the day is worked backwards from, so the card it opens is the day's own card. You reported it again, so I measured it rather than explaining it again - and it was real. A bar you can see and tap but cannot hold is missing the one thing every other bar on your chart has, and the day's own card does not need that to be true.

6. What that bar does now. The card underneath is unchanged: it is still the day's own card, still headed the end of your first batch, still naming the day's first batch, the day's finish and every module on it. What is new sits on top of it - the four readings every other bar carries (which batch it is, the delta t, the cycle, the pans), the same pair of presses, + 5 minutes and - 5 minutes, and Back onto the line. A hold there is a hold on that whole module, exactly as a hold on batch 1 of any other module is. Measured on your own day at 4:00 am: five minutes on the last bar moved that module five minutes later. Its row read "starts 4:05 am", the tag on its bar read B1 delta t=+5, and its card read Batch 1, delta t = +5 min, 4:05 am to 4:17 am, where it had been 4:00 am to 4:12 am. Not one module above it moved a minute. Nor was any of them given more room: every one of their latest starts still read 4:00 am, which is where they already were, and the day card's own line - that every module of your line is already as late as it can go - was as true after the hold as before it. The day's own finish did not move either: 6:03 am both times, because the day finishes with whichever module ends last, and on your day that is The rests and the stretch and folds, which the packing never delays. The pans a day did not change at all: 6 before and 6 after. Pressing Back onto the line put the batch back to 4:00 am and the tag back to B1 with no delta t, and the round trip left your stored data byte for byte the same.

7. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Nothing here is stored at all: the shade is worked out from the day each time the screen is drawn, and the only thing that decides whether it is drawn is whether you have filled the hour boxes in. Measured live, after typing hours into a person's card, taking them out again, and opening and shutting the cards: every field of the day read exactly as it had before, with one trace left behind, and it is named here rather than glossed over. The first time an hour is typed, the day grows an empty shifts key, and taking the hours back out leaves the key standing at nothing. It is twelve characters, it reads as no hours at all - as does an absent key, which is the one thing it can never mean anything else - and it is the same shape your own stored data already carries for merges. So the day is not byte-identical to itself after a round trip; it is identical in every field, plus those twelve characters of nothing. The copy taken beforehand was then put back and compared: 13,888 characters, byte for byte, and your three saved days read 24, 12 and 24 pans.

8. Every rule that stands on it was proved load-bearing, not assumed. Ten faults were put back in all, and each was watched failing by a test that names it, then restored byte-identically. With the rule dropped so everybody is shaded again it fails reading that a person whose hours you typed is still wearing a shade; with the shade dropped everywhere it fails the other way, reading that a day with no hours typed anywhere lost its shade; with the rule inverted so only a typed person is shaded it fails both; with the row drawn from the work again, past the rule, it fails reading that the two bands are back; with the words silenced along with the drawing it fails

reading that the fact left the row; and with the shade made to answer to the whole day rather than to the person it fails reading that one person's hours took the shade off somebody else's row. Of the last bar: with its card taken back to the day's card alone it fails reading that the card no longer says which batch it is; with the day's own reading dropped off it, the other way; and with the two pairs of presses stopped saying which one moves the whole day it fails reading that the card no longer names the pair. The suite is 1298 passing with none failing.

9. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. You do not need to type your hours again on the other phone: the shading is decided by the hours you already had.

23 Sep 2026 - engine v179, WORKING HOURS AND THE DELTA T (no database step). Each person's line is now shaded across the hours they are actually working - drawn from their own jobs, from their first of the day to their last, with nothing for you to type. And the delta t is back across the whole chart: a press on batch 1 of a module that keeps its own time is a hold again, with its own way back. Your widest scale gives up the delta t before it gives up the batch number, on your own rule.

1. What you asked, in your own words. "can you shade their working hours in their line?" - and, asked which of two ways you meant it, you chose the one that needs nothing typed. And then, separately: "i ask for delta time, that function is not worker across the chart. And the return to original button is missing?" Both were real, and neither was your phone.

2. The shade, drawn from the day itself. Each person's row carries a wash from their own first job of the day to their own last, the gaps between jobs included - somebody who is back in an hour has not gone home. It is worked out from the day rather than typed, so it is drawn on every day you have ever built, including the days with no hours set on anybody. Where you HAVE typed the two hours on a person's card, those stay exactly as they were: the typed hours are the outer band and the working shape sits inside them.

3. Measured live, and one thing to expect on a row. On your own day at 1.6 pixels a minute, Wei's shade was drawn at 0 and 149 pixels wide, which is 93 minutes of a 4:00 am day - 4:00 am to 5:33 am - against her five jobs. A job shorter than four pixels is still drawn four wide so you can see it and tap it with a thumb, so on a row whose last job is one minute the last bar reaches a couple of pixels past the shade. Measured live, Wei's last job is the fold, one minute, drawn at 147 for 4 pixels and ending at 151, where the shade - which reads the minutes themselves - ends at 149. The shade is the true length of the work and the bar is the true size of your thumb, and the difference is only ever that.

4. The delta t, and why it had gone. A press on a batch used to be a hold measured from where the line puts that batch - a delta t. Since v154 that had been narrowed to batches below batch 1, and to batch 1 only of a module set to follow the one above. Every module of your own day is set to "Its own time", which is the one case the narrowing left out, so batch 1 of every module wrote an absolute start time instead: no hold was ever written, no delta t was drawn anywhere on your chart, and the card's own way back - which exists only while there is a hold to take off - disappeared with it.

5. It is back, and the way back is called "Back onto the line". That is its name from v151, and it can be renamed in one line if you meant the other wording. Measured live on your own day at 1.6 pixels a minute, pressing + 5 min on batch 1 of Dimple and top: the card read Batch 1 and delta t = +5 min with 4:05 am to 4:11 am, the bar's own tag read B1 delta t=+5, the note read "Dimple and top held 5 minutes later than the time you gave it - every batch of it moved with that.", and the three modules below it did not move a minute. Pressing Back onto the line put the batch back to 4:00 am and said so: "Every batch of Dimple and top is back at the time you gave it." Going there and back left your stored data byte for byte identical at 13,888 characters, because the hold is kept in the delta itself and none of your stored start times was rewritten.

6. The badge now knows when to keep quiet. Your own rule: "if the scale is too wide to show batch no. and delta t then forgo delta t". So the batch number - the thing that tells one batch from another - is never given up, and the delta t is printed only above the widest stop. Measured live at the widest reading, 72 pixels an hour: the tag read exactly B1 while the bar kept its own tint, which takes no room and still answers which of the batches is held, and the tip on that tag still named the hold in full.

7. Why the last module's batch 1 is different - the question you asked while this was being built. It

behaves differently on purpose, and it has since v152: tapping batch 1 at your LAST module does not open an ordinary batch card at all. That batch is the moment the whole day hangs from - it is where the day is worked backwards from - so its card is the day's own card, headed "the end of your first batch". Measured live, tapping batch 1 of Cutting and packing read "Cutting and packing - the end of your first batch", named the day's first batch ending at 4:12 am and the day finishing at 6:03 am, and listed all eight modules. On the FIRST module it is the other way round: batch 1 there has no module above it to be held back from, so a press is simply that module's own start time. Everything in between is the ordinary batch card with the delta t and its way back.

8. And where "Put my start times back" lives, since you looked for it. It is on that day card and it is a one-shot undo: it appears only once you have pressed "Pull them back to their latest start", because there is nothing to take back until something has been pressed. Measured live on your own day, every one of the eight modules was already ticked as late as the line allows, so the card said so - "Every module of your line is already as late as it can go - there is nothing to pull back." - and offered neither press, which is why you could not find it there.

9. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Measured live, after working through the batch card, both presses and the widest scale, your stored settings were put back from a copy taken beforehand and compared: byte for byte the same 13,888 characters, and your three saved days still read 24, 12 and 24 pans.

10. Every rule that stands on it was proved load-bearing, not assumed. Thirteen faults were put back in all, and each was watched failing by a test that names it, then restored byte-identically. The shade: with it drawn from the whole day instead of from that person's own jobs it fails reading whose work it covers; with the bars painted before the shade it fails reading that the bars no longer cover the wash; with the shade given a stacking order it fails reading that the hours climb over the module titles, which is the fault this screen already had once. The delta t: with the hold narrowed back to the v154 rule it fails reading that a press on batch 1 wrote a start time instead of a hold; with the module's other batches left behind it fails reading that the hold did not take them with it; with the widest scale still printing the delta t it fails reading that the batch number was given up for it; and one fault in the tests' own stand-in screen was found and fixed on the way. The suite is 1294 passing with none failing.

11. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. Both facts live inside the scenario, which already syncs whole, so they travel to your other phone with the rest of your plan.

23 Sep 2026 - engine v178, REASSIGNMENT (no database step). "Share them out" now really hands the whole day back. It used to clear every module and leave the stretches you had placed by hand exactly where they were - and a hand-placed stretch is read first, ahead of the module's own person, so the press said one thing and the chart drew another. Both full-day presses now take those stretches back, and they say how many they took back, because taking back a decision of yours is a change you are told about.

What you asked, in your own words. "how to rerun job assignment after we set the more details?" - and then, once that question had been measured and answered, "build v178, Name it as REASSIGNMENT". The answer to your question is the People box, and measuring it is what found the fault below.

1. The fault: the press promised a day it did not draw. The People box holds three answers, and two of them redraw the whole day's hands at once. "Share them out" cleared every module - but a stretch you had placed by hand is kept on the module as its own entry (that is what makes a hand-over move one stretch and not a whole row), and it is read BEFORE the module's own person. So the day kept your hand-placed stretch standing while the press said it had handed everything back. Measured on your own day: after the press the toast said the day had been shared out, and the box's own closed words still read "Your own" - and your Dimple and top was still sitting on Wei, where you had put it.

2. What it does now. Both full-day presses - "Share them out" and "One a module" - take every hand-placed stretch back along with everything else. The stretch is not deleted quietly either: the press names how many it

took back, and a day with none adds nothing to the sentence, because a count of zero is not news.

3. Measured live on your own day, starting at 4:00 am at 1.6 pixels a minute, with your one hand-placed stretch on Wei. Pressing "Share them out": the box went from "Your own" to "Sharing them out", your hand-placed entry was gone from the record, and the toast read "Sharing them out - every job reassigned to as few hands as can cover the day. 1 stretch you had placed by hand goes back on the day's own arrangement." Dimple and top moved off Wei's row onto the row the day picks for itself, and every other job on the day was untouched. Pressing "One a module" on the same day: six people, one to the six modules that hold hands, the box reading "One to a module", and the same sentence naming the one stretch it took back - and the stretch no longer overrides the person that press had just given that job.

4. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario was changed by building this, and nothing here blocks a sale or reads an order. The only writes are the ones your own press makes. Measured live: after both presses and a full pass of opens and closes, your stored data was put back from a copy taken beforehand and compared - byte for byte the same, 13,888 characters - and your three saved days still read "No fridge, 1 person" 24 pans, "My sister proposal 21/9/2026" 12 pans and "One baker day" 24 pans.

5. Every rule that stands on it was proved load-bearing, not assumed. Eight faults were put back in all, and each was watched failing by a test that names it, then restored byte identically. With the stretch left standing, the test fails reading that the old arrangement is still on the chart; with the count left out of the sentence, it fails reading what the press said; with the count worked out after the day had already been cleared, it fails the same way; with "One a module" left holding the stretch, it fails reading that one job is short of the person the press gave it; with the empty entry written back as an empty map instead of no entry at all, it fails reading that nobody has one spelling; and with the count written always in the plural it fails on the first hand-back. The suite is 1287 passing with none failing.

6. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. Everything here is inside the Scenario planner, under More.

23 Sep 2026 - engine v177 (no database step). Each person now has two facts of their own, and both are yours to state once: the hours they are here, drawn as a band along their row and a real bound on who the day may hand work to, and which modules they are trained for. A job nobody is named for goes to somebody who can work it, inside their hours, and between two of those the day asks the most specialised person first. Nothing is refused and nothing is hidden: a job that lands outside somebody's hours or training is shown and said in words.

What you asked, in your own words. "back to the scenario planning. 1. I need each person to have the start work time and end work time, 2. which module they are dedicated for, and only taking over others job, as priority 1,2, or3." And then, while this was being built, the two sentences that settled the shape of it: "a person we should be able to specify their skill, by module, can be more than one, by module", and "module should select the one that specialised". Of the takeover order you said the thing that made it simple: "the person card superseded the module person assignment. maybe no priority rank is needed to be simple". So there is no 1, 2, 3 rank anywhere: the card says what each person can do, and the day picks by itself.

1. Their hours, counted from your day's start and not from the clock. You asked for a start time and an end time, and told me how to read them: "the working hours is relative to the chart, not exact hours, end also relative to the chart". So the two boxes are clock times you can read at a glance, and what is stored is the minutes from the start of your own day - which means moving your day's start slides every person's hours with it and nothing is retyped. Measured live on your own day, which starts at 4:00 am: typing 05:00 and 09:00 stored 60 and 300, which is exactly five and nine o'clock less four.

2. An empty box is a real answer, and it is "here all day". Leave both empty and that person is here all day, which is what every person on every day you have ever built already is. The hint under the boxes says so in one line, and it also says which clock the times are read against, because two empty boxes with no explanation would read as two boxes you had not filled in. Typing the whole day by hand - 4:00 am to midnight - is stored

as no entry at all, so "here all day" has one spelling rather than two that mean the same thing and could ever drift apart.

3. And a pair that comes out back to front is refused and SAID. An end earlier than the start is a mistake, and the app now puts both boxes back to the hours the day is actually using and says which way round they go. A box that took a number and quietly did something else with it is the kind of thing that reads as a fault for weeks, and this app does not do that.

4. What the hours are for. Your own aim, in your own words: "the person that have the most consistant work is the person handling oven, colling and packing, he only help do others when he has free slot, by one person not doing too many functions, it is easier to train that person as will. Probably can employ part-timers." So the hours and the training are what the day reads when it hands out a job nobody is named for - and they are a PREFERENCE, never a refusal. A job nobody qualified is free for still goes to a free pair of hands, and the row says why.

5. Which modules each person can work, and the specialist is asked first. The person card now carries a tick for every module that actually holds hands, so there is never a tick against a module switched off or against the proofer, which would be a tap that does nothing. Only six of your eight modules hold hands on your own day, and those six are exactly what the card offers. Tick nothing and that person can work anything, as they do today - the same absence rule this app already uses for automatic cycles and merged jobs. Tick one or more and the day only ever gives them those.

6. And between two people who can both do a job, the more specialised one is asked. Your own rule, in your own words: "say one person have all skill, person2 have have the skill, person3 have only 1 skill, module should select the one that specialised". Somebody trained for one module is asked before somebody trained for five, and somebody who can work anything is asked last of all - which is what leaves the specialist on their own job and the generalist on what is left over.

7. Measured live, on your own day. With everybody untrained and no hours set, your day packs to four people and 58 minutes of hand work, and the three stretch-and-fold slots all sit on Wei. Ticking Person 4 for "The rests and the stretch and folds" and nothing else moved all three folds off Wei and onto Person 4, even though Wei is the lower number and was free at that minute - the specialist was asked first, exactly as you asked for. Wei's row fell from five jobs to two and the other rows were untouched. Unticking put the day back to where it was, to the job.

8. Their hours, drawn on their own row. Each person's hours are painted as a wash along their own row, under everything the day draws on it, so you can see at a glance when somebody is here without reading a number. It carries no stacking order of its own on purpose: a band given one would climb over the pinned module names, which is the fault this screen already had once. Measured live at a phone's width on your own day: a shift of 05:00 to 09:00 was drawn at 96 pixels along the row and 384 wide, which is 60 minutes and 240 minutes at your own 1.6 pixels a minute. It takes no taps at all, so it can never swallow one meant for a bar over it.

9. And when a job cannot fit somebody's hours, the card says so in words. This is the part that matters most to the way you work, because it is the same rule this app has always kept: a plan that does not fit is shown, never quietly rearranged. Measured live: with Wei set to be here 4:00 am to 4:05 am while the dimple runs to 4:06 am, the card read "Wei is set to be here 4:00 am to 4:05 am, but Dimple and top runs to 4:06 am - a job that cannot fit inside those hours." The same sentence appears on their row's tip and on the Production line board, all three printed from one funnel so they can never disagree.

10. And a job outside somebody's training is said in the same place. Somebody given a job they are not trained for is told so in the same words and the same place: the job is named, the times are named, and the note says it is said here rather than the job being quietly moved. The card also reads back what they are trained for - "Trained for 1 module: The oven swap and the bake" - so the answer you typed is the answer you can see.

11. On the Production line, the worker's own hours. The board a worker reads already says what they do next and how long they have; it now also says when they are here, and names any job that cannot fit inside

those hours. Withholding that while printing their jobs would be perverse - when you are here is the first question of a working day. The board is still a READING and writes nothing: measured live, a tap on a person's row on the board left your stored settings character for character identical, and the board offers no hour boxes and no ticks.

12. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day was changed by this version. Measured live: after a full pass of presses on the planner - the two hour boxes typed into and cleared, ticks set and taken off, cards opened and shut - your stored data was put back from a copy taken beforehand and compared, byte for byte the same 13,888 characters, with the two new keys absent again because writing them is something only a press you make can do. Merely opening the screen adds nothing. Your shelf itself is untouched: No fridge, 1 person still reads 24 pans, My sister proposal 21/9/2026 still reads 4, and One baker day still reads 24. And the two chart windows kept their pan at 300 across two of those presses, so a card write no longer throws either window back to the start of the day.

13. Every rule that stands on it was proved load-bearing, not assumed. Eighteen faults were put back in all, and each was watched failing by a test that names it, then restored byte-identically. Two of the eighteen were faults in earlier drafts that the fault sweep found, and both were unfalsifiable as first written - a test that could not fail - which is the same lesson this file has carried before about a stand-in being more forgiving than the real thing. Among the rest: with the band given a stacking order the test fails reading that the hours climb over the module titles; with the band drawn last it fails reading that the bars no longer cover the wash; with the boxes storing a clock reading instead of minutes from the day's start it fails reading where the band was drawn; with an untick leaving an empty list behind it fails reading that unticking is an answer and not a missing entry; with the copy of a day not carrying the two keys it fails reading that the hours of the day you left are still on the day you opened.

14. No database step. Not one stored field is added or changed on the server, so there is nothing to run in Supabase. Both new facts live inside the scenario, which already syncs whole, so they travel to your other phone with the rest of your plan and need no new table and no new SQL.

23 Sep 2026 - engine v176 (no database step). The Production line is now a board a worker reads, not a form you fill in: the day chart from your plan, a clock line showing where the real time falls in it, and a strip saying what each person does next and how long they have. The screen used to ask you for 21 numbers; it now asks for 12, and 8 of those are folded away.

What you asked, in your own words. "can we simplify the production portion, it ask for alot of key in" - and then the shape of what you wanted: "maybe one function for production line, 1. copy the chart into Prodcution line, but should not be editable, it become a dashboard for worker, give the worker good inform of what next for them, how long, and others, you can be creative". And of the boxes on the screen you said the thing that made this simple to build: "must of the question are answered in scenario plan". They are, and that is where the board now reads them from.

1. The day chart is carried across, not drawn a second time. The board and the Scenario planner are the same drawing, produced by the same code, so the two screens can never disagree about a day - not its modules, not its batches, not its times, and not who is on what. What the board drops is everything that changes the day: the Scale buttons, the People box, the backwards chip, the undo chip, the Start the day now button, New module, and Save changes. Measured on your own day: the planner wears 8 buttons above the chart, and the board wears none of them - with the two cards shut the whole screen carries two buttons, and those two are the fold heads at the foot. What the board keeps is everything that only reads: you can pan the day with a right press, the clock follows your finger along the ruler, and the tips on the module names still open. A board that could not be panned would be a worse board on a long day.

2. A clock line, drawn with no press at all. A board nobody has touched yet still has to say what time it is, so the line is placed the moment the screen opens and kept in place once a second. It is drawn exactly where the real clock falls in your day: at 96 pixels an hour on your own day, a press of nothing at all puts the line at 156 pixels for your 4:00 am start, and the clock it reads out matches the ruler under it. Measured live on your own day: at 5:00 am the line stands at 252 pixels and reads 5:00 am - which is 156 plus exactly one hour of 96. At

4:30 am it stands at 204 and reads 4:30 am. Before your day begins it parks at your day's own start and says "day starts 4:00 am"; after your day has ended it parks where the day ended and says "day ended 6:03 am". It never wraps around to the other end of the ruler, which is the tempting wrong answer and would have put the line in the middle of a day that is not the one on the screen.

3. What is next, in clock times and not countdowns. Above the chart the board says, for each person, the job they are on and the one after it, with the clock time each starts and how many minutes it runs: on your own day at 5:00 am that reads "Wei - the rests and the stretch and folds, 5:01 am, 1 min, then the rests and the stretch and folds at 5:32 am". Clock times and not a countdown on purpose: a countdown means rewriting a box somebody is reading every single second, and this app has a rule about that - a redraw must never move what you are looking at. On a day that is not shared out to people, the strip falls back to one list of what the line does next.

4. And the strip answers AGAIN as the clock runs, without touching the chart. This one was found while measuring, not reported: the strip was drawn once when the screen opened, so a worker opening the board at 4 am would still be reading "next: mix at 4:01" at eight in the morning. It now looks again whenever the clock has changed the answer, and only then - inside the working day it redraws at the exact minute a job starts or ends, and never in between. The chart itself is not redrawn, which is the planner's own rule and the reason a worker's scroll is never lost. Measured live across a job boundary: at 9:00 am the strip names Cutting and packing, at 9:10 am it names the oven swap, and the number of times the strip has been rebuilt does not change between them.

5. Three things about a day the clock has already gone past. The first was a real fault, found by opening the board in the evening. The strip had one fallback sentence for two different empty lists, so a day that finished at 6:03 am was being described as "Nothing on this day needs hands" - which is a different thing entirely, and the opposite of true. A day whose jobs have all been and gone now says so: "The day is finished - every job on it has been and gone." The second is the offer to make a sound. The board rings people one minute before each job, and it rings by working out from the clock which jobs are due - but on a day that is already over there is nothing left due, so switching calls on would ring for nobody. That offer is now withdrawn on a finished day rather than left promising a sound it cannot make. Measured live at 7:16 pm on your own day: no Start calling button, and the line under the strip reads "The day is finished". The third is the other half of that, and it matters more: calls you turned on earlier can ALWAYS still be turned off. Stop calling stays on a finished day, so a sound started before cannot be left ringing with no way to stop it.

6. What left the screen, and where the numbers went. Twenty-one boxes became twelve. Nine of them are now read off the day you built under More -> Scenario planner - how many pairs of hands, the pans a tub makes, oiling and weighing out, the dimple, the packing, the oven and the bake-and-swap - so those questions are asked once and answered on both screens. The twelve left are the ones genuinely yours: four behind the first fold (how long you will bake for, the pans you own, what your proofer holds, and the pace between batches you would like), and eight behind the second fold (the rests, the fold, the two proofs, the cooling, the clock you want the first batch at the oven, and the minutes of soft tolerance around a start). Both folds start shut and open in place, so opening one never redraws the chart above it. The line card still answers while it is shut, in one line: measured on your own day, "44 pans today - the proofer is the wall that holds it back", worked out from the day rather than typed anywhere.

7. And the timetable card is folded, which is where it belongs. The backwards timetable - the last moment each stage may start - was left on this screen on your word ("you decide, that is your baby") because it is worker information: it is the one thing on the screen that says when the dough has to be in the tub. But it is a page of advice, not a glance, and it was sitting out on the board in front of a worker who came to read what is next. It now lives behind its own fold head. Measured at a phone's width with both cards as they come: the timetable card is 64 pixels tall, its head and nothing else, where it was 943 while the timetable sat outside it. That one change alone takes the whole screen from 2005 pixels to 1126.

8. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day changed by this version, and nothing here blocks a sale or reads an order. The board is a READING of your day, and every

number on it is computed rather than stored - the box you type a number into writes the one number you typed and nothing else. Measured live, going through a full pass of presses on the board: a bar tap, a module tap, a right press, and both cards opened and shut twice each, with your stored data compared before and after. Your data reads 13,888 characters, and it is byte for byte the same characters it read before the first press.

9. Every rule that stands on it was proved load-bearing, not assumed. Nine faults were put back in all, and each was watched failing by name. The strip: with the finished-day sentence taken off, the test fails reading that a finished day falls back to the sentence for a day with no hands on it; with the calling offer always drawn, it fails reading that a finished day offers a sound it cannot make; with the foot dropped whenever the day is over rather than only when calls are off, it fails reading that a sound already started cannot be turned off; with the strip never rebuilt, it fails reading that the strip still names a job the clock has passed; and with the live minute taken off the strip's own signature, it fails reading that the parked sentence no longer names the hour it is standing in. Two more put a null where the parked sentence goes, and each was caught by the word "null" appearing on the screen. The fold: with the timetable put back outside its fold, the test fails reading that a shut card is already showing the clock times its caret says are behind it; and with the test's own reader taken back to one that walks straight through a hidden box, the same test fails - the reader had been reading text the browser never paints, which is the same lesson this test file already carried about a stand-in being more forgiving than the real thing. Every one was restored byte-identically afterwards. The suite is 1266 passing with none failing.

10. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The day the board draws is the one already synced, and the four numbers behind the first fold were already synced with the rest of your production plan.

CHANGELOG and the guide both carry all of this, the app's own More screen reads Engine v176, and the Production line names itself as the board it now is.

23 Sep 2026 - engine v175 (no database step). Two things you reported: the scenario windows no longer jump back to the start when you press something, and the right-button drag now works while a card is open on the chart. Both inside More -> Scenario planner.

What you reported, in your own words. "Few problem of screen jump here and there. One obvious one is the scenario windows, when i click, window reset. And right button drag dont work". Both were real, and neither was your phone.

1. Why a press threw the windows back to the start. When you press something that changes the day - a batch card's + 5 min, the People box, the scale, opening a saved day - the screen is drawn again from scratch. That has always been how it works, and it is what keeps the whole panel agreeing with itself. But a box drawn again is a NEW box, and a new box starts its own scroll at nothing. So every one of those presses put both windows of the day chart back at the far left and the top, and would have put the list of saved days back at its first row with them. Measured live at a phone's width with the modules' window panned 300 pixels: one press of + 5 min left both windows reading 0.

2. What it does now, and where you were is put back only after the measuring is done. Where each of the three sliding boxes was - the two windows and the list of days - is read before the screen is drawn again, and written back afterwards. The order is the whole of it: the list of days is measured to work out its height, and a row's foot moves when the list under it slides, so measuring a list that had already been slid back down would come out short by exactly the amount it moved and would cut your fourth day in half. Nothing is slid until the measuring is finished, and the measuring all happens before anything is drawn on the screen, so you never see the day sitting at its start either. Measured live on your own day: panned 300, one press of + 5 min, and the window reads 300 on the day drawn after it - and both windows are put back by their own numbers, never one number shared, because they pan together and a redraw is not the place to find out they had drifted.

3. Why the right-button drag did nothing with a card open. A card is drawn on a full-screen sheet that sits over the whole chart, so with a card open your right press landed on the sheet, not on the day behind it - the day you were actually looking at could not be moved at all. Measured live at a phone's width with a batch card open: a right press 200 pixels into the chart, dragged 60 to the left, left the window reading 200 exactly where it started.

The sheet carried no gesture of its own, so nothing was being taken from anything else.

4. What it does now. The sheet answers the same right press the two windows do, and which window moves is worked out from where your hand is - the upper window for a press on the upper half of the chart, the lower one for the lower. Your own card is left alone: a right press on the card, or on any of its buttons, moves nothing, so the day cannot slide about while you are reading the card that describes it. The left button is untouched and still opens and works cards. Measured live with a batch card open: a right press 60 pixels to the left moved the modules' window from 200 to 260, the same press on the people's window moved that one from 260 to 310, the two windows came back into step at 310 on the next frame, and the same right press on the card itself moved nothing.

5. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before. Your stored data was put back from a copy taken before the measurements and compared: byte for byte the same, and your three saved days still read 24, 4 and 24 pans.

6. And the rule that stands on it was proved rather than assumed. Nine faults were put back in all and each was watched failing by name, then restored byte-identically: with the putting-back left out entirely the test fails reading that the press threw the windows back to the start; with only the modules' window put back it fails the same way; with the two windows collapsed into one number it fails; with the putting-back done before the measuring it fails reading that the shelf was capped short of where the fourth day ends; with the class name matched dot and all it fails (this is the fault that was in the first draft of the fix, and the test caught it); with the sheet never wired it fails reading that a right press over an open card did not move the day; with the sheet always panning the upper window it fails on the lower press; with your own card not spared it fails reading that a right press on the card dragged the day behind it; and with the left button let through it fails the older drag test, which says the left press is the one that opens cards. One more thing was found and fixed on the way: the tests' own stand-in screen did not move a box when a box it sits inside was scrolled, which is the opposite of what a window does, and that gap had been hiding whether the measuring happened before or after the sliding. The suite is 1220 passing with none failing.

7. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v174 (no database step). Two things you asked for: the modules' window no longer lets the day's writing come up between the module titles when you pan it right, and the Your scenarios card now shows four days at a time with the rest reached by sliding the list. Both inside More -> Scenario planner.

What you asked, in your own words. "the windows when panning right, some of the chart writing shown in between module titles" - and "now i had 8 scenario, we have to make the secenario just shown 4, the rest shown by slider". You were right about both, and neither was your phone.

1. The writing between the titles, and why panning right brought it there. The module titles on the left of each window are pinned: they stay put while the day slides under them, which is the whole reason you can read which row is which at any hour. What was wrong was the ORDER they were stacked in. Every B-number, every band of your hands and every line down the bars was drawn at or above the title's own level, and when two things sit at the same level the one drawn later wins - and the day is always drawn after the titles. So nothing had to go wrong for a number to appear over a name; the day simply took its turn on top of it. Measured live on your own day at a phone's width, with the chart panned 25 pixels: nine B-numbers and six bands of your hands were the topmost thing at their own position inside the 156-pixel title column.

2. What it does now. The three things that can be stacked are written down once, on the chart itself, as three levels: the titles lowest, the clock strip above them, and the readings you take highest. The clock strip has to outrank the titles because the titles scroll under it, and a reading has to outrank the clock strip because the reading's own top line lands on the strip. Everything the DAY draws - its marks on your bars, the B-numbers and the names on them, and a module's own note - is deliberately under the titles, so no pan position can bring one of them up over a name. Measured live on your own day across six pan positions, 1,400 points sampled inside

the title column: not one of them reads as the day, while the same eight B-numbers and five bands are still drawn exactly where they were, under the titles where they belong.

3. And the clock you take is refused where the day is not in front of you. The same fault had a second half. Pressing the clock strip over the title column could still read a minute and draw the hairline and its number across the titles - because the strip's own ruler scrolls with the day, so once you have panned, the strip's left edge has travelled behind the titles and the arithmetic no longer knows where the titles end. It now measures where the day actually starts rather than trusting the strip's edge. Measured live with the day panned 25 pixels: pressing at 140 pixels across - inside the titles, which end at 183 - reads nothing and draws nothing, while the same press 17 pixels further on reads 4:26 am as it always did. On a day that is not panned nothing about this changes at all.

4. The shelf now shows four days and slides to the rest. On your eighth saved day the card grew a row per day, and a shelf of eight days plus the buttons under it ran past the bottom of your phone. It is four days at a time now, with the list itself sliding up and down for the rest, and a line under it that says so: measured live with eight days, "Showing the first 4 of your 8 days. Slide the list itself up for the rest." The headings stay put as the list slides, so a row you have scrolled to is still under the heading that names it. A card with four days or fewer is exactly the card it was before - no scroller, no height, no note.

The height is the part worth being careful about, and it is measured rather than assumed. A long day's name wraps to a second line at a phone's width, so the fourth row is not always as tall as the first three, and a shelf capped by arithmetic would have cut the bottom off the fourth day. Measured live on the eight days: three rows of 59 pixels, then a fourth of 73 because its name took three lines, and the shelf stops at exactly 290 pixels - the fourth row's own foot. The fifth row begins at 290, so nothing of it shows. That is the reading this was built from: capped by arithmetic it would have been 236 pixels and the fourth day would have been cut in half.

5. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before. Your stored data was compared before and after every measurement: 13,888 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans. The eight-day card used for that measurement was put on a copy of your data and taken off again, with your own file restored byte for byte and nothing left behind.

6. Every rule that stands on it was proved load-bearing, not assumed. Eight faults were put back in, one at a time, and each was watched failing by name, then restored byte-identically. The ladder: with a B-number put back above the titles the test fails reading that panning right puts the day's writing between the module titles; with the title cell taken off its own level it fails reading that the levels can no longer be moved in one place. The reading: with the guard taken out it fails reading that a reading was taken with the pointer over the module titles. The shelf, four ways: capping it off the last day instead of the fourth fails reading where the fourth day ends; writing no height at all fails the same way; measuring it as a multiple of one row instead of reading the fourth row's foot fails reading 40 pixels where the fourth day ends at 220; and taking the slide away fails reading that the days past the fourth are unreachable. The suite is 1217 passing with none failing.

7. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v173 (no database step). Two faults in the "Minutes from one batch to the next" box, both on a day that runs more than one batch: on a module that runs only one, the box did nothing at all and said nothing about why; and on a module whose batch two had been nudged closer by hand, a smaller pace could send that batch in front of batch one, and the day then drew both on the same minute. All inside More -> Scenario planner.

What you reported, in your own words. "similar problem happen. 1st module, time between each batch, set to 21, but no effect, or the chart dont workout" - and then, when asked which of the two you were seeing, "both batch start the same time". You were reading a 12-pan day, and the first module of that day is "Mixing by hand in the tub (set the minutes)": one batch, and the same pace it has always had. Both halves of your report were real, and neither was your phone.

1. On a module that runs one batch, the box does nothing - and now says so. A pace is the minutes

between one batch and the next. A module that runs a single batch in the day has no second batch for a pace to sit between, so whatever you type there cannot move anything. That is what your first module is: measured on your 12-pan day, "Mixing by hand in the tub" runs one batch of four pans at a pace of 20 minutes, so a 21 typed into it had nothing to act on and the chart was already telling the truth. The fault was that the box said nothing at all about that - it looked exactly like a box that was broken. It now explains itself in your own words: measured live on that very module, the line under the box reads "This module runs one batch in the day, so there is no second batch for a pace to sit between yet. How many batches it runs in the day is the box above that sets it - make that two and these are the minutes between them." Make it two and the line goes.

2. And on a module that does run more than one, a smaller pace could put batch two in front of batch one. This is the "both batch start the same time" half, and it is the more serious of the two.

A saved day remembers the exact minute each batch starts, one module at a time. A module that carries that list is re-spaced at the pace you type, so a batch you have nudged by hand keeps the minutes you nudged it by rather than being pulled back onto the rhythm. That is right for a batch you nudged LATER - and exactly wrong for one you nudged EARLIER, because its deviation from the rhythm is a negative number, and adding a negative to the smaller rhythm lands it in front of the batch above it.

Measured on your own first module, in the shape that does it. A nudge of batch two by hand with the box still reading 87 leaves the module carrying the times 4:00 am and 4:21 am - and that batch is then 66 minutes earlier than the rhythm, so its deviation is -66. Typing 21 sent it to minute -45: three quarters of an hour before your day begins. The day cannot draw a batch that starts before it does, so both batches were drawn on minute 0 - both at the same time, exactly your words. This is also where the "no effect" reading of the chart comes from on a two-batch module: the box reads 21, and the two bars sit on top of each other and will not come apart.

3. What it does now, and it is your own app's rule rather than a new one. Two rules, and both were already the app's own. A hold is only ever a HOLD - every other press on this screen already obeys that, so a batch can be held back off the rhythm and can never be pulled in front of it. And no batch is left at or behind the batch above it: where a hold has become impossible in the smaller rhythm, the batch takes the rhythm rather than a minute that is not its own. Measured live on that same module, through the same card: typing 21 leaves the two batches at 4:00 am and 4:21 am, drawn 21 minutes apart with the box reading 21, and the second batch is no longer able to overtake the first.

4. And a third fault, found while proving the second: the Auto box did not save. This one you had not reported and would have met eventually. "Minutes from one batch to the next" carries an empty box with the word Auto in it, meaning each batch starts the moment the one before it ends. Emptying the box moved the flag and nothing else.

The zero that means Auto was being written through the same writer every other box on this screen uses; that writer refuses any number below the box's own minimum; and this box's minimum is 1 - so the refusal came back BEFORE the save and before the repaint. The choice was made in memory, nothing on the screen moved, and the next reload threw it away. Measured on your own mix with two batches: emptying the box left both batches 30 minutes apart and the screen did not move at all; the card reopened with its empty Auto box drawn over the very same unchanged chart; and a reload forgot the choice altogether. The zero and the save are now written directly, so Auto is an answer the app keeps: measured live, emptying the box redrew the two batches 20 minutes apart - the module's own cycle, which is what Auto means - and a reload kept it.

5. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before. Your stored data was compared before and after every measurement: 13,888 characters, byte for byte the same, and your three saved days still reading 24, 12 and 24 pans.

6. Every rule that stands on it was proved load-bearing, not assumed. Six faults were put back in all and each was watched failing by name, then restored byte-identically. The hold rule: with the clamp taken out the test fails reading that a batch pulled closer than its line carried a hold it does not have. The forward sweep: with the sweep removed the test fails reading that a batch was left in front of the batch above it. The one-batch note: with

it gated off the test fails reading that a pace box with nothing to space does not say so. The Auto half, three ways: with the 0 written through the writer again the test fails reading that Auto did not write the 0 that means Auto; with the save and the repaint left out it fails reading that the Auto press saved and repainted nothing; and with the list not deleted it fails reading that Auto left a list of times behind, so the cycles no longer decide. One fault in the tests' own stand-in screen was found and fixed on the way - it had no firstChild, so the "did the screen repaint" assertion was comparing nothing with nothing and could never fail - which is the same lesson that file already carries about a stand-in being more forgiving than the real browser. The suite is 1215 passing with none failing.

7. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v172 (no database step). Four things: handing one stretch of work to somebody else no longer moves the people's rows, their names or their jobs; the "Minutes from one batch to the next" box reaches your batches again, and says so when your oven is what is holding them apart; the Your scenarios card is a table you can run your eye down; and the question a card asks before deleting something is drawn over the card that asked it instead of behind it, with the row of buttons at the foot of a card given the shape it never had. All inside More -> Scenario planner.

What you reported, in your own words. First, and the largest thing in this version: "i have 4 persons, 1st Jien, 2nd Wei, 3rd and 4th. I saw Jien is heavy loaded, so i click one of Jien session, and select to switch that session to Wei, where he is free. BUt what happen is Jien disapper, and Jien name chage to Wei, and the originally Wei sessions disappeared". Then: "im now on my 36loaf scenario, pls check why im lowering the minutes from one batch to the next, cannot reflex even i reduce it to 13, look like something prevent it from loweriing below 25". And, separately: "can you make the Your Scenario, reorganised it to a better visual clarity, details in table, action button below?" And then, while that was being built: "please check the process of delete a module and the process of deleting a scenario, it dont look right".

1. Handing one stretch of work to somebody else moved everybody on the chart. This was your report, and it was the largest thing in this version.

All three of the things you saw were real, they were one fault, and none of it was your doing.

A day that is sharing them out has nobody written on it. The rows you see are worked out fresh every time the chart is drawn: the first job nobody is on goes to person 1, the next to person 2, and a new row appears only when everybody is busy. That is why a shared-out day reads as perfectly sensible people - and it is also why its row numbers mean nothing at all. They are not yours and they are not on the modules.

Your names, though, are kept by number, and every day you have shares them. So the moment one stretch was given to Wei, a row for Wei existed before the free jobs were handed out - the packing began at Wei's row instead of at the first row, every invented number moved up by one, and the row that had been person 1 became person 2. Jien was person 1, so Jien was gone; the row he had been on now wore Wei's name; and a fifth row appeared at the top with what was left over. Exactly the three things you saw.

2. What that press does now, and what it was measured doing. Before a stretch changes hands, the arrangement already on the screen is written down: each job is given the person the chart is already showing at it - one name for a module whose batches all sit on one row, one name per line for a module worked as lines, and a note on the single stretch where a module's batches really are split between two people. A stretch you had already placed by hand, including one you had handed back to whoever is free, is left exactly as it is. From that moment the numbers on the rows are facts rather than the packing's invention, and your names go on meaning what you see.

Measured live on your own day at a phone's width, going through your exact press. The card read "Move this slot off Jien" and its button read "Move it to Wei". Before the press the chart drew four rows - Jien with five jobs on him, Wei with one, and two more below them. After it, the same four rows were in the same order with the same names: Jien had lost the one stretch you moved and nothing else, Wei had gained that one stretch and nothing else, and the two rows below were untouched. The People box stopped saying the day was sharing them out and read "One to a module", which is what the card had told you it would do before you pressed it. And

the chart held there: the page was read again from scratch and drew the same four rows, the same names and the same jobs - where before this fix that very re-read was what moved them.

3. And one sentence on that card was describing a day that was not in front of you. Found by measuring the fix on your own day rather than on a made-up one. The card carried a note explaining that the press settles the whole day, and it said so in the words "This day is sharing them out" - on a day whose People box plainly read "Your own", because one stretch on it had been placed by hand. Two claims, about a day that was making neither. The note may now say only what is true of the day on the screen: the part that says the press settles the whole day is always said, and the part about the box stopping saying it is sharing them out is said only when the box is saying it.

4. The box that would not move, found and fixed. Two different things were true at once, and only one of them was a fault.

The fault: a saved day keeps, for each module, a list of the times its batches start. A list is read as the whole truth about where every batch sits, so a module that carries one ignores the box entirely - the box could be set to 13, or to 87, or to anything, and the batches would not move. Measured on the engine, not reasoned about: a module with the list 0, 25, 50, 75 and the box at 13 still ran 25 minutes apart; set to 87 it still ran 25 minutes apart; with no list at all, 13 ran 13.

How the list got there: the press that moves one batch's time wrote the module's whole list back onto it, and so did the day card's "work the day backwards", which does it for every module in the day at once. So one press of one button used to freeze the pace box of the entire line - which is why it looked like something was preventing it, rather than like a box that had quietly stopped working.

5. What changed. Three things, all small. A press that moves a single batch no longer writes a list onto a module that never had one - batch one of a module IS the module's start time, so writing that one number is the same answer the pace already gives, and leaving no list behind is what keeps the pace alive. A module that already carries a list has it re-spaced to the pace you type, with every batch keeping what you did to it by hand: each batch moves by the minutes the rhythm moved, so a batch you deliberately held off a collision is still held off by the same minutes after. (23 Sep 2026, added at v173: this re-spacing was unbounded in v172, and a batch that had been nudged EARLIER by hand carried a negative deviation that the smaller pace then added on, putting that batch in front of the one above it - the day drew both on the same minute. See the v173 entry above: a hold is now only ever a hold, and no batch is left at or behind the batch above it.) And the box now carries a live note under it, so a pace the oven cannot reach says so where you typed it instead of doing nothing.

6. The other half of it, which is not a fault and now says so. A module that takes one batch at a time cannot start two of them closer together than one batch lasts - that is your oven, not the app. Measured on your own "Mixing the dough in the tub": its cycle is 20 minutes, so typing 13 still runs the batches 20 minutes apart, and the note under the box says exactly that, names the 20 minutes, and names both ways to run closer: give the module a second production line under "How many production line you have", or switch on "Allow multiple production line" if those minutes are the dough's time and not the machine's. Measured live on your own frozen day: typing 13 turned that module's stored list from 1, 88, 175, 262 into 1, 14, 27, 40, and the day then ran 1, 21, 41, 61 - 20 minutes apart, its own cycle. Typing 40 ran the day 1, 41, 81, 121, and the note went back to being a plain explanation.

7. The Your scenarios card is now a table. One row for each saved day: the day's name, the pans a day it turns out, and how many modules it has, with the pencil at the end of its own row and Save changes and Save a copy below the table rather than mixed into it. The day the planner is currently plotting is filled in and carries a line down its left edge, and every name starts at the same place so the days read as a column.

8. And one measurement that made the table worth having. At a phone's width the card has 315 pixels to spend, and two headings were taking 170 of them - "PANS A DAY" on one line claimed 91 pixels to sit over a two-digit number and "MODULES" claimed 78 to sit over a one-digit one - which left the name 92 pixels. Every name you have is wider than 92, so every row wrapped, each to a different number of lines, and the card read as three ragged paragraphs. The heading is written over two lines now instead, which hands its slack back to the

name: measured at 375 by 812 the name goes from 92 pixels to 138, all three rows come back to the same 59-pixel height instead of 59, 72 and 59, and every name starts at the same x. The mark that said which day was open used to be a dot typed in front of its name, which pushed that one name 17 pixels right of the other two; it is a line down the row's left edge now, which costs the row nothing.

9. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before - a row opens that day, and the pencil beside it opens the rename card without opening the day as well. Your stored data was compared before and after: 13,888 characters, byte for byte the same. The one press in this version that does write to your day is the hand-over above, and it writes only what the chart was already showing - the same people, on the same rows, at the same jobs. Its own round trip was measured on your day, starting from your day as it stood and putting it back byte for byte afterwards: 13,888 characters before it and the same 13,888 after it. Your three saved days are untouched, and each still asks for the pans a day it always asked for - 24, 12 and 24 - and still carries the same 9, 8 and 8 modules. One reading in the table is worth naming, because it is not the same number: "Pans a day" in the table is what each day's line turns out, which is exactly the number the app already puts in a day's own summary line. For two of your days that is the 24 you asked for; for "My sister proposal 21/9/2026" the one tub passes 4 against the 12 you asked it for, and the shelf says 4 for the same reason its summary line already does - so the two screens cannot disagree.

10. The delete button opened a question you could not see. This was the fault you found, and it was not a wording problem or a small one: the question was being drawn underneath the card that asked it.

Both of these deletes are asked from inside a card - Delete this module from the module card, Delete this scenario from the card the pencil opens - and the card stays on the screen behind the question, which is deliberate, so you can see which module you are about to lose. What was wrong is the order the two are drawn in: the layer the question goes on sat at 60 and the layer of a card sat at 65, so the card was painted over the question every time.

Measured on your own day at a phone's width, not reasoned about. With the question open, the topmost thing at the question's own Delete button was an input belonging to the module card underneath, and the topmost thing at the question's own centre was a label from that same card. So the sentence could not be read, neither button could be pressed, and the only thing you saw was the screen dimming a second time with nothing appearing on it. Tapping anywhere did what the card underneath does instead, and the question sat there unanswered until something else happened to close the card, at which point it surfaced on its own - a delete box appearing when you had stopped asking for one. Both flows were the same fault, and both are measured fixed: at the question's Delete button the topmost thing is now the Delete button, at Cancel it is Cancel, and at the centre of the box it is the box.

The question's layer now sits at 72, above every card and above a toast, and only the app-password lock outranks it. Nothing about the flow changed otherwise: Cancel leaves the module and leaves the card open where it was, and Delete takes the module out and closes the card with it - both driven and watched on a live day, a module going from eight on the line to seven with the day answering without it, and a saved scenario going from three to two with the shelf redrawn and no name left standing on the card. And the question now names the module in quotes, as every other delete in the app names what it is about. Unquoted, a module called "The proofer again" read as one sentence with the name lost inside it: "Delete The proofer again?".

11. And the row of buttons at the foot of a card had no shape at all. The same report's other half, and a smaller thing, but the same kind of thing: a class used in 26 places across 12 files that never had a rule of its own.

Without one, its buttons flowed as plain inline blocks - flush against each other and hard against the left edge of the card. Measured on the module card at 375 by 812: "Delete this module" and "Done" were 0 pixels apart, so a thumb reaching for Done was touching the button that deletes the module, with 88 pixels of dead card to their right. The scenario card was the same shape the other way up - its Delete stacked directly above its Save at exactly the same left edge - and both read as two buttons run together rather than as a row of choices.

The row is now one rule, written once: the card's own edge, 8 pixels between the buttons, and a line's worth of air above them. wrap is part of it rather than a detail, because the module's three buttons need 437 pixels and have 303 - a row too wide for its card has to go onto a second line, not off the card's edge. It is the shape the label printer's own sheet already carried in full and the shape three other cards had each patched in by hand, a little differently every time; those three patches are gone, because one answer in one place is what stops a fourth being added. Measured after, on the same card: Duplicate this module on a line of its own ending at the card's edge, then Delete this module from 94 to 264 and Done from 272 to 345 - 8 pixels apart, Done sitting where a thumb expects it at the card's own right edge.

12. Every rule that stands on it was proved load-bearing, not assumed. Nineteen faults were put back in all, and each was watched failing by its own name before being restored byte-identically.

For the hand-over: with the arrangement left unwritten the test fails reading that the day has stopped being able to say what it is, which is the fault put back; with a person written onto a module without the crew that carries it, the same test fails with your own three symptoms - five rows where there were four, renumbered 2, 3, 4 and 5; with the People box allowed to read a crew of zeroes as an arrangement, the hand-back test fails reading that a day which has not moved no longer says it is sharing them out; and with the card's note gated off, it fails reading that the card does not say what the press settles.

For the pace: with the re-spacing taken out the test fails reading "the stored list did not follow the pace she typed"; with the note about the oven taken out, "a pace the machine cannot reach is not explained"; with every batch pulled back onto the rhythm instead of keeping its own offset, "a batch held off by hand was pulled back onto the rhythm"; and with the press left writing a list again, "Mixing the dough in the tub came out of the press carrying a list of times, so its pace box is dead".

For the question: with the layer put back under the card it fails reading "under the card that asked it"; with it put above the app-password lock, "drawn over the app-password lock"; with the module's name left unquoted, "does not name the module in quotes"; with the destructive button put first, "answered by the safe button and then the destructive one"; with Cancel made to run the delete, "Cancel deleted the module"; and with the card left open behind the delete, "the card stayed open after the module went".

For the foot row: with the rule declared a second time it fails reading "declared 2 times"; with the buttons put back to inline blocks, "flow as inline blocks"; with the 8-pixel gap taken away, "sits flush against"; with the row back at the left edge, "not laid out against the card's own edge"; with wrap taken away, "cannot wrap"; with one card patching its own row by hand again, "lays its own foot row out by hand"; and with the foot's buttons reordered, "does not carry its three buttons". The two faults found while measuring the note on your own day are put back the same way: with the note always claiming the day is sharing them out it fails reading that the card claims a box that is not saying it, and with the note taken off the card altogether it fails reading that the card does not say what the press settles. The test suite is 1215 passing with none failing.

13. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v171 (no database step). Three things you asked for while looking at the day chart: a right press and hold moves a window's day about without opening anything, the clock balloon rides two inches clear of your pointer, and the batch card is brief. All inside More -> Scenario planner.

What you asked, in your own words. First: "And can i drag the module window up/down, left/right by right click and hold? Dont let this action open up the card". Then, while that was being built: "put the clock balon 2 inches higher than. cursor". And: "the batch pop up, make it as brief as possible".

1. A right press and hold now takes hold of a window's day and moves it, both ways at once. Press the right button anywhere on either chart window, keep it down and move: the day comes with your hand, sideways and up and down together, exactly as if you had taken hold of the paper and pulled it. The window under your hand moves and the other one comes with it sideways, so the two can never disagree about a minute - that part was already there and is untouched. Measured live at a phone's width (375 × 812): a right press on the modules' window and a 60-pixel drag moved its day from 0 to 60, and the people's window to the same 60. The window marks itself while your hand is down and unmarks the moment you let go.

And it opens nothing, which was the other half of your sentence. The gesture is the RIGHT button, and a browser fires no ordinary click at all for a right press, so a press for panning can never be read as a press for opening. The ruler's own press, which drags the clock reading, now ignores the right button too - one press, one gesture. Proven live, not assumed: a right press on a bar dragged the day and left the card shut, and a left click on that very same bar opened it. Your own hand is on the safe side of it either way - a finger cannot set the right button at all, so nothing here is reachable from your phone and the way you touch the chart is exactly as it was. The browser's own menu, which a right press would otherwise put on the screen in the middle of the gesture, is stopped over the chart.

2. The clock balloon rides two inches clear of your pointer. It used to sit 8 pixels above your finger, which on a phone is under it. It is lifted two CSS inches now - 192 pixels, and named in the unit you used so the number is one this can be held to rather than a feeling. Above your pointer wherever the window has the room, and where it does not, the same two inches BELOW it - which is the side that matters, because the clock strip you press with a finger is at the top of the modules' window, and a balloon that came to rest at the window's top would be lying on your hand. Measured live at 375 × 812: the full 192 pixels of clearance in both directions, and never nearer than 175 pixels anywhere on the chart. It is still held inside the window you are looking through rather than to the day, so a tall day cannot carry it off the screen.

3. The batch card is brief, and its own reading is no longer cut off. What left the card: the paragraph under each pair explaining what a move means in that module's own terms, and the second line that repeated the two times the line above it already reads out. What is on it now: one line of four readings - which batch, how it sits against the line, the two times, and what it costs your hands - then the two pairs, five minutes and one minute, and the way back where there is one. Measured at a phone's width, the card came down from 338-366 pixels tall to 265-288. The explanations are still true and are still there, on the press itself as its accessible name, and every press still answers in words when it lands.

And one fault found while measuring that card, fixed here. Four readings on one line ran off the card's right edge - "Batch 2 on the line 8:51 am -> 8:5" - with the end time, the clock you opened the card for, cut off. A reading is set never to be squeezed, on purpose, so a time can never be shrunk to nothing; a line one word too wide therefore overflows rather than shrinks. The card's readings now WRAP to a second line instead, which only ever happens where the line would otherwise be clipped. The same fault was possible on every card that shares that line, and it is fixed on all of them by the one change.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before. Your stored data was compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

Three rules now stand on it, and every one was proved rather than assumed. One says a right press and hold drags a window both ways and that a right press over a bar opens nothing, while a left press on the same bar still does. One measures the balloon against the palest bar and pins the two inches of clearance, above or below, with the reason the window's own top is not good enough written into the test. One pins that a card's readings wrap and that they still refuse to be squeezed - the two settings asserted together, so neither can be changed alone. Each was proved by putting the fault back and watching it fail by name, then restored byte for byte: no downward fallback for the balloon read "a balloon with no room above her pointer sits -50px below it, not the two inches she asked for"; no upward preference read "the balloon sits 232px above her pointer, not the two inches she asked for"; and removing the wrap read "a card's readings are back on one line, so a row too wide for the card is clipped at its edge again". The test suite is 1209 passing with none failing.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

CHANGELOG and the guide both carry all of this. Both PDFs rebuilt and read back with PyMuPDF, the app's own More screen reads Engine v171, and your phone is untouched by nothing here.

23 Sep 2026 - engine v170 (no database step). The frame around every bar is drawn quietly now, so a block of work still reads as a block without every bar being a box drawn round it. Nothing else moved.

All inside More -> Scenario planner.

What you reported, in your own words. "the border look like over emphasized". You were right, and it was my own v169 line being too loud. Nothing about the day, the bars or your data was wrong - the frame was simply drawn heavier than it needed to be.

What was too much. The frame went in at a line one and a half pixels wide, drawn at full strength in the bar's own solid colour. Round a pale bar that is not an edge, it is a box drawn on: every bar on the chart was outlined that way, so the day read as a page of boxes rather than as a day with edges on it.

What changed, and it is only the frame's own two numbers. The line is one pixel wide rather than one and a half, and the whole frame is laid down at 55 parts in 100 rather than at full strength. The colour is untouched and is still the bar's own tone - the colour your hands already wear inside that bar, and, for a person, the colour of their card's stripe - so the frame still says at a glance which module a bar belongs to. Measured live on your own day at a phone's width: the frame reads the bar's own colour, a 1-pixel line with the layer at 0.55, on a batch bar, on a lane bar that shares its row, and on a person's time slot alike.

Your B-numbers and your people's names are untouched, and no line crosses a label, exactly as at v169: the frame is still drawn below them.

One number softens every frame together. The strength is declared once, on the chart's own wrapper beside the marker's transparency, so the modules' window and the people's window and all sixteen tones are all softened by the one edit and cannot be tuned apart. Lower it towards 0 for a whisper, raise it towards 1 for v169's line back again - it is one number either way, so say the word if this is still not the weight you want.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before, and the frame still takes no taps at all. Your stored numbers were compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

And the rule that stands on it was extended rather than left standing on the old numbers. The frame test now carries your report in its own name, pins the 1-pixel line, pins the layer's strength and refuses a strength of 0 or 1 - so a later edit cannot quietly put the shout back - and refuses a second declaration of it. Each was proved by putting the fault back and watching it fail by name: the 1.5-pixel line back made it read "the frame is no longer a 1px line"; full strength back made it read "--frame is 1: one at 1 is v169's shout back again"; and the strength declared twice made it read "the frame's strength is declared in more than one place, so the two windows can disagree". Each was restored exactly as it was afterwards. The tests are 1,201 passing with none failing.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v169 (no database step). Every batch bar and every person's time slot now wears a frame in its own colour, so a block of work reads as a block and not just as a patch of colour on the paper. And the clock's lines have gone to the back of the day while every bar became see-through, so the ruler reaches inside them the way you asked instead of sitting on top of them. All inside More -> Scenario planner.

What you asked, in your own words. First: "Can the batch and person time slot having a more highlighted frame?" Then, while that was being built: "its scale should be at background, the marker should be ontop, with some degree of transparency". Those two together are this release, and the second one turns round the way the ruler was drawn at v165 and v166.

Every marker now wears a frame, and it is drawn in its own colour. A batch bar and a person's time slot are the very same shape in the app - one rule frames both, so a batch and a slot cannot end up looking like two different things. The frame is the bar's own solid colour: the same colour the hands inside that bar already wear, and, for a person, the same colour as the stripe down the side of their card. That matters more than it sounds - a frame in a colour of its own would be one more thing to keep in step, and it would drift the first time a tone changed. Nothing about the bar's shape moved for this: a batch bar is still 16 pixels tall and a person's slot still 11, measured, and the frame is a line laid inside the bar's edge rather than a border added round it. A border was tried and thrown away - it would have squeezed the inside of the bar, which is where your hands and the

cycle shades are placed, and an inner shadow was no better because those same shades are painted over it.

The ruler is at the back and the markers are over it, as you asked. Since v165 the lines were drawn as a layer of their own on top of everything. They are now the paper's own ruling again - the surface the bars are drawn on - and every bar is a wash over it rather than a solid block. The bar is 72 parts in 100 opaque, which is the "some degree of transparency" you asked for: enough that the clock reads through it, not so much that the day looks faded.

What that cost, honestly, and what was done about it. A line seen through a bar loses about a quarter of its strength, because the same wash covers both the line and the paper beside it. Left alone, the ruler would have gone faint again - and faint is the fault you reported twice. So both lines were drawn heavier: the hour line at 60 parts in 100 of full strength where it was 34, and the finer grid at 45 where it was 20. Measured through one of your batch bars, the hour line now differs from the bar beside it by 29 parts in 255 and the grid by 22, against the 20 that your own v166 rule demands. The honest other side of that: on the bare paper between the bars the ruling now reads about twice as strong as it did at v168. That is the price of drawing it behind the markers, and it was paid deliberately.

And no line crosses a label. Your B-numbers and the name of the person on a stretch of their day are laid above the frame, so thickening the frame cannot put a line through a number you have to read.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before - a batch its Batch card, a cycle its own card, a person's stretch the hand-over card; the frame takes no taps at all. Your stored numbers were compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

Three rules now stand on it, and every one was proved rather than assumed. One says the ruling is the paper's own background and that no second copy is drawn over the markers - two grids, one in the gaps and one across a bar, is how a reader ends up seeing two different grids. It also pins the transparency to one number declared once, so the modules window and the people's window can never be tuned apart. One says every batch bar and every person's slot wears a frame in the tone its own hands or card already wear, and that no bar grew a border or a shadow of its own. One is the arithmetic above: measured, not trusted, that both lines still differ from a marker by at least 20 parts in 255. Each was proved by putting the fault back and watching it fail by name, then restored exactly as it was. Five faults were put back in all, and one of them is the point of the whole release: with the grid's ink back at v166's strength, the test fails reading "differs from the marker beside it by only 9.8 parts in 255, which is invisible on a phone - the fault she reported twice at v165 and v166". The tests are 1,201 passing with none failing.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

Softened at v170: the frame described above went in at a line one and a half pixels wide and at full strength, and you reported it as "the border look like over emphasized". It is now a one-pixel line laid down at 55 parts in 100, in the same colour as before. Read the v170 entry above for the rest.

23 Sep 2026 - engine v168 (no database step). The extra slider between your two chart windows is gone, so each window's own sideways bar is the pan. The modules window is taller. The "People at once" row is out of the people's window, and an overlap is still shown in red. All inside More -> Scenario planner.

What you asked, in your own words. First: "since the both windows have their own slider, additional slider is redundant. Remove that" - so the slider that sat between the two windows is removed. What you asked for when the two windows were built still stands and is untouched: the two windows stay in step, so any minute is in the same place on both, whichever one you drag. Then: "can you increase the module window further?" - the modules window is taller. And: "im thinking of remove the people at once?", settled by "you can just show when there is overlapping highlight in red box, like previously have" - the row is out, and the overlap is still shown in red.

The extra slider is gone, and the two windows still move as one. Measured on your day at a phone's width, 375 by 812: nothing sits between the two windows any more but their own two edges - the chart is the modules'

window and the people's window, in that order, with nothing in between. Dragging one window's day sideways moved the other to the same pixel, and the window under your finger took no write back at all, so a drag is never fought by the window that is following it. On the two windows being able to reach the same place: they could not quite, and it was one pixel. The ruler's last hour sits at the very end of the day, and its line was being drawn one pixel past it, which made the modules window 2,461 pixels wide against the people's 2,460. Both windows now read 2,460 and both pan to 2,139. That last hour's name is also written to the left of its own line now, where a ruler with an end on it has always put it, and it reads in full inside the day - measured, all 25 of the ruler's names sit inside the day, the last of them ending 3 pixels short of its end.

The modules window is taller, and here is exactly what that buys. It takes the larger half of the chart's ceiling now: 64 of the 84 parts where it had 50 of the 80, and 590 pixels where it had 470. On your phone at 812 tall this changes nothing you will see - your day's modules already fitted the window exactly (403 pixels of modules in a window that was already wide enough for them), so the taller window is room there rather than a difference on your screen. Where it shows is a shorter screen: at 667 tall the old window held 333 pixels and made your day scroll, while the new one holds 427, so your whole day sits in front of you with nothing to scroll. It is also the room the next two module rows will need as you add them.

The "People at once" row is out of the people's window. You were right that the day already says it without that row. An overlap is still shown in red, in the places that were already saying it: the red outline on the colliding bar in the people's window, the colliding stretch of that person's own row, the red line in their card, and the collisions written out in words under the diagram - "two jobs at once", or how many. Nothing about an overlap is hidden by the row going, and it is one line to bring back when there is a second pair of hands to weigh up.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. No bar changed colour, and every tap still opens exactly what it opened before - a batch its Batch card, a cycle its own card, a person's stretch the hand-over card. Your stored data was compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

Four rules now stand on it, each proved load-bearing rather than assumed. The ones from v167 still hold, with the slider's part of them narrowed to what is left: the two windows are adjacent with nothing between them, and a drag on either moves the other while the window under the finger takes no write back. Three are new. One asserts that the "People at once" row leaves nothing on the chart while a collision is still marked in red. One asserts the ruler's last name is drawn inside the day, and that the ruler's own track clips, so the two windows cannot disagree by even a single pixel. One asserts the modules window holds the larger half of the chart's ceiling, 64 of the 84. Every one was proved by putting the fault back and watching it fail by name, then restored byte-identically.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v167 (no database step). The day chart is now two windows sharing one horizontal slider - the process above, the people below, one ruler's lines down both - and the blank space under your people's rows is gone. This release also carries two smaller faults you reported while it was being built.

What you asked for, in your own words. First: "the window for person stay over size", which you confirmed means the empty space under the people's rows. Measured at a phone's width, 375 by 812: the chart was one window, fixed at 650 pixels tall, the day's own content was 473, and your people's rows ended at 457 - so 177 pixels of empty window sat under them. Then your proposal, which was better than a tweak: "instead of consider then one window, why not create 2 windows, and let the ruler sync in the 2 windows. As the business grows, the persons will grows, processs might not, we need to have a better way to manage, what you say?" And on how the two should sit: "can the 2 window share the horizontal slider, place between the 2 windows, make the 2 windows as close as possible", with the slider going straight into the people's bars underneath it.

What is on the screen now. Two windows, one above the other, with a single slider lying between them and touching both - measured gap of 0 pixels either side. The upper window holds the clock and every module,

batch and cycle; it scrolls up and down on its own when you add more modules. The lower window holds your people's rows and nothing else, and it scrolls on its own too, so as your people grow the second window grows and the first one does not have to. The ruler's lines reach down both. The slider between them is the only horizontal control: drag it, and both windows move together; drag either window's own day sideways, and the other follows and the slider follows with it. Measured on your day at 375 by 812: the slider reads a travel of 2,139 pixels, a drag on one window moved the other to the same pixel, and the window under your finger received no write back at all - the echo is dropped rather than sent, which is what would make a phone stutter mid-drag. On a day short enough to fit, the slider dims and cannot be dragged, because a control you can grab that does nothing reads as a fault.

The blank space. It is gone. The chart is no longer forced to a fixed 650 pixels; it is now exactly what its two windows need. Measured on your day: upper window 403, slider 20, lower window 70, totalling 493, with your people's last row sitting exactly on the window's floor - 0 pixels of empty window under it, where there were 177. The height ceiling from v157 survives as a ceiling: the chart is never taller than it was, it is only allowed to be shorter when your day is shorter.

The first of the two smaller faults: a deleted saved day left its name on your card. In your words: "why the deleted scenario stil listed?" You were right, and it was reproduced rather than explained. Deleting "My sister proposal 21/9/2026" left the card still offering "+ My sister proposal 21/9/2026 · Mix by hand, no chiller", dressed exactly like a day still saved, while your storage held only two. The cause was that the card offered each ready-made day by name whenever that name was not in your list, so deleting one put the name straight back. Now the card offers one plain row, "+ Add a ready-made day", which names nothing; opening it lists whichever ready-made days you do not already have, and one tap puts that day on your own list, where it is yours to change like any other. Proven live: after deleting that day the card named it nowhere, and the chooser handed back exactly that one day.

The second: on a computer, a person's name card was cut. Hovering a person's name on the chart opens a card of what they do that day. In the new lower window that card was taller than the window it opened into - measured 180 pixels of card inside a 69-pixel window, 96 pixels of it below the window's own bottom, with no scroll position that could show it whole. It is now laid on the screen rather than inside the window, so it reads in full wherever the name is, and it is pulled back from the screen's edges rather than running off them. The "People at once" row's card, which sits in the same window, is placed by the very same function, so the two rows cannot behave differently. Measured on a computer at 1280 by 900: both cards fully on the screen, with 8 pixels to spare at the foot.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. No bar changed colour, and every tap still opens exactly what it opened before - a batch its Batch card, a cycle its own card, a person's stretch the hand-over card. Your stored data was compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

Three rules now stand on it, each proved load-bearing rather than assumed. One asserts the two windows are adjacent with the slider between them, and that a drag on either moves the other and the slider while the window under the finger takes no write back. One asserts that a ready-made day you have deleted leaves no name standing on your card, with the offer row itself excluded from the check. One asserts that both a person's card and the tally row's card are placed beside their name and pulled back from the screen's edges. Every one was proved by putting the fault back and watching it fail by name, then restored byte-identically.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

Added at v168: the single slider shared by the two windows, described above, has since been removed - each window now carries only its own sideways bar, and the two still move as one. Read the v168 note at the top of this file for what is on the screen now.

23 Sep 2026 - engine v166 (no database step). The ruler's lines are now drawn strongly enough to be seen on a phone - including across your batches, your cycles and your people's occupied slots.

What you reported. In your words: "At ver165, you still fail to reveal the ruler at person's time slot". You were right, and this release is that same fault finished rather than a new idea.

What v165 got right, and what it got wrong. v165 did draw the ruler over the bars - that part worked, and it is untouched. What it got wrong was the strength of the ink. The grid line was drawn at about one part in six of full strength, and over one of your slots it came out as a colour that differs from the slot by about ten parts in 255 - a hairline the eye cannot find on a phone in normal light. So the ruler was drawn, and it was invisible, which on a screen is the same as not being there. That is why you came back with the same words a second time.

What changed, and it is only the ink. Both lines are drawn in a strength measured against the bar they cross rather than against the cream paper between the bars - the hour line at about a third of full strength and the finer grid at about a fifth. The lines themselves are in the same places as before, at the same intervals for the scale you are reading at, and the ruler at the top of the chart still draws none of them over its own ticks. Nothing moved; the day simply reads as one ruled sheet now, so the edge of a marker has a line beside it wherever you look.

And a line never crosses a label. Your B-numbers on the bars, and the name of the person on a stretch of their day, are now laid over the ruling rather than under it. A gridline drawn through a B2 is a number you have to guess at, and the whole point of the ruling is to make the day easier to read, not harder.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. No bar changed colour, and every tap is untouched - the ruling takes no taps at all. Your stored numbers were compared before and after and are byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

About the height of the planner window, which you also raised. That is not in this release. You are right that it is a real thing, and what you proposed - two windows, one for the process and one for the people, sharing one ruler - is the honest fix rather than another tweak. It is written up as its own piece of work, so it is decided on its own rather than half-done here.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

Changed at v169: the ruler's lines are no longer drawn over the bars, as described above. The ruling is the paper's own background again and every bar is a see-through wash over it, so the clock reads through a marker rather than across it. The goal is the same one named here - the clock must hold inside a slot you are working - reached the other way round. Read the v169 note at the top of this file.

23 Sep 2026 - engine v165 (no database step). The ruler's lines are now drawn over your batches, your cycles and your people's occupied slots. Before this the lines stopped dead at the edge of every bar, so inside a slot you were working there was no ruler at all.

What you asked. In your words: "i requested a ruler to draw into person's line time slot, maybe you cant visual my complaint". Asked how the lines should reach that slot, you said: "Just like the ruler draw over the batches and cycles". You were right that I had not pictured it, so I went and measured instead of answering.

What was actually happening, and it was worse than "the lines are only on some rows". Your person's row did carry the ruler's lines - it has since v160. But every marker, batch and cycle is a solid coloured bar, and a solid bar sits ON TOP of the lines behind it. So the ruler vanished wherever a bar was. Measured on your own day: of one module row's ruler lines, 5 fell inside your batches; of the slots on one person's row, 14 fell inside them. Every one of the 14 was covered. Pointing at a ruler line inside one of your slots, the topmost thing on that pixel was the slot itself. That is why the ruler looked present on an empty stretch of the day and absent wherever you were actually working.

And it was the same on the module rows. So the lines were not reaching over your batches and cycles either - what you were seeing was them showing in the gaps between bars, with your eye joining them up. This release makes that true for real rather than only in the gaps.

What changed, and it is one thing. The ruler is now a layer of its own, drawn OVER the bars - over the batches, over the cycles and over your people's occupied slots alike. The lines themselves are the same lines as

before: the same hour line, the same step for the scale you are on, at the same positions. Nothing about the ruling moved; it is only no longer hidden. The whole day now reads as one ruled sheet, and a slot is ruled exactly like the empty day beside it.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. No bar changed colour, and every tap on a bar is untouched - a batch, a cycle and a person's own stretch are still tapped and still open the same cards. Your three saved days were read before and after: One baker day, No fridge 1 person and My sister proposal 21/9/2026 each still on the target they had, stored numbers untouched.

Two rules were proved load-bearing rather than assumed. One test now stands on that there is exactly one ruling - if the lines were ever drawn both under the bars and over them, the translucent grid would read darker in the gaps than across a bar. Another stands on that the ruling is drawn at a height above the bars, that it lets every tap through, and that the ruler's own strip at the top turns it off so it cannot put lines between the ticks you read the clock off. Each was proved by putting the fault back: removing the ruling's height made the first fail, and putting a second ruling back on the track made the other fail.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

Changed at v169: the lines described above are no longer drawn over the bars as a layer of their own. The ruling is the paper's own background again and every bar is a see-through wash over it, so the clock reads through a marker rather than across it. What has not changed is why: the lines must reach into a slot you are working, and the ruler's own strip at the top still draws none of them over its ticks. Read the v169 note at the top of this file.

23 Sep 2026 - engine v164 (no database step). A number you type into the cycles of a module now stays where you put it. Before this, typing into a second box put the first box back to what it was when the card opened, while the first box still showed the number you had typed.

What you asked. In your words: "why some parameter that i save can change by itself", and then, when I had guessed wrong about where: "it was on stretch and fold module, the cycle1,2 , loading, unloading. It is not like you explain". You were right and my first answer was not. I stopped explaining and went and did it on your own fold module until I could see the number move on its own.

What was really happening, and it is exactly where you said. Inside a module's card there is a list of cycles, and each cycle has its own boxes: the name, Minutes, Load, Unload. The card does not redraw itself while you type, and that is on purpose - it was your own instruction at v142, because a redraw would throw the box out from under your finger. But the code was keeping a copy of the whole cycle list from the moment the card opened, and every box wrote that old copy back. So the first box you used was still carrying the numbers as they were when the card opened, and the moment you touched a second box, the first box was put back to those old numbers. The screen kept showing what you had typed, the day behind it kept the old value, and the two disagreed. That is what "changing by itself" was.

Measured on your own day, so it is not a theory. Your fold module runs four cycles. Cycle 1's Minutes was 31; typing 40 into it stuck. Then touching cycle 2's Unload put cycle 1's Minutes straight back to 31, while the box under your finger still read 40. The same thing happened on one row: Load 5 stuck, then typing 44 into the Minutes beside it put Load back to 0 while its box still read 5. Renaming a cycle stuck only because nothing else was waiting to be undone.

The fix, and it is small. A box now changes only its own value, on the list as it stands at the moment you touch it, instead of writing back the copy from when the card opened. Nothing else in the card can be dragged back with it. The comment beside the code names your words and the exact numbers above, so the next person to touch it cannot put the copy back.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Your three saved days were read before and after and compared byte for byte: One baker day, No fridge 1 person and My sister proposal 21/9/2026 each still make what they made, their stored numbers untouched.

A test now stands on it. It opens your fold module's card, types into cycle 1's Minutes and then cycle 2's

Unload, and asserts the first number survives the second box - the exact fault. It also covers two boxes on the same row, and a cycle renamed after a number was typed. Proved load-bearing: with the fix taken back out it fails, reading "the second box she touched put cycle 1's Minutes back to where it was when the card opened".

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v163 (no database step). The clock above your people's rows has been taken back out. The clock is drawn once again, at the top of the chart where it has always been, and the ruler's lines still run down through your people's rows exactly as they did.

What you asked. In your words: "i dont want the clock, i just want the ruler to draw into person row". That was said after v162 had put a second clock above your people's rows, so this release takes that second clock back out. Nothing else about the chart moves.

The ruler's lines were never the thing missing, and they are still there. The lines that run down your people's rows came in v160 and they are untouched: every row of the day carries the faint grid, so a marker's edge still lands on a line you can follow up the chart to the top. That is the part you asked for, and it is the part that stays.

What comes back. The chart window is about 35 pixels taller again, because the block pinned at the foot of the chart no longer carries a clock row inside it. On your own day at a phone's width the foot block goes back from 105 pixels to 70.

One rule is kept, and it is worth naming. There is now a test whose whole job is that the clock is drawn once and that no clock is ever drawn above the people's rows again. It was written because the second clock came from a shared builder, so putting it back is one line - and a thing that can come back by accident should have something standing on it.

What has not moved. Not one module, batch, start time, cycle or saved scenario of yours changed, and nothing here blocks a sale or reads an order. Your three saved days are the same three days, each on the target it had, their stored numbers byte for byte the same.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

The operations guide (section 23) carries the same change, and the tests are 1188 passing with none failing.

23 Sep 2026 - engine v162 (no database step). The times along the clock are written a second time, standing directly above your people's rows, so a marker at the foot of the day can be read without tracing it 368 pixels up to the top.

What you asked. In your words: "why cant you draw the ruler into person's lines?" I went and measured it rather than answering from memory, because the answer turned out to be "half of it already is". The ruler's *lines* have reached your people's rows since v160 - every row of the day carries the faint grid, so a marker's edge already lands on a line you can follow up the chart. What had not reached them is the *clock*: the times written along the ruler, which lived only at the top of the panel. On your own day the top of the chart stood 368 pixels above your people's rows - the whole height of your eight module rows - so reading the clock and reading a marker meant looking in two places far apart. Asked which of the two ways to close that gap you wanted, you chose this one: "yes, draw the clock above the people's rows".

The clock is now drawn twice. Once at the top of the panel where it has always been, and once more standing directly above the people's rows, inside the block that is already pinned to the foot of the chart - so it rides the same pin as the rows it was asked for, and the day scrolling can never leave it behind. The foot one sits on the first person's row with no gap at all between them, measured on your own day.

The two clocks cannot disagree, because they are the same clock. They are built by one and the same function, from the same minute, with the same step and the same labels - so the top of the chart and the foot of it can never say two different things about where a minute is. Measured on your day at your usual reading: 97 ticks each, every label the same, and the 8:00 am tick at exactly the same pixel on both, whether the day is scrolled or not.

And the foot clock is read exactly as the top one is. Point at it, or drag your finger along it, and the same

hairline follows with the same reading in the same words and to the same minute, snapped the same way. It wears a crosshair under a computer's pointer, as the top one does. Two clocks that look alike now behave alike; there is no second way to work the foot one.

What it cost, measured rather than promised. About 35 pixels of the chart window. On your own day at a phone's width the block at the foot grew from 70 pixels to 105, and your day still fits inside the window with nothing to scroll - so those 35 pixels came out of empty space under your bars, not out of your day. At a wide screen the window is 720 tall holding 705 of content, and the day still fits.

One fault was found by testing on the real app, not by a test. The code that finds a clock's own track asked the wrong kind of list for the answer, and a real browser's list does not have that question to ask - so the screen failed to open at all with "This screen couldn't load", while every test stayed green, because the test's stand-in list is more forgiving than a real one. That is the standing lesson about stand-ins, and it is written into the code beside the line so the next person to touch it cannot repeat it. Fixed, and the tests now cover the wiring that hid it.

What has not moved. Not one module, batch, start time, cycle or saved scenario of yours changed, and nothing here blocks a sale or reads an order. Your three saved days were compared before and after: One baker day makes 24 pans still, No fridge, 1 person makes 24 pans still, and My sister proposal 21/9/2026 makes 4 pans still, each on the target it had, their stored numbers byte for byte the same.

What is not in it, so you are not left looking for it. The other way to reach the markers - moving the clock times around rather than drawing a second one - stays available if you ever want it instead. And the per-batch person, letting one batch go to somebody else while its neighbours stay, is still a release of its own because it needs a stored field and a one-time step in Supabase.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

The operations guide (section 23) carries the same change, and the tests are 1188 passing with none failing.

23 Sep 2026 - engine v161 (no database step). Your people's rows are the app's own white now, and a tap on one of their markers moves THAT STRETCH and nothing else - the batch it belongs to is untouched.

What you asked, and it was two things. In your words: "I want to have white background for person's line", and "The reassign job to next person is not whole day, it is that slot only". Asked what one tap should take with it, you told me: "we dont change the batch. Say a labour slot belongs to person1, clicking that slot, will offer to swap it to others, this basically to balance work load".

Your people's rows take the app's own white. They wore the cream band that every soft strip in the app wears, which reads as a highlight sitting under the markers rather than as the plain paper the day is read against. They are now the same white as the panel behind them and as every module's name cell, so the foot of the chart is the same paper as the rest of it. What still marks the block as the foot of the day is untouched: its own line above it, its own shadow, and the hairline between one person's row and the next.

A tap now moves one stretch, not the whole job. Until this release, tapping a marker handed over the whole module on that line - every batch of it. Your fold loop folds three times in one batch, so a tap anywhere near it took three of your markers with it. That is not what you asked for and it is not what balancing a day needs. What changes hands now is the single stretch your finger is on: the one batch, and which part of that batch. Its neighbours stay exactly where they were, and the batch itself is never touched.

The card says so before you press it, with the number. If the batch you tapped has more than one stretch on that person's row, the card counts them and says only the one you tapped moves. You are told what is about to happen rather than finding out from the chart afterwards.

And what you can hand a stretch to. The list is the other people who have a row on this day, named as you have named them - plus "whoever is free". That last one is not a person; it is what the day does when nobody is named, and without it a stretch on a module nobody was put on could be handed to somebody and could never be handed back. Pick it and the press under the list changes to say so, before you press it.

Your module's own card keeps up. A hand-over is now one stretch of a module rather than the whole of it, so "Who is at this module" is the module's default and no longer the whole answer. The card now carries a note naming what has been handed on, and to whom, with the count - so a card reading Person 2 while a marker of that module sits on Person 1's row cannot happen. On a day where you have handed nothing over, there is no note at all.

One thing the earlier note got wrong, said plainly. v160's entry said that moving the person on one batch of a module would need a new stored field on the module and a one-time step in Supabase. The stored field was indeed needed - it is there now - and the Supabase step was not. Your planner travels inside the app's own settings record, so nothing had to be run.

Nothing of yours is rewritten. The only thing that changes is the one stretch you hand over with your own press. Your day was read before and after: One baker day still makes 24 pans, and its one person's row still carries the same 32 stretches across the same clock. No module, no batch, no start time and no cycle moves by being read, and there is nothing to run in Supabase.

Measured on the real app rather than assumed. A stretch was tapped on your own day and handed over: the toast named the module and the clock, the marker left one person's row and arrived on the other, and the module's other three batches stayed where they were. Tapping in a gap still opens that person's own card, exactly as it always did.

Measured rather than promised. The suite is 1187 passing with none failing. The two tests that stood on the module-level move are now three that stand on the single stretch - that only the one key is written and the module's own person, its crew and its times are untouched; that a second hand-over keeps the first; that a stretch handed back to whoever is free survives as its own answer; and that the chart puts exactly the tapped stretch on the other row. The app's own More screen reads Engine v161.

23 Sep 2026 - engine v160 (no database step). In the Scenario planner, the ruler's lines now carry on down through the people's markers, and a tap on a stretch of somebody's day hands that job to another of your people.

What you asked, and it was two things. In your words: "can the ruler extend down to peoples marker area?", and "can the personX marker be click to change it job to personY, by a drop down person selector". Asked which way you wanted the ruler to reach the markers, you chose **a faint grid down every row** - the clock itself stays where it has always been, at the top. Asked what the tap should do, you told me: "click on the person's occupied time slot, a drop down list, list the other people available".

The ruler's lines reach the markers. Every row of the day has always drawn a faint line on the hour, so a marker could be traced up to an hour and no closer. That grid now carries on down through the module rows and through the people's rows, so the edge of a marker always lands on a line you can follow up the chart to the clock. On your day the lines are 36 pixels apart at the widest reading, 24 at the standard one, and 12 and 16 at the two close ones - so the grid you read a marker against is the ruler's own step everywhere the ruler can be read at all.

Why the grid has a step of its own. At the closest reading a single minute is 3.2 pixels wide. A line every 3.2 pixels is not a grid, it is a wash of grey. So the grid takes the smallest step that is both at least the ruler's own - which is what makes every grid line also a tick on the ruler, so the two can never disagree about where a minute is - and at least twelve pixels of your day. The clock strip at the top keeps its own ticks and is untouched; a module you have switched off keeps the grid rather than losing the only thing its dimmed bars are read against.

It costs nothing on the screen. The grid is painted, not drawn one line at a time, so your day is exactly as fast to open and to scroll as it was - which matters, because the ruler alone draws one tick for every minute at the closest reading.

Tap a stretch of someone's day and hand it over. Tap a person's marker on the timeline and a small card opens naming the job it is about - the module, its line where it has one, and the clock it runs from and to, in the same words the marker says under your pointer. Under it is a list of the other people on this day, and one press - "Move it to ...", named with whoever you picked - hands that job over. The two rows swap it the moment you

press, and anything that now collides is drawn red, exactly as before.

What the card tells you before you press it. A module carries one person for all of its batches, so if the module you tapped runs four batches on that line, all four move together - and the card says so, with the number, rather than letting you find out afterwards. The list offers only people who have a row on this chart, because somebody who is not standing anywhere on your day is not an answer the day can give. If you have only one person on the whole day, the card tells you that instead of opening an empty list.

The card goes with the job. The moment you press, the job leaves that person's row and the card closes with it, so what is left on the screen is the day as it now stands. A tap on the marker again - now on the other person's row - opens a fresh card naming whoever holds it. A card left standing would have gone on offering to move a job off somebody who no longer has it, and the press under it would have done nothing.

The tap is read by the minute under your finger, not by the marker. A person's marker is 11 pixels tall in a 34-pixel row, and at the widest reading one minute of work is 1.2 pixels wide - so a fingertip lands beside a job far more often than on it. The tap takes the whole height of the row, and it takes the nearest stretch of work within about six pixels, which is what makes the narrow end of the scale usable at all. A tap where that person is not working is not swallowed: it still opens the person's own card, which is what every tap on a person's row did before.

Nothing of yours is rewritten. The only thing that changes is the one job you hand over with your own press. No module, no batch, no start time, no cycle and no saved scenario moves by being read, and there is nothing to run in Supabase - the whole of this release is drawn from the planner you already have on your phone.

What is not in it, so you are not left looking for it. Moving the person on one batch of a module, leaving its other batches where they are, cannot be done without a new stored field on the module and a one-time step in Supabase. It is a release of its own whenever you want it. And a clock repeated above the people's rows was the other way to reach the markers - you chose the grid, and the other way stays available.

Measured rather than promised. The grid's spacing at each of the four readings, its reach and the card's own list are all covered by tests, as is the card closing the moment the job moves: the suite is 1186 passing with none failing, six tests new for this release and one of those standing on the card going with the job. The app's own More screen reads Engine v160.

23 Sep 2026 - engine v159 (no database step). Every card that opens over a screen can now be pushed aside by its own title bar, and the Scenario planner's chart window is one fixed height that no longer changes as you add people.

What you asked. In your words: "whenni click on batch to adjust the start time the pop up windows should be allow to drag" - and then, while I was working on it: "the chart window became shorter and shorter as no. of people grows. the window height should be fix".

What you settled, in three answers. You asked about the batch card and told me it should be **every** card, not only that one. You told me the position should **always reset** - a card opens where cards have always opened, and nothing is remembered. And you asked for **a faint grip line** on the title bar so it reads as something to hold rather than something to discover.

The batch card sat over the very bars it moves. The card that adjusts a batch's start time opens over the day it is about, and on a phone it covers the bars you are deciding between. That is what the drag is for: push it down out of the way, keep working, and it stays where you put it until you close it.

It is every card, not only the batch one. All of the cards that open over a screen share one title bar, so all of them now move the same way: the batch clock, a module, a person, an order, a customer, a product, a supplier, a unit and the rest. A card that looked the same but refused to move is a difference you would only find by trying it twice.

A card always opens where cards have always opened. The drag is a way to push a card aside while you work - it is not a setting and it is not remembered. Open the next card and it is where every card has always been.

The grip is what says so. There is a short faint bar near the top of every card's title bar, and the pointer over that bar becomes a hand. Measured, the bar is 34 pixels wide.

A card is held inside the screen, and that is deliberate. A card can be pushed until its own edge comes close to the edge of the screen, and no further. This matters more than it sounds. A card that could be dragged past the edge would hand the screen behind it something to scroll that it never had - and you could then scroll your own close button off the top, where no finger reaches it. Holding **the whole card** inside the screen means no button on it can ever be parked out of reach. On the phone a card is already nearly the full width, so what you get is vertical travel, which is the direction you want.

Measured rather than promised. On your One baker day the batch card for the mixing module is 337 pixels tall. Dragged down by 300, it moved by exactly 300 and its whole box stayed on the screen, and the screen behind it gained nothing to scroll. Dragged by 900 - far further than a finger can reach - it moved by 8 pixels and stopped, because 8 pixels is all the room it had. A module card, 784 pixels tall, moves 8 pixels; a person card, dragged 2000 pixels across and down, stopped with its whole box still inside the window; a wide card on a phone, 347 pixels wide in a 375-pixel screen, moves 8 pixels sideways, which is all the room it has. A drag on the card's close button is a press on that button and nothing else: measured, the card does not move and the cross still closes it.

The buttons keep working after a drag. Pressing the plus-five-minutes button on a batch card that has been dragged aside moves the batch and leaves the card where you put it - the card redraws its contents without going back to the top.

A drag on the day itself still scrolls the day. The title bar asks to be handled by hand; nothing else on the screen does. Dragging the chart - or the page - moves the page exactly as it did before, and the page does not move while you are dragging a card.

The chart window is one fixed height now. This is the second thing you asked for. The panel showing the day used to be a rubber band: its height was whatever the day's rows added up to, held between a floor and a cap. Measured at 375 pixels with one person it was 473 pixels tall and with a taller people list it was 648 - so the window grew as you added people, and once it had grown to its cap the pinned people list at its foot carried on growing into the days above it. At the same phone size your eight-module day left eight rows of modules visible with one person and four with a crowd. The panel is now **one height** - 80 percent of your screen, up to 760 pixels - and it does not change whatever you add. On your phone it is 649 pixels at 375 by 812 and 720 pixels at 1280 by 900. Your own eight-module day draws 473 pixels inside it, so **all eight module rows are visible at once**, and even your tallest people arrangement - seven person rows, which draws 648 pixels - still fits in the window with nothing to scroll.

Nothing of yours was rewritten. Your three saved days were compared byte for byte before and after this release: One baker day, No fridge 1 person and My sister proposal 21/9/2026 still stand on the targets they had, and not one module, batch, start time, cycle or saved scenario changed. Nothing here blocks a sale and nothing here reads an order.

No database step. The whole of this is drawn from the scenario already on your phone, so there is nothing to run in Supabase.

22 Sep 2026 - engine v158 (no database step). The Scenario planner's people rows come down to one line each and hold still, everything about a person moves into a tip on their own name, the "PEOPLE" heading is gone, and the clock ruler now reaches a single minute where you are zoomed in the closest.

What you asked. In your words: "can i have all the peoples cards noted in tool tips as well?", "i dont need the PEOPLE title", and "make the ruler resolution to 1min". Every one of those is below. Where a thing left the screen you told me where it should go instead: every person's own notes into the tip on their own name, and the ruler stepped so a minute is drawn where a minute can be drawn.

The people rows are one line each, and their height can no longer move. A person's row used to print its own notes as text on the row - how much work they are carrying, how many places they have to be, and a red line when two of their jobs collide. That made it the only part of the chart taller than the bars beside it, and the

only part that changed height as you worked: measured at 375 pixels, a row was 56 pixels on a clean day and 70 pixels when someone had a collision, so the whole strip measured 91 pixels one moment and 105 the next. **The strip sits pinned at the foot of the window**, which is what you asked for at v154 - and a pinned block that grows and shrinks as you scroll is a block that moves the very thing you are reading it against. The notes have left the row, so a person's row is **35 pixels** now, **on a clean day and on a day with a collision alike**, and the strip measures a flat 70 pixels either way. Those 35 pixels are what the row is: the same name cell the modules wear, and nothing else.

Everything about a person is in the tip on their own name. Point at a person's name on a computer and their own box opens beside it, carrying all of it: who they are (the whole label, including " (with 2, 3)", which the row shortens), how much work they are carrying, how many places they have to be, the red collision line in the words it already used, and **What they do today** - their jobs, each with its clock, up to six of them and then a count of the rest. It is the same list their own card carries, from the same builder, so the tip and the card cannot drift apart and say different things. The tip opens from the **title** alone, exactly as the module tips do since v156: pointing anywhere else on the row does nothing, and the tint that used to follow the pointer across the whole row is gone. That was your own v156 instruction - "only the note should shot as a too tip" - now applied to the last rows on the chart that had not had it.

Told rather than hidden: the tip is bottom-anchored, because the last row sits on the bottom edge. A module tip is centred on its row, which is what keeps a first or a last module's box off the edge of the window. The people strip is different - its last row **is** the bottom edge - so a centred box would be cut in half. A person's tip now sits on its row's own bottom line and grows **upward** into the day. Measured at the window's tightest, on a person with a collision and six jobs, the tallest possible tip is 196 pixels tall and it clears the clock strip above it with room to spare.

The widest-person trap, measured rather than promised. A person's tip is taller than a module's, so its lines have to be allowed to fold - and a folding box inside a 156-pixel name cell collapses to the width of its widest single word: measured live at 66 pixels wide and 923 pixels tall, one word to a line. The box is now told to be as wide as its own content, and the rule carries a note saying why, because the fault looks like a styling detail and is not. Measured after: 272 by 180 pixels, sitting on its row's bottom line, inside the window.

The "PEOPLE" heading is gone. It named what every row under it already names - each one wears a person glyph - so the strip is now its rows with nothing over them. The strip keeps its line and its shadow, so it still reads as the foot of the day.

The tally row's three facts are in its own tip, and on a phone they are unreachable. The **People at once** row is the one you asked for: the most hands the day needs at once, the busiest stretch by name, and the minutes paid for twice. Those three facts are printed nowhere else in the app, and that row has no tap of its own - so once they are in a tip, a phone has no gesture that can open them. **You chose this, with the consequence in front of you** ("tip only, no card"), and it is written here rather than left to be found: on your phone those three facts are now unreadable. On a computer nothing is lost. Giving that row its own card is one line whenever you want it, and this note is the place that says so.

The collision is not lost by the row giving up its red line. Under the day there is still a note for every collision, in full: who it is, **two jobs at once** with the minute it happens, both jobs named on lines of their own, and the sentence saying what to do about it - move one along the day, or combine with another person and accept it. The red outline on the bar itself is unchanged. So the row's red line leaving costs nothing you were reading.

The ruler now reaches a minute where a minute can be drawn. The day is drawn at four scales, from 1.2 to 3.2 pixels to a minute, and a minute hairline is physically undrawable at the two wide ones - so the tick step follows the scale, which is what you chose ("a minute where it can be drawn"). **Wide** ticks every 30 minutes, 36 pixels apart, which is the ruler you had. **Standard** ticks every 15 minutes, 24 pixels apart. **Close** ticks every 5 minutes, 12 pixels apart. And **Closest** ticks every **single minute**, 3.2 pixels apart. The hour is still solid and labelled at every one of the four, and the half hour is still the dashed tick it always was - the one you judge a 15-minute pass by. Everything finer is a new thin hairline at reduced weight, so a ruler at Standard does not read as a ruler made of dashes. The count is honest: at Closest your 24-hour window draws 1441 ticks, and a full

redraw takes 2.8 milliseconds against 2.3 at Wide - half a millisecond for the lot, so the ruler is drawn as elements rather than as a painted pattern.

And the cursor reads the exact minute. The hairline that follows your pointer down the day used to snap to the five minutes a batch's buttons step by. That was stricter than the screen it reads: a batch can be nudged by **one** minute, so the cursor could not name a time the chart was already able to set. It reads to the minute now - measured, a pointer 101 minutes into the day reads **5:41 am** where it used to read 5:40 am - and it still refuses to name a time outside your day, and still hangs its label to the left at the far end of the evening.

What has not moved, measured rather than promised. Your three saved days were compared before and after: **One baker day** still reads 24 pans across 8 modules, **No fridge, 1 person** still reads 24 pans, and **My sister proposal 21/9/2026** still reads 4 pans, each with the target it had. Their stored numbers were compared byte for byte and are untouched - the whole of this release is drawn from the scenario already on your phone. Nothing here blocks a sale and nothing here reads an order.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

CHANGELOG and the guide (section 23) both carry all of this, both PDFs rebuilt and read back, the app's own More screen reads Engine v158, and the tests are 1174 passing with none failing - eight of them new: that no person row prints its own notes and the tip carries them all, that the tip and the person's own card agree word for word, that a collision changes nothing about a row's shape, that the tip is a sibling of the name and not inside it, that the two tip rules both survive as file text, that the tally's three facts are in its tip and off its row, that the "PEOPLE" heading is gone, and that the ruler's step follows the scale at all four stops.

22 Sep 2026 - engine v157 (no database step). The Scenario planner's list of controls is much shorter and the modules window is taller: the module tags have moved into the tip, the scale is two presses around the stop it is on, the people are one box, and "The line" has left the row for the module's own card.

What you asked. In your words: "i want to make the modules window taller. The module tag, can you put them in the tip tips? Make the scale with+ -, People: make it drop down option, The line: this name is very unclear,, we just need the button, WORK BACKWARD. The idea is to make this list shorter and make room for a bigger modules windows." Every one of those six is below, and where a thing left the screen you told me where it should go instead: every tag into the tip, the line's own settings into the module's card, and the two paragraphs above the chart removed.

Two paragraphs left the day. The two sentences that used to stand between the "The day" heading and the first bar are gone: the one teaching how to read the chart, and the one naming the moment your day hangs from. At 375 pixels they measured 330 pixels between them, 210 and 120, and they were the whole of what you scrolled past before you reached a bar. With them gone the first bar sits 278 pixels below the heading where it sat 707. Nothing is lost by it: how to read the chart is section 23 of the operations guide, and the button that moves the day now says its own name in the row under the chart, so the move is still findable without a sentence pointing at it.

The modules window is taller, and now it fits. The window has a floor of 500 pixels and a cap of 760 where before it had only a cap of 560: the floor is what makes it a window rather than a box that resizes itself every time you edit a module, and the raised cap is what lets a long day use the room the rest of this release gave back. On your seeded day at 375 pixels the window is 616 pixels tall where it was 536, and the day's own drawing is 616 - so before this release the window hid 80 pixels of your day and now it hides none of it.

The tags are off the row and into the tip. Every tag a row used to wear is gone from the row: how many of a module you have, how many of its lots are in it at once, whether it waits above, whether it runs itself, and whether you have asked for more passes than a day holds. Point at the module's own title and the tags are the first line of the box, joined with a dot, with the two notes under them. Measured at 1280 wide on your day, the box is 42.6 pixels for the two notes alone and 58.9 with the tag line, and all nine boxes measured inside the panel left and right.

The tip is now taller than the row it belongs to, and here is why that is all right. Three lines of tip is 58.9

pixels against a 45-pixel row, so the box no longer fits inside its own row the way v156's did. What keeps it inside the panel is that it is centred on its row: half of it goes up into the clock strip's row above your first module, and half of it goes down into the rows below your last one. Every tip was measured inside the panel at both ends of the day, and nothing on the chart moves while one is open.

The scale is two presses. The four stops are kept exactly as they were, with the same pixels behind them: Wide, Standard, Close, Closest. What changed is that they are no longer four buttons. There is a round minus and a round plus, and the stop you are standing on is the word between them, so you can see where the dial is without opening anything. At either end of the dial the press that would go past the end is off rather than doing nothing when you tap it.

The people are one box. Three chips became one drop-down. What the box says when it is closed is the arrangement in force right now - Sharing them out, One a module, One to a line, Your own, or Combined when you have combined two people yourself - so the box names what it has done rather than hiding it. The four lines inside it are the arrangement in force and the two things it can do to your day, One a module (a person for each job, one job each) and Share them out (as few hands as can cover the day), and Combine two people on the end.

"The line" left the row, and its name with it. You said the name was unclear and that its place is the module's own card, and it is there now. The control row reads Scale, People, Walk the day and nothing else. On the module's card, under the words **Waiting on the module above**, is the same sentence the old card opened with - which of your modules wait on the one above, by name, and which keep the times you placed - and the same two presses it always offered, Chain every module and Take the waiting off. One module's own three answers, As the one above finishes, Never before the one above finishes and Its own time, are untouched and still on that same card, which is where they have been since v154.

The button alone. You said we only need the button, so the label that used to stand beside it is gone. It read "Your day backwards · first batch out 9:09 am", which at 375 pixels wrapped to two lines: 54 pixels of a phone screen spent explaining a button that already says what it does. What is left is the button, **Work the day backwards**, in the row under the chart.

What the row came back to. At 375 pixels your control row is 222 pixels tall where it was 305, so 83 pixels came back to the chart. Together with the two paragraphs that left and the taller window, the run from the "The day" heading to the first bar is 278 pixels where it was 707, and the day itself is not scrolled to at all: the window is 616 and the drawing is 616.

Nothing on a phone is left without a way to read it. A tip opens where there is a pointer that can hover at things, and a phone has not got one - measured at 375 pixels with a touch phone emulated, no tip opens at all. Everything a tip carries, the tag line and the two notes, is written out word for word at the top of the module's own card, in the row's own tag colours, and that card is what a tap on the row has always opened. There is nothing new to learn and nothing new to find.

What has not moved, measured rather than promised. Your three saved days were compared before and after this release. One baker day still reads 24 pans across 8 modules, its first batch out at 9:09 am and its day finishing at 1:30 pm. No fridge, 1 person still reads 24 pans. My sister proposal 21/9/2026 still reads 4 pans. Not one module, batch, start time, cycle or saved scenario of yours changed in this release. Nothing here blocks a sale and nothing here reads an order.

No database step. The whole of this is drawn from the scenario already on your phone, so there is nothing to run in Supabase.

The app's own More screen reads Engine v157. The tests are 1166 passing with none failing, eight of them new: that no tag is left on any module row and that every tag the tip carries is the same tag the module's card carries word for word, that a module you have two of is carried the same way, that the taller tip is still held inside the panel it is drawn in, that the scale is exactly two presses around the name of the stop it is on and that each press is off at its end of the dial, that the people are one drop-down whose closed text is the arrangement in force, that the line's own group has left the control row for the module's card, that the window is taller and wider-capped

and the two paragraphs above it are gone, and that the day-backwards group is the button and nothing else.

22 Sep 2026 - engine v156 (no database step). The module's notes now open as a tip over the day when you point at its title, and the row keeps its height while you read them.

What you asked. In your words: "i prefer not to expend the module with mouse hoover. Remain the height. When mouse hoover to the module title, only the note should shot as a too tip." That is a correction to v155, and it is a fair one: v155 brought the notes back by making the row taller, which moves every row below it the moment you point at it, so the bar you were reading is no longer where you left it. Reading a number should never move the thing you were reading it against.

The tip, and where it comes from. Point at a module's own title and its two notes open in a small box beside the name: the time the module starts and how many batches it runs, then what a batch costs in minutes and in your own hands. It is drawn over the day rather than in it, so nothing on the chart moves while it is open - not the row it belongs to, and not one row under it. Point away and it is gone.

Measured rather than promised. On your One baker day, with the tip open on the mixing module, the row reads 46 pixels before and 46 pixels after, the row underneath starts on the same pixel it did, and the day's own total is 527 pixels either way. The box is 43 pixels tall, which is inside the row it belongs to - that is deliberate, because it means a tip on the first or the last module of a long day cannot be cut off by the edge of the panel. The whole box measured inside the panel on the first module and on the last one.

Only on a computer, and only on the title. The tip opens where there is a pointer that can hover at things, which a phone has not got: on your phone pointing is a finger and a finger is a tap. There the notes are written out word for word at the top of the module's own card, which is what a tap on the row already opens, and that is unchanged from v155. Pointing at the row anywhere other than the title does nothing, because you asked for the note and not for the row to react.

What has not moved. Everything v155 did to the row is kept: a row is still exactly as tall as the bars it draws, the name still takes the row's first line with the badges on the line beneath, the name column is still 156 pixels, and your whole day still fits the window. Your three saved days were compared byte for byte before and after: One baker day still reads 24 pans across 8 modules with its first batch out at 9:09 am and its day finishing at 1:30 pm, No fridge, 1 person still reads 24 pans, and My sister proposal 21/9/2026 still reads 4 pans. Nothing here blocks a sale and nothing here reads an order.

No database step. The tip is drawn from the scenario already on your phone, so there is nothing to run in Supabase.

This change history and the operations guide (section 23) both carry all of this, both PDFs rebuilt, the app's own More screen reads Engine v156, and the tests are 1158 passing with none failing - one of them new: the tip is off the row, laid out absolutely so the row cannot grow, centred on the row it belongs to, opened by the title alone, and its two lines are held to a height that fits inside a row.

22 Sep 2026 - engine v155 (no database step). Every module row is now as tall as its own bars, so your whole day fits the window with no scrolling; the two note lines come off the row and come back when you point at it on a computer, and they are always written out on the module's own card.

What you asked, and you were right. In your words: "The window for scrolling module became too small, why not reduce the height of each module, the notes can just shown up upon mouse hoover." You were right that the rows were too tall, and the cause was not the bars. On your One baker day the eight module rows measured 912 pixels of drawing inside a window about 536 pixels tall, so only about four of them could be seen at once and the rest had to be scrolled to. The bars themselves need only 45 pixels; it was the words beside them that made the rows tall. The name column is 138 pixels wide, and a name like The rests and the stretch and folds wraps to seven lines in it, with each of the two note lines wrapping to two more.

Every module row is now the height of its bars. A row measures 46 pixels on its own day, and the whole of it - the clock strip at the top, your eight modules and your people at the foot - now measures about 502 pixels inside the 527 the window gives it, so the day is read in one look and nothing is scrolled to. Every number in it is the number it was.

The two notes leave the row, and two things bring them back. On a computer, pointing at a row brings both notes back and tints the row, and the row grows taller to hold them rather than covering what is beside it, so a note is never cut in half. On your phone there is no pointing, so a reveal that needed one would be a reveal your phone could never reach; the notes are written out instead at the top of the module's own card, which is what a tap on the row already opens. Nothing new to learn and nothing new to find. On your mixing module the card now opens with what the row used to whisper: the time it starts and how many batches, then what a batch costs in minutes and in your own hands.

The name takes the row, and the badges drop below it. This was the real damage of the old row, and it is worth naming. Your longest module name is 273 pixels wide, and it was sharing its line with a small badge, which left just 35 pixels of the name to read - about two words of it. The name now has the row's first line to itself and the badges sit on the line beneath, and four of your eight module names now read whole where one did before. The name column is 18 pixels wider as well, from 138 to 156, which is the difference between one and three more of your names reading whole. A name still too long for the row is shortened with three dots, and the whole of it is on the module's card and under the pointer on a computer.

A fault found while building it, worth naming. A module drawn as two production line and set to wait on the module above it could not be drawn at all: the drawing asked for the waiting badge's words in a place where they did not exist, and that is not a blank space, it is an error that stops the screen. Two taps on your own module card would have found it. None of your saved modules is drawn as two production line, which is the only reason you never met it. It is fixed, and a test now stands on it so it cannot come back.

What has not moved, measured rather than promised. Your three saved days were compared byte for byte before and after this release. One baker day still reads 24 pans across 8 modules, its first batch out at 9:09 am and its day finishing at 1:30 pm. No fridge, 1 person still reads 24 pans. My sister proposal 21/9/2026 still reads 4 pans. Not one module, batch, start time, cycle or saved scenario of yours changed, and the change is in how a row is drawn, not in what any row says. Nothing here blocks a sale and nothing here reads an order.

No database step. The row is drawn from the scenario already on your phone, so there is nothing to run in Supabase.

This change history and the operations guide (section 23) both carry all of this, both PDFs rebuilt, the app's own More screen reads Engine v155, and the tests are 1157 passing with none failing - three of them new: one on the notes being carried by the row and printed by the card, one on the name taking its own line ahead of its badges, and one on the module drawn as two production line that waits on the module above.

22 Sep 2026 - the documents, not the app (no engine version, no database step). The guides and this change history print their own italics now, so the stray asterisks are gone.

Every one of these documents is built from a markup in which one asterisk on each side of a word means "print this in italics". The builder stopped understanding the single asterisk at some point and reads only the pair, so a word meant to lean quietly came out wearing its own asterisks instead: *'why no button for work backward', Cutting and packing, status New*. Not one sentence read wrongly, which is why it went unnoticed for so long. There were 116 of them in this change history, 44 in the operations guide, and 9 more across the money and profit guide, the alert guide and the safety guide. All of them now use the pair of underscores the builder does understand, so they print as the italics they were always meant to be.

The asterisks that were meant to be read stay exactly as they are: the one marking a required field on the order form, and the fill-in lines on the weekly note sheet. Those are marks, not markup.

Nothing in the app changed. More still reads Engine v154, every screen is exactly where it was, and no test moved. All five documents were rebuilt and read back to check: this change history prints no asterisk at all now, and the only two left in the operations guide are the two that are meant to be there.

22 Sep 2026 - engine v154 (no database step). A module now says how it takes its start - three answers where there was one switch - so the packing begins the minute the cooling ends; every batch's own card offers the move again; and the people's rows are held below the modules so you can read a slot against them.

What you asked. Your three reports, in your own words: "the cut and packing batch din follow the earlier batch end", and then your own correction, "cutting and packing sit below cooling down, so cutting and packing batch start should follow cooling down batch end". "The delta t disappeared, before this we have it. I want each batch start time to be adjustable, like the 1st module. Just need to show delta on the batch 1st offset only, then following module of that step dont have to show the delta because it follow the previous module tightly." And "I want to freeze the persons card, so that by scrolling thru modules i can see exactly where that slot of that person tie up to and searching for opportunity to move some batch start time to reduce the number of person needed." You also set the rule the first of them had to be built by: "the behaviour has to base on configuration, not a hard wired".

1. **What was actually wrong, and it was words against behaviour.** The old switch was called **Waits for the module above**, and its hint promised a follow - *2nd module 1st cycle should follow 1st module 1st cycle completion*. What the day really did with it was a **floor**: a batch could be pushed later than the module above, never pulled back onto it. So a packing step sitting 8 minutes after the cooling ended could never come back onto it, however the switch was set. Your report was right and it was the most useful kind: the sentence on the card and the arithmetic under it disagreed.

2. **A module now says how it takes its start, in one of three answers.** Open any module and the card asks **How this module takes its start**, with three choices: **As the one above finishes** (batch 1 of this module starts the minute the module above finishes its own batch 1, batch 2 on batch 2 - a tight follow with no gap); **Never before the one above finishes** (a batch here can never start before the module above has finished that same batch, but a later time you set yourself stands - this is the old switch, kept exactly as it behaved); and **Its own time** (the module above has no say over this one). Each one carries a sentence saying what it means before you tap it.

3. **The choice is on the module, not on a module's name.** This was your rule and it is the way it is built: nothing in the day knows a module by what it is called. Set **As the one above finishes** on your packing and it lands on whatever really sits above it - cooling, or anything else you put there, in whatever order you keep your modules. There is a test that renames every module of a day and confirms the same answer, so this cannot quietly become a rule about cooling.

4. **Nothing of yours was re-laid by being read.** A saved module keeps the behaviour it already had: a module that was waiting on the one above it opens on **Never before the one above finishes**, and a module that stood on its own opens on **Its own time**. A new module arrives on **Never before the one above finishes** too, because **As the one above finishes** on a module you have not yet told what feeds it would move your day the moment you added it. Your three days were compared before and after: 24 pans on One baker day, 12 on the seeded line, 4 on your sister's, and not one start time of yours changed.

5. **Every batch's card offers the move again - this is your second report.** A batch below batch 1 of a module set to wait used to be **refused** the two button pairs, with a sentence telling you to go and change the module instead. Every batch at every module now offers **Five minutes at a time** and **One minute at a time**, and the last paragraph of this note says plainly what a move means in each place so the two cannot disagree.

6. **The offset is read out on batch 1, and only on batch 1 - exactly as you asked for it.** On batch 1 of a module that is not the first, the card reads how this module takes its start, in words: **follows the module above**, **never before the module above**, or **on the line**. The batches below batch 1 do not repeat that, because they ride the module above at this module's own pace - there is no offset of the module's to read on them. A batch you have **held** on purpose is the exception, and it is a deliberate one: it reads **delta t = +5 min** on its own card, because a batch that is off the line must never look like one that is on it.

7. **A move below the first module is a hold, and a hold can only ever put a batch later.** That is not a limitation of the buttons: a batch cannot start before the dough it is made of exists. What makes it safe rather than a wall is that a hold is an **offset measured from where the line puts the batch**, not a time - so move the module above it, or change how many batches it runs, and the held batch comes with it, keeping its place in the queue. The card of a held batch carries **Back onto the line**, which takes the hold off, and taking a hold all the way off now says so - *Cutting and packing is back on the line* - where it used to say *held 0 minutes behind*,

which is a sentence about nothing happening.

8. **A press on batch 1 of a module that follows moves the whole module.** A module set to **As the one above finishes** has no start time of its own to write: it begins where the module above ends, whatever a stored start says. So a press there writes the module's own hold instead, and **every** batch of it comes with it - which also keeps a press on batch 1 meaning the same thing on every module, rather than being a button that appears to do nothing on one of them.

9. **The people are held below the modules - this is your third report.** The module rows scroll inside the day's own panel and the **People** rows and the **People at once** total are pinned at its foot, so scrolling down through eight modules never takes the person you are reading away from the bar you are reading it against. This is the move the whole report was for: hold a slot in view, look for the batch you can move to empty it, then go and move that batch. The red now-line of the walk through still runs over the people's rows, because that line is the day happening rather than a reading of it.

10. **One fault found while testing, and it is worth writing down.** The test that stands in for the browser kept an attribute a real browser also exposes as **dataset** - and the chart reads which batch you tapped off that. So every simulated tap had been reading as **batch 1**: a whole class of taps the tests could not see, which is exactly why "the delta t disappeared" could be true on your phone while a test suite said all was well. The stand-in is unforgiving now, the way the browser is. Two tests written against it failed immediately and both were real: the wording of a hold taken off, and the day's own card opening where a batch's card should.

11. **No database step, and it is on every phone.** Nothing to run in Supabase. The planner travels with the app's own settings row, so this new per-module answer goes up with everything else the planner holds and comes down onto your other devices the same way - the choice you make on your phone is the choice your tablet reads.

12. **What has not moved.** One baker day still reads 24 pans across 8 modules with its four stored rhythm times untouched; the seeded line still makes 12 with the chiller as the wall; your sister's line still makes 4. v152's backwards press and v153's button on the row are unchanged, and neither this release nor those reads a single order or blocks a single sale.

The CHANGELOG and the guide (sections 23 and 26) both carry all of this, both PDFs are rebuilt, and the app's own More screen reads Engine v154. The tests are 1154 passing with none failing - 14 of them new: 8 on the model (the three answers and how a saved module is read through them, the packing landing on the cooling's end, the follow being a choice on the module and not a rule about a name, a switched-off module stepped over, a module above with fewer batches answering for the batches it has, a hold riding the module above and never pulling a batch earlier, and her three days reading the same times) and 6 on the screen (the three pills on the module card, the packing landing on the oven, the pairs on every batch, the offset read on batch 1 alone, a press on batch 1 of a following module moving the whole module with its way back, and the people's rows held below the modules). No database step - nothing to run in Supabase.

22 Sep 2026 - engine v153 (no database step). The day's own row works your day backwards - one press, and it is a button that acts instead of one that only tells you where to tap.

What you asked. Your words: *'why no button for work backward, this is to reposition the batches latest start time'*. You were right, and it was about a control that only talked: v152 put the move behind a tap on one bar and left the row carrying a chip that named the moment and then told you which bar to tap. A thing shaped like a button whose whole effect is to send you somewhere else reads as a dead control, and that is what it was.

1. **The button does the work now.** The row under the chart reads **Your day backwards · first batch out 9:09 am**, and the one pressable thing beside it is **Work the day backwards**. One press and the day is worked back on the spot - no card in the way, the same seven modules pulled to their latest starts, and the same message saying what moved and what it cost your hands.

2. **It is the card's own press, not a second one.** Same writer, same words, same snapshot, same way back - so the two doors cannot drift apart. Press from the row, then open the card on batch 1 of *Cutting and packing*, and the card is showing you the day the row just made, with the way back on it.

3. **Pressed on a day already worked back, it says so and writes nothing.** This was your own choice and it is the rule the button lives by: a button that comes and goes reads as a fault, so it stays in the row on every day. On a tight day the press answers *Every module of your line is already as late as it can go - there is nothing left to pull back.*

4. **The way back is beside the button that moved the day.** While there is something to put back, the row carries **Put my start times back** next to **Work the day backwards** - the pattern Stop beside Start already uses. The card keeps its own copy, so the undo is either one tap away or three, and never missing. It is still one-shot and still lasts only while the screen is open: leave the screen and the day you pressed is the day you have.

5. **A line with no chain says so, rather than drawing a dead button again.** A day of a single module is not a chain - its first batch is simply that module's own start time - so the button is drawn, carries no moment in its label, and says in words that there is no module above the last one to work back through. The fault you found is the one thing this release refuses to repeat.

6. **The sentence above the chart now names the button.** It used to tell you to tap batch 1 of *Cutting and packing* and stop there, which was the only door v152 had. It now names both: the button in the row does it in one press, and tapping batch 1 opens the same move with the whole chain laid out before you press it, which is the one to read first.

7. **One fault found while building it, and it was in this release's own new code.** The undo works off a snapshot of the day kept in the screen's own memory and never stored, and v152's card took it from the right place. The row's new press first took it from the app's saved state instead, which put that snapshot into your data - and a reload would then have offered you a way back to a day you had already left. It was found by measuring the stored data after a press rather than by reading the code, it is fixed, and there is now a test that fails if the snapshot ever reaches your data again. Nothing of yours was changed by it.

8. **What has not moved.** Your three shelf days were compared byte-for-byte before and after: *No fridge, 1 person* still reads 24 pans, *My sister proposal 21/9/2026* still reads 4, and *One baker day* still reads 24 pans with the first batch out at 9:09 am and the day finishing at 1:30 pm. A press followed by the way back leaves the stored data byte-for-byte identical to what it was - measured, not promised - and the snapshot is gone when the screen is. Nothing here blocks a sale and nothing here reads an order.

CHANGELOG and the guide (section 23) both carry this, both PDFs rebuilt, the app's own More screen reads Engine v153, and the tests are 1140 passing with none failing - five of them new: the row press moving the day exactly as the card does, a finished day keeping its button and saying it is done, the way back sitting beside the button that moved the day with the card offering it too, a line with no chain saying so instead of drawing a dead button, and the one that pins the snapshot out of your saved data. No database step - nothing to run in Supabase.

22 Sep 2026 - engine v152 (no database step). The planner can work your day backwards - the day hangs from your first batch, and every module gets its latest start.

What you asked. Your words: *'how to make the calculate backward works?'* The backward count already existed on More -> Production line (section 24), where it counts back from the minute the first pans must be at the oven. The Scenario planner had none: it could only ever push a day later, never pull it earlier. So the planner has one now, on the planner's own day.

The moment your day hangs from is your own last module, not the oven. The anchor is the end of your first batch at the last switched-on module of your line - whatever that module happens to be. Add, remove or reorder a module and the anchor moves with your line. On your own *One baker day* that module is *Cutting and packing*, so the day hangs from 9:09 am.

Where it is, and how you find it. Tap the B1 over the first bar of that last module - the card you already get by tapping a batch - and it is now the day's own card. It reads the module, the end of your first batch and your day's finish, then every module with the time it sits at and the time it could sit at. On your line it reads: *Mixing the dough in the tub 4:01 am to 4:43 am, The rests and the stretch and folds 4:21 to 5:03 am, Oil the pans and*

weigh the dough out 6:24 to 7:06 am, Into the proofer 6:39 to 7:21 am, Dimple and top 7:24 to 8:06 am, The proofer again 7:30 to 8:12 am, The oven swap and the bake 8:00 to 8:42 am, and Cutting and packing 8:57 am with a tick. A tick means that module is already as late as the line allows. So the control is findable rather than hunted for, the day card also carries a chip reading First batch out 9:09 am with the sentence saying which bar to tap.

The rule is measured off your own day, not subtracted from the anchor. This is the one place I departed from what I first planned, because the arithmetic was wrong on your own numbers. Counting backwards is not the sum of the cycle minutes: on your 24-pan line the eight modules' minutes add to 853 while the chain you actually run spans 1107 minutes, because your modules do not sit end to end - the fold loop and the proofer are wider than the few minutes a batch takes in them. Take 853 off your anchor and the mix is handed a LATER start than the one it already has, which is the opposite of what you asked for. So the calculation measures the room between each pair of modules as your stored times really stand, and pulls each module back by that room. On your line it gives the mix 4:43 am, and the 42 minutes it gains are the slack your day really has.

One press, and a second press does nothing. Pull them back to their latest start moves the seven modules that can still come later and leaves every other one exactly where it is. A module that is switched off is stepped over, exactly as the day already steps over one. A press never moves a module EARLIER than it already stands, and pressing it again finds nothing left to take.

It is a real change, and the way back is on the same card. Said plainly rather than quietly: this rewrites the start times of your switched-on modules and saves them into your scenario. It is not a preview. So while you are on that screen the card carries a second button, Put my start times back, which puts every module the press moved exactly where it was and then goes. It is one-shot, and it lasts only while the screen is open, because it is never stored - the same reason the walk through's own clock is not. Leave the screen and the day you pressed is the day you have.

Below that are the same two pairs of buttons you already know, and here they move the whole day. Five minutes at a time and One minute at a time, each with an earlier and a later button. Press one and every module moves together by that amount, so the day keeps its shape: on your line all eight modules move five minutes and every gap between them is the number it was. That is the move for reading what an earlier or a later start does to the whole day without re-drawing it.

A tighter day can need another pair of hands, and the app says so instead of hiding it. This is the honest cost of taking the slack out. Squeeze the day to its shortest span and two of your jobs can land in the same minute, so a press can raise the people the day needs - on your One baker day from one person to two. Nothing blocks and nothing is refused: the toast reports it in words, the People area below shows the collision in red, and the button under it takes it back.

The clock you point with now follows your finger, and the clock strip stays where you can read it. Your four asks of that day chart, all in this release. The reading follows the pointer DOWN the day as well as across it - it used to sit at the top of the chart, which is the right place only while the top of the chart is on screen, so pointing at a bar four rows down meant reading a clock a screen away. It rides 8px clear of your finger rather than under it, because a reading under your own finger hides the bar you are holding it against - and near the top of the chart, where there is no room above, it drops below the pointer instead.

Lifted at v171: the 8 pixels above your finger described here are two CSS inches now, because 8 pixels is under a fingertip rather than clear of it. The rule itself is unchanged and this entry still describes it: above your pointer where the window has the room, and below it where it does not. Read the v171 entry above for the rest.

The chart is a panel of its own now: the clock strip is pinned at the top of that panel and the rows scroll under it, so the hours stay beside the bars you have scrolled to. The panel is capped, so it never grows past what you can see at once.

And one fault found while building it. The red now-line of the walk through was being drawn in the chart's name column - 138px, about an hour and a half, to the LEFT of the minute it was naming, because its position was measured when the chart was built and a chart that is not yet on the page has no position at all. Measured

off the clock strip's own column instead, it now stands on the minute it names. Nothing about the walk through changed except that the line is where it says.

What has not moved. Your three shelf days were compared byte for byte before and after: No fridge, 1 person still reads 24 pans, My sister proposal 21/9/2026 still 4 pans, and One baker day still reads 24 pans across 8 modules with the first batch out at 9:09 am, the day finishing at 1:30 pm and the stored rhythm of 4:01, 5:28, 6:55 and 8:22 am untouched. Not one module, batch, start time, cycle or saved scenario changed, because nothing moves except by your own press. Nothing here blocks a sale and nothing here reads an order.

CHANGELOG and the guide (section 23, cross-referenced to section 24) both carry all of this, both PDFs rebuilt, and the tests are 1136 passing with none failing - nine of them new: four on the calculation (the exact latest start of every module on One baker day, that it is measured off the day rather than subtracted from the anchor, that a switched-off module is stepped over, and that a press never moves a module earlier nor does anything a second time) and five on the card and the drawing (the card opening on batch 1 at the last module and nowhere else, one press leaving the day's own numbers alone, the way back putting every start time exactly where it was, the step pairs moving the whole day with its shape kept, and the chip signposting the moment the day hangs from). No database step - nothing to run in Supabase.

22 Sep 2026 - engine v151 (no database step). The module card answers to you, every batch is numbered, and the day calls your people.

What you asked for, in your own words. Eleven things in one message, all of them inside More -> Scenario planner: the scenario's name first; the module card asking only what you alone can answer; a B number over every batch; moving a batch at any module the same way you move one at the first; naming Person 1 and Person 2; a marker that reads as a person; and the day calling you a minute before each job. It all ships together, as you asked. One of the eleven was not a wish but a fault: a person's card could not be opened at all before this, which is why naming Person 1 and Person 2 had no way in.

The scenario's name comes first. The card at the top of the planner now reads What this scenario is called, then My day starts at, then Pans a day you want from it. Going to that screen to rename a scenario used to mean reading past two numbers you did not come for.

The module card no longer asks you two things you do not think in. Which job of your line is this has gone. How you place each batch's start time has gone as well, and you said an alternative is coming - so the batches' stored times are kept exactly where they are, and nothing reads differently for it. What taking those two blocks out costs you is written at the end of this note, honestly.

A new module now arrives ready to use. One cycle of 20 minutes with all 20 of them as your hands at the start, Pans in one batch 1, waits for the module above switched on, allow multiple production line switched on. Before this a new module arrived at 15 minutes with both switches off, and you had to type five things before it was a module at all.

Batches follow the one above, unless you say otherwise. Minutes from one batch to the next and How many batches it runs in the day both carry an Auto. Auto is written as an empty box with the word in it, so you can always tell a number you set from a number the day worked out. The batch count is a box on the first module - that is where your day's number comes from - and on any module that has a number of its own. A module on Auto says its number out loud and names what it is following, and beside it is one button, Give it its own number, which keeps the count it is running now so taking the wheel never moves your day a minute. Auto spacing is not new behaviour: the day has always read a module with no pace of its own as one batch following the one before it end to end. It has a name now.

How many production line do you have - the number is unchanged, the words are yours. You said it yourself: "that 2 refer to 2 production line, if the meaning is for that purpose, a better remark needed for that. no functional changes needed." So the behaviour is exactly what it was, and the remark is rewritten around production line instead of the word lines. It also says plainly what happens when Allow multiple production line is on: then the module is one production line with one person, and this number is not in force. And Let its batches overlap is called Allow multiple production line now, everywhere on the card and in the note under the day.

Duplicate this module. At the foot of the module card, beside Delete, there is now Duplicate this module. The copy lands directly under the one you copied and carries its cycles, its numbers and its person exactly - a copy is a copy and none of the new-module defaults reach it. You asked for this because you can add modules as you wish now, and two ovens of yours are two of the same thing.

Every batch is numbered, on every module and every production line. A small B1, B2, B3 above each bar, in a band the row grew for it, on every module including one drawn as two production line. The number carries the batch it names, so tapping B3 opens batch 3's own clock - a number you use to tell one batch from another is no use if it is not also a way to pick one.

Move a batch at any module, not only at the first. Tapping a bar or its number opens that batch's clock, and the card now says the batch number and its delta t beside each other. On the first module a move is the batch's own start time, which is what it always was. On every later module a move is a delta t - a hold measured from where the production line puts that batch - so the batch rides the chain: move the module above it and this batch comes with it, keeping its offset. A delta t can only ever hold a batch back, because a batch cannot start before the dough it is made of exists, and the tag over the bar shows delta t = +5 whenever there is one. When your batch is held back on purpose the card also offers Back onto the line, which takes the hold off, so a nudge is never a one-way door.

Your people have names, and the card that was unreachable opens. This is the fault you reported: "in the person card, now person card is not accessible". The person's row had no handler of any kind, so no gesture anywhere opened their card. Tapping a person's row below the day now opens it: a box for what you call them, a switch for whether the day calls them, and the jobs they are on today. The name is kept in the app's settings rather than in the scenario, so you type it once and every scenario uses it - person numbers start again at 1 in each scenario, so the name you give person 1 shows wherever person 1 is working, and the card says so. The card's own heading and its "Call them a minute before their next job" line follow the name as you type it, so the card never goes on calling them Person 1 while the row behind it already says their name. The box itself is never rebuilt while you are typing in it, which is what keeps your cursor where you left it.

A person's marker is theirs now, not the module's. The People rows used to be tinted by the module each stretch of work came from, so one person's row was eight colours and read as nothing in particular. Each person now has their own colour, with their name on the stretch where it is wide enough - so a row reads as one worker's day and two rows are two people. A collision still wears its red outline over the top.

The day walks, and it calls you. A Start the day now button in the Walk the day row, with Stop beside it while it runs, and a red now-line that follows the real clock down the chart. The call comes one minute before each job a person owns - one minute and not the minute of it, because your fold is a one-minute job and telling you to fold as you should already be folding is telling you too late. A call card appears over everything: their name, their own colour, the job, and the clock time it is for, with one OK. OK takes it away. If you do not tap it, the next call simply replaces it rather than queuing behind it, which is your own rule. Each person's call is a different note, so two people called in the same minute are two sounds. The card and the guide both say the limit plainly: it runs only while this screen is open and awake, and a phone that locks or a browser that closes stops it. It is a call for the day you are walking through, not an alarm that survives the app being shut.

The call names the cycle only when you have named it. A module's own name can be a mouthful ("The rests and the stretch and folds"), and the thing you can actually go and do is the cycle inside it, so the call carries the cycle's name as a line of its own when that cycle has one. An unnamed cycle gives no line at all. That last part is a fault this release nearly shipped: the line was drawn through the same number-formatting helper the rest of the chart uses, and on a cycle with no name that helper answers "0", so the card went out printing a bare 0 above the clock. There is a test on it now, and it is written so it fails if the line ever comes back.

Your day start is never rewritten by it. Starting the day sets the run's own minute zero to the minute you press, and nothing is saved - so a 4:01 am plan started at 9 am calls at 9 am, and your scenario's own day start is exactly where you left it.

What this release takes away, said plainly rather than quietly. Three things are gone with the two blocks

removed from the module card. The job list is gone, so the way to CORRECT a module whose name the Production line does not recognise is gone with it - a module still crosses over by what you call it, exactly as before, and every module of yours already crosses. The per-batch start boxes are gone from the card, so the one control that put every batch back on its own even spacing, Space them evenly again, is gone too; tapping a bar gives you the same five-minute and one-minute moves, but not that reset. Say the word and any of the three comes back in one line.

What has not moved. Your two saved lines and the seeded starting line were compared byte-for-byte before and after this release: No fridge, 1 person still reads 24 pans, My sister proposal 21/9/2026 still reads 12, and the seeded starting line still reads 12 with the chiller as the wall. Not one module, batch, start time, cycle or saved scenario of yours changed, because every new default is for a module you have not filled in yet - a stored number or a stored start is yours and is never rewritten, which is the same rule the cycles migration used. Nothing here blocks a sale and nothing here reads an order.

CHANGELOG and the guide (section 26) both carry all of this, both PDFs rebuilt, and the tests are 1127 passing with none failing - fourteen of them new: five on the model (a new module's own defaults; Auto following the one above while a number of your own is never overruled; a delta t riding the chain and never pulling a batch in front; reading a line twice still a no-op; all three seeded lines pinned batch time by batch time; and the call landing one minute before the job) and nine on the drawing and the card (a batch number on every bar including two production line, a copy carrying its source's cycles, the person's name on the row, the card's own heading following it, a named cycle on the call and no line at all for an unnamed one, and the four halves of the live clock). No database step - nothing to run in Supabase.

22 Sep 2026 - engine v150 (no database step). The word is cycle - and a batch of four now reads as four.

What you told me, in your own words: "there is not steps term in this project, i want the cycle shown in batch, and the cycle labour shown visually". You were right on the first point, and the app was the thing saying it. I measured the day chart before changing anything: nothing had disappeared. On your One baker day the fold is one batch of four cycles and its bar draws four shaded bands with your three folds on them, and the oven is one batch of two cycles with your two-minute swap on the second - 24 cycle bands and 32 stretches of your hands across the day. The feature was built and it was drawing. What was wrong was the word on it.

Why your own two saved lines look plain. Both the line you call No fridge, 1 person and My sister proposal 21/9/2026 have exactly one cycle in every module, because both were saved before cycles existed and were given one cycle each so the bar's shape would not move. One cycle is drawn as one plain bar. You settled that yourself, and the app agrees with you: each batch has at least one cycle, and a batch with one cycle is that cycle. So nothing of yours was rewritten - both lines are correct as they stand, and not one module, batch, start time or saved scenario was touched.

The cycles box no longer calls a cycle a step. It used to say "Each cycle is one step of this module", which is exactly the sentence that made you doubt the feature was there. It now says "Each cycle is one piece of this module's work", in the order you work them. The hint over the day chart says the same thing in the same word: the paler bands inside a bar are the separate cycles of that batch. Three more places where the app said step while it meant in sync now say what they mean - batch counts that differ from the module before them, and the line under the day naming any that no longer match. The batch's own card used to say "the step the day is read in"; it now says the amount the day is read in, which is what it is.

The Production line says stage, not step. That screen walks your bake day in order, and it called each part of the day a step. It now calls each of them a stage. The word job is left alone there on purpose: on that screen job already means a hand-job - one of the six jobs you do - and the proofer is not one of them, which the screen says in its own sentence. Calling the proofer a job would have put two contradicting sentences on one screen.

A batch of four now really does read as four. This is the one fault in the drawing, and it was a real one. The shades the cycles are drawn in said "a batch of four reads as four steps", but the ladder only had three shades and started again at the fourth - so on a four-cycle batch the first cycle and the fourth wore the same palest

band, and the palest of the three was so faint over the pastel bars that it barely showed at all. There are four shades now, spaced evenly, and the palest has been lifted off the floor so it still reads as a wash. A fifth cycle and beyond reuse the darkest rather than falling back to the palest, so a long batch never ends on the faintest band.

What has not moved. Not one minute, position or labour band on the day chart. Your own saved line still reads 24 pans across 9 modules, the seeded starting line still 12 pans with the chiller as the wall, and your sister's still 4 pans in 8 modules. Nothing here blocks a sale and nothing here reads an order.

The tests are 1113 passing with none failing - three of them new, and one of them fails if a four-cycle batch is ever drawn with two cycles sharing a shade. No database step - nothing to run in Supabase.

22 Sep 2026 - engine v149 (no database step). The way down - taking batches off, so a day that makes more than you asked for can come back to your number.

What you found, in your own words: "This scenario makes 36 pans, and you want 24. this does not agrees?" You were right, and it was two faults in one card, not one. Your One baker day made 24 pans. You typed 36, pressed **Use these numbers**, and the ladder did what a ladder does - it raised all eight modules from 4 batches to 6 and the day made 36 pans, which is what you had just asked it for. Then you put 24 back in the box. The card underneath said "This scenario already makes your 24 pans" while the line directly above it said the day makes 36. Both sentences were drawn from the same card, one of them had to be wrong, and the ladder that could have settled it only ever went one way.

The climb only adds. Every rung of it raises something: more batches in the day, or a second machine. Not one rung has ever taken a batch away, so there was no way on that screen to bring 36 back to 24 - no amount of pressing would have done it. That half simply did not exist, and no wording change could have covered for its absence.

There is now a way down. When the day makes more than the number you asked for, the card is headed **The way down** and carries one rung: every module the day is waiting on, and the batches to take off each of them. On the day you were looking at it reads: all eight modules holding the day at 36 pans, 6 -> 4 batches in the day at 6 pans a batch, the same change in each of them - that takes the day from 36 to 24. The same **Use these numbers** button applies it, on the heading and at the foot of the card.

Why it is one move and not a ladder. The day is the **least** any module turns out, so every module sitting on that least is holding the day there. Bring one of them down alone and the day does not move at all, because one of the others takes its place as the limit. They have to come down together, and then the day lands on your number in a single move. That is why this half is a settling move rather than a rung-by-rung climb.

The batch times you have set by hand are left exactly where you put them. This is the one place the way down deliberately differs from the climb. Raising a count makes a new rhythm, so the climb re-spaces a module's batch times from its start. **Lowering** a count only takes batches off the **end** of the day - so if you have dragged a batch to a time that suits your morning, that time stays. Nothing on this move rewrites a minute you chose.

When a batch is bigger than the number you asked for, it says so. Six pans a batch and one batch a day, and you ask for three: no count of batches gets under one batch, so there is no move to offer and no button to press. The card says it plainly instead - one batch is 6 pans, so this line can never come below 6 - because a card with nothing on it and no explanation reads as broken.

The two numbers can no longer contradict each other. The branch that reads "already makes your 24 pans" is now only reached when the day really does make 24. When the day makes more, you get the way down and the honest reading of the gap: "This scenario makes 36 pans, and you want 24. That is 12 pans more than you asked for."

What has not moved. Not one module, batch time, cycle or saved scenario. Your own saved line still reads 24 pans across 9 modules, the seeded starting line still 12 pans with the chiller as the wall, and your sister's still 4 pans in 8 modules. Nothing here reads an order and nothing here blocks a sale.

Engine 149. Tests are 1110 passing with none failing, eight of them new - four on the move itself (the 36 -> 24 settling move on all eight modules, a day already at your number offering nothing, a dragged batch time surviving the move untouched, and a batch too big to come down past) and four on the card (its heading and button, the two numbers never contradicting each other, the button actually bringing the day down, and the plain answer when a batch is too big). No database step - nothing to run in Supabase.

22 Sep 2026 - engine v148 (no database step). A cycle's labour drawn where you work it, and the ladder's button found.

A cycle's labour is now drawn where you work it, and wide enough to see. This is your own sentence from 22 September: "cycle should be drawn with its labour shown at the appropriate timing." It was being drawn at the right minute already - the fold at the end of its rest, the oven swap after its bake - but it was drawn too small to find. A one-minute fold is under two pixels at the scale the day is read in, and where it fell it landed exactly on the seam between two rests, so what you saw was the seam and not your hands. Two changes, and neither of them moves a minute:

- **The width has a floor.** A stretch of your hands inside a bar is now never drawn thinner than 4 pixels, which is the same floor the person row below already gives the same stretch of your day. So the bar and the person attending it now agree about what a one-minute job looks like - before this, the fold was 1 pixel on the bar and 4 pixels on the person row, which is two drawings of one fact disagreeing with each other. **Each end of it wears a fine light edge**, so a band that lands on the seam between two cycles still reads as a band rather than as the join.

Only the width is floored. **The position is the model's own minute**, so the picture still says WHEN: the fold sits at the end of its rest and the oven swap sits after the bake. Hover or long-press a bar and it now names the cycle and the exact minute, so the true length of a floored band is never lost - it reads, for example, "1 min of you at 4:51 am".

The button that applies the ladder is now on the ladder's own heading. You told me you could not find **Use these numbers**. You were right, and the reason was arithmetic: it sat at the very foot of The climb card, and a day of nine modules makes a nine-rung ladder, so on a phone the button was about four screens below the heading it belongs to. It is now on the heading as well, beside the words The climb, and it is still at the foot of the card. Both press exactly the same thing.

When there is nothing to apply, there is no button - and the card says so. If your scenario already makes the number you asked for, the card reads "This scenario already makes your 24 pans" and offers nothing to press, because there is nothing for it to change. That is the other half of why the button could not be found: on a scenario that already reaches your number it does not exist at all. Raise the number you want above what the line makes, and the ladder appears with its button on the heading.

What has not moved. Not one module, batch time, cycle or saved scenario. The day still reads the same numbers it read in v147: One baker day is still 24 pans in 8 modules, one person, 3 h 52 min of hands, 4:01 am to 1:30 pm. Nothing here reads an order and nothing here blocks a sale.

Engine 148. Tests are 1102 passing with none failing, five of them new - they pin the fold's drawn minute, the bar and the person row agreeing about a one-minute job, the oven swap being drawn after the bake, and the apply button being on the heading. No database step - nothing to run in Supabase.

22 Sep 2026 - engine v147 (no database step). Module, batch, cycle - the three words on the Scenario planner, and the model that moved with them.

What you asked. In your own words: *"let us realign our terms use in the production line"* - a module instead of a brick, batches inside a module, and cycles inside a batch, with a batch's time moved by buttons instead of a drag. The words and the model under them were changed together on purpose: renaming first and building the cycles later would have left the word **cycle** meaning two different things while you were learning it.

The three words, and what each one is. A **module** is the equipment plus the hands tending it - what this app has called a brick until now. A **batch** is one lot: the six pans of one oven load, one tub of dough. A module can run several batches in a day, and how many is yours to set. A **cycle** is one step of that module's work, and it is

the cycle that owns the minutes and the labour. Your own sentence settled it: *"batch is not a process but cycle is."* So a batch carries no minutes of its own - the time it holds is the sum of the cycles inside it. A batch with a single cycle is drawn on the day as one plain bar, exactly the way a module has always been drawn.

Labour now sits at its real end of a cycle: its load, or its unload, or both. Open a module and under **Cycles in one batch** every cycle has a name box, **Minutes**, **Load** and **Unload**. Load is your hands at the start of that cycle, unload is your hands at the end, and the day draws them where they really happen. A cycle cannot start until the cycle above it ends, so a cycle is never offered a start time of its own - its position is set by its minutes.

Two of your eight steps were in the wrong place, and both are corrected. Your 1-minute fold was drawn two rests early; it is now the **unload of each 30-minute rest**, so the rests run back to back inside the batch - 30 minutes, fold, 30 minutes, fold, 30 minutes, fold, 30 minutes - and the fold lands at the end of the rest, where your hands really are. And your **oven swap is a cycle of its own after the 13-minute bake** rather than labour before it, so the first swap now reads 8:13 am, thirteen minutes after the pans went in. Nothing else about the eight steps moved: the tub is still 4:01 am and the first six pans are still at the oven at 8:00 am, the minute the whole day is counted backwards from.

Moving a batch is buttons now, and the drag is gone. Tap a batch on the day chart and a small card opens with that batch's clock and its own two pairs of buttons: **Five minutes at a time** and **One minute at a time**, each with an earlier and a later button, plus the batch's own times written out. Five is the step the day is read in and the step the time line reads out; one is the nudge for lifting a batch off a collision. This replaced dragging a bar completely, and it was your call: a bar being dragged across the day and the hairline of the time cursor were fighting over the same finger. Nothing about a press does arithmetic of its own - a pressed time and a typed time give the same answer - and when the day's own rule puts the batch somewhere else the card says which rule it was. A cycle drawn inside a bar is not a separate control: tap it and you get that batch's card, because a cycle's start is its batch's start plus the cycles above it.

A batch that waits on the module above offers no buttons, and says why. Where a module has **Waits for the module above** switched on, a batch's time is not its own - batch 3 of that module IS batch 3 of the module above - so there is nothing to step and nothing to mislead you. The card says exactly that, in as many words, and names the switch that would free the time. The row itself carries a **waits above** tag and the words **waits on** and the module it is waiting for, so a chained batch is never silent about why it is where it is.

Batch counts follow the one before, and the day says when one does not. Editing **How many batches it runs in the day** carries down to the modules after it that were running the same number, and stops dead where a module has a number of its own - it is never quietly overwritten. A standing line under the day then names each module whose batch count differs from the one before it, both numbers written out, with a plain sentence that it is a note and not a rule: the day runs on the numbers you have put in. On the seeded starting line, for instance, it reads **Wash, oil and fill - 6 batches, The fold loop - 4**. Nothing is locked by it and nothing is refused.

The tile that said Cycle time now says Minutes a pan. Same quantity it always was - the minutes one pan takes off the slowest module, with that module named under it - said in the words this app already uses elsewhere, because cycle now means a step inside a batch. No arithmetic moved with the name.

A module can have as many lines as the day needs, and lots go to the free one. How many of these do you have is that - a mixer, a sink, two ovens. At the even spacing the app works out, the lots take their turns down the lines exactly as before; where you have moved a start, a lot leaving one module now goes to whichever line is actually free rather than the one whose turn it is. No figure on an untouched day changes because of it.

One baker day was re-derived, and it is an hour later than the last version of it said. Correcting where your fold and your swap really happen costs the day a work slot, so the pace one pair of hands can hold is **87 minutes** between batches, not 81: the tubs now go in at 4:01, 5:28, 6:55 and 8:22 am, the packing starts at **8:57 am**, the last batch is packed at 1:18 pm and the day ends at **1:30 pm**, which is 9 h 29 min. The 81 no longer holds, and the reason is checkable rather than a matter of opinion: at 81 the second rest's fold lands on minute

82, the very minute the second tub's dough has to be mixed, so the app puts a second person on it. It still reads 24 pans a day in 8 modules, one person, and 3 h 52 min of hands, and nothing collides. Guide section 25 carries the whole story, including what the sweep of every other pace found.

What has not moved. Your own saved line still reads 24 pans across 9 modules and exactly the day you asked for; the seeded starting line still reads 12 pans with the chiller as the wall; your sister's still 4 pans in 8 modules. Not one saved module, batch time or scenario was altered - the words you read and the way you move a time changed, and the days themselves did not.

22 Sep 2026 - engine v146 (no database step). One baker day - your own bake day set on the bricks, with a latest start on every batch.

What you asked. In your own words: *"i need you to create one scenario and save it as One baker day. Set the bricks for me, with latest start time each brick batch."*

Where it is. More -> Scenario planner. Under **Your scenarios** there is one new row, **+ One baker day**, reading **One pair of hands, set for you**. One tap puts it on your shelf and opens it. From then on it is an ordinary saved scenario, exactly like your own line and your sister's proposal: it stays in your list, it opens when you tap it, and the row that offered it is gone because you already have it. Delete it the way you delete any scenario, and it stays deleted - it is never dropped back in on you.

The eight bricks are your eight steps, in your own order. Mixing the dough in the tub; the rests and the stretch and folds; oiling the pans and weighing the dough out; into the proofer; dimple and top; the proofer again; the oven swap and the bake; cutting and packing. Nothing invented, nothing left out. The fold and the two proofer bricks carry no line step against them, because the Production line has no field for them - and the screen names them rather than writing their minutes onto a field that means something else.

Every brick's first cycle is a latest start, which is what you asked for. Open the day and the eight bricks read:

- Mixing the dough in the tub - **4:01 am** (20 minutes, all of it yours) The rests and the stretch and folds - **4:21 am** (123 minutes, 3 of them your hands) Oil the pans and weigh the dough out - **6:24 am** (15 minutes of you) Into the proofer - **6:39 am** (45 minutes, no hands) Dimple and top - **7:24 am** (6 minutes of you) The proofer again - **7:30 am** (30 minutes, no hands) The oven swap and the bake - **8:00 am**, the anchor: the minute the first 6 pans have to be at the oven (2 minutes of you) Cutting and packing - **8:27 am** (12 minutes of you)

The four batches sit **81 minutes apart**, so the second tub starts at 5:22 am, the third at 6:43 am and the fourth at 8:04 am. The last batch of the day is packed at 12:30 pm and the day is done at 12:42 pm.

One of the eight is not where your chain puts it, and that is the whole finding of the scenario. Your own chain would start the cutting and packing at **8:15 am**, fifteen minutes after the oven. Alone, you are still folding the last tub at 8:27, so **8:27 is the first minute that job needs nobody else**. It is not a preference: start it a single minute earlier, at any minute from 8:26 backwards, and the day asks for a **second pair of hands**. The window you can move it in without buying anybody is **8:27 to 8:33 am**, and the seeded day sits at the start of it.

Why the batches are 81 minutes apart and not 40.5. This is your own proofer figure, from the same morning: a batch is in the cabinet 81 minutes - 45 in, 6 out to be dimpled, 30 back in - so the cabinet could take a new batch every **40.5 minutes**, about 8.9 pans an hour. But 40.5 assumes somebody is free to carry a batch out and the next one in the moment the cabinet empties, and with one pair of hands there is no such minute: the hand-work of a single batch is 58 minutes by itself. Rather than guess, every rhythm from 40.5 upwards was swept through the app's own arithmetic, and **81 minutes is the pace one pair of hands can hold**: every hand-window in the day then tiles exactly, each one beginning where the last one ends, and the proofer is never asked to hold two batches at once. It still works at 80.5 minutes; at 80, or at 40.5, the very same day reads **2 people**. So the two figures you have are both right and they answer different questions: the card's 40.5 is what the cabinet allows when you are not alone, and this scenario's 81 is what one baker can actually feed it.

What the day reads. 24 pans a day in 8 bricks - four batches of the six pans you bake. **One person, 3 h 52 min of hands** across the whole day, and a run from **4:01 am to 12:42 pm**, which is 8 h 41 min. There are **no red notes** under the day, because nothing collides: one pair of hands, no job doubled up. The brick setting the

pace is the **mixing**, and seven others pass the same 24 pans - so the card says they are one limit between them, and relieving only one of them will not move the day.

It loads onto the Production line as your own numbers. Under **Start from a scenario** at More -> Production line, tapping **One baker day** shows you everything it would change before anything moves: minutes to mix one tub 20, oiling the pans and weighing the dough out 15, dimple and top 6, take 6 pans out and put 6 in 2, minutes the pans are in the oven 13, pans in one load 6, cutting and packing 12, the target 24 pans, and 1 person. The fold and the two proofer bricks are listed by name, because that screen has no field for them. Nothing is written until you press the button.

Every number is yours to move. The day is seeded, not frozen: drag a bar, open a brick and type a number, change how many pans a day you want, or move the time your day starts, and the whole thing answers again the way every other scenario does.

What has not moved. Your own saved line still reads **24 pans across 9 bricks** and exactly the day you asked for, the seeded starting line still reads **12 pans with the chiller as the wall**, and your sister's still **4 pans in 8 bricks**. One baker day arrives as a row you choose: not one brick, start time, merge or saved scenario of yours changed, and nothing here blocks a sale.

CHANGELOG and the guide (section 25) both carry all of this, both PDFs rebuilt, and the tests are 1090 passing with none failing. No database step - nothing to run in Supabase.

22 Sep 2026 - engine v145 (no database step). Your real bake day on the Production line - and the same day read backwards from the oven, so no dough is ever mixed too early.

What you asked. In your own words: *"work backwards, from end of process, the previous process would have a latest start time, by going this way, we prevent preparing dough too early and prevent dough from over fermented. But we can fine tune each process to start earlier to better manage manpower utilisation."* You had corrected your own step order that morning, after checking it with your sister: *"i think my self confused, i double check with my sister and now understood the right flow. we cannot mix the dough too early, it will over fermented."* And you had already worked your ceiling out yourself and asked about it: *"6pan sit in it for 81min, proofer able to hold 12pans, so effectively the poofer can produde 6 pan/40.5min? So the proofer became the bottle neck?"* Yes - 40.5 minutes exactly. This release gets that figure onto the screen out of your own numbers rather than out of a guess.

The flow on the Production line was not your day. It read mix, then the chiller, then washing and oiling the pans, then weighing the dough out, then the oven. Against your day that put the chiller where you have none, transposed the oiling and the weighing out, and - the worst of it - had no fold in it at all. The map now runs your order, top to bottom: mix the dough in the tub; the rests and the stretch and folds; oil the pans and weigh the dough out; into the proofer; dimple and top; proofer again; the oven; then cutting and packing. The proofer is drawn twice because your day uses it twice, with the dimple in between.

The wall was a chiller you have not bought. The screen watched a chiller of 12 trays over 5 hours and made that your ceiling - 12 pans a day - which is a cap on a cabinet that is still only a what-if on your planner, kept switched off. The station you actually own is the proofer, and it is the honest ceiling: a batch sits in the cabinet for the whole 81 minutes (45, then the dimple, then 30), so 12 pans means one batch every 40.5 minutes and no faster. Change **Pans your proofer holds at once** to 18 and the screen answers 27 minutes instead, and says which of your limits is the wall now. It never tells you to buy one: at one pair of hands a bigger cabinet buys you nothing, and the screen says so.

Your day, read backwards. A new card beside the flow, **Your day, backwards**, gives every step the last moment it may start, worked back from the oven - the end of the process, which is where you asked to start from. For the first six pans to be standing at the oven at 8:00 am, the dough goes into the tub at 4:01 am:

- Mix the dough in the tub - 4:01 am The rests and the stretch and folds - 4:21 am Oil the pans and weigh the dough out - 6:24 am Into the proofer - 6:39 am Dimple and top - 7:24 am Proofer again - 7:30 am The oven, bake and swap - 8:00 am

That is **239 minutes** from the tub to the oven and **254** from the tub to the pans coming out, which is your own

arithmetic and nothing else. One batch is **46 minutes of your hands**: 20 mixing, 3 folding, 15 oiling and weighing, 6 dimpling, 2 swapping.

The clock is a field, not a rule. The minute the first batch must be standing at the oven is yours to set - it arrives at 8:00 am - and every start time below is worked back from it. Move it and the whole day moves with it.

Five minutes early is allowed, and nothing is locked. Every step also carries a five minute band above it: a step may begin up to five minutes early to suit your hands, which is the room you asked for to manage the people. It is advice drawn on the screen and it is never a gate - no time here stops you baking, hides an order, or overrules your own eye. You watch the dough; the timetable watches the clock.

Five numbers changed what they were asked, and each one is named here rather than changed quietly. Your saved numbers move across by themselves the first time your phone opens this version, so you will see these the moment you load it - and every one of them can be typed back in one tap if your own timing says otherwise:

- **Minutes to mix one tub of dough**: was 6, now 20. The old box asked how long it took you to load a mix; this one asks how long the whole mix takes, which you gave as twenty minutes for one tub.
- Pans one tub of dough makes**: was 28, now 6. The old box asked what your mixer's bowl holds. Your cycle is one tub of dough into the pans of one oven load - six - and the mixing and folding are spread over that tub, so this is the number the chain needs.
- Oil the pans and weigh the dough out, for 6 pans**: was two jobs under two names (your washing and oiling read 20, and the weighing out was never timed), now one job at 15. Your own count of the hand-work has it once, which is the proof they were one job: with the two merged, the app's 46 minutes a batch is your 46 to the minute.
- Minutes to dimple and top 6 pans**: was 8, now 6 - your one minute a pan. The 8 was measured with the topping in it.
- Minutes to take 6 pans out and put 6 in**: was 4, now 2 - "2min covers out and in", as you put it. One turn of the oven is the baking and the swapping together, 15 minutes as before.

Everything else of yours is untouched: your two pairs of hands, your 5 hours, your 24-pan day, your 12 pans, 6 pans a bake, 15 minutes a turn, the 30-minute cooling - and your 12 minutes of cutting and packing, which you had timed and which is still yours.

A step with no minutes on it is still named, never counted as free. Cutting and packing is the one job that arrives with an empty box - every other step of the day comes seeded with your own figure - so on a fresh phone the flow says **not timed yet** and the backwards card says in a sentence that it is not counted in the 254 minutes, rather than quietly treating your time at the table as free. You had already timed yours at 12 minutes for 6 pans, so on your phone that box holds your 12 and the card's heading reads **After the bake** rather than **Still to time**. A heading left reading **Still to time** over a figure you had given would read as the screen not having noticed it. Either way the step is drawn on the backwards card as sitting **after the bake, outside the 254 minutes**, because those minutes happen while you are waiting on the next batch. Nothing was reset.

One thing on the join between the two screens reads differently, and it is a correction. When you tap a saved scenario under **Start from a scenario** on the Production line, the list of numbers it would change is now the corrected list: **oil the pans and weighing the dough out** is one row where it used to be two, because it is one job under two names, and a **retard** brick is named as something the scenario said but the line has no field for rather than quietly writing a tray count onto a field that no longer exists. The bricks themselves are untouched, and nothing is written until you press the button. What did **not** ship is a button on the Scenario planner to push a planned day's own start times onto the bricks: your bricks are timed from a day that starts at 3 pm, and the bake day here is a morning that has to be ready by 8 am, so writing those times in would have moved every saved brick. The timetable is drawn, and it is not pushed anywhere yet.

What has not moved. The Scenario planner, brick for brick. Your own line still reads 24 pans across 9 bricks and your sister's still 4 pans in 8 bricks, and both are stored exactly as they were - not a brick, a start time or a merge has changed. The chiller bricks stay where they are, switched off, because the fridge is still worth planning for.

CHANGELOG and the guide (section 24) both carry all of this, both PDFs rebuilt, and the tests are 1084 passing with none failing. No database step - nothing to run in Supabase.

21 Sep 2026 - engine v144 (no database step). The People row counts the places one person is in, a line of a brick counting as a place - and a collision is written so it can be read.

What you asked. In your own words: *"fix the hats count and the collision sentence. And what is the concept of hat, i dont understand, why name them by numbers?"*. Both were real faults, and the question was fair - "hat" was my word, and it was a poor one. It is gone from the app and from this book; what it meant is said plainly below.

What a "hat" was. It was my shorthand for **the places one person has to be in the day**. You put it better yourself: *"one person would not need to wear too many hats"* - one worker asked to stand at six different jobs in one day is being handed six hats. That is the thing the People row is there to count, so that is what it now says: **in 6 places**.

The count was wrong. Each person's row read **at 6 bricks** - the number of different bricks, not the number of places - so a worker on **line 1 and line 2 of the same fold** counted as being in **one** place. That is exactly the doubling-up the number is supposed to show you, and it was the one thing the row said nothing about. A line of a brick is a place of its own now, so that same worker reads **in 2 places**. On your own line the row has gone from **at 6 bricks** to **in 7 places**, because your fold is drawn as two lines. The number only appears when it is more than one, exactly as before.

The collision sentence was unreadable. It was one long sentence - *Person 1 is at X and Y at the same time, ...* - and your bricks have long names: *Mixing by hand in the tub (set the minutes)* is forty characters, twice in one sentence, with the word *and* between them lost in the middle. It is now three parts on four separate lines. Line one, who and when: *Person 1 - two jobs at once, 4:00 pm to 4:05 pm*. Lines two and three, each job on a line of its own with a bullet in front of it: *Mixing by hand in the tub (set the minutes), line 1* and *Mixing by hand in the tub (set the minutes), line 2*. Line four, what to do about it: *Move one of them along the day, or combine with another person and accept the collision*. Two lines of one brick still say which line they are, because without that they read identically. And if a day has more than three of these, the last block **counts the rest** instead of quietly stopping - a list that stopped at four read as though the day had four problems when it had nine.

Why the people are numbers. Because a number is what the app stores: who is on a brick is a number, combining two people is joining two numbers, and the same person on three bricks is the same number in three places. That is what makes "put person 1 on both lines" work in one tap. It also has nothing to do with you - **Person 1, 2, 3** are places in the day, not names of people, and the app never shows them to a customer. Naming them (Mei, Kak, Adik) is a small change on top and I have not built it - say the word if it would help you plan.

What has not moved. The day itself. Not one minute, one pan or one figure. Your own line still reads **24 pans across 9 bricks**, the starting line still **12 pans with the chiller as the wall**, your sister's still **4 pans in 8 bricks**. Only the words on one row and the shape of one red note changed.

CHANGELOG and the guide (section 23) both carry all of this, both PDFs rebuilt, and the tests are 1062 passing with none failing. No database step - nothing to run in Supabase.

21 Sep 2026 - engine v143 (no database step). A time cursor on the day chart: point at the chart and a line follows your pointer down the whole day, reading out the clock, so you can line two bricks up against each other.

What you asked. In your own words: *"i need a time x axis cursor on the time chart"*.

Where it is. The day chart in **More -> Scenario planner**. One change, in one screen.

How to use it. Run the pointer across the chart and a fine line follows it, down through every row - the ruler, every brick, every person, the total. A little brown label at the top of the line reads the clock at that point: *5:05 pm*. Move the pointer and the reading moves with it; move the pointer off the chart and it goes away.

On a phone. A finger cannot hover, so on a phone the **clock strip is the handle** - the one band along the top of the chart where the hours are written. Put your finger on that strip and slide it left or right and the line follows, reading out the time. Everywhere else on the chart your finger still scrolls the day the way it always did, so you

have not lost the swipe.

Why it only counts in fives. The line reads the clock to the same five minutes a bar snaps to when you drag it. That is deliberate: the cursor can never name a time the chart could not also set, so you are never shown a time you cannot ask for.

At the ends of the day. Point before the day starts or past the end of it and the line goes away rather than naming a time that is not in your day. At the far end of the evening the reading hangs to the left of the line instead of the right, so it is always written out in full and never half cut off.

Two bricks lined up. This is what the line is for. Point at where a brick starts, read the time off the label, then point at where the next one starts and read that - no counting squares. It reads the chart and it writes nothing: no brick, no time and no number changes because you pointed at something.

Nothing you have has moved. Not one figure. Your own line still reads 24 pans across 9 bricks, the starting line still 12 pans with the chiller as the wall, and your sister's still 4 pans in 8 bricks.

CHANGELOG and the guide (section 23) both carry all of this, both PDFs rebuilt, and the tests are 1061 passing with none failing. No database step - nothing to run in Supabase.

21 Sep 2026 - engine v142 (no database step). Typing a number into a long card no longer throws the card - the box keeps your cursor and the card stays where you were reading.

What you asked. In your own words: *"why the display jump and it is very confusing"*. It was a real fault, and this is it fixed.

What was actually happening. The brick's card is a long card - fifteen boxes, a box per line, a row for every cycle - and **every single number you typed into it was rebuilding the whole card underneath your finger**. That is the jump you saw: the box you were typing in was thrown away and a brand new one put in its place, so the cursor went with it and the keypad shut. On a phone that reads as one digit per tap, on a card that had moved somewhere else. It was worst on the three boxes that change how tall the card is - typing a 6 into *How many of these do you have* grew the card by a line's worth of boxes on every keystroke.

What it does now. Typing a number writes the number and **leaves the card alone**. The box keeps your cursor, the keypad stays open, and the card stays exactly where you had scrolled to - checked box by box: the brick's own numbers, the people box, and every line box under *How many of these do you have*. Three boxes still change the card's shape, because their answer really does add a row or move a time - **How many times it runs** (one more row per cycle), **How many of these you have** (one more line box), and **Minutes in, when its first pass starts** (cycle 1's own time). Those three now redraw **once, when you leave the box**, instead of once per digit.

And it is not only the brick's card. The jumping was in the shared pop-up that every screen in the app opens, so the same fix covers all of them: any card long enough to scroll keeps its place while you work down it, and a card you open fresh still starts at the top where it should.

Nothing you have is affected. This changes only when the screen redraws, never what it redraws: every figure, every scenario and every saved line reads exactly as it did in v141. Nothing new to run in the database.

21 Sep 2026 - engine v141 (no database step). A brick you have two of is now drawn as two lines, and each line carries its own person - so one worker stands at one job instead of wearing every hat in the day.

What you asked. In your own words: *"i want is just better drawing.. Since now one brick having 2 lines, each line can have its own person, and few line can have a combine person sharing to different lines. This will make training of worker easier, because one person would not need to wear to many hats"*. So this is the **drawing**, and **who is on each line** - nothing about how the day adds up has moved.

A line, and why a brick you have two of is two of them. Everything on this screen was already worked out as if a brick you have two of were two streams taking the lots in turn - lot one and lot three on one, lot two and lot four on the other. The screen was drawing those two streams flattened into a single row. Now each one is a row of its own: the brick's own name, badges and minutes on the first line, and a row per further line underneath, tied

together by an accent down the left and a dotted seam between them.

Who is on each line is yours to set. Open the brick and, once *How many of these do you have* says 2 or more, you get a box per line - **Line 1 - who is on it, Line 2 - who is on it.** 0 means whoever is free. Put the same number on two lines and that is one person covering both, which is exactly what you asked for; if those two lines really do need them in the same minute, their row in the People list goes red and names the minute and which line collided with which. The line boxes and the brick's own *Who is at this brick* are the same number, so the two can never disagree.

The People list now counts the places one person has to be. Each person's row reads how many different places that one person is in - and every job on it names its line ("Stretch and fold in the tub, line 2"). That number is the whole point of this release: it is the one that tells you when one person is being asked to be in too many places at once.

*(This entry was first written as "counts hats", as **at 3 bricks**. Both the word and the count were wrong, and v144 fixes them: the row now reads **in 3 places**, and a line of a brick counts as a place. See the v144 entry at the top.)*

Your starting point in one tap. One a line (it was *One a brick*) gives every line its own person, so a brick you have two of takes two people rather than one. It says *One a line* only when a brick in the scenario really is drawn as lines; otherwise it reads exactly as it always did. Its name and the message it leaves say which of the two it just did.

What has not moved. With no lines drawn and nobody set - which is every brick you have - the day is computed exactly as before, and every figure you have already been given reads identically: your own scenario still 24 pans and exactly the day, the seeded line still 12 pans with the chiller as its wall, your sister's still 4 pans in 8 bricks. Two of a brick still buys cycles out of the pans you planned; it does not bake pans you did not plan.

One thing it deliberately does not do. With **Let its cycles overlap** switched on, a brick's lots no longer take turns, so there are no lines to draw and the brick goes back to one row with one person. The people you set per line are **kept**, not thrown away, and the brick's card says so - switch overlapping back off and your lines come back, each with the person you gave it.

Nothing to run in the database. This is a screen and a drawing; no SQL, no new table.

21 Sep 2026 - engine v140 (no database step). A brick can now be told that its cycles may overlap - one switch, with the rule for when to use it written beside it - and the two lots are drawn as two bars.

What you asked. In your own words: *"i also want to allow each brick cycle to overlap, this will improve the visualisation"*, and then, as I was building it: *"allow overlap button and the overlapping criteria"*. So it is a **button** in the brick's card, **Let its cycles overlap**, and the criteria sit under it in plain words rather than in a rule you would have to work out.

What a brick did before, and still does unless you switch this on. A brick held its own cycles apart: lot two could not begin until lot one was out of it. That is the right rule when the dough is physically *in* the thing, and it is what stops you folding dough that is still in the mixer. Where you switch it on: when the minutes in a pass are the **dough's** time and not a machine's - dough resting between folds, a second bin on the go. Then every cycle sits exactly where you put it.

How you can tell it is on. The brick's row wears a **2 at once** tag naming how many lots are in it together, and the row grows taller so the overlapping lots are drawn as separate bars in their own lanes instead of painting over each other. A brick that is not overlapping looks exactly as it did - same bar, same height, same place.

Three things it deliberately does not do. It never lets a lot start before the dough exists: a brick that waits for the brick above still waits, so switching this on is not a way around the chain. It does not put pans in the day that you did not plan. And it does not quietly double you up - if two of your own cycles need the same person in the same minute, that person's row still goes red and names the minute.

And it is not the same as a second brick. **How many of these do you have** means one lot in per brick you have - with 2 of them, lot three may start while lot one is still folding, but lot four waits for lot two. **Let its cycles**

overlap means every lot at once. The switch is the wider of the two.

One correction to v139, written down so it never looks like a bug. v139 said a brick whose passes come round faster than one pass takes "now has its passes held apart instead of drawn overlapping - that day could never be worked". That is still what happens **with the switch off**, which is how every brick arrives. With it on, those passes are drawn overlapping, in lanes, which is the whole point of the switch.

Nothing else in the app moved, and both of your scenarios read exactly as they did - the sister's line still 4 pans in 8 bricks. Tests are 1048 passing with none failing, and the guide (section 23) carries all of it.

21 Sep 2026 - engine v139 (no database step). The scenario planner now builds a real line: a brick can be two of them, every pass in a brick has its own start time you can drag or type, and a brick can be set to wait for the brick above it - lot 10 waits for lot 10.

What you asked. In your own list: *"1. allow me to set the same brick to 2 or more, if that module became bottle neck.", "2. in a brick, let say we might need 8 cycles, each cycle should be able to have its own start time, drag-able.", "3. 2nd brick 1st cycle should follow 1st brick 1st cycle completion, and so on.", "4. Each brick, cycle10 only starts after one brick earlier cycle 10 ends, etc.", and "5. objectives are: less man, less manhours thru better planning, shorter hours."* Then, as I was building the screen: *"from your layout, the stretch and fold should have 2nd, 3rd brick as the person who do that process and start another cycle without having to wait for the 1st cycle to finish"* - which is what **How many of these do you have** does: a second fold starts the next lot without waiting for the first to finish.

What changed. Everything below is inside More -> Scenario planner. Nothing else in the app moved, and every figure you have already given reads exactly as it did.

Every pass in a brick is drawn on its own, and dragging one moves only that one. Before this version a brick's passes were a single rhythm - a gap and a count - and every pass sat where that rhythm put it. Now each pass is **its own bar on the timeline** and can be dragged by itself: cycle 3 moves and cycles 1, 2 and 4 stay where they were. The toast when you let go names the one you moved - *"Cycle 3 -> 4:15 pm"* - so there is no doubt which pass changed. For a small nudge, a number beats a careful finger: the brick's card now has a **row per cycle**, each with a box for the exact minute and the clock reading beside it, so a five-minute correction is typed rather than dragged. And **Space them evenly again** puts every pass back on the even rhythm from the first one, once you are done experimenting.

A brick can be two of them - or more. A new box in the brick's card, **How many of these do you have**, and the row wears a **2 of them** tag when you raise it. What a second one really buys you, and what it does not: it lets the brick hold **twice as many passes in a day**, and it runs at **twice the pans an hour** - but it does **not** change the pans you planned, because a second mixer standing idle does not bake anything. To use it you raise the passes. That is the honest reading of buying one: it removes a wait, it does not invent work. With **2 of them**, that brick's own passes are allowed to overlap, because two machines can run at once - which is the whole point of having two.

A brick can wait for the brick above it - and that is the chain you described. A switch in the brick's card called **Waits for the brick above**, and it does exactly your points 3 and 4: lot 1 of this brick cannot start until lot 1 of the brick above has finished, lot 2 waits for lot 2, and so on down the line. Turn it on for the bricks that really are one line and a slow pass upstream holds every later lot behind it - which is precisely why a second fold, or a second mixer, is worth buying. Two things worth knowing. **A chained bar cannot be dragged**, because where it sits is decided by the brick above: those bars are drawn faded and wear a **waits above** tag, so nothing on the screen looks like a control that does nothing. To move a chained lot you move the cycle it is waiting on, and the whole line moves with it - which is the lesson the screen is for. And **if two bricks have different numbers of passes**, the screen says so instead of guessing: passes past the end of the brick above follow its *last* pass, and the row names it. **The chain arrives switched off**, so nothing you have already measured moves until you turn it on yourself.

The line, in one tap - and your three objectives as three numbers. You named the goal as three things, so the screen now says all three in one line under the day: **2 people · 13 h 40 m of hands · a 16 h 20 m day.**

People is the most hands needed at once (*less man*), **of hands** is every person-minute of the day added together (*less manhours*), and **the day** is first job to last (*shorter hours*). Drag a cycle, move a start, add a second brick, and all three answer again at once - so a change either improves all three or you can see which one it cost. Beside the timeline a chip reads **3 waits above**, or **Not chained**: tap it and it explains the rule, names which bricks are waiting, shows those same three numbers, and offers the whole thing in one tap - **Chain the whole line** or **Take the waiting off** - so you never set eight switches by hand.

The climb can now offer a second one, not only more passes. The ladder used to relieve the wall by raising that brick's passes, and stopped where a day could not hold more of them. That stopping point is exactly where a second brick is the real answer, so a rung can now read **2 of them** as well as **more passes**, with the same before-and-after pans figures. A rung that adds a second one says what it costs in as many words - **that is a second one to buy** - because that rung is money where the others are time. A brick that needs both is offered both, in order.

One arithmetic change worth writing down, so it never looks like a bug. A brick whose passes come round faster than one pass takes - a 5-minute gap on a 28-minute fold - was always a drawing of something that cannot happen: the same hands cannot start pass 2 while pass 1 is still on them. Those passes are now held apart to one at a time, when the brick is the only one of it and is not chained. None of your own bricks is built that way and every figure you have given reads exactly as it did, but if you ever type a gap shorter than a pass, the screen will hold the passes apart rather than draw a day you cannot work.

21 Sep 2026 - engine v138 (no database step). The scenario planner is yours to build in: bricks you make, move, give to a person and drag along the day, a ruler and a start time, scenarios you save and name - and the scenario you choose can now be loaded onto the Production line.

What you asked. In your own list: *"1st i cannot change the time scale, say i want a strat at 8am. And i need an time scale ruler", "I need to move the sequence of the lego brick", "I want to be able to create, edit and delete lego brick", "i want to be able to save and name a scenario. Edit or delete", "how to have person ??, I would want one person for each module, And a Total person showing overlapping persons time slot if any", "The lego brick should be able to restart? if choosen to, and how, can it be drag to move the start time of each brick", and "The fold brick should show 28min, and its person should show 2 minutes, when i adjust the brick until mininum total person overlap, that will be best scenario for manpower planning."* Then, as I was building it: *"can allow to save a scenario? scenario and production line has some function overlapping, the chooses scenario should be able to be loaded to production line"*.

What changed. Everything below is inside More -> Scenario planner, plus one card on More -> Production line. Nothing else in the app moved.

The day starts when you say, and there is a ruler to set it by. One box at the top sets **the time your day begins**, and every clock reading on the screen is counted from it - so *"say i want a strat at 8am"* is one field rather than ten start times to retype. A ruler now runs under the timeline: a tick every half hour with the o'clock ones named, so a bar can be read against the clock instead of guessed at. And the **time scale** is yours: four chips - **Wide, Standard, Close, Closest** - change how wide a minute is drawn, from the whole day on one phone screen to close enough to tell a 15-minute bake from the one before it. The hour lines on the timeline follow whichever you pick, so grid and ruler always agree.

A brick is now a thing you own, not a row you read. + New brick adds one, and it arrives **switched off** so an empty brick cannot become the wall and answer your day with a nonsense figure. Open any brick and you can rename it, change its picture, change every number in it, and - this is new - **delete it** or move it **up Earlier / down Later**, because the list order is the day's order: the timeline is read top to bottom and a tie for the wall goes to the earlier brick. So *"move the sequence of the lego brick"* is two buttons, and the day re-sorts around it.

One person for each module, and the Total person you described. Every brick can be given to a person: **0** means whoever is free, **1, 2, 3...** names someone. Two chips do it in one tap - **One a brick** gives each job its own person, **Share them** hands them all back to whoever is free. Under the timeline there is now **one row per person**, each showing the jobs that person is at and how much work they carry, and where two of their jobs

overlap the bar wears a red outline with a sentence naming **both jobs and the minute they clash**. Nothing blocks a plan that cannot work - the whole point is that it is *shown*. The **Total person** row stacks person 1, person 2, person 3 and the rest on top of each other, which is what you asked for: *"with that i can see, which time period utilise more manpower"*. It reads out the most people needed at once, the busiest stretch by name, and how many minutes are being worked two at a time. Bring the bricks closer together and that row flattens - that is the manpower planning.

Combining two people, which was your own example. Combine two... asks two questions: the person who keeps the job, then whose bricks move to them. Person 1 then person 3 makes the row read **Person 1+3** - *"to generate person 1_3"*, in your words - and whatever then collides is exactly what that combination cannot pay for. Give one brick back in its own editor and the **1+3** label prunes itself the moment person 3 has a brick again, so a label can never claim somebody who is not really there.

A brick restarts, and a brick can be dragged. Same number, two ways to move it. Type the gap and the number of runs and the brick restarts itself later in the day - your fold, four times. Or take the bar on the timeline with your finger and **drag it**: the start moves in 5-minute steps, the row's times update as you drag, and letting go will not open the brick's card. Drag to move, tap to open. Your pans do not change when you drag; only your people do.

Your fold brick now reads the way you described it. It holds for 28 minutes, it is **2 minutes of you**, and it restarts 30 minutes after the last fold, four times. One honest note: if your phone's fold brick still reads **30 in minutes one pass holds it**, that is the copy your phone saved from the first version of this screen. Type **28** into that box and it is right. I have not overwritten it for you, because your numbers are never changed by an update - that is a rule of this app, not an oversight.

Save a scenario, name it, open it, rename it, delete it. Save changes keeps the one you are in; **Save a copy...** makes a second named one and carries on in it - which is how you keep the line you have and the line you are thinking of buying side by side. **Your scenarios** lists everything saved: tap to open (the row you are in wears a filled dot), or tap the **pencil** to rename or delete it. Deleting a saved scenario never touches the bricks you are working on.

Your sister's line is in there too, built from her own words. Next to your saved scenarios there is one more row - **+ My sister proposal 21/9/2026** - and one tap puts it on the shelf and opens it. It is the line she wrote out, and it has **no mixer and no chiller in it**: everything mixed by hand in one big tub, folded in that same tub for 2 minutes every 30 minutes three times, tipped into four pans, an hour in a proofing cabinet, then dimple and oil in the pan, half an hour of rest, the oven, and cooling and packing. As she described it that is **4 pans a day in 8 bricks**, and the thing stopping it is **mixing by hand in the tub** - not the oven. Bring the bricks closer together and the Total person row will show you where the three folds collide with filling the pans.

Two of its numbers are mine, not hers, and both say so on the brick. She never gave a mixing time - she mixes by hand, so there is no machine time to give - and I have seeded it with the length of your own mixer pass so the brick is not blank. The **Into the pans** brick is worked out from **your** measured 3 minutes a pan, which is 12 minutes for four. Everything else is her own numbers, and every number on every brick is editable. The tub mixes **4 pans in one go**, so if her tub is bigger, change the batch on that brick and the whole day re-reads itself.

One repair under this. Opening a saved scenario used to leave the box at the top of the screen still reading the scenario you were in *before* it - the name, the start time and the pans you want. It now repaints with the scenario you actually opened, so the top of the screen and the row you just tapped always agree.

And the chosen scenario loads onto the Production line. A new card at the top of More -> **Production line - Start from a scenario** - lists everything you have saved. Tapping one shows you **every number it would change before anything moves**: what it was, what it would become, and which numbers it would leave exactly as they are. So that the two screens can agree without retyping, every brick's card now carries one more line - **which step of your line is this** - where you say whether it is the mixer, the wash, weighing the dough out, the topping, the oven, the cooling or the retard. A brick's **minutes** and its **batch** are what cross over; a brick that is

not a step of the line (your fold, loading the chiller) is named as left out rather than silently dropped. **How many pans you own** and **how many hours you will bake for** stay exactly as they are, because a scenario does not know them and the screen says so instead of filling them with a nought. Nothing is written until you press the button, and the button counts the numbers that would really move - so a scenario you have already loaded says *"nothing here changes"* rather than promising eleven.

21 Sep 2026 - engine v137 (no database step). More -> Scenario planner: build your line out of modules, watch the day as a timing diagram, and raise the pans you want until something stops you.

What you asked. "can i have a apps function that visualise the flow... something for me to optimise line balancing", and then, once we were talking about it: a screen you could **configure**, a **timing diagram**, and - this was the important one - *"not optimised for mobile, for a clearer interface"*, so you could "toase about all possibility to configure the production flows". You described the parts yourself: a module around the oven, a module around the mixer, a module around a fermentation, each with its **own cycle time**; the folds as a **loop**; ingredients going in "from time to time, not one shot"; and the fridge and the bins as **half-finished goods** sitting between steps. And you set the rule for the whole exercise: *"alll this are design, and require proofing, and before the proofing, the fugures have to be right first, at least logically."*

What changed. A new screen, **More -> Scenario planner**. It is the one screen in the app that is **wider than the rest** - up to 900px instead of 540 - because a day laid out across a timeline needs the room. Everything else about the app is exactly as it was.

A module is a Lego brick: a machine and the pair of hands that tends it, counted as one thing. Each module has its own cycle time - how long one pass holds it - how many pans that pass deals with, how many minutes of you it takes, and how often it repeats. Your fold loop is just a module that repeats more often, so a loop needs no special anything. Ten modules come seeded: the mixer, the fold loop, wash/oil/fill, load the chiller, retard overnight, unload, dimple and top, the oven, cool and pack - and the fridge, which arrives **switched off**.

Scenario 1 is the line without a fridge, and the fridge is a switch, not an absence. It sits in the module list switched off and reads *"not in this scenario"*, so the screen answers without it. Switch it on and it appears. That is deliberate, because you said it plainly: *"if the output fulfilled my need now, then fridge would not be the 1st thing to buy for me"* - so the buying question and the design question had to be the same screen, one tap apart.

The day is a timing diagram. One row per module running across the day, with the clock along the top every two hours. A **solid block is you standing there**; a pale one is the module running without you - the retard and the bake have no solid block at all because they need nobody. Below the modules is a row **per person**, carrying the jobs that person is at. There is one person on your seeded day, doing 4 h 56 min of work. Swipe sideways; the day is wider than the phone.

The climb - this is the part you asked for. You described the method yourself: *"i will slowly increase the output needed per day until 1st limitation hit, tweaking the 1st limitation and expecting the output to grows, until the next limitation surface"*. So: type how many pans a day you want, and a ladder appears. Each rung names the one module that stops you at that number, exactly what to change about that module, what the day becomes, and which module becomes the wall next. **One tap puts the whole ladder into the modules.**

And the ladder has to be honest, which is where it earns its keep. Your seeded day makes **12 pans** and it is the chiller that stops it. Raise the target to 36 and the ladder climbs to **24** - then stops, and says why in plain words: *"A 12 h pass fits 2 times in a day, and it is already running that often. So this one cannot be run more often - it has to take more pans at once."* Twelve trays, twice a day, is twenty-four pans and no amount of wanting makes it more. **Thirty-six pans needs eighteen trays.** It is the same trays-to-a-day arithmetic from the Production line, now holding this screen to account - and it means this screen can talk you **out** of a purchase as well as into one.

Every figure is yours to move. Tap any row - or any brick in the "Not in this scenario" list - and one card opens with that module's numbers, each label naming its own unit: minutes one pass holds it, pans in one pass, minutes of you, minutes from one pass to the next, passes in the day, minutes in when the first pass starts, people this pass needs. Nothing here reads an order and nothing here blocks a sale; it is a planner you type into.

One more honest note. A day can only hold so many passes of something: a 12-hour retard fits twice in 24 hours, and the screen will not count a third however many you type. Your number is kept exactly as you typed it - the row just wears an *"a day's limit"* tag and counts the passes that are really there.

20 Sep 2026 - engine v136 (no database step). The Production line now opens with a flow map: your bake day top to bottom, in the order it actually happens.

What you asked. "can you do a flow map so that i can visualise better, make a intuitive one". The four bars answered *how fast* each part is, but nothing on the screen showed the **order** - that the dough is mixed first, rests overnight, and only then goes pan by pan through the oven and back round again. Numbers side by side don't tell you what happens next.

What changed. A new card called **The flow** now sits at the top of the screen, above the bars. It is the bake day written down the page, one step under the next with a numbered dot and a line running between them, so your eye follows it like a recipe:

> Weighing in and loading the mixer -> Retard overnight in the chiller -> Wash, oil and > fill -> Weighing the dough out into pans -> Dimple and top -> The oven swap and bake -> > Cooling and packing -> The day.

Each step carries its own cost: the hand steps say **minutes a pan**, the chiller, the oven and the pans going round say **pans an hour**, and a step you have not timed says **"not timed yet"** rather than quietly reading as free. The last node states the day the line actually makes. The four bars you already use are still there, unchanged, directly below the map.

The step that sets your pace is the red one. The same red the bars use, so the two agree. Whichever step is the wall wears a **"the slow one"** tag and the one sentence that explains why - at your current numbers that is the chiller, holding 12 trays; fill in the three untimed steps and the mark moves to the wash, and the traffic goes to your hands instead.

One honest note about the hand steps. The six hand jobs are one shared pool - you do them in turn - so no single one of them has a pace of its own. The map therefore shows each as **minutes a pan of work** and never as a rate, and when the wall is the hands the "slow one" mark goes on the heaviest job while the sentence beside it says the minutes are shared across your pairs of hands. Only the chiller, the oven and the pans are given a rate, because those are the only three that are not you.

20 Sep 2026 - engine v135 (no database step). The mixer box on the Production line now explains itself, and its number is right: 25 kg is about 28 pans, not 25.

What you asked. "Most dough in one mix, in pans - What is this?" Fair question, and the fault was mine. Your form asked for your mixer in **litres or kg** ("Most dough your mixer takes in one go - look for the bowl size on the mixer or in the manual"), and the Production line asks for it in **pans**, because every other number on that screen is in pans so it can be compared to the rest. I converted your answer without telling you, and never checked which unit you had meant. So the box read 25 and quietly meant "25 pans" - a figure you had never given.

What changed. You told me your mixer takes **25 kg** of dough. One pan takes about **900 g** of dough - that is what bakes down to the 800 g you sell, the rest being water that bakes off and oil that goes in the pan - so 25 kg is about **28 pans**. That is now the figure the box is seeded with, and the hint under it says the whole thing in one line, so the conversion is on screen rather than in my head: *"Your mixer's bowl size, counted in pans: one pan takes about 900 g of dough - that is what bakes down to your 800 g - so a 25 kg bowl is about 28 pans."*

One caveat, said plainly. A 25 kg mixer means the bowl holds 25 kg, but a wet focaccia dough climbs the hook and will overflow a bowl filled to the brim, so what you can actually mix in one go is probably less. The hint says so too - lower the number if you never fill it. You know your mixer better than the arithmetic does.

If your box still reads 25, type 28. The figure the app remembers for you is the one you last typed, so your own phone keeps showing 25 until you change it - the new seed only reaches a fresh phone. This is the one small thing to do by hand.

What it affects. Two things, and both are small at your scale: how many mixes a day takes (the day's pans divided by this), and - once you time "Minutes to weigh in and load one mix" - the mixing minutes divided

across each pan. At 28 pans a day of 36 is two mixes.

20 Sep 2026 - engine v134 (no database step). Every time box on the Production line now says what it wants: minutes.

What you asked. "the field input unit is not shown?" You were right, and it was worse than a missing label - it was six missing labels. Every other box on that screen carries its unit in its own name ("Minutes per bake", "Pans per bake", "Hours you'll bake for"), so the pattern was already there and the time boxes did not follow it. You saw "Wash, oil and fill 6 pans" above a box holding 18 and had to guess whether that meant minutes or seconds.

What changed. All six hand-work boxes now lead with the unit, so the box and its number read as one thing:

- Minutes to wash, oil and fill 6 pans Minutes to dimple and top 6 pans Minutes to take 6 out and put 6 in Minutes to weigh in and load one mix Minutes to weigh the dough out into 6 pans Minutes to cool and pack 6 pans

The two card headings say it as well - the times you measured are "each for 6 pans, in minutes", and the rest of the kitchen work is "the hand-work around the bake, in minutes".

Nothing else moved. No number changed, no arithmetic changed, no figure on the screen reads differently. This is the labels only, so if you had already worked out what went in each box, nothing you typed has been affected.

Nothing to set up, and no database step. Open the app and the boxes say minutes.

20 Sep 2026 - engine v133 (no database step). The Production line now counts the whole bake day, not three jobs of it: weighing in and loading the mixer, weighing the dough out into pans and cooling and packing each have their own box to time, and any job you have not timed is named on the screen rather than quietly counted as free.

What you asked. "no time allocation for handling the mixer? and weighing the dough, no other process that need manpower, or should still need some resources?" You were right, and it was a real hole in the screen. It was timing three jobs - wash oil and fill, dimple and top, and the oven swap - and treating those three as the whole day's work. The mixing, the weighing out and the packing were not in the arithmetic at all, so the day it promised you was a day nobody could actually have delivered. Worse, the plan I wrote for this screen had named cooling and packing; the screen I built dropped it.

What changed. "The rest of the kitchen work" is a new card on the Production line with three more boxes to time, the same shape as the three you already timed: weighing in and loading the mixer, weighing the dough out into pans, and cooling and packing. The list of jobs now reads all six, each with its minutes a pan or the words "not timed", and every pair of hands is placed against all six and not just three.

A job you have not timed is named, never assumed free. Leaving a box blank is allowed - you may already have that time inside your wash, oil and fill figure, and there is a line under the box saying so. But a blank box does not silently count as nothing: the screen lists the jobs still untimed and tells you plainly that the day above looks longer than it really is until you time them. A time with nowhere to go - minutes on the mixer but no batch size for it to spread over - is treated the same way, rather than being divided by nothing.

The mixer is timed per mix, not per pan. Its minutes are spread over the batch it makes, because that is the honest reading and it means a bigger mixer really does cost you less work for every pan. Twenty minutes of weighing in is one seventh of a pan's work in a 25-pan mixer and half that in a 50-pan one.

And the answer changes when you time them. This is the part worth knowing before you do. At your own numbers with the three new boxes filled in - 20 minutes a mix, 3 minutes to weigh out 6 pans, 12 minutes to cool and pack 6 - the whole day's hand-work is 8.3 minutes a pan, so one pair of hands does about 7.2 pans an hour instead of the 12 the screen used to claim. Five hours then makes **36 pans, not 60**. The wall moves too: it stops being the chiller and becomes your own hands. Every number on the screen is yours to correct, and the day figure moves the moment you type.

Nothing to set up, and no database step. Open the app and the three new boxes are there, blank and waiting to be timed once.

19 Sep 2026 - engine v132 (no database step). Your customer's own track page now keeps up with the order: anything you change about an order reaches their page when you save it, and an order you have just added has a page behind its link from the moment it exists.

What you asked. "add new order and edit order din sync?" They were, on your own phones - an order is one record, so an edit made on one phone is on the other within seconds. What was NOT in step was your customer's track page, and that was a real defect. Their page was only rewritten when the delivery day moved, the tracking number changed or the courier charge changed - so if you corrected the items, the price, the address or the customer's name and saved, their page went on showing the order as it used to be: the old items, and the old total, in the very message you had already sent them the link in. Adding an order wrote no page at all, so a link you sent straight away could answer "not found".

What changed. Every way of saving an order - adding one, editing one, the Note / tracking box, and deleting the courier's Delivery & fuel line on Money - now hands your customer's page the whole new version of the order, and the page itself decides whether anything it shows has moved. Where the app used to hold a short list of the things it believed that page showed, it now compares the entire row, so a figure on that page can never be left out of the decision again.

An added order has its page at once. Adding an order publishes your customer's page there and then, instead of waiting for the first time you happen to change its status, so the track link in your confirmation message always has something behind it.

A save that changes nothing writes nothing. The page compares itself with what you last published, so saving an order without changing anything they can see writes no page at all - and if a write is ever refused (a phone with no signal, a SQL script not yet run), the next save tries again rather than counting it as done.

Nothing to set up, and no database step. Open the app and it is already in step.

19 Sep 2026 - engine v131 (no database step). A Production line screen: your bake day drawn as a row of bars, the station that is holding you back named and made the only red thing on the screen, and a ranked list of the changes you could make with what each one actually buys you.

What you told me. "can i have a apps function that visualise the flow, bottle neck highlight and something for me to optimise line balancing, so that i can have an optimise manpower usage?" Then, when I asked how it should help you decide: "Both - ranked levers and what-if", and the numbers should come from "A planner I type into". You asked for all four levers to be yours: **People, Pans and trays, The task times, and Days and hours.**

The line, drawn. Open **More -> Production line** and your day sits on one screen: one bar each for **your hands, the chiller, the oven and your pans**, every one measured in the same unit - **pans an hour** - so they can be compared honestly. The shortest bar is the one setting your pace. It wears a red wash, is named in a sentence just underneath, and it is the only red thing on the screen, so your eye lands on the one thing that matters. Each bar carries a plain line saying what that station is: "1 pair of hands · wash, oil, fill, top, oven swap", "12 trays of dough over 5 hours", "6 pans every 15 min".

What your own numbers already say. The three tasks you timed for every 6 pans - 18 minutes washing, oiling and filling, 8 dimpling and topping, 4 for the oven swap - are 30 minutes of work for 6 pans, which is **5 minutes of work for every pan**. One pair of hands therefore finishes **12 pans an hour**; your oven does **24**. And the chiller, holding your trays of dough overnight at one tray to one pan, turns 12 trays through a 5-hour day at only **2.4 pans an hour**. So the screen names **the chiller as the wall, not the oven** - your day comes to **12 pans**, 48 short of the 60 you want - and it tells you that **one more pair of hands buys nothing at all** while the trays are already full.

The third pair of hands buys nothing. That is the manpower answer, and the screen says it in words rather than leaving you to work it out: once the chiller is out of the way, the **second pair doubles the line** (12 pans an hour to 24) and the **third buys nothing**, because the oven stops at 24. Spending on help before fixing the wall is money standing around.

Where the hands go. With two or more pairs the screen splits them by how much work each job is, so two pairs reads "1 on the wash, oil and fill, and 1 on the dimple and top", with the oven swap taken by whoever is free.

Every job finishing at the same moment is the whole trick to balancing a line, and it is the same arithmetic at two pairs of hands or at six.

What to change, ranked. Six moves are scored by the pans they actually buy on a day like yours and sorted best first - **one more pair of hands, six more trays, six more pans, a wash 10% faster, one more hour, one more pan per bake.** Each row says what the day becomes and which wall is left standing after it. At your numbers only **six more trays** gains anything (+6 pans, taking the day to 18) - and the other five keep their place and say so in as many words ("No help - the chiller still sets the pace"), so nothing on the screen ever looks like a button that does nothing.

A planner, not a gate. Nothing on this screen reads an order, nothing here blocks a sale, and nothing here asks you to keep anything in step with it - the numbers are yours to correct right on the screen, and the screen's answers move as you type. It is seeded from the kitchen form you filled in, so it says something true the first time it opens. The planner shares between your phones like the rest of your settings, so both of you plan from the same numbers.

One number worth settling once. Your form answered **8** trays where you have also told me **12**. The tray count is your day's hard ceiling, so it decides everything above - the field says so, and it is yours to set.

19 Sep 2026 - engine v130 (no database step). The courier charge can now be changed from Edit as well, and deleting the courier's line on the Money screen takes the charge off the order with it.

What you told me. "q2, yes make deleting it clear the charge too", and "edit a courier related now only at one button, we should make edit an order able to alter details for courier."

Deleting the courier's line now clears the charge. A charge you paid has two halves: the **Delivery & fuel** row in your books, and the charge on the order. The Money screen only shows you the row, so deleting it used to take that half alone - the order still wore the charge, and the next Save in the order's own courier box quietly wrote the row straight back. Now the row and the order go together: the amount, who bore it and the COD mark all come off. The confirmation says so before you agree, naming the order, and the row is labelled **Courier (order #A3F9C2)** so you can see which order you are deleting it from. Deleting an ordinary expense still touches no order at all.

The courier charge is under Edit too. The amount, who paid the courier, how they settle it and how you paid are now asked in the Edit order pop-up as well as in the **Note / tracking** box - one shared set of controls, so the two forms cannot word it differently or save a different charge. Edit's own order total moves as you type the amount, so you can check the figure before you save rather than after.

Nothing was taken away. The Note / tracking box keeps all of it and is still the quick way in; Edit simply gained the same questions, so an order can be corrected end to end without opening a second box. Nothing changes on its own: a charge is only written when you save, and saving an unchanged form writes an unchanged charge.

19 Sep 2026 - engine v129 (no database step). Still to collect now counts the courier charge the customer pays with the order, so the row promises the same money their own message asks them for.

What you told me. "q1, it should reflex rm72." RM72 was RM64 of bread plus the RM8 the courier was charging, and the Money screen was showing RM64 - the items alone.

What it does now. When the customer bears the courier charge and pays it with the order, **Still to collect** counts what they will actually hand over: the items plus that charge. So the figure on your Money screen is the same RM72 their own message asks them for, and you are never left counting a charge you are about to be paid.

A COD charge stays out of it. When the courier collects the charge at the door, that money was never yours to collect - so it is not in this row either, exactly as it is not in the total you ask them for.

A charge you bore is not in it. Nothing changes there: the bread is what they owe you, and the postage is your own cost.

Your takings do not move. What has already come in - Cash in, TNG in, the net, every journal - still counts the items alone, because a charge the customer pays is pass-through: it arrives and leaves in the same breath,

so it was never profit. Only **Still to collect** changed, because that is the one figure about money still on its way to you.

19 Sep 2026 - engine v128 (one database step - run supabase/courier_cod.sql BEFORE you deploy). A courier charge the customer bears can now be a Courier COD charge: the courier collects it at the door instead of it sitting in the total you ask for.

What you told me. "courier charges can be collect, that means customer pay courier upon collect." A charge the customer bears was always being added to the total you ask them for. That is right when they pay it with the order - but when the courier collects it at the door, asking for it by TNG as well means the same money is asked for twice, once by you and once by the courier.

What it does now. The Note / tracking box asks how they settle it. Under **Who paid the courier**, answering **The customer paid it** now opens a third box, **How they pay it**, with two choices:

- **With their order (in the total)** - the charge sits inside the total you ask for, exactly as it always has. Nothing you have already recorded changes, because this is where that box starts. **COD - the courier collects it on delivery** - the charge comes out of the total you ask for, and is named everywhere as payable to the courier when the order reaches them.

Nobody is asked for the same money twice. The confirmation, the payment reminder and the shipped message all quote the items alone as the total, and name the charge underneath as COD with the instruction spelled out - "Courier charge: RM 8.00 - COD, pay the courier when your order reaches you". The customer's track card says the same. COD is the word the couriers themselves use for a parcel the receiver pays for; it is also called a reverse charge, but COD is what your customers will understand.

What you see. The box tells you what the customer owes as you type, and it now reads differently for COD - "the customer owes RM 30.00 - items total RM 30.00, plus RM 8.00 collected by the courier on delivery" - so what you are asking for and what the courier is asking for are never confused. The Edit pop-up's order total follows the same rule. And the order's row is tagged **Courier RM 8.00 · customer · COD**, so you can see at a glance which charges the courier is collecting rather than looking for that money in your tin.

Your books are untouched, either way. A charge the customer bears was already pass-through, and COD does not change that - it still never reaches your takings. Only a charge you bore yourself becomes a Delivery & fuel expense, exactly as before. (v129 above changes where that charge is counted under **Still to collect**.)

One database step. Run **supabase/courier_cod.sql** in the Supabase SQL editor BEFORE you deploy this build. The app publishes the whole tracking row in one call, so if that column does not exist yet the call is rejected as a whole and the customer's tracking page stops updating for EVERY order - not only the ones with a COD charge.

19 Sep 2026 - engine v127 (no database step this time). Deleting a courier charge really deletes it - the tag comes off the order and the charge leaves the customer's track card.

What you told me. "why i delete courier charges and the tag is not remove?" You were right, and it was a real fault, not a display quirk.

What was wrong. The charge is two things: the amount, and **who paid it**. Putting **Who paid the courier** back to **Not recorded** felt like deleting the charge - and it is the obvious way to do it, because that picker is what decides where the charge goes - but it only deleted the answer to *who*, and left the amount behind. The order's row kept a tag reading **Courier RM 8.00 · customer** for a charge nobody had assigned, and nothing you could do in that box would take it off. Worse, it named the customer as the payer when nobody had said so, so a row could claim they owed money they did not.

What it does now. A charge is the amount AND the payer. Put the payer back to **Not recorded** and both go: the tag comes off the row, the customer's total goes back to the bread alone, and the expense row goes out of your books. Clearing the amount still does the same, as before. And a charge with no payer recorded - one restored from an older backup, say - is tagged with nothing at all rather than being labelled as the customer's.

Clearing a charge now reaches the customer's track card too. It did not before: the card was only

republished when a charge was added, so a customer could keep looking at a courier line you had just deleted. Now the card is republished whenever the charge changes, in either direction.

Where to see it. Orders -> an order with a courier charge -> **Note / tracking**, then set **Who paid the courier** back to **Not recorded** and Save. The tag goes from the row.

One thing left as it is, so you know. Deleting the **Delivery & fuel** expense from your Money screen does *not* take the charge off the order - that row is your own record of money going out, and the order is the record of what happened, so the two are deliberately separate. If you would rather that deleting the expense cleared the order's charge as well, tell me.

19 Sep 2026 - engine v126 (no database step this time). The messages that ask your customer for money now show how the total is reached, and the app shows you the same arithmetic.

What you told me. You pasted the confirmation for order #7E51BC - 4 × Focaccia 9"x12" 800g, self collect, RM8 courier, Total RM 72.00 - and said: "the message need to show the add up for rm72. And it should be the same for APP."

What was wrong. The RM72 itself was right. What was missing was any sign of where it came from. The line above it named the bread without a price, so the RM64 of foccacia and the RM8 courier charge were both invisible, and the total read as a figure picked out of the air. That is the worst place to leave it, because this is the message that asks them for the money.

What it does now. When the customer bears a courier charge, the messages no longer state the answer on its own - they show the sum, in the same order every time, with the total underneath:

- Items: Focaccia 9"x12" 800g x4 Items total: RM 64.00 Courier charge: RM 8.00 Total: RM 72.00

The app says the same words, in the same order. One piece of code works out those two figures and their total, and the confirmation, the payment reminder, the shipped message and your own screens all quote it. Two surfaces cannot drift apart when there is only one sum.

On your side of the app. The **Note / tracking** box's "The customer owes..." line and the **Edit** box's order total both now read the breakdown as **items total ... + courier charge ...** - the same two words the customer sees, so the figure in front of you and the figure they were sent are provably the same arithmetic.

An order with no courier charge is untouched. Its messages come out byte for byte as they did before, and the tests hold that.

Where to see it. Orders -> an order with a charge the customer pays -> **Send confirmation, Remind for payment** or **Send shipped message**. Your side: that same order's **Note / tracking** and **Edit** boxes.

One thing left for you to decide - not changed. The Money screen's "still to collect" counts the items only: RM64 on this order, not the RM72 the customer will hand over. That is deliberate, because the line measures what stays in your purse and the RM8 arrives and leaves again in the same breath. If you would rather it showed the full amount you are about to be handed, tell me and I will change it.

19 Sep 2026 - engine v125 (no database step this time). The courier charge now appears in the two totals YOU read, and in the one customer message that was still missing it: the confirmation.

What you told me. Straight after deploying v124: "if customer were to paid courier, the total is not shown to me and to customer." You were right, and it was three places, not one.

What was wrong. v124 did put a charge the customer bears into the payment reminder, the shipped message and the track card, and named it in each. But the **confirmation** - the very first message, the one carrying the TNG QR that asks them for the money - still quoted the items alone. And on your side of the app, the **Note / tracking / courier** box never showed what the customer owed, and the **Edit** box's order total left the charge out. So both totals you read were lower than the figure the customer had actually been asked for. A customer paying what the confirmation said would have come up short, and you would have had no way to see why.

What it does now. The **confirmation** names the charge on its own line just above the total, exactly as the payment reminder already did - so the message that asks for the money and every message after it agree. The

shipped message, which named the charge but never said what the order came to, now ends on the total as well. In your app, the **Note / tracking / courier** box tells you what the customer owes as you type the amount, and it moves the moment you say the customer paid it - showing the arithmetic, items plus courier, so the figure explains itself. The **Edit** box's order total counts the charge the same way.

A charge you bear still appears in none of the customer's totals and in none of the figures they are shown - that part has not changed, and it is still deliberate. Your own books are untouched by a charge the customer pays, exactly as v124 had them.

Where to see it. Orders -> open an order -> **Note / tracking** (the total sits under the question of who paid). **Edit** -> the Order total line. The customer's side: all four messages, and their track card.

19 Sep 2026 - engine v124 (one database step: run supabase/courier_fee.sql in Supabase BEFORE you deploy this build). The Note / tracking box now records the courier charge itself, and who bore it - and the two answers behave differently on purpose: a charge the customer pays is added to their total and named where they can see it, and a charge you pay comes off your profit under Delivery & fuel.

What was wrong. You looked at the box behind an order's **Note / tracking** button and said the **Paid by** picker there was for courier charges. It was not. It records how the CUSTOMER paid YOU - Cash or TNG - and it is what puts an order's money into the right book on your Money screen. It has been there since v101, from your own words about needing to reconcile cash against the TNG app. Your reading was the more useful one though, because the thing you were looking for did not exist anywhere in the app: there was no field for what the courier charged, and none for who paid it.

What it does now. That box carries the **Courier charge** - what it cost to send that order - and, under it, **Who paid the courier**.

When the customer pays it, the charge is added to what they owe. It is named on its own line above the total in their payment reminder and their shipped message, and on their track card, so a total that runs a little more than the bread explains itself rather than looking wrong. Your books are not touched at all - and that is deliberate, not an oversight: in your purse that money arrives and leaves again in the same breath, so counting it as takings would only make a good day look better than it was.

When you pay it, a third box asks **How you paid the courier**, and the charge becomes an ordinary **Delivery & fuel** expense dated the day you paid it. It then shows up on your Money screen, in that account's own journal, and as a line on your profit statement - so it comes off your profit like any other cost, which is exactly what you asked for: off profit only if you paid it. Recording the same order twice updates that one expense rather than adding a second, and clearing the amount takes it away again.

Either way, the order's row is tagged with the charge and whose money it was - amber when it was yours, quiet grey when it was theirs, sitting beside the Cash / TNG tag.

The box you misread is renamed. The picker that was called **Paid by** now reads **Paid by the customer**, and it sits below the courier charge rather than above it, so the two questions cannot be confused again: one is what THEY paid YOU, the other is what YOU paid the courier.

Where to see it. Orders -> open an order -> **Note / tracking**. The customer's side: the payment reminder, the shipped message, and their track card.

Before you deploy. Run **supabase/courier_fee.sql** once in Supabase (Dashboard -> SQL -> New query -> Run). The app publishes each order's tracking row in a single call, so if that column is missing the call is refused whole and the customer's track page stops updating for every order - not only the ones with a charge. Run the SQL first, then deploy.

Nothing else moved - your prices, customers, delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v123 (no database setup needed). The product dropdown's middle section now groups sold-out and not-sold-that-day together, and a product the shop does not sell on a day is named by the days it IS sold instead of being given a count for a day it was never on.

What was wrong. The dropdown's middle section was called **Sold out** and read its count for the day you were adding to. But a product the shop does not sell on that day at all - a Saturday-only loaf on a Wednesday - was sitting in **On the shop**, and worse, could be labelled "sold out" from a count for a day it was never on the menu. Neither told you anything true.

What it does now. The middle section is **Unavailable**, and it holds both kinds together: a product that has **sold out** for the day you are adding to, and a product the shop **does not sell on that day**. The second kind is named by the days it IS sold - a Saturday-only loaf on a Wednesday reads "(only Sat & Sun)". The headings are now **On the shop, Unavailable, Taken down**.

Why they are one section, not two. They are the same thing to you: an active product you can still choose. You bake to a plan of your own, and as you put it - every day the kitchen may make 12 focaccia, ordering closes the day before at 6pm, and the 12 are not necessarily sold off, so some go to a walk-in or into the fridge. **The picker is a guide while you sell, never a gate:** nothing is blocked, and the order-by deadline you set per product is deliberately not counted, because that deadline stops a stranger ordering - it says nothing about what you may sell by hand.

Where to see it. Orders -> **+ New order** -> Items, and Orders -> open an order -> **Edit**.

Nothing else moved - your prices, customers, delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v122 (no database setup needed). The product dropdown now shows the three kinds inside the list itself: each kind sits under its own heading, and on a newer phone each section is tinted in its colour.

What changed. Adding items to an order - in **+ New order** and in **Edit** - the product list now comes in three labelled sections instead of one long run of names: **On the shop** first, then **Sold out**, then **Taken down**. The headings appear on every phone, because they are part of the list itself.

The colour. On a phone with the very latest system - a new iPhone or a recent Android - each section is also tinted inside the open list: **green** for on the shop, **amber** for sold out, **grey** for taken down. The menu then reads as three blocks you can tell apart without reading a word.

On an older phone. The open list is drawn by your phone itself, and an older phone will not let a page colour it. You get the same three headings in the same order, and the closed box still wears the colour of whatever it holds - you lose the section tint and nothing else. Nothing is broken, and nothing has moved; only the colour is missing.

This replaces the v121 note below. v121 put the colour on the closed box alone, because the open list could not be reached. Newer phones can be reached, so the list carries the colour too.

Where to see it. Orders -> **+ New order** -> Items, and Orders -> open an order -> **Edit**. The list is otherwise unchanged: the same products, the same "(hidden)" marker, the same "12 left" count.

Nothing else moved - your prices, customers, delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v121 (no database setup needed). The product dropdown now shows you which kind of product each choice is, in colour, and lists them in the order you would look for them.

What changed. Adding items to an order - in **+ New order** and in **Edit** - the product list is now sorted into the three kinds you already know, in the order you would reach for them: what is **On the shop** first, then anything that has **sold out** for the day you are adding to, then whatever you have **taken down** (still marked "(hidden)"). A product still in Draft stays out of these lists entirely, as it always has.

The colour. The closed box wears the colour of whatever is picked - **green** for on the shop, **amber** for sold out, **grey** for taken down. On an order with several lines that means one glance down the list tells you which lines need a second look, without reading each one.

Why only the closed box. The open list is drawn by your phone itself, and iOS does not let a page colour the

rows inside it. So the colour sits on the box you see while the list is shut, which is where you read it anyway.

Where to see it. Orders -> **+ New order** -> Items, and Orders -> open an order -> **Edit**. The list itself is unchanged: the same products, the same "(hidden)" marker, the same "12 left" count.

Nothing else moved - your prices, customers, delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v120 (no database setup needed). One customer is one row. The app now recognises a phone number by its digits, so the same person can no longer appear twice just because their number was written with a + in front of it - and the duplicate already in your app joins itself into one the first time you open it after updating.

What was wrong. A customer used to be identified by their number written in exactly the same way every time. So **+60123456789** and **60123456789** looked like two different people, and the Customers list showed the same customer twice. Your messages always went to the right person - only the list disagreed. Numbers are now read by their digits, which means **012-345 6789**, **+60 12-345 6789** and **60123456789** are one person, the same way "012345" already found "012-345" in the finder.

It tidies itself once. The first time you open the app after updating, it joins any customer who was split this way, carrying their dog's name and photo, what they like and avoid, and their note across to the one that remains. There is nothing to do - you will simply find one row where there were two. If you are the careful sort, take a cloud backup first from **More -> Backups**: that is the way back if anything looks wrong, and it costs one tap.

New: "Join with another customer". Open a customer's history and you will find the button under their number. Tap it, pick the duplicate from the list, and confirm. Reach for it when two rows are split for some other reason - a misspelt name with no number beside it, say - which the tidy above cannot guess at. It moves the other person's orders under the name and number of the customer you were looking at, so open the row you want to KEEP. This one cannot be undone inside the app; your cloud backup is the way back, and the confirm box says so before it acts.

Two things worth knowing. Update both phones - a phone still on v119 can save a number the old way from its Profile card, and only the updated build does the tidy. And on the Profile card itself, the number you type is now stored in the one form every label, message and link already uses, so a "+" or a space can never split a customer again.

Nothing else changed - your prices, your delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v119 (no database setup needed). Start typing a customer's name in an order form and the people you have already served appear underneath - tap one and their name and WhatsApp number fill themselves in.

How it works. Type two letters - "aun" is plenty - and a short list appears right under the name box. Each row shows their number, how many orders they have placed, and the item they buy most, like **012-345 6789 · 5 orders · usually Focaccia**. Tap the right person and the name and the number both go in; the delivery day, the items and self-collect-or-courier stay yours to set. One letter on its own is not a search - it would offer you half your address book - and a name you have never served offers nothing at all. Typing a number without its dashes finds them too, so "012345" works the way "012-345" does.

Both order forms do it - the **+ New order** card and the **Edit** pop-up, the same way.

Two things worth knowing. It draws on the orders already in the app, so on a brand-new setup it has nothing to offer until your first order is saved. And on the **Edit** form, saving a changed name has always carried that name to every other order belonging to the person it was - that is the one-saved-name rule from v68 and it has not changed, but picking a name is now much easier than typing one, so it is easier to do by accident. If you ever pick the wrong person while editing an order and save it, tell me and I will add a confirmation that asks first.

Nothing else changed - the Customers screen, its finder, your prices and the shop are all exactly as they were.

18 Sep 2026 - engine v118 (one SQL step in Supabase; nothing to re-install). The order alert on your

phone now arrives as a real notification instead of a wall of raw text. The ping that reaches you when a customer orders was showing the whole message as unreadable computer text - curly brackets, quote marks, the word "title" - instead of a notification. The cause was the way the message was handed to the ntfy service: the channel name has to travel inside the message for ntfy to lay it out as a title and a body. It was missing, so ntfy did the only thing left and printed everything as plain text.

What you see now. A notification titled **New order - Jien Luv 2 Bake**, with the delivery day written out in full (Sat 19 Sep 2026, not 2026-09-19), the total as money (RM 10.00), a bread icon, and a higher priority, so it is more likely to make a sound and pop up rather than wait quietly. A courier order also carries the address; a self-collect order no longer shows a delivery address it does not have.

A tidy-up came with it. Your channel name used to be typed inside the rule that sends the pings, and three leftover test rules from setting this up carried copies of it too. It now lives in one small table, and the three leftovers have been removed - so the name that guards your alerts exists in exactly two places: your database, and the ntfy app on your phones.

Nothing about ordering changed, and no phone needs re-subscribing.

17 Sep 2026 - engine v118 (no database setup needed). The Paid step now stays in its place on every order, wearing an X until the money is in - and turns into a green tick when it is. This replaces what v117 did, on your instruction.

Your reasoning, and you are right. The customer's track page is the design: they must see the same steps in the same places every time, and your own app has to read the same way. v117 left the Paid step **off** an order that skipped it - five dots instead of six - which meant the two lines did not line up, and a missing step is itself a thing to puzzle over.

What happens now. A regular who pays when she collects still never passes through Paid. But that step keeps its place in the line, and it wears an **X** - amber, deliberate - instead of a tick, and it never turns green while the money is outstanding:

> New tick · Confirmed tick · **Paid X** · Baked tick · Packed tick · Collected / Shipped dot

Press **Paid · Cash** or **Paid · TNG** when she hands the money over - at whatever stage the order has reached - and **the X becomes the tick**: the ordinary green tick, in the same spot, with the **Cash** or **TNG** tag beside the row. The order itself stays where it was: taking the money never drags it backwards.

Both lines change together. Your row's map and the **customer's own track page** draw the same six steps, so the X is on their page too, and it turns green there the moment you record the money. A step that is *deliberately gone past but not paid* now looks different from a step that has not been reached (grey) and from a step that is done (green) - which is exactly the distinction that was missing.

17 Sep 2026 - engine v117 (no database setup needed). A regular who pays when she collects no longer has a Paid step on her order - and the Paid · Cash / Paid · TNG buttons stay on until the money is actually recorded, wherever the order has got to. From your note about close customers paying by TNG or cash at pickup.

What was wrong. Move an order from Confirmed straight to Baked - no payment - and the row's map still showed **Paid with a green tick**. It was claiming money you had not taken. The app assumed "past Paid means paid", which is fine for her old orders but false for a regular who settles up at the counter.

What happens now.

- **A bypassed order has no Paid step.** Her route reads New -> Confirmed -> Baked -> Packed -> Collected / Shipped: **five dots instead of six**, and no tick for a payment that never happened. **The Paid · Cash / Paid · TNG buttons stay on** at every stage from Paid onwards - Paid, Baked, Packed, and Collected / Shipped - until the money is recorded. So the moment she hands over the cash at the counter, wherever the order has got to, you press the button right there: **the order stays where it is** (marking an order paid never drags it backwards) and the money is stamped with the day it landed, which is the day the Money screen counts it. **And the Paid step comes back once it is paid** - a green tick, in its proper place in the line, with the **Cash** or **TNG** tag beside

the row. So the map ends up telling the truth about that order either way. The buttons disappear the moment the payment is recorded, and they never appear before the order reaches the Paid stage. **The customer's own track page follows the same rule** - the same five dots, no Paid tick for a payment that has not happened, and the step appears there too once it is recorded.

One rule behind it, worth knowing: moving an order into Paid **or any later stage** without having pressed a Paid button is taken as *this order owes money*. That is what your older orders are protected from - nothing already in the app changes, because only orders you move now are marked.

17 Sep 2026 - engine v116 (no database setup needed). A Day one form: your opening balance - the cash in the tin, the money on your phone and what is on your shelf - entered once, in one place. From you telling me *"we need to enter opening balance"*.

What an opening balance is here. Four answers about the morning your books begin: what is in the tin, what is on your phone, what is on your shelf, and who still owes you. Until now the app could take all four, but in three different places and one ingredient at a time, which is a poor way to start.

The form. More -> Money -> the new **Day one** line (under **Books**) -> **Set**. It asks for:

- **Cash in your tin** and **Money on your phone (TNG)** - two boxes, and the day your books begin, which takes any date but should be the day you are starting from. **What is on your shelf** - your ingredients, one box each, already in the unit you use for that ingredient (kg for the flour, g for the salt - no converting). Ingredients you have marked as *not something I buy*, like labour, are left out: they have no shelf to count. **Who still owes you** needs no box - those orders stay unpaid and show under **Still to collect**.

Nothing is half-saved. Every box is checked before anything is written: if one has something that is not a number, nothing saves at all until you fix it. And **a box you leave empty is left exactly as it is** - so the form is safe to reopen later just to correct one figure, and safe to run again without wiping what you already had.

What it writes. The tin and the phone become ordinary money-in rows dated that day (so from then on the Money screen's **Net** is what you should really hold, and the weekly check works from the first week), and the shelf becomes each ingredient's **On hand** - the figure your shopping lists subtract, so your first list buys only what you are genuinely short of.

What it deliberately does not do. Money you **owe** - a loan for the oven, a supplier you have not paid - has nowhere to go: the app keeps a loan as a way of paying, not as a debt, so record what it pays for as it happens. Your own money in shows under **Capital you put in** on Profit, never as income. And the tin must not be entered as a sale: a sale has a customer behind it.

Section 3 of the **Money and Profit** manual has been rewritten around the form - it is attached.

17 Sep 2026 - engine v115 (no database setup needed). The month arrows on Profit work both ways again. From you telling me *"the profit month can move earlier but cannot move later"*.

What was wrong. Open Profit - it starts on this month, where the **right** arrow is correctly switched off, because there are no numbers after today. Press **left** to step back to August, and **right** stayed switched off too, so you could not get back to September. You could walk backwards through the months but never forwards again: to return you had to leave the screen and open it again. The arrow's state was worked out once, when you opened the screen, and then reused on every step - so it kept answering for the month you started on rather than the month you were looking at.

Fixed. The arrows are worked out fresh on every step, so **left** and **right** now move both ways between the months you have, and **right** switches off again only when you are back on this month. Verified by stepping back and forward repeatedly. The same fault did not exist on the other calendars - Orders, Delivery dates and the date pickers all recompute their arrows each time they draw.

17 Sep 2026 - engine v114 (no database setup needed). Every spending line in Profit & Loss now opens - including the ones reading 0.00 - and the lines are twice as tall, so they are easy to tap. From you asking *"in profit the expenses is not clickable, is that a bug?"*

You were right, and it was my mistake. The lines that open their journal were the ones with money in them. A

line showing **0.00** was deliberately left dead, because there is no journal to show - but I made it look **exactly** the same as a live one: same colour, same font, no hint. So if your month has most categories at 0.00, nearly every line you tapped did nothing, and "not clickable" was the only sensible conclusion. A line that looks alive and does nothing is worse than no line.

What it does now. Every spending line opens: *Packaging, Utilities, Delivery & fuel*, the categories you have added, and **Total expenses**. A line with money in it shows its journal as before. An empty one opens and **says so in your own words** - *"Nothing recorded under Rent in September 2026"*, In / Out / Net at zero, and a line telling you it will fill up on its own as you record spending under that category. The statement's own totals (**Sales, Cost of sales, Gross profit, Net profit**) stay figures rather than doors - you did not ask for those, and they are not spending.

Two more things found while fixing it.

- **The lines were only 17 pixels tall** - a small target for a finger, and easy to land in the gap between two lines and hit nothing. Tappable rows everywhere (the Profit statement, the Money screen's cash rows, the Books list) are now **36 pixels**, comfortably thumb-sized. **The Total expenses journal did not say which category a row belonged to.** It read *"2 Sep · boxes · Cash"* with no way to tell what "boxes" was for. It now names it - *"2 Sep · Packaging - boxes · Cash"* - while a single category's journal still reads short, since its title already says which category it is.

17 Sep 2026 - engine v113 (no database setup needed). Every way you pay now has a book you can always open - including a pocket that has been quiet - behind a new Books line on the Money screen. From you asking *"where can i find pocket journals"*.

What was wrong. The journal rows live on the money card, and that card only shows a pocket's row **when that pocket moved money in the stretch you are looking at.** The card opens on **Today**, so a pocket that paid for something last week had no row today - and no way into its book at all. Cash and TNG never had the problem: their four rows (Cash in, Cash out, TNG in, TNG out) are always there. The pockets were the only books that could go missing.

The Books line. Under the two money cards there is now a second line beside *Categories & ways to pay*:

> **Books** > Every way you pay · 5 books, each opening into its own rows [Open]

Open lists **every** way you pay - Cash, TNG, Loan, Personal Pocket Kean, Personal Pocket Suan, and anything you add later - each with what moved by it in the stretch on screen, each opening into its own book **under the line you tapped**, so the list stays in front of you and you can step from one book to the next. The first method with anything in it opens ready; a quiet pocket says so plainly - *"Nothing moved this way in this stretch"* - with its In / Out / Net at zero.

Two things this keeps honest: a method you have since **renamed or deleted** still gets a line, because its money is still in the books and a figure you cannot open is a figure you have to take on faith; and every figure in the list comes off the Money screen's own rows, so the list and those totals cannot drift apart.

A small thing the screenshot caught while building this: the first version printed the word "null" under each closed line. Fixed, and there is now a test that fails if it ever comes back.

17 Sep 2026 - engine v112 (no database setup needed). You can pay a personal pocket back out of the till in one go, and the category list on the expense form is no longer cut off. From your question - *"can my own withdrawal payback to the cash register like Personal pocket Kean or Suan?"* - and your note about the categories.

Pay back a pocket. On **More -> Money**, beside **+ Put money in**, there is now **+ Pay back a pocket**. It is for the case where a pocket of yours - **Personal Pocket Kean, Personal Pocket Suan**, a loan, or one you add later - paid for something, so the money left you rather than the till; that pocket's line on the Money screen then reads **-RM 40.00**, which means *the till owes it 40*. Open the form and it comes up **already on the pocket that is owed**, with the amount filled in - Kean's RM 40 sitting in the box, and a line saying so. Say whether the money came out of **Cash** or **TNG**, check the day, and press **Pay back**.

It writes both halves at once, which is the whole point: the **till goes down** by that amount (it shows in Cash out and in the Net, and in the Cash journal with your note), and the **pocket's line comes back to zero**. One pocket can be paid back without touching another. Because it is your own money going back to you, the till's side is recorded as a withdrawal - so it never counts as a cost and your **profit does not move**. On the pocket's own list it reads *"Paid back by the till"*, not *"From my pocket"* - the two are different things and now say so. Paying it back in part is fine: pay RM 25 of the RM 40 and the pocket reads -15 after.

The category list was being cut off - fixed. Thank you for catching this. On the expense form, the row of categories ran off the side of the screen: it was laid out as one single line, so at phone width everything past about the sixth category - *Salary (you), EPF / SOCSO, Marketing, Equipment & tools, Other, My own withdrawal*, and the **+ New category** chip - was off the screen and impossible to reach. Measured on a 375-pixel phone, the row was 744 pixels wide inside a 343-pixel box. The pills now **wrap onto as many lines as they need**, the same fix applies to the ways-to-pay row and to the new pay-back form, and the whole list is visible and tappable. Nothing about the categories themselves changed.

17 Sep 2026 - engine v111 (no database setup needed). Every figure on the books is now a door: a journal behind each line of spending in Profit & Loss, and a book for every way you pay - **Cash, TNG, Loan, Personal Pocket Kean, Personal Pocket Suan, and anything you add later**. From two notes close together: *"the expenses items in Profit & Loss should reveal its journals"*, then *"each CASH, TNG, LOAN, Personal Pocket Kean, Personal Pocket Suan, and others that might be added in future need a journal"*.

Profit & Loss - tap a spending line and see what made it up. On **More -> Profit**, the running-cost lines (Packaging, Utilities, Delivery & fuel...) are now tappable, and so is **Total expenses**. Tapping one opens its journal for the month on screen: every expense behind that figure, oldest first, each reading *"<day> · <what it was for> · <how it was paid>"* and ending on the total the statement itself shows. So *"Packaging RM -58.00"* opens two lines - the boxes paid from a pocket and the bags paid in cash. A line with nothing in it that month is not tappable; there is nothing behind it. The header now says **Running costs · tap a line to see the spending behind it**.

Money - one line and one book per way of paying. The single *"Paid by loan / other"* row is gone. In its place, every way of paying that is not cash or TNG gets **its own line** with what moved by it, and its own journal behind it: **Paid by Loan, Paid by Personal Pocket Kean, Paid by Personal Pocket Suan, Paid by Bank OD** - whatever you have called them. Money of your own you put in that way reads as a plus on that line; what was paid out of it reads as a minus. Cash in, TNG in, Cash out, TNG out and the Net are exactly as they were, and none of these other ways is in the net - it never came out of your purse.

Why they appear by themselves. The lines are read off your own transactions, not off a fixed list, so a method you add tomorrow needs nothing from me: record one payment with it and it has a line and a book. And a method you have since **renamed** keeps a line for the name its older rows were written with - those rows are still in the books, and a figure you cannot open is a figure you have to take on faith. Each book ends on **In / Out / Net** for that method.

One small correction while here: the note at the foot of the Profit screen about your own hours now matches v110 - either pay yourself a **Salary (you)** expense (with EPF / SOCSO as their own category), or mark Labour as a not-bought ingredient and put the hours into the recipes. Count them one way, never both.

17 Sep 2026 - engine v110 (no database setup needed). An ingredient you never buy - **your own labour, electricity, gas** - can be marked as a cost, and it then never appears on a shopping list again. From your note: *"certain ingredient we dont purchase, in ingredient we can set that as a non purchase item, like labour and electricity"*.

What it does. Open **More -> Ingredients** and tap **Edit** on the ingredient (or fill in the New ingredient card at the top). Above the cost there is now a switch: **Not something I buy - it only ever costs. It still counts in a recipe's cost, but it never appears on a shopping list or a purchase order**. Turn it on for Labour, Electricity, Gas, your own time.

Where you see it.

- The ingredient's **card** stops showing an On hand and a Keep-at-least line, and says instead: *Not bought - a cost in your recipes, never on a shopping list*. Nothing else about the card changes, and the cost you typed is untouched. Every **shopping list** (the preview, the list you save, and an "orders changed" follow-up) leaves that ingredient out - and says so at the foot of the list: *Not on this list: Labour, Electricity - marked as not something you buy. Their cost still counts in the products that use them*. So it reads as deliberate, not as something the app forgot. The **recipe cost** is unchanged: a recipe using Labour at RM 8 an hour for 15 minutes still adds RM 2.00 to that product's cost, and that runs on into the Profit screen's cost of sales.

Switching it OFF restores the ingredient exactly as it was - the app removes the mark rather than storing a "No", so nothing that exists today changes unless you turn this on.

And a word on the books, since this touches Profit: if you mark Labour as not-for-purchase because it is your own unpaid time, it is already inside each product's cost of sales through the recipe - so do **not** also record it as a Salary expense, or the same hours get counted twice. The **Salary (you)** category is for when you actually pay yourself money out of the till.

16 Sep 2026 - engine v109 (no database setup needed). A journal for each way the money moves - the cash book, the TNG book - and a bug that was hiding in the TNG column. Two things, both from you asking how to see the journals.

A bug first, because it matters. Every order you marked **Paid · TNG** since v106 was landing in "Paid, no method" instead of the TNG column. The app was checking for the old way of writing it ("tng") while the paid buttons write the list's own label ("TNG") - so the TNG in figure read RM 0.00 with the money in your phone, and the same money showed as unaccounted for. Fixed, and there is now a test that would have caught it. If your TNG column has looked wrong since v106, that was it.

Tap a figure, see its journal. On the Money screen, **Cash in / TNG in / Cash out / TNG out** (and the loan line) can be tapped. Each opens that method's book for the stretch: every order paid that way with its code and customer, everything you spent out of it with its note, anything of your own you put in - in date order, ending on In / Out / **Net**, and a line saying what that is meant to be (your purse, your phone, or money that never went near either).

16 Sep 2026 - engine v108 (no database setup needed). Where the two lists live, and how to change one. One fix and one addition, both about the categories and the ways to pay.

Why. v106 put them in a Settings card and you could not find it, and there was no way to change an entry once it existed - "dont put the setting separately, it should be at where it suppose to be."

They are on the Money screen now, as a line under the money lists reading "Categories & ways to pay · 11 categories · Cash, TNG, Loan" with an **Edit** beside it. Tap Edit and the lists open: every category with its kind (ingredients / running cost / your own money) and every way to pay, each line saying "edit".

Tap a line and it turns into the form: rename it, change what kind it is, or delete it - in place, with the rest of the list still in front of you. Renaming moves what you have already recorded with it, so a slip of the thumb does not split your history into "Packaging" and "Packing".

And the + chips sit where you are working: the Add an expense form and Put money in each end their pills with **+ New category / + New way to pay**, which opens the same small form in place - the amount and the note you had already typed stay put.

No window inside a window. These forms expand in place instead of opening a pop-up of their own, the same rule the date field follows: the app has one pop-up layer, so a second one would wipe out the form underneath it.

16 Sep 2026 - engine v107 (no database setup needed). The item line on an order reads properly again. One fix, on the two forms where you type an order's items.

What was wrong. When the selling price joined the line (v101) it took most of the row. Measured on a 375px phone: the price box was 183 pixels wide and the product dropdown was squeezed to **2** - so the name of what you had picked was pushed onto a line of its own and unreadable, and the whole line looked broken.

What it does now. Each item line is two clean lines: **the product across the top**, full width so its name reads,

and its controls under it - how many, the selling price on the right, and the X to remove. The same shape in the + New order card and in the Edit pop-up, since both are the same row.

16 Sep 2026 - engine v106 (no database setup needed). A note on every transaction, and two lists you shape yourself. Three small things, all on the money screens.

A note on each one. Add an expense now takes a note ("Mydin run, 2 boxes") - the Put money in form already had one - and both lists on the Money screen show it beside the amount, so a row reads as what it was, not just a figure.

The categories are yours. More -> Settings gains a **Categories & ways to pay** card. Every category shows what it means to the accounts (ingredients / running cost / your own money), each has a X to delete it, and you can add your own - "Baking class", "SSM licence" - choosing which of the three it is. Deleting one never touches what you have already recorded: those rows keep their label and still count, printed at the end of the statement's cost list rather than dropped.

And the ways to pay are yours too. The list is Cash, TNG and **Loan** - the third choice you asked for - and you can add your own, a **Bank OD** or a cheque, so a shopping run nobody would pay for out of the till can still be recorded honestly. The same list appears on the Bought prompt (where the receipt is still in your hand) and on both money forms.

One thing that follows from it. Money that never came out of your purse - a loan, an overdraft - no longer sits in the net. The Money screen leaves it out of the cash figures and names it on its own line, "Paid by loan / other", with a sentence saying why. Net now means exactly what it says: what should be in your purse and on your phone.

16 Sep 2026 - engine v105 (no database setup needed). The books: a profit and loss account, a month at a time. One new screen, More -> Profit.

Why. Money told you what was in your purse. It could not tell you whether the baking made money - a RM250 flour-and-butter run makes one week look like a disaster while the bread it went into sells the week after.

What it shows. More -> Profit, month by month (arrows either side), reading down like a set of accounts: Sales, Cost of sales, Gross profit, the running costs one line each, Total expenses and Net profit - with the gross margin beside the order count. Then, kept apart from all of it, your own money: capital you put in, drawings you took out, and what is left in the business.

The one rule that makes it accounting rather than a cash total. Ingredient cost is what the baking cost to make, taken from your RECIPES as the bread sells - not the packs you bought. Buying and using sit in different places on purpose: a shopping run is money out on the Money screen and stock on your shelf, and it becomes a cost of sales here as the bread made from it goes out. Nothing is counted twice, and a big stock-up week stops looking like a loss.

Your own money is neither income nor a cost. What you put in is capital; what you take out is drawings (including the "My own withdrawal" category). Both move cash - Money counts them - and neither changes profit. If you pay yourself a proper salary, record it as **Salary (you)**, with **EPF / SOCSO** as its own category, and those DO count as running costs.

The category list is now a chart of accounts: Ingredients & shopping (stock), Packaging, Rent, Utilities, Delivery & fuel, Salary (you), EPF / SOCSO, Marketing, Equipment & tools, My own withdrawal, Other - the same list the Add an expense form offers, so every row lands on the right line of the statement.

Two things it will not pretend to know. Your own unpaid hours are not a cost - if your time is worth paying, pay yourself a salary and it becomes one. And a product whose recipe prices to nothing is counted as costing nothing; the screen says how many lines that was, so it is visible rather than flattering.

16 Sep 2026 - engine v104 (no database setup needed). Money you put in yourself, and taking it back out. One change, on More -> Money.

Why. Some spending happens before any order money arrives - a packet of flour paid from your own purse, a float of change for the day - and there was nowhere to put it. The Money screen only counted what customers

paid.

Put money in. More -> Money now has a **Put money in** button beside Add an expense: how much, the day you put it in, cash or TNG, and a short line saying what it was for. It counts into the Cash in / TNG in rows, because your own money really is in your purse and those rows are what you check the purse against - and one line underneath says how much of the money in was yours: "of the money in, RM 100.00 was your own". The entry is listed in its own **Money in** card, so you can see where it came from, and take it off again.

Taking it out again. Add an expense has a new category, **My own withdrawal** - money you take back for yourself rather than a cost of baking. It leaves through the same money-out list as everything else, so the net, the day lines and your backups count it with no new machinery. Your own money in and out is shared between your phones, like every other list.

16 Sep 2026 - engine v103 (no database setup needed). Money out: what you spend, beside what you take in. One change, on the shopping side of the app.

Why. v101 gave you the money coming in. Spending had nowhere to go: a purchase order worked out what a shopping run should cost, and tapping Bought put the packs on your shelf, but nothing recorded what you actually paid.

A shopping run now asks. Tap **Bought** on a saved list and, once the packs are on your stock, a small box opens: **What did you pay?** - pre-filled with the list's own total, with **Cash** or **TNG** beside it. Accept the guess with one tap, or type what the receipt really said. **Skip the money** leaves the stock added and records nothing, exactly as the app behaved before.

Everything else goes in by hand. More -> Money now carries an **Add an expense** form for what a purchase order never sees - packaging, delivery and fuel, utilities, equipment, or a market top-up - each with the day you paid it, so a receipt found in your bag later still counts on the right day.

The Money screen adds it up both ways. Cash in, TNG in, Cash out, TNG out and **Net** - what should be in your purse and on your phone for Today / This week / This month - with that stretch's spending listed underneath, each line showing what it was for and how you paid. What is still to collect stays out of the net: it is money owed, not money held. Expenses are shared between your phones, and counted in your backups, like every other list.

One fix along the way. A payment taken just after midnight was being counted as the day before, because the app read the date off the stamp in UTC. It counts on your own day now.

16 Sep 2026 - engine v102 (no database setup needed). Swipe across the calendar to pick a run of delivery dates. One small change, on More -> Delivery dates.

You asked whether the delivery-date calendar could work the way a product's availability calendar does. It now does: **drag across days and the whole run fills in under your finger**, then **Add selected** puts them on the calendar. It is the same gesture as dragging a sell period on a product. A single tap still picks one day (or puts it back), and a day that is already a delivery date is left alone by a drag - it comes off by tapping it, the way it went on.

That makes four ways in: tap dates, swipe a run, tap a weekday letter for every one of that day in the month, or **Generate the next dates** on your Settings pattern.

What I deliberately did NOT copy: the From / To boxes. On a product's availability card those two boxes are how a SEASON is said - a sell period genuinely has two ends. A delivery date is one day, not a stretch, so a From/To here would only ever have been a slower way to do what swiping now does.

16 Sep 2026 - engine v101 (no database setup needed). A price you can change on an order, and the money recorded as cash or TNG. Three things, all about money.

A price you can change on an order. Every item line on the + New order form and in the Edit pop-up now carries its selling price. It opens on what the product costs; type over it and THAT order is sold at your price - the confirmation, the payment and pickup reminders, the receipt, the customer's own track page and every money number in the app follow it. Your menu price is untouched, and a menu price you change next week

never rewrites a sale already made. Blank keeps "whatever the product costs", which is how an unpriced product has always behaved. Both forms also show the items total as you go.

Cash or TNG, written down when the money lands. The Paid button is now **Paid · Cash** and **Paid · TNG** - one tap each. The order remembers which, and when, and shows it beside its status. The Note / tracking box (now **Note / tracking / payment**) carries a **Paid by** box as well, for an order you marked paid before you could tell, or one to correct later. None of this reaches the customer: their page just shows Paid.

Two places to check it against. The delivery day's own header on Orders now carries a short till - "Cash RM 95 · TNG RM 120 · 1 to collect". And More -> **Money** is a screen of its own, with **Today / This week / This month**: cash collected, TNG collected, paid but no method recorded, and what is still to collect with the number of orders. Collected money is counted by the day it LANDED, so a transfer that comes in today counts today even when the order delivers on Friday; what is still to collect is counted by delivery day, because that is the day you hand it over. An order paid before any of this existed shows under "Paid, no method" rather than being guessed at.

16 Sep 2026 - engine v100 (no database setup needed). Making delivery dates for a weekday is one tap again. One change, on More -> Delivery dates.

Why. You found the function missing: adding your delivery dates for, say, every Monday meant tapping each Monday by hand. Nothing had been removed - the app's only "do the dates for me" button lived on Home, inside the "No delivery dates yet" message. That message shows only while you have no upcoming dates, and pressing the button creates the very dates that hide it: it was a one-shot button that took itself off the page the moment it worked. So once you had dates, which is every day since your first week, it was never there again.

Three ways in now. On More -> Delivery dates: tap dates and **Add selected**, as before; or **tap a weekday letter** - tap **M** and every Monday in the month shown is picked for you, tap it again to put them back, and the letter underlines solid once that whole weekday is picked; or press **Generate the next dates**, which is always there, follows the delivery days in your Settings, and adds the next six dates that are not on the calendar yet. Generate moves the calendar to the month the new dates start in, so you can see what it just did.

One wrong line fixed with it. Home's "no dates yet" message used to say "the next Mon/Wed/Fri dates" whatever your delivery days really were. It names your own days now (and so does the line under Generate).

15 Sep 2026 - engine v99 (no database setup needed). The customer can tap the "Next available" line and be given that day. One change, on the order page.

Why. A product you keep on the shop, on a day it cannot be ordered, tells the customer when it can be had - "Next available: Sat 19 Sep". Until now they had to go and find that date on the calendar themselves.

What it does now. The line is a control: tapping it takes that delivery day, so the product in front of them can be ordered straight away, and the page moves up to the calendar, where the chosen day is written out in words. It keeps its gentle pulse, and wears a small arrow so it reads as something to tap.

The one time it does not take the day. If the customer already has something in their basket, the line stays a plain label and tapping it answers instead: "Your basket is for Fri, 18 Sep. To order for another day, choose it on the calendar above." Taking a day then would silently move their whole order to a new date - and drop anything that does not fit there - which is not something a tap on a product line should ever do. Changing the day is what the calendar is for.

15 Sep 2026 - engine v98 (no database setup needed). The last status reads Collected / Shipped, and a placed order gains a Note / tracking button. Two small changes on the Orders screen.

One label for the last stage. v97 named it per order - Collected on one the customer fetches, Shipped on one you post. You asked for the pair itself instead, so the stage now reads **Collected / Shipped** everywhere it is named: the row's status list, the row's own journey map, the day's status filter, and the map on the customer's track page (in all three languages). Which message the row offers is still decided by the delivery method - a courier order shows **Send shipped message**, a self-collect one **Send pickup reminder**.

A way in for the two fields you reach for most. Beside Edit, every order now carries a **Note / tracking** button.

It opens a small box with exactly two fields - the order's **Note** and the courier's **Courier tracking number** - and a Save. Send, and the note lands on the order, while the tracking number also goes onto the customer's track card and into the shipped message. Nothing else moved: **Edit** still opens the full form (the delivery day, the customer, the address, the items, the quantities) whenever you need it. The tracking box that sat on the row itself in v97 is gone - the button is the quicker, tidier way to the same field, and it works on any order rather than only a courier's.

15 Sep 2026 - engine v97 (one database line to run once). The last status is named for how the order leaves, and a posted order can carry its courier's tracking number. One change, across the order row, the WhatsApp messages and the customer's track page.

The last status is Collected or Shipped. The final step used to say "Delivered" for everything. It now reads **Collected** on an order the customer fetches from you and **Shipped** on one you post - the same step in the journey, named for what actually happened. That is true everywhere the status is named: the row's status list, the row's own journey map, the day's status filter, and the map the customer sees on their track page (in all three languages). Nothing about the step's behaviour changed, and every order already sitting at that step keeps its place - only the word is different.

A tracking number, typed where you post the parcel. A courier order shows a **Courier tracking number** box under its details from Packed onwards, and a **Send shipped message** button beside the status. Type the number the courier gave you, press the button, and WhatsApp opens with the message already drafted: the order number, what was sent, and the tracking number. The box is in the Edit pop-up too, for a number you need to fix or read back over the phone.

The customer sees it. The number is published with the order, so a posted order's track page shows it on its own line under the delivery details - they can read it to the courier without messaging you. Self-collect orders show no such line. The shipped message is only offered on courier orders: a self-collect order keeps its **Send pickup reminder**, which is the message that fits that hand-over.

What you need to run once. The tracking number needs one column added to the tracking table - supabase/track_no.sql (a single ALTER, safe to re-run). Until that runs, everything else works; the number just reaches the message and not yet the customer's page.

15 Sep 2026 - engine v96 (no database setup needed). "Set day's availability" now shows how the day adds up. One change, on the pop-up you already use.

Why. v95 made the day's total count only the products you sell that day, but the pop-up that sets those numbers still showed nothing about where the total came from - and a product you do not sell that day sat in the list looking as if it were part of it.

What it does now. Under the product rows, the pop-up ends with **How the day adds up:** one line per product on sale that day with what it contributes, and a final line with the total - **the same number the order page uses for that day.** Type a + or a - and the lines and the total move as you type. Underneath, the pop-up names what it left out ("Not counted: Saturday loaf - not sold on this day, so no order can go on them here") and says what is already booked, so you can read how much the order page can still take. A product not sold that day keeps its row so you can still set its numbers; it simply does not enter the total.

One number, one place. That sum and the day's capacity are computed by the same function the shop's numbers come from, so what you read on screen and what your customers get cannot drift apart.

15 Sep 2026 - engine v95 (no database setup needed). A day's total now counts only the products you actually sell that day. One change, on your side of the app, with one knock-on for customers.

What it was doing. The "x/y" on a day - the chip beside the date on Orders, the same number on the day in the month calendar, and the dials on Home - was adding up the daily limit of every product on your menu, whether or not that product can be ordered on the day you were looking at. So a day with one product on it (a limit of 12, say) still read **1/42**: thirty of those units were limits belonging to products that could never take an order that day.

What it does now. Only the products on sale that day count. A Saturday-only loaf adds nothing to a

Wednesday, and a product kept on the shop as **Unavailable** adds nothing to a day it is not sold on. The same day now reads **1/12** - what that day can really take. Products with no sell marks are unchanged (they sell every day, so they always count), and a day where nothing on sale has a daily limit falls back to the **default capacity** in Settings, exactly as before - it does not drop to 0, which would make the app call the day Sold out.

The knock-on worth knowing. This same number is what the app tells the order page about a day, so a day now turns FULL at the moment every unit you are actually selling that day is booked, rather than at some larger number padded out by products that were never on sale. Past every real slot being taken, the order page stops offering that day - which was always the intent, and now happens at the right moment rather than late.

Nothing else moved. Your over-capacity warning, the "N ordered · N left" line on a delivery date, the purchasing plan and each product's own "N left" stamp all read the same number they always did for products that are on sale.

15 Sep 2026 - engine v94 (no database setup needed). One sentence on the customer's greyed card, reworded. Wording only - nothing behaves differently.

On a product you keep on the shop, the line saying why it cannot be ordered today now reads **Only available on Mon.** where it read "Only sold on Mon." The old wording was the seller's voice on a customer's card, and "available" is what the rest of that same card already says ("Only available for delivery from 19 Sep"). The Chinese and Bahasa Malaysia say the same thing as before, in their own words.

15 Sep 2026 - engine v93 (no database setup needed). The track link now lands on the track card, lit. One change, on the customer's side of the shop.

Why. Your confirmation message ends with a Track your order link. Tapping it opened the order page at the very top, with the tracking card somewhere below - the customer had to scroll and hunt for the very thing they had just tapped, on a page they have never seen before.

What it does now. The page comes up with the Track your order card already in view AND gently glowing, so it cannot be missed. The glow does not fade on a timer: it goes when their pointer or finger reaches the card, so it is still burning while they are looking around for it. That is the same rule your own app uses when you tap an order in the New orders box and it flashes the row - the flash is there until you arrive.

Nothing else changed. Typing a code into the box on the page behaves exactly as before (the card is already in front of them, so nothing moves), and a customer who simply opens the shop sees no glow at all.

15 Sep 2026 - engine v92 (no database setup needed). The New product card on Products folds away. One change, on one screen.

Why. The card that adds a product sat open at the top of the Products screen, and open it is a long form - it pushed your three lists (On the shop, Draft, Hidden) right off the bottom, so every look at what you sell began with a scroll past a form you were not filling in.

What it does now. The card arrives folded to a single line reading **+ New product**. Tap it and the form opens; tap the title again, or tap anywhere else on the screen, to fold it away. It stays open after you add a product - the toast still says it was saved as a draft - so you can type the next one straight away, and it starts folded again the next time you come to the screen. Nothing about adding a product changed: the form inside is exactly as it was, and it is still the same card that stays on the page while an Edit pop-up is open over it.

15 Sep 2026 - engine v91 (no database setup needed). The Orders calendar now answers a tap it cannot act on. One change, on one screen.

Tapping a day you do not deliver on. The month calendar at the top of Orders rings your delivery days in green with how booked each one is, and draws every other day of the month plain. A plain day never did anything before - tapping it was silent, which left you guessing at the reason. Now the calendar answers: a warm line appears under the grid naming the day and saying it is not a delivery day, and that delivery days are added in More -> Delivery dates. A marked day says its usual name above the day at the same time, so a holiday that is not a delivery date now tells you both what it is and why no order can go on it. Opening a real day, or tapping another plain one, replaces the line. Nothing can be ordered on a plain day, and nothing is created by

the tap - the calendar simply says so now instead of staying quiet. A plain day already gone stays silent: there is nothing left to add to it.

Where it appears. The month calendar at the top of Orders, the day picker inside the + New order card, and the Delivery day picker in an order's Edit pop-up - every calendar in the app that offers only your delivery days.

Nothing else changed. Delivery dates are still made in exactly one place, More -> Delivery dates (or **Generate next delivery dates** on Home). Tapping a calendar elsewhere in the app does not create one, and a tap that could open a day before still opens it.

15 Sep 2026 - engine v90 (no database setup needed). A product can stay on the shop on a day it cannot be ordered, with the date a customer can next have it. One new switch, off unless you turn it on.

Keep it on the shop. A product's Availability card now opens with a switch at the very top, above the sell-day calendar. It is for the two or three hot items customers come back looking for - if it is not listed they wonder whether you still make it - but leave it on for everything and the menu fills up, so it is a decision you make one product at a time. The card's own title carries the state too, so you can read it while the card is still closed.

What the shop shows with it on. On a delivery day the product is not sold on, the card stays instead of vanishing: greyed, stamped **Unavailable** (a day you never sell it is not the same as a day it ran out), with the reason it already writes ("Only sold on Mon") and one more line - the next date it can be ordered. On a day it has sold out the card is exactly what it was, **Sold out** stamp and all, and gains that same line. The line reads **Next available: Sat 19 Sep**, and names how many are left that day when a daily limit is set; it pulses gently so it reads through the dimming, and sits still - simply lit - if your phone asks for reduced motion. Nothing on a greyed card can be ordered: the + and - do nothing, it never enters a basket, and a line already in a basket still leaves with a note when the day changes. The switch never overrides **Draft** or **Hidden** - those still take a product off the shop entirely.

With the switch off, nothing at all changes. Off is the default and no product you already have carries the setting, so every product behaves exactly as it does today until you switch one on. One consequence worth knowing: on a day where every product is unavailable, a shop with kept items shows those greyed cards instead of the "nothing is on the menu" line.

14 Sep 2026 - engine v89 (no database setup needed). The Delivery Dates screen trades its list for the calendar: a ticked day comes off when you tap it again, and the past dates fold away into one group. Three small changes, all on one screen.

Tapping a ticked day takes it back off. A delivery date that was already on the calendar used to come off only from the list underneath it. Now tapping that green day takes it off on the spot - the same tap that put it there undoes it, so the calendar is the one place you both add and remove dates. If that day already has an order on it the app asks first ("2026-09-16 has 2 order(s) on it. Delete the date? The orders are kept in your delivery history."), because that is the one removal worth a second look; a day with nothing on it just goes, with a small "Delivery date removed" note. The date is all that is ever removed - the orders stay in your history either way.

The list of dates still to come is gone. The calendar already draws every delivery date as a green tick, and now it lets you take any one of them back, so the list below was repeating what was already in front of you. That is most of why this screen reads so much lighter. Nothing else moved: the green tick, the coloured holiday tags and the "N ordered · N left" line are all as they were.

The past dates are one group, folded away. Dates already gone keep their green tick but cannot be tapped, so nothing on a past day can be changed by a stray touch. They now sit together under a **Past dates (12)** heading that starts folded - tap the heading to open the list, tap it again to fold it back. Every past date is in there now: the old list showed only the ten most recent and left the rest out of reach. Each one still opens its orders or can be deleted exactly as before, with the same "Del" button and the same question when orders sit on it.

14 Sep 2026 - engine v88 (no database setup needed). The holiday tint sits on the day numbers at the Orders screen, and a mouse resting on a marked day names it. Two small things, both about the Orders screen and the bubble.

The tint now sits on the line of dates. On the Orders calendar each delivery day carries its booking count, and that count sat *under* the date number, which pushed the number 5.5px above the middle of its row. The occasion wash is centred on the row, so on this one screen the tint sat lower than the numbers it was meant to be washing - and the dates themselves did not line up with each other. The count now sits at the foot of the day, every date number sits on one line, and the tint runs through that line. Nothing else about the calendar moved: the green ring, the chosen day's fill, and the counts all read exactly as before.

On a computer the name comes up on hover. v87's bubble needed a tap. With a mouse, resting the pointer on a marked day now shows its name on its own, exactly as the customer's page does. This is deliberately only for a pointer that can hover: a touch screen keeps the tap, so a finger sweeping a run of days across the Availability or Delivery Dates calendars never drags a bubble along with it.

14 Sep 2026 - engine v87 (no database setup needed). Tap a marked day and it tells you its name - on every calendar in the app. v85 drew your marks on all five of the app's month calendars, but only as colour: a wash and a band, with nothing to say which holiday it was. The customer's own shop page has answered this since its calendar arrived - tap a marked day and a small dark bubble above it says the name - so the app now does the same thing, with the same look.

A tap names the day and still does that day's usual job. The bubble appears above the day ("Malaysia Day") and goes away as soon as you tap anywhere else. Nothing else about the tap changed: on the Orders calendar it still opens the day, on a product's Availability calendar it still marks or unmarks the sell day, on Delivery Dates it still starts the range you are about to mark. A day that sits inside two marks takes the name of the shorter one - the same rule the colours already follow, so the name and the colour always belong to the same mark. A day already gone is never named, because no mark is drawn on a past day either.

Two kinds of day could not be read at all before, and now can. On the Orders calendar a marked day you do not deliver was a quiet number with nothing to tap, and on the Delivery Dates calendar a marked day that is *already* one of your delivery dates could not be tapped either - its tick is removed from the list below the calendar, never by tapping its square, so its tap was free. Both now say their name when tapped, and nothing else about them moved: a Delivery Dates square still never ticks or unticks a date.

The Order date boxes name their day above the date, not above a day of the month. Those two fields fold their calendar the moment a day is picked, so a bubble drawn on the day would vanish with it. The name sits above the date the box is showing instead, which also means an order recorded on a public holiday says so for as long as that date is on the field.

Nothing was added, moved or removed, and the shop is not touched: this is the same bubble the shop's calendar has been showing, arriving on your own calendars.

14 Sep 2026 - engine v86 (no database setup needed). On the Orders screen your marked days now sit on their days, at the same depth as everywhere else. v85 put your marks on every calendar in the app. On the Orders screen the wash did not line up with the days underneath it. Two things were behind that, and both came from the same fact: the days on that one screen are taller than the days anywhere else, because a delivery day there carries its booking count ("3/12") under its number.

A holiday band was stretched to fill its week row. Where every day is the app's usual height, a band came out the same 30px sheet a single-day box is, so the two shapes matched - which is how the rest of the app has always drawn them. On the Orders screen the rows are taller, so the band there came out deeper than a single-day box, and deeper in the weeks that held a delivery day than in the weeks that did not, so one holiday's own band changed depth from row to row inside the same month. **A band is now a sheet of one depth, centred on its row** - the same depth a single-day box has, on every calendar, whatever height the days around it happen to be.

A plain day kept the app's usual height inside a taller row, so it sat at the top of the row rather than on its middle. A single-day mark was therefore drawn 6px higher on a day you do not deliver than on a day you do: the same holiday, on a different line, decided by whether you deliver that day. **Every day of the Orders month is now drawn at one height,** delivering or not, so a mark sits on the same line as the day it covers, beside the

booking counts.

One more thing fell out of that. Today's soft glow is drawn around a whole day on the app's shorter calendars, and around a taller day it would have become a box again - the very thing that glow replaced. **Today on the Orders screen now glows around its number**, the same round glow a delivery day's number carries, so today reads the same whether or not you deliver that day.

Nothing else changed. The shop's calendar and your Delivery Dates calendar look exactly as before (their days are all one height, so they never had this), and no mark was added, moved or removed: this is the drawing only.

14 Sep 2026 - engine v85 (no database setup needed). Your marked days are now drawn on every calendar in the app, not just the one you marked them on. Until now a holiday only appeared on the customer's calendar and on the Delivery Dates screen you marked it on. Every other calendar in the app drew a plain month: the month calendar at the top of the Orders screen, the little **Order date** calendars in the new-order form and in the Edit-order pop-up, the delivery-day picker inside that pop-up, and a product's **Availability** calendar where you mark the days it sells. So a holiday you were planning around was invisible on every screen you actually plan on.

Now all of them draw your marks, in the same two shapes and the same colours as the Delivery Dates calendar - a see-through band across a run of days, deeper the shorter the run, and a box for a single day. A mark on a day you do not deliver is drawn too, on a dimmed day, since that is the only way a holiday on a quiet day can be seen at all. Nothing about the marks themselves changed: they are still only ever something to see, and a mark still never adds or removes a delivery date or changes what a product sells. Two small pieces of the drawing were tidied while moving it: the green tint on a day a product sells is now a see-through wash rather than a flat fill, so a holiday band running behind it still shows through instead of the mark disappearing on exactly the days you sell (the tint itself looks the same as before); and a marked day that happens to be today now keeps the soft glow that says "this is today", the way the shop's calendar has always drawn it.

14 Sep 2026 - engine v84 (no database setup needed). A day you marked is now the same gentle wash everywhere, whichever calendar you are looking at. A marked day was drawn one of two ways, decided by a single question: is this a day you deliver for? On a day you deliver for the mark went pale, so the green "you can order" ring stayed the first thing the eye read. On a day you do not, the mark was drawn at full colour, as a solid block with the number turned white. September's Malaysia Day sat on a delivery day and read as a quiet tint; October's World Animal Day fell past the delivery days published so far and read as a loud solid block - so the same kind of holiday looked like two different things, and a mark quietly changed its own appearance as the delivery window moved towards it.

Now a mark is only ever a see-through wash of its own colour, and its depth follows how long the mark runs: a one-day holiday the deepest, a long school break the palest. The customer's calendar and your own Delivery Dates calendar read that depth from the same place, so a day wears exactly the same mark on both sides of the shop and in every month. Nothing is ever a solid block of colour any more, which also means the date number and the green delivery pill always read on top of a mark, whether or not you deliver that day. Only the look changed: marking, removing, renaming and loading standard occasions all work exactly as before, and a mark still never adds or removes a delivery date.

14 Sep 2026 - engine v83 (no database setup needed). Marking a holiday now tells the shop straight away. A day you mark on the Delivery Dates calendar is drawn on the customer's own calendar (engine v79 to v81). But saving a mark only ever sent your delivery dates to the cloud - it never sent the customer's page anything - so the marks reached the shop only by accident, whenever you next happened to save a product or a setting. Mark the holiday, open the shop, and nothing had changed. Found while you were looking for a holiday on the shop and finding none.

Now the same save does both. Marking a day, removing a mark, loading a set from the standard list, or renaming or recolouring one all republish the customer's page within a couple of seconds, exactly the way editing a product does. Nothing else about a mark changed: it is still only ever something to see while planning, it still never adds or removes a delivery date, and a name you typed yourself still never leaves your phone.

One tap to catch up the marks you already have. The fix sends the marks from now on; the ones already on your calendar have never been sent at all, so open More -> Settings -> Storefront and tap **Publish now** once. After that the shop's calendar follows your marks by itself. Nothing to run in Supabase.

14 Sep 2026 - engine v82 (no database setup needed). A product's selling days are now marked on a little calendar inside the product's own screen - weekend, certain days, a season, or any mix of them.

Until today a product could say only two things about *when*: a notice period, and one from-date and to-date for a season. You asked to set it freely instead, so the two date boxes have become a card you mark.

Open a product and tap the Availability card. It sits under Daily limit and stays folded until you tap its title, exactly like the product-text card you already know. Folded, its title tells you what is marked - "Every day", or "Sat & Sun", or "1-24 Dec 2026" - so the Products screen alone tells you a product's selling days without opening anything.

Inside, it is a month calendar and four gestures. Tap a weekday letter (M, T, W...) and every one of that weekday **in the month shown** is marked: tap S twice and the product sells on Saturdays and Sundays. Tap one day to mark just that day. Slide your finger across several days and the whole run is marked. Tapping or sliding over a day that is already marked takes it back. Everything you mark is added together - mark Saturday and Sunday, then a week in the middle, and the product sells on all of them.

Nothing carries over into the next month, which is the point of it: marks belong to the days you marked and no others, so a December-only set stays December-only and January starts with nothing marked. Page to another month with the arrows and mark that one too if you want to sell then. And a product whose card you have never opened sells on every delivery day exactly as before, so nothing on your live shop moves until you mark it.

A period that crosses a month or a year is one mark with two ends you can stretch. The card lists your marks under the calendar; tap one and its **Starts** and **Ends** dates appear above it, and you can push either end as far as you like - a Sat & Sun run from 1 Dec to 4 Jan is one line, not two marks. Leave an end empty and the mark simply has no bound on that side ("from here on", or "up to here"). Type the two dates the wrong way round and they are put in order for you. A small x on a listed mark removes it whole.

On the order page, a day a product is not sold for is not a note - the product is simply not there. No card, nothing to want and not have. If a customer already had it in their basket and then switches to a day it is not sold for, it leaves the basket and a short line says why. The one rule that keeps its note is the notice period: a product needing 14 days IS sold that day, it only has to be ordered earlier, so it stays on the menu reading "Orders close 14 days before delivery - pick a later date" - the two are different questions, and only the second one hides anything.

Two things worth saying plainly. The old From/To season boxes are gone; a product that had them opens with that period already sitting in the card as one mark, so nothing you set before is lost. And a product whose selling days have all gone by keeps them, on purpose: your own rule is "if I have not indicated a selling date, it is not selling", so a mark you leave on says exactly that, and the way to put a product back on sale every day is to take its marks off with the x. No database setup, and delivery dates, capacity, value packs, ingredients, suppliers and everything else are untouched.

14 Sep 2026 - engine v81 (no database setup needed). A holiday you marked is now drawn on the shop's calendar exactly the way your own calendar draws it: a pale band across a run of days, a solid box for a single day. The marks were already reaching the shop; it was only their shape that differed, and a mark that looks different on the customer's page than on yours is a mark your customer has to learn twice.

On the shop, a marked day now wears the mark your own calendar gives it. The shop had been drawing each marked day as its own flat pale square, so a holiday running over a week read as a row of separate blocks where your own calendar shows one continuous shape. Now a holiday running over **several days** is a **pale see-through band** across those days - one unbroken rounded band per week row it crosses, in its own colour, and **deeper the shorter the run**, so a one-day holiday still stands out inside a long school break. A **single-day** holiday is a **solid box** in its full colour on its one day, a touch narrower than the band so a longer band still

peeks out at its sides. The eight colours are your own calendar's eight, and the strength steps are the same three, so a marked day reads the same to your customer as it does to you.

The one place it gives way: a day you deliver for. Green is the colour the delivery ring uses, and a solid box would swallow that ring - so on a day the bakery delivers for, the day a customer can actually order, the box goes pale there and the **green ring stays the first thing the eye reads**. The mark is still there on that day, and the day is still named on a tap. This is the one guard on the shop's calendar, and it covers every mark colour rather than green alone, so no colour can ever hide the days that can be ordered for.

Nothing else moves. Tap a marked day - or, on a computer, rest the mouse pointer on it - and the same small bubble names it, e.g. "Malaysia Day"; tapping anywhere else puts the bubble away, and still nothing is listed under the grid. Today still breathes its soft glow. A run still stops at today rather than colouring days already gone, exactly as your own calendar leaves past days alone. And your own typed-in days still never reach the shop: only a day that is, by name and date, one of the built-in standard days can be published, exactly as before. No SQL, no schema change. The calendar the shop draws with is still a copy of your app's own, and a test now pins the two copies' mark rules together as well as their month grids, so a day can never again look like one thing here and something else there.

14 Sep 2026 - engine v80 (no database setup needed). Today glows instead of wearing a box, a holiday you marked is a soft tint you can tap to name, and the Orders screen's sideways date strip becomes the same month calendar the shop shows. Three small things, all pulling the same way: a calendar should read the same on both sides of the shop, and nothing drawn on it should look like something it is not.

On the shop, today no longer wears a little red rectangle. It always meant to be a ring around the date, but on a day you do not deliver the app had no circle to draw that ring on, so it traced the day's number as a box instead - which is why it looked like a red square on the quiet days and only looked right on the days you deliver for. Today is now a **gentle glow** that breathes around the number, with no box at all, and it sits happily on a day you deliver for too: the **green ring** says you can order for this day, the soft brown glow says it is today. Under **prefers-reduced-motion** (a phone set to calm its animations) the glow holds steady instead of pulsing, so it is still there to find. The same fix reaches the app: the today marker on the Delivery Dates calendar, and inside every date picker, was a hard outline and is now the same soft glow.

A holiday you marked is now a soft wash, and its name comes out on a tap. The shop's calendar used to mark a standard day with a small coloured dot and then list the month's names in a "Holidays" line under the grid - which crowded the very line that line was there to keep clear, and made a holiday compete with the days a customer can actually order for. A marked day now wears a **pale wash** of its own colour over the whole day. The wash is deliberately neither a ring nor a filled dot, because a ring is what "you can order for this day" looks like, and the two must never be read for each other. Tap a tinted day - or, on a computer, rest the mouse pointer on it - and a small dark **bubble** above it names the day, e.g. "Malaysia Day". Tapping anywhere else puts the bubble away. The "Holidays" line under the grid is gone. Your own typed-in days still never reach the shop: only a day that is, by name and date, one of the built-in standard days can be published, exactly as before.

In the app the same holiday is named **beside the day** instead of in a list: open a delivery day and, if it carries a mark, its name sits under the date in the mark's own colour. Nothing at all is drawn on a day you have not marked.

The Orders screen's date strip is now a month calendar. The row of date chips at the top of Orders grew a little longer every week. It is now the same **month calendar** the customer sees - one month at a time, with left and right arrow buttons either side of the month's name. The days you deliver for are ringed in **green**, and under each one is **how booked it is**: the same "3/12" the Home dials show, or **FULL** when the day has no room left. A day whose orders have already closed shows its count in **red**; a day already gone is dimmed but still opens, because you backfill and look back at old days. Every other day of the month is plain and cannot be tapped. Tap a green day and the app's view moves to it - that day is filled in on the calendar and its orders appear below - and tapping a day in the list below moves the calendar with it, so the two always agree. The arrows reach only the months your delivery days fall in, so paging never strands you on an empty month.

The + New order card arrives folded, and leads with the calendar. It used to sit open on every delivery day, which is not what that screen is for day to day. It now starts **shut**, showing just its title, "+ New order" - tap the title to open it, and tap the title again (or tap anywhere else on the screen) to fold it away. Opened, it reads in the order you think in: **the month calendar first**, the same one the shop shows, already sitting on the day you are looking at, so tapping another green day moves the whole screen before you have typed anything; **then the customer's details** - name, WhatsApp number, self-collect or courier, a note; **then the items**, with "+ Add another item" and "+ Add order" at the end. It stays open after you add an order, so you can type the next one straight away, and coming back to the screen later starts it folded again.

The Edit pop-up now reads the same way. Its **Delivery day** control was a button you had to tap to reveal a calendar; the calendar is simply there now, already open on the order's current day. The pop-up then reads the day, then the customer's details, then the items - the same order as the card. Every day still to come is tappable, plus the order's own day even if it has passed, so an old order always shows where it is, and the soft notes re-read underneath the moment you pick a different day. Moving an order is unchanged in every way that matters: still one delivery day per order, and capacity, the confirmation and the payment reminder, the track page and the WhatsApp messages are all exactly as they were. The button-style day picker the app used in these two places has been deleted along with its tests, because nothing calls it any more.

14 Sep 2026 - engine v79 (no database setup needed). The customer picks a delivery day on a calendar, and so do you in the app. The shop used to show its delivery days as a row of date chips laid out sideways. Now, under **Pick a delivery day**, the customer sees a **month calendar** - one month at a time, with the arrow buttons either side of the month's name to look ahead or back. The days you are taking orders for are picked out in **green**; every other day of the month is plain and cannot be tapped. They tap a green day, and the day they chose is written out in words just underneath - **"Your delivery day: Wed, 16 Sep"**. That is the one line of text below the calendar, exactly as asked: no dates floating outside the grid, and still **one delivery day per order**. The first open day is already chosen for them when the page opens, so ordering can never be blocked by forgetting to tap. Nothing else about taking orders changed: a day you have cut off is still not shown at all, and a day that is full still greys out with a red **Sold out** badge and cannot be tapped.

Only the standard days you have marked reach the customer. If you have loaded days from the built-in list onto your own Delivery Dates calendar, those days now wear a small coloured dot on the customer's calendar, and a short **Holidays** line under the grid names the month's ones (e.g. "16 Sep · Malaysia Day"). They follow the colour you gave them. A day **you typed yourself** - a birthday, a promo, a school break - never appears: only a day that is, by name and date, one of the built-in standard days can be published, so your own notes stay private to your phone. Mark nothing in the app and the shop's calendar simply carries no dots and no caption.

In the app, the same little calendar now stands wherever you name a date. Three places changed. In the **Edit** pop-up, **Delivery day** is no longer a drop-down list: it names the order's current day, and tapping it opens a small month calendar beneath it - every day still to come is tappable, plus the order's own day even if it has passed, so an old order always shows where it is. The **Order date** box in that pop-up is a calendar too, with a **Today** shortcut for an order you are typing in right now. And the **+ New order** card has gained a **Delivery day** picker at its top: tap another green day and the whole screen moves to that day, so the product list you are choosing from is always the one that is actually sellable for the day you picked. The soft notes that help you decide - the change/cancel window, a warning, a closed day, too little left - all still sit under the calendar, and none of them stop you.

The calendar expands in place under the button rather than opening in its own window, because the app's pop-ups share one layer and a calendar opened inside one would wipe out the Edit pop-up it was opened from; expanding in place also reads better on a phone. No database setup needed - nothing about how orders are stored, counted, bought for or messaged changed.

13 Sep 2026 - engine v78 (no database setup needed). The two confusing "Add" buttons on the Delivery Dates calendar now say what they do. On the Delivery Dates screen, **Mark an occasion** used to carry a small button reading **"+ Add occasion"**. The name did not tell you that it opens a ready-made list of Malaysia's holidays and fun days, all ticked and waiting - so you had to open it to find out. It now reads **"+ Load standard**

occasions", and the window it opens is titled the same.

Inside that window there were **two buttons both called "Add"** - the big one at the bottom that files the days you ticked, and a smaller one in the **My own day** strip just above it, which files a day you name yourself (a birthday, a one-off promo). Reaching for "Add" and hitting the wrong one gave you **"Type a name first"**, for a day you never meant to name. The smaller one is now spelled out as **"Add my own day"**, so the only plain **Add** in the window is the one that adds the ticked days.

Nothing else moved. Ticking a row and pressing the big **Add** still files exactly what you ticked and nothing you did not; **Untick all / Tick all** and the heading tick boxes work as before; the **My own day** box still takes any name and any date. No database setup needed.

13 Sep 2026 - engine v77 (no database setup needed). The greyed line in the translated-text card now says what happens if you leave it blank. Engine v76 turned the product screen's Mandarin / Bahasa Malaysia text into one folding card, and each empty line showed its translation on offer. That greyed line read "e.g. <the translation> - blank keeps English". It now reads **"e.g. <the translation>.....if blank, it will be filled with English"** - the same words you asked for, saying plainly what the customer gets rather than asking you to work it out. The line in the card's own explanation says the same. Deleting what is in a line brings the greyed line and its arrow straight back, exactly as before. Nothing else moved: the card still folds away when you tap outside, the -> still takes the words, the circular-arrow button still translates that line again.

13 Sep 2026 - engine v76 (no database setup needed). The translated text on a product is now one quiet card that folds away, filled in for you one line at a time. The **Product text for your customers** part of the New/Edit product screen - the same words in **Mandarin** and **Bahasa Malaysia** - was the busiest corner of the whole app. It was always open, so a long screen got longer before you needed any of it; eight lines each carried its own **"Translate this one"** button; two **"Fill all"** buttons sat above them; and a little **"auto"** tag marked the machine-made lines, in a word you had to remember the meaning of.

Now the card starts **shut** and stays shut - it is the setup part of the screen, not the part you touch every day. **Tap its title to open it.** Opened, each line already has its translation worked out and waiting as the same pale grey suggestion you know from every other box, with the little -> at its right edge: **tap the -> and the suggested words drop in** - exactly the gesture you already use. The moment you take them, the arrow is replaced by a small **circular-arrow button** - tap it and that one line is translated again on the spot, which is what the old "Translate this one" button did. **Type your own words over any line and the circular-arrow button goes away** for that line: it is yours, and it is never overwritten. **Leave a line blank** and that language keeps the English, the same as before. And when you are done, **tap anywhere outside the card and it folds away again.**

The **"Fill all Mandarin" / "Fill all Bahasa Malaysia"** buttons are gone, the eight **"Translate this one"** buttons are gone, and there is no **"auto"** tag to read any more - a line showing the circular-arrow button is machine text, and a line with no button is yours. Nothing about what your customers see has changed, and nothing needs setting up: a product is still translated for the shop the first time you save it and again when you Publish, exactly as before, so a product you never open this card for still reads in all three languages on the order page.

12 Sep 2026 - the TNG QR box no longer claims to be on the order page (no new engine; your phones did not change). The box in **Settings -> Storefront** holding your Touch 'n Go QR image link said it was **"shown on the customer's track page"**. It never was - your payment QR reaches the customer in the **WhatsApp confirmation** and the **payment reminder**, and nowhere else. The help text now says that, and the order page's own settings no longer carry the QR at all, so it cannot be drawn there by accident. Nothing about how you take payment changed, and nothing about how the app works changed.

12 Sep 2026 - engine v75 (no database setup needed). The glow on an order you have just jumped to now stays until you reach it. Tapping a row in the **New Orders** inbox (or a result in **Find an order**) takes you to that order's delivery day and marks the order with a green glow - but the glow used to fade after less than two seconds, which is no time at all when your eye is still travelling down a long day, and you were left hunting for a row that was no longer lit. **The glow now keeps pulsing until you get to the row.** It stops the moment your

finger or the mouse arrives on that order, or when you tap it open - so it is waiting for you however long you take, and it never lingers as a puzzle later. (If a phone has **Reduce Motion** turned on in its accessibility settings, the same green outline holds steady instead of pulsing - it still stays until you reach the row.)

12 Sep 2026 - engine v74 (no database setup needed). The little arrow in an empty box now keeps working. Engine v73 added the small right-arrow that sits at the right edge of a box showing a greyed example: press the **-> key** on a keyboard, or tap that arrow, and the example drops in as real text. On a phone that tap worked the first time and could then go dead - once you had tapped into a box, tapping the arrow on the next one did nothing.

Two things were behind it, and both are fixed. The tap was measured against the **page**, and the phone's keyboard shifts the page under your finger as soon as any box is focused, so the arrow was no longer where it had been. And taking a suggestion **opened the keyboard itself**, so the page jumped after the very first tap. The tap is now measured against the **box itself**, taking a suggestion no longer opens the keyboard at all, and the strip you can tap is about half as wide again - so a thumb that lands a little wide of the arrow still counts. A tap on the greyed text stays an ordinary tap that just puts your cursor there. **The arrow looks exactly as it did**, and nothing else changed.

12 Sep 2026 - a new order now opens on the order itself (no new engine; nothing to set up). Tapping a row in the **New Orders** inbox took you to that order's delivery day - and stopped there. On a busy day, with a long list of orders, you then had to read down the list to find the very order you had just tapped, and it was easy to open the wrong one.

The app now goes the whole way: it opens the delivery day, **scrolls the order to the middle of the screen and flashes it** for a moment, exactly as the **Find an order** search box already did. If the list had been narrowed with the status filter so the order was hidden, that filter is cleared first, so the order is always there waiting for you. Nothing else on the Orders screen changed.

12 Sep 2026 - engine v73 (no database setup needed). Five changes, all about the shop's promise to the customer and the baker's control over it.

Your rule is that an order is **not cancellable and money is not refunded, but it can be moved to another day**, and a customer may ask until a set number of days before delivery. That number now lives **per product**. Open a product in **Products** and, under its daily limit, a new box **"Changes or cancellations (days before delivery)"** - type 2, say, and that product's card on the shop reads **"Change or cancel up to 2 days before delivery."** It only ever *tells* the customer; it never blocks you, and you still move any order by hand. Leave it blank (or 0) and no note shows for that product. When one order holds several products with different numbers, the customer sees the **strictest** one - the largest, i.e. the earliest cut-off - so a mixed basket always gives one clear figure. The same sentence appears on the green **"Order received!"** card just after they place the order.

Second, **you can now move an order to another delivery day in the app**. Open an order with **Edit** and, in the pop-up, a new **"Delivery day"** box lists every day still to come, soonest first. Pick the new day and **Save changes**: the whole order moves as one - items, customer and the price it was sold at all stay put - and the customer's **"Track your order"** page follows to the new day straight away. Under the box are quiet notes to help you decide: the order's own change/cancel window, a reminder when the new day falls **inside** it, a note if that day's orders have already closed, and a note if that day does not have enough left for an item. None of them stop you - they only tell you - and the capacity check now counts the order on the **new** day, so you can see at once if that day would go over. This is the **only** way to change an order's day: never delete an order and re-type it, because deleting reads as a cancellation. A day left empty stays in your list - to remove it, use **Del** on the Deliveries screen as before.

Third, **a Policies box** for your own cancellation and refund wording. It sits in **More -> Settings -> Storefront**. Type it once in English, tap **Translate**, and the Mandarin and Bahasa Malaysia boxes are filled for you (edit either and it becomes yours). It appears on the shop **under Track your order**, in the customer's own language and with the line breaks you typed, and it stays hidden until you write something.

Fourth, and everywhere in the app: **an empty box showing a greyed suggestion can now be accepted with**

one tap - the arrow. A small right-arrow sits at the right edge of such boxes; press the **right-arrow key** or tap that arrow and the suggestion drops in as real text you can edit - the same gesture as accepting a suggestion in an AI chat. It is offered only where the grey text is a genuine recommendation; the sign-in and cloud boxes never offer it, so a made-up key or password can never be accepted by accident.

No database setup and no new SQL - the new product box and the Policies wording travel with your existing shared data.

9 Sep 2026 - the wish list's automatic email went live (no new engine; the phones did not change).

Adding a wish on **More -> Software wish list** now emails the **full wish list** - every wish, ticked or not, newest first, with the app version and the date - to the developer address(es) you set under Settings -> **Website & developer**, all on its own, whenever the phone is signed in and sharing data. Before today that automatic email needed a behind-the-scenes service that was not switched on yet, so only the **"Email the full wish list"** row sent it by hand; that row is still there as a backup any time.

10 Sep 2026 - the homepage opens faster (no new engine; the phones did not change). Two things about the public homepage were quietly costing every visitor time, and both are fixed.

The homepage's six photographs were stored *inside* the page itself, so the page was **249 KB before a single picture had even been asked for** - and every visitor downloaded all of it, on every visit, before they could see anything. The photos are now six separate files that the browser fetches and **keeps for next time**, which brings the page itself down to **22.7 KB**. The pictures are byte-for-byte the same ones, so the homepage looks exactly as it did, and each one now states its size up front so nothing jumps around as the page loads.

Second, the "What customers say" section used to fetch **every published review photo up front** - including the reviews a visitor never waited around to see. A photo is now fetched **as its review comes around**, one step ahead of time so it is ready before it slides in, and the reviews a visitor never reaches are never downloaded at all. On a phone that means less data and a quicker page, and it changes nothing about how the section looks or behaves.

A review photo is also shrunk a little harder before it is stored - to 1000 pixels across instead of 1600. The card on the homepage shows a picture no more than 340 pixels high, so the smaller copy looks just as sharp while being about **a third of the size**, which is kinder to the customer's mobile data and keeps your stored photos light. This applies to photos left from now on; the reviews already on your page keep the picture they were stored with.

Your sign-off line - **Jien Luv 2 Bake | Small Batch. Big Love.** - has moved from the very bottom of the homepage footer, where it sat shoulder to shoulder with the developer credit, up to the **top of the footer**, above your contact details and set off by a thin line. It reads a touch larger now, and in whichever language the visitor has chosen, like the rest of the footer. The "Website by ..." credit stays where it was and now stands on its own.

10 Sep 2026 - the homepage's language buttons can no longer go dead (no new engine; the phones did not change). On some phones the homepage's **EN / Mandarin / BM** buttons could stop responding - the page would sit in English and tapping Mandarin did nothing - often after the page had been left alone for a while, or refreshed. The culprit was a leftover background worker from when the backoffice was served from the top of the website instead of from /admin/. (A background worker is the machinery the app installs so it can open without a connection.) That old worker had cached the whole website, and whenever something failed to download it answered with whatever it had - including handing the browser **a web page where it had asked for a script**. A page that receives a web page instead of its script never runs it, so the code that wires up the language buttons never ran; the English words you saw are written into the page itself, which is why the page still looked perfectly normal.

The old worker is now cleared away wherever it is still installed: a small cleanup file sits at the same address the old worker used, so the next visit replaces it, empties everything it had cached, and removes itself. The backoffice's own offline cache is deliberately left untouched, because the app needs it to open without a connection. The app's worker has also been taught, for good, never to answer a script or a picture with a web page - so the same fault cannot come back on any part of the site.

v72 - The shop's sold-out notes read in the customer's language (10 Sep 2026)

The shop page switches language in place, but two corners of it had stayed English no matter which language a visitor chose: the short note on a product that **cannot be ordered for the chosen day** (that it only opens from a certain date, only runs up to one, or that orders close so many days ahead - with the date written in the page's language), and the notes that appear above the menu when the page **has to change a basket** - an item that just sold out, or a quantity trimmed to what is actually left.

Both now read in **English, Mandarin or Bahasa Malaysia** along with everything else. The reason a product is closed is now handed to the page as a plain fact - which kind of rule, and the day or number of days behind it - and the page writes the sentence itself, so no English can slip into a Chinese or Malay page. An English visit is word for word what it was.

v71 - The shop's Mandarin / BM buttons answer instantly (10 Sep 2026)

On the shop page, tapping **EN / Mandarin / BM** used to reload the whole page. That re-downloaded the page and fetched the menu, the "left today" numbers and your saved storefront settings all over again, so each tap sat there for a moment before anything changed - and the slower the phone's connection, the longer the wait. A reload could also wipe a basket that was already half filled.

Now the shop switches **in place**, exactly like the homepage already did: the words change the moment you tap, and nothing is downloaded again. Whatever the customer has already done is left untouched - the items in their basket, the delivery day they picked, the name, WhatsApp number and address they typed - and if they are looking up an order, that card re-reads in the new language too. Nothing to set up.

v70 - An order keeps the price it was sold at (10 Sep 2026)

Until now an order only remembered **which product** was bought, not what it cost or what it was called. Everything - the amount on the customer's confirmation, the total on their own "Track your order" page, their lifetime spend on the customer list, the Home estimate - was worked out from **today's** menu. So the moment you re-priced or renamed a product, every past order silently moved with it: last month's order suddenly read RM22 instead of the RM15 the customer actually paid, and a renamed loaf rewrote what people had already bought.

Now an order is a record of a sale. When an order comes in - from the shop or typed in by hand - it keeps the **name and the price it was sold at**, and those are what every message, the tracking page and the customer's spend use. Rename or re-price a product today and yesterday's orders stay exactly as they were. A line you deliberately swap to a different product when you **Edit** an order does pick up that product's current price - that is a new sale - but a line you leave alone keeps its old price.

Existing orders (taken before this version) are stamped once with today's values, which is exactly what they are showing right now, so they stop drifting from here on. Nothing to set up; the phones pick it up with the next sync.

v69 - Either place updates the customer (10 Sep 2026)

Yesterday's fix put a customer's name and number back in one place - but it made the **customer card** the side that always won, so if the card and an order disagreed, the card's name was the one shown. That is no longer the rule: now it is simply **whichever you edited last**. Correct a customer's name with **Edit** on one of their orders and that is the name the list, the history pop-up and the "Hi {name}!" greeting use; open their card, type the name and save, and the card's is the one used. Either place updates the customer - neither is the boss. They normally read the same anyway, because each save writes through to the other side; this only decides the rare case where an old copy and a newer one are both sitting there. Nothing to set up; the phones pick it up with the next sync.

v68 - One name and one number per customer (10 Sep 2026)

A customer's name and WhatsApp number used to be held in two places at once - on their orders, and on their customer card - and the two did not always agree. If you opened a customer and typed their name on their card, the customer list could carry on showing the old one, because it read the name off the order. And renaming someone who has **no WhatsApp number** on file could lose them: the name is what ties a numberless customer to their history, so the new name slid off their orders and appeared nowhere at all.

Now the name and number you save on a customer are the ones that count, and fixing them once fixes them everywhere: saving their card writes the name and number onto **every order they have**, so the labels, the WhatsApp messages and the customer list all follow. It works the other way round too - correct a wrong number with **Edit** on an order and that customer's card updates as well, so the two can never drift apart again. Correcting a number carries that customer's bring-a-friend credits across with it, and renaming someone who has no number no longer loses them.

Leaving a box empty means "leave this as it is" - it never wipes a number your confirmations depend on. And if you renamed anyone before today and it never stuck, the app puts their orders right the first time this version opens. Engine v68, guide v68.

v67 - Naming a "No name" customer (10 Sep 2026)

A customer who reached you without a name - a storefront order that carried none - used to stay **"No name"** in your customer book even after you opened them and typed a name. Now it sticks: open them, tap **Add details**, type the name (and the dog's name, what they like, and so on) and save, and their row shows that name straight away - in the list, the finder and the 'Hi {name}!' WhatsApp greeting. A tidy-up; nothing else about your customers changes. Engine v67, guide v67.

v66 - Auto-translate product text, plus Draft / Publish / Hidden states (9 Sep 2026)

Product **description**, **selling unit word** (the "RM15.00 / loaf" bit), **serving tip** and now the **name itself** read in English, Mandarin or Bahasa Malaysia automatically - a free translator fills the Mandarin/BM boxes for you the first time you save or publish a product. A small "auto" tag marks a machine translation; if one "sounds weird", open the product's edit screen, type over that box (it then becomes yours and is never overwritten), or tap a "Fill all Mandarin" / "Fill all Bahasa Malaysia" button or the circular-arrow button next to a single line to have it translated again now. Good translations are simply left alone, and you never hand-fill a whole product.

Products also gain a **third state**. A new product now starts as a **Draft** - fully built with its recipe, but **not on the shop and not orderable** until you **Publish** it. The Products screen shows **three separate lists**: **On the shop**, **Draft - not on the shop yet** (each with a Publish button) and **Hidden - taken down** (with an Unhide button). Hiding a product still keeps its history, recipe and purchase orders.

The **"Copy follow-up"** bring-a-friend message gains a small **EN / Mandarin / BM** choice, so the whole check-in you paste into WhatsApp - including the product's serving tip - can be in your customer's language. No SQL this version.

v65 - Reviews preview + "waiting to publish" cue (9 Sep 2026)

More -> **Reviews** now shows each review **exactly as customers see it** on the homepage - the photo, the stars, the message, then the reviewer's name and the language and date, in the same warm card - and they play **one at a time**, like the homepage's "What customers say" carousel. The reviews waiting for your **Publish** come first; once you have worked through them, the ones already on the homepage follow. It never slides on by itself while you are deciding - a swipe, the arrow buttons, or a dot under the card moves it on, and Publish / Take down /

Delete moves you straight to the next review. A line at the very top of the screen says how many are waiting ("3 reviews waiting for you to publish"). The **Reviews** row on the **More** menu shows how many are waiting too, and the app's **Home** screen gains a "**N new review(s) to publish**" card that opens the Reviews screen in one tap - so a new review is never left waiting unseen. No SQL this version.

v64 - Review carousel, EN/Mandarin/BM site, developer contact + wish email (9 Sep 2026)

The homepage "**What customers say**" reviews now play as a **carousel** - one review fills the section at a time, then slides on every few seconds. Visitors can swipe or use the left / right arrows and the dots below to move through them, and it pauses while someone is reading. One review, or none, still shows simply.

The whole **homepage and storefront now read in English, Mandarin or Bahasa Malaysia** via an **EN Mandarin BM** switch near the top of each page (English is the default, and a visitor's choice is remembered next visit). Every product can carry its own **translated shop names** - Products -> Edit -> **Shop names** - so shoppers in Mandarin/BM see a proper translated name on the order page; a blank box keeps the English name, which always stays the real name in the app, labels and history.

The thank-you after sending a review is warmer and honest: it reassures the sender that their review and photo reached the baker and will appear once approved, and a critical review (3 stars or fewer) gets an extra line of thanks for honest feedback. A reviewer now attaches their photo **two ways: Take photo** (camera) or **Choose photo** (gallery).

Developer contact is now in the app: under More -> Settings -> **Website & developer**, type the developer's name and one or more email addresses. A small **Website by ...** credit links from the bottom of the homepage, the storefront footer and More -> About (each address tappable). Add the developer's optional **WhatsApp number** in the same card and that link opens WhatsApp (Business) with a ready "Hi!" - message the developer straight from the page - with the email links staying underneath. And adding a **software wish** on the More screen emails the developer the **full wish list** - every wish, ticked or not, newest first, with the app version, the project address and the date - to every address you listed; an **Email the full wish list** row on the card can send it from your own mail app too. No SQL this version.

v63 - Customers tab, wish list, holidays fix (8 Sep 2026)

Customers is a real tab of its own, swapped with **Purchase Order** (now under **More -> Purchase Order**). The tab stays your automatic customer list - who ordered, how much, their favourite - and now every person can carry a **profile** (the **Edit** button in their history pop-up): their dog's name and a small photo (auto-shrunk to a thumb), what they like, what to avoid, and a note. Profiles save into a synced customers collection, so both phones know the same people, and they build up quietly for a future AI chat about your customers. A **finder** box at the top of the tab filters the list as you type - name, number, dog's name, likes, notes - and says how many match, just like "Find an order" on Orders. A saved profile shows as a paw or photo beside the name and a small line under the row.

On **Home**, the "Upcoming holidays" card is fixed: it used to vanish whenever nothing was marked on the Delivery calendar. It is now always there - up to three marks at a glance, a scroll for any more, and a friendly empty state that still opens the calendar so you can add days.

On **More**, a small green **Engine v63** pill (in the header card) shows which build this phone runs, a **Software wish list** sits at the bottom and behaves exactly like your to-do list (add, tick, reword, remove - a tick stays ticked and never touches the weekly routine), and **Full change history** opens this history as a PDF kept on your website. No SQL this version.

v62 - "Not sharing" strip + read-only View (8 Sep 2026)

When a phone is not on the shared cloud - shared data off, never set up, or signed out - a thin amber strip at the very top of every screen reads "**Not sharing right now**" with a one-line reason and a **Fix it** tap straight to Settings -> Shared data. It never locks the app, and it disappears on its own the moment sharing is back on. (Set up once per phone: no SQL - the shared-data setup from v53 does everything.)

Also, every backup copy in Settings -> **Backup & safety** now has a **View** button: a purely read-only look inside that copy - its orders grouped by delivery date (customer, product, quantity, status), the product price list and ingredient stock at that time. Looking at a June price confirms the detail without moving today: no rewind, nothing written, ever.

v61 - Cloud backups (8 Sep 2026)

Every time the app opens while signed in, it quietly saves a full **Daily** copy (first open each day), a **Weekly** copy (Mondays) and a **Monthly** copy (the 1st) to your own Supabase cloud, keeping the newest 7 daily / 4 weekly / 3 monthly. A "Back up to cloud now" button, and one "Before restore" copy saved before every restore, stay until you delete them. Each copy can be **Restored** (steps the phone - and the shared cloud - back to that copy; a "Before restore" copy is always saved first, so it is never one-way), **Downloaded** as jienluv2bake-backup-YYYY-MM-DD-KIND.json (which re-imports like an Export), or **Deleted**. Lives under More -> Settings -> "Backup & safety". The copies never contain the per-phone app-login email/password or the app password. Setup: run supabase/backups.sql once.

v60 - Homepage customer reviews (8 Sep 2026)

Customers can leave a **What customers say** review on the homepage - name, 1-5 stars, a message in English / Mandarin / Bahasa Malaysia, and an optional photo. Every review lands *unpublished* in More -> Reviews, where **Publish** shows it on the homepage, **Take down** hides it, and **Delete** removes it - nothing becomes public until you tap Publish. Reviews appear on the homepage only. Setup: run supabase/reviews.sql once.

v59 - Occasion import tidy-up (8 Sep 2026)

Added **World Animal Day** (4 Oct) to the fun/pet-day import, plus whole-group ticking - tick a heading box, or use **Untick all** / **Tick all**.

v58 - Occasion import (8 Sep 2026)

A one-tap + **Add occasion** import of Malaysia's days plus fun days (pet, baking, kindness), and a **My own day** quick-add.

v57 - Private notes per product ingredient (8 Sep 2026)

Each ingredient line in a product can carry a short private description *for that product only* (e.g. which flour that product uses). Typed and seen on the Products screens only - never on the shop, a label, or the product cards.

v56 - "Keep at least" reserve (8 Sep 2026)

Ingredients can set a **Keep at least** level. The order list tops an ingredient back up to that level in whole packs when your stock drops more than 10% under it, and low items ride along even when today's bakes don't use them. The Ingredients screen shows a red low-stock strip. Guide synced to v56.

v55 - Time & length units (7 Sep 2026)

min/hr and **cm/m** are pre-loaded in Units (under More), like g/kg - so bake times and pan sizes can be measured and converted (1 hr = 60 min, 1 m = 100 cm). Baking stock is guarded so a time or length unit can never distort a gram count.

v54 - Ingredients on hand (7 Sep 2026)

Ingredients now have **On hand** stock. The PO buys only what you don't have (rows you already have drop out), "Bought" on a saved list adds its packs once, and marking an order **Baked** subtracts its ingredients (undo restores them).

This history covers the two-phone cloud era (v54+). Earlier versions (pre-v54) were never recorded version by version, so they are not listed here rather than invented. Each new version is added here as it ships.