BAKESTER · HOME BAKERY · PENANG

Jienluv2bake - change history

Every version of the backoffice app, from v54 to v438 - newest first. The green 'Engine v##' pill on the More screen tells you which build a phone is running.

About this history: it honestly begins at version 54, when the two-phone shared-cloud era started (2 September 2026). The app grew through earlier versions before that, but they were never written down version by version, so they are left out rather than invented.

What changed in each version of the backoffice app, newest first. Each version number is the "Engine" you can see on the app's **More** screen, so you can always tell which build a phone is running.

12 Oct 2026 - engine v438, A TAP ON THE MAP CANNOT MOVE THE PIN - AND THE FOURTH NUMBER IS THE SAME BALL (no database step).

Your words: *"the same problem clicking on a point of a map, the pin jump to that point, customer might just want to drag the map, or zoom. But the original pin had been moved unintentionally"* - and then *"the number 4 should not be inside the the Place order button, it should sit beside, same size with other number to be consistant, make the place order button smaller to make roon for the 4 ball"*

1. AND YOU HAD ALREADY TOLD ME THIS ONCE - ABOUT THE OTHER MAP. Your words at v264, about the map in your own backoffice: ***"click on the map should not move the pin, only dragging the pin will."*** **It was fixed there, and the shop's map never got it. The same person, the same fault, the same fix - and only one of the two maps was ever told.** That is the fault worth naming here, more than the tap.

2. SO A TAP CANNOT MOVE A PIN. It only does one thing now: on a map with NOTHING on it yet, it drops the first pin - because there is no pin to move, and dragging a small dot with a thumb is fiddly. **From that moment a tap does nothing at all**, and the pin moves only by dragging it, which nobody does by accident.

3. AND v264's OWN REASONING IS THE WHOLE OF IT: *"every touch on this box put the pin wherever the finger landed - **two answers to one accidental contact.**"* **A phone is where this bites: a pan begins with a finger going down, and a sloppy one lands as a tap.**

4. I CHECKED EVERY MAP IN THE APP, NOT JUST THIS ONE. There are two, and all three of their tap bindings now carry the guard - the shop's was the only one that did not.

5. AND THE FOURTH NUMBER IS THE SAME BALL AS THE OTHER THREE. You are right that it was not consistent: **it was drawn inside the Place order button** as a small white disc. It now **stands beside the button**, at the same 30px as , in the same place in the sequence - after the details, before the press. **And it dims with the button when the basket is empty**, so the two never show different states of one thing.

6. AND THE BUTTON IS NARROWER FOR IT, which is the room you asked to be made - 213px at 375, down from about 245, with its label still on one line. On a 320px phone in Bahasa the label wraps to two lines, and the bar grows by 14px - **the space under it is measured from the bar itself, so nothing hides behind it.**

Measured: tests **3,320 -> 3,322**, and **verified live** - a first tap placed a pin, a second tap left it exactly where it was. **No database step.**

12 Oct 2026 - engine v437, THE MAP DOES NOT OPEN ITSELF, AND THE PIN IS AN OFFER RATHER THAN A STEP (no database step).

Your words, with a picture of the map sitting over Place order: *"see how the map overlap the place order button? We should not let the map drop drown automatically, it should upon clicking only. Pin a map can be confusion that is not mandatory to be faced by customer, especially when we have auto complete. Pin helps on certain scenarios that we understand, maybe we just fade that function"*

1. YOU NAMED THE CAUSE EXACTLY: THE MAP WAS OPENING ITSELF. When the address lookup found

your door, **tapping that row unfolded the whole map underneath** - which pushed the page down and put it over Place order. You caught the same thing once before with the privacy notice.

2. SO IT NOW WAITS FOR THE PRESS. Nothing opens it but the customer. Two gains: **it can never cover Place order again**, and a customer who never adjusts the pin **never downloads the map code or its tiles at all** - which is most of them.

3. AND THE PIN IS STILL SET when he taps one of the lookup's rows. Only the map stopped appearing unbidden; the line under the button still says **"Pin set for this address"**, so nothing he chose is lost.

4. AND THE PIN IS AN OFFER NOW, NOT A STEP. It used to be **a whole paragraph and a full-width button at the same weight as the address field above it**, so it read as the next thing to do. Closed, it is now **one quiet line**:

Door hard to find? Pin on the map

5. AND THE EXPLANATION MOVED INSIDE THE MAP - the sentences that say **skipping is perfectly fine**, and when pinning actually helps (a condo block, a guard house). **That is where he is moving the pin and where it can be read properly**, instead of sitting in front of someone who has already given us his address.

6. AND THE LINE GOES AWAY WHILE THE MAP IS OPEN. The press does not close the map, so a line still offering it while it is open would be **a control that does nothing at all.** It comes back the moment the map shuts.

Measured: tests **3,320 -> 3,322**. **No database step.**

12 Oct 2026 - engine v436, THE SHOP ASKS BEFORE IT CHANGES HOW YOU GET YOUR ORDER (no database step).

Your question: *"if customer select a collection point, and also keying in her address in courier delivery section, how the software will handle?"* - then *"as both also can be collapsed"* - then *"a confirmation pop up, asking whether customer intend to change to delivery? the point selection will be cleared, vise versa"*.

1. FIRST, A CORRECTION I OWE YOU, BECAUSE I TOLD YOU SOMETHING THAT IS NOT TRUE. I said a courier order left holding a collection point would **send a driver to the point's pin**. **It does not.** The app checks the way an order is going **the moment the order arrives**, and a delivery order is stored with no point at all. So the run, the label and the customer's own tracking page were never wrong - **I read the part that uses the point without reading the part that stores it.** Your shop was not damaged.

2. BUT YOUR QUESTION FOUND SOMETHING REAL, AND IT IS ABOUT THE CUSTOMER, NOT THE DATA. If he picks a point and then switches to **Courier delivery**, the page used to leave **the point still selected**, with nothing saying whether it still applied to his order. **That is the same fault you caught once before: the page answering its own question.**

3. SO NOW IT ASKS. Changing how he gets his order - either way - **asks first**, and the question names what is about to be lost:

> **Change to delivery?** > You chose to collect from **Bayan Lepas, Sg Ara, Desa Ria, Lengkok Kenari.** Changing to delivery clears that, and we bring your order to your address instead. > **Keep Self collect** · **Yes, change**

4. AND NOTHING MOVES UNTIL HE ANSWERS. "Keep" leaves everything exactly as it was - his place, his way, and his third step circle. **The circle is only greened when the change is actually made**, which is the rule this picker has followed since v415: *the page does not credit him with a choice he did not make.*

5. IT ASKS ONLY WHEN SOMETHING WOULD ACTUALLY BE CLEARED. If he has chosen nothing, tapping the other way just switches - no question. A question that guards nothing is noise on a page that already carries several things to read.

6. A TAP ON A COLLECTION POINT IS THE SAME CHANGE, SO IT ASKS TOO. Tapping a place **is** choosing to collect, so it takes the same road - and **the place is only chosen once he answers**. Choosing it first would have left a customer who pressed "keep delivery" holding a place on a delivery order.

7. AND TWO SMALL THINGS FOUND BY LOOKING AT THE SCREEN. The point's own name ends with a full stop, so quoting it inside the question printed **"...Kenari.. Changing to delivery"** - the stop is trimmed **where it is quoted**, never in your record. And the two buttons came out one label wrapped across two lines, so they now stack instead of squeezing.

8. AND ONE WORD OF YOURS, TAKEN: *"not change to collecting but change to self colect"* - **"Self collect" is what your shop calls it**, on the button he taps and everywhere else, so the question uses your word rather than a third name for one thing.

Measured: tests **3,314 -> 3,320**. **No database step.**

12 Oct 2026 - engine v435, FOUR STEPS NOT THREE, AND THE PAGE FINALLY SAYS WHAT A BAKE DAY IS (no database step).

Your words: *"it is actually a 4 steps ordering process. not 3. the 4th is place order."* - and then the one that matters most: *"I think customer 1st confusion is that we never explain what is a bake day to them."*

1. IT IS FOUR STEPS, AND THE FOURTH WORE NO NUMBER. The page numbered **1, 2, 3** and then **Place order** sat in the bar with **no number at all** - so the line above the steps promised three while the customer did four things. It now reads **"You'll do the 4 steps, ordering 1, 2, 3, 4 - and then we take over"**, and **the Place order button wears a 4**, so the promise and what he actually does agree.

2. AND THE REAL ONE: WE NEVER TOLD A CUSTOMER WHAT A BAKE DAY IS. You said it yourself. **The words were used FIVE TIMES on that page and explained none of them** - the rules card, the step heading, the line under the calendar, the Self collect blurb. **And the only hint about when the bread is ready was buried inside the Self collect option, so a courier customer never met it at all.**

3. SO THE PAGE NOW SAYS IT, UNDER "1 PICK A BAKE DAY":

> **A bake day is the day we bake.** Your bread is made that morning, so it is at its freshest that day - and **the earliest it can be ready for you is the afternoon**, not the morning.

It sits exactly where the question is asked - he is choosing a day, and that is the moment he wonders what one is. **Every customer sees it, courier or collect.**

4. AND ONE WORD I CHANGED, ON PURPOSE. You said "late afternoon or afternoon", and I wrote "the afternoon". The **late** afternoon is the window we collect at a **Point** - the bread is ready earlier than that, and promising that window here would promise less than the bakery can keep. Say the word and I will put "late afternoon" back.

5. THE THREE LANGUAGES ALL CHANGED TOGETHER, and the Chinese and Bahasa are my own wording rather than a native speaker's - still owed the same check as the rest of this session's shop sentences.

Measured: tests **3,311 -> 3,314**. **No database step.**

12 Oct 2026 - engine v434, THE SHOP'S THREE STEPS: ONE COLOUR, AND THE PAGE SAYS WHICH DAY IT PICKED FOR HIM (no database step).

Your words: *"the 1,2,3 in the same round ball"* - and then, better than my own answer: *"say, moonday was selectted for you, just click on the prefered date if you have one"*

You were right, and your second sentence is the fix.

1. THE THREE CIRCLES WERE NOT THE SAME, AND THE REASON MATTERED. The circles were already the same round ball - same size, same shape. **The first one was a different COLOUR**, because the shop picks the next bake day **for** the customer so an order is never lost by forgetting to tap - and step 1 counted that as a step **he** had done. On a fresh page, before he had touched anything, circle 1 was green.

2. YOUR OWN RULE, AND THE CODE ALREADY SAID IT. Step 3 was built the other way on purpose, and its comment reads: **"the page does not credit him with a choice he did not make."** Circle 1 was breaking exactly that rule.

3. SO ALL THREE NOW ASK ONE QUESTION EACH: has HE done this? On load they carry a single colour. Each turns green as he does it - and that is the whole point of the colour, from the day it was built.

4. AND THEN YOU SAID THE BETTER THING: TELL HIM. Under the calendar the page now says:

"Mon, 12 Oct was selected for you - just click the date you prefer, if you have one."

Because a page that does something on his behalf must **say so**, rather than let a highlight imply he did it. The day he picks himself reads plainly - **"Your bake day: Wed, 14 Oct"** - and the circle turns green at the same moment.

5. AND THE LINE ABOVE THE STEPS, IN YOUR WORDS:

"You'll do the 3 steps, ordering 1, 2, 3 - and then we take over."

6. THE THREE LANGUAGES ALL CHANGED TOGETHER, and the Chinese and Bahasa are my own wording rather than a native speaker's - counted as still owed, the same as the rest of this session's shop sentences.

7. AND THE DAY IS NOW SPELLED OUT. You asked for **"Your chosen bake day: Monday, 12 Oct"** - so the label changed, and every date in the shop now reads in full rather than as "Mon": **"The next day you can have bread is Monday, 12 Oct."**, "Order before 6pm on Sunday, 11 Oct", "Next available: Monday, 7 Sep". **A date spelled out is easier to read at a glance than a shorthand.**

8. ONE PLACE DELIBERATELY KEPT SHORT, AND TELL ME IF YOU DISAGREE. The **"Bake days: Mon, Wed, Fri"** card is still abbreviated. That card is a standing list of the days you bake, not one date - and **"Bake days: Monday, Wednesday, Friday"** is a mouthful on a phone, where it would push onto a second line. It is also the card you picked off the sheet an hour earlier. **Say the word and it becomes full too.**

9. AND THE COLLECTION AREA NAMES ARE BLACK NOW - AND NO LONGER SHOUTED. You spotted **"PENANG ISLAND"** and **"PRAI"** in the brand red. On this shop that colour is the voice for things a customer has to **notice** - an instruction to tap, a day that is closed. **An area name asks for nothing; it is a signpost saying where the places under it are.** In black it still leads its group, without competing with the places themselves.

And then you asked why it was capitalised. The capitals were never in your data: the area is stored as "Penang Island" and the page was drawing it in capitals. So it now reads the way you write it - **"Penang Island", "Prai"** - and the small size, the weight and the wide spacing are what keep it a quiet heading, none of which needed capitals to work.

10. AND YOUR SENTENCE REPLACED THE OLD INSTRUCTION UNDER IT: "Tap to choose which Self collection Point you want to collect your orders from." It is **still in the red**, because you named the area for the colour and gave the instruction new words - say the word if you want that line black too.

11. AND "SELF COLLECT" IS A TITLE NOW. You said it should be **"bigger font being a title"**. The answer was not a new size invented for it: every other heading on your shop is the same one - **16px, bold, in the serif** - so it now wears exactly that, and reads like the product names above it. **"Courier delivery" gets it too**, being the other half of the same choice.

Measured: tests **3,307 -> 3,309**. **No database step.**

12 Oct 2026 - engine v433, THE SHOP'S "BAKE DAYS / ORDER BY" CARD NOW READS AS TWO SENTENCES (no database step).

Your words, with a picture of the card: *"instead of a line, customer might confused that it is 2 group of word, one left one right"*

You were right, and it was one line of styling.

1. THE WORD AND THE ANSWER WERE PINNED TO OPPOSITE EDGES. The row was built to push its label to the far left and its answer to the far right with **nothing joining them** - so a customer could read straight down the left, then straight down the right, and pair them up wrongly. **And this is the one card on the shop that states the two rules an order depends on.**

2. YOU PICKED THE SHAPE OFF A DRAWN SHEET: one line, one sentence. It now reads:

Bake days: Mon, Wed, Fri Order by: 6pm the day before

3. AND IT IS NOW IMPOSSIBLE TO SPLIT IN TWO, because there is only one group of words on each line.

4. THE THREE LANGUAGES ARE UNTOUCHED. The colon goes in the page itself, **not** inside the translated words - so the Chinese and Bahasa read exactly as they did, and none of them has to carry punctuation it never had.

5. WHERE IT IS: the shop page only. I checked - this card is on /store/ and nowhere else, so nothing else in the shop moved.

Measured: tests **3,303 -> 3,307**. **No database step.**

12 Oct 2026 - engine v432, AN ORDER MOVES AS ONE - ALL OF IT, NOT JUST THE LINES ON THE DAY YOU ARE LOOKING AT (no database step).

Your words: *"why a new order already change to confirm yet appear in new order at the top of order page?"*

You were right, and I found it - then I built it and watched it go wrong, then fixed it.

1. AN ORDER CAN HAVE ITS LINES ON TWO DIFFERENT BAKE DAYS. When that happens, the Orders screen for a day only shows - and only *changes* - **the lines on that day**. So when you set it to Confirmed, the lines you could see changed, and a line on another day was **left at New**.

2. AND THE LIST AT THE TOP SCANS EVERY DAY. So it went on finding that leftover line - and because the lines belong to the same order they carry **the same order code**, which is exactly why it looked like the order you had just confirmed. **The list was right about what it could see; the status change was writing only half the order.**

3. THE FIX: THE WHOLE ORDER MOVES. Change an order's status and every line of it moves with it, wherever those lines sit. **Nothing of yours is rewritten** - no saved order is touched. An order already split stops causing trouble the moment you touch it.

4. AND I HAVE TO BE HONEST ABOUT ONE THING: I DO NOT KNOW HOW YOUR ORDER CAME TO BE SPLIT ACROSS TWO DAYS. I looked for it and could not find it. If you remember moving an order's bake day at some point, that is very likely the clue - **tell me and I will chase it.**

Measured: tests **3,298 -> 3,303**. **No database step.**

11 Oct 2026 - engine v431, THE EDIT CARD NOW LETS YOU REACH A PAST BAKE DAY (no database step).

Your words: *"when i edit an order should it allow me to select the past date,?"*

You asked a question and the answer was "it should, and it did not".

1. THE EDIT CARD HID EVERY BAKE DAY THAT HAD GONE BY. Open an order with **Edit** and the "Bake day" calendar was handed a list with **every day before today quietly removed** - so one of your own bake days came out as a **dead grey square, the same as a day you never bake**, and you could not move an order onto it. The only past day it kept was the one the order was already sitting on.

2. WHY THAT MATTERED. You take an order by hand on Friday and enter it on Sunday. The card offered only **future** days - so **Friday's takings read empty and Monday read busy**, and the ingredients came off the wrong bake.

3. AND THE REST OF THE APP NEVER DID THIS. The **Orders screen's own calendar** and the **New order** card have always shown every day you have set, and a day already gone opens there on purpose - the code says it plainly: *"she backfills and reviews old days"*. **This one picker was the odd one out**, and it is the one you asked about.

4. YOU SAID YES, AND NOTHING IS BLOCKED. Every bake day you have set now appears in that calendar, past or future, **exactly as it does on the Orders screen**. A day you **do not** bake is still a plain grey square - this did not turn the whole month into buttons. **And it sat against your own rule: no rule may stand in the way of a sale you would take by hand.**

5. AND PICKING A PAST DAY STILL WARNS YOU. The app's own sentence - **"That bake day is already**

past." - now appears under the calendar when you pick a day that has gone. You are told, and you may still do it, the same as everywhere else. **That sentence was already written; until now it could only ever appear on an order that was already sitting on an old day, never on one you had just picked.**

Measured: tests **3,294 -> 3,297**. **No database step.**

STILL OUTSTANDING FROM v429, AND IT DOES NOT SHIP BY PUSHING: the order-alert SQL (supabase/order_alerts.sql) must be **run again in Supabase**, or your phone alert keeps the old word.

11 Oct 2026 - engine v430, "PAID" ON A CUSTOMER'S CARD NOW MEANS THE MONEY IS ACTUALLY IN (no database step).

Your words: *"look like there is a discrepancy for Peggy's order, at Peggy's customer card, has that her order is paid, but it is not??"* - and then: ***"pls check any other discrepancy"***

You were right. I built Peggy's exact case and the card said "Paid". The test that catches it fails on that very sentence.

1. THE APP KEEPS TWO FACTS APART, AND ONE SCREEN WAS READING ONLY ONE OF THEM. "Paid" is the stage the order has reached - it is what the dropdown moves. **Whether the money arrived is a separate record**, and it is only written when you press **Paid · Cash** or **Paid · TNG**. Peggy's card was printing the stage word by itself, so it said **Paid** over money nobody had handed over.

2. EVERY OTHER SCREEN ALREADY READ BOTH. The Orders list's Cash / TNG tag, the day's till, the invoice, the receipt and the customer's own **Track your order** page all read the money record as well as the stage. That is why only this one card disagreed with everything else - and it is why nothing else needed changing.

3. I CHECKED FOR OTHERS, AS YOU ASKED. THERE WAS ONE MORE, AND I LEFT IT ALONE. The **View** on a backup copy also prints the stage word - but that screen is a list of **what a saved copy contains**, so there the stage is the right thing to say. Everything else was already right.

4. THE SHAPE YOU PICKED, FROM THE SHEET. The stage chip stays exactly where it was, and the money gets a chip of its own beside it:

- **Not paid** - amber - while the money is still out (the stage reached paying, or went past it) **Cash** or **TNG** - once it is in, the same words the Orders list already uses **Refunded** - if the money came in and went back out

SO A ROW NOW READS #CAFE02 · Self collect · Delivered · Not paid - the one you would chase: the bread has gone, the money has not come.

5. AND NOTHING OF YOURS MOVED. An order from before this record existed carries no flag at all, and it reads exactly as it always did - no chip appears on it. An order that has not reached the paying stage gets no chip either. Only the orders that really are owed something changed.

Measured: tests **3,285 -> 3,294**. **No database step.**

STILL OUTSTANDING FROM v429, AND IT DOES NOT SHIP BY PUSHING: the order-alert SQL (supabase/order_alerts.sql) must be **run again in Supabase**, or your phone alert keeps the old word.

10 Oct 2026 - engine v429, THE BAKE DAY IS NOT THE DELIVERY DAY - IN SEVEN PLACES, NOT ONE (SQL STEP - the order-alert function must be run again, then push).

Your words: *"when i entered a specific delivery date favour by customer over bake date, the print label still took the bake day?"* - and then, when I started looking: ***"pls check throughoutly"***

You were right, and it was seven places. I checked every surface an order's date reaches a person.

1. THE LABEL NEVER READ THE DELIVERY DATE AT ALL. It built its date from the bake day and nothing else - in **every** style that prints one. And the **mailing** label - the one that goes **on the outside of the parcel** - literally read **"Deliver ..."** followed by the bake day, so a box baked Friday and sent Saturday was addressed **"Deliver Fri 4 Sep"**.

AND YOU CHOSE THE SHAPE: BOTH, EACH NAMED. A van order with its own day now reads **Bake Fri 4**

Sep · Van Sat 5 Sep, so whoever packs a bag sees **when to make it and when it leaves**. **With no delivery day typed the label is byte-for-byte what it always was** - a day you have not decided cannot be named. And a **parcel** label's "Deliver" is the day it **arrives**.

2. AND ONE OF THE SEVEN WAS A CUSTOMER STANDING OUTSIDE A CLOSED SHOP. The **pickup reminder** told a customer to collect **on the bake day** - but **a Place order's bread still arrives by van**, and the run writes the van's day onto it. So a run that reached the shop the morning after the bake told the customer to come **on the wrong day**. That one is fixed; the rest are a word on a screen or a label.

3. THE OTHER FIVE. The **customer history** row said *"deliver"* followed by the bake day. The **CSV column headed "Last delivery"** carried the bake day. The **Track your order** card printed the bake day **unlabelled, immediately before the words "Courier delivery"** - which is the exact shape that caused the original *"is it Wednesday or Thursday?"*. The order card said **"Delivering"** followed by the bake day, while a **"Courier delivery date"** box sat directly under it. And **your own phone alert** said **"Delivery:"** followed by the bake day - which is how the two came to look like one thing.

4. AND ONE OF THEM CANNOT BE FIXED WITH A DATE, WHICH IS THE PART WORTH KNOWING. The phone alert fires **the moment an order lands** - and **the van is booked afterwards**, so at that instant **there is no delivery day to name.** So that one does not gain a date; it **stops calling the bake day the delivery day** and says **"Bake day"**. **Naming a delivery day it cannot know would be worse than the fault** - the same reason the confirmation has never stated one.

5. AND WHAT WAS ALREADY RIGHT, WHICH IS MOST OF IT. The four customer messages read one seam (courier_job.js) built for this exact fault in an earlier version: a van order's line is labelled **"Baking day"** and the van's day gets a line of its own. The Delivery run, the run badge on Bake days, and the courier booking all read the van's day. **Those were not touched** - touching them would have been the regression.

Measured: the suite is 3,285 tests, all green (4 new). Five bites, all fired by their own message, every file restored byte-identically. **And the no-delivery-day case is asserted byte-for-byte**: the common order, whose van day you have not filled in yet, prints exactly as it did before.

THE SQL STEP. The alert message lives in supabase/order_alerts.sql, which runs in the database - **so this one is not shipped by pushing alone.** Run that file again in Supabase, then push.

No database step for the app itself - everything else is app code.

10 Oct 2026 - engine v428, THE RED IS HELD LONGER (no database step - pushing this one is the whole of it).

Your words: *"the red is too short time"* - and then, after the first increase: ***"still too short"***

What it was. The red was on for only **215 milliseconds** of each blink - the window was 43% of a half-second - which is **a flicker rather than a warning. You cannot read a word in a fifth of a second**, so the flash was drawing your eye to the line and taking the line away before you could use it.

AND YOUR SECOND TELLING IS THE ONE THAT TAUGHT ME THE SIZE OF IT. I raised it to 0.9 seconds a blink - nearly three times as long - and **you said still too short. So the eye wants more than half a second to take a line in AND use it.**

AND THEN YOU GAVE ME THE TIMING ITSELF: *"on 1sec, off 0.5sec cycle"* - so that is exactly what it does. The red is held for **one second**, then the line goes pale for **half a second**, and it repeats three times. **Measured on your own shop, sampling one cycle:** red at 0.2s and at 0.9s, pale at 1.1s and 1.4s, and red again at 1.6s - **her cycle, to the millisecond. And the test now carries YOUR figures rather than a floor I chose** - a full second of red, and an OFF stretch of half a second. It also pins that the red **ends before the pale begins**: the two must not overlap, and the thin crossfade between them is deliberate, because a truly square edge at this size reads as a rendering glitch rather than a light going on and off. **And the cycle OPENS on the red** - "on for a second" means it is already on when it starts, and a fade-up would spend part of your second arriving.

AND IT WAS A ONE-LINE CHANGE, WHICH IS THE POINT. Because v427 made the screen **ask the**

stylesheet how long the flash is, the timer followed on its own - **no number to update in a second place, and nothing to drift.**

Measured: the suite is 3,281 tests, all green. Three bites, all fired by their own message, the stylesheet restored byte-identically. **Verified on your own shop page**, sampling a single blink at five points: the line is **red from just inside the start to just before the end, and pale on either side of it.**

AND ONE THING YOU SHOULD KNOW: this version is not on your shop yet. I checked what is deployed, and **v427 is what is live there** - 500ms with the narrow window, which is the very flash you called too short. **So when you looked and it was still too short, you were looking at the one you had already told me about.** The change is in the files, waiting to be pushed.

No database step. Pushing this one is the whole of it.

10 Oct 2026 - engine v427, THE WARNING BLINKS RED THREE TIMES (no database step - pushing this one is the whole of it).

Your words: *"can you flash the red twice or trice"*

It blinks three times now, half a second each - a second and a half in all.

AND ASKING FOR MORE BLINKS EXPOSED A COUPLING WORTH KNOWING ABOUT. How long the flashing runs lived in the stylesheet, and the moment the mark comes off lived in the screen's own timer - two numbers, in two files, describing one fact. And they could only ever disagree in ONE direction, because **the timer is the thing that fires while the animation is switched off** - which is exactly what a reader who has asked for less movement gets. **Three blinks is 1500ms and the timer said 1200** - so asking for three would have **cut the flash off in the middle of the second one**, and it would have looked like the app ignoring what you asked for.

So the screen now ASKS the stylesheet how long the flash is rather than carrying a number of its own, and a test pins that it asks - the two numbers can no longer drift apart. And the arithmetic allows for the **last** blink finishing rather than being clipped.

Measured: the suite is 3,281 tests, all green (1 new). Three bites, all fired by their own message, both files restored byte-identically. **Verified on your own shop page:** the animation reports **three iterations of half a second**, and the line is red when caught mid-blink.

No database step. Pushing this one is the whole of it.

10 Oct 2026 - engine v426, THE WARNING FLASHES RED - AND TWO WORDS FROM YOU FOUND A NUMBER I HAD GOT WRONG (no database step - pushing this one is the whole of it).

Your words: *"can you flash the only available on friday in red color?"* - and then, reading it back: ***"it is abit too fade"***

1. THE WARNING NOW FLASHES RED. The shop had no red at all before this - and that is why it needed its own rather than a borrowed one: **the whole page is terracotta**, so a "red" taken from the palette would have flashed and read as the **brand colour arriving**, which is the opposite of standing out. It has its own red, and it clears the app's own bar at **5.35:1** - a flash nobody can read is not a warning.

2. AND YOUR SECOND SENTENCE FOUND A NUMBER I HAD WRONG - IN A VERSION ALREADY PUSHED. v425 faded the warning line by 70% **on top of a colour that is already the faded tone.** So the warning was at **2.94:1**, not the **5.21** I wrote down. **I measured the colour and then folded a second fade over it, and the two compound. Both the changelog and the guide said 5.21, and both were wrong.** The line now carries **no fade at all** - --muted at full strength is **5.40:1**, which was the right answer all along. A dated correction is written **inside the v425 entry**, so the wrong number does not sit in your history looking authoritative. **The name and the ordinary description still fade hard**, exactly as you chose - that part was right.

3. AND ONE THING DID NOT WORK, WRITTEN DOWN RATHER THAN HIDDEN. The first red animated **var(--muted) to var(--red)** - and **the red never arrived**: the animation ran, its name was on the element, and

the colour did not move an inch. The two colours are written out as plain values now, and a test ties them back to the tokens so they cannot drift apart. **A var()-to-var() pair is not something every engine will blend.**

Measured: the suite is 3,280 tests, all green. Three bites, all fired by their own message, the stylesheet restored byte-identically. **Verified on your OWN shop page**, where a Friday-only place is already parked: the line reads **Only available on Fri**, it turns **red** while it flashes, it comes back, **the page does not move**, and the line sits at full strength instead of the double fade.

No database step. Pushing this one is the whole of it.

10 Oct 2026 - engine v425, A PLACE YOU CANNOT PICK IS FADED TO THE EYE - AND THE ONE LINE THAT SAYS WHEN IT IS OPEN STAYS READABLE (no database step - pushing this one is the whole of it).

Your words: *"the unavailable point grey out is not the background grey out but wording grey out"* - and then, when I had done exactly that: ***"the grey not obvious"***

AND YOU WERE RIGHT THE SECOND TIME TOO. My first cut greyed every line to the app's own muted tone. **That tone is a WARM BROWN**, and at that size beside the body copy it reads as ordinary text - so the row did not look unavailable at all. **So I made you a sheet of four treatments and you picked one** - a colour is not something to settle in a sentence, and the sheet is still there if you want to look again.

An unavailable place now FADES, by different amounts for different lines:

- **The name and the description go to 42%** - measured against the card, that is **2.40:1**: unmistakably pale. They are not the thing the customer has to read. **The line carrying the reason goes to 70%** - **5.21:1**, against normal text's 13.45:1. Clearly paler than an ordinary row, **and still comfortably readable** - because **a reason nobody can read is not a reason.**

CORRECTION, added with v426: THAT "5.21" WAS WRONG, AND THE 70% WAS WRONG WITH IT. That line's own colour is **already --muted**, the faded tone - so folding a further 70% over it gave an effective **2.94:1**, not 5.21. **I measured the colour and then applied a second fade on top of it, and the two compound.** The line now carries **no fade at all**: --muted at full strength is **5.40:1**, which is what this should always have said. **You caught it by reading the page** - *"it is abit too fade"*. The name and the ordinary description still fade hard, exactly as you chose.

AND THE REASON IS NOT THE SAME LINE IN BOTH CASES, which is the fiddly part: a place not served today says it on its own line, while a place his basket has not reached says it on the **description** line, because that sentence replaced the description. Both keep the readable fade, and the rule that fades a description is written to leave the reason line alone.

AND NO BACKGROUND FILL. It filled with the **page** cream, so inside the card's own tint the row became a flat patch of a different cream - **the same "hole" you caught at v413**, in a place I had not looked at yet. The row keeps the card's surface and is set apart by its words and its dashed frame: **the dashes are the shape and the fade is the tone**, and either one alone would be a cue made of colour only. **Both parked reasons changed together** - a place not served that day, and one his basket has not reached - because two rows that look alike must behave alike.

Measured: the suite is 3,280 tests, all green (1 new). Six bites, all fired by their own message, the stylesheet restored byte-identically. **Verified by looking** at 375px on the real shop page: Farlim's name and description are plainly pale, **Only available on Wed and Fri** is plainly readable, the frame is dashed, and there is no fill.

No database step. Pushing this one is the whole of it.

10 Oct 2026 - engine v424, THE WORDS FLASH AND NOT THE ROW - AND A STUCK MARK I SHIPPED YESTERDAY (no database step - pushing this one is the whole of it).

Your words: *"i mean just flash the word Only available on Wed, Friday"*

1. THE WORDS, NOT THE ROW - and your "just" was the instruction. v423 pulsed the whole row. You meant **the line**: the row is already greyed and dashed to say *not this one*, and what a customer needs told is

when it IS open. So **Only available on Wed and Fri** now blinks - dims and comes back, once, in about a second - and nothing else on the page changes at all.

ONLY THE LINE THAT CARRIES THE REASON, and the two parked reasons carry it in different places: a place not served today says it on its own line, while one his basket has not reached says it on the **description** line, because that sentence replaced the description. Blinking the description for *both* would flash a sentence that says nothing about why he cannot pick it - the opposite of pointing at the warning.

2. AND IT CAUGHT A FAULT I HAD SHIPPED IN v423 - and this is the part worth your attention. v423 took the mark off the row on the **animation's own end event**. But for a reader whose device asks for **less movement**, the stylesheet switches the animation **off** - **and an animation that never runs never ends.** So the ring would have stayed on that place **for good**, and tapping a second one would have left the first one lit as well. **The reader it would have hit is the one who asked for less movement, which is exactly where nobody thinks to look.** The mark now comes off on a **timer**, which always fires.

It was found by reading the rule back, not by anyone hitting it - and v423 was already pushed, so the fault was live for a few hours on a device set that way. The reader who asked for less movement still gets the warning drawn to his eye: the line simply **holds bold** instead of blinking.

Measured: the suite is 3,279 tests, all green (2 new). Three bites, all fired by their own message, every file restored byte-identically. **Verified live at 375px:** the row reads **Only available on Wed and Fri**, that line's opacity goes **1 down to 0.2** while nothing else moves, **the scroll position and the row's height are both identical**, and the mark is gone once the timer finishes.

No database step. Pushing this one is the whole of it.

10 Oct 2026 - engine v423, A PLACE HE CANNOT PICK NOW FLASHES INSTEAD OF MOVING THE PAGE (no database step - pushing this one is the whole of it).

Your words: *"when customer click on unavailable point, just flash the warning Only available on, xxx. The screen jump is not desired"*

What it was. Tapping a greyed place opened the message box further up the page - **and that box scrolls itself into the middle of the screen every single time it is shown.** So the page **leapt away from the list he was reading**, at the exact moment he was reading it. It was never this screen's decision: the scroll is inside the box, and it is right for the **order guards**, where he has pressed Place order and has to be told why not.

Now the row itself pulses for about a second, and nothing moves. And nothing new is said, because the warning is already on the row - **Only available on Fri** for a place not served that day, and the smallest-basket sentence for one his basket has not reached. The tap draws his eye to the line he already had, which is the whole of what he needs. **The box is still there for the order guards, untouched.** It is simply no longer used for a place he merely touched.

THE PULSE IS A SHADOW, NOT A THICKER BORDER, and that is the whole craft of it. A border that thickened would grow the row by a couple of pixels for a second and nudge everything under it - **a smaller version of the very jump this replaces.** Measured live at 375px: the scroll position is **identical** before and after the tap, and the row's height does not change by a single pixel.

AND BOTH PARKED REASONS FLASH - a place not served today, and one his basket has not reached - because **two rows that look alike must behave alike**, and only one of them moving the page would be the fault this app treats as a bug. And for a reader whose device asks for less movement, the pulse becomes a still outline rather than nothing at all.

Measured: the suite is 3,277 tests, all green (2 new). Three bites, all fired by their own message, every file restored byte-identically. **AND PUTTING THIS IN BROKE A TEST THAT HAD NOTHING TO DO WITH IT**: the offer strip's reduce-motion check found its block **by position**, so a second reduce-motion block moved its subject out from under it - it finds it **by content** from today. **A test that finds its subject by position breaks when an unrelated rule is appended.**

No database step. Pushing this one is the whole of it.

10 Oct 2026 - engine v422, A WHOLE ROUTE IS TAKEN IN ONE PRESS (no database step - pushing this one is the whole of it).

Your words: *"my question if the day should call 3 route and 2 doors, how to go about?"*

And the honest answer to it was: badly. Three routes on one day meant ticking stops by hand for **each van**, and the only bulk press there was - **Untick them all** - is for the **whole day**, not for one route. So a three-van day was three rounds of careful ticking, on the screen where you spend the money.

So every route heading now carries its own press. It reads **ROUTE A · 2 STOPS** with **Take this route** on the right, and the press ticks everything on that route. **It turns over to "Untick this route" the moment the whole route is on**, so it always says what it will do next rather than what you did last - the same spelling the day's own press already uses.

AND IT TAKES ONLY ITS OWN ROUTE. A stop you TUGGED into two routes is drawn under its **first** one, and it is taken with that one - **the block IS the route**, so the press and the block under it can never disagree about what a route holds. And the prices go back to unasked when the tick list moves, exactly as one row's own tick makes them: **a price for a list you have since changed describes a journey you are no longer taking.**

This is that day, end to end - three headings, each with its own press, and two doors of their own:

```
Fri, 16 Oct - 6 stops

Farid's house          x2      <- a doorstep, on no route
Gina's house           x1      <- a doorstep, on no route

ROUTE A - 2 STOPS      [ Take this route ]
  Sg Ara
  Farlim, Air Itam

ROUTE B - 1 STOP       [ Take this route ]
  Chai Leng Park, Prai

ROUTE C - 1 STOP       [ Take this route ]
  Bukit Mertajam
```

So a van is now one press, three times over - press Take this route, price it, book it, and go on to the next. The two doors stay where they are, at the top, until you tug them onto a route.

Measured: the suite is 3,276 tests, all green (1 new). Three bites - the press being there at all, it taking only its own route, and it turning over - **all fired by their own message**, every file restored byte-identically. And putting the press on a heading **broke two older tests**, which is the check doing its job: a heading's textContent had always been its own words, and it now contains a button's too. They read the heading's **label** from today.

No database step. Pushing this one is the whole of it.

10 Oct 2026 - engine v421, TUG A CUSTOMER'S DOORSTEP INTO A ROUTE, AND ONE WRONG CHINESE SENTENCE (no database step - pushing this one is the whole of it).

Your words: *"how do i tug in any other point to a route already named for that day?"* - and then, twice, correcting what you meant: *"not any other point but any other customer delivery?"* and *"not to belong but i manually tug into a route"*

1. A CUSTOMER'S OWN DOORSTEP CAN BE TUGGED INTO A ROUTE. Until now a route held **Places only**: a customer's delivery always drew at the **top** of a run day, with no heading over it. Now every doorstep carries a small **tag** - **Route A** when it is on one, **+ Route** when it is not - and **the tag is the press**: tap it and pick the route.

AND "TUG", NOT "BELONG", IS THE WHOLE DESIGN. Your own words were *"for that day"*, and you corrected me when I wrote it up as a thing a doorstep *belongs* to. **So it is not a membership: it is an action you take on a run day**, which is exactly why the app remembers it **on the order** and **never on the customer**. A route written onto a customer's own record would be a **permanent** change made by a passing arrangement -

and nothing of yours is rewritten by a plan. A **Place** is the opposite case, which is why a Place's routes live on the Place. **One idea, two homes, and the difference is yours.**

2. AND A ROUTE CAN BE MADE THERE AND THEN, WHICH IS THE WHOLE POINT. My first cut showed the tag **only once you already had a route** - and you told me why that is a dead end: *"if there is no collection point route, i would choose to create manual route, tag doors into same manual route"*. **So the tag is on every doorstep always, and with no routes it offers to make one.** Name it, press once, and this doorstep is on it; every other door is then two presses - the tag again, and the same route. **A route made on a run day is an ordinary route**: the same list as the rest, so it can carry Places too and it is there again next week. There is no second kind of route to learn. And a doorstep whose route you have since **deleted** goes back to the top, exactly as an unrouted one does - **nothing is lost when a route goes.**

3. AND ONE MANDARIN SENTENCE WAS WRONG, AND IT WAS MINE. The promo line on your shop told a Mandarin reader, in so many words, that **the code would be "delivered" along with the order**. **A code is not delivered anywhere** - it is *attached to the order* - and it now says that instead. Found by reading every string I had written in Mandarin and Bahasa rather than trusting them.

4. AND THIS HISTORY CAN NOW PRINT CHINESE - IT COULD NOT BEFORE, AND FIVE ENTRIES HAD ALREADY LOST WORDS TO THAT. While trying to quote the Mandarin sentence I had just fixed, it came out of this document as a **gap**. The cause was one line in its builder that deleted every character outside the Latin alphabet. **It had been erasing real content for months**: **v407**'s bake-day wording, **v355**'s and **v354**'s privacy-notice wording, **v349**'s notice **TITLE** (which survived only as a bare dash), and **v323**'s bring-a-friend offer. **All five read properly now, and this line is the proof of it:** 烘焙日 · 你的资料 · 个人资料保护声明 · 朋友推荐优惠. The document's own look has not changed - every line without Chinese is set exactly as it was.

5. AND THE REST OF THE WORDING CHECK GOES TO YOUR CIRCLE, NOT TO ME. The shop uses **two different pairs of words** for the same two things - one pair for collecting, another for delivery - and Bahasa shows the same split. **I added to the mixture**, so some of that drift is mine. Both pairs are understood in Malaysia, but a page that switches between them for one thing reads as two different things - the same fault you caught with "delivery day". **Which pair a Malaysian Mandarin and a Bahasa reader would rather see is their call, not mine, and I will not guess it. The sheet I have made for them has every sentence a customer can see, in all three languages** - it is a page of your own, not part of the app.

Measured: the suite is 3,275 tests, all green. Sixteen bites across v420 and v421, all fired by their own message, every file restored byte-identically. **Verified by looking**, at 375px: the unrouted doorstep leads with **+ Route**, the tugged one sits under **ROUTE A · 1 STOP** with **Route A** , the row is not squeezed - and on a run day with **no routes at all**, both doorsteps offer **+ Route** and the press opens a pop-up that makes the route and puts the doorstep on it in one go.

No database step. Pushing this one is the whole of it.

10 Oct 2026 - engine v420, YOUR OWN PLACES COME BACK TO YOUR ORDER CARDS, AND A PLACE NOT SERVED TODAY IS GREYED (no database step - pushing this one is the whole of it).

Your words: *"the point i setup is not available for selection in +order, edit order and other"* - and *"selecting a bake day the point is not active will grey out that point"*

FIRST, THE FAULT, AND YOU WERE RIGHT - YOUR OWN PLACES WERE BEING HIDDEN FROM YOUR OWN SCREENS.

The **Collect from** picker on the **+ New order** card, the **Edit** card and the order screens only ever offered your kitchen - never the Places you had made. It has been that way since v406, and the reason is worth writing down: a rule was put in the one place that builds that picker, on the belief that it was the list a **CUSTOMER** chooses from. **It is not.** That list is built in one file, views/orders.js, and the shop builds its own from the published Places - so the rule was in the wrong place and did the exact opposite of what it was written for: it hid your Places from you, while the customer's page went on offering every one of them.

And it put a gate where you have told me never to put one: nothing may block or hide a sale you take by hand - you could not even key in an order collecting at a Place you had just opened.

What it was testing for was wrong too. The rule asked a Place for its **pin**, and refused it without one. **A pin is for both ends** - the van that has to be given coordinates, and the customer being shown where the place is - **but a Place has to work without one.** You put it plainly: *"a pin can be for both, and i will put in one later, without pin it should still function"*. And it does: a Place is offered, chooseable and orderable with no pin at all. The pin is needed only at the moment a **van** is actually booked, which is the courier's own requirement - and that is where it is enforced, on the **Delivery run**, which says *"this Point is not pinned yet"* and offers **Pin the Point** beside it. Nothing was dropped - it went back to its own home.

Verified on your own account, at 375px: the **Collect from** box on a New order now reads **My kitchen · Bukit Mertajam · Farlim, Air Itam · Georgetown · Chai Leng Park, Prai** - all four of your Places, not one of which is pinned.

SECOND, THE GREY-OUT YOU ASKED FOR. Once a customer has his bake day chosen, every Place that is not served that day is **drawn dashed and greyed**, and tapping it says why rather than doing nothing: **Chai Leng Park, Prai - Only available on Mon. Nothing is taken off the page** - the Place is still listed, still named, still under its own area. Greying is not hiding, and it is not a gate.

A choice already made falls back to your kitchen. If he picks a Place and then moves to a day that Place is not served on, the choice is released rather than kept. That is the **third** reason a Place can stop being open - paused, basket too small, and now not served that day - and all three behave the same way.

Nothing else moved. A Place with no days ticked is served every day, so it is never greyed. And it costs no new wording in any language: the reason a tapped Place gives is **the very same sentence the row already carries underneath it**.

Measured: the suite is 3,271 tests, all green. Six bites, all fired by their own message, every file restored byte-identically - one for the pin gate coming back, and one for the likeliest way to get the grey-out wrong: reading a missing day list as **"never served"** instead of **"every day"**, which would have taken every untroubled Place off your shop page in one go.

No database step. Pushing this one is the whole of it.

10 Oct 2026 - engine v419, DELIVERY ROUTES, AND THE DAYS A PLACE IS SERVED (no database step - pushing this one is the whole of it).

Your words: *"now i need few delivery routes, for one day delivery run day. so point will carry its route info, and day of the week choices, day of the week can be any day, up to 7 days, configurable. No of route configurable"*

And your four answers, every one of which changed the shape of it:

- Where the days live - *"we don control by route, we control by point"*
- Whether the shop should filter by them - *"every point listed in store, just that if that point only for friday, then mention it"*
- Whether a Place could be on several routes - *"a point might belong to 2 or more routes"*
- And what that means for its days - *"say one point that is on 2 or 3 routes, that it get more day probably"*

So there are three small things here, not one big one.

1. Delivery routes - a new screen under More, in the Logistic group. A route is one van's run round a few nearby Places. It is deliberately NOT the same thing as an area: an area is a geography a customer finds himself in (Penang Island, Sg Ara), and a route is the handful of doorsteps one driver does in one go. You make as many as you have vans, name them what you like, rename one, delete one, and drag them into your own order.

2. A Place carries its routes, and the days it is served. Open any Self collection Point and you now see **Delivery routes** - tick one, or several if more than one van serves it - and **Which days this Place is served**, where you can tick any day up to all seven. Which Places make a smooth run is your call: nothing in the app measures a distance, suggests a grouping, or stops you putting one Place on two routes. And the day list lives on the POINT, never on the route, so there is exactly one place a day is ever typed.

3. The shop lists every Place exactly as it does now - and a Place served on only some days SAYS SO. Your words again: *"every point listed in store, just that if that point only for friday, then mention it."* So nothing is hidden, greyed or filtered: Farlim still sits under Sg Ara with the others, and under its name it now reads **Only available on Fri**. And that line costs no new word in any language - it borrows the sentence the shop already uses for a product that is not on every day, so it reads correctly in English, Chinese and Bahasa Malaysia from today.

On the delivery run, the day's stops are grouped by route - a heading like **Route A · 3 stops** over the Places that van is going to, so you can tick one van's work and book it. A Place on two routes is still ONE row: it draws under its first route and names the other on its own line, because on that screen **a row is one stop** and every tick and every price rests on that. A stop on no route comes first, with no heading over it, the same way an unfiled Place draws at the top of the shop.

And nothing of yours moved. A Place you have not touched is served every bake day and is on no route, and it publishes **byte-for-byte** the row it published yesterday - I read the live payload back: of your four Places, only Farlim carried anything new, and the other three came out unchanged to the character.

No database step. The suite is **3,267 tests, all green** (21 new). **Nine bites, all fired by their own message** - and two of them caught real faults: a null in a half-synced day list was silently becoming **Sunday**, in the app and again in the shop's own copy of the same check. Also pinned: a route named on one phone reaches the other, and a rename can no longer put a Place on no route or open it every day of the week.

10 Oct 2026 - engine v418, THE HANDLE NOW MOVES THE DROP-DOWNS TOO (no database step - pushing this one is the whole of it).

Your words: *"after i reposition with the handle, the drop drop list have to organise is my preference, pls check all"*

YOU WERE RIGHT, AND IT WAS WORSE THAN ONE LIST - IT WAS EVERY PICKER FED BY A DRAGGED LIST. I checked all five screens that drag - **Ingredients, Products, Categories, Points, Areas** - and every drop-down they feed was listing the **raw record order**, with one exception. A drag writes a sort field (and an order inside a heading); **it never touches the array**, so the handle rearranged the shelf and left every picker as it was.

Four were wrong. All four are fixed:

- **Collect from** on an order card - the one you reach for most. **And this one needed more than a sort**: your order is **per area**, so sorting one flat list just interleaves the areas - a Point you put first under *Sg Ara* would sit level with one first under *Prai*, and **neither of them is first.** So the picker now reads **the same grouped order the Points screen draws** - your areas in your order, each with its Places - and the Places list moved into the file that owns the grouping, so a picker and the shelf it feeds cannot disagree. **The product box on an order card** - now in the order your Products screen shows them. **An ingredient box inside a product's recipe** - now in your Ingredients order. **The product box on a manual shopping list** - same order as the order card's.

And two were already right, and are now pinned so they cannot quietly break: the **Area** box on a Point and the **Sits under** box on a category both read the tree, which has always carried your order.

I also checked the two things this does NOT apply to. Suppliers, Units and Parcel couriers have no handle - nothing to follow - and **the shop** already drew your order from v410.

No database step. The suite is **3,246 tests, all green** (4 new). **Six bites, all fired by their own message** - and two of them proved a helper correct but wired to nothing, **which is why the rule is pinned in the source for the two pickers no test can open.** And one caught a half-fix of mine: I had made the picker follow your sort **flat**, and reading the live app showed it **interleaving your areas** - so it reads the grouped order instead.

10 Oct 2026 - engine v417, WHAT SELF COLLECT ACTUALLY SAYS (no database step - pushing this one is the whole of it).

Your words: *"You come to us - no delivery charges. We message you the exact collection point, and time window for collection., it is normally late afternoon on the bake day."*

THAT IS NOW THE LINE, WORD FOR WORD - in all three languages.

Self collect

You come to us - no delivery charges. We message you the exact collection point and the time window for collection. It is normally late afternoon on the bake day.

It answers BOTH questions a collecting customer has, and it answers them in your terms: **where** - *"the exact collection point"* - and **when** - *"the time window for collection"*, which used to be missing entirely.

AND "NORMALLY" IS DOING REAL WORK, SO IT STAYS. The window is **yours to set, per Point** - and a pickup time is not a promise. **Take that one word out and the sentence turns into a commitment the page cannot keep.** A test now pins it, with that reasoning beside it.

AND A BLEMISH IN THIS VERY DOCUMENT IS FIXED. Twenty-seven stray asterisks were printing across these pages - every entry that had been written with a single *like this* came out with the marks on the page, because the PDF's markdown reader only knows ... for italics. **It converts the pair now, and the count is down to one - the asterisk inside a real SQL count(*), which is meant to be there.** Older entries are repaired by it too, which is why this PDF has more changed pages than the entry above would suggest.

No database step. The suite is **3,242 tests, all green**. Three bites, all fired - including the words going back to the old line, and *"normally"* being removed.

10 Oct 2026 - engine v416, NEITHER WAY IS PICKED FOR HIM (no database step - pushing this one is the whole of it).

Your words: *"you need a 'click to choose' after the title, and self collect should not be default"* · *"for delivery as well"*

NEITHER SELF COLLECT NOR COURIER IS CHOSEN ANY MORE. The card used to open with Self collect already edged - which is the page answering its own question, and the very thing you were looking at. **It now opens with both plain and its own small arrow, and the edge only ever means a real tap** - the same rule the places got in v415, applied to the question above them.

AND THE QUESTION NOW SAYS IT IS ONE TO ANSWER: *"Tap to choose one."* With neither card edged, the two of them read as **two descriptions** rather than two things to pick - so that line stands under **"How will you get your order?"** until he picks one, and goes the moment he does.

AND A CONSEQUENCE THAT NEEDED A DECISION: WHAT IF HE NEVER PICKS? An order has to carry whether a rider is coming or the bread is being kept back, so there is **no honest default to fall back on**. **Pressing Place order with neither chosen now asks** - "Please say how you will get your order. Tap Self collect or Courier delivery above" - **which is the same stop the page already makes for a missing WhatsApp number, and for the same reason: it is the one thing the order cannot do without.**

THE BUTTON IS NEVER DISABLED. It stays live and pressing it **points at the two cards** - your standing rule, that no website rule may hide or block a sale. **This is a question, not a switched-off control.**

No database step. The suite is **3,242 tests, all green** (2 new). Every new rule proved by putting the fault back - including one that found **my own assertion was not testing arrival at all** (it set the value itself first, so it would have passed with Self collect still the default), and the rule was pinned in the markup instead.

10 Oct 2026 - engine v415, NOTHING IS CHOSEN UNTIL THE CUSTOMER CHOOSES (no database step - pushing this one is the whole of it).

Your words: *"the self collect have to have message follow, click to choose point, without click it should not highlight"* · *"only by clicking either self collect or delivery, the number 3 change color"*

BOTH OF YOUR POINTS ARE THE SAME RULE, AND IT IS A GOOD ONE: THE PAGE DOES NOT CREDIT HIM WITH A CHOICE HE DID NOT MAKE.

NO PLACE IS HIGHLIGHTED UNTIL HE TAPS ONE. The card used to open with a place **already lit** - the kitchen, or the Point you marked as one - and that is exactly what your v413 picture was showing me: **a place wearing the panel reads as a place the customer PICKED, and he has picked nothing.** It opens plain now,

and the highlight only ever means a real tap.

AND A LINE STANDS WHERE THE HIGHLIGHT USED TO: *"Tap the place you will collect from."* Without it the card would open on a list of names with nothing saying any of them is a thing to press. **It goes the moment he taps one** - a page must not go on telling him to do something he has just done.

AND THE 1 · 2 · 3 CIRCLE TURNS GREEN ONLY WHEN HE SAYS HOW HE IS GETTING IT. This one could not simply ask "is a way chosen", because the page sets that to Self collect the moment it loads - so the rule would have been true before a customer had touched anything. **It asks whether HE chose**, which is the only honest thing to ask. A name and a number on their own are no longer enough.

AND AN ORDER PLACED WITHOUT CHOOSING STILL WORKS, EXACTLY AS IT ALWAYS HAS. The empty place means "collect from the bakery", which is how every order has said it since long before there were Points, and the card already tells him the exact place is messaged to him. **This is a page that stops pretending, not a page that asks a new question** - nothing is refused and nothing is blocked.

AND TWO OLD FAULTS CAME OUT OF BUILDING IT.

1. **Tapping a place while already collecting left the instruction under his thumb** - the obvious place to clear it is the repaint, and **the repaint does not run when the way is already correct**, which is exactly that tap. Found by asserting it, which was the only way it was ever going to show. 2. **The area HEADINGS were being handed the highlight.** They sit in the same list as the places and carry no id, so "nothing is chosen" matched **them** too - **and it drew nothing wrong, because nothing styles a heading as active.** It had been there since v410 and only appeared when the browser was read directly. **A wrong class with no styling is a wrong class nobody sees.**

AND THE THREE CIRCLES ARE TESTED NOW, WHICH THEY NEVER WERE. They turn green through classList, and the shop's own test stub had an inert classList - so **a rule the customer watches change was being held up by the changelog alone.** The stub is real now, and nothing else moved: all 75 of the shop's existing tests passed with it.

No database step. The suite is **3,240 tests, all green** (2 new, and the two faults each proved by putting them back).

10 Oct 2026 - engine v414, "USE MY LOCATION" IS GONE, AND THE NOTE BOX SHOWS A TIME (no database step - pushing this one is the whole of it).

Your words: *"i think we should take out 'use my location' button totally" · "the e.g message in store for Notes for collection/delivery, change it to 'collect 6pm /prefer deliver 6 to 8pm'"*

"USE MY LOCATION" IS OFF THE PAGE - the button, everything behind it, and every sentence it could produce. The map press beside it does the same job better: a customer can SEE the door he is choosing and move it, where a phone's own guess at where he is standing cannot be checked at all. The typed address and the map are the two ways that remain, and skipping both is still fine - the driver goes to the address he typed.

AND THE SENTENCES WENT WITH IT, because a string with no control behind it can only rot: *"Finding you..."*, *"Location sharing is off for this site"*, *"Finding your location took too long"*, *"This browser can't share a location"*, *"Your location couldn't be found just now"*, and the warning about a spot that is only accurate to so many metres - **all six, in all three languages**, so nothing can ever print a message about a button that is not there.

AND THE NOTE BOX NOW SHOWS A REAL EXAMPLE. *"collect 6pm / prefer deliver 6 to 8pm"* in the empty box under **Notes for Collection/Delivery** - **where it used to say "Gate code, landmark, what to look for", which describes ONE of the two ways an order leaves.** A time is what both of them want, and an example she would recognise beats a description of the category. Translated into all three.

AND TWO OF YOUR TESTS MOVED RATHER THAN BEING THROWN AWAY. The pin's own tests placed a pin by pressing the button that just came out - and **not one of their assertions was about that button: they are about what the ORDER carries**, including the rule that the next customer does not inherit the last one's

front door. They are driven through the **map** now, in the file that already has a map to drive, with every assertion word for word as it was.

No database step. The suite is **3,238 tests, all green**. **And I broke one of my own files while testing this and had to restore it from the last commit** - a script of mine opened a file for writing before it had finished reading what it needed, which empties the file the instant it opens. **Nothing was lost, the file was restored, and the whole suite was re-run to prove it** - but it is written down, because a testing script that can damage a file is a testing script that needs to be built differently.

10 Oct 2026 - engine v413, THE PLACE YOU PICKED IS NOW THE ONE YOU CAN SEE (no database step - pushing this one is the whole of it).

Your words, with the picture: *"you see the white, that is a hole when the self collect card unfold. It confuse user of actually which is the selected"*

THE COLOURS WERE THE WRONG WAY ROUND, AND THAT WAS THE WHOLE FAULT. Every place row was drawn **white**, and only the **chosen** one was tinted - **and the card itself is tinted.** So inside an opened card, **the place you had picked dissolved into the background, while every place you had NOT picked stood out as a bright white slab.**

A customer reads the loudest row as the selected one - and it was pointing at the wrong place. That is exactly what you described: the white is the hole, and it is what "confuse user of actually which is the selected".

IT IS THE OTHER WAY ROUND NOW. The **chosen place is the one wearing the panel** - white, with the orange edge - and every other place sits flat on the card. Nothing on that list is louder than the one you picked.

AND EVERY OTHER PLACE STILL KEEPS ITS OWN THIN FRAME. *"can the unselected be frame as well with a very thin frame, so it is clear that custoerr is selecting from only 2 available choices"* - **and the frame was already there, which is the point:** the line colour is a **cream chosen to sit on white**, so on the card's own tint it was invisible and an unchosen place read as bare text rather than as the other thing you may pick. It is a soft brown hairline now, so **every row reads as one OF the choices** while the picked one stays the loudest.

AND I OWE YOU A NOTE ABOUT HOW LONG THIS TOOK. You said the word "hole" **three times**. I fixed one real thing (the card taking the form's white) and reasoned about the colours twice from numbers - and **the moment you sent the picture it was obvious in one look. A picture from your phone costs you nothing and would have saved all of that; I should have asked the first time.** That is written down now.

No database step. The suite is **3,237 tests, all green**. One bite, fired by its own message, and the rule is **pinned in the stylesheet** because nothing in this suite can see a colour.

10 Oct 2026 - engine v412, BOTH WAYS FOLD, THE HOLE IS GONE, AND A POINT CAN SAY ITS OWN WORDS (no database step - pushing this one is the whole of it).

Your words: *"i want the self collect card at store to be folded as default"* · *"even after i expend it, there should not be hole"* · *"courier delivery and address should be one piece"* · *"make the point description custmable in point card"*

BOTH WAYS NOW ARRIVE FOLDED. On your shop, a customer sees **"Self collect"** and **"Courier delivery"** as two short cards, each with its own small caret - and nothing else until they tap one. Tapping a card **chooses that way and opens it**; tapping it again puts it away. **Only one can be open at a time**, so the fields of the way the customer did NOT take are never standing open under it.

AND THE ADDRESS IS PART OF THE COURIER CHOICE, THE SAME WAY THE PLACES ARE PART OF COLLECTING. That was the mirror I had missed: at v410 I moved the places *into* the Self collect card and left the delivery address sitting below the whole picker as a separate box. Her words - *"courier delivery and address should be one piece"* - are the same rule twice, so it is now the same code twice.

AND THE CARD NO LONGER TAKES THE FORM'S OWN WHITE. *"even after i expend it, there should not be hole"* - the picker sits inside the white "Your details" form and each card had the same white, so an opened card was **white on white: a void with thin outlines floating in it**. The cards now take the page's own colour, so an open one is a panel.

AND THE HOLE WAS NOT ACTUALLY GONE - SEE v413. That was only half of it, and **not the half that mattered**: the colours of the place rows themselves were the wrong way round, which is what made a customer unable to tell which place he had picked. It is fixed in the next entry, and it took a picture from her to see.

AND A POINT CAN SAY ITS OWN WORDS. A new box on the Point card - **"What customers read"** - and whatever you type there is the line under that Point's name on your shop. **Leave it empty and nothing changes at all**: the standard sentence stands in, and the box's grey hint IS that sentence, so you can see what a customer would read before you write anything. **Your words outrank even the kitchen's own line**, because they are about your own front door. A Point with a smallest basket the basket has not reached still says WHY on its own line - that is the one thing you can act on and your description cannot know it.

AND ONE MORE OLD WRONG THING, FOUND WHILE LOOKING FOR SOMETHING ELSE. After an order is placed the shop resets itself to Self collect for the next customer - and it was doing that **by hand**, writing the highlight onto the words rather than onto the card. It had been wrong since v410 moved the edge to the card, so **a second order from the same phone lit the wrong element and left the fold wherever the last customer left it.** It now goes through the one place that owns that decision.

No database step. The suite is **3,237 tests, all green** (6 new). **12 bites, all fired by their own message - and two of them proved my own new assertions were too weak and had to be tightened**: one only checked that the address came *after* the courier card opened, which a box moved out below the picker satisfies perfectly; the other, earlier, checked a redraw against a card list it had just built instead of the one on the page.

Still owed: the Chinese and Malay versions of the new sentences are **mine, not a native speaker's**.

10 Oct 2026 - engine v411, A POINT CARD ARRIVES FOLDED - AND A LAYOUT FAULT I SHIPPED IN v410 (no database step - pushing this one is the whole of it).

Your words: *"can the point fold up by default?"* · *"why the point card are in such orientation?"*

FIRST, THE ONE YOU SPOTTED, WHICH WAS MY FAULT AND SHIPPED IN v410. *"why the point card are in such orientation?"* - because I broke it. Adding the drag grip made each card a row with the grip on the left, and **the wrapper I put round the words only caught the first block of them.** The three detail lines - the minimum, the collection window, the pin - and the four buttons were left as **brothers of that wrapper inside a flex row**, so every one of them became a column standing side by side: **the text was squeezed to 12 pixels wide** and the card read as a stack of slivers. **Every test passed**, because nothing in this suite has a layout engine - and I checked the SHOP page in a browser but only ever drove this screen by reading its text. **You saw it and I had not.** It is fixed: one wrapper round the whole body.

AND THE SAME FAULT HAD A SECOND HALF. The card's body is 44px narrower than the card, because the grip takes that off it - so the four buttons, which were fine before, no longer fitted, and **"Put the pin on the map" folded onto five lines.** Fixed with the same rule your Money screen already carries for exactly this: buttons sized by their own labels, and the row wraps when it has to.

AND THE CARDS FOLD, WHICH IS WHAT YOU ASKED FOR. A Point card now arrives **shut** - the name, its chips and a caret - with everything else one tap away. **A folded card still tells you what you need without opening it:** who hands the bags over, what the Point costs you per order, and how many orders have gone there. **A list of names alone would be a list you had to open row by row.** And what you open **stays open across a redraw**, so pausing or editing one Point does not snap the card shut under you.

The whole Points screen went from 2868 pixels to 2144 on a phone - with the same four Points and the same headings.

No database step. The suite is **3,230 tests, all green** (1 new). Three bites, all fired by their own message - **including one that found an assertion of mine proving nothing:** my "stays open after a redraw" check rendered into a *second* card list and left the first one alone, so it passed whatever the code did. It now redraws the same list.

Still owed: the Chinese and Malay versions of the new sentences are **mine, not a native speaker's**.

10 Oct 2026 - engine v410, YOUR COLLECTION POINTS GROUP BY AREA, THEY CAN BE DRAGGED,

AND EACH ONE DECIDES WHETHER ITS ADDRESS SHOWS (no database step - pushing this one is the whole of it).

Your words: *"i need something like category for product on POINT. 1st level 2nd level, 3rd level....example top level is Penang Island & Prai. Under Penang Island will be Area like Sg Ara, Balik Pulau, Farlim, Georgetown"* · *"The store under Self Collect, the chioces should be able to list clearly, in point card, a switch to on/off address to be shown or not"* · *"design the handle too"* · *"there is 2 main selection, select self collect and a drop drown appear, should not be in present interface"*

YOUR POINTS NOW SIT UNDER AREAS - a tree as deep as you like. *"something like category for product"* was exactly right, and that is what it is: the same pattern as your product Categories, so it works the same way and there is nothing new to learn. **More -> Collection areas**, add *Penang Island* and *Prai*, then nest *Sg Ara* and *Balik Pulau* inside the island. **Any depth** - a third level costs nothing. Then on each Point's card there is a new **Area** box, and filing a Point is saving it.

AND A HANDLE, SO YOU SET THE ORDER. *"design the handle too"* - the same grip your Ingredients, Products and Categories already carry, on the left of every row, on the areas AND on the Points. **Drag reorders among the brothers and nothing else** - moving a Point to a different area is the Area box, not a sideways drop, because *"reading a sideways drop as 'move it somewhere else' is a rule you have to be taught"*.

AND EACH POINT DECIDES WHETHER ITS ADDRESS IS PUBLIC. A new tick box on the Point card: **"Show this address on the shop"**. Off - as every Point is today - and your public page shows the name only. On and your page prints the street under it, so a customer knows where he is going before he orders.

AND IT IS THE SHOP ONLY. YOUR WHATSAPP MESSAGE IS UNTOUCHED, AND THAT IS DELIBERATE. A customer who has **paid** must be told where to walk or the order cannot be collected, and the two are not the same kind of thing: **the shop is a public page anyone can read; the confirmation is a private message to one person who has already chosen.** A test now pins that, precisely because switching it off in the message is the tempting change and it is the one that strands somebody.

AND THE PLACES ARE PART OF CHOOSING TO COLLECT, NOT A BOX THAT APPEARS. Your words - *"select self collect and a drop drown appear, should not be in present interface"* - were right: it was **two steps for one decision**, and a control that appears once you have chosen something reads as a second question. **The places now live inside the Self collect card**, with the border round the whole piece, and **tapping a place chooses Self collect as well as the place.** There is no longer a state where somebody has picked a place without saying they are collecting, or the reverse. The **"Collect from"** box and its line are gone from the page and from all three languages.

AND THE HEADINGS ARE HER WORDS, UNTRANSLATED - *Penang Island*, *Sg Ara* - the same rule a Point's own name already follows: your data, and a customer reading in Mandarin sees "Sg Ara" today exactly as he does now.

THREE FAULTS FOUND WHILE BUILDING THIS, AND TWO OF THEM WERE MINE.

1. **The Categories drag handle was not styled at all**, and had not been since v399 - a comment and a rule had been inserted **inside a selector list**, so the browser read it as .cat-handle, .ing-card-row .card-row { flex-wrap: wrap } and the grip never got its size, its cursor, or **touch-action: none** - the one that matters, because without it the page scrolls under your finger and a category cannot be dragged at all. **Fixed**, and the new Point grip joins the one rule rather than getting a copy. 2. **"Penang Island" did not draw on the shop.** The page only drew a heading in front of a Point that belonged to it, so a parent heading holding no Point of its own **vanished** - and **Sg Ara came out indented one step with nothing to be indented from. Found by LOOKING at the page, not by a test: my fixture gave every heading a Point of its own, so every test passed.** The app's own grouping had it right and the shop did not. 3. **My own drag would have thrown in the middle of your drop** - I called the redraw from inside the drop, and the helper moves the row *after* that returns, so it would have moved a row already detached from the page. Caught by reading the helper, and fixed by not redrawing: **a reorder changes no card's text, only where it sits.**

No database step. The suite is **3,229 tests, all green** (35 new). **21 bites, all fired by their own message** -

including the two that were caught by an older test first and had to be rewritten until only the new rule could catch them, and one that found a test of mine proving nothing because its fixture's dates happened to make two different orders look identical.

Still owed: the Chinese and Malay versions of the new sentences are **mine, not a native speaker's**.

10 Oct 2026 - engine v409, YOUR KITCHEN NOW SAYS IT IS YOUR KITCHEN (no database step - pushing this one is the whole of it).

Your words: *"yes, fix the kitchen line"*

YOUR KITCHEN WEARS THE KITCHEN'S OWN LINE AGAIN. v407 dropped the invented "Our kitchen" row - correctly - and never gave the real one its words, so your kitchen went on reading *"Self collection Point - we message the exact spot and time once your order is confirmed"*: **true, and useless.** It told a customer collecting **at the bakery** exactly what it would tell him about a friend's shop. It now says *"Sungai Ara, Bayan Lepas - where we bake"*, which is what the row it replaced used to say.

And it was found by LOOKING at your live shop, not by a test. The two lines were **identical in every fixture I had** - because a test fixture is not a bakery. Your shop is where **Sg Ara is marked as the kitchen**, so it was there to be seen. That is twice in two days that opening the real page found something the suite could not.

No database step. The suite is **3,194 tests, all green**. One bite, fired: put a Point's line back on the kitchen and the test fails by its own message - *"the kitchen says it is the kitchen - it read: 'Self collection Point - we message the exact spot...'"*.

AND ONE THING WORTH KNOWING, BECAUSE IT NEARLY MISSED YOU. This fix was finished **before** you pushed v408 - and the push went out from the edit that was already saved, **leaving this one sitting uncommitted behind it.** Nothing was lost: it is this release, and it is the reason this is a version of its own rather than part of v408. **If a fix ever seems to have gone missing, that is the first thing to check** - a push between two of my messages can leave the newest change unwritten to GitHub while the message above it goes up.

10 Oct 2026 - engine v408, THE NOTE FIELD NAMES BOTH WAYS, AND A CODE IS SAID AGAIN AT THE BUTTON (no database step - pushing this one is the whole of it).

Your words: *"change the Delivery Notes(optional) to Notes for Collection/Delivery"* · *"If there is promo coupon attached, say it on top of Place Order, word like: Your promo coupon will be attached together to the order for validation, actual payable will be confirm by Whatsapp. You can better the phrase"*

THE NOTE FIELD IS NOW "Notes for Collection/Delivery (optional)." The same fault a third time, and the second one you caught yourself: the label named **one** of the two ways an order leaves. A customer collecting his own bread read *"Delivery note"* and skipped a box that asks him for the **gate code or the landmark** - the thing that actually gets him his bread. **The placeholder moved with it** ("delivery time" is the wrong question to ask someone who is coming to fetch it).

And I changed the same label on THREE of your own cards too - **New order**, **Edit**, and **Note / tracking**. It is the same field, with the same words, on cards that take both kinds of order, and **two screens disagreeing about what one box is for is how the next confusion starts.** If you meant the shop only, say so and I will put yours back.

AND A CODE IS NOW SAID AGAIN, RIGHT ABOVE PLACE ORDER:

Code FRESH10 goes with your order. We'll check it and confirm the final amount on WhatsApp.

Your phrasing, tidied - and one word of it matters. You said *"attached together to the order for validation, actual payable will be confirm by Whatsapp."* **That is exactly right, and the page already half-said it** - the line under the code box reads *"We'll take it off when we confirm your order"* - **but that is up the page.** By the time a customer reaches the button it is off the screen, and they are looking at **one** figure: the total on the bar, **which is the goods and never has the code taken off it.** So the promise is repeated where the money is.

THE LINE NAMES YOUR CODE BACK TO THEM - *"Code FRESH10 goes with your order"* - because "your code" is abstract ten minutes into a page. **And it promises no figure**, which is deliberate: whether the discount holds is yours to decide at the confirmation (a code can be one-per-customer, or first-order only), so the line says the amount is **confirmed**, never that it is already correct.

It only appears when a code is actually on the order. An ordinary order's bar is the bar it has always been, byte for byte - hidden and empty, adding nothing to its height.

AND THE FIXED BAR NO LONGER RELIES ON A NUMBER MEASURED BY HAND. The bar is stuck to the bottom of the screen, and .bar-spacer is the only thing stopping it swallowing the last row of the shop. That spacer was a **measured 126px** - right for a two-row bar, and **wrong the moment a third row appears.** It is now read off the bar itself on every repaint. **Proved on the worst case, at 320px wide in Bahasa Malaysia, with a code on:** the bar reaches **179px** and the spacer follows it exactly - **the old fixed 126px would have left 53px of the page hidden behind the bar.**

No database step. The suite is **3,194 tests, all green** (3 new). Nine bites, all fired - the line never shown, the code never named, a ringgit figure promised, the label drifting back in the shop, in the markup and in your admin cards, and the placeholder going back to "delivery time". **Two of them were caught by an older test first, so they were rewritten until only the new rule could catch them** - and one found that **my admin label change was not pinned by anything at all**, so a test was written for it.

Still owed, and named so it is not forgotten: the Chinese and Malay versions of the new sentences are **mine, not a native speaker's** - your circle should read them.

10 Oct 2026 - engine v407, THE TWO WAYS TO GET YOUR BREAD NOW SAY WHAT THEY ARE, AND THE DAY IS CALLED A BAKE DAY (no database step - pushing this one is the whole of it).

Your words: *"do the self collect and courier lines"* · *"from the store i can see the bake day is term as delivery day, that might be what confusing. Change it to bake day"* · *"it should be 1. pick a bake day"* · *"order can be collected late on the bake day, a message will be send to you when your order is ready"*

SO THE TWO WAYS NOW SAY WHAT THEY ARE, BEFORE YOU TAP EITHER ONE. The picker was two bare words side by side - *Self collect* and *Courier delivery* - and nothing on the page said what either one meant or what it would cost. **That is the other half of what went wrong for your customer**: he chose a courier when he wanted to collect, and the page had given him nothing to choose with.

It is now one full-width row per way, a bold name over a plain line:

- **Self collect** - *"You come to us - no delivery charge. We message you the exact place, and once your order is ready you can collect any time on your bake day."* **Courier delivery** - *"A rider brings it to your door. Any delivery charge is told to you on WhatsApp - the total shown is for the bread only."*

BOTH LINES ARE TRUE BEFORE ANYTHING IS TAPPED. That is the whole of it: he can compare the two without tapping, and cannot tap his way into a choice he did not understand. When he does tap one, it takes the orange edge the collect-from rows already use, so which one is on is never a guess.

AND COLLECTING NOW ANSWERS THE SECOND QUESTION: WHEN? Your words - *"order can be collected late on the bake day, a message will be send to you when your order is ready"*. **The answer is the message, never a clock** - a pickup time is not a promise (v340), and the bread is ready when it is ready. **And it still names the place, which is not decoration**: while you have no Self collection Point open, **the collect-from list is empty and hidden**, and that line is the **only** place the shop can tell a customer where to come.

AND THE DAY IS NOW CALLED A BAKE DAY - IN ALL THREE LANGUAGES. Your catch, and it was the same fault as the two bare words. The page said *"Delivery days Mon, Wed, Fri"* and *"1 Pick a delivery day"* - naming the day after the thing that happens to it **for some customers**. The one person it misled was the one **collecting his own bread**: he did not recognise the day he was choosing. **It is the day the bread is baked**, which is true of every order however it leaves.

So it is a bake day everywhere a customer meets it - the row under the answer, the heading **1 Pick a bake day**, the line on the calendar, *"Your bake day: Fri, 9 Oct"*, the notes that say when a product cannot be ordered,

and the change/cancel window. **English, Mandarin and Bahasa Malaysia together** - 烘焙日 and *hari membakar*.

AND I FOUND A REAL FAULT OF MINE WHILE DOING THIS, SO IT IS FIXED HERE. v406 made your kitchen **a collection Point you tick** - but **the shop was never told**, so it went on printing its own invented *"Our kitchen"* row above the list. **The moment you ticked a kitchen, a customer would have seen your kitchen TWICE**: once invented, once under the name you gave it, as two choices that look like two different places.

Now the flag travels (your app publishes it, the shop keeps it) and **the invented row is the fallback it was always meant to be** - it stands in for your kitchen only while you have not ticked one. **It cannot simply be deleted**: an order taken before v406 carries no Point at all, and that empty id **is** the kitchen, so the row that posts it has to stay.

No database step. The suite is **3,191 tests, all green** (5 new). Every new rule was proved by putting the fault back: the invented kitchen row returning, a ticked kitchen leaving nothing chosen, the flag dropped on arrival at the shop, a row losing its explanation, and the day drifting back to "delivery day" - **two of those bites were caught by an older test first, so they were rewritten until only the new rule could catch them.**

Still owed, and named so it is not forgotten: the Chinese and Malay versions of the new sentences are **mine, not a native speaker's** - your circle should read them (the page is already live).

10 Oct 2026 - engine v406, THE THREE STEPS TURN GREEN, AND YOUR KITCHEN IS A COLLECTION POINT (no database step - pushing this one is the whole of it).

Your words: *"the 1,2,3 turn green after customer done with it" · "I want all self collection thru collection point, not from the kitchen" · "at self collection point, add a switch whether that collection is a kitchen. Allow only one collection point as kitchen for the time being. NO pin is needed if it is kitchen" · "a point need a pin" · "Make the trial an official one now"*

THE TRIAL IS OFFICIAL. The shop page a customer opens **is** the new one - no address flag, nothing to opt into. **The flag is gone rather than left switched on**, because a flag that is permanently true is a branch nobody will ever test the other half of. **The self collect / courier lines are still not in it** - that is still the next piece.

THE 1, 2, 3 TURN GREEN AS THEY ARE DONE. A day chosen, something in the basket, name and number filled in - each circle goes green the moment it is true, so *where am I* is answered without reading anything.

AND THAT WAS MY FAULT, SO IT IS WORTH SAYING. The green state was **drawn** in the first version and **never wired** - nothing ever set it, so the circles stayed the same colour through the whole order. **A step marker that never moves is decoration**, which is exactly what it was meant to replace.

AND YOUR KITCHEN IS NOW A COLLECTION POINT. The **New Self collection Point** card and the **Edit** card each carry a tick box - **"This is the kitchen"** - and **ticking it unticks whichever Point held it**, so only one ever can. **A Point marked that way needs NO pin**: you are not going to pin your own front door, and a customer collecting there is given your address. The Point's card shows **The kitchen** so you can see at a glance which one it is.

AND ONE DELIBERATE RULE OF YOURS HAS BEEN OVERRULED - SO WHAT IT WAS PROTECTING IS KEPT. The Points file said, in its own words: *"THE KITCHEN IS NOT A POINT... **a Point that quietly became the kitchen would inherit a fee she does not owe and a life she cannot end."*** **That was written for good reason and it is real**: a Point carries a fee (what you pay whoever receives), a smallest-basket rule, and a **Pause**. **So your kitchen is read as FREE, with NO minimum, and it can never be paused** - whatever is stored on the row. It is the one place that must always exist, because it is where the bread is.

AND A POINT DOES NEED A PIN - THE KITCHEN IS THE ONLY ONE THAT DOES NOT. A Point without one is a name and nothing else, and a customer is being told where to **walk**. **The rule is kept at the doorway rather than at the Save button**, and that is deliberate: a new Point cannot have a pin when it is made - the map press is on its card afterwards - so requiring one to save would make a new Point unsaveable. A Point with no pin is simply **not offered to a customer** until you put it on the map.

AND NOTHING OLD IS BROKEN. A Point is only the kitchen if you tick it. **Until you do, a customer is still**

offered "My kitchen" exactly as before - and every order taken before today, which carries no Point of its own, still means the kitchen.

No database step. The suite is **3,186 tests, all green** (9 new). Every new rule was proved by putting the fault back: the green never wired, a second Point made the kitchen at once, a Point with no pin offered to a customer, and the kitchen inheriting a fee and a Pause.

10 Oct 2026 - engine v405, THE SHOP NOW ANSWERS THE CUSTOMER'S QUESTION (trial only - see the link below).

Why. Your first customer ordered without your help, and it went wrong three ways - a wrong day, courier instead of collecting, and underneath both of them: **he wanted bread today and did not know ordering closes the day before.**

AND THE PAGE'S OWN FAULT WAS THAT IT STATED THE RULES AND NEVER THE ANSWER. "Delivery days Mon, Wed, Fri" and "Order by 18:00 the day before" are both true, and neither tells a stranger whether HE can have bread, or when the earliest he can get it is. He had to add up today's date, three weekdays, the clock and "the day before" **in his head, on a phone, while not being confident online.** So he picked a day that looked fine and it was wrong.

SO THE PAGE NOW SAYS IT OUT LOUD, ABOVE EVERYTHING ELSE:

For a future session: build_changelog.py reads a > at the start of a line as ordinary text, so this printed a bare > down the page. Written out plainly now.

TODAY'S BAKE HAS CLOSED

The next day you can have bread is Mon, 12 Oct.

Order before 6pm on Sun, 11 Oct, and it's yours.

When there is still time the same box opens *"You're in time"* and turns green. **The rules stay underneath for anyone who wants them** - nothing has been taken away.

AND IT TELLS YOU THERE ARE THREE THINGS TO DO. One line - *"You'll do three things - and then we take over"* - and then a big numbered circle on each part: **1 Pick a delivery day · 2 What would you like · 3 Your details. The circle turns green as each one is done**, so *where am I* is answered without reading anything. That sentence is doing real work for a nervous customer: **it promises the page ends.**

AND IT SAYS 6pm, NOT 18:00. 18:00 is how your admin stores it and that is right there; on a customer's page it is a small puzzle, and a small puzzle is exactly where he went wrong.

AND NONE OF IT IS LIVE. IT IS AT A TRIAL ADDRESS, ON YOUR WORD:

https://jienluv2bake.com.my/store/?trial=1

Open that address and you see the proposed page; open the ordinary shop address and you see exactly what is there today. Everything above - the answer box, the three steps, 6pm - is behind that one flag. **Send the trial address to your circle.** It is a flag rather than a second copy of the shop on purpose: a copy would be a second shop to keep in step with this one, and it would put an unlaunched ordering page at a guessable address - **you already have one public /test/ you have been meaning to deal with, and this must not become the second.**

ONE THING TO KNOW BEFORE YOU SEND IT, AND IT IS THE ONLY RISK HERE. The trial page is your REAL shop with different words on it, so **an order placed from that address is a REAL order** - it lands in your Orders and takes a receipt number. If your circle should place throwaway orders, tell me and I will mark them so you can spot and delete them.

AND TWO THINGS NOT DONE YET. The self collect / courier lines - the other half of what went wrong - are **not in this**; you chose to see this first. And the Chinese and Malay versions of the new sentences are **mine, not a native speaker's** - pick() falls back to English rather than blanking, so nobody meets a hole, but a Chinese customer reading English is a downgrade and it wants a pair of eyes before the flag comes off. **Your circle is exactly who should check it.**

No database step. The suite is **3,181 tests, all green** (4 new). Every new rule was proved by putting the fault back: the ordinary page showing the trial box, the clock reverting to 18:00, and the cut-off being ignored when working out the next day - **that last bite found my own test was too weak to catch it, and it was rewritten before it was trusted.**

10 Oct 2026 - engine v404, TAKING A BAKE DAY OFF NOW ASKS FIRST (no database step - pushing this one is the whole of it).

Your words: *"it is just a click to take away baking day, it should be with confirmation"*

IT ASKS NOW - EVERY TIME. On the **Bake days** calendar, tapping a green day that carried **no orders** took the day away **on the spot**, with no question. It now asks, exactly as it already did for a day that has orders on it.

AND THE HOLIDAYS MADE IT WORSE THAN IT SOUNDS. Tapping a marked day does two things - it names the day AND removes it - so a tap meant to do one did both. Now the naming still happens, and **the day only goes when you confirm.**

AND THE OLD BEHAVIOUR WAS DELIBERATE, WHICH IS WHY IT IS WORTH SAYING PLAINLY. The code said so in its own words: *"taking back a date she just added is the reverse of the tap that added it, so it happens on the spot."* **That was a reasonable rule and it was wrong**, because the calendar is a **month grid that scrolls under your thumb** - the one screen in this app where a control that acts on a single stray touch does real damage. **Your call overrules it, and the old reasoning is replaced in the code rather than left standing.**

AND THE QUESTION NAMES THE DAY THE WAY YOU READ IT. It used to name it in the form the app stores dates in - *"Delete bake day 2026-10-13?"* - which you have to decode. It now reads **"Delete bake day 13 Oct 2026?"**. A confirmation you have to stop and work out is one you learn to dismiss, which is the same as having none.

AND ONE THING IT DELIBERATELY DOES NOT DO. The question for a day that holds orders says so, and says what happens to them: *"...has 1 order(s) on it. Delete the date? **The orders are kept in your delivery history**."* Without that, the question reads as though deleting the day deletes the orders - and a confirmation that leaves you guessing at the worst is one you cancel out of habit.

NOTE ON THE COMMIT. This release touches none of the files in the test-only commit before it (fd23de0, the word "null" guard). They are separate on purpose.

No database step. The suite is **3,178 tests, all green** (2 new assertions). Every new rule was proved by putting the fault back: the tap removing the day with no question, and the reassurance about the orders dropped from the message.

10 Oct 2026 - engine v403, MONEY AND PROFIT CAN LOOK AT ANY DAY, WEEK OR MONTH (no database step - pushing this one is the whole of it).

Your words: *"At money, we have 3 choice to look at the money, today, this week, this month. I need something more vasatile, like what we have for invoice. I fact there is several place same nature that we should improves, pls search it"* - and then *"same thing for Profit"*.

MONEY NOW ASKS TWO QUESTIONS, LIKE THE INVOICE DOES. Four choices - **A day · A week · A month · All** - and then **left / right arrows** to step to the one you want, with the window named above them:

[A day] [A week] [A month] [All] left Week of 4 Oct - 10 Oct right

THE OLD THREE COULD NOT DO THIS. They picked the *kind* of window and stopped there - "This week" was always the current week, and **last week was unreachable**. There was no way to look at a week that had already finished.

AND PROFIT GOT THE SAME CHOOSER. It used to be a month at a time, so a week's profit was a question you could not ask. Now it has the same four choices and the same arrows. **Three screens now share one builder**, so a week means the same seven days on the invoice, on Money and on Profit - they cannot drift apart.

AND YOUR APP HAD THREE DIFFERENT DEFINITIONS OF "A WEEK", AND MONEY WAS THE ODD

ONE OUT. Home's "This week" and the invoice both start on **Sunday**; Money started on **Monday**. So two screens offered something called "this week" covering **different days**. Asked which to use you chose **Sunday, to match the rest** - **and this DOES move today's figures**: your week now starts a day earlier, so Sunday's takings are inside "This week" where they were not before.

AND MONEY STILL COUNTS UP TO TODAY. A week or a month **still running** is counted to now, never into the future - a till has no takings from a day that has not happened. A period that has **already finished** is shown whole, because every day of it really did. The invoice shows whole periods instead, which is right for a document about what you supplied.

AND THE SEARCH FOUND A FAULT I HAD ALREADY HALF-FIXED - WHICH IS WHY IT IS WORTH TELLING YOU ABOUT. "All" means *no dates at all*, and the code that checks whether a day is inside a window **silently matched nothing** when handed no dates: "2026-10-01" <= "" is **false**, so an unbounded window excluded everything. **Profit's "All" read Sales RM 0.00 beside a month of real orders, and Money read every figure as zero - with no error anywhere.** That comparison had been written out **three separate times**, in three files, and fixing one left the other two broken. **It now lives in one place.** Found by opening the screens and looking at the figures, not by any test.

AND WHAT THE SEARCH FOUND, since you asked. Only **two** things were pinned to "now": **Money** (now fixed) and the **Customers** filter's "Ordered in the last 30 days". Everything else - Profit, Bake days, Products' availability, the Orders calendar, Home's delivery weeks - could already step. Home's "This week", "Coming 4 weeks" and promo's "last 28 days" are fixed *cards* rather than choosers, which is a different thing. **Customers' 30-day filter is left alone** - say the word if you want it changed too.

NOTE ON THE COMMIT. This release and the one before it (v402 - the van screen becoming the kitchen pin) share the same version and changelog files, so they land in ONE commit. Both are described here.

No database step. The suite is **3,177 tests, all green** (4 new). Every new rule was proved by putting the fault back: the range rule reverting to the plain comparison (breaking "All" on both screens), Money's week going back to Monday, and Profit's "All" reading zero.

10 Oct 2026 - engine v402, THE VAN SCREEN IS NOW THE KITCHEN PIN (no database step - pushing this one is the whole of it).

Your words: *"instead of send a van, change the term to Location pin for Kitchen"*

SO IT IS CALLED THAT NOW. On **More -> Logistic** the row reads **Location pin for Kitchen**, with the line *"Your kitchen's own door, the pin every trip starts from, and the van service that collects there"*. The mark changed from the van to the pin, because the pin is what that screen actually holds.

AND THE REASON IS THE POINT. That screen was named after the van - v309's reasoning was *"naming this screen after the van keeps one word for one idea"*, because an order card's **Send a van** is one of the two ways an order LEAVES. **The idea it was keeping one word for was the wrong one.** This screen does not send anything: it holds **the one point every trip starts from - your own door** - and the van is one of the things that uses it.

AND THE ORDER CARDS KEEP SAYING "Send a van". You were asked which one to rename and chose the screen alone, which is the right call: on an order card those words name **how the order leaves**, and "Location pin for Kitchen" would be the place rather than the trip. **The two are now named for their own jobs rather than matching each other.**

AND NO ADDRESS MOVED. The screen still lives at the same place in the app - a rename has never moved a link, and the guide's own rule is that every address is what it always was, so nothing you bookmarked stops working. Only the words changed.

AND THE NAME LIVES IN THREE PLACES, which is a trap worth knowing about: the menu row, the title at the top, and the heading on the screen itself. Renaming one and missing another would have left you with **two names for one screen**. A test now reads all three together and fails if they disagree - the same rule that already checks every "More -> something" signpost against the screens that really exist.

No database step. The suite is **3,173 tests, all green** (1 new). The rule was proved by putting the fault back four ways: the menu row keeping the old name, the title keeping it, the address moved along with the name, and the order card renamed as well.

09 Oct 2026 - engine v401, A SHOPPING LIST YOU CAN WRITE YOURSELF (no database step - pushing this one is the whole of it).

Your words: *"I want to add a manual PO issuing, the rest of the po process follow what we already have for po processing. I think a thing only different is it dont tie to specific bake date"* - and, asked what goes on it, *"Just pick products to bake"* with *"then i can use the amend to adjust the details"*.

THE PO SCREEN HAS A NEW PRESS: WRITE A LIST BY HAND. Pick the products and how many, and the app works the ingredients out of your recipes and prices them in whole packs from your cheapest supplier - **exactly as a bake day's list does**. Nothing here is on a bake day.

AND THE WHOLE REST OF THE PROCESS IS UNCHANGED, because it is the same machinery. The saved list appears in PO history like any other, with the same per-shop groups, the same **Buy / Not buying**, the same **Undo**, the same **Print this shop**, the same **Amend** for adjusting any detail afterwards, and the same stock and money behaviour when you buy from it.

AND THIS IS THE PART THAT MATTERS: SAVING IT NEVER MARKS A BAKE DAY AS SHOPPED. A normal list remembers which bake days it covers; this one **records no day at all** - so it can never tick a day off your list, and deleting it can never un-tick one. **That is the entire difference**, and it is an absence rather than a rule, which means there is nothing anywhere that can forget it.

AND A FORGOTTEN CORNER I FOUND WHILE BUILDING IT. With **no bake days** at all, that screen used to stop at one sentence - so the new press would have been **invisible exactly when it is most useful**. A list that needs no bake day has to be reachable when there are none. The empty screen now offers it too.

AND TWO FAULTS MY OWN CHECKS CAUGHT, WHICH IS WHY THEY ARE NOT IN THIS RELEASE. The press first referenced something that did not exist where it was drawn - **the smoke harness that presses every control on every screen found it**, and no rendering test could have. And the product picker's answer was being read from the wrong place, so **choosing a product changed the picker on screen and did nothing at all** - found by driving it, not by reading it.

No database step. The suite is **3,172 tests, all green** (2 new). Both new rules were proved by putting the fault back: the hand-written list recording a day list, and a list genuinely claiming both bake days and ticking them off.

09 Oct 2026 - engine v400, THE BUTTON SAYS BUY, AND IT BECOMES BOUGHT (no database step - pushing this one is the whole of it).

Your words: *"can the button initially say buy, after buy it become bought?"* - and yes, that is exactly what it does now.

A SHOP YOU HAVE NOT BOUGHT FROM SAYS BUY. ONCE YOU HAVE, IT SAYS BOUGHT. Press **Buy**, answer what you paid, and the press is replaced by the record - *Bought tick 9 Oct, 11:42* - with **Undo** beside it:

Mydin · RM 96.00 · [Buy] [Not buying] Yen Grocer · RM 34.00 · Bought tick 9 Oct, 11:42 · [Undo]

AND THE TICK MEANS SOMETHING NOW. "Bought tick" was doing two jobs, and only one of them well. On the **press** the past tense plus a tick read as *already done* - but since v389 pressing it only opens "What did you pay?", and **nothing moves until you answer**. The tick belongs to the record, where it is true: *Bought tick 11:42* can only appear once the packs really are on your shelf.

AND THE TWO PRESSES ARE NOW A PAIR. Buy and **Not buying** are the two answers to one question - did you buy at this shop? Before, the row mixed a verdict with a decision ("Bought tick / Not buying"), which read as though one had already been settled.

AND NOTHING ELSE MOVED. The record still says **Bought**, the stock card still says *Bought - Mydin (12 Oct)*, and the card still counts *Bought 1 of 2 shops - still to buy: Mydin*. Only the press changed its word.

ONE THING YOU MIGHT WANT TO CHANGE TOO. The shopping list's table has a column headed **BUY** (the

quantity to buy) and now a **Buy** button sits beside that shop's name. They are different things - a heading against a button - but if the two reads as confusing on your phone, say so and I will rename one.

No database step. The suite is **3,170 tests, all green** (2 new). Both new rules were proved by putting the fault back: the press reverted to "Bought tick", the record word removed, and the pair broken by renaming one of them.

09 Oct 2026 - engine v399, YOU ARRANGE YOUR OWN INGREDIENTS (no database step - pushing this one is the whole of it).

Your words: *"can you give me a handle to reorganise the ingredient card?"* - and asked what you wanted to move, *"the order of the cards"*.

EVERY INGREDIENT NOW HAS A GRIP. The little dotted grip on the left of the card is the same one your **Products** screen has always had - this screen had simply never been given it. Hold it and move the card up or down, and your ingredients sit in **your** order: by aisle, by how often you use them, flour first. Nothing else about the card changed.

AND THE ORDER IS YOURS FOR GOOD. It is written onto the ingredients themselves rather than being a position in a list, because **your two phones sync whole records**. A position in a list would never travel between them: you would arrange your shelf on one phone and the other would keep the old order. Written onto the record, both phones agree.

AND NOTHING MOVED UNDER YOU. An ingredient you have never dragged has no order set, so it keeps exactly the place it had - **the list reads the same the moment you update as it did yesterday.** And an ingredient you add afterwards goes to the **end**, without disturbing anything you have arranged.

ONE THING I GOT WRONG AND HAD TO FIND BY LOOKING. The class that marks a card as draggable was already in use elsewhere in the app for something quite different, so your ingredient cards silently picked up a four-column grid and the text was squashed into a column about a third as wide as it should be. **Every test passed** - the test suite has no layout engine, so a squashed card and a correct one look identical to it. It only showed up on a real screen at phone width.

AND THE CARD IS NOW ROOMIER THAN IT WAS. Adding the grip took 21 pixels off the words and 25 off the buttons, which made "Keep at least" fold onto **three lines** for two words. So on this card the buttons now move onto their own line and the name and the two strips get the full width - measured back to **one line** each. The change is scoped to the ingredient card only, so **no other screen in the app moves.**

AND AN OLD OVERFLOW IS GONE WITH IT. The card was already running about 19 pixels wider than a phone before any of this; measuring it is how the class mix-up was found in the first place. It now fits with **nothing** hanging off the edge.

No database step. The suite is **3,168 tests, all green** (2 new). Both new rules were proved by putting the fault back: the list drawn in storage order instead of yours, and a hidden ingredient drawn as one of the rows you arrange.

09 Oct 2026 - engine v398, MORE THAN ONE OWNER (no database step - pushing this one is the whole of it).

Your question: *"investment can be from few owner, how to differentiate"* - and, asked whether money going out should name its owner too, you chose yes.

EACH OWNER NOW HAS THEIR OWN SHARE, ON ITS OWN. The Your investment account keeps a section per owner - their own rows, and **their own balance** at the foot of it:

IMPRESSIVE DIRECTION 3 Oct · Put in RM 2000.00 Impressive Direction still in RM 2000.00

KEAN 1 Oct · Put in RM 1000.00 5 Oct · Took out -RM 300.00 Kean still in RM 700.00

...closing on the totals for everybody together.

AND YOU ALREADY HAD THE LIST FOR IT. An owner is just a name on the **From** list - the one you shape yourself - so it can be **Kean**, **Suan**, **Impressive Direction**, or a pot like **Savings**. Nothing new to learn: the

same list feeds both halves.

AND MONEY GOING OUT NAMES ITS OWNER TOO. When you record a **withdrawal**, the form now asks **"Whose investment does this come out of?" Without that, one owner's balance is impossible** - the app would know what Kean had put in and never what Kean had taken back, so it could only ever show what each owner paid in and one combined figure for what went out.

AND THAT QUESTION IS ASKED OF A WITHDRAWAL ONLY. An ordinary cost - packaging, rent - is not anybody's money going back to them, so it is never asked to pick an owner.

AND MONEY WITH NO OWNER NAMED IS STILL SHOWN. A withdrawal recorded before you had named owners, or one you simply did not attribute, is kept in a section of its own called **Not named**. It is a fact about your books, and quietly dropping it would make the owners' shares add up to more than the total - **which is exactly the kind of figure that looks right and is not.**

No database step. The suite is **3,166 tests, all green** (4 new). Every new rule was proved by putting the fault back: the account keeping one total instead of a line per owner, money with no owner silently dropped, the owner question put to an ordinary cost, and the owner not saved onto the withdrawal - **that last one was found by one of the new tests, and the fault was real: the form asked, and then threw the answer away.**

09 Oct 2026 - engine v397, WHEN YOU PAY FOR SOMETHING WITH YOUR OWN MONEY (no database step - pushing this one is the whole of it).

Your words: *"When i enter an expenses, like buying equipment, should i allow to credit investment account?"*

A MIXER BOUGHT WITH YOUR OWN MONEY IS NOW YOUR CHOICE, ON THE EXPENSE ITSELF. When the way you paid is one of your own pockets, the expense form asks **"Whose money paid for it?"** with two answers:

- **The bakery owes me** - what the app has always done. That pocket's line shows the bakery owes it back, and **Pay back a pocket** settles it when the till has the money. **It is my investment** - the money stays in the bakery. It counts as money you have **put in**, and **nothing is owed back to you** for it.

AND THE TWO CAN NEVER BOTH CLAIM THE SAME RINGGIT. If that money is your investment, the pocket **stops** saying the bakery owes it - otherwise the same RM500 would be counted twice, once as a debt to you and once as money you had put in, and you would appear to be owed money you had already invested. A row that is an investment is therefore left out of the pocket's line entirely, while an ordinary one still shows the debt exactly as before.

AND IT IS NOT ASKED WHEN THE TILL PAID. Paying from Cash or TNG is the bakery's own money - neither owed back nor invested - so the question only appears for a pocket.

09 Oct 2026 - engine v396, YOUR INVESTMENT, AS AN ACCOUNT (no database step - pushing this one is the whole of it).

Your words: *"SO investment is an account, a source of money, we are able to generate investment account.?"* - and, asked whether it should show money taken back out as well, *"yes, with taken back out too"*.

THERE IS A CARD CALLED YOUR INVESTMENT ON THE MONEY SCREEN. It reads what you have **put in**, what you have **taken back**, and what is **still in the bakery** - and pressing **Open** shows every row behind those figures, as a book you can Print or Share like all the others:

put in RM 800.00 · taken back RM 50.00 · still in RM 750.00

AND IT IS DELIBERATELY NOT TIED TO THE TODAY / THIS WEEK / THIS MONTH BUTTONS. Every other figure on that screen answers "what moved this week". This one answers **"where do I stand with the bakery"** - a **balance**, which only means anything counted from the beginning. A week's version would say you were RM200 in on a week you had taken RM900 out.

TWO KINDS OF ROW ARE LEFT OUT OF IT, and both would have made the balance a lie. They are rows the app writes itself under your withdrawals category - which is what keeps them off your Profit statement, and which is exactly why they would otherwise be read as your money leaving:

- **Money moved between your own pots** never left the business. Counting a transfer would say you had taken out money you had merely moved from TNG to Cash. **A pocket being paid back** is the till settling a debt it already owed. That is not you putting money in afresh.

AND IT FOLLOWS YOUR WORD, NOT A FIXED ONE. It counts money taken out **by category class**, so if you rename "My own withdrawal" to anything you like, the account still counts it - a rename must never quietly stop a figure from working.

NOTE ON THE COMMIT. This release and the one above it share the same version and changelog files, so they land in ONE commit. Both are described here.

No database step. The suite is **3,163 tests, all green** (7 new). Every new rule was proved by putting the fault back: the pocket still claiming the investment as owed, what you paid for yourself not counted as invested, the question asked when the till paid, the choice not saved on the row, money between your own pots counted on both sides of the balance, the transfer counted as money taken out, and a renamed withdrawal category no longer counting.

09 Oct 2026 - engine v395, MOVING MONEY BETWEEN YOUR OWN POTS (no database step - pushing this one is the whole of it).

Your words: *"if its for sometimes we want to transfer money from tnG to cash, or to bank"*

THIS IS NOT MONEY IN, AND COUNTING IT AS MONEY IN WOULD HAVE BEEN WRONG. Putting money in says money of your own **entered** the business, and it adds to the Money screen's Cash in and TNG in. TNG to Cash is the **same money moving** - nothing entered and nothing left. Filing it as a money-in would have **counted your own RM200 twice** and inflated the figures you read your takings from.

SO THERE IS A NEW TRANSFER PRESS, beside **Put money in** on the Money screen. Pick the pot it comes **out of**, the pot it goes **into**, how much, and the day - and it writes **both sides at once**:

- the pot it came from goes **down** the pot it went to goes **up the Net does not move at all**, because nothing entered or left the business

WHY BOTH SIDES IN ONE PRESS. Doing the two halves by hand is how a busy morning goes wrong - the same reason **Pay back a pocket** writes its two rows together. And the outgoing half is filed under your own **withdrawals** category on purpose: an ordinary category would put *moving your own money* on your Profit statement as a **cost of trading**, and it is not a cost of anything.

TWO POTS ARE NEEDED TO TRANSFER. The press only appears when you have at least two ways to pay, and the form will not let you send money from a pot to itself.

09 Oct 2026 - engine v394, A "FROM" FIELD ON PUT MONEY IN (no database step - pushing this one is the whole of it).

Your words: *"should have additional field : from xxx"* - and, asked whether it should be typed or picked, *"Which pot it came from - picked"*.

PUT MONEY IN NOW ASKS WHERE IT CAME FROM. A third list, yours to shape like the other two: **My own pocket** and **Savings** to start with, with a **+ New source** chip, and every pot yours to rename or remove on the same lists screen.

IT IS A DIFFERENT QUESTION FROM "PAID IN AS", AND THE FORM KEEPS THEM APART. *Paid in as* is Cash, TNG, Loan - **how** the money went in. *From* is your own pocket, savings, a person - **where** it came from. Money can come out of savings and still go in as cash, so both are asked, and the row carries both.

AND IT REPLACES A PHRASE THAT COULD NOT BE RIGHT. A money-in row has always read **"From my pocket"** whatever your money actually came from. It now reads **"From Savings"**, or whatever you called it - and a row written before this version still says exactly what it always said.

AND THE WORD IS YOURS: THIS IS AN INVESTMENT. Her question - *"can i say put money in is an investment"* - and yes, so the form now says so. **What was NOT taken is the other half of that question**: taking money out is **not** called a profit distribution, because that is only true in a month the bakery actually made

a profit. In a month you made a loss it is still a **withdrawal**, and a label claiming more than your books can back is exactly what this app refuses to print. That category is yours to rename whenever you like, and renaming it moves what you have already recorded.

AND THE LISTS CARD IS NAMED FOR ALL THREE NOW. It read *Categories & ways to pay* - which stopped being true the moment a third list joined them, and a heading missing one of the things under it is how you fail to find that thing.

NOTE ON THE COMMIT. This release and the one above it share the same version and changelog files, so they land in ONE commit. Both are described here.

No database step. The suite is **3,156 tests, all green** (6 new). Every new rule was proved by putting the fault back: the transfer writing only one side, the transfer filing itself as a cost, the sources list never carried to her other phone, renaming a source leaving her recorded rows on the old name, and the pot she picks not being saved onto the row.

09 Oct 2026 - engine v393, UNDO ONE SHOP'S BOUGHT (no database step - pushing this one is the whole of it).

Your words: *"that delete is for deleting the whole po, what if i only want to delete one bought only"* - and, on what to call it, *"call it undo is more appropriate than delete"*.

YOU FOUND A REAL GAP, AND IT WAS MINE. Once a shop was marked bought its packs were on your shelf and its money was on your books, and there was **no way back** - only the list's Delete, which removes everything. I built **Not buying** so that press asks first and cannot be taken back, and I left **Bought** just as one-way without ever giving you the way out.

SO A BOUGHT SHOP NOW CARRIES AN UNDO. Beside *Bought tick 9 Oct, 11:42* there is an **Undo** press. Press it and you are asked first, with the figures in front of you:

"Undo Mydin? The packs it added come back off your stock and the money comes off your books. Your stock: Strong flour 7.5 kg to 0 kg, Butter 400 g to 0 g. RM 48.90 comes off your books. This cannot be undone."

AND IT IS CALLED UNDO RATHER THAN DELETE, WHICH IS YOUR POINT AND A GOOD ONE. A delete removes a record; **an undo puts something back the way it was** - and this reverses **both halves** of what the press did. **A name that promises less than the act is how a control gets pressed by mistake.**

IT REVERSES BOTH HALVES IN ONE PRESS, WHICH IS WHAT YOU CHOSE. The packs come back off your shelf **and** that shop's money row comes off your books. Undoing a shop you did not mean to buy should leave no trace of it, and leaving the money behind would make the undo only half true.

AND IT ONLY TOUCHES THAT ONE SHOP. Undo Mydin and Yen Grocer stays bought, its stock untouched and its money where it is. **Nothing else on the list moves.**

AND THE STOCK CARD IS TOLD. The packs coming back off are written onto the card as their own line - *Un-bought - Mydin (9 Oct 2026)* beside the *Bought - Mydin (9 Oct 2026)* it cancels - so the card goes on explaining the figure at the top of it rather than quietly disagreeing with it.

AND STOCK NEVER GOES BELOW ZERO, SO THE UNDO TELLS YOU THE NUMBER FIRST. A bake may have used those packs since you bought them; if taking them back would go below nothing, the reversal stops at zero and **the box names that lower figure before you agree**, not after.

TWO LISTS DELIBERATELY DO NOT OFFER IT. A list you bought **before v389** carries the whole thing as one flag and may have no money row at all, so there is nothing there that can be put back with confidence - and this app does not invent a reversal it cannot account for. And a shop you marked **Not buying** moved nothing in the first place; that press stays as it was, asking first and then staying put.

NOTE ON THE COMMIT. This release and the one before it (v392 - the buttons that were running off the side of the screen) touch the same version and changelog files, so they land in ONE commit. Both are described here.

No database step. The suite is **3,150 tests, all green** (7 new). Every new rule was proved by putting the fault

back: the Undo press removed, the money left on her books, the stock not reversed, the reversal written down nowhere, the clamp dropped, the undo taking the other shop with it, and the Undo offered on the two lists that must not have it.

09 Oct 2026 - engine v392, THE BUTTONS THAT WERE RUNNING OFF THE SIDE OF THE SCREEN (no database step - pushing this one is the whole of it).

Your words: *"i cant find the delete button"*

YOU WERE RIGHT, AND IT WAS NOT YOUR PHONE. The row of buttons at the top of a saved shopping list - Back, Amend, Print, Regenerate, Delete - **did not wrap onto a second line.** And a button cannot squeeze itself narrower than the word inside it, so once "Regenerate" had taken its share there was nothing left to give, and everything after it was simply pushed off the right-hand edge of the screen. On a 360px phone - a very common one - that row measured **403px across on a 360px screen**: the page ran 108px wider than the phone, and **Delete sat entirely off the edge**, reachable only by sliding the whole page sideways, which nobody would think to do.

AND IT WAS THE SAME ON EVERY SCREEN WITH A ROW OF BUTTONS. This was never about the shopping list; it was the shape every row of buttons in this app is built from. Any screen where the buttons were too many or too wide for your phone had its last one off the edge.

SO THEY NOW WRAP. Too many buttons for the width and the extras drop onto the next line, at their full size, instead of leaving the screen. On your phone the shopping list now reads **Back · Amend · Print** on one line with **Regenerate · Delete** beneath it - with Delete on a line of its own, which for a button that removes something is no bad thing.

AND NOTHING ELSE MOVED. The wrapping only engages where the buttons would otherwise overflow, so every screen where the row already fitted measures exactly as it did before.

AND I HAVE WRITTEN DOWN WHAT THIS COST. No test in this app could have caught it: there is no layout engine in the test suite, so a row that runs off the screen looks identical to one that fits. What is pinned now is the **rule itself** - the button row declares that it wraps, and a test reads the stylesheet and fails if that is ever removed. **A fault nothing can measure is a fault that comes back.**

AND ONE THING I GOT WRONG MYSELF. In the instructions I gave you for testing v391 I said to press Delete "at the bottom" of the list. It is at the **top**. That is on me, not on you - you went looking where I told you to look.

No database step. The suite is **3,143 tests, all green** (1 new). The rule was proved by putting the fault back: the wrapping removed from the stylesheet, and the guard failing by its own message.

09 Oct 2026 - engine v391, DELETING A SHOPPING LIST NOW TELLS YOU WHAT IT LEAVES BEHIND (no database step - pushing this one is the whole of it).

Your words: *"when we delete a po, money paid dont reverse out?"*

YOU WERE RIGHT, AND THE ANSWER IS THAT NOTHING SAID SO. Deleting a saved shopping list has never touched the money you recorded through it. Each **Bought** press writes its own money row tagged to that list, and nothing removes those rows when the list goes - so deleting a list you had bought on left its money sitting on your books, with the list gone. And since v389 a run over three shops writes three rows, so one delete could leave three of them behind.

SO THE DELETE BOX NOW SAYS WHAT IT LEAVES. If there is money behind that list, the box names the exact figure before you agree:

"...The RM 61.40 you recorded on it stays on your books, and its packs stay on your stock - remove the money on the Money screen if you want it gone."

AND THE MONEY STAYS ON PURPOSE - you chose it, and it is the right call. Tidying a document must not rewrite your books. Money that left your purse is a fact about your business, in exactly the way the packs on your shelf are - and your stock has never reversed on a delete either, for the same reason (and because a

bake may have used those packs since). **What was wrong was only that the box stayed silent**, so the money looked like it vanished along with the list. A consequence you have to find out about afterwards is the fault; naming it first is the fix. The row is still yours to remove: the Money screen has its own X on it.

AND IT IS NAMED ONLY WHEN THERE IS MONEY. A list nobody has bought from prints exactly the sentence it always did - a warning about RM 0.00 on every ordinary delete would be noise you learn to skip.

No database step. The suite is **3,141 tests, all green** (2 new). Both new rules were proved by putting the fault back: the money sentence removed, and the delete quietly wiping the money rows. And the second of those is the one that matters - **nothing had ever asserted that deleting a list leaves her accounts alone.**

09 Oct 2026 - engine v390, THE STOCK CARD IS A REAL TABLE, AND YOU CAN TAKE A MISTAKEN LINE OFF IT (no database step - pushing this one is the whole of it).

Your words: *"a stock card should be as clear as a table with row and column" · "certain listed i might want to delete after testing" · "can optimise for desktop as well"*

THE STOCK LIST IS NOW THREE COLUMNS WITH HEADINGS: WHEN · WHAT HAPPENED · CHANGE. Every movement reads straight down its own column, the way your other tables do:

WHEN WHAT HAPPENED CHANGE 9 Oct Baked - Rosemary Focaccia ×2 -500 g 8 Oct Bought - Mydin (8 Oct 2026) +8 kg

AND THIS IS A DELIBERATE REVERSAL, WITH THE REASON RECORDED. At v385 the stock rows were tried as columns and **a phone rendered them unreadable** - the column widths in the app are set for the filing page's five columns, so a stock row came out as "1 Oct..." and "Stockta...", cut off. The fix then was to write each row as one sentence instead. **That avoided the cause rather than fixing it.** The stock card now has **its own column widths and no others** - so this card is sized for these three columns and the filing page is untouched - and the middle column **wraps onto a second line instead of being cut**. A long line now reads "Baked - Rosemary Focaccia" and "×2" underneath it, rather than losing its end.

AND EACH STOCK LINE CAN BE TAKEN OFF THE CARD. A small X sits at the end of every stock row. Press it and you are asked first, and **you are told both figures before you agree**: your flour is 7.5 kg now, and taking this movement back leaves 8 kg. **Removing the line and reversing the movement are one act** - delete a test Bought and the packs it added come back off your shelf. Doing only one of the two would either leave your shelf holding packs you had just deleted the record of, or leave a card whose rows no longer explain the figure at the top. And stock never goes below zero, so if a reversal would, **the confirm says the lower figure before you commit** rather than surprising you after.

IT CANNOT BE UNDONE, which is why it asks. The list is the record of what happened, so a line taken off is taken off.

AND IT IS BIGGER ON A BIG SCREEN. Her words: *"can optimise for desktop as well"*. On a wide screen this card opens at 760px instead of being held to the phone column the rest of the app uses - a table is the one thing that gets better with width, and on a desktop every row now sits on a single line. **On your phone nothing changes**; the card keeps the width it has always had.

ONE THING I FOUND BY LOOKING AT THE RENDERED CARD, NOT BY A TEST. The X button was wired straight to the delete, so **drawing the card ran the delete once per row** - it would have opened a confirm over every line the moment you opened the card, and left no button on any row. Nothing in the test suite could see it, because a function that does its work when it is *built* looks identical in the data. It has its own test now.

No database step. The suite is **3,139 tests, all green** (7 new). Every new rule was proved by putting the fault back: the X running instead of building, a X offered on the printed sheet, the line's own figure dropped from the paper and the message, the removal leaving the shelf alone, and a reversal not clamped at zero.

NOTE ON THE COMMIT. This release and the one before it (v389 - pay for a shop and only then does its stock go in) touch some of the same files, so they land in ONE commit. Both are described above and below.

09 Oct 2026 - engine v389, YOU PAY FOR A SHOP AND ONLY THEN DOES ITS STOCK GO IN (no database step - pushing this one is the whole of it).

Your words: *"the po when bought pressed, it push all ingredient into stock immediately, before i enter how much to pay and by what method, it should not like that, it is a lump sum total, it should allow individual supplier bought and after pay only push into stock"*

WHAT WAS HAPPENING, AND IT WAS WORSE THAN AN ORDERING. Pressing **Bought** on a saved list put **every ingredient on the whole list** into your stock first, and *then* opened the "What did you pay?" box. So the shelf moved before you had said what you paid or how - and if you closed that box and walked away, **the stock was already written and no money had been recorded at all**. And because a shopping list is one lump across every shop, that single press claimed you had bought from **all** of them - Mydin, Yen Grocer and the rest - even if you only went to one.

SO IT IS ONE SHOP AT A TIME, AND THE SHELF WAITS FOR YOUR ANSWER. Every shop's heading on a saved list now carries its own **Bought** press. Press it and it asks what **that shop** cost - and **only when you answer does that shop's packs go into your stock**. Close the box without answering and **nothing has happened at all**: no stock, no money, nothing to unpick.

AND EACH SHOP RECORDS ITS OWN MONEY. What you type becomes its own money-out row, and the row now names the shop - so three shops on one day read as three rows you can tell apart, in your Money screen and in the per-method journal, rather than three identical "Shopping run (PO)" lines.

AND YOU CAN SAY "NOT BUYING". A shop you did not buy from - out of stock, or cheaper elsewhere - can be marked **Not buying**. Nothing comes into stock and nothing is recorded, and the list can then **finish** instead of sitting open for ever waiting on a shop you have decided against. **That press asks you first**, because unlike Bought it takes effect straight away and cannot be taken back; a stray finger on a phone must not settle a shop for good.

AND THE STOCK CARD SAYS WHICH SHOP. A bought row used to read *Bought - 12 Oct 2026 list*; it now reads *Bought - Mydin (12 Oct 2026)*, because with several shops on one day the old wording would have been three identical lines that answered nothing.

AND ONE SENTENCE THAT HAD BECOME UNTRUE IS GONE. The card used to say, once you had bought, *"these packs were added to your stock"* - which is simply false when a shop was skipped or not yet bought. It now tells you where the run actually stands: what is still to buy by name, or, when everything has landed, how many shops were bought and how many you passed on.

ONE THING I FOUND BY LOOKING AT THE FINISHED SCREEN, NOT BY A TEST. Bought is now **safe to mis-tap** - it opens a box that must be answered, so a stray finger changes nothing. Not buying was not, so it asks first.

AND YOUR OLD SAVED LISTS ARE UNTOUCHED. A list you bought before this version was one flag for the whole thing, so it is read as every shop on it having been bought - it shows what happened and offers you no press that would put the same packs on your shelf a second time. Amending a list is now offered until the first shop lands, bought **or** passed on, because an edit rewrites the whole list and one shop's packs are already on the shelf.

No database step. The suite is **3,132 tests, all green** (11 new). Every new rule was proved by putting the fault back: the stock push moved back onto the press, the press buying the whole run again, the money row losing the shop's name, the pay box pre-filling one shop with the whole run's cost, a shop with nothing to buy still offering a press, an old bought list offering to buy again, "Not buying" putting stock in, the note claiming everything was added, the stock row losing the shop's name, the table never drawing the presses, a list of loose estimates losing its press, and the confirm on "Not buying" removed.

09 Oct 2026 - engine v388, THE PO PAGE STOPS JUMPING TO THE TOP WHEN YOU TICK A DAY (no database step - pushing this one is the whole of it).

Your words: *"the po page when click, the page jump, rerender"*.

WHAT WAS HAPPENING. On the shopping-list screen, every time you ticked or unticked a bake day, the whole page threw itself back to the top and rebuilt itself. On a short list you would barely notice; on a long one -

your nineteen days - ticking anything near the bottom dumped you back at the top, and you had to scroll all the way down again to tick the next one.

WHY IT DID THAT. Ticking a day did two things: it redrew the list, and it wrote the day into the address bar so the link could be shared or reopened later. It wrote that address the way a browser reads as **"go to this page"** - which told the app to navigate, so the app obeyed: it emptied the screen, rebuilt the whole thing, and left you at the top. The list was right afterwards either way; **what changed was where you were standing.**

SO IT NOW WRITES THE ADDRESS WITHOUT LEAVING THE PAGE. The address still shows exactly the days you have ticked - so a saved or shared PO still reopens on the right days - and pressing a tick no longer moves you at all. **You keep your place.**

AND THE TWO THINGS THAT LOOK ALIKE IN THE CODE ARE NOT. Writing the address, and navigating to an address, take one line each and produce a nearly identical URL - which is exactly why this was easy to get wrong and invisible in a screenshot. The tick now uses the same mechanism the **Orders** screen already uses for the same reason. A test now tells the two apart, because the screen **looks correct either way** - a test that only checked the list would have passed straight over the fault.

Measured, on the real page: ticking with the document scrolled to 2,500px left it at **0** before the fix, and it stays at 0 once the screen rebuilds - so it was not a flicker you could scroll away from. And the line in the app that looked like the culprit, scrollTop = 0, **was doing nothing at all**: that element is not the thing that scrolls. Emptying the page was what moved you.

No database step. The suite is **3,121 tests, all green** (2 new). Every new rule was proved by putting the fault back: the tick navigating again, the address naming the wrong days, the address left blank where it should say "none", the list not redrawn, and the address written twice.

08 Oct 2026 - engine v387, TAKE ONE SHOP'S LIST TO THAT SHOP (no database step - pushing this one is the whole of it).

Your words: *"now the PO, lump together all supplier in one po is not practical"* - and, asked which problem you meant, *"i want a separate list per shop"*.

AND WHAT I FOUND FIRST, BEFORE CHANGING ANYTHING: THE PO ALREADY SPLIT BY SHOP. Each shop had its own section, its own heading, its own subtotal, and its own **Copy order** and **Message supplier** buttons. I rendered a three-shop run and looked at it rather than assuming. **So the one exit still shared was the printer** - Print gave you every shop on one sheet - **and that was the whole of your complaint.**

SO EACH SHOP'S LIST NOW PRINTS ON ITS OWN. Every shop's heading carries a third button, **Print this shop**. Press it and **that shop's list - and nothing else - goes to the printer**: the other shops' sections are left off the sheet, and **so is the run's total**, because the whole run's figure is not Mydin's figure and each shop's own subtotal is already on its heading. Take Mydin's page to Mydin, print Yen Grocer's on Friday, and hand each shop its own list.

AND THE DAY'S ORDINARY PRINT IS UNTOUCHED. The Print button at the foot still prints every shop together, for the run you do in one go. That one matters more than it sounds: the printing code leaves a marker on the page for the instant it is printing, and **a marker left behind would have blanked your ordinary PO print** - later, somewhere else, with nothing on screen to say why. It comes off two ways, on the print finishing and on a timer, and a test now drives the press and checks it is gone. Nothing in this app tested that before.

No database step. The suite is **3,119 tests, all green** (4 new). Every new rule was proved by putting the fault back: the per-shop press removed, the marker never cleaned up, the browser never asked to print, the other shops not hidden, the run's total printed on a one-shop sheet, the hiding rule left unscoped, and the shops merged back into one section.

AND ONE OF THOSE FOUND A HOLE OF MINE. Merging the shops back into a single section broke no test at all - the buttons were still there and the hiding rule still matched nothing, so everything looked fine. **The one-section-per-shop structure IS the feature**, and nothing was checking it. It is pinned now.

ONE THING I DELIBERATELY DID NOT DO. A saved list is still **one** record, not three. Its lines already

remember which shop each came from, and History already shows them grouped. Splitting the SAVE would change your PO history, the day's \"saved\" memory and the Bought press that adds stock - and you chose the *document* problem. **If you want to tick off one shop as bought and leave the others for the next trip, that is a different job and I will ask you before touching it.**

08 Oct 2026 - engine v386, IT IS CALLED THE STOCK CARD NOW, AND IT IS ON EVERY INGREDIENT (no database step - pushing this one is the whole of it).

Your words: *"i want stock card"* - after *"cannot find it"* twice.

AND YOU WERE RIGHT, AND I WAS WRONG. The thing you were asking for **already existed** - it is what v385 built. It was just **called "Journal"** and **hidden until something had moved**, so on most of your ingredients there was nothing to press at all. **A card you cannot find is a card that does not exist**, and that is what I had given you.

SO IT IS CALLED STOCK CARD, AND IT IS ON EVERY INGREDIENT. Not just the ones with history - every single one. Open an ingredient and the press is there.

AND STOCK COMES FIRST ON IT. It used to lead with the price. You asked for a **stock** card, so the shelf is the first thing on it: the amount on hand at the top, then **STOCK** - every movement with the date, what moved and how much, and **what moved it** - and the price history below that.

AND AN EMPTY CARD NOW SAYS WHAT IT IS FOR. It used to be hidden because an empty page reads as a fault. Instead of hiding it, the card says plainly that it fills up as the ingredient moves - baking takes the recipe off, un-baking puts it back, Bought adds what you bought, a stocktake or Day one sets the real amount. **So an ingredient you have never baked with still tells you what its card is for**, rather than offering nothing.

ONE THING I FIXED AFTER LOOKING AT IT. On an empty card the explanation appeared **twice** - once in the empty message, once in the note at the foot - which is a wall to wade through. The note is for a card with rows, where it explains the columns. An empty card now explains itself once.

No database step. The suite is **3,115 tests, all green**. Every new rule was proved by putting the fault back: the press hidden again until something moved, the press called Journal again, the price section moved back to the top, the empty card turned back into a bare one-liner, the card's own title losing your word for it, and the foot note put back on an empty card.

TWO OF THOSE BITES MISSED THE FIRST TIME, AND THAT FOUND TWO REAL GAPS IN MY TESTS. The section-order rule had only ever been checked on a card with **one** section - so swapping the two sections changed nothing it could see. **A rule about the order of two things cannot be proved with one of them missing.** And the card's own printed title was not checked at all. Both are now pinned properly.

08 Oct 2026 - engine v385, THE JOURNAL NOW SAYS WHERE YOUR STOCK WENT (no database step - pushing this one is the whole of it).

Your words: *"i dont see a journal button, maybe there is never price movement, but why only show when there is price movement, qty movement cannot?"* - and the honest answer was worse than you expected: **quantity movements were never recorded at all.**

WHAT WAS WRONG. Five things move your stock, and every one of them simply overwrote the number on the shelf: baking takes the recipe off, un-baking puts it back, Bought adds what you bought, a stocktake sets the real amount, and Day one sets your starting stock. **None of them was written down.** So the moment you asked "why is my flour at 2 kg?", the app could only reply "that is the number".

NOW IT ANSWERS. The Journal press on an ingredient opens **two parts**: PRICE, exactly as before, and **STOCK** - every movement with the date, what moved and how much, and **what caused it**:

```
1 Oct 2026 · Stocktake  -  4.5 kg to 2 kg · -2.5 kg
6 Oct 2026 · Bought  -  6 Oct 2026 list · +5 kg
8 Oct 2026 · Baked  -  Rosemary Focaccia ×2 · -500 g
```

AND THE TWO FIGURES ARE AT THE TOP: the price your recipes use, and the amount on the shelf. A list of movements without the figure it explains leaves you doing the sum.

THE ROWS ADD UP TO THE NUMBER, AND THAT IS THE WHOLE POINT OF A JOURNAL. I checked it rather than assuming it: over a run of day one -> bake -> buy -> bake -> un-bake -> stocktake, the rows summed to exactly the figure on the shelf. That needed care, because **stock never goes below zero** - so a bake against an empty shelf moves nothing, and a row claiming it moved the recipe's 250 g would have made the journal stop explaining its own number.

AND THE PRESS NO LONGER HIDES. It used to need a PRICE move, so an ingredient whose price had never changed offered nothing at all even if its stock moved every week - which is exactly what you hit. It now appears when **either** part has something in it.

STOCK ROWS SHOW A DASH WHERE MONEY WOULD BE. A stock movement is not money, and the journal's figure column is formatted as ringgit everywhere - so 500 g would have printed as "RM 500.00". The dash is the app's existing way of saying there is no figure in that row, and it is true. **Grams dressed as ringgit would be worse.**

ONE THING I CHANGED AFTER LOOKING AT IT, NOT AFTER A TEST. The first version drew the stock rows as aligned columns. On a phone that was **unreadable** - the column widths are set for the filing page's five columns, so a stock row came out as "1 Oct..." and "Stockta...", the very words you need, cut off. A test asserting a cell's text cannot see a column that truncates it. The rows are one sentence each now, which is also exactly what prints on paper.

IT STARTS EMPTY, AND NOTHING BEFORE IT CAN BE RECOVERED. Every movement that has already happened is gone; there is no way to rebuild it. The journal records from today and I will not invent a history that never existed.

AND THE LOG IS CAPPED AT 50 MOVEMENTS PER INGREDIENT, so the very oldest fall off. I measured the cost rather than guessing: **with every one of thirty ingredients filled to the cap, that is 149 KB - under 3% of what a phone allows the app.** With the ten ingredients that actually bake, it is about 49 KB.

No database step. The log rides on the ingredient, the same way the price log has since v285, so pushing this one is the whole of it. The suite is **3,115 tests, all green** (13 new). Every new rule was proved by putting the fault back: baking logging the recipe's amount instead of what the shelf gave up, baking and stocktaking writing no row at all, a row no longer naming the order or the list, Bought logging the planned packs, a stocktake logging the amount typed rather than the change, a movement of nothing written as a row, the cap removed, the stock section dropped, a stock row's grams put in the money column, and the press made to need a price move again.

08 Oct 2026 - engine v384, A REVIEW'S PICTURE IS PRIVATE UNTIL YOU PUBLISH IT (a database step you run once, then push).

Your words: *"plan the photo fix too"*, then, asked how far, *"Gate it properly, and harden."*

WHAT WAS WRONG. A customer's picture went straight into a part of your storage marked PUBLIC, so it could be fetched by anyone who had its link **the moment it uploaded - before you had approved the review.** The rule on that folder said only approved pictures were readable; it was not true, and had not been from the day it was written. **And neither folder had a size or a file-type limit at all**, so anyone with your shop's public key could have used your storage as free hosting for any file of any size.

THE PICTURE IS NOW PRIVATE UNTIL YOU SAY YES. It is uploaded into a part of your storage that **no anonymous reader can reach - not even with the exact link.** When you press **Publish**, the app moves the picture across to the part your homepage reads, in the same breath as publishing it. Nothing unapproved is on the internet at all.

AND THE OBVIOUS FIX DOES NOT WORK, WHICH IS WHY IT IS TWO STEPS. A picture on a page cannot carry a key - an image tag sends no password. So the homepage can never read a private file, however the rules are written. The only ways to keep a plain picture working and make unapproved ones unreachable are to **move the file at the moment you approve it** (what this does) or to hand out links that expire. Moving it was the choice, because an expiring link would break a picture that had been fine for a year.

IF THE MOVE FAILS, THE REVIEW IS NOT PUBLISHED, AND THE CARD TELLS YOU WHY. The other way

round - publishing and hoping - would put a review on your homepage with its picture missing, and you would have approved something you could not see the whole of. The review simply stays waiting.

AND TAKING A REVIEW DOWN NOW REALLY TAKES IT DOWN. Hiding it stopped the homepage showing it, but the file itself stayed reachable by anyone who had kept the link. Taking it down now deletes the public copy as well - and because the private original is untouched, publishing it again brings the same picture back **without asking the customer for it twice.** If the picture cannot be deleted, the review still comes down and you are told the file is still there: a review you cannot take down is a worse fault than a file left behind.

AND IN YOUR OWN REVIEWS CARD, A PICTURE THAT CANNOT BE SHOWN SAYS SO rather than leaving a blank where a photo should be. That card exists so you can judge the picture before you publish it - a silent blank would read as "no photo" and you could publish something you never saw.

AND BOTH FOLDERS NOW REFUSE ANYTHING BUT A PICTURE, AND ANYTHING OVER 2 MB. That is the hardening, and it is worth having whatever else happens. Your shop shrinks a photo to a small JPEG before it uploads, so the 2 MB ceiling never turns a real one away.

What to run, and in this order. In the Supabase SQL editor, run supabase/reviews.sql **again, the whole file** - it is safe to re-run - and **then push this**. Run it first because it is what creates the private folder: until it exists, a new review simply posts without its picture (the form is built to carry on), and the moment you push, every phone moves to the new behaviour by itself.

ONE THING THIS DOES NOT REACH, SAID PLAINLY. A picture already uploaded sits in the public folder today, and moving a file between folders is not something SQL can do safely. Published ones are fine and stay exactly as they are. **An existing UNPUBLISHED review with a picture is already public and stays public until it is dealt with.** So before running the SQL, count them:

```
select count(*) from reviews where not published and photo <> '';
```

If that says 0, there is nothing to do. If it says more, tell me the number and I will tidy them into the private folder from your app - there is no way to do it from SQL alone.

AND A SECOND SMALL STEP, TO CONFIRM THE MOVE. After running the SQL, this should return true - proving the public folder no longer accepts an anonymous upload, which is the one line everything else rests on:

```
select exists (
  select 1 from pg_policies
  where schemaname = 'storage' and tablename = 'objects'
    and policyname = 'customer uploads a review photo'
    and with_check like '%review-photos-pending%'
    and with_check not like '%''review-photos''%'
) as uploads_are_private;
```

The suite is 3,102 tests, all green (13 new). Every new rule was proved by putting the fault back: the upload sent to the public folder again, the upload handing back a public address, a failed move publishing anyway, the picture's address left out of the publish, a take-down leaving the file up, a take-down forgetting the private name, a refused delete swallowed, the delete helper ready to delete anything, the card never asking for a link, the fallback removed, the anonymous upload put back on the public folder, the private folder declared public, do update weakened to do nothing, the limits dropped, and the public read widened to the private folder. **One of those bites found a weak check of my own**: it matched the *name* of the limits column, which still appears when the value is empty, so it passed with no limits set at all.

08 Oct 2026 - engine v383, A PHONE THAT IS BEHIND NOW UPDATES ITSELF BEFORE THE PASSWORD SCREEN (no database step for this part - pushing it is the whole of it).

Your question: *"how come some app user after keying in pin, but login to an old version app?"*

THE APP HAS ALWAYS KNOWN, AND YOU COULD NEVER SEE IT. When the app opens it checks which build your website is serving, and if this phone is behind it puts an amber strip on screen reading **New version ready**, naming both builds. That check runs deliberately before the password screen - a comment in the code says so: *"so a phone sitting on the lock or sign-in screen still learns it is behind."*

AND THE PASSWORD SCREEN WAS STANDING ON TOP OF IT. The strip sat at layer 25; the password screen is a full-screen opaque panel at layer 80. So the one message that says this phone is behind was being painted **underneath** the screen the person was looking at. I measured it on a real screen: asking what a tap at the centre of that warning actually reaches, the browser answered **the Unlock button**. So someone typed the password, went straight into yesterday's app, and never saw the warning that would have told them.

SO IT RELOADS ITSELF FIRST, BEFORE THE PASSWORD. If the phone is behind when the app opens, it reloads into the current build **before** the password screen appears. You type your password **once** and you are already on the newest build. You never land in the old one at all.

WHY BEFORE THE PASSWORD RATHER THAN AFTER IT. A reset after the PIN was your idea, and the instinct was right - but it has one trap: **the app locks on every open**, so a reload after the password would bring the password back and make everyone type it **twice, every time**, and the only way round that is to remember the unlock across the reload, which weakens the lock. Doing it first costs nothing and **leaves the lock exactly as it was.** Your password is still asked for once per open, as before.

AND IT CANNOT TRAP A PHONE IN A LOOP. It reloads **once per build** and then gives up. That matters, because a website genuinely stuck on an old file - a wedged deployment, a proxy handing out a stale copy - would otherwise reload for ever and the phone would never reach a password screen again. After one attempt it stops, and the strip explains why.

AND THE STRIP NOW SITS ABOVE THE PASSWORD SCREEN. Layer 90, over the 80 of the password panel - so even in the stuck case you can see it and press **Update now** before typing anything. It is only ever drawn when the phone is actually behind; the rest of the time it is hidden and costs nothing. This is the same amber bar, doing at last what it always said it did.

ONE RISK I INTRODUCED AND CLOSED IN THE SAME BREATH. The check now runs on the way in, before the password screen - so a request that hung would have left a phone with no password screen, no app and nothing to press, which is a worse fault than the one being fixed. **The check has a deadline**: after 2.5 seconds it gives up, says nothing, and the app carries on exactly as before. A test drives a request that never answers and proves the app still opens.

No new database step for this part. The suite is **3,089 tests, all green** (8 new). Every new rule was proved by putting the fault back: the anti-loop rule dropped, the attempt not recorded, the reload never firing, the deadline removed, the strip sunk back under the password screen, **the reload deleted from the boot**, and **the reload moved to after the password** - the last two being the wiring itself, which nothing covered until the bites found it.

08 Oct 2026 - A ONE-LINE DATABASE FIX: A REVIEW CAN NO LONGER ARRIVE ALREADY PUBLISHED (a database step you run once - there is NO new app version, because nothing in the app changed).

How this came up. You asked how secure Supabase is after the September report of about 16,000 databases left open to the public. That report was misconfiguration by those owners, not a break-in - so I went through every one of your own access rules and checked. **Your customer tables are all correctly locked** (only your signed-in app can read them) and **nothing was exposed**. But the sweep found one real hole, in your reviews.

WHAT WAS WRONG. Your reviews are meant to be approve-first: a customer posts, and nothing appears on the homepage until you press Publish. The row did start unpublished - but the rule only set the STARTING value, and it did not stop a visitor from sending their review already marked as published. Anyone with the shop's public key (it is in the homepage's own script, as it must be) could post a review straight to your database with that flag on, **skipping you entirely**, and it would go live on the homepage.

WHAT IT WAS NOT. Not a leak. Nobody could read your orders, your customers' names, numbers, addresses, your backups or your receipts through it - I checked each of those and they all refused an anonymous reader. What it broke was your approval step: a spammer, or a competitor, could have put words on your homepage without you ever seeing them.

What to run. In the Supabase SQL editor, run supabase/reviews.sql **again** - the whole file. It is written to be safe to re-run, so this is one paste and Run. The one changed line is the rule on who may add a review; it now

also requires the review to arrive unpublished.

How to confirm it took. Run this and read the single cell it returns:

```
select exists (
  select 1 from pg_policy
  where polrelid = 'public.reviews'::regclass
    and polname = 'customer leaves a review'
    and pg_get_expr(polwithcheck, polrelid) like '%published = false%'
) as insert_is_locked;
```

You want true. Anything else and the fix did not land.

YOUR HOMEPAGE FORM IS UNAFFECTED, AND THAT IS PROVEN RATHER THAN HOPED. The form sends five things - the name, the stars, the words, the language and the photo - and never the published flag, so every genuine review still arrives unpublished exactly as it always did. A test now pins both halves: that the rule requires an unpublished review, and that the form never sends that flag (so a later edit cannot quietly break your homepage).

ONE THING I DID NOT CHANGE, SAID HONESTLY. A photo attached to a review sits in a bucket marked public, so a picture is reachable by anyone who knows its link the moment it uploads, before you approve the review. Its name is a timestamp plus six random characters, so it is not something anyone would find by accident - but the rule that says only approved photos are readable is doing less than it looks. The proper fix is a private bucket and timed links, which changes how the homepage and your admin card load the picture. **Say the word and I will plan it**; I have not sneaked it into this change.

No new app version. Nothing in the app changed, so version.js still reads 382 and no phone will be told to update for this. The suite is **3,081 tests, all green** (2 new), and both new checks were proved by putting the fault back: the rule loosened again, the rule removed, the rule pointed at the wrong role, and the homepage form made to send the flag.

08 Oct 2026 - engine v382, A CUSTOMER'S ORDER HISTORY IS A DOOR TOO (no database step, no SQL - pushing this one is the whole of it).

Your words: *"make the customer history one a door too"*.

OPEN A CUSTOMER AND PRESS AN ORDER NUMBER IN THEIR HISTORY. It opens that order's card on its own bake day, exactly as the Receipt register and the filing page do - and each row carries its own order, so a history of three orders is three different doors.

AND THIS ONE CLOSES THE CUSTOMER'S CARD AS IT GOES. This is the only one of the four that sits inside a card rather than on a screen of its own: without that, the Orders screen would open underneath while the customer's card stayed on top of it, and the press would look as though it had done nothing. I proved that on a real screen rather than assuming it.

THIS WAS LEFT OUT ON PURPOSE LAST TIME, AND THE REASON IS WORTH ONE LINE. A press that cannot be driven to its outcome is a press nobody knows works, and this card had no test of its own - so at v381 it stayed plain text rather than shipping something I could not prove. **The test came first this time, then the door.**

No database step and no SQL. The suite is **3,079 tests, all green** (4 new). Each new assertion was proved by putting the fault back: the door removed, the card left open over the order it opened, and every row given the same order. **Two of my own bites were malformed and said so** - one changed nothing, the other broke the file - and reworking them found a branch in the code that no caller could ever reach, which is gone.

08 Oct 2026 - engine v381, AN ORDER NUMBER IS NOW A DOOR: PRESS IT AND THE ORDER OPENS (no database step, no SQL - pushing this one is the whole of it).

Your words: *"can make the order number clickable to bring us to the order so i can admen it, or look at it detail"*.

PRESS AN ORDER NUMBER AND THAT ORDER'S OWN CARD COMES UP. On the **Receipt register** (More -> Money) and on the **Consolidated invoice**, the order code - #A3F9C2 - is now a press. It takes you to

the order's own bake day, brings its row to the middle of the screen and lights it, and opens the order's card over it: price, Cost each, notes, the lot. Close the card and you are on the right day with the row still lit.

THE RECEIPT REGISTER USED TO TELL YOU TO GO AND SEARCH THE CODE YOURSELF. That sentence is gone with the press that replaces it. **But a receipt whose order is not on this phone still has plain text and no press** - there is nothing to open, and a press that landed on the Orders screen and opened nothing is exactly the kind of control this app treats as a fault.

AND NOTHING CHANGES ABOUT WHAT THE REGISTER IS. It is still read-only: no press on that screen touches a number, and the run of numbers stays untouchable - including by you on a bad evening. A door to the order is not a door to the register.

ON PAPER IT IS STILL PLAIN TEXT. The filing page's Order column is pressable on the screen and ordinary text on the sheet that goes to the printer - a filing page is read on paper, and a link printed in the middle of one is meaningless. It is also **why the paper and the screen cannot drift**: both are drawn from one description of the page, and only the screen layer is allowed to add the press.

A NUMBER WHOSE ORDER WAS REMOVED HAS NO DOOR AT ALL. Those rows are the whole reason the removed numbers are on the page - there is no order to open. They keep their number and stay plain text.

ON THE ORDER'S OWN ROW, AND IN A CUSTOMER'S HISTORY, THE CODE STAYS A PLAIN TAG - and both are deliberate. On the row, it is the label of the thing you are already looking at, with Edit beside it. In a customer's history it is left alone for a blunter reason: **that screen has no test of its own**, and a press nobody has driven to its outcome is a press nobody knows works. Say the word and I will build the test and make it a door too.

AND THE CARD CANNOT COME BACK ON ITS OWN. Once the link has been followed, the address is put back to the day's own address, so the rebuild that happens when the cloud answers, when you save, or when you come back to the app does not reopen the card in your face. I proved that one on a real screen rather than assuming it.

No database step and no SQL - the link is an address in the app, nothing more. The suite is **3,075 tests, all green** (13 new). Each new assertion was proved by putting the fault back and watching it fail by its own name - the link ignored, the address not put back, the row not lit, an unknown code landing in silence, an orphaned order refused, the register's press removed, a void number given a door, and the printed sheet made pressable. **Two of my own assertions were caught being weak by those bites**: one checked that a *wrong* claim was gone without checking the *right* one was made, and one asserted the register's links were all order doors using a check that is TRUE of an empty list - so it would have passed if every door had vanished. Both are fixed.

08 Oct 2026 - engine v380, THE COST OF A LOAF IS FROZEN ONTO THE ORDER, SO AN OLD MONTH STOPS MOVING (no database step, and no SQL - pushing this one is the whole of it).

Your question was: *"when i change the ingredient cost, for age orders, will its COS change?"* - and the answer was **yes**, and that was the fault. Every figure on your Profit screen worked the cost out again, from scratch, out of **today's** ingredient prices. So correcting one flour price quietly rewrote every month you had already closed. **Your words: *"yes, freeze the cost onto the order"* and *"and allow me to adjust it,"*.**

SO THE COST IS NOW FROZEN ONTO THE ORDER, THE SAME WAY THE PRICE ALREADY WAS. An order remembers what it was sold for and it now remembers what it cost you to bake, taken at the moment you take the order. **Change an ingredient price tomorrow and last month's profit does not move.** The price half of that has been there since v70; this is the other half.

AND NO FIGURE ANYWHERE CHANGES ON THE DAY THIS SHIPS. What it writes onto your orders is exactly what they were already showing - the live recipe is where the books were reading the number from a moment before. I measured this rather than assuming it: **sales, cost of sales and gross profit came out identical either side of the freeze, for the day, the month and the year** (I proved it on a seeded set of orders carrying both a single recipe and a four-loaf set built from another product).

AND ONE THING I WILL NOT PRETEND ABOUT. The figure locked onto an order from an earlier month is

today's recipe cost - not what that order cost you at the time. The app has never recorded what a loaf cost on the day it was baked, and there is nothing truer to write. So the Profit screen now says so under Gross profit, in plain words, rather than letting a pinned number look like a history it is not.

AND YOU CAN ADJUST IT. Open any order and press **Edit**: under the price on each line there is now a **Cost each** box. It opens on what the recipe says, and you can type over it - for a batch that cost you more than the recipe says, or a one-off. Beside it, while you are using a figure of your own, a press reads **Recipe: RM 4.10** and puts yours back to the recipe. **The word "Cost each" is printed on screen**, because two number boxes on one line and one of them unlabelled is how a cost gets typed into a price. A line whose recipe is still empty shows no cost at all until you build one.

THE COST NEVER REACHES YOUR CUSTOMER. It rides on the order, and the customer's own track card is built field by field rather than copied from the order - so what a loaf costs you stays yours. That is pinned by a test of its own, because your margin is not something a customer or a competitor should be able to read off a tracking page.

AN ORDER YOU HAVE NOT COSTED STILL MOVES WITH THE RECIPE. A line with no frozen cost, and a line whose own cost you have cleared, goes on following your ingredient prices exactly as it always did. This is a freeze on the sales you have made, not a freeze on the app.

AND A REFUND STILL DOES NOT REFUND THE COST. The ingredients went into the bread whether or not the bread came back, which is the rule your books have always used and this version does not touch.

No database step, no SQL, no upload. The suite is **3,062 tests, all green** (18 new). Each new assertion was proved by putting the fault back and watching it fail by its own name - the frozen cost ignored, a zero locked onto a line, the catch-up overwriting a figure you had adjusted by hand, the cost leaking onto the customer's card, and the undo going missing or doing nothing. **The Profit screen's own note was rewritten too**: it used to tell you that editing a price moves months that have already closed, and that sentence is no longer true.

THIS COMMIT ALSO CARRIES v379 (the sortable filing page) - it was never committed, and both versions ride the same files.

08 Oct 2026 - engine v379, PRESS A COLUMN TITLE AND THE FILING PAGE SORTS BY IT (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"can you allow me to sort the column by its title"*.

Press any heading - Date, Order, Invoice, What, Customer - and the page sorts by it. Press the same one again and it turns round. A small arrow beside the heading says which column is sorting and which way, so a page that has been sorted never looks like one that has not.

AND IT STILL OPENS IN INVOICE ORDER, BECAUSE THAT IS WHAT THE PAGE IS FOR. It is a filing document, read against the serial, and a page that opened in some other order would have to be re-sorted before it could be filed. Sorting is something you choose; Invoice is where it starts and one press takes you back to it.

AND TWO ROWS THAT TIE STILL READ IN INVOICE ORDER UNDERNEATH. Sort by Customer and two orders from the same person keep their filing order within that name, rather than shuffling between two draws - a page that reorders itself when nothing has changed reads as a bug.

THE HEADINGS ARE THE SHEET'S OWN. The sort is wired onto the header row the document already draws, rather than a second set of headings put on the screen for the purpose - so the paper and the screen can never come to promise different things. **On paper the headings are ordinary text**, as they should be.

No SQL, no upload. The suite is **3,042 tests, all green**.

08 Oct 2026 - engine v378, THE FILING PAGE NOW SAYS WHY A NUMBER IS MISSING (no database step, nothing to upload - pushing this one is the whole of it).

Your words, reading your own printout: *"why inv 0001 dont show? it should showing the reason"* - and it should. The page listed 0002, 0003 and said nothing about 0001, which reads as a document with a hole in it -

and **a gap in a receipt run is exactly what an auditor asks about.**

WHY IT WAS MISSING. That page is built from your **orders**, and 0001's order was removed - so there was never an order to draw a line from. **The numbers do not live on your orders; they live in the receipt register.** The page now reads it.

A number belonging to no order gets its own line, in its own place in the run:

```
Date    Order     Invoice    What            Customer   Amount
5 Oct   #AAAA01   #000001    order removed    -            -
5 Oct   #052A1B   #000002    Focaccia ×2     Aunty Bee  RM 32.00
```

AND IT CARRIES NO AMOUNT AT ALL - not RM 0.00. The money for it never existed, and a figure beside it would be a number in a filed document that nobody ever paid. It is drawn muted, so it reads as a cancelled number rather than as a sale, and **it adds nothing to the Total.**

THREE THINGS IT IS CAREFUL ABOUT:

- **"Removed" is judged against EVERY order you have**, never against the period you are looking at. An order that is simply in another month is elsewhere, not gone - and calling it removed would put a false statement on a page you file. **A void number belongs to the period it was ISSUED in**, on your own clock - a number issued at 2am in Penang is not pushed back onto the previous day, which at a month boundary is the difference between it being in this month's filing or not. **If the register cannot be read, the page SAYS SO** rather than being quietly short: *"any number belonging to an order that has been removed is NOT shown here."* A filing page that is quietly short is worse than one that admits it could not check.

SO THIS PAGE NOW NEEDS THE INTERNET TO BE COMPLETE. It draws your own orders first, so you never wait to see your own sales - the removed numbers arrive a moment later and the page redraws. With no connection it is still a complete list of your **sales**; only the removed numbers are absent, and it says that they might be.

No SQL, no upload. The suite is **3,038 tests, all green**.

08 Oct 2026 - engine v377, THE FILING PAGE LEARNS TO SAY WHAT WAS ORDERED, AND TO USE A DESKTOP (no database step, nothing to upload - pushing this one is the whole of it).

Your words, looking at your own October list: *"why no description?"* - and you were right. The page listed numbers without saying **what** any of the orders were, and the item names are how you recognise the order you are about to file against. **So the WHAT column is back**, and the page now reads:

```
Date    Order     Invoice    What           Customer            Amount
5 Oct   #052A1B   #000002   Focaccia ×2    Ms Yeap Lian Sioh    RM 39.90
7 Oct   #F4470C   #000003   Sourdough ×3   Mr Tan Eng Boo       RM 105.00
2 Oct   #E3C190   none yet   Focaccia ×2    Neoh Choo Leong      RM 32.00
```

AND IT BROUGHT BACK THE "no price" MARK with it - a line nothing can price says so on its own row again, the way every other screen does, instead of only being mentioned in the note underneath.

AND YOUR DATA SHOWED A SECOND THING THE PAGE WAS GETTING WRONG. Four of your rows read **none yet** while only **one** of them was still to collect. **"none yet" means there is no invoice NUMBER - it does not mean the money is missing** - and three of those four were paid long ago, before you started issuing receipt numbers. The page now says so:

1 of these has no invoice number yet but is already PAID - they came before the numbering started. Open each order's Invoice and it takes the next number.

THAT LAST SENTENCE IS ALSO THE FIX FOR THEM. Those orders can be numbered: **open each one's Invoice** and it claims the next serial, exactly as designed in v360.

AND THE PAGE NOW WIDENS ON A DESKTOP. Your words: *"that page can be optimise for desktop brouwser"*. **The whole backoffice is capped at a 540-pixel phone column, on every screen at every size** - right for the screens you tap through, wrong for a filing page you read and print at a desk. **The cap is lifted for this page alone**, and put back the moment you leave it, so no other screen moves.

A4 IS FINE, AND THE PAPER IS PROTECTED FROM THE PHONE. The printed sheet is already capped at **180mm** - A4's text width - and every one of the six columns has room inside it. The smaller type the phone uses is now written screen and, so **paper can never inherit it**; it stayed off paper by arithmetic before, and a rule that only stays right by arithmetic is one margin change away from shrinking a document you are filing.

No SQL, no upload. The suite is **3,033 tests, all green**.

08 Oct 2026 - engine v376, THE CONSOLIDATED INVOICE IS NOW A FILING LIST (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"in the consolidated invoice, can we have a column for order no. and a column for invoice no, sort it to inv will allow us to printout for filing purpose"*.

AND THE INVOICE NUMBER YOU WANTED ALREADY EXISTED. The app has been issuing one since v360 - the receipt serial, #000001, #000002, ... - the moment you record money. **It is a real, unbroken, never-re-used series, which is exactly what a filing folder is read against.** So this needs no new numbering, no counter and no SQL. The **order code** sits beside it because that is what finds the order again.

The page now reads:

```
Date    Order    Invoice     Customer      Amount
5 Oct   #C2FDA5  #000001     Aunty Bee     RM 32.00
5 Oct   #A1B2C3  #000002     Mei Ling      RM 18.00
20 Oct  #D4E5F6  #000003     Aunty Bee     RM 36.00
20 Oct  #998877  none yet    Uncle Tan     RM 16.00
```

In invoice-number order, under a header that names each column. An order you have **not recorded as paid has no serial** - a receipt is for money received - so it keeps its place on the page, says **none yet**, and **sorts to the end** where it cannot be mistaken for part of the numbered run.

TWO THINGS THIS DELIBERATELY TOOK AWAY, both because you chose the flat layout when shown both:

- **The per-customer grouping and its subtotals are gone.** The per-customer statement is still one press away - the **Whose orders?** box narrows the page to one person and the Total at the foot becomes theirs. **The item column is gone.** The layout you drew is Date · Order · Invoice · Customer · Amount, and the detail lives on the order's own invoice, one order code away. **And the "no price" warning moved with it** - an order nothing can price still adds RM 0.00 to the Total, so the note now says how many those are and to check the product's price, rather than leaving a confident nothing unexplained.

AND THE PHONE NEEDED THREE TRIES, FOUND ONLY BY LOOKING. The columns were sized from arithmetic twice and got it wrong twice - the first pass squeezed the **Customer** column to **0 pixels** (a header reading "Custo..." and no name on any row), and the second cut the **Invoice number itself** to "#0000...". **On a page whose whole purpose is filing, an unreadable serial is the one thing that must not happen.** The columns are now sized for the serial first, with the type one step smaller on a phone - and **the printed page is unaffected**, because A4 is far wider and none of it is cut there.

No SQL, no upload. The suite is **3,031 tests, all green**.

08 Oct 2026 - engine v375, THE FORGET QUESTION WAS COUNTING ITEMS, NOT ORDERS (no database step, nothing to upload - pushing this one is the whole of it).

FOUND BY LOOKING AT IT, NOT BY TESTING IT - and that is the point. Asked to confirm the new Forget press on screen, the question read **"...from 2 orders"** for a customer whose own card, right behind it, read **"1 order · 3 units"**. The two items were two **rows** of **one** order, and the count was counting rows.

A NUMBER THAT CONTRADICTS THE SCREEN IT SITS ON IS WORSE THAN NO NUMBER - and this is the sentence you weigh before an act that cannot be undone. It now counts **orders**, so the question and the card agree.

AND THE TESTS WERE GREEN WHILE IT WAS WRONG. The suite asserted the number, the number was the row count, and the row count was what the code returned - so everything agreed with itself and nothing

agreed with you. **The live look is what caught it**, which is why it is worth doing even when every test passes.

No SQL, no upload. The suite is **3,027 tests, all green**, and the fix is bitten: count rows again and the test fails by name.

08 Oct 2026 - engine v374, A REFERENCE ON THE CONSOLIDATED INVOICE, AND A REAL "FORGET ME" (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"build the invoice number and the delete control"* - the two things I put to you as **the ones a business gets ASKED for**.

THE CONSOLIDATED INVOICE NOW CARRIES ITS OWN REFERENCE: CI-MONTH-2026-10, CI-WEEK-2026-10-04, CI-DAY-2026-10-05, CI-ALL, and a short tag when it is one customer's. It leads the line under the title, so it is on the screen, on the paper, in the shared message and in the PDF.

IT IS DERIVED FROM THE DOCUMENT, NOT COUNTED - and that is the whole design. Re-printing October's statement must carry the same reference. A counter would turn one statement into two different documents for the same money the second time you opened it.

AND IT IS A REFERENCE, NOT A TAX SERIES, SAID PLAINLY. A unique, sequential invoice series is what an e-invoice filing asks for, and **the rules for that are not something I will guess at.** Your per-order invoice number is still the order's own code, deliberately, since v293. This makes the document **referenceable**; the series question stays open until those rules are checked.

AND THE "FORGET" PRESS NOW WORKS FOR EVERYONE - WHICH IS A PROMISE YOU WERE ALREADY MAKING. Your privacy notice tells customers the sales record is kept and **their details are deleted on request**. Until now that was done **by hand in Supabase**, because the press was deliberately withheld from anyone who had ever ordered: deleting their record alone would have thrown away their reward and their note while leaving their name standing on every sale - **and a button that does half of what it says is worse than no button. That is no longer the case, because the orders are cleaned too.** So the press is offered on every customer, and it does the whole of what it says.

What it removes: their name, number, address, any dropped pin, any note and any line note, from every order of theirs - plus their coupons and hand-outs, and their name off anyone else's records.

What it keeps - and this is the point: the sales record. The amounts and the dates stay, and their orders stay together as **one row reading "Details removed"**, so your takings, your register, Profit and a per-customer statement all read exactly as they did before.

THREE THINGS IT DELIBERATELY DOES NOT TOUCH, EACH FOR A REASON:

- **A coupon that is already coming off an order.** Only their *unspent* coupons go, because a spent one IS the discount on a promised order - removing it would put that order's price back up. **The "referred by" number on their order.** That is the **referrer's** number, not theirs - clearing it would change your referral counts and make the Give-coupon press vanish. **The sales record itself.**

AND IT NAMES WHAT IT CANNOT REACH, because a promise kept only as far as the app is not the promise: your **dated cloud backups** (until they age out), any **backup file you have downloaded**, the copy in your **Supabase project** (until the next sync), and any **courier booking**, which sent their name, phone and address to the courier's own servers and cannot be recalled. The confirmation says all of this before you press it.

THE ONE THING A CONTROL CANNOT CLEAR, and it is yours to run. The shop sends the customer's name, number and address into your Supabase project with every order, and that raw copy is kept for orders that never imported cleanly. To see whether any are there:

```
select id, created_at, status, data
from public.incoming_orders
where status = 'new'
order by created_at;
```

To blank one, with the id from that list:

```
update public.incoming_orders set data = '{}' where id = '<paste the id>';
```

No SQL for the app itself. The suite is **3,027 tests, all green**.

08 Oct 2026 - engine v373, AN "ALL" SCOPE ON THE CONSOLIDATED INVOICE (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"pls add a selection ALL, on top of A DAy, A week, a month"*.

THERE IS NOW A FOURTH SCOPE, AND IT COMES FIRST: All · A day · A week · A month. All is everything she has ever sold, with no window on it at all - so a customer's whole history is one document, and so is her own record of the lot.

AND IT CARRIES NO STEPPING ARROWS. There is nothing to step when you are looking at everything, and a press that cannot do anything reads as a broken screen - so the row keeps the word **Everything**, which still says what is on the page, and simply has no arrows on it. The three periods keep theirs.

"ALL" IS NOT "A VERY WIDE PERIOD". It does not widen the window; it removes the window. **The three periods are untouched** - a month is still exactly a month, and a test pins that so a later change cannot quietly turn every scope into everything.

AND ITS EMPTY STATE DOES NOT OFFER ANOTHER PERIOD, because it is not a period. With nothing sold at all it says **"Nothing has been sold yet."**

No SQL, no upload. The suite is **3,016 tests, all green**, and the new rules were bitten - widen the range instead of skipping it, draw the arrows anyway, or move All back to the end of the row, and each fails by name.

08 Oct 2026 - engine v372, ONE INVOICE OVER A PERIOD - A DAY, A WEEK, A MONTH, OR ONE CUSTOMER (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"i need a month consolidated invoice printing page, selectable individual, daily, monthly"* - and asked what one document should cover: *"per day, per week, per month, per customer as well"*.

WHERE IT IS: More -> Money -> Consolidated invoice. Pick **A day · A week · A month**, step back and forward with **left / right**, and pick **whose orders** - everybody, or one customer off the list. The document lists every order in that scope, grouped under each customer with a subtotal, and ends on one Total. **Print** and **Share** are the same two presses every other book in the app wears, so the PDF that goes to WhatsApp is the page you were reading.

IT IS AN INVOICE, WHICH IS A DIFFERENT NUMBER FROM YOUR TAKINGS - ON PURPOSE. Your takings leave out the courier charge, because that money is passed to the courier and was never yours. An invoice includes it, because the customer was billed it. **So this document uses the invoice's own figure, not the takings figure** - otherwise it would disagree with the sum of your own individual invoices by every courier charge in the period, which is the one thing a consolidated invoice must not do.

AND IT FOUND A REAL FAULT IN THE SHARED SHEET RENDERER. Every document the app prints has always supported a **heading** and a **subtotal** - the paper, the shared text and the PDF honour them. **The screen did not.** So a heading would have shown on your phone as *"AUNTY BEE RM 0.00"* while printing correctly underneath it. It was invisible until a document first had a heading in it, which is this one. Fixed, and proven: the same sheet now draws the same rows on the screen, on the paper and in the message.

FOUR THINGS IT SAYS RATHER THAN HIDING:

- **An order with no bake day** cannot belong to any period. It is **counted and named** at the foot of the document rather than quietly left out. **An order refunded in full** is off the rows and **said** in the same line - a document that is quietly short is worse than one that admits it. **A line nothing can price** is marked **"no price"**, not printed as a confident RM 0.00. **A walk-in with no name and no number** - every one of those used to be its own "customer" on a list. They now share one heading.

AND THE WORD "WEEK" MEANS THE SUNDAY WEEK HERE, COUNTED BY BAKE DAY. The app has two weeks: the Sunday one the Home screen names (*"Week of Sun, 4 Oct"*) and the Monday one the delivery runs use. This document uses the Sunday week, counts everything by **the day the order is FOR**, and **says so on the paper**. That is also why its figure will not match the week's figure on Home - Home counts orders by the

day they were **placed**. Both are right; they answer different questions.

No SQL, no upload. The suite is **3,012 tests, all green** (up twenty-one).

08 Oct 2026 - engine v371, THE REFUND CARD NOW SHOWS WHAT WAS CHARGED (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"can the refund also show the original charges amount as well??"*

AND IT CAN, because the app already knows them. The refund card now opens with the order's own receipt - **Items total · Courier charge · any Promo · any Bring-a-friend discount · Total** - drawn in the same soft band the rest of the app uses for "How it adds up", and **above** the tick list so you decide against the real figures rather than from memory.

AND IT IS THE SAME RENDERER THE INVOICE AND THE CUSTOMER'S OWN MESSAGES USE, not a second sum drawn on this card. Two renderings of one order's money is two figures that can disagree - and the screen you decide a refund against is the last place that should happen. So what the customer was quoted and what you refund against cannot drift apart.

THE DISCOUNT ROW IS THE POINT OF IT. Items less the coupon **is** the money that came in - which is what makes the limit at the bottom of the card make sense. On a RM32 order carrying a RM3 coupon the box now reads **"How much back? (up to RM 29.00)"**, and you can see why.

No SQL, no upload. The suite is **2,991 tests, all green**, and the new assertion was bitten: take the charges off the card and the test fails by name.

08 Oct 2026 - engine v370, A REFUND YOU CAN CHOOSE: HOW MUCH, AND WHAT (ONE SQL STEP FIRST, see the top of the entry).

Your words: *"refund should not be a full without choice to how much to refund, what to refund"* - and when asked which shape, you picked **tick the items, then adjust the figure**.

WHAT THE REFUND PRESS DOES NOW. It opens a card instead of a two-line question:

Refund this order? Which came back? Focaccia ×1 of 2 ..... RM 16.00 Sourdough ..... RM 18.00 Courier charge ..... RM 8.00 Ticked so far ..... RM 24.00 **How much back?** [RM 24.00] adjust it if you want to Why (optional) [ ] [Cancel] [Refund]

Tick what came back and the figure fills itself in. Then type over it if you want to - a rounded figure, goodwill, or just the postage. **Once you have typed, the figure is yours** and moving the ticks will not overwrite it. On a line of two, a **quantity** appears so you can return one of them.

IT WILL NOT LET YOU GIVE BACK MORE THAN THEY PAID. That sounds obvious and it is the whole reason for the next paragraph: the limit is **what the customer actually handed over**, not what the bread was priced at.

AND THAT FORCED A SECOND FIX, WHICH IS THE IMPORTANT PART OF THIS VERSION. Your takings were counted at the goods' **full price, with the discount still in them** - so a RM16 order with a RM3 coupon read as RM16 in the till even though the customer handed over RM13. **While a refund took the whole order, that did not show**, because the order was simply skipped. **The moment a refund can be a part of an order, it stops cancelling** - and a full refund of a discounted order would have left a phantom RM3 sitting in your takings.

You chose "count what the customer actually paid", knowing it would move your past figures. So a discounted order now counts at what came in, and **a full refund of one reaches nothing** - where it would have left the discount behind.

WHAT MOVES IN YOUR PAST FIGURES, said plainly: any order that carried a **bring-a-friend coupon or a promo code** now counts at what the customer paid rather than the bread's full price. Orders with no discount are **unchanged to the cent**. Money, Profit, the week's figures on Home, and each customer's lifetime spend all move together, so the screens still agree with each other.

AND A REFUND IS VISIBLE IN THE JOURNAL. Tap **Sales** on Profit and the refund and the discount appear

as their own lines, so the rows still add up to the figure above them - the one thing this app calls a bug in every book it draws.

THE ONE STEP, AND IT MUST BE DONE BEFORE YOU PUSH: run supabase/receipts.sql in your Supabase SQL editor again. It adds a column and replaces three functions. **Without it, a partial refund on one phone would arrive on your other phone as a full one** - the same order reading RM5 back on one handset and RM22 on the other. It is safe to re-run and changes nothing already recorded.

WHAT I DELIBERATELY DID NOT DO. A promo code's **usage count** is untouched: a code still counts as used, and still counts what it gave away, even on an order that has since been refunded. Re-opening a spent code because bread came back would surprise you, and it is a different question from this one - say so if you want it changed.

No new Supabase step beyond that one file. The suite is **2,990 tests, all green**, including all five refund tests from v361 **unchanged** - which is what proves a refund you made before today still reads exactly as it did.

08 Oct 2026 - engine v369, THE ACCENT COLOURS NOW CLEAR THE SAME FLOOR AS EVERYTHING ELSE (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"build the contrast fix, brand colour stays"*.

HOW IT WAS FOUND. By measuring every screen in a real browser for the first time - the same access you granted so I could look rather than guess. The app set itself a **4.5:1 floor** in v313 and applied it to the small grey lines. **That fix covered one colour only.** Everything else - the amber, the green, the red - had never been measured where it actually sits.

WHAT WAS UNDER THE FLOOR, and what each one measures now:

| where | before | now | |\---|\---|\---| | Ingredients' empty-state line (amber on amber) | 3.78 | **4.96** | | Promo's "point of no return" badge (white on amber) | 4.23 | **4.91** | | The **engine pill**, a valid promo chip, a "Cool" chip | 4.29 | **4.83** | | The Production "mixer" chip, a Scenario note | 4.31 | **4.55** | | The "Email the full wish list" link | 4.41 | **5.92** |

AND THE WORST OF IT WAS NOT IN THE BACKOFFICE AT ALL - IT WAS IN YOUR SHOP. The shop had written **its own greys** rather than using the shared one, and three of them were genuinely hard to read on a phone:

| where | before | now | |\---|\---|\---| | The customer's own **tracking page** step labels, at 9.5px | 2.40 | **5.40** | | The line naming the Act, closing the **privacy notice** | 3.24 | **5.40** | | The small line under **Place order** - the privacy panel's own trigger | 4.08 | **5.40** |

ALL THREE ARE THINGS YOU RULED MUST BE READABLE, and the notice is the one you were most particular about. They had drifted because a colour written by hand cannot be seen by a test that checks a shared colour - which is exactly the fault v313 set out to end, surviving in the one file it did not cover.

AND ON YOUR INSTRUCTION THE BRAND TERRACOTTA DID NOT MOVE. White on it measures **4.48:1**, just under the floor - and that is now the **only** pair left under it, deliberately. **The fix was not to change your colour; it was to stop using it for TEXT.** Eighteen places were drawing text in the brand fill colour; they now use the darker terracotta the app had already been using for text in about twenty-five other places. So the brand is untouched and the words are darker.

AND ONE THING I FOUND THAT WAS NEVER A STYLE CHOICE AT ALL. The WhatsApp number inside the privacy panel asked for a colour called --terra - **which does not exist in the shop.** It fell back to the brand fill colour (4.41:1) and read as if someone had chosen it. It now draws the brand text colour (5.92:1), and a test refuses any stylesheet here that reads a colour belonging to another document.

WHAT I DELIBERATELY DID NOT TOUCH, and why - a faint colour is sometimes the point. A delivery day already gone, a sold-out product's title, the board's eight person-colours, the review stars, and the printed label sheet all sit below the floor **on purpose**: they are states and marks, not sentences, and the note written at v313 warns specifically against re-deriving the past-day colour. **None of them was changed.** If any of them reads badly to you on the phone, say so and I will treat it as a fault rather than a style.

ONE THING I COULD NOT MEASURE, SAID RATHER THAN GUESSED. The inactive language chips (Mandarin, BM) sit on the shop's **gradient** header, and a gradient cannot be measured the way a flat colour can. My first pass reported them as unreadable and **that number was wrong** - my own instrument, not your app. They need your eyes: if the small Mandarin looks washed out against the brown, tell me and I will deal with it directly.

Measured: every screen re-measured live after the change - **the backoffice now has nothing below the floor on any of the 26 screens**, and the shop's privacy notice reads 5.40 where it read 3.24. The suite is **2,980 tests, all green**. Five bites, all restored byte-identically.

08 Oct 2026 - engine v368, EVERY SCREEN NOW BUILDS ITSELF ON EVERY CHANGE (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"if you can build check list why not check yourself"* - and you were right to ask.

WHAT I CAN AND CANNOT DO, SAID PLAINLY. Opening your admin app in a browser has been refused to this session, and that refusal is not mine to work around - so **the one thing I cannot do here is LOOK at a screen.** But almost everything a checklist would ask a person to do - does it open, does it crash, does it draw anything at all - I can do by building the real screen and pressing it. **So I have, for every one of them.**

ALL 26 SCREENS, TWICE EACH, ON EVERY CHANGE. Every screen in the app - Home, Orders, Products, Customers, More, and all twenty-one behind it - is now built for real **on a fresh install** and again **with real data in it**, and two things are asked of each: that it does not crash, and that it **draws something**. A screen that opens blank and a screen that opens empty look identical to you and mean opposite things, and now neither can ship unnoticed.

AND THE LIST OF SCREENS IS THE APP'S OWN, NOT MINE. The test reads app.js's route table and fails by name if a screen exists in the app but is not covered - so a screen added later cannot quietly miss the pass. A checklist I wrote by hand would have drifted the first time we added a screen.

WHAT THIS DOES NOT PROVE, AND SAYING SO IS PART OF IT: it proves a screen CONSTRUCTS, not that it LOOKS right. Spacing, contrast, whether a button is reachable by a thumb, whether the words are ones you would use, whether it is what you MEANT - none of that is visible to a machine. Those still want your eyes on a phone, and this release does not pretend otherwise.

AND IT IMMEDIATELY EARNED ITS KEEP - TWICE OVER. First, the very first run "failed" three screens, and **all three failures were the test's own fault, not the app's** (a stand-in that answered "not found" for things a real page would have found, and an element with no style object). That is worth knowing: a harness that cannot express what a screen does invents faults, which is as useless as one that hides them - so the stand-in was made capable, and it now THROWS on a selector it does not understand rather than quietly finding nothing. Second, it was then **bitten twice to prove it works**: the exact bug that nearly shipped on v367 (a screen that lost one import and would have crashed on your phone) and a screen made to draw nothing at all. Both failed by name, both restored byte-identically.

Measured: the suite is **2,976 tests, all green** (up fifty-three). **No SQL, no upload.**

08 Oct 2026 - engine v367, CLEARING OUT THE COUPONS THAT ARE ON NO ORDER (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"not sure, i think for now you can remove all coupon first"* - said about a stale **Apply coupon** on an aged order.

WHERE IT IS: More -> Bring a friend -> Clearing up -> "Coupons not on any order". It says how many are sitting there, how many still work, and what they are worth together, and a press clears them.

AND IT DELIBERATELY DOES NOT REMOVE *ALL* COUPONS. THIS IS THE IMPORTANT HALF. A coupon that is already coming off an order **IS** the discount on that order's Total - so deleting it would **put the price of an order you have already promised back UP**, from a cleanup screen, with nothing on screen saying so. That is not a cleanup; it is a fault wearing a cleanup's clothes. So this clears **only the coupons that are on no order at all**, which is exactly the set causing a stale **Apply coupon** to appear - and the card says so, in your own words:

"Coupons already coming off an order are NOT touched, so no order's Total changes."

What it removes and what it keeps. A coupon the app can currently find on a live order is kept, by the app's own definition of "on an order" - the same function that prices the order, so the two can never drift apart. Everything else goes: an unspent reward, a friend's discount whose order was removed, a coupon you added by hand and never used.

AND IT ASKS FIRST, NAME BY NAME. The question says how many, what they are worth, that a customer holding one will no longer have it honoured by the app, and that it cannot be undone from here - and the press is dressed as a dangerous one, under its own heading, away from the boxes you use every week. **With nothing to clear it is quietly inert and says so** rather than offering to remove nothing.

AND THAT IS ALL THIS VERSION IS. It is a cleanup, not a diagnosis: I could not reproduce your coupon on an aged order, and I will not guess at a fix. The clearing takes the stray coupon away so it stops misleading you, and the question of HOW two screens came to disagree is still open, waiting on one fact from you.

No SQL, no upload. The suite is **2,923 tests, all green** (up nine). Two bites on the rules, both fired and restored byte-identically - the naive "remove everything" (which fails the test that proves an order's price does not move) and "nothing is ever recognised as on an order"; plus one on the screen itself, removing the question so the press acted silently - and the view test **caught a real bug of mine before it shipped**: a variable in that file named cur means the settings object, not the currency, so the confirm printed *"worth [object Object] 3.00"*.

08 Oct 2026 - engine v366, THE RECEIPT REGISTER - and a hole in the run closed (ONE SQL STEP FIRST, see the top of the entry).

Your words: *"build the receipt register screen"* - asked for the day after you removed a paid order and watched #000001 stay spent.

WHERE IT IS: More -> Money -> Receipt register. It is the whole run of receipt numbers, oldest first, one line each: the number, the order's code and customer, the moment it was issued, and a **refunded** chip where there is one. **Receipts #000001 to #0000NN - N issued** heads it, and under it comes the one sentence that matters: whether the run is unbroken.

AND IT FOUND A REAL HOLE - ONE YOU WOULD NEVER HAVE SEEN ANY OTHER WAY. The claiming function draws a number with nextval *before* it knows whether the row will land. If your two phones ever asked for the same order's number in the same instant, the second one drew a number, the row was refused, and **that number was simply lost** - a missing receipt in a run whose entire legal purpose is to have no missing receipts. v360 wrote that down as a known trade. **It did not have to be one, and it is now closed:** the function takes a lock on the ORDER CODE first, so one order cannot be claimed twice at once, and no number can be drawn and then dropped. **This is why the receipt register exists** - nobody could have spotted that from an order screen, or from a total.

THE ONE STEP, AND IT MUST BE DONE BEFORE YOU PUSH: run supabase/receipts.sql in your Supabase SQL editor again. It replaces the claiming function. It is safe to re-run and changes nothing already recorded - every number already issued keeps its number. Until it is run, the register still works and nothing is broken; only the lock is missing.

A MISSING NUMBER NOW WEARS A BANNER OF ITS OWN, naming the numbers, and says plainly that they cannot be filled in afterwards - a missing number is one the sequence drew and never issued, and using it would put two receipts on one number. **A clean run does not wear that banner**, because a warning shown where it does not apply is noise.

AND THE VOID RECEIPTS ARE VISIBLE AT LAST. A number whose order is no longer in this phone's orders is shown with its number intact and marked *"not in this phone's orders"* - the case you found by removing a paid order. The screen says the receipt stands, the money moved, and the number stays spent. It is worded as **what was seen on this phone**, not as a verdict: the register is shared, your orders are not, and a phone that has not finished syncing must not be told it has a hole it does not have.

AND IT IS READ-ONLY, WITH NO CACHE. There is no button on it that changes a number - a register you

can edit is not a register - and it re-reads every time it is opened, so it can never show you yesterday's run while a receipt was issued this morning. **If it cannot be read, it says so and says nothing was changed**; it never draws an empty register, because an empty run and an unreachable one mean opposite things and this is the one screen where that mistake would matter.

AND IT READS STRAIGHT FROM THE TABLE THAT ALREADY EXISTS - no new Supabase step for the screen itself. The read permission has been in receipts.sql since v360. The only new SQL is the lock above.

Measured: the suite is **2,914 tests, all green** (up fourteen). Five bites, all fired and restored byte-identically: an unreadable register drawn as an empty one; the gap banner removed; the void-receipt note removed; missing numbers before the first row not reported; a failed read reported as a success.

08 Oct 2026 - engine v365, A REMOVED ORDER'S RECEIPT NUMBER IS NOW SAID OUT LOUD (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"what happen with old order before today"* and *"i happen to enter a new order and pay on it then delete the order, whay happen, it took inv 001, then i go back to aged order and click inv, it then took inv002, even the 1st inv already deleted?"*

FIRST: what you saw is CORRECT, and it is the whole point of a numbered series. A receipt number is issued the moment you record the money - not when the order is taken, and not when the paper is printed. Removing the order afterwards does not give the number back, because re-using a serial would put **two different sales under one number**, which is the one thing a numbered series exists to prevent. So **#000001 stays in your receipt records as a void receipt, and the next receipt takes #000002.** That is what an auditor wants to see. Nothing is broken and nothing is lost.

WHAT WAS WRONG IS THAT NOTHING TOLD YOU. Removing a paid, numbered order asked a bare *"Remove order ...?"* and said nothing about the number - so the one thing on that screen you cannot undo was the one thing the screen did not mention. **Now, when the order you are removing has a receipt number, the question names it and says what happens:** that the number was issued when the payment was recorded, that removing the order does **NOT** cancel it, that it stays in your receipt records, and that the next receipt takes the next number. **An order with no receipt number is removed exactly as before, with no mention of receipts** - a warning shown where it does not apply is noise, and noise is what makes a real one unreadable.

AND SECOND: what happens to an order paid before receipts were numbered. Nothing was backfilled and nothing will be - those orders simply sit **outside the series**, which is not a gap in it. They are already handled, and here is the whole of it: **an old paid order gets its number the first time you press Invoice.** But its unnumbered note was **wrong**, and this release fixes the wording. It used to read *"this order has not reached the bakery's records"*, which is false twice over - the order is right there on the screen, and for an order paid before the feature existed the records have nothing to do with it. The thing that has not happened is the **drawing of the number**, so that is what it now says, and the advice it gives is the one that works: open the invoice again once the phone is back online.

Still open, and said rather than hidden: there is **no screen that shows you the receipt register** - the void receipt from a removed order is recorded in your Supabase project but you cannot look at it from the app. Worth building before you pass RM150,000; ask when you want it.

No SQL, no Edge Function, no upload. The suite is **2,900 tests, all green** (up two), and all three new rules were **bitten**: put the old wording back and one test fails by name; make the remove question forget the number and one fails; show the receipt sentence on an order that never had a number and one fails.

08 Oct 2026 - engine v364, THE APP NOW SAYS SO WHEN A SHOP ORDER IS WAITING (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"build"* - the half that was still missing after v363.

The half that was missing was not the fault, it was the SILENCE. A shop order the app could not read is left waiting in the queue on purpose, and it retries every 30 seconds - **but it appeared on no screen at all.** The customer waits, and **nothing anywhere told you.** v363 removed one cause of that; **this removes the**

blindness, whatever the cause.

What changed. On the **Orders** screen, above the New-orders inbox, a small amber note now appears whenever the shop has sent an order this app cannot read. It says **how many**, **what each one is for**, and the two things you need:

1 order from the shop is waiting and could not be read - · it is for Pizza, which is not in your Products Nothing is lost: it stays in the queue and is tried again every 30 seconds. Add what it names under Products and it will come in on its own.

And it tells you the honest reason, in words - never a code. The three it can give: no bake day on it / nothing on it / it is for something not in your Products.

AND IT CLEARS ITSELF. The moment the order can be read, the note goes. **A warning that cannot clear is worse than no warning at all**, so the screen now watches for the notice appearing *and* disappearing - before this it only repainted when something was *imported*, which would have meant the note never appeared when it mattered and never left when it stopped mattering.

One deliberate piece of stubbornness: if the queue cannot be reached, the note STAYS. "Nothing is waiting" is a positive claim, and it must never be made on a request that never came back - the same rule the label's open count follows. A failed check leaves the last honest answer standing rather than replacing it with a comfortable one.

And nothing of yours was touched to do it. The note is a memory of the last poll, not a setting - nothing was added to your data, nothing synced, and **no new field went near the cloud** (which is the thing that quietly breaks publishing for every order).

No SQL, no Edge Function, no upload. The suite is **2,898 tests, all green**, and all three halves were **bitten**: stop drawing the note and one test fails by name; let an unreachable queue clear it and three fail; skip a stuck row without remembering it and three fail.

08 Oct 2026 - engine v363, AN ORDER IS NO LONGER THROWN AWAY OVER ONE ITEM YOU NO LONGER SELL (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"build the import fix now"* - after the other shop sent word of it the same night.

THIS IS THE ONE THAT WAS LOSING A CUSTOMER'S ORDER, AND LOSING IT SILENTLY.

What was wrong. When an order comes in from the shop, your app checks it against your products before taking it - **and it demanded that EVERY line match.** So **one line the shop had sold which your app no longer knows** - a product you **paused, renamed or deleted while a customer's page was still open** - **threw away the whole order, including the items you DO still sell.**

And the throw-away was silent and permanent. A refused order is left waiting in the queue, retrying for ever, **and nothing on any screen says a word.** On the other shop this produced a real order that sat unserved for a day while the shop's own counts had already moved - the two sides disagreeing, with a customer's order in between.

What changed. One item you still sell is now enough to take the whole order. The lines that could not be matched are **written onto the order's own note**, naming **the item, the quantity, and the price the shop charged** - because your order total is added up from the lines it holds, so **without the figure you would be looking at a total the customer never paid.**

And the honest refusal is kept. An order where you sell **nothing at all** still waits rather than being taken - there is nothing to make it out of, and taking it would leave an empty order and hide the problem.

ONE THING STILL NEEDS YOU, AND IT IS THE POINT OF ALL THIS: the fix stops FUTURE orders being lost. It cannot find the ones already lost. Run this in **Supabase -> SQL Editor**:

select created_at, (data::jsonb)->>'date' as day, (data::jsonb)->'lines' as items from public.incoming_orders where status = 'new' order by created_at desc;

Any row older than a day or two is a customer who has been waiting. Their order is sitting in your database

and never reached your app. Tell me what it returns and I will take it from there.

No SQL, no Edge Function, no upload. Your orders, customers, prices and days are untouched. The suite is **2,893 tests, all green**, and all three halves were **bitten**: put .every back and two tests fail by name; stop collecting the unmatched lines and one fails; stop writing the note and one fails.

And v362 went in first, in its own commit - so this one stands on its own. **The Engine on your More screen reads 363.**

07 Oct 2026 - engine v362, A DAY'S ORDER COUNT COULD READ SHORT, AND NOTHING WENT RED (no database step, nothing to upload - pushing this one is the whole of it).

Where this came from, because it was not you who found it: your other shop did. munchies reported that a day showed one quantity to a customer and a different one in the app, and its session traced it to **code the two shops share**. They left a note for this side. **The note was right - the same fault was here, in the same file.**

What was wrong. A day's booked count is asked for on the calendar's own chip, in the week view and on the dashboard. It was being read from **the costing routine** - the one that works out how much flour a day needs.

And that routine SKIPS ORDERS. It passes over an order whose product has been **deleted**, and - **silently, with no warning anywhere** - over one whose product **has no recipe** yet. Which is right for costing: there is nothing to explode. **But it is wrong for counting.**

So a day you had really sold could count LOW, or count as EMPTY, while the shop and the "Only N left" numbers - which read the orders - went on saying it was booked. **Two numbers about one day, disagreeing, with nothing on any screen going red.**

What changed, and it is one line. The correct counter was **already in the same file, four lines below the wrong one** - it counts the orders and nothing else. The day's own count now reads that. **And the silent skip speaks up:** a product with no recipe now says so, so a line that cannot be costed is a line you are told about rather than one that quietly disappears from the totals.

Nothing you see changes on a day where every product has a recipe - which is every day so far, in all likelihood, because the fault needs a product with no recipe or one you have deleted. **That is exactly why it went unnoticed: it only bites on the days that are already unusual.**

No SQL, no Edge Function, no upload. Your orders, customers, prices and days are untouched. The suite is **2,891 tests, all green**, and both halves were **bitten**: make the day read the costing routine again and two tests fail by name; make the recipe-less skip quiet again and one fails.

And the note it came from is written down properly - the two shops keep a shared bridge note, and this side has now answered it.

07 Oct 2026 - engine v361, THE REFUND BUTTON (RE-RUN supabase/receipts.sql BEFORE YOU PUSH).

Your words: *"yes build the refund button"*.

Where it is. On a **paid** order's row, beside the money buttons: **Refund**. It appears only once the money is in - there is nothing to give back until it has arrived - and **it asks first**, naming the amount. **Giving money back is not a stage to nudge like Paid · Cash**, so it is never one tap on a row you press all day. Say no and nothing has happened.

And when you confirm, two things move together:

1. **The order is marked refunded**, and its row wears a **Refunded** tag instead of Cash / TNG. 2. **The sale stops counting.** A refunded order is **neither takings nor owed** - the money came in and went back out - so it is skipped by **the Money screen, every journal and Profit**. Put plainly: it is not money you have, and it is not money anybody owes, and the app does not make you choose.

And it is deliberately NOT also written down as an expense. The sale is gone; subtracting the refund a second time would take it off your profit twice. **One event, one effect.**

The receipt keeps its serial number. It is **marked**, never renumbered and never deleted - money that moved

is money that moved, and **a receipt that vanished would leave exactly the gap in the sequence that numbering exists to prevent.** The paper reads *"Receipt #000123 - refunded"*.

And you can undo it. A **Refund** pressed on the wrong order offers **Undo refund** - the order goes back to paid and the receipt stops being marked. **The number is never touched by either press**, so nothing about the sequence moves in either direction.

And when the receipts register cannot be reached, UNDO REFUSES RATHER THAN HALF-DOING IT. An order put back to paid on your phone while the register still says refunded is **your books and your paper telling two different stories** - so if the register cannot be told, **nothing changes**, and the message says why.

THAT IS THE ONE THAT MATTERS TO YOU TODAY - and it is why running the NEW script matters. The undo needs unrefund_receipt, which is only in the **current** supabase/receipts.sql. On the older one: **refunds still work** (that function was already there), and the sale still stops counting - but **Undo will refuse, and say so**, rather than leave the register disagreeing. **Paste the current file into a new query and run that.**

And a refund itself is never blocked by the register. The money really has gone back - that is a thing you did, not a thing the app decided - so the sale stops counting the moment you confirm, whatever the register says. **If the register could not be marked, the message tells you so**, because an unmarked receipt for a refunded sale is something you need to know about rather than discover on paper.

HOW TO DEPLOY THIS ONE.

1. **Re-run supabase/receipts.sql in Supabase** (Dashboard -> SQL -> New query -> Run). It is safe to run again, and it now also carries the function that takes a refund mark back off. 2. **Then push.**

If you push first nothing breaks - the refund still marks the order and still comes off your takings - but the receipt would not be marked in the register until the SQL is run.

No Edge Function. The suite is **2,885 tests, all green**, and every load-bearing rule was **bitten**: stop the till skipping a refund and two tests fail by name; count a refunded sale in Profit again and two fail; let the trading journal count it and one fails; offer Refund on an unpaid order and one fails; make the refund skip its question and one fails.

Two things worth your eye, and both are honest limits rather than defects:

- **The customer's track card still shows a refunded order as it was.** The card is theirs, and what it says about a refund is a decision yet to be made. **Your shop's Policies text still says orders are not refundable.** Now that you can refund from the app, that line is worth another look.

07 Oct 2026 - engine v360, EVERY RECEIPT NOW HAS ITS OWN SERIAL NUMBER (THERE IS A DATABASE STEP THIS TIME - RUN THE SQL BEFORE YOU PUSH).

Your words: *"i want to get ready for anual sales more that rm150k"* - and then *"system will need to be prepare for refund"*.

What this is for. Above **RM150,000 of gross takings in a twelve-month period**, a business must issue **serially numbered receipts** (Income Tax Act 1967, s.82(1)(b)). Your orders already carry their own code, but a code is not a **series** - it is deliberately not sequential, and it says nothing about whether any receipt is missing. **This gives every paid order its own serial, from a run that cannot contain a gap.**

Where the number comes from - and why it could not simply be counted. A run of receipt numbers is a **shared counter**, and this app has never kept one on a phone, on purpose: you and your sister both take payments, and a counter living on one phone would either clash with the other's or leave a hole. **A hole in a receipt sequence is exactly what an auditor asks about.** So the counter lives in your database, where a number can be handed out exactly once.

The number belongs to the ORDER, not to the press. It is claimed **once**, when you record the money as received - and every later read returns the same number and consumes nothing:

| what you do | what happens | |\---|\---| | Press **Paid · Cash** or **Paid · TNG** | the order claims its number | | Print the receipt | the same number | | Print it again | the same number | | Open it on your other phone | the same

number |

So pressing Paid twice, or printing ten copies, cannot burn numbers or leave gaps. An accidental press is harmless by design.

And you were right to ask for refunds. A refund **marks** the receipt and **keeps its number** - it is never deleted or re-used, because money that was taken and given back is still money that moved, and a receipt that vanished would leave the very gap this exists to prevent. The paper reads **"Receipt #000123 - refunded"**.

Where you see it. On **Invoice**, under the customer's name: **Receipt #000123 · Order #A3F9C2** - both numbers, because each does a different job: **the serial proves the sequence, and the order code is what finds the order again.** The printed sheet carries the serial under its title as well.

What it says when there is no number yet, and this matters. It never prints a blank, and never invents a zero:

- **paid, but the number has not been claimed** - *"No receipt number yet - this order has not reached the bakery's records."* (This is what a phone with no signal leaves behind; open the receipt again and it claims it.) **not paid** - *"a receipt number is issued when the money is recorded as received."*

HOW TO DEPLOY THIS ONE.

1. **Run supabase/receipts.sql in Supabase first** (Dashboard -> SQL -> New query -> Run). It is safe to run twice. 2. **Then push.**

If you push first nothing breaks - the number is simply not claimed, and the receipt says so - but orders paid before the SQL is run would sit without a number until you open each receipt. **Running the SQL first avoids that.**

One thing still to come, and it is on purpose: there is no refund BUTTON yet. The number survives a refund and the paper can say so, which is what "prepared for refund" means - but nothing in the app records one yet. That is its own piece of work.

And one thing about your shop's own words. Your Policies text on the shop says orders are **not refundable**. Prepared for a refund or not, that is what a customer reads today - worth a look when you next open Settings -> Storefront.

No Edge Function. The suite is **2,878 tests, all green**, and the four load-bearing rules were **bitten**: remove the guard that stops an unpaid order claiming a number and a test fails by name; remove the "already has one" shortcut and another fails; stop padding the serial and four fail; make a refund wipe the number and one fails.

07 Oct 2026 - engine v359, THE VAN'S PRICE NO LONGER ASKS YOU THE SAME TWO QUESTIONS TWICE (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"in a edit order, after send a van, we entre pickup time, pressing get a price, again it ask for pickup time again?"* - and then your own answer to what to do about it: *"why not remove the one inside get a price. the logic is i entered all details that is relevent, i just want a price."*

You were right on both counts, and your reasoning is better than mine was.

What was wrong. Choosing a van on an Edit order drew the **same two facts twice**, in two pairs of boxes that did not talk to each other:

- the order's own **Courier delivery date** and **Pickup time**; and, inside **Get a delivery price**, a second pair - **The day the driver collects** and **The time the driver collects**.

You filled in the order's. The price half was still holding the value it had taken **when the screen was drawn** - before you typed - so it looked like it was asking you again. **And it wasn't only confusing:** the price was asked for using **that** box's value, so a van could be priced for your Settings dispatch time and then come at a different hour.

What changed. The two boxes inside Get a delivery price are gone. The price simply reads the order's own **Courier delivery date** and **Pickup time** - the same two you typed - and the line above the button says where they come from. As you put it: you enter all the details that are relevant, and you want a price.

The value is read when you press the button, not when the screen was drawn. That is the whole repair: what is priced and what the van will collect at can no longer drift apart.

And the half worth knowing about, because it is about money: booking still books **the trip that was priced**, so changing the order's date or pickup time **after** a price does not move a booking. The line under the button now says exactly that, in those words.

No SQL, no Edge Function, no upload. Your orders, customers, prices and days are untouched. The suite is **2,867 tests, all green**, and both halves were **bitten** - make the price read the Settings time instead of yours and the test fails by name; put a box back into that half and it fails by name too.

One thing I cannot check for you, and always your side of the line: the admin app sits behind your PIN and I am not allowed to open it, so **nobody has looked at this screen with a human eye**. What is proven is that the two boxes are gone and the price is asked for your own time. **Open an Edit order with a van and check it reads the way you want.**

07 Oct 2026 - engine v358, APPLYING A COUPON NO LONGER LOOKS LIKE IT DID NOT TAKE (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"imention about coupon applying that feels like taking two time but actually one coupon apply"* - and, when I asked you to pick which of four things it was, *"some problem, i dont know how to discribe"*.

SO I STOPPED ASKING AND SHOWED YOU. A page was built that runs **the app's own coupon code**, with the order row drawn exactly as your screen draws it, a row of live figures underneath, and nothing pretend. You pressed it, and said: **"yes that's it"**.

What was wrong - two halves of one fault.

1 · The screen did not look like it had taken. After ONE press the order row said three things at once:

- *" Bring-a-friend - RM8.00 already off this order's total." " RM8.00 coupon ready for this order"* and it offered the **Apply coupon** button again.

Two of those contradict the first. One says the coupon is off; the next says it is still *ready* and hands you the button. So the only honest reading of that screen was *my press did not work* - and pressing again was the obvious next move. Your words for it were exactly right: **it feels like it takes two times.**

2 · And pressing again burnt a coupon for nothing. The second press spent the **next** coupon, stamped it on the **same** order, and moved **no money at all** - because an order can only ever take one coupon off. The page said it plainly: *"2 coupons were spent. RM8.00 came off the order. The other 1 went nowhere."*

What changed.

- **The button is gone once the order has its coupon.** Everything that block had to say is already said by the line above it, which names the coupon and the money it moved. **The row now reads as applied, because it is. And the engine refuses it too.** Marking a coupon used now checks whether that order already has one and **refuses**. That half is deliberately not about the screen: **a second phone holding a screen drawn before the press can still press**, and no screen-level fix can reach a phone that has not repainted. **Three answers, not two.** A refusal used to say *"Nothing to apply"*, which was a lie when the truth was *"this order already has its coupon"* - they now read differently.

PROVED BY DRIVING IT, NOT BY READING IT. With the fix in, the page shows: one press -> the button gone, one coupon spent, RM8.00 off; then the hidden button forced from a stale screen -> **the guard held, still one coupon, still RM8.00 off.** Four new tests, and **both halves were bitten** - take the engine guard out and two tests fail by name; let the screen offer the button again and one fails.

No SQL, no Edge Function, no upload. Your orders, customers, prices and days are untouched. The suite is **2,866 tests, all green**.

ONE THING TO DECIDE, AND IT IS YOURS. If any order already had two coupons spent on it, those coupons are spent and gave you nothing. They can be found and put back - say the word and I will go looking. Nothing is changed until you ask.

07 Oct 2026 - engine v357, THE NOTICE OPENS ABOVE THE BUTTON, NOT ON TOP OF IT (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"can the pop up be open higher not to block the Placae Order button"*

What changed, and it is where the panel hangs from. The panel opened directly above the *line* it belongs to - which put it **on top of the row**, and that row is **Place order**. It now opens above the **whole bar**:

- its parent is the **bar** rather than the words, so it hangs off the bar's top edge; **Place order stays in sight** while the notice is read - which matters most at the moment a customer is deciding whether to tap it; and it still **floats** - the bar does not grow, so the button does not move either.

One thing this made me change, and it would have been a bug. "Click away to close" was decided by asking whether the press landed outside the *words* wrapper. The panel is no longer inside that wrapper, so **a press on the WhatsApp link inside the panel would have counted as "away"** and shut the notice the customer was reading. "Away" now means outside **both** the words and the panel.

It also gained a ceiling, measured rather than picked. At 375px wide the notice is English 405px, Chinese 371px, **Bahasa Malaysia 518px** - so the panel fits all three whole instead of scrolling, and a viewport term stops it running off the top on a short phone.

AND ONE THING TO KNOW ABOUT THE v356 COMMIT YOU PUSHED. It contains the first half of this change by accident - the markup that moves the panel, without the code that goes with it - because the edit was already on disk when the commit was made. **In that pushed state the notice closed itself whenever you tapped anything inside it, including the WhatsApp number**: the "click away" test still believed the panel was part of the words. Nothing else broke, and this version is the fix.

FOR NEXT TIME: say so when you are about to commit, or commit *before* asking for the next change. My edits land on disk as I make them, and a commit made in between sweeps up whatever is half-finished.

No SQL, no Edge Function, no upload. Nothing in the backoffice changed. The suite is **2,864 tests, all green**.

07 Oct 2026 - engine v356, THE NOTICE NOW SAYS HOW LONG THE DETAILS ARE KEPT (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"add how long we keep them? What is the suiatble answer?"*

Two clocks, not one - and that is the answer. The third point never said how long anything was kept, which is one of the things a notice is asked to state. It does now:

Your details stay with your order in our order book, on a secure cloud service and on our own phones behind a PIN. **We keep the sales record for seven years, as the tax office asks** - but your name, number and address are not part of that record, so ask us any time and we will delete them.

Why seven years, and why it is not seven years for everything. Seven is the Malaysian figure, checked rather than guessed: LHDN's own wording is "Dokumen percukaian mesti disimpan paling kurang 7 tahun" - tax documents must be kept **at least** 7 years - and it comes from **s.82 of the Income Tax Act 1967**. But a sales record is **money and dates**. A customer's name, number, address and door pin are **not needed for it**, so they do not have to ride along for seven years, and the notice now says as much. That is a better answer than one blanket period, and a truer one.

AND THE HONEST NOTE, recorded so it is not discovered later: there is **no control anywhere in the app that deletes an order or a customer's details**. The only deletes in it are old calendar rows, imported incoming orders, and reviews. So *"ask us any time and we will delete them"* is a promise **you keep by hand**, in Supabase - not by pressing anything. If it should become a button - a "forget this customer" that clears the name, number, address and pin while leaving the sale in the figures - that can be built.

In all three languages. Mandarin and Bahasa Malaysia both carry the seven years and the distinction between the record and the details.

No SQL, no Edge Function, no upload. Nothing in the backoffice changed. The suite is **2,864 tests, all green**.

07 Oct 2026 - engine v355, THE FIRST LINE SAYS "YOUR DATA" (no database step, nothing to upload -

pushing this one is the whole of it).

Your words: *"use data instead of particular"*

What changed, and it is one phrase. The opening line of the notice named the items it was talking about. It does not any more:

- **was:** We use **your name and your WhatsApp number** only to make your order, get it to you, and talk to you about it. **now:** We use **your data** only to make your order, get it to you, and talk to you about it.

AND THE NOTICE STILL NAMES EVERYTHING, which is why this is safe. The Act asks the notice to name the personal data involved - and it does, twice over, in the same panel:

- the **address** and the **pin on the door** are named in the second half of that very line; the **name, number and address** are named again in the next point - *"the driver is given your name, number and address"*.

So nothing went unnamed; only the first sentence stopped repeating the list. **That check is the one to make again if this line is ever shortened further** - a notice that names nothing is a different thing from a notice whose naming sits one line down.

In all three languages. Mandarin says 你的资料, Bahasa Malaysia says data anda. The heading, the Act line and the contact number are untouched.

No SQL, no Edge Function, no upload. Nothing in the backoffice changed. The suite is **2,864 tests, all green**.

07 Oct 2026 - engine v354, THE NOTICE SAYS "YOUR DETAILS" AGAIN (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"you din use the details?"* - and then *"your details"*

What changed, and it is two words back. When v353 stopped the notice claiming *"nobody else sees your details"* - which was **not true** - the phrase "your details" went with it, and the panel ended up saying *"Nobody else sees them"* and *"They stay in our order book"*. You asked for it back:

- **was:** ...Nobody else sees **them**, and we never sell them. **now:** ...Nobody else sees **your details**, and we never sell them. **was: They** stay in our order book... **now: Your details** stay in our order book...

And it is honest again, because the sentence before it now does the work. v353 put the people who actually handle an order in front of it - the driver, the alert, WhatsApp - so *"nobody **else**"* means what a reader takes it to mean. Standing alone before v353 it was simply false; now it is a closing reassurance with the truth already said one line up.

In all three languages. Mandarin says 你的资料, Bahasa Malaysia says maklumat anda.

No SQL, no Edge Function, no upload. Nothing in the backoffice changed. The suite is **2,864 tests, all green**.

07 Oct 2026 - engine v353, THE NOTICE SAYS "WE USE", NOT "WE COLLECT" - AND STOPS CLAIMING NOBODY ELSE SEES YOUR DETAILS (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"as asian, we are sensitive to when you said you collect, that mean you keep it, can we say like 'We use your name..... that the only'"*

1 · "Collect" is gone from the whole panel. The first point opened *"We collect your name, your WhatsApp number and, for a delivery, your address and door pin."* - and "collect" is data-protection jargon that lands, in this part of the world, as *"they are taking it and keeping it."* It now opens on **use** and carries your own limit:

- **was:** We collect your name, your WhatsApp number and, for a delivery, your address and door pin. We use them only to make your order, get it to you, and contact you about it. **now:** We use your name and your WhatsApp number **only** to make your order, get it to you, and talk to you about it. For a delivery we also need your address and the pin on your door.

Nothing legal was lost. The Act wants the notice to **name the data** and **say why it is used** - it does not require the word "collect". The name, the number, the address and the pin are all still named, and so is the purpose. (What would understate it is dropping the *naming* too - "we use your details" - and that is the line to

keep.)

2 · And the second point was not true. It read *"The driver is given what they need to deliver. Nobody else sees your details, and we never sell them."* **Three other parties were seeing them**, all of them traceable in the code:

- **the courier** is given the customer's **name, phone number and address** to book the trip (admin/js/courier_job.js); **the alert that pings your phone** carries the customer's **name, WhatsApp number, items, delivery day, total, note - and for a delivery, their address** - as plain text through a push service (supabase/order_alerts.sql); **WhatsApp** carries the confirmation and the payment QR to them.

So it now says who, in plain words and with **no company name in it** - which is the **class** of third parties the Act asks you to state:

To bring your order to you, the driver is given your name, number and address. Your confirmation reaches you on WhatsApp, and an alert carrying your order pings our phone. Nobody else sees them, and we never sell them.

In all three languages, as the guidance asks. The heading, the Act line and the contact number are untouched.

No SQL, no Edge Function, no upload. Nothing in the backoffice changed. The suite is **2,864 tests, all green**.

07 Oct 2026 - engine v352, THE NOTICE LINK IS NO LONGER RED, AND IT OPENS ON A CLICK (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"privacy notice, can be underline, but dont put red color."* - and then *"make it expend only when i click on it, away, colapse."*

1 · The link is no longer red. The words **privacy notice** in the line under Place order are what opens the notice. They were underlined **and coloured in the app's terracotta**, which - sitting under a sentence of grey small print - read as a red link. They now take **the small print's own colour**, so the two can never drift apart:

- **the underline stays** - it is what says the words are pressable; **the weight stays** - bold, exactly as before; **on hover the underline thickens** rather than changing colour, so nothing here goes red again; **and it gained a focus ring it never had.** appearance:none with no border left it relying on whatever the browser felt like drawing when a keyboard reached it. It now carries the shop's own ring, in the small print's colour.

Why the colour mattered more than it sounds. The line is the only place the notice is given, and it sits inches from a big terracotta PLACE ORDER button. A red link under grey type pulls the eye to the legal small print and away from the button.

2 · It opens on a CLICK now, not on hover, and it closes when you click away. That is what you asked for - it had opened on hover for one version (v351) and you changed it the same day. **It is also the better behaviour, and not only because you asked:** hover does not exist on a touchscreen, so a hover-opened notice needed a second, hidden path to be usable on the phones this shop is ordered from. **One press that behaves the same on a laptop and on a phone has no second path to keep in step.**

- **Clicking anywhere outside the notice closes it** - including a click on the shop around it. **A click INSIDE it does not close it**, so the WhatsApp number in the panel still works and the paragraphs stay put while they are being read. **The keyboard works the same way**: Enter or Space on the words opens and closes it, and Escape closes it.

3 · And it fixes a bug I shipped in v350 - the press did nothing on a freshly loaded shop. When the notice moved under the button (v350), its press was wired up inside a function that **runs again on every render and on every language switch**. Binding a *toggle* there means binding it again each time it runs, and **two bindings cancel each other out**: pressing the words did **nothing** on a shop that had just loaded, and opened the notice only after a language switch had made the number of bindings odd.

Driving the press once - which is what I did when I checked v350 and v351 - passed by luck. Pressing it **twice in a row** is what exposes it, and that is what found it today. The wiring now lives at module scope, bound once, beside the language pills.

No SQL, no Edge Function, no upload. Nothing in the backoffice changed. The suite is **2,864 tests, all green**,

and the new one is an **invariant** rather than an example: **renderStatic must bind no event listener at all**, because anything toggled there is toggled twice. It was bitten by putting a single binding back inside that function, and the file was restored byte-identical.

07 Oct 2026 - engine v351, THE NOTICE OPENS ON HOVER (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"i want to delete this 'we use your details only to bake and deliver'"* - and then *"no need to press, it show like a tooltips, move mouse away it colapses."*

Two changes, and the second one reshaped it.

1 · That phrase is gone. The line is now just the agreement:

By ordering you agree to our privacy notice.

Which left **"privacy notice" twice back to back** - "...our privacy notice. Privacy notice". So **the words IN the sentence are now the trigger**, the way Amazon's own checkout does it: nothing repeats, and nothing was added.

2 · It opens on HOVER, and it collapses when the pointer moves away. That is what you asked for. **But hover does not exist on a phone, and this shop is ordered from phones** - so the tap is kept, and it is not optional: a tap opens it and a tap closes it there. The keyboard works too - tab to it and focus opens it, Escape closes it - which is what the accessibility rules ask of anything that appears on hover. **Four ways in, one panel.**

And the panel no longer pushes the bar taller. It floats above the line now, so opening the notice **never moves Place order**. Until this version, showing the notice grew the bar and lifted the button away from the customer's thumb at the exact moment they were reaching for it.

The bar got shorter, and the spacer had to be re-measured from two directions, not one. Dropping the phrase made the closed bar smaller - at 375px wide: English 100px, Chinese 101px, Bahasa Malaysia 100px. **But at 320px wide Bahasa Malaysia wraps to a second line and reaches 118px.** So the spacer is set from the tallest language **on the narrowest screen**: **126px**. Set from the 375px reading, it would have hidden the last row of the shop on a small phone.

Nothing else about the notice changed. The same four facts, the same contact number read from Settings -> Storefront, the same Act and the same reviewed date - and the same rule: **a notice, never a gate.**

No SQL, no Edge Function, no upload. Nothing in the backoffice changed. The suite is **2,863 tests, all green**, and the spacer test was **bitten**: cut to 90px it failed by its own message naming the 320px measurement, and the file was restored byte-identical afterwards.

07 Oct 2026 - engine v350, THE NOTICE MOVED TO THE BUTTON (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"i think the warning is too obvious"* - and then *"we just choose one."*

What changed, and it is WHERE the notice lives rather than what it says. The notice was a block sitting in the middle of the details card (v348, v349). You said it read as too obvious, so it is now **one quiet line under Place order**, and the four facts open from a panel behind it:

By ordering you agree to our privacy notice - we use your details only to bake and deliver. Privacy notice

Lifted at v351: that is not the line any more. **"we use your details only to bake and deliver" is gone** (your call), and the separate **Privacy notice** press went with it - the words **"privacy notice" in the sentence** are the trigger now, so nothing repeats. It also opens on **hover** and collapses when the pointer leaves, and the panel **floats** instead of pushing the bar taller. The heights in the paragraph below are v350's and were re-measured at v351.

Why the button. That is where the large platforms put theirs - Amazon's whole consent line at the till is ten words - and it is the one place on the page a customer is already reading before they commit. **The panel opens in place, never on a page of its own**, because the Commissioner's own guide says a notice "hidden behind too many links" is a notice nobody reads.

And the law now CLOSES the notice instead of opening it. The heading was **"Personal Data Protection Notice"** and the first line was **"Under the Personal Data Protection Act 2010, we must tell you how we use the details you send us."** Both are gone. Malaysia's own template for small businesses opens in plain words - "We take your personal data seriously" - and naming the Act first made the notice read as a legal document rather than something written for a customer. **The Act is still there**, and so is the reviewed date, in the line at the foot of the panel:

Under Malaysia's Personal Data Protection Act 2010 (Act 709). Last reviewed: 7 October 2026.

The line is 12.5px and must not get smaller. It is now the ONLY place the notice is given, and the same guide warns against a font "so small that it results in the data subject not reading" it. It is the same size the old block used. In all three languages, as before.

And the bar got taller, which costs a little of every screen. The bar is fixed to the bottom, so its height is paid on every page. Measured at 375×812: **English 118px, Chinese 119px, Bahasa Malaysia 136px** - its line wraps one line longer than English's - against **78px** for the old one-row bar. The spacer under the page, which is what stops the bar covering the last row of the shop, is set from the **longest** language: **140px**. Set from English's 118px, it would have hidden the last row of the shop from every customer reading in Bahasa Malaysia.

Nothing else about the notice changed. The same four points, the same contact number read from Settings -> Storefront, and the same rule: **a notice, never a gate** - no tick box, and nothing in it can refuse an order.

No SQL, no Edge Function, no upload. Nothing in the backoffice changed at all - your orders, customers, prices and days are untouched. The suite is **2,863 tests, all green**, and the new one was **bitten three times**: the line set to 11px, the spacer cut to 90px, and the panel renamed so the press named a panel that was not there - each failed by its own message, and both files were restored byte-identical afterwards (checked by sha).

07 Oct 2026 - engine v349, THE NOTICE NOW NAMES THE LAW IT IS GIVEN UNDER (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"can you change the word to something more professional like Under privacy act XXX"*

What changed. v348 added the notice; this changes what it is CALLED and opens it by naming the law.

- The heading was **"Your details"**. It is now **"Personal Data Protection Notice"**. Under it sits a new first line: **"Under the Personal Data Protection Act 2010, we must tell you how we use the details you send us. Here is all of it:"**

Lifted at v350: both of those are gone. The notice now sits under **Place order**, not in the details card, and it **closes** with the Act rather than opening with it - the heading is **"How we use your details"** and the Act moved to the foot of the panel. Naming the law first was your own call at v349 and it was right at the time; what changed is that it made the notice read as too obvious in the middle of the form. **The Act is still named, and the date is still recorded** - only the order moved.

Why it reads better this way, and it is not only polish. A heading like "Your details" describes the section a customer is standing in. **"Personal Data Protection Notice" says what the block IS** - a notice, given under a law - and naming the Act tells a customer this is something the bakery **has to** say, not something it has chosen to say. That is the difference between house rules and a notice.

The Act is the Personal Data Protection Act 2010 (Act 709) - the same Act whose section 7 is where the duty to give this notice comes from.

In all three languages, which is what the guidance asks for:

- **English:** Personal Data Protection Notice **Mandarin:** 个人资料保护声明 **Bahasa Malaysia:** Notis Perlindungan Data Peribadi

Nothing else about the notice changed. The same four points, the same date, and the same rule about the contact number: the sentence is translated, the number is read from your Storefront setting. **It is still a notice and never a gate** - no tick box, and nothing in it can refuse an order.

No SQL, no Edge Function, no upload. Your orders, customers, prices and days are untouched. The suite is

2,862 tests, all green, and the new line was **bitten**: take it out of one language and **three tests fail by name**, because the languages must all carry the same keys.

07 Oct 2026 - engine v348, THE SHOP NOW TELLS CUSTOMERS HOW THEIR DETAILS ARE USED (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"the privacy notice"* - and then *"make it a small one in store?"*

What was missing. The shop asks a customer for their **name, their WhatsApp number, and for a delivery their address and the pin on their door** - and until now it said almost nothing about what it does with any of it. One line under the phone field said "we'll never spam you", and that was the whole of it. **A customer was handing over their details with no notice at all.**

What is there now. A **small notice at the point of collection** - inside the card where the details are typed, not tucked away at the foot of the page. **It is written in English, Chinese and Bahasa Malaysia**, so a customer reads it in the language they are ordering in. It says:

- **what is collected** - name, WhatsApp number, and for a delivery the address and door pin; **why** - to make the order, get it to you, and talk to you about it; **who sees it** - the driver gets what they need to deliver, nobody else, and it is never sold; **how long** - kept with the order book; ask to delete and it goes, apart from the sales record the tax office requires; **how it is held** - a secure cloud service and your own phones, behind a PIN; **your rights** - ask any time to see, correct or delete it, with **your WhatsApp number**; **when it was last reviewed** - dated, which the guidance asks for.

It is a notice, never a gate. There is no tick box and nothing in it can refuse an order. A rule that blocked a sale you would take by hand is the one thing this app does not do - a customer who taps send before reading is still a customer.

And the contact number is deliberately NOT typed into the three translations. The sentence is translated; the number comes from the **same setting the order button builds its link from**, so changing it in Settings -> Storefront moves both together. A number written by hand into a translated string would look right today and be wrong the first time you changed it - and a test now holds that.

Why this is the shape it is. Malaysia's PDPA requires a **written notice** of what is done with personal data, and the Notice and Choice principle is about a customer being able to make a choice **before** they send it - which is why the notice sits beside the fields rather than in a policy page they would never scroll to. **The penalty for not having one is up to RM300,000 or two years, or both.** The Commissioner's guidance also asks for the notice in **both English and Bahasa Malaysia** - this one is in both, and in Chinese as well, because your shop already speaks three languages.

Two things worth your own eye, and neither is a defect:

1. **The guidance suggests naming the person in charge of data questions.** This notice names the **bakery** and gives your **WhatsApp number**, not your personal name. Putting your own name on the shop page is your decision, not mine - say the word and it goes in. 2. **The Chinese and Bahasa are mine.** I wrote them plainly and a test holds that all three languages are really translated, but **have someone you trust read the Chinese and Bahasa once** before a stranger does. If you would rather write them yourself, they are six short strings.

No SQL, no Edge Function, no upload. Your orders, customers, prices and days are untouched. The suite is **2,862 tests, all green**, and the three new guards were **bitten**: leave a language in English and a test fails by name; type a number into the translated sentence and two fail; break the number's grouping and one fails.

07 Oct 2026 - engine v347, THE SHOP'S SPARE MENU NOW LISTS ONLY WHAT YOU SELL (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"trim it to what i sell"*

What changed. v346 fixed the shop's **spare copy** of your settings - the one it uses when the published settings cannot be reached. It left one thing alone deliberately and told you why: the spare copy's **menu** still carried two products, one of them a **Sandwich you do not sell**. You have now said to trim it, so it lists **exactly what you sell: one focaccia, RM16 a loaf** - the same name, price and unit that Settings -> Storefront publishes.

Why this is worth a version rather than a shrug. That spare copy is what a customer sees at the one moment it matters: when the published settings cannot be reached, which is also when their order cannot be placed. **A menu that offers something nobody can order is worse than a shorter one** - the customer picks the Sandwich, the order arrives, and you have to refuse it or improvise in front of the person you were trying to keep.

And this is what v346 warned about: it cost more than two lines, because four tests had been borrowing the sample menu as their own fixture. They were not testing the menu - they were testing **behaviour that needs two products**:

- one wanted **two products with different change/cancel windows** so the receipt could be checked for the strictest; one wanted **a product the baker had not switched on** so it could prove no note is offered; one wanted **per-product availability stamps** on two products in one day; one wanted to **fill a basket from two products and deplete them differently** on the next poll.

Each of those now writes down the menu it needs instead of borrowing yours. That is a better test: a test that reaches into your live settings for its fixture goes wrong the moment you change what you sell, which is exactly what happened here. **One test had been asserting the old placeholder number itself**, and another spelled out RM15.00 six times; both now read the figure off the config, so they cannot go stale silently again. A test helper that swaps a menu in and puts it back was added and **bitten**: with it disabled, six tests fail by name.

Nothing you see changes, and nothing about your published menu moves - this is the spare copy underneath it. **No SQL, no Edge Function, no upload.** Your orders, customers, prices and days are untouched. The suite is **2,859 tests, all green**.

07 Oct 2026 - engine v346, THE SHOP'S SPARE NUMBER WAS SOMEBODY ELSE'S (no database step, nothing to upload - pushing this one is the whole of it).

The second half of the soft-launch pass. Nothing you can see changes on any screen you use.

What was wrong. The shop carries a **spare copy** of your settings - the number to message, the menu, the delivery days - for the case where the published settings cannot be reached. **The spare copy's WhatsApp number was not yours.** It held *60123456789*, which is the example number written in the notes above it - a real number, belonging to a stranger.

You would never have seen it, and that is the problem. The published settings override the spare copy, so the shop has been showing your real number all along - I read the live shop's own settings to be sure. **The spare copy is what a customer gets at the one moment it matters: when the published settings cannot be reached, which is also the moment an order cannot be placed.** A customer whose order had just failed was being handed a stranger's WhatsApp number instead of yours.

And the price in that spare copy was wrong. Focaccia sat there at **RM15** while your homepage and the live shop both say **RM16**. A customer who landed on the spare copy would have been quoted a figure you do not charge. Corrected to RM16.

What is deliberately NOT changed, and why. The spare copy's menu keeps **two products**, and one of them is a Sandwich the shop does not currently sell. That is left alone on purpose: this file is described in its own header as **the starting point** as well as the offline fallback, so its menu is a **sample that shows the shape** rather than a live copy of your published one. **If you would rather it listed exactly what you sell today - one focaccia - say so and it is a two-line change**, but it changes what four tests in the shop's suite use as their fixture, so it is worth doing deliberately rather than as a side effect. The shop's name, tagline, delivery days, 6:00PM cut-off and day capacity in that spare copy all **already match what is published** - each one was checked against the live shop rather than assumed.

Nothing else moves. No SQL, no Edge Function, no upload. Your orders, customers, prices and days are untouched, and the shop's live settings are unchanged - this is the spare copy underneath them.

07 Oct 2026 - engine v345, THE TRACK PAGE STOPPED HANDING OUT EVERY CUSTOMER'S NAME AND ADDRESS (THERE IS A DATABASE STEP THIS TIME - RUN THE SQL BEFORE YOU PUSH).

Your words: *"what housekeeping needed before i soft launch the app"*

This one is not housekeeping. It is the answer to your question, and it was the first thing found.

What was wrong. Your order-tracking table - the one the customer's "Track your order" card reads from - carried a read rule that said *every row is public*. Not "a customer may read their own order": **every row, to anyone who asked.** The key that reads it is public on purpose - it has to be, it ships inside the shop's own page - so a stranger who opened your shop and looked at its source could ask for **the whole table at once and receive every customer's name and delivery address**, along with the courier driver's name and phone number, which sit on the same row.

Your shop only ever asked for one order, by its code, and that part was built correctly. But **asking is not enforcing.** Nothing stopped a caller from simply not asking for one - and the page you wrote cannot be the thing that protects the data, because anyone can write a different page.

Why it could not be fixed with a rule. A database read rule can look at a **row**, but it can never look at **what the caller came to ask for** - and "only if you know the code" is a statement about the question, not about the row. So the read has moved into a small **function**: it takes the code, matches it, and returns **one row**. The table itself is now readable by **nobody**. There is one door left, and it only opens on an order whose code you already have.

HOW TO DEPLOY THIS ONE - the order matters. This is the first version in a while with a real database step:

1. **Run supabase/tracking.sql in Supabase first** (Dashboard -> SQL -> New query -> Run). It is safe to re-run - it creates the function and drops the old public rule. **This is the half that actually closes the door.** 2. **Then push.** The shop's half switches over to the function.

Between the two, tracking will not work - so do them one after the other rather than leaving a gap. **Nobody is affected while you do it**: the shop has never been shown to an outsider, which is exactly why this is the right week to change it.

What you will and will not see. Nothing on your own screens changes at all, and the customer's card looks and reads exactly as it did - same status, same delivery line, same courier details, same money. The only difference is that the row reaches the card through the function instead of off the table.

Nothing else is touched. Your orders, customers, prices and days are untouched. No Edge Function needs uploading. The suite is **2,858 tests, all green**, and the new guards were **bitten**: put the old public rule back and a test fails by name saying the table is public again; take a column out of the function's list and the card's own tests fail, because the card really does go without it.

And the changelog's own asterisks are corrected in the same pass. Seven places in this document - the newest being v344's own entry, at the top - wrote emphasis with a single asterisk, which the PDF prints **literally** rather than as bold. Found by reading the built PDF back, not by looking at the markdown, which is why it survived. **Fourteen marks in seven places, now bold.** The words are unchanged; only how they print.

07 Oct 2026 - engine v344, HOUSEKEEPING: A DEAD LINE REMOVED (no database step, nothing to upload - pushing this one is the whole of it).

Nothing you can see has changed, and this entry is here so the version list stays complete rather than leaving a gap.

What went: the pickup reminder still held a line that said *"Packed and will be sent for delivery on <the bake day>"*. **It could not be reached by any order** - it has been dead since v340, when you decided a courier customer should be told nothing at Packed at all - and it also named the **bake** day, which is the thing v337 to v343 spent a day and a half taking off every other surface. **It is deleted, so it cannot mislead anyone reading the code later, and a test now holds the rule it belonged to:** the pickup reminder says the order is READY and never that it is being sent - that belongs to the shipped message, at Collected / Shipped.

Your data is untouched. No SQL, no Edge Function, no upload. The suite is **2,857 tests, all green**, and the new guard was bitten: put the line back and it fails by name.

06 Oct 2026 - engine v343, A RUN IS ON THE DAY THE VAN COMES (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"the delivery run should not be on bake day only, for the case of 7th bake day order deliver 8th, his order should be appear only on date 8th. SO after his order, other order not specifing specific delivery will be on bake day 6th"*

1 · A run's day is the day the VAN comes. It used to be the bake day, always - so an order baked on the 7th with the van typed as the 8th still sat on the **7th's** run, which is the same confusion the last three versions have been unpicking everywhere else. **Now it is the day you typed on the order, and the order's bake day only when you have typed none.** Your example, exactly: an order baked the 7th and delivered the 8th appears on the **8th's** run and on no other, and an order with nothing typed stays on its bake day.

2 · A van day that is no bake day gets a run of its own - the "not on bake day only" half. Type a van day on a morning you are not baking and it appears in the day list as its own run, with its own stops.

3 · The run's own field says what it is. It read *"The bake day"*; it is now **"The day the van runs"**, because it is the one control deciding which day's work you are looking at.

4 · The Run (N) badge on a bake day counts the same way, so a card promising three stops opens a screen holding three. Where the van's day is no bake day there is no badge for it - a van day has no bake-day card - but it is still reachable from the run screen's own day list.

AND YOUR CORRECTION ON v342 IS IN THIS VERSION TOO.

Your words: *"for the v342, you miss underrstood me, what i want is the courier booked details like the one we see after pressing GET A DELIVERY PRICE, i shows LALAMOVE is on this order, with check the trip, status, bookedz:3:25pm and lalamove link"*

You were right twice over. v342 unfolded a summary of the **order** - who it is for, the items, the money - which is not what you open a booked row to read. **And the thing you actually wanted was hidden by my own caution:** the trip's card lives inside the price section, under a flag I had turned off to be safe, so pressing the section's own button drew no trip at all. Two fixes:

- **See this order** is now **See the trip**, and what it unfolds is **the courier's own record of that booking**, saying the courier is on this order, with **Get a delivery price** right there. Open that and the trip's card is waiting - the vehicle, when it was booked, where it has got to, the customer's share link, **Check the trip** and **Cancel trip**. **The order summary is gone**, because it was the wrong thing.

Nothing on that panel can be edited, and nothing is asked of the courier until you press - a repaint of the screen never spends a quote.

Your data is untouched. No SQL, no Edge Function, no upload. The suite is **2,857 tests, all green**, with two more rules bitten: key the run by the bake day again and the day test fails by name; put my "safe" flag back and the trip card's guard fails.

06 Oct 2026 - engine v342, THE RUN'S WARNING OPENS THE ORDER (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"on delivery run, it indicate an order which its lalamove already active, can clicking a button inside the red ribbon, drop down its full detail, as what is shown in edit order, get a delivery price."*

Done - and you chose for it to unfold in place, under the row rather than in a pop-up.

The button. A row whose courier trip is already booked wears the warning it always has, and that block now carries a second press: **See this order**, beside *"Call off the trip and add to this run"* rather than instead of it. It folds the order out underneath its row; **Hide the order** folds it back.

What drops down:

- **who** it is for, and their number; **what** they ordered - each line as it was sold, with the quantity and the price that line went out at; **where** it goes: the address, or the Point it collects from; **when** - the bake day, the day the van comes, and **your pickup time**, each on its own line; **their note**, when they left one; **what it comes to** - the same rows the Edit card and the invoice show; and at the foot, **Get a delivery price**, which asks the courier what that one order would cost.

Not one figure on it is worked out fresh. The money is drawn by the same renderer, from the same money function, that the Edit card, the Note / tracking card, the invoice and - most importantly - the customer's own message and tracking page all read. **A summary written only for this screen would have been a second figure, and a second figure is one that can disagree with the message you have already sent.** The days use the same wording those messages use, so the bake day and the van's day are said apart here exactly as they are there.

Nothing on it can be edited. This is the screen you work on while a van is out, and every control that changes an order already lives in one place - the Edit card this panel deliberately mirrors.

Your data is untouched. No SQL, no Edge Function, no upload. The suite is **2,854 tests, all green**, and the two rules that protect it were bitten: move the detail inside the row and it fails naming the trap; let its open state live in the drawing rather than beside it, and it fails too.

06 Oct 2026 - engine v341, A PICKUP TIME IS NOT A PROMISE (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"for courier lalamove, there is no delivery window open and delivery window closes promise, it is more for self collection point, and for delivery run that courier have to go many stop, and delivery time have to be a window reasonably. Pickup time is something user should specify."*

You are right, and this corrects what I built this morning. A single van delivery is not promised to a customer as a window. A window is the right promise only where a van has **many stops** (the run) or where a **Point's own hours** are the promise - and both of those already exist. What you fix on a courier order is the **pickup time**: when the van collects from you.

So the courier half of the card now asks both, and they are kept apart:

- **Pickup time** - new. The time the van collects from **you**. It fills in the **Delivery run's** own "the day and time the driver collects", so that screen no longer opens on a time the app thought of. **Delivery window (opens / closes)** - kept, and now plainly the **customer's** promise. It is set for one of two reasons: a **run with several stops** needs one, or **the customer themselves has told you the hours they are available** - *"from 2 to 5"*, say. Left blank, the message promises to confirm the time.

The pickup time never reaches a customer. A van collected from your kitchen at nine is at the customer's door later than nine, so nine is not theirs to be told. It is stored on the order, remembered, and used by the run - deliberately not published in the confirmation, in the messages or on the tracking card. A test pins exactly that, and it was bitten: let it leak and the test fails showing the customer being told , 09:00.

All three doors to a courier order ask it - **+ New order**, **Edit**, and **Note / tracking** - so whichever one you open, the same question is there and the same answer comes back.

Your data is untouched. No SQL, no Edge Function, no upload. The suite is **2,851 tests, all green**.

06 Oct 2026 - engine v340, A COURIER ORDER IS TOLD NOTHING AT PACKED (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"to make it simpler, can we do like this, when status changed to PACKED ... 2. for courier order, we take away the SEND SHIPPED MASSAGE. When status change to SHIP/COLLECT - Send ship message button, by that time, courier already collected, lalamove link can be shared, no more confusion."*

Done - and it removes a message that was saying something untrue. At **Packed** nothing has left your kitchen, yet the row offered a courier customer **"Send shipped message"**: *"your order is on its way"*, with a share link, before any driver had taken the job. That is the same confusion the tracking card was fixed for in v338.

Packed now offers:

- **self collect** - **Send pickup reminder**, unchanged. The bread really is ready, and this is what says so and where to fetch it. **courier** - **nothing at all.** There is nothing true to tell them yet.

Collected / Shipped is where the shipped message lives now - the one place it is offered, for a van order

and a posted parcel alike. By the time you mark it there the courier has the order in hand: the message is true, its Lalamove link works, and the customer's own tracking page shows that link too.

Your data is untouched and no screen moved. No SQL, no Edge Function, no upload. The suite is **2,846 tests, all green**, and the new rule was bitten - put the button back at Packed and it fails by name.

06 Oct 2026 - engine v339, BAKE DAYS MOVES TO THE SHOP (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"i think the bake days should not be at logistic, it should be in the Shop."*

You are right, and the reason is worth writing down. Logistic is the list of ways an order LEAVES your kitchen. A bake day is not a journey - it is **the day the shop is open**: the day the storefront's own calendar offers a customer, and the only days it will take an order for. So **Bake days** now sits under **The shop**, at the top of the group, because which days you are open comes before what you are advertising.

Logistic keeps the four that belong to it: Delivery run, Send a van, Self collection Points, Parcel couriers.

Nothing else moved, and no address changed - every bookmark, #/deliveries included, still opens the same screen. The suite is **2,844 tests, all green**, and the guard that catches a screen going missing in a regroup was bitten: take the row out and it fails, naming the screen that can no longer be reached.

06 Oct 2026 - engine v338, THE VAN GETS ITS OWN DAY AND TIME (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"and after we fix the courier delivery date and time in +add order or edit order, the card should remember"* - and, about the tracking page: *"the lalamove link should not be there because the driver might not be confirming, it only create more confusion if they were to click the link. Lalamove link and Delivery: Finding a driver, should not be send at this stage."*

This finishes the job that began with your customer asking whether his delivery was Wednesday or Thursday.

1 · The order now carries the van's own day and time, typed by you. + New order and **Edit** each gain three boxes in the courier half: **Courier delivery date**, and a **delivery window** (opens / closes). **They always open EMPTY.** Nothing is ever worked out from the bake day - the bake plan is a plan and not a schedule, as you told me in September.

2 · And they are remembered. This is the half that was missing. Until now an order had nowhere to write the day the van comes, so the app assumed the bake day. The day and the window are now stored on the order itself, and reopening **Edit** brings both back exactly as you left them. Blanking a box takes it back off the order rather than leaving an empty word behind.

3 · The customer's tracking card stops gluing the van's window to the bake day. For your case - baked Wednesday, van Thursday morning - it used to read Wed, 7 Oct ... 9-11 am, **a time on a day the van does not come.** It now names both days and puts the window with the van's own day; the message line is also called **delivery** rather than "pickup window", which is what your run screen has always called that value. **An order with no van day on it reads exactly as it always has**, so nothing already in your records changed.

4 · The Lalamove link and "Finding a driver" are not sent at that stage any more. Your words, and you were right - a link the driver may never take up only creates confusion. While a booked van is still **finding a driver**, the customer's card now says nothing about the trip at all: no link, no status. From the moment a driver is on the way, both appear as before. **Nothing is thrown away:** the order keeps its link and its status, your own screens are unchanged, and only the sending stops. A parcel is untouched - its consignment number is exactly what a customer needs.

5 · The Delivery run no longer hands you the bake day. Opening a run used to fill the **"day the driver collects"** box in for you with the bake day - so for Mr Tan's Thursday van it read **Wednesday**, written in as though you had chosen it, and that is the day that would have been booked. It now comes from **the day you already typed on the order**, and is empty when you have not typed one. The price box inside an order does the same.

One thing deliberately left alone: the run's **time** box still opens on your own dispatch time from Settings. That time is **when the van collects from you**; the time on the order is **when the bread reaches the customer** - different by the length of the journey - so seeding one from the other would send the van to your kitchen before the bread had left it. If you would rather it opened empty each time, that is a one-line change.

Your data is untouched. No SQL, no new column, no upload: the day and window ride on the order, and the customer's card carries them inside the text it already publishes. The suite is **2,843 tests, all green**, and three of the new assertions were bitten - put the old behaviour back and they fail by name.

06 Oct 2026 - engine v337, THE DAY IS CALLED THE BAKE DAY NOW (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"our bake day is Wednesday 7th Oct, customer want a lalamove delivery of Thursday morning 9am, so we book lalamove for them. The the delivery date i choose is 7th. When we send confirmation message to him he was confused, then asking whether the delivery date is wed or thurday."* Then *"1st i think we should change the orginal delivery date to bake day"*, and when I asked how far the new name should reach - *"Everything you read, screen name too"*.

Three things were wrong. The first one reached a real customer.

1 · The confirmation message called the bake day a delivery. A courier order's day is **the day you bake**; the van goes whenever you book it. So a van booked for **Thursday morning** was announced in the message as *Delivery: Wed, 7 Oct* - that is the sentence your customer read, and it is exactly why he asked you whether it was Wednesday or Thursday. **The message now says the two apart:**

- a **courier** order reads *Baking day: Wed, 7 Oct - Courier delivery*, and on its own line *Your courier delivery time will be confirmed separately* - because at that moment the app does not know the van's window yet, and it was stating one anyway; once you **have** booked the van, its window goes on **its own line** - never glued to the bake day; a **collection** order is unchanged (*Delivery: Wed, 7 Oct*), because your kitchen or a Point is genuinely where and when it changes hands.

2 · Every screen you read said "delivery date". It says bake day now. The More row and the screen title (**Bake days**), the Home dials' caption, the Orders calendar's answer when you tap a day you do not bake on, the Edit card's day picker, the PO's day ticks, "per bake day" on the products, the shared-data card, the backup copy, the delete and "Added N" messages, the Settings card - **22 files, 79 lines.**

3 · One signpost pointed at a screen that no longer exists. Found by building the guard, not by looking. The Orders calendar used to say *"Add it in More -> Delivery Dates"* while the row it was sending you to had already been renamed - **a dead end, and no test could see it**, because the message and the row were only ever read apart. It now reads **"Add it in More -> Bake days"**, and a new test reads **every** *More -> X* in the app against the rows and headings that really exist, so the next rename cannot leave one behind. **Proved by biting it:** put the old name back and the test names the file and the dead signpost; then restored byte-identically (sha256 checked).

Two things I deliberately did NOT change, and you should know which:

- **The customer's own words stay.** Your shop page, its calendar, its "Your delivery day: Wed, 16 Sep" line, and the notice rule ("Orders close (days before delivery)") are all unchanged - a stranger reads those, and this rename was about **your** screens. **One money line still says "delivery day"** - *"what is still to collect is counted by delivery day, because that is when you hand it over."* **Deliberate:** on a courier order the money is collected when the van delivers, which is not the bake day. Renaming that word would have made the sentence untrue.

STILL OPEN, and it is the same complaint: the customer's tracking card. The link in every confirmation opens a page that still glues the van's window to the bake day. **Known, not fixed** - it is the next piece of this work.

Your data is untouched. Nothing was renamed in the cloud: deliveryDateId, the day records, the #/deliveries address and every synced record keep their old internal names - **only the words on screen changed.** No SQL, no Edge Function, no redeploy. The suite is **2,826 tests, all green**.

06 Oct 2026 - engine v336, YOUR OWN REWARD COUPON NOW COMES OFF YOUR ORDER (no database step, nothing to upload - pushing this one is the whole of it).

Your words: *"when i refer a friend and get a coupon, but redeem that coupon will not reduce my order price"* - **and you were exactly right, again.**

What was wrong. There are **two kinds of coupon** in the app, and only one of them was ever being read:

- **The friend's** - RM3 off *their* first order. That is the one I fixed in v322. **Yours** - the RM3 you earn for bringing someone in. **Nothing ever took that one off anything.**

So when you pressed "Apply coupon" on your own order: the coupon was used up, the app said *"already taken off this order"* - **and the price did not move.** The coupon was spent and gave you nothing.

Why the two behave differently, and it is a real difference. The friend's coupon is earned and spent **on the same order**, so the order is written on it. Yours is earned on **the friend's** order and spent on **one of yours, later** - so the order number written on it is the wrong one to look for. **The app was looking for the friend's kind of coupon and never for yours.**

Now fixed. Pressing **Apply coupon** on your own order marks the coupon *and records which order you spent it on* - so the money function can find it, and **the Total comes off.** Your receipt and the customer's message also now call it a **bring-a-friend reward** rather than "the one you were sent", because it is your own coupon on your own order.

Careful about two things it must NOT do, and both are tested: a coupon you are **holding but have not applied** comes off nothing - it is yours to spend when you choose; and a coupon you spent on **one order** cannot come off a second one.

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,825 tests, all green**.

05 Oct 2026 - engine v335, THE SELF-CHECK NOW COVERS EVERY SCREEN (no database step, nothing to upload - pushing this one is the whole of it).

I said yesterday that the new check left three screens out. This closes them - and finding out why taught me more than the check itself.

It was blind to three kinds of control, not three screens:

1 · Tick boxes and switches. The bring-a-friend on/off, the day ticks on the Purchase Order, every product toggle. **The check walked straight past them.** A switch whose handler breaks is exactly the kind of fault it exists to find.

2 · Number boxes. The amount fields, the daily limits, the discounts. Same - skipped.

3 · Screens that fill themselves in. The Reviews screen asks the cloud for the waiting reviews and draws its Publish and Delete presses **only when the answer arrives.** The check looked the instant the screen was drawn, saw an empty card, and reported "nothing to press" - **when what it was really looking at was a screen that had not answered yet.** It waits now. **That is the one worth knowing about: it is a whole class - the reviews, the promo counts and the tracking card all arrive that way.**

What it walks now: all twenty-five screens, about 5,800 presses, nothing skipped - and it still costs about four seconds. **And it still catches yesterday's fault**, which I put back to be sure: *revealed -> BUTTON "X" ... g is not defined*.

And the lesson, which is the same one twice over: my check was reporting on its own blind spots and I nearly read it as the app being fine. It says what it did NOT walk rather than passing quietly - and that is the only reason I found any of this.

Your data is untouched. No SQL, no Edge Function, no key. **No app code changed** - this is the check itself. The suite is **2,822 tests, all green**.

05 Oct 2026 - engine v334, THE APP NOW CHECKS ITSELF BEFORE IT REACHES YOUR PHONE (no database step, nothing to upload - pushing this one is the whole of it).

You asked: *"you work is not checked, how can you improve?"* - so I built the answer.

What it does. Before every version leaves my hands, the app now **opens all twenty-five of its screens one after another, presses every button and link and dropdown on each, and types into every text box** - then presses whatever those presses reveal, including a confirmation's Yes. **Anything that throws is a failure, and it says which screen, which control, and what it was doing when it broke.**

And the proof it works: I put yesterday's fault back to see whether the check would catch it. It did - and this is what it said:

/orders · revealed -> BUTTON "X" (after INPUT "Find an order - name, #code, phone...") · g is not defined

That is yesterday's dead button, named, on the screen it was on, with the press that revealed it and the error it threw - the same fault you found by asking. It is now found automatically, every time.

WHAT IT DOES NOT CHECK, said plainly. It is not a browser: it cannot see whether a thing *looks* right, whether it is in the right place, or whether a motion works - that still needs eyes on the real phone. **It catches a control that breaks when used. It cannot tell you a button is ugly, or that a number is wrong.** And three screens come up empty to it today (Purchase Order, Bring a friend and Reviews) because the test data I feed it does not give them anything to work with - **it says so rather than passing quietly, and I will close those three.**

What it found on its first run, in my own test setup rather than the app: four faults that were the stand-in's, not the app's - which is exactly the class of thing it exists to sort out.

How fast: it walks all twenty-five screens and makes about two and a half thousand presses **in about one second**, so it runs every time with the rest of the checks and costs you nothing.

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,822 tests, all green**.

05 Oct 2026 - engine v333, THE REMOVE BUTTON I ADDED YESTERDAY DID NOT WORK (no database step, nothing to upload - pushing this one is the whole of it).

You asked the question that mattered: *"have you tested it really can delete?"* - and the answer was no. I had checked that the X was *there*, and that the row opens. **I had never pressed it through to a deletion. So I did - and it was broken.**

What was wrong. The remove button I put on a search result was wired to a variable name that does not exist in that part of the code - the name belongs to the New-orders box next door. Pressing it **threw an error and removed nothing at all.** A button that does nothing is the thing this app treats as a bug everywhere else, and I had shipped one.

Now fixed and properly tested - by pressing it. The test now **presses the X, answers the confirmation, and checks the order is actually gone**, from **both** doors: the New-orders box and a search result. It also fails with the old mistake put back, so it is a real check and not a hopeful one.

And it was worth asking for a second reason: the New-orders box's own X has been there since the box was built and had **only ever been checked for existing** - never pressed. It works, and it is now covered too. **Two buttons call the same removal; testing one proves nothing about the other.**

On v332's other half - nothing was lost. Opening an orphaned order was working, and still is. **No order was ever removed by the broken press**, because a press that throws removes nothing.

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,821 tests, all green**.

05 Oct 2026 - engine v332, AN ORDER WITH NO DELIVERY DAY CAN BE OPENED - AND A SEARCHED ORDER REMOVED (no database step, nothing to upload - pushing this one is the whole of it).

Two of your reports, and they were the same dead end met from two sides.

What an "orphan" is. When you delete a delivery day that still has orders on it, the app asks first and promises: **"The orders are kept in your delivery history."** You agreed, so they are kept - **but the day they belonged to is gone, and until now that left them stranded.** The row could not be opened (no arrow, no jump), so **the one thing you could do with one was delete a real sale** - and if you did not want to delete it, you could do nothing

at all.

Both halves are fixed:

1 · The row now opens the order. Tap an orphan in the New-orders inbox, or a search result, and its **own Edit card** opens - the one with the **Delivery day** calendar in it. Pick a day, press Save changes, and the order is back where it belongs. **Nothing is deleted.** And that row now says what it is: *"its delivery day was removed - tap to put it on one"*.

2 · A search result can be removed. Your words: *"C2FDA5 i search this order, but no button to delete it."* **A search result had no controls at all** - no arrow, no X - so the one place you go to find an order by its code was the one place you could not act on it. **It carries the same X the New-orders inbox has always had**, and it still asks before it removes anything.

You chose this route yourself, and it is the safe one: **every orphan can be opened, put back on a day, or removed one at a time - and nothing leaves your sales or your Profit statement unless you remove that one order yourself.**

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,819 tests, all green**.

05 Oct 2026 - engine v331, A CODE SEARCH FINDS THAT ORDER AND NO OTHERS (no database step, nothing to upload - pushing this one is the whole of it).

Your report: *"C2FDA5 why when i type this 17 order found?"*

What was wrong. The search box has a second way of matching, added for a phone number typed with dashes or spaces - "*012-345 6789*" cannot match a number stored as "*60123456789*" character for character, so the box strips everything that is not a digit and looks for those digits in the order's number. **An order code was falling into that same path.** "*C2FDA5*" carries two digits in it, so stripping the letters left "*25*" - and the search then found **every order whose phone number contains 25**. In a Malaysian book that is most of them: **seventeen, on yours.**

And it is worse than a long list. The one row you asked for is buried inside it, so a search that returns seventeen is a search that has stopped working.

Now fixed. A query with a letter in it is matched as text and nothing else - and a code never needed the number path anyway, because "*#C2FDA5*" is already in what the search reads. **Typing a code now finds exactly that order, with or without the #, in any case.**

And the number path still does its own job - "*012-345*", "*0123456789*", "*60123456789*" and "*016 555 7777*" all still find their order, exactly as before. **Both halves are pinned as a test**, so neither can be broken by fixing the other.

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,818 tests, all green**.

05 Oct 2026 - engine v330, THE BRING-A-FRIEND DISCOUNT SHOWS IN THE WORKING (no database step, nothing to upload - pushing this one is the whole of it).

Your report, and you were right about both halves: *"the discount dnt show in the total adding in edit, probably other place?"*

What was wrong, and it was worse than a missing line. v322 took the bring-a-friend discount off the order's Total and named it on the customer's message - **and left it out of the Edit card and the Note / tracking card altogether.** Those two do not read the order's Total; **they work the total out themselves** from the lines you are editing. So **the Total on the Edit card was the total BEFORE the discount** - and the same figure was missing from your receipt and from the invoice.

That is the figure you read while editing an order and quote to a customer. It now reads:

Items total - RM 30.00 · Bring-a-friend discount - -RM 3.00 · Total - RM 27.00.

Both cards, the receipt and the invoice now show the discount as its own line, right above the Total it moves - and the Total itself has it taken off.

The other place, answered. I swept every screen that shows an order's money. The discount was already named on the customer's WhatsApp message and on their tracking page. **The Edit card, the Note / tracking card, the receipt and the invoice were the ones missing it - and they were the ones that list the money as a sum you can check by adding it up.** The two pop-ups now price the discount through the same single rule the rest of the app uses, so **a figure she reads while editing and the message the customer gets can never be two different figures.**

AND THE SWEEP FOUND A SECOND FAULT, which I have also fixed. If a coupon was worth more than the order - the rare case, and a hand-given one could do it - the receipt said it had taken off the whole amount while the Total stopped at zero, **so the lines did not add up and the figure claimed more off than the order was worth.** It now takes off only what is really there: **a coupon can never make an order cheaper than nothing, and the line and the Total always agree.** That fix reaches the customer's message and their tracking page too, because all of them read the one money function.

And the rule is now a test, not a hope. The receipt's own check **adds up every line and compares it with the Total** - on a plain order, with a code, with a courier charge, with a COD charge and with a discount - so **the next time anything comes off an order and is not given a line, the suite goes red instead of the number quietly not adding up.** That is the check that found the second fault above, within seconds of being written.

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,816 tests, all green**.

05 Oct 2026 - engine v329, THE DAY THAT WOULD NOT OPEN: THE ACTUAL CAUSE, FOUND AND FIXED (no database step, nothing to upload - pushing this one is the whole of it).

The message you read out was the whole answer: "cur is not defined".

What it was. When I added the bring-a-friend discount on **v322**, one line that names the discount on the order's row used a value called cur - **and I never declared it there.** It was declared in the two blocks *next to* it, not in the one that used it. So the moment a day held an order that had been given a bring-a-friend coupon, drawing that day's panel **crashed**.

And that is why it looked like what it looked like. The calendar square is repainted a moment before the panel is built, so a crash while building the panel left **the square red on the day you tapped and the panel still showing the day before.** Exactly your recording, and your three words for it - **"the day like hang"**.

And it was never about 5 October or 7 October. It was about **any day holding an order that had been given a bring-a-friend coupon.** Only 7 October had one, which is why only 7 October refused. **Every day with such an order had the same fault** - and now none of them do.

I have to be straight with you about the road here. My first two answers (v326, v327) were guesses about your two days sharing an internal id. **They were wrong, and they cost you two pushes.** What finally worked was **v328 making the app report the error itself** - one tap and you handed me the exact line I could not see from the outside. **That is the lesson I am keeping.**

What stands from v326 and v327, and it is kept because it is good on its own: the day you are on now travels in the page's address (so a bookmark or a shared link opens the right day), the screen remembers which day each id meant, and a day whose records ever do share an id is separated at load, losslessly. **None of that was your fault** - but none of it is harmful either, and two of the three are small improvements in their own right.

And v328 stays, and it earned its place: if a day ever cannot be drawn again, the square will not turn red and the reason will name itself on the card, in one line you can read out. **It is the reason this fault is fixed at all.**

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,813 tests, all green**.

05 Oct 2026 - engine v328, THE DAY THAT WILL NOT OPEN NOW SAYS WHY (no database step, nothing to upload - pushing this one is the whole of it).

Why, and it is an honest one. You pushed v326, then v327, and told me both times that 7 October still would not open. **Two versions of fixing it from the outside got nowhere - and that is on me, because I was guessing.** What I could see was the shape of the fault from your recording: the day turns red and the panel underneath never changes. What I could not see was **why**, and the reason is in that one day's own records,

which I cannot look at.

So this version stops guessing and makes the app speak.

The day is now built BEFORE the red mark moves. If it cannot be built:

- **the red mark does not move** - because a red day sitting over a different day's panel is exactly the fault you have been looking at, and it should never have been possible; and **the reason takes the day's place**, in the app's own words, with a line to read out to me.

What you will see. Tap 7 October on the new version and you will get one of two things:

1. **The day opens normally** - a plain square with its heading, its count and its orders. Nothing to do; it is fixed.
2. **A card headed "This day could not be opened"**, naming the day and ending with **"Tell the baker's helper this message: ..."**. **Read that sentence out to me and I will have the answer in one step** instead of another round of guessing.

Whichever it is, nothing is changed and nothing is lost - the card says so itself. 7 October stays in your list either way, and no order, product, price or bake day is touched.

If you get the second one, the message is what I need - not a description of it. The words after "this message:" are the whole answer.

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,812 tests, all green**.

Added 05 Oct 2026, the same day: you tapped 7 October and it said "cur is not defined" - and that line was the answer. See v329 above: a missing declaration added in v322, crashing the day's panel whenever an order had been given a bring-a-friend coupon. This version is what surfaced it.

05 Oct 2026 - engine v327, THE DAY STAYS PUT WHEN THE SCREEN REBUILDS ITSELF (no database step, nothing to upload - pushing this one is the whole of it).

Why. You pushed v326 and told me plainly: *"the 7/10 still not working right."* **You were right, and my fix had a hole in it.** I found the hole by writing the test I should have written the first time, and it failed with your bug in one line: expected 'Thu, 10 Sep 2026' / actual 'Mon, 7 Sep 2026'.

What I missed. v326 made the day you tap travel with the tap - that part was right, and it is why the calendar square turns red correctly. But **the address the app writes carried only the day's internal id**, and the Orders screen **rebuilds itself from that address** every time the cloud answers, a pull lands, or the app regains focus. On that rebuild there is no tap and no remembered day - only the address. So the screen came back to the id, the id is shared by two of your days, and the panel opened the other one again.

And that is why you never saw it flicker: the rebuild happens in the same instant as the tap, before the phone has drawn anything. **The wrong panel was the only picture that ever reached your screen** - which is exactly what your recording showed: the red mark moves, and the panel does not change by a single pixel.

Now fixed twice over:

1 · The address carries the day, not only its id - #/orders?date=...&day=2026-10-07. A rebuild reads it back and opens the day you tapped. **A bookmark or a link you share now opens the right day too.**

2 · The app remembers which day each id meant, for as long as you are on the screen - so every other rebuild does the same: changing an order's status, saving a day's availability, and anything else that redraws the screen. **Thirteen separate places rebuild it; all thirteen are covered by that one memory.**

And this fix does not depend on the v326 repair having run - it holds whether or not your two days still share an id.

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,811 tests, all green**.

Added 05 Oct 2026, the same day: this did not fix it either, and the cause was never what this entry says. The real fault was v329 - a missing declaration from v322 crashing the day's panel whenever an order had been given a bring-a-friend coupon. Read v329 above. **The id-repair and the day-in-the-address described here are kept because they are small improvements in their own right, not because they were the answer.**

05 Oct 2026 - engine v326, A DAY YOU TAP IS THE DAY YOU GET (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your report: *"why the order calander not able to select 7/10/26?"* and then, exactly: *"that day shown 2/12, clicking that date, the date turn red, but the SET day's avaibility not changing to 7/10/26."* **You were right, and your recording showed it happening - four times in thirty seconds.**

What was happening. Each day in the calendar is drawn from that day's own record, so the **red mark followed your tap** - it always did. But the panel underneath was looked up a different way: by the record's **id**, and that lookup answers with the **first** record holding it. **Two of your delivery days are carrying the same id**, so tapping one of them turned the day red and then opened the *other* one. Pressing **5 Oct** worked and **9 Oct** worked because they are not the pair - **7 Oct** was the day it happened to.

Two fixes, and you get both:

1 · The panel now opens the day whose cell you tapped - the day's own date rides along with the tap, so the calendar and the panel under it can never be two different days again.

2 · The two days are given their own ids when the app loads - nothing is deleted, every day and every order and price survives, and **each order follows its own day** (its own saved delivery date is what tells the two apart). Your **7 Oct becomes openable, editable and deletable again** - and so does anything else the shared id was quietly affecting.

AND IT IS FIXED THE SAME WAY ON EVERY PHONE. The replacement id is worked out from the old id and the date rather than rolled fresh, so both of your phones reach the **same** answer. Two different answers would have put the pair back with the next sync.

How it got that way I cannot say - I could not find anything in the app that makes two days share an id. What I can say is that the fault it caused is now impossible to repeat from your side, and if it ever happens again the screen will simply follow your tap.

Your data is untouched. No SQL, no Edge Function, no key. The suite is **2,809 tests, all green**.

Added 05 Oct 2026, the same day: this version did not fix it, and v327 is the rest of it. Making the day travel with the tap was right, but the app writes only the day's id into the address and then **rebuilds the screen from that address** - and on that rebuild the day was lost again. **v327 carries the day in the address and remembers it, and the fault is closed. Read v327 above.**

05 Oct 2026 - engine v325, THE FRIEND'S NUMBER COMES OFF THE NEW CUSTOMER'S ADDRESS (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your question, and it was a good one: *"for new customer clicking link from his friend, after he place an order have you remove his page linking his friend phone number?"* **I had not.**

What was happening. A referral link is /store/?via=60123456789 - **the friend's own phone number.** The shop reads it once, stamps it onto the order as *"this one came through their link"*, and then **the number just stayed there**: in the new customer's address bar, in their browsing history, and in anything they copied out of the address to pass on. After the order, it is only a phone number being carried around.

Now it comes off the address the moment the order is placed.

AND THE MOMENT MATTERS, which is why it is not simply done when the page opens. A customer can arrive by the link and browse for ten minutes before ordering - taking the stamp off on arrival would lose the referral altogether. So it goes when **the order really landed**, which is the same branch the app already uses to remember the order was placed. **An order that fell back to WhatsApp keeps its stamp on purpose** - nothing was recorded, they may try again, and the message they send carries it.

Only via is removed. If the same customer is also looking at an old order, track is a separate key handled by its own rule, and each removes only itself.

Your data is untouched. No SQL, no order, product, price or bake day moved. The suite is **2,805 tests, all green**.

05 Oct 2026 - engine v324, A TRACKED ORDER STOPS OWNING THE PAGE, AND YOU CAN FORGET A STRAY CUSTOMER (no database step, nothing to upload - pushing this one is the whole of it).

Two things you asked for, and they ship together.

1 · THE ORDER CODE COMES OFF THE ADDRESS BAR. Your words: *"when a customer track his order, his store version became associated with that order code. Anyway to take it off automatically?"*

The tracking link opens as /store/?track=CODE, and **that code used to stay in the address bar.** So the next time that page was opened - a bookmark, a history entry, a link re-shared - it landed back on **that one order's card** instead of on the shop. **The page belonged to an order rather than to the bakery.**

Now the code comes off the address the moment the card appears, and the card stays up for that visit. The link still does exactly what it is for. The cost, said plainly: **a refresh now lands on the shop** rather than on the card - which is the direction you asked for, and their code is in their WhatsApp if they want it again.

AND ONE THING THAT WOULD HAVE BROKEN IN SILENCE: via is left alone. That is the friend's referral stamp, and a customer can arrive **by** a referral link and **then** track an order - stripping the whole address would have quietly broken bring-a-friend for exactly that person. Only track is removed, and that is a test.

2 · FORGET - for the strays. Your words: *"i need a button to delete a customer as well, i found there is few stray customer."*

READ THIS BIT, because it is why the button is where it is and not everywhere. Your customer book is **built from your orders.** Someone you added by hand has never ordered, so **their row *is* the record you typed** - forget it and they are gone, name, number, reward and note, with a question that names exactly what goes and says **"nothing else in your book is touched."**

A customer who HAS ordered cannot be deleted from here, and it is not an oversight. Their row is their **order history**. Removing the record would leave the row standing - it comes straight back from the orders - while quietly throwing away their reward, their note and their dog's name. **A button that did half of what it said would be worse than no button**, which is your own rule about dead controls: two rows that look alike must behave alike. **For two records of one person, use Merge** - that is what it is for.

Your data is untouched. No SQL, no order, product, price or bake day moved. The suite is **2,804 tests, all green**.

*(One note: this commit carries two version numbers - v323 and v324 - because the tracking-card line and these two were built back to back and both were still uncommitted. The Engine on your More screen reads **324**.)*

05 Oct 2026 - engine v323, THE TRACKING CARD SAYS WHY THE TOTAL IS LOWER (no database step, nothing to upload - pushing this one is the whole of it).

Why. You asked for it after the last version: the customer's own tracking page was showing the **correct lower total with nothing saying why.**

What they see now. On the tracking page, under the items and the courier charge, a line in the customer's own language:

Bring-a-friend you were sent: -RM 3.00

- sitting where the promo line sits, so the figures above it still add up to the Total below. In Mandarin it reads 朋友推荐优惠：-RM3.00, and in Bahasa Malaysia *Bawa rakan yang menghantar anda: -RM3.00*. It is the same sentence your WhatsApp confirmation uses, so the two agree word for word.

And it needed no database step, which is worth explaining because it nearly did. The obvious way was a new column on the tracking table - and the app's own code warns against exactly that: *"a missing column kills publishing for EVERY order silently."* **So it reuses the discount column that is already there**, which is honest because **a code and the friend's coupon can never both apply** - that column holds *the* discount on the order, whichever it is.

The small piece of cleverness, and why it is safe. The card tells the two apart by the code being **absent**: a discount with no code beside it can only be the friend's, because the app refuses to let a promo code be labelled

without a name. If that ever stopped being true, a code would start wearing the friend's words - **so that rule is a test.**

Your data is untouched. No SQL, nothing to upload. The suite is **2,799 tests, all green**.

05 Oct 2026 - engine v322, THE NEW CUSTOMER'S DISCOUNT ACTUALLY COMES OFF NOW (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your two reports: *"the bring a friend discount used but not really create a discount for that new customer"*, and then *"there store front copy still hold the discount in cache."*

You were right on both, and the first one was worse than a wording slip.

What was happening. When you pressed **Give coupon**, the app wrote the friend a coupon - *"First order - via Jienny's link"* - and then **nothing ever read it**. The order's Total, the confirmation and the three later messages are all worked out by one function, and that function knew about promo **codes** and nothing else. So the coupon was recorded, the order was priced as though it did not exist, and pressing **Apply coupon** told you *"already taken off this order"* about a figure **nothing had ever taken off.** An app telling you it did something it did not do is the worst kind of fault, and I wrote that sentence.

What you see now.

- **The RM3 comes off the order's total by itself.** The customer pays the lower figure, so your money and your message agree. **The message says so, in its own line** - *"Bring-a-friend you were sent: -RM3.00"* - sitting with the items and Courier charge so the lines add up to the Total beneath them. A discount the customer cannot find in the message is one they will ask you about. **The order says it too** - *" Bring-a-friend - RM3.00 already off this order's total"* - so the lower figure is never unexplained on your own screen either. **It cannot be used twice.** The coupon is tied to the one order it was born on, so their next order is priced normally no matter what any cached copy of a page happens to hold.

And your second report was a real gap of its own. Everywhere else in the app, anything that changes what a customer sees republishes their tracking page. **The Give coupon press did not** - so the customer's own card went on showing the price they were quoted *before* the discount. That is fixed: giving a coupon now republishes their card, exactly like changing the items or the address does.

Her own rule is honoured, unchanged: one discount per order - and if a customer typed a **code**, **the code wins** the total, because that is the one they typed and can see. **When that happens the friend's coupon is not spent** - it stays theirs for the next order, rather than being quietly burned.

One thing left, and I would rather say it than let you find it. The customer's **own tracking page** now shows the correct lower Total, but it does not yet say *why* - because adding that line there needs one small database step. **Their WhatsApp confirmation, which is where they are asked for money, already explains it in full.** Say the word and I will do the tracking-page line as its own version, with the one-line SQL for you to run first.

Your data is untouched. No SQL needed for this version, no order, product, price or bake day touched. The suite is **2,796 tests, all green** - including the one that would have caught this on the day it was written: not "does the coupon say the right thing" but **"does it actually come off the total"**.

05 Oct 2026 - engine v321, THE OFFERS WERE NEVER BROKEN - YOUR MAC ASKED FOR NO ANIMATION (no database step, nothing to upload - pushing this one is the whole of it).

THE ANSWER, AND IT CAME FROM YOUR RECORDING. Six seconds of your screen, one turn pulled out frame by frame at a tenth of a second: the message changed **between two frames**, with no half-way picture in between. The offers were changing the whole time. **What was missing was the animation - and the reason is a setting on your Mac.**

Your Mac has "Reduce motion" turned on. That is an accessibility setting in *System Settings -> Accessibility -> Display*. When it is on, the app is not allowed to animate, so it was showing you each new message **instantly, with nothing to see** - which is exactly what "no effect" looks like. The message *was* changing every 2 seconds; it just arrived with no movement at all.

AND THAT IS WHY I KEPT GETTING IT WRONG - I HAVE TO TELL YOU THIS PROPERLY. The very same setting is switched on in the browser I test in. So every check I made was made under your exact condition, and every one of them said "correct" - because the app *was* correct. I was reading where things ended up and calling it verified, and the one thing your setting removes is precisely the thing I was never able to see. **Four versions went by on that mistake.** I am sorry.

What changed - and it means you get an effect either way. The house rules say: when someone asks for less movement, take away the **movement**, and allow **a short fade - up to a fifth of a second - instead of nothing.** The app was taking away both. Now the new message **fades in over 0.18 seconds** rather than appearing in a single frame. No travel, no spinning, nothing to make anyone dizzy - but you can see it change.

IF YOU WOULD RATHER HAVE THE FULL SLIDE, where one message scrolls up and the next follows from below:

1. Open **System Settings** 2. Go to **Accessibility** -> **Display** 3. Turn **Reduce motion** off

The offers will then slide exactly as designed. **With it on you now get the fade instead of nothing**, so both ways round you will see something.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,789 tests, all green**.

And one thing said plainly rather than buried: I still cannot *watch* the animation from my side. What I can say for certain is what the code computes - the panel that arrives is in its place and fades from nothing to full - and that the setting which was hiding it is named above. **You are the one who can see it, and you have the ?debug=offers line if anything ever looks wrong again.**

05 Oct 2026 - engine v320, A FINGER CAN NO LONGER STOP THE OFFERS FOR GOOD (no database step, nothing to upload - pushing this one is the whole of it).

First, plainly: I owe you an apology on method. I told you twice that this was fixed, and both times I had never actually watched it move. Every check I made was on a browser pane that reports itself as **hidden**, and **a hidden tab runs no animation at all** - so all I could ever read was where things *ended up*. That is how two versions went to you broken. It will not happen again in this particular way, because there is now a way to ask the strip itself (below).

AND I HAVE FOUND THE ONE THING IN THE CODE THAT CAN STOP THE OFFERS FOR GOOD.

The strip holds still while your pointer rests on it - that was your own ask, so you can finish reading a message. The app remembers *"the pointer is on it"* as a little flag, and it clears that flag when the pointer **leaves**. **On a touch screen, that "leaves" may never arrive.** A finger landing on the strip sets the flag, and if the matching leave never comes, the strip is held **for the whole life of the page**: the timer keeps ticking, everything thinks the pointer is still there, **and the messages never change again.** That looks exactly like a broken strip - and it is the only mechanism in the code that can stop the turning dead, which matches what you are seeing.

What changed. A pointer can now hold the strip only where a pointer can hover - that is a mouse. **On a phone the strip simply carries on**, because on a touch screen there is nothing to rest and so nothing to hold. On your Mac, resting the mouse on it still holds it, exactly as before.

AND A WAY FOR YOU TO ASK IT DIRECTLY. If the offers are ever not turning again, add **?debug=offers** to the end of your shop's address, like this:

`https://jienluv2bake.com.my/store/?debug=offers`

One small line appears under the strip and tells us exactly what is going on - how many offers it can see, whether the timer is running, whether it thinks your pointer is resting on it, whether your phone can hover at all, and whether the tab is hidden. **No customer ever sees it: it only appears with that ?debug=offers on the link.** Next time, you read it to me and I will know, instead of guessing.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,788 tests, all green**, including two new ones: that a touch can no longer hold the strip, and that the debug line exists and can never reach a customer.

05 Oct 2026 - engine v319, THE FOLLOWING MESSAGE WAS BROKEN ON A TWO-OFFER SHOP (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your report, three words: *"there is no effect."* **You were right and I had broken it** - and the way it was broken is worth reading, because your shop is exactly the case it broke on.

WHAT WAS WRONG. When a message is replaced, the app has to know WHICH one is leaving so it can send it up through the top. It worked that out by counting backwards from the one showing: *the previous one*. That was inherited from the old 3D flip, where two panels leaving by opposite doors made sense, and it is fine when **three or more** offers are running.

WITH EXACTLY TWO OFFERS, "THE PREVIOUS ONE" IS THE MESSAGE THAT IS ARRIVING. So the message about to come in was parked **above** the window - where the one that had just left sits - and it then travelled **back down** into place. **Nothing was ever left waiting below**, so "one message following another" could not happen at all. A two-offer strip is not an edge case; it is the ordinary one, and it is what yours runs.

How it hid, and that is the uncomfortable part. With three offers running it looked perfect, so the page I checked and the harness I built both showed the effect working. **And 2,784 tests were green while it was broken on your phone.** Nothing was watching the two-offer case, which is the case that matters most. **That is now the one the test exists for.**

MEASURED, ON THE REAL SHOP, AT 375, TWO OFFERS:

	at rest	after a step
before	[0, -183] - **nothing waiting below**	the arriving message came from **above**
now	[0, +183] - **the next waits below**	it comes **up from below**, following the one leaving

And the second half of the fix. A message that has just left through the top has to be back at the bottom to arrive from below next time - and moving it there *with* an animation would slide it the whole way **through the window**, sweeping the message you just read back across the strip in the wrong direction. So every waiting message is now put back in its place **instantly, with no animation**, before the movement starts. A jump paints only where it lands.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,786 tests, all green** - two of them new, and the one that matters asserts the two-offer case by name.

05 Oct 2026 - engine v318, EACH MESSAGE IS ITS OWN BOX AND THEY SCROLL UP (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your words, and they were the whole spec: *"Maybe you box up each message, when 1st message start to scroll up, the 2nd message is following, So effectively you see 2 message, one follow by another, it scroll up, stop 2sec, scroll again until mouse over."*

What I had wrong. v316 sent them **sideways** - I read "scrolling off" as off to the side. You meant **up**. Both readings of *"box up each message"* went onto a page side by side, you picked the one where **every message is its own box**, and that is what is built.

What you see now.

- **Each message is its own amber box.** The box used to belong to the strip - one box that never moved. It has moved onto the messages: the strip is now just the **window** you look through, with no box of its own, so nothing is drawn twice. **They scroll UP.** The message being replaced leaves through the top and the next follows from below, with a **small gap passing between them** as they go - that gap is what makes it read as one message following another rather than one block of text sliding. **Two messages are visible during the move**, which is exactly what you described. **Then it stops for your 2 seconds**, and **rests while your pointer is on it** - unchanged.

And the rule that keeps your page still survives again. Measured on the live shop at 375: **206.6 pixels, identical at every step.** (It was 210.6 before - four pixels shorter, because the box moved off the strip and the strip stopped adding padding of its own.)

AND THREE MOTIONS HAVE NOW COME AND GONE, ALL THREE MINE: a fade (v292), a 3D flip (v296), a

sideways slide (v316) - each with reasoning I wrote into the code about what "reads as a glitch". You are the one who watches this every day. The code and the test now say so, so none of them quietly comes back.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,784 tests, all green**, and the rewritten test holds every part of this: the motion is up, each message carries its own box, the strip carries none, the gap between the two, and no fading.

05 Oct 2026 - engine v317, EACH OFFER HOLDS FOR 2 SECONDS (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your words, right after the scroll went in: *"maybe the scrolling is too fast and hardly see the results."*

YOU WERE RIGHT, AND IT WAS NOT MARGINAL - I measured your own shop before changing anything:

- your long offer runs to **38 words** (the code's line **plus your own sentence in the chalk hand**) - about **11 seconds** to read at a comfortable pace the short one (*AZQKJ*) is **14 words** - about **4 seconds** each of them was getting **1.5 seconds**, of which **0.4 is the move itself**

So the long one had about **a tenth** of the time it needed, and the short one about a quarter.

BUT THE SCROLL WAS NOT WHAT WAS TOO FAST. The move is 0.4 of a second, which is already the **top of what the house rules allow** (250-400ms for a change like this) - slowing the animation would have made it draggy and bought you **not one extra second of reading**. What was too short was the **stillness**: 1.1 seconds before, **1.6 seconds now**.

The pace is 2 seconds an offer - your number. It was 1.5s, chosen back when the line was mostly a short code; your offers have grown since, which is what changed the arithmetic.

One thing said plainly rather than left for you to find: 2 seconds is still under the 11 your long offer wants, so it will still turn before you have read all of it - and the move now takes a fifth of each turn. **Watch it on your phone with two codes running.** If it still turns too soon, **it is one number and nothing else needs touching** - say the word and I will put it wherever you like.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,784 tests, all green**, and the test that pins the pace now records that 2 is *your* number, chosen with the arithmetic in front of you.

05 Oct 2026 - engine v316, THE SHOP'S OFFERS SCROLL INSTEAD OF TURNING (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your words: *"the promo code, can it be like scrolling off, and new code follow, that kind of animation."* The offer line under your bakery name used to **turn** - a quarter-turn flip, the old offer tipping away over the top while the next rose from below.

What you see now. The offer being replaced **slides off to the LEFT** and the next one **follows it in from the right**, at the same 1.5 seconds an offer. The strip is clipped, so nothing shows outside it. Everything else about it is untouched: the dots, the pointer resting on it to hold it, the tab going quiet when you leave it, and **nothing turns at all when only one code is running**.

AND THE FLIP WAS MY CHOICE, NOT YOURS - that is worth saying plainly. Back in v292 I picked a fade, then turned it into a 3D flip, and I wrote the reason into the code: that sliding one line of message sideways "reads as a glitch". You are the one who looks at this every day, and you asked for the scroll. The argument loses.

ONE HONEST TRADE, so it is not a surprise. During a scroll **two offers are briefly on screen at once** - the tail of the one leaving and the head of the one arriving. That is exactly what makes it read as a scroll rather than a swap. The flip never showed two, so if a long offer like *"Free delivery within Penang on RM200.00 and above - use code FREEDEL"* looks busy mid-move, that is why.

AND THE ONE THING THAT MUST NEVER REGRESS WAS RE-MEASURED. v295's whole point was that the strip is exactly as tall as your tallest message and never changes height, because the page under it used to

jump up and down every 1.5 seconds. **Measured on the live shop across three turns: 210.6 pixels, 16 samples, never once different.** A sideways slide is a transform, just as the turn was, so the guarantee carries over exactly.

Two corrections to the guide while I was in it. It described the old turning motion, and it said *"on a phone a tap holds it still for about twenty seconds"* - which **stopped being true in v297**, when we removed exactly that because it trapped anyone who touched the strip. Both now match what the shop actually does.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,784 tests, all green**, including the rewritten strip test: the flip coming back, a fade creeping in, an unclipped rotor, and a move slower than the house 250-400ms band each turn it red.

05 Oct 2026 - engine v315, TWO MORE CARDS LEAVE SETTINGS, AND ONE OF THEM IS RENAMED (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your two instructions: *"move Courier(Lalamove) and Mailing labels(courier), both to Logistic?"* and *"move Message Style to the shop"*. Both are the same idea as v314 - a thing belongs where its work is, not under "defaults, backup, transfer".

What moved.

- **Send a van** is now its own screen under **Logistic** - your own door for the driver, and which Lalamove account this phone is talking to. Same card, same pickup pin, moved whole. **Message style** is now under **The shop**, beside Bring a friend, Promo codes and Reviews. How your words reach a customer is the same subject as the offers and the reviews.

Two words I did NOT use, and you should know why. The first choice was to call the new Logistic row **"Courier"** - and it would have sat directly above **Parcel couriers**, which means something else entirely. A parcel is something you **post**; a van is a trip you **book**. v309 already gave those two their own words on the order cards, so the screen is called **Send a van** and the card inside still names **Lalamove**, asked from the registry rather than typed here.

AND THE THIRD CARD I DID NOT MOVE - because the name was the real problem, not the address. "Mailing labels (courier)" holds **one address**, and I went and checked what actually reads it. **Four things do:**

- the **FROM block on a parcel label** the **door the courier collects from the letterhead on every invoice** the **from line on the wish-list email**

So it is not a courier setting at all - it is one address typed once that four screens read. Filed under **Logistic**, an invoice with the wrong heading would have sent you looking under deliveries. You chose **"Rename it, keep it in Settings"**, and that is what it is now: **Your address**, with the four jobs named in the card itself, and a way to it from the **Send a van** screen so the trip is still one tap when it is a delivery you are fixing.

What Settings keeps: Delivery settings, App password, **Storefront**, Your address, Website & developer, Live availability, Shared data, Backup & safety, Danger zone. **Storefront stays there deliberately** - your call, and recorded so no later version proposes moving it again.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,784 tests, all green** - and the v312 route guard caught **both** new screens by name when I briefly took them out of the menu.

05 Oct 2026 - engine v314, BRING-A-FRIEND MOVES TO THE SHOP, AND THE REWARD IS A COUPON (no database step, nothing to upload - pushing this one is the whole of it).

Why. Two of your own, in one message. First: *"can be brought to The Shop, rather than in Settings."* Second: *"can we make to more seamless with other promo?"* - and when I asked what seamless meant, you chose **"One place, read as a family."**

What you see now. More -> The shop carries three rows instead of two: **Promo codes**, **Bring a friend**, **Reviews**. The scheme itself has not changed at all - same switch, same three numbers, same links on a customer's card, same Give-coupon press on an order. It simply lives where you would look for a customer offer,

instead of beside "defaults, backup, transfer".

And each half now names the other. The Bring a friend screen carries a card called **"A link, or a code"**, which says the difference in the terms you gave me: a link **costs nothing to issue and travels** - forward it and it still works; a code and a label is for a partner who prints brochures. Promo codes says the same thing from its side and points back.

BUT THE IMPORTANT PART OF THIS VERSION IS YOUR SECOND SENTENCE. You said: *"if we state only credit of ringgit, there might be confusion of how much credit to apply, but we can state, only one coupon apply for each purchase."* **You were right, and it was worse than you described - it was two places, and both of them argued with themselves:**

- On a customer's card holding **two** RM3 coupons, the chip above read **"2 ready"** while the line beneath it read **"unused = you still owe RM 3.00 off an order"** - because that figure was taken from the **first** coupon only. Two numbers, disagreeing, on one line. On an order with two ready, the line read **"RM 3.00 credit available on this order"** beside a button reading **"Apply credit (2)"** - and the press spent exactly **one**. So the screen showed one coupon's money next to a count of two, and did something else again.

The fix is not arithmetic. It is your rule, written down. No screen shows a running balance any more. The customer's card now reads **"2 ready - one per order"**; the order reads **"RM 3.00 coupon ready for this order - one per order, 2 more after it"**; and the message a customer forwards to their friend now promises *"a RM 3.00 coupon for a future order - one coupon per order"* so nobody has to guess. The rule is stated in all four places you would meet it.

AND WE NOW CALL IT A COUPON, IN YOUR WORDS. Eighteen strings changed - Give **coupon**, Apply **coupon**, **Coupons**, Add **coupon**, the customer's card, and all three WhatsApp messages (English, Mandarin, Bahasa Malaysia). **Your decision on the two schemes is unchanged and I have kept to it: both stay, and the difference between them is what you hand over. AND ONE THING DELIBERATELY DID NOT MOVE WITH THE WORD: the stored name.** Every coupon already given out lives under state.credits, and renaming that would have orphaned all of them - so the words changed and the storage did not. It is a test.

ALSO NOT RENAMED, because they are a different thing entirely: "Insufficient Credit" and the EasyParcel wallet line on the parcel block. That is the parcel money, not your reward.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,783 tests, all green** - and one of the new ones is the v312 route guard, which **caught this new screen by name** the moment I briefly took it out of the menu, exactly as it was built to.

05 Oct 2026 - engine v313, THE SMALL GREY LINES ARE READABLE NOW (no database step, nothing to upload - pushing this one is the whole of it).

Why. Every small second line in the app is drawn in one grey - the line under a menu row, under an order row, the hint under a box, a collection Point's address in the shop. You looked at it against the darker candidate and said do it.

IT WAS UNDER THE FLOOR, AND YOUR OWN APP HAD ALREADY ADMITTED IT ONCE. Measured against the card behind it the old grey came to **3.50:1**. The standard asks **4.50:1** of text. And in v309, when I drew the **Post a parcel / Send a van** headings, I deliberately dodged this exact grey for the line beside them and wrote the reason into the code: *"at 12px it is only about 3.5:1 - under the 4.5:1 floor that applies to text."* So the app knew. It had only ever fixed that one spot.

AND THIS IS THE PART WORTH READING: THE COLOUR IS SET BY THE WORST PLACE IT LANDS, NOT THE BEST. The value you approved on the page measured a comfortable **5.02:1** on a plain card - and only **4.24:1** on the soft grey behind a used promo chip or a "courier paid" tag. Five more backgrounds - the striped band, the warm block, the green, amber and red chips - came in between **4.27 and 4.47**, all under the floor. So the colour is a step darker than the one on the page: **#7b6552**, which is **4.56:1 on the worst background there is** and **5.40:1 on a card**. One colour has to clear the worst thing it is drawn on, or you have only moved the problem to the chips.

What changed, exactly: one value, in two files. The backoffice and the shop each keep their own copy of the grey, and both moved together - a shop whose quiet line is a different grey from the backoffice's is the same fault as two rows that look alike and behave differently.

AND ONE THING DELIBERATELY DID NOT MOVE. The 2px rule inside the **Post a parcel / Send a van** headings is an *edge*, not text, and an edge answers to 3:1 rather than 4.50:1. Left at the new darker text colour it would have stopped being a join inside a card and started reading as another box. It keeps the old, lighter value under its own name - **measured on the real card: still exactly rgb(154, 132, 113), unchanged, 2px.**

And the calendar's quiet day stayed put too. That colour was built as a halfway mix between the old grey and the ink, precisely because the old grey was faint. The grey is not faint any more, so re-deriving it would have pushed a quiet day *toward* a delivery day and narrowed the very gap it exists to keep - for no gain. It stays at 6.71:1 against a delivery day's 13.45:1.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,772 tests, all green**, including a new guard that measures this grey against **every background it can land on** and fails by name if it drops under the floor - the test that would have caught 5.02-on-a-card.

05 Oct 2026 - engine v312, THE MORE SCREEN GROUPED BY THE WORK (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your words: *"when i work on Products, i have to alway go into Others to find, ingredient, unit, category"* - and *"Everything about delivery should be group under logistic."*

You were right on both counts. Seventeen screens sat in ONE list under a single heading called "Manage", in roughly the order they were built. Nothing was wrong with any row; the problem was that its neighbours told you nothing. Categories, Units and Ingredients were at rows 13, 11 and 14 - two of them separated by other screens - and the four delivery screens were at rows 7 to 10, mixed in among the money screens.

What you see now. Seven headings, in the order you reach for them:

- **Logistic** - Delivery run · Delivery dates · Self collection Points · Parcel couriers
- **Products & ingredients** - Categories · Units · Ingredients
- **Buying** - Purchase Order · PO history · Suppliers
- **Money** - Money · Profit
- **The kitchen** - Production line · Scenario planner
- **The shop** - Promo codes · Reviews
- **Settings & this app** - Settings · Change history · Developer contact

Logistic is your word, and it carries all four ways an order leaves the kitchen. **Products & ingredients is the one your sentence asked for**: the three things you reach for while working on a product, together, under the name of the screen they belong to.

And the trip out of the screen is gone. The Products screen now carries the same three at its top, under the same heading - Categories, Units and Ingredients, one tap from where you are editing. It sits BELOW the + New product fold, so the thing you came there to do is still the first thing on the screen.

"About" is gone as its own section, because Settings and the change history are the same subject - the app itself - and two headings for one subject was one too many.

Nothing moved that has to be found twice. All seventeen addresses are exactly what they were, so an old bookmark still lands where it always did, and every row keeps the words it had underneath it; only the heading above it changed. **Two of those second lines now say the moment rather than the mechanism** - "Which days you deliver, and who is on each" instead of "Set and manage delivery dates", and PO history reads "Your saved shopping lists".

AND THE ONE THING WORTH KNOWING: A ROW DROPPED IN A REGROUP IS SILENT. The screen still exists and its address still works - it has simply become unreachable from the menu, and nothing goes red. So

the new test reads the app's own route table and walks it against this menu, and fails by name if a screen is missing. It was watched going red before it was trusted, and it caught a real fault on the way in: the first cut passed the rows as raw lists and drew four empty headings with nothing under them.

Your data is untouched. No SQL, no order, product, price or bake day touched, and **no database step**. The suite is **2,769 tests, all green**.

Measured at 375 wide, same phone, same data: 1,994 px became 2,263 px. Grouping costs height - about **+269 px**, a third of a screen of extra scrolling. Finding beats scrolling here, but it is a real trade and it is worth knowing about.

05 Oct 2026 - engine v311, THE ADDRESS BOX SAYS WHY IT ISN'T SUGGESTING (ONE STEP FOR YOU, BELOW - the app side is a push, the reason needs a redeploy).

Why. Your words: *"why sometimes in add order or edit order, the address is not auto complete, and the map dont show?"* - and then, precisely: *"the suggestion list never appear"*. **It was true, and the app never said why. It knew, and stayed quiet.**

WHAT WAS HAPPENING. The suggestions are a Google feature that has to be switched on, and the server deliberately answered a switched-off one with **a plain "no" and no error** - its own words: *"this reaches her as a box that simply does not suggest, which is exactly what it did before this version existed."* That reads well until the feature *does* exist and **never works**: a box that never suggests and never explains itself is indistinguishable from a broken box.

What you see now. When the suggestions cannot work, the box says so, in the place the list would have been:

Address suggestions are switched off - Google refused the request. The Google key needs Places API (New) enabled and allowed.

AND IT SAYS IT ONCE, THEN STOPS ASKING. A line under a field she is typing in must not flicker, and a phone must not spend a request per keystroke to be told the same thing. And the box now **marks the difference between the two kinds of failure**: something you must fix (not signed in; the Google key not allowed to use Places) is **said out loud**; a signal dropping out is **not**, because that one passes on its own and the suggestions come back by themselves.

THE ONE THING THIS NEEDS FROM YOU: the *reason* travels in the server's answer, so the courier function has to be uploaded again. The exact line is in the commit notes below. **Until it is uploaded the box stays exactly as it is today.** Nothing else changes, and nothing is broken by uploading it.

AND THE SCREENSHOT FOUND A SECOND FAULT THAT HAD ALREADY SHIPPED - a real one, mine, and this one is worth reading. Looking at the real card, the EasyParcel block was printing the **literal word "null"** under the weight box and **three more beside its buttons**. The block builds its own contents as a list with optional parts in it, and **replaceChildren(null) does not skip a null - it inserts a text node reading "null". Every one of the 2,758 tests passed while it was on your screen**, because the test shims quietly filter nulls for you. **A forgiving stub hid a real fault** - the same lesson as v226, word for word. It is fixed, and the real-browser check that would have caught it is now in place and was watched going red.

Your data is untouched. No SQL, no order, product, price or bake day touched. The suite is **2,758 tests, all green**, and the real card now runs **41 checks** on the Orders screen and **9** on the address box.

04 Oct 2026 - engine v310, THE SAME DIVIDE ON EVERY CARD (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your words, right behind v309: *"the drawing of line between parcel and lalamove should be consistent over the app"*. You were only half-served - v309 put the two headings on the **Edit** card and nowhere else.

THREE CARDS CARRY BOTH KINDS, SO ALL THREE NOW SAY THE SAME TWO THINGS. The **Edit** card, the **Note / tracking** card, and the **+ New order** card each draw **Post a parcel** over the parcel fields and **Send a van** over the van fields - **from one definition**, so the words cannot drift apart. **Two headings on each is not the check; the same two is**, and that is exactly what the new test compares.

AND THE WAY I GOT IT WRONG FIRST IS WORTH KNOWING, because it is the kind of thing that would have looked fine on paper. I tried putting the heading *inside* the parcel block, so that no card could ever forget it. **That put it BELOW the consignment number on the Edit card** - the number is a field the card itself draws just above that block - so the heading landed in the middle of its own group and left the number stranded above it. The test caught it. **The words can be shared; the position cannot**, because only the card knows where its own fields are.

Your data is untouched, and there is nothing to run. No SQL, no upload, no key. No order, product, price or bake day is touched, and no field moved on any card. The suite is **2,756 tests, all green**, and the real Orders screen now runs **40 checks** - including that each heading sits above its own fields in document order, that neither one says "kind 1", and that **all three cards show the same two.**

04 Oct 2026 - engine v309, THE TWO WAYS AN ORDER LEAVES ARE SAID APART (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your words: *"what i see is easy parcel is kind 1, lalamove is kind 2, but the interface din draw a clean border between them, so it is quite confusing for user"*. **The first half of that sentence is the proof.** You had the two the wrong way round - which you could only do because the card never told you which was which. Four labels sat next to each other - **Parcel carrier**, **EasyParcel**, **Courier charge**, **Get a delivery price** - with nothing saying what belonged to what.

What changed. The card now draws a **rule with a name on it**, twice:

Post a parcel *Nationwide, a few days* - the consignment number, the parcel carrier, and EasyParcel

Send a van *Today, inside Penang* - the courier charge, who paid it, and **Get a delivery price**

AND THE NAME SAYS WHAT IT DOES, NEVER "KIND 1" OR "KIND 2". Those are this app's own words for the two kinds of courier and they mean nothing on your screen. What you need to know is that one of them is a van today and the other is a parcel over a few days - so that is what it says. **Nothing moved:** every field is exactly where it was, and each heading went above the fields it already sat on, so a heading can never end up announcing the other one's controls.

TWO THINGS I MEASURED RATHER THAN GUESSED, and both would have failed quietly. My first rule used the card's own line colour - **1.25:1 against the surface**, a faint tint that disappears on a phone in Penang daylight, which is exactly where you read it. **A boundary you cannot see would not have answered you.** It is now **3.5:1**, which clears the 3:1 the standard asks of an interface edge. And the small grey beside the name - *"Nationwide, a few days"* - was **3.5:1**, **under the 4.5:1 that applies to any text**; it is now **5.69:1**. Both were measured in the browser by reading the rendered pixels, not worked out on paper.

Your data is untouched, and there is nothing to run. No SQL, no upload, no key. No order, product, price or bake day is touched, and no field moved on any card. The suite is **2,756 tests, all green**, and the real Edit card now runs **36 checks** - including that each heading sits above its own fields in document order, and that neither one says "kind 1".

04 Oct 2026 - engine v308, THE EASYPARCEL BLOCK READS AS OFF, NOT AS BROKEN (no database step, nothing to upload - pushing this one is the whole of it).

Why. You looked at EasyParcel's top-up offer and found the thing that decides this: **only an account topped up by RM500 enjoys integration.** For a bakery that posts the occasional parcel that is the wrong money, so **you decided not to sign up** - and you were right. This version is the consequence: the EasyParcel block now says so plainly instead of sitting there looking broken.

What changed. The block asks the server once, when it opens, whether EasyParcel is set up. If it is not, you now get one short paragraph rather than a set of controls that can only ever fail:

EasyParcel - Not set up yet - and nothing here is needed to post a parcel by hand. Record the carrier above and type the consignment number, exactly as before. EasyParcel is not set up yet - its key has not been saved on the server.

AND IT OFFERS NO PRESS AT ALL IN THAT STATE. No weight box, no "Check the price", no "Book it". A

button whose only possible answer is the same sentence every time is the shape of a dead control, and this app treats a dead control as a bug - so when it is off, it is off.

The important half is the reassurance, not the apology: nothing is missing. Posting a parcel by hand is how this app has always posted one - record the carrier, type the consignment number - and that is completely unchanged. **The API was only ever going to save you the retyping.**

AND A REAL FAULT THIS FOUND, WHICH IS THE PART WORTH KEEPING. The first attempt had the block call the pop-up's own repaint when its answer came back. **That repaint arrived at an unpredictable moment while she was working further down the same card - and stranded the door block's own pending look-up.** A test caught it (a pin the card had just found was thrown away). **Every other block in this app paints ITSELF and leaves the card alone** - paintCourier, paintParcel, paintPoint - and this block now does the same. One repaint of the whole card, on the booking, because that one changes the tracking box above it.

Your data is untouched, and there is nothing to run. No SQL, no upload, no key, no account. No order, product, price or bake day is touched. The suite is **2,756 tests, all green**, and the real Edit card now runs **29 checks** - including five that can only be answered with the key absent, which is how it will be for you.

04 Oct 2026 - engine v307, THE EASYPARCEL SEAM (no database step in the app - but ONE one-time step for you, below).

Why. You asked to compare Ninja Van and SPX, and the honest answer turned out to be that **you do not integrate with a carrier - you integrate with an aggregator.** Ninja Van's own API is not self-serve (you must have an account, ask their sales team by email for even a sandbox, and production keys come only after an "integration audit" with sample orders). SPX has no public direct API at all for someone who is not a Shopee seller. **EasyParcel carries both** - along with J&T, Pos Laju and DHL - and its API **is** self-serve, with a free demo environment. That is the same choice you made back on 28 September: *"record first, the booking API after / EasyParcel behind that API"*.

WHAT YOU GET: ONE PRESS, EVERY CARRIER, PRICED FOR THIS PARCEL. On an order's **Edit** card there is now an **EasyParcel** block. Type the parcel's weight, press **Check the price**, and every carrier they use comes back with a price for **that parcel, to that postcode, at that weight** - cheapest first, with the delivery time and whether they collect or you drop off:

J&T Express - RM6.20 · J&T Standard · 3-5 working day(s) · drop off · they collect · cheapest **SPX - RM7.10** · SPX Express · 3-4 working day(s) · drop off · they collect **Ninjavan - RM9.40** · Ninja Van Standard · 2-3 working day(s) · they collect

That is the comparison you asked for, done per parcel instead of guessed at from a blog post. The cheapest is marked but **never chosen for you** - which carrier to use is your call, and the reason to show several is that the cheapest is not always the one you want.

AND YOUR WALLET IS WARNED BEFORE, NOT AFTER. EasyParcel is **prepaid**: a booking with too little credit fails **with the parcel already packed**. So the balance is shown beside the price, and when it is short it says so plainly - *"Your EasyParcel balance is RM4.20 and this parcel costs RM6.20 - top it up before booking, or the booking will fail with the parcel already packed."* When a booking does come back refused, **their own words are shown** - "Insufficient Credit" - rather than a generic failure.

BOOKING SAYS THE PRICE AND THEN PAYS IT. Pressing Book asks first, naming the courier and the amount - *"Book this parcel with SPX for RM7.10?"* - and only then books and pays in one call. **The consignment number is SAVED AT ONCE**, not left waiting for a Save that may never come: the same rule your booked trips already follow, because a parcel that has been paid for must not be discardable by closing a card. It lands in the tracking box, where every screen already reads it.

WHAT THIS DOES NOT DO, said plainly rather than discovered.

- **It does not write the consignment number onto the customer's card by itself** - it goes in the tracking box, exactly as a parcel you booked by hand does. **It does not pick your carrier record for you.** EasyParcel's name for a service is not your own list's entry, so booking fills in the number and leaves the carrier box alone -

one press above if you want it named. **It is not required.** Posting a parcel by hand still works exactly as it did - record the carrier, type the number. This is offered beside that, never instead of it. **Dry and sealed only, unchanged.** Fresh focaccia and anything frozen are still not parcels.

ONE THING ONLY YOU CAN DO, and it is a signup, not a key in chat. Sign up at easyparcel.my, complete the account verification, and register the key for API access - their three steps. Then the key is stored **on the server** (never in the app, never in a browser) with two Terminal commands, which are written at the top of the new function. Until that is done the block says so in words and nothing else changes. **Set it to DEMO first** - the demo host books a parcel nobody collects, which is how you can try the whole thing before a sen is spent.

Your data is untouched, and there is nothing to run in the app. No SQL, no upload. No order, product, price or bake day is touched, and every screen behaves exactly as it did. The suite is **2,756 tests, all green**, including **thirty-six** new ones - the wire format against EasyParcel's own 55-page document, the address reading, and the seam end to end on the real card.

04 Oct 2026 - engine v306, THE SMALLEST BASKET A POINT WILL TAKE (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your ask: *"the per-point minimum order, this should be switchable"*. It is exactly that - **per Point, and switchable**, on the Point's own card.

THE SWITCH, AND WHERE IT COMES FROM. On a Point you now get **Minimum order** with two choices - **No minimum** and **Only on a basket of at least** - and the second one opens a **Smallest basket (RM)** box. That is the SAME switch the Promo codes screen has used since v269, because a minimum is a minimum and learning a second shape for one idea is how two screens come to mean two different things by one word. **You picked ringgit**, out of the two units offered, so "a basket of RM30" means the same thing everywhere in the app.

EVERY POINT STARTS WITH NO MINIMUM, which is where they already are and what you asked for when this whole feature was being discussed - *"keep it as simple as possible, say no minimum for self collect order"*. The card says so plainly: **No minimum order - one loaf still goes.**

On your card. Each Point's row now reads **Minimum order RM30.00**, or the sentence above.

In your shop. A Point whose smallest basket the customer has not reached is **PARKED, NOT HIDDEN** - it stays on the page with the reason in its own line:

Farlim, Air Itam Needs a basket of RM30.00 or more - yours is RM15.00 so far

Tapping it says the same sentence rather than silently doing nothing, and **the moment their basket reaches RM30 the Point opens up - on the same repaint, with no reload.** Add a loaf and watch it; take one back and it parks again, **and a Point they had already chosen falls back to your kitchen**, because the shop cannot post an order to a Point whose basket is not met.

THIS IS YOUR RULE, SO THE SHOP HONOURS IT - and that is not the same as a gate. The thing your standing instruction forbids is a rule **the app invents** (a closed day, a sold-out line) standing between you and a sale. A smallest basket is a rule **you typed on your own Point**, so enforcing it is the shop doing what you asked. **Nothing on your own side is blocked:** an order you take over the phone for one loaf at Farlim is yours to take, and the app will not argue.

AND ONLY WHAT THE SHOP NEEDS LEAVES YOUR APP. The Point's smallest basket is now published along with its id and name - **the receiver, their phone, the fee and the address still never do**. A smallest basket is the opposite of private: it is exactly what the customer has to know *before* choosing, and without it the page could only take an order the Point does not want.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. Every Point you already have carries no minimum, so **every one of them behaves exactly as it did yesterday** - proved rather than asserted: the shop's own v299 checks still pass, 20 of 20, with Points published that have no minimum at all. The suite is **2,720 tests, all green**, including eight new ones, and every one was watched going red with the fault put back.

04 Oct 2026 - engine v305, THE HOURS ON THE RUN ROW (no database step, nothing to upload -

pushing this one is the whole of it).

Why. Your ask, and it is the one number the Delivery run was missing. The collection hours you set on a Point decide **when the bread has to BE THERE and handed over** - so a trip booked for the wrong part of the day should be visible on the screen where you spend the money on a van, not discovered a day later.

What you see. A Point's row on the run now reads:

Farlim, Air Itam - 2 orders collecting here · Lebuhraya Thean Teik · **collect 2-6 pm** · Focaccia x2 · Focaccia x1 · Focaccia x3

Where, then when, then what - the address, then the hours, then the bread.

And it says nothing when you have not set any. The card already tells you a Point has no collection window; repeating it on every run row would be noise on the screen you read while working. A customer's own doorstep never claims hours either - a doorstep is not a place with opening times.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. A Point with no hours behaves exactly as it did yesterday, a doorstep's row is unchanged, and the load line is untouched. The suite is **2,712 tests, all green**, including three new ones, and every one was watched going red with the fault put back. **One of them taught something worth keeping:** the check that a doorstep never shows hours stayed GREEN under the first fault I tried - because reading a null Point already gives nothing - so the fault was not the one I had chosen. Re-pointed at the mistake that would really do it, *taking the hours from whatever Point happens to be first rather than from this row's*, it turns red by its own name.

04 Oct 2026 - engine v304, THE COLLECTION WINDOW (no database step, nothing to upload - pushing this one is the whole of it).

Why. This is the piece we left open, and you chose it from three: **the window belongs to the PLACE.** You type it once on the Point - *"Farlim: collect 2-6 pm"* - and every order collecting there is promised it. Not the van's arrival window, which is a different thing entirely.

What you see, on More -> Self collection Points. Every Point's card gains **"Customers can collect from"** and **"and until"** - two time boxes, the same pair the Delivery run fills in for the van, so a window means one thing in this app and is read by one piece of code. And the row itself now says which of the two it is: **Collect 2-6 pm**, or **No collection window - customers are told the day only.**

What your customer is told. The confirmation, the payment reminder and the pickup reminder all name the place **and the hours**: *"Self collect at Farlim, Air Itam, collect 2-6 pm."* Leave both boxes empty and they are told the day and nothing else, which is a promise you can keep rather than one that reads as open all day.

THE ONE THING TO GET RIGHT, and it is your judgement rather than the app's: SET THE HOURS FROM WHEN THE BREAD IS THERE, NOT FROM WHEN THE SHOP OPENS. The van arrives during the round, so a window starting at opening time can have a customer standing at the counter before their order has been delivered. The app deliberately does **not** work this out for you - the same way it never works out the fee - because a time the app derives is a time it can get wrong.

AND THE VAN'S OWN WINDOW IS NEVER QUOTED TO THEM. When you book a round, the trip's window is stamped on every order it carries - and for a collection that is when the **bread reaches the Point**, which is your business, not the customer's. Telling them both would be telling them two different times in one message. So a collecting customer gets the place's hours or nothing at all, and never the van's.

A hole that fixing this found, and it was mine. The **pickup reminder** said only *"Packed and ready for pickup on ..."* and named **no place at all** - so a customer collecting at Farlim was told their order was ready and never where to go, while the confirmation, the payment reminder and the shipped message all named it. v299 claimed all four later messages named the Point; three of them did. It now reads **"Packed and ready to collect from Farlim, Air Itam on ..."**, and a collection from your own kitchen keeps the words it has always had.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. A Point with no hours behaves exactly as it did yesterday, and every order already placed is unaffected. The suite is **2,709 tests, all green**, including eighteen new ones - the model, the one place that decides whose window a

customer is told, the card's own time boxes, and the three messages a collecting customer actually receives, driven through the real builders rather than through the helper behind them. **Every one was watched going red** with the fault put back.

Also in this version, and it is housekeeping you will never see: the window stopped being the courier's own and became a small module both the courier and the Points card read, so **a window cannot mean one thing on one screen and another somewhere else**, and its own tests moved with it.

04 Oct 2026 - engine v303, A POINT YOU SET UP AND HAND OUT BY HAND (no database step, nothing to upload - pushing this one is the whole of it).

Why. Two things you found while watching v301 and v302 land, and they are the same gap: everything about Points worked for a **customer** ordering from your shop, and almost nothing worked for an order **you** take.

1. + New order and Edit now let you say where it collects from.

Your words: *"how about + new order, and add edit order?"* You were right, and it is the important half - you take a great many orders over the phone and in chats, and until now an order you keyed in yourself **could not be a Point order at all.** It could not go on the Delivery run, the Point's fee was never counted for it, and the customer was never told where to collect.

- **A "Collect from" picker**, under **Fulfillment** - **My kitchen** first, which is the default and is what every order you have ever taken already means, then your open Points in your order. **It appears only when there is something to choose.** With no Point open, or on a Courier order, it is absent rather than sitting there doing nothing. **The Point's name is frozen onto the order** the moment you take it, so renaming or ending a Point later leaves every order that already went there still saying where it went. The same rule that keeps a sold price on an order. **Choosing your kitchen clears it.** An order collecting from your kitchen carries no Point at all, which is exactly what it has always meant.

2. A Point's address box now asks Google as you type.

Your words: *"there is no address auto complete for collection point?"* - and then the better question, *"why not make the point consistent with the customer card?"* Both fair. The suggestion box was written for the order's delivery address and had never been given to this one, so the address a **driver** is sent to - and the address the pin is looked up from - was the one address you had to peck out in full on a phone.

- **The same behaviour as the order's address box, because it is now literally the same code**, moved into one place so the two boxes cannot drift apart. **Nothing is ever blocked by it:** a lookup that fails shows nothing at all, and whatever you typed is what gets saved.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. An order taken before today carries no Point, which is what "collect from my kitchen" has always meant, so nothing needs migrating and no order is rewritten. The suite is **2,691 tests, all green**, including twelve new ones - nine driving both order cards, and three driving the Point's own address box through the real suggestion channel. **Every one of them was watched going red** with the fault put back before it was called done.

04 Oct 2026 - engine v302, AN ORDER COLLECTED AT A POINT REACHES THE RUN (no database step, nothing to upload - pushing this one is the whole of it).

Why, and it is a fault I found while starting the collection window you asked for. The run was not carrying Points at all. Not rarely - **never**, and it said so quietly: a day where customers were collecting at Farlim read **"Nothing to run yet"**.

WHAT WAS WRONG. Choosing a Self collection Point in your shop **never changes the Self collect / Courier choice** - that is by design, a Point IS a collection. So the order is stored as a *collection*, and the run screen used that word to decide what needs a vehicle and **skipped every collection**. Every Point order was thrown away **before** the row work v301 added could ever see one. v301 built the right rows and they were unreachable.

The rule now, and it is the one that was always meant: a courier order needs a vehicle, and a collection AT A POINT needs one too - the bread still has to reach the Point. A collection from **your own kitchen** still needs none, because you hand those over yourself. That half is the one a careless fix breaks, and it has its own

test.

What you see now. A day of collections at Farlim **offers the run**, the day's own line reads **"1 stop"** rather than "2 courier orders", and the Delivery dates screen's **Run (N)** button appears on that day counting the same stops. **A Point is ONE stop however many customers collect there**, so two customers at Farlim still read **1 stop - 6 items**.

AND ONE RULE NOW LIVES IN ONE PLACE, which is why it went wrong. The run screen and the Delivery dates screen's **Run (N)** button both have to answer *"does this order need a van?"*, and they each answered it themselves. **Two readings of one rule is how they came apart**, so the rule and the stop's own name are now asked of one function in courier_job.js and read from there.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. An ordinary courier day behaves exactly as it always has, and a collection from your kitchen behaves exactly as it always has - **a run with no Points on it is unchanged**. The suite is **2,679 tests, all green**, including eleven new ones: six on the rule itself, three driving the run screen with a **real** Point order, and two on the Run badge. **Every one of them was watched going red** with the fault put back - the gate, the kitchen half, the shared stop name and the badge - before it was called done. One of them is worth naming: the test fixture itself had been modelling a Point order as a courier order, **which the shop has never produced**, and that is the reason nothing caught this. **A fixture that cannot happen is not a test.**

04 Oct 2026 - engine v301, THE DELIVERY RUN CARRIES POINTS (no database step, nothing to upload - pushing this one is the whole of it).

Why. This is the piece you asked for: the Delivery run carrying Self collection Points.

A POINT IS ONE STOP, AND THAT IS THE WHOLE VERSION. A courier charges **a base fare plus a fee for every extra stop**, so the run screen's oldest rule is *one stop per customer* - a customer who bought three things is one doorstep, not three. A **Self collection Point is the next version of the same idea**: four customers collecting at Farlim are **one place the van goes**. A run built one-stop-per-customer would send the same driver back to the same shop and **bill you a stop fee each time**.

What you see on the run now.

- **A Point gets a row of its own**, listing **how many orders are collecting there** - *"2 orders collecting here"* - and its own address. The customer's name is deliberately **not** on it: the bread is going to Farlim and the customer is meeting it there. **The load line counts STOPS**, so two customers at one Point read **"1 stop · 6 items"**, not two stops. Every line of bread is still counted. **A Point you have not pinned offers "Pin the Point"** right there, because the run cannot price a trip to a place with no coordinates. **A Point and a doorstep on the same run are two stops** - the mixed run you described.

Three things it is careful about, and each one is money.

- **THE VAN GOES TO THE POINT, NEVER TO THE CUSTOMER'S HOUSE.** A customer who chose to collect at Farlim is **not at Farlim** - and she may still have a doorstep of her own pinned from an earlier delivery. That door is **ignored** for a collection order. Sending a driver to her house with four other people's bread would be the most expensive way to be wrong on this screen. **ONE DROP ON THE WIRE, and this is the one that shows up on your bill.** The price request is read back in the test on the **bytes that would leave your phone** - because a trip built one-drop-per-customer is priced for a journey you are not taking. Two customers at one Point: **one drop. A POINT YOU PAUSE STILL SENDS THE VAN.** Pausing decides what is **offered**, never what an order already promised - the customer was already told to go to Farlim, and their bread still has to get there. And a Point you have **deleted** has no pin to give, so that order falls back to the customer's own door rather than a van sent to coordinates nobody has any more.

And the double-booking guard still sees everyone. A Point is treated as already on a trip if **any** of its customers is - so a Point can never be quietly swept onto a second van, which is the fault you reported in v242, at a place instead of a door.

AND A POINT YOU ADD NOW ACTUALLY REACHES YOUR SHOP. You reported this while the run was

being built - *"the point added still not able to appear on store?"* - and **you were right, and the fault was mine.** Your Points travel to the shop inside the **storefront row** the shop reads, and **the Points card was not republishing that row.** Every other screen that changes something the shop shows - Promo codes, Products, Categories, Settings - ends its save with a republish; the Points card only ever saved to your phone. So a Point you added **stayed on your phone and never went out**, and your shop went on offering only the Points it had last been told about. The test proves it on the **request itself**: adding a Point **publishes**, pausing one **republishes**, and **pinning one does not** - because the pin is deliberately never published, so republishing for a drag would be a request that changes nothing.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. **A run with no Points on it behaves exactly as it always has** - every order that is not collecting at a Point keeps a row, and a stop, to itself. The suite is **2,668 tests, all green** - including two new ones that fail if a Point change ever stops reaching your shop. Proved live on the real Delivery run at a phone width, and on the wire: **two customers collecting at one Point price ONE drop, at the Point's own address**, and the mixed run prices the Point and the doorstep as two. The Points card was driven at a phone width too - **40 checks, all green** - including the three that read the publish request itself, and every one of them was **watched going red** with the fault put back before it was called done.

04 Oct 2026 - engine v300, EVERY POINT GETS A PIN ON THE MAP (no database step, nothing to upload - pushing this one is the whole of it).

Why. You asked for the Delivery run to carry Points, and said a Point should **"get like what customer is getting location pin"**. This is that pin - and it is the piece the run cannot work without, which is why it comes first.

A courier is not given an address. It is given a point. Lalamove wants **"5.41405,100.31408"**, and "Farlim, Air Itam" is a guess about one - the courier itself says so when it cannot place an address. So **a Point with no pin is a name you can read and a place a van cannot be sent to**, and the run cannot price a trip until every end is a real point.

What a Point has now, on More -> Self collection Points.

- **A line saying where it is** - **Farlim, Air Itam · 5.41405, 100.31408**. **Both the name and the two numbers**, because a name alone cannot be checked and a pin in the wrong place is only ever noticed by looking at the numbers. **"Put the pin on the map"** under every Point, opening **the same map a customer's doorstep is placed with**. It is the same act, so there is nothing new to learn - and your Point's own address is handed to the lookup, so if you have already typed where it is you may not have to drag anything. **An unpinned Point says so plainly**: *"Not pinned yet - a van cannot be sent to a name alone."* It never pretends to have a door. Once it is pinned, the press becomes **"Move the pin"**.

Two things this version is careful about.

- **Correcting a Point never un-pins it.** The pin is not a field of the form - it is placed on a map - so an edit had to be made to carry it across deliberately. Without that, fixing a spelling would have silently taken the door away, and you would have found out when a driver was sent nowhere. There is a test that edits a pinned Point and checks the pin is exactly where it was. **The pin never leaves your app.** The shop is still told only the Point's **name** - a customer chooses a Point by name, and the driver is the only one who needs the door.

And a half-typed pin is not a pin. Anything malformed - one number and no other, a latitude of 91 - reads as **unpinned** rather than as a point in the sea. A row half-synced between your phones can never send a driver somewhere absurd.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. Points you have already defined simply start unpinned. The suite is **2,656 tests, all green**. Proved live on the real card at a phone width: **37 checks**, including that an unpinned Point says so, that the press opens the map picker, that a pinned Point shows its numbers, and that correcting a Point leaves the pin where it was.

Next, and it is the thing you actually asked for: the **Delivery run** carrying Points. It is its own version because the run screen is the one that spends your money, and it needs changing carefully - a Point's orders have to

become **one stop** on the trip rather than one stop each, and the screen that ticks them has to say so.

04 Oct 2026 - engine v299, THE CUSTOMER CAN COLLECT FROM A POINT (no database step, nothing to upload - pushing this one is the whole of it).

Why. v298 gave you the card that defines your Self collection Points. This is the other half: **your customers can now choose one**, your order carries which one, and **the customer is told where to go**.

What your customers see, in the shop.

- **"Collect from"**, under the Self collect / Courier choice. **Your kitchen is first** - it is not a Point and needs no record - and **your open Points follow, in your order. The whole thing only appears once you have a Point open.** With none, the shop is **byte-for-byte the shop it was**, which is how you said you would do it: one at a time. **The chosen row takes the brand's orange edge**, the same "this one is on" the buttons already use. **A Point paused or deleted while a customer has it chosen falls back to your kitchen.** The shop cannot post an order to a place you have stopped offering - and the kitchen is always there.

What your customers are told. The confirmation, and all four later messages, say **"Self collect at Farlim, Air Itam"** and give the **address of the place**. That wording is written **once** and read by both message builders, so a customer cannot be told one thing in the confirmation and another in the reminder.

Three things this version is careful about.

- **Only the Point's NAME ever leaves your app.** The receiver's name, their phone, the fee and even the address **stay in your app**. The shop is a public page with no login, and publishing the receiver's number would put a private person's phone on a page anyone can read. The address is withheld for a plainer reason: the message that tells a customer where to go is built from **your** copy. **The name on an order comes from YOUR record, never from the page.** The shop is public, so a name it sent could be anything at all; the order is matched against **your own Points** and an id you do not have falls back to the kitchen. This is the same rule the promo codes follow. **A Point you PAUSE is still honoured for orders already placed.** Pausing decides what is **offered**, never what an order already promised - the customer was already told where to go. Exactly the rule that keeps an ended promo code coming off the order it was placed on.

And your own screens say it too. The order row and the packing label read **"Self collect · Farlim, Air Itam"**, so whoever is packing a bag can see where it is going. A kitchen collection reads exactly as it always has.

Your data is untouched, and there is nothing to run. No database step, no upload, no key - the Points travel in the storefront row your shop already reads. **Every order already placed is unaffected**: an order with no Point IS a collection from your kitchen, which is what it has always meant, so nothing needs migrating. The suite is 2,650 tests, all green. Proved live in the real shop at a phone width: **20 checks**, including that the kitchen is first and carries no id, that a paused Point cannot stay chosen, that with no Points open the shop is unchanged, and that a long Point name does not push the page sideways.

04 Oct 2026 - a correction to this entry. This entry says your Points travel in the storefront row your shop reads, and they do - **but only once your phone has actually republished that row**, and the Points card was **not** republishing it. Adding a Point kept it on your phone and **never sent it to the shop**, so the shop went on showing only the Points it had last been told about. You reported exactly this - *"the point added still not able to appear on store"* - and it is fixed in **v301**: every change to a Point now republishes, which is the same call the Promo codes, Products and Categories screens have always made. **Pinning a Point still does not republish**, on purpose - the pin is never published at all. Nothing else in this entry changes.

04 Oct 2026 - engine v298, SELF COLLECTION POINTS - THE CARD (no database step, nothing to upload - pushing this one is the whole of it).

Why. You asked to get the system ready for collecting from places that are not your kitchen - Sg Ara, Farlim, Chai Leng Park, Bukit Mertajam - and to discuss it first. We did, and this is the first piece of it: **the card**, which is the piece you scoped yourself.

Her words: *"we just need to have a card for points"*

What you see now, on More -> Self collection Points.

- **A "New Self collection Point" card**, and your Points listed under it. **Five things per Point:** its name, the address, **who receives**, **their phone**, and **the fee per order** - what *you* pay whoever receives there. **Pause · Edit · Delete** on every row. A paused Point sinks to the bottom and goes pale, and **Resume** brings it back. **Every Point is its own record**, so opening one changes nothing about the others - which is how you said you'd do it: *"open collection point one by one... and not likely will open all point one go."*

Three things the card is careful about, and each is one of your rules.

- **Your kitchen is NOT a Point.** Collecting from Sg Ara is what your shop already offers - free, no minimum, always there. It gets no record, no fee and no provider, because it has none of those. A Point is a *third* thing beside it. **Delete never rewrites where an order went.** An order keeps the Point's **name frozen onto it**, so an order that went to Farlim still says Farlim after the Point is deleted - the same way an order keeps the name and price a product was sold at. **Pause is a normal ending, not a failure.** You said most Points you open will end this way, so pausing is one press and completely reversible.

WHAT THIS VERSION DOES NOT DO YET, said plainly rather than discovered. This is the **card only**. Your shop does **not** offer Points to customers yet, an **order does not carry which Point** it went to, and the **Delivery run does not include them**. Those are the next pieces. So today: you can define your Points and see them, and nothing else changes.

And the fee is yours to set, with no arithmetic done for you. You said *"allow me to manually set it fee x2 or x3, but default x1"* - so the app never measures an order's size. There is no product-size field and no volume sum anywhere.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. One new list is added and starts empty; no order, product, ingredient, price or bake day is touched. The suite is 2,646 tests, all green. Proved live on the real card at a phone width: **26 checks**, including that a nameless Point is refused with a reason, that pausing one Point leaves every other alone, that deleting one leaves the order still saying where it went, and that nothing scrolls sideways at 375px.

04 Oct 2026 - engine v297, THE OFFERS REALLY FLIP, WITH DOTS, AND A CLICK NO LONGER FREEZES THEM (no database step, nothing to upload - pushing this one is the whole of it).

Why. Three things in one message, and the third was a fault of mine. Her words: *"I dont like the flip, the flip should be 3D flip, and also i came with a flaw, once we put mouse over it or click it, the flip stop, there should be 2 dot if there is 2 message, 3 dot if 3 message. move the mouse outside the window, the flip should be back"*

1 - The flip is properly 3D now. v296 turned both panels through the **same** arc - the old one out one way, the new one in from where it had just left - which mirrors them the whole time and reads as a vertical **squash**. Now the offer being replaced tips **away over the top** while the next rises from **below**. Two panels turning through opposite arcs is a card turning over; that one difference is the whole of it.

2 - One dot per offer. Two messages, two dots; three messages, three dots; **one message, no dots at all** - the same rule that already leaves the turn unarmed, because one offer is a statement. Pressing a dot goes straight to that offer, and **does not stop the turn**.

3 - The flaw you found, and it was mine. v292 paused the flip for **twenty seconds on any click**. It was meant as the only pause a phone had, but it was a **clock, not the pointer** - so clicking the strip trapped it for twenty seconds, and moving the mouse away could not release it. That is exactly the symptom you described. **The pointer being over the strip is now the whole of the pause**, so moving the mouse out - of the strip, or out of the window entirely - always starts it again.

4 - I read your own design skill, which I should have done before touching this, and it caught four things I had wrong:

- **The dots were 22 pixels across. WCAG 2.5.8's floor is 24.** They are 28 now - the dot you see is still the same size as the ones on your homepage, and it is the button around it that carries the size. **They had no focus ring**, so anyone tabbing to them could not see where they were. They have one. **The turn used the same easing both ways.** Your skill says *ease-out coming in, ease-in going out* - that is what it does now, and it is the

difference between a turn that settles and one that just stops. **The turn was 420 milliseconds**, just outside the 250-400 band your skill names for a change like this. It is **380** now.

One thing I did NOT change, and it is yours to call. Your skill asks for body text to be pushed toward **7:1 contrast** because your customers read it on a phone in Penang daylight. The offer line sits at **5.3:1** - comfortably past the 4.5:1 that is required, but short of that goal. Getting to 7:1 means darkening the amber text noticeably, which is a change to how the strip looks. Say the word and it is one line.

What I could not check from here, said plainly. The preview pane was off-screen the whole time, and **a hidden page does not run animations at all** - so I could verify *where* each panel ends up (it is the opposite side, which is the 3D fact) but **I could not watch the turn itself**. Please look at it on your phone. Everything else was measured: the strip holds **169.4 pixels** and the page below it **380.1**, identical on every single turn, so the dots have not brought the jumping back.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. The shop's page and nothing else. The suite is 2,635 tests, all green. Proved live on the real shop page at a phone width: **26 checks**, including that a click no longer freezes the turn, that a pointer leaving the window starts it again, that pressing a dot goes to that offer without stopping the turn, and that the two panels sit on opposite sides.

04 Oct 2026 - engine v296, THE OFFERS FLIP RATHER THAN FADE (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your question, straight after v295 landed: *"can the flip be an animation"*. Until now the offers cross-faded into each other. Now they actually **turn**.

What you see. Each offer is a panel that **turns a quarter-turn into place**: the one leaving tips away from you while the next swings up to face you. Straight on when it settles, so it stays perfectly readable - it is the change between them that moves, never the text you are reading.

How the two halves stay out of each other's way. Both panels travel through the same angle as they swap, so if they faded at the same rate you would catch them **both half-turned and half-visible in the middle** - which reads as a smudge rather than as a turn. So the one leaving fades out quickly, in about a sixth of a second, and is gone before it is half-way round; the one arriving holds its fade back until the turn is nearly finished, then comes up. Two timing rules, no timer in the script.

The height fix from v295 is untouched, and I checked rather than assumed. A turn is paint-only - it cannot change how tall anything is - but the same measurement was run again with the animation in place, taken **mid-turn** rather than at rest: the strip holds **135.4 pixels** and the page below it stays at **346.1**, exactly as before. So it flips without the page moving.

If a customer has asked their phone for less movement, the offers still change - they simply arrive facing them, with no turn. That setting has been respected here since v292 and still is.

One thing to watch, and it is yours to tune. The turn takes about four tenths of a second out of every 1.5 seconds, so the words sit still for about a second. Whether that is the right balance is much easier to judge on the real shop than in a description - if it feels busy, say so, and either the turn gets quicker or the 1.5 seconds gets longer.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. The shop's page and nothing else: no order, product, ingredient, price, bake day or promo code is touched. The suite is 2,631 tests, all green. Proved live on the real shop page at a phone width - 14 checks, all passing, measured mid-turn - and the turn itself was frozen part-way round in a browser and looked at, to be sure it is a real turn and not a squash.

04 Oct 2026 - engine v295, THE SHOP'S OFFER STRIP STOPS MOVING THE PAGE (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your report, two codes running: *"when the message switch, the page is like jumping up and down repeatedly... The window should be fix, base on the tallest message."* You were right, and it was mine to fix - **v292 caused it.**

What was wrong. The strip is only as tall as the message it is showing. When one of your codes carries a sentence of your own and the other does not, those are **two different heights** - so every time the strip turned, it grew or shrank and **everything below it moved with it**. Measured on the fault, on a real browser at phone width: the strip swung between **40.8 and 135.4 pixels** and dragged the whole page up and down by almost a hundred pixels, every second and a half, for as long as a customer stayed on the page.

What it does now.

- **The strip is exactly as tall as your TALLEST message, and never changes.** Every running offer is drawn at once and stacked under each other; only the current one is lit. The box is sized by the longest, so a short offer simply sits inside it. Measured again with the fix: the strip holds **135.4 pixels** and the page below it stays put at **346.1**, through turn after turn. **It turns every 1.5 seconds**, your number. **Pointing at it still stops it**, and a press on a phone still holds it - both unchanged. **The slight dimming is gentler** now that the offers come round twice as fast, so the words are readable for most of the time they are up.

One thing worth knowing about the 1.5 seconds. A two-line message - the offer plus a sentence of yours - gives a customer about a second and a half to read it before it goes. That is quick. It is exactly what you asked for and it is one number to change, so if it turns out to be too fast once you watch it with real customers, say so.

One more thing it fixes, quietly. The lit offer now says which one it is to a screen reader, and the others are marked as behind it - previously, with all the offers in the strip, a reader could have taken all of them in turn as though they were one message.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. This is the shop's page and nothing else: no order, product, ingredient, price, bake day or promo code is touched, and no code's terms change. The suite is 2,630 tests, all green. Proved live on the real shop page at a phone width: 14 checks, including that the strip's height is identical on every turn, that the page under it never moves, and that the box measures the same whichever offer is lit. The fault was then put back on purpose - the strip immediately swung 94.6 pixels and the checks went red - which is what proves the measurement can see it.

04 Oct 2026 - engine v294, THE INVOICE NUMBER IS THE ORDER'S OWN CODE (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your words, the day v293 shipped: *"for the invoice, i think we can use the order code as invoice number"*. You were right, and it turns out to be **simpler and safer** than the running number v293 gave you.

What an invoice is numbered with now. The order's own code - **Invoice #A3F9C2** - which is the same six characters already on that order's row, in the confirmation message, in the payment reminder and on the customer's tracking card. **One order, one reference, everywhere.**

Three things that got better, and they are the reason this is a real improvement and not just a preference.

- **No two phones can ever take the same number.** The order code is unique by construction. The running number v293 used had one case where that was not guaranteed, and it had to be explained and accepted. **Pressing Invoice now writes nothing at all.** v293 stamped a number onto the order the first time you opened its invoice. There is no number to stamp now, so the press only reads - nothing on the order moves, and there is nothing to keep in step between your two phones. **There is no number for you to look up.** A reprint is the same invoice because the number was never stored in the first place.

The trade-off, said plainly rather than discovered later. The code is **not sequential**: nothing on an invoice says how many you have issued, or which of two invoices came first. And two different orders could in principle draw the same six characters - that is the same risk the order tag has carried everywhere since the shop was built, and it is unchanged by this version.

Please read the v293 entry below with this in mind. Everything in it still holds except the number: the letterhead, the items at the price they were sold at, the courier charge, the code taken off, the Total, Print and Share. Where that entry describes a running number that goes up one each time, **this entry replaces it.**

Your data is untouched, and there is nothing to run. No database step, no upload, no key. This version **removes** the two fields v293 could write onto an order, and writes neither of them - an order you invoice is not

touched at all, and no order, product, ingredient or price is rewritten. The suite is 2,625 tests, all green. Proved live on the real Orders screen at a phone width: 21 checks, including that the invoice number is the order's own code, that making an invoice writes nothing onto the order, that opening it twice is the same invoice, and that the paper carries your letterhead with the bakery's name said once.

04 Oct 2026 - engine v293, AN INVOICE FOR A CUSTOMER (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your words: *"And customer need an invoice"*. Asked what it should carry, you chose *"One order, one invoice"* and *"Yes - name, address, a number"*.

What you see now, on Orders -> any order -> Invoice (beside Edit and Note / tracking).

- **A card opens with the invoice on it**, and **Print** and **Share** underneath - the same two presses every other book in the app has. Share hands it over as a **PDF file**, so it goes into WhatsApp as a document the customer can keep. **It carries your real address at the top.** That address already exists - it is the one you typed for the Mailing labels in Settings, the same block that prints as FROM on a parcel. So an invoice and a parcel can never show two different addresses for one bakery, and **there is nothing new to type**. **One order, one invoice.** The number is given the first time you open it and is never given again: open the same order next year and it is the same invoice, with the same number and the same date. ***(Superseded by v294, 04 Oct 2026: the number is now the order's own code - Invoice #A3F9C2 - so there is no number to give, nothing is written onto the order, and the two-phone caveat below no longer applies at all. Read the v294 entry at the top of this page.)*** **The order's own code prints beside the number** - *Invoice 0007 · Order #A3F9C2* - which is what a real invoice does, and it means two invoices could never be confused for one another. ***(v294: the code is no longer beside the number - it IS the number.)***

What the paper says. Your name and address, then **Invoice 0007 · Order #A3F9C2 · the date the order was placed**, then one row per item as **4 × Focaccia** with the line total in the money column, the **courier charge** when the customer bears it, the **code taken off as a minus row** named after the code, and finally the **Total**.

Why the figures can never disagree with anything else. Every figure on the invoice is read from the one function your confirmation message, your tracking card and your order rows already read, and the item names and prices are the ones **frozen onto the order** - so an invoice for an old order never shows a product you have since renamed or today's price. An invoice cannot state a sum the rest of the app contradicts.

The one limit, said plainly rather than hidden. Your two phones hold one shared copy of your business, and a running invoice counter kept in Settings would be **overwritten by whichever phone saved last** - numbers would repeat or skip without warning. So the number is written **on the order itself**, which is a record of its own and cannot be lost. The consequence: **two phones issuing an invoice in the same instant could take the same number.** Nothing is lost when that happens - both invoices exist, each on its own order - and the order code tells them apart. With one baker and two phones it is vanishingly unlikely, and the alternative would be an invoice you cannot write without the internet, which is worse for a bakery. ***(Superseded by v294, 04 Oct 2026: the order's own code IS the number now, so this whole limit is gone - there is no counter to share, and pressing Invoice writes nothing at all. Read the v294 entry at the top of this page.)***

What is deliberately NOT on it. Your customer's per-item note is not printed - it is a production instruction, not a line item. Say the word if you would rather it were, and it is a one-line change.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. Two fields are added to an order **only when you press Invoice on it** - an order you never invoice is not touched at all, and no order, product, ingredient or price is rewritten. The suite is 2,627 tests, all green. Proved live on the real Orders screen at a phone width: 22 checks, including that the Total is the customer's own total to the cent, that the number is written on every row of the cart, that opening the invoice a second time is the same invoice, that Print really reaches the printer, and that your address block never says the bakery's name twice.

04 Oct 2026 - engine v292, THE SHOP'S OFFERS TURN INSTEAD OF HIDING (no database step, nothing to upload - pushing this one is the whole of it).

Why. Your words: *"the public code shown in shop, if more than one hide behind each other, we need a carousel,*

where it message turn, put the mouse over it stop rotate." And you were reading it exactly right. The shop's offer line took the **first** live public code and quietly dropped every other one - so a week when you were running three offers, your shop was advertising one of them and hiding two behind it.

What your customers see now, at the top of the shop.

- **Every code you have running takes its turn.** About six seconds each, then the next one - the same pace as the reviews carousel on your homepage. **The words fade; the strip does not.** The amber box stays exactly where it is and only the line inside it changes. The whole box blinking off and on would read as a broken page rather than as a second offer. **Point at it and it stops.** It waits while the pointer is over it and starts again when you move away. **On a phone, a tap holds it still for about twenty seconds** - long enough to read a long offer - and then it carries on by itself. A tap can never leave the strip stuck on one offer. **A background tab stops it**, and it starts again when the tab comes back. **With only one code running, nothing turns at all** - no movement, and literally no timer running. One offer is a statement, not a one-slide carousel. **And if a customer has asked their phone for less movement**, the words still change so they see every offer - they simply change without fading. The setting is about movement, not about hiding things from them.

Which code shows first, and what counts as running. The first one you published. A code is advertised only if it is public, switched on, inside its dates and not used up - asked through the same rule the code box already uses, so the standing line and the box can never disagree about what "still running" means. A **personal** code is never advertised; being unadvertised is the whole of what personal buys. If you pause or end a code while someone is looking at the shop, the strip picks up the change on its own and does not leave them reading an offer that has gone.

Your data is untouched, and there is nothing to run. No database step, no upload, no key. This version changes the shop's page and nothing else - your orders, products, ingredients, prices, bake days and promo codes are exactly as they were, and no code's terms are changed by it. The suite is 2,611 tests, all green. Proved live against the real shop page: 50 checks across three states - three codes running, one code running, and three codes with reduced motion - including that the turn wraps at the end, that hovering and pressing both stop it, that it starts again by itself after a press, that a hidden tab stops it, and that the amber strip itself is never faded.

04 Oct 2026 - engine v291, A REWARD YOU CAN ACTUALLY HAND OVER (no database step, nothing to upload - pushing this one is the whole of it).

Why. You asked the question that made this version: *"how do we exercise their reward, if the reward is only written text?"* Until now the reward was a sentence on the customer - you read it and settled up by hand, and the app neither reminded you nor remembered what you had given. Nothing said **when** a reward had come round.

What you see now, on More -> Customers -> open a customer.

- **The reward is a card of its own on their record**, in amber so it does not read as the green bring-a-friend ledger below it. **It says what they brought in, and what is due.** *"9 brought in · every 5 · 1 due"* - the count is from your own orders, recounted every time, so nothing can double-count. **A Given button records the hand-over.** Their card then says *"1 given"* with the date, so the next time you open them you can see what you have already settled. **Undo last takes a press back**, if you tapped it twice or they returned it.

Their reward, and its number, are two separate boxes. Your sentence stays exactly as you wrote it - *"a free loaf for every five friends"*. Beside it is a number box, **"Given every ... customers brought in"**. Nothing is ever read out of your sentence: a parser that misread "every five" would tell you a partner is owed a loaf you never agreed to. Leave the number empty if it is not a set figure - the app still counts what they brought in, it just never calls one due.

What counts as "brought in" is one number, whichever way they work. A friend with a share link and a partner with a code are counted the same way, and a person can be both - so the count is the union of the two, never their sum. A cart that carried both a link and a code counts **once**. And a friend who had **already ordered from you** is not a new customer here, exactly as the Give-credit button has always judged it.

Every hand-over is its own record, and that is deliberate. Both your phones hold one shared copy of your business. A "rewards given" tally kept as a single number would be **overwritten by whichever phone saved last** - one phone's hand-over would vanish in silence. One record per hand-over cannot lose an update. It is the same reason your credit ledger is a list.

Nothing is ever blocked. The Given button is offered whether or not the app's arithmetic says a reward is due - you may settle a favour early, or hand one over for a reason the app cannot see. The count is information, never a gate.

Your data is untouched, and there is nothing to run. No database step, no upload, no new key. This version adds one new list and touches nothing you already have: your orders, products, ingredients, prices, bake days and saved plans are exactly as they were. The suite is 2,607 tests, all green. Proved live at a phone width: 29 checks, including that a cart carrying both a link and a code counts once, that a second hand-over is recorded rather than replacing the first, that Undo takes back exactly one, and that the number box really reaches the customer's record.

04 Oct 2026 - engine v290, YOU CAN ADD A CUSTOMER YOURSELF, AND READ YOUR REMARK ON THE LIST (no database step, nothing to upload - pushing this one is the whole of it).

Why. Two things, and they arrived together. **A partner who has never ordered from you had nowhere to live** - until now a person existed only by placing an order, so a partner you recruited to hand out labels could not be named as a promo code's owner. And you wanted to **read your note about someone without opening them**.

What you see now, on More -> Customers.

- **A "New customer" card.** Add someone by hand - a partner, or a friend who sends people your way. **A name or a number is enough**; everything else (their reward, their note, a photo) is written exactly where it always was. **They appear under their own heading: "Added by hand - no orders yet"**, at the bottom of the list, with a count. Their card opens like anyone else's, and fine-tuning their reward works there too. **Your remark now shows on the row itself**, so you can read it without opening anyone - in your own words, on one line, whichever is longest clipped with the full note still one tap away.

The important part, and why they are in the same list rather than a separate one. Your customer list stays what it was - **people who have ordered**, with their spend and their last delivery - and the people you added simply sit underneath, grouped. They are counted, they export to the CSV, they can be messaged in bulk, and **they are offered by the name box when you take an order**. That last one matters most: if they had been kept in a list of their own, the day your partner finally ordered the app would have offered you nothing, you would have typed the name and number by hand, and that one person would have become **two records** - one keyed by name, one by number, for good. **The moment they order they move up into the list proper, by themselves.**

Two things that would have read as faults. A hand-added person has no orders, and their row would have said **"0 orders · 0 units · about RM 0.00"** - the shape this app uses for a broken screen. It now says what is true: *"Added by hand - no orders yet"*. And opening their card said *"This customer's orders were removed"*, which is untrue of someone who never had any, and reads as lost data. It now says they have not ordered yet.

No database step and nothing to upload - admin/ only.

04 Oct 2026 - engine v289, A NAMED REWARD, AND A CODE THAT NAMES ITS PERSON (no database step, nothing to upload - pushing this one is the whole of it).

Why. Two schemes, one missing half each. The **bring-a-friend** link is for a casual advocate and costs you nothing to hand out - but its reward could only ever be a flat **RM3**, the same for everyone, set once in Settings. And a **promo code** was anonymous: its orders and its label's opens were counted, but not *whose* they were, so a partner's label couldn't be told from anyone else's.

Both stay, and the difference between them is the thing you hand over - a *link* for a friend, which travels and costs nothing; a *code and a label* for a partner who prints brochures and runs their own marketing.

What you see now, in two places.

- **On a customer's card** - a line you write yourself: **Reward**, e.g. *"a free loaf for every five friends"*. Not a number, because a partner may be owed a loaf, a favour, or an arrangement of their own. It sits in the profile block, above the bring-a-friend block, so it shows **whether or not that person has a WhatsApp number** - the casual friend-to-friend advocate is exactly the one who may not. The RM3 and 90-day settings are untouched and still the defaults. **On a promo code** - a new choice: **Whose code is this**. Pick a customer and their name appears on the code's row as *Aunty Bee's code*, beside "public - shown in the shop". Their label now tells itself apart from anyone else's, and the code's own count and opens are their tally.

THE THING TO KNOW: their name is never published. The data your shop is given is readable by anyone holding its public key. A code's person is kept in your own app - **never** sent to the shop, and **never** printed on the label. The label stays impersonal, which is right for a piece of paper handed to whoever walks past.

What the reward is, and is not, said plainly. It is a **label you apply by hand**, exactly like the credits. The app names it and counts what that person brought in; it does **not** total what you owe, and it does not track what you have already given. That is the trade for being able to write *"a free loaf"* instead of a figure.

No database step and nothing to upload - admin/ only.

04 Oct 2026 - engine v288, HOW MANY TIMES EACH LABEL WAS OPENED. THIS ONE HAS A DATABASE STEP - run supabase/promo_visits.sql once, see the box below.

Why. You could already see what a code **sold**, worked out from your own orders. What you could not see is whether the label was **picked up at all** - and those are different problems needing different answers. *"Nobody followed the link"* means print more cards, or hand them out somewhere else. *"Forty people followed it and two bought"* means the card is fine and the offer needs work.

What you see now. On **More -> Promo codes**, each code's row shows **how many times its link was opened**, and under that a strip of the **last 28 days, one bar a day** - so a label going cold is visible at a glance rather than something you have to work out from a date list. A day with no opens is a faint stub rather than a gap, so "quiet" never reads as "no data". A code the app *did* get an answer for and which nobody has opened yet says **"Not opened yet"** in words - that zero is real and worth knowing.

What the number is, and what it is not - please read this part.

- It counts **opens, not people**. You chose that: a reload counts again, and a phone that leaves the page open counts once. It is a **pulse for alive-versus- cold**, not a headcount. **Sharing the link in WhatsApp or Facebook adds an open with nobody behind it** - those apps fetch a link to draw the preview. And your own testing counts. **A customer who types the code at the shop, with no link, is not counted.** This measures the CARDS, not the code in general. The app will not show a figure it does not have. If Supabase cannot be reached, the row is left **exactly as it was** rather than showing a zero - because a zero here is a claim ("nobody opened your label") and it must not be made on the strength of a request that never came back.

The one step, and it is yours: run supabase/promo_visits.sql once. Open **Supabase -> your project -> SQL Editor -> New query**, paste the whole file, press **Run**. It is safe to run more than once. **You can run it before or after you push** - if you push first, nothing breaks: the shop's page simply fails to record opens until the table exists, quietly, and the Promo screen shows no count. (Your own browser's console will show a **404** while the table is missing. That is the missing table, not a fault, and no customer ever sees it.)

What was built, and why it is shaped this way. One row is stored per open. The shop page records it in the background and never waits for it - a visit that cannot be recorded must not slow a customer down or change a word on the page. **The shop counts an open once per page load**, which matters more than it sounds: that page re-reads its settings every 30 seconds while it is open, so without that rule a single customer leaving the tab open would have added an open every half-minute and a label that sold nothing would have read as a triumph. Measured live: the settings were re-read **7 times in one minute and the open was recorded exactly once**.

No SQL runs on its own and nothing is uploaded - the Edge Functions are untouched.

04 Oct 2026 - engine v287, EVERY CODE HAS A LABEL, AND PRINTING NO LONGER FREEZES IT (no

database step, nothing to upload - pushing this one is the whole of it).

Why. In your words: *"instead of printed card, i think something like what we have in shop and code is more useful, we have QRs, some active some retired, when a promo code come together with a QR, when you tab on label, you are allow to copy, print."* Until now the only QR in the app was drawn on a **separate printed-card page**, reached by a **Print it** press that was a deliberate point of no return - it froze the offer permanently. There was no copy, and no QR anywhere on screen.

Every code now has a label, on its own row. Under **More -> Promo codes**, each code carries **its own QR**, so the list reads as a set of labels with the life chips saying which are still going. **Tap the label** and it opens: the QR drawn large enough to hold another phone up to, the link in words, **Copy link**, and **Print it**. The square is built from the very address the printed card uses, so a label and a card always point at one place - and **Copy link** pastes straight into WhatsApp, where the shop opens with the code already in the box.

A retired code keeps its label. An **ended** code still shows its QR and still offers **Copy link** - a dead link is worth being able to look at - but it is **not offered a print**, because a label with an ended code in it would not work. It says so rather than simply hiding the button.

Printing no longer freezes anything. Until this version, the first print pinned the offer for good: the amount, who it is for, the smallest basket, what it cannot sit beside and the name all stopped moving, an end date could only be moved later, a ceiling could only be raised, and the code could not be deleted at all. **All of that is gone.** Print and copy as often as you like, change the offer whenever you like, and **retire a label by ending the code** - which is what ending was always for: it stops new uses and leaves orders already placed with what they were promised. There is no new thing to learn; **End** is the retirement, and **Pause** is the reversible version of it.

The one thing to know, and the app says it on the label itself. A label already in someone's hand is honoured at whatever the offer says when the customer **orders**, not when they picked it up. So if you print a hundred cards and then change the offer, those cards give the new offer. That is the trade, and it is better to read it here than to find out from a customer.

What has NOT changed. A code still needs a **name** and a **cost ceiling** before it can be given a label - a label carries no number and no end date, so the ceiling is the only thing bounding what it can cost you. Editing a code still cannot reach back into an order that already used it: every order keeps the code as it was written when the customer typed it.

No database step and nothing to upload. This is admin/ only - no SQL, no Edge Function, no key. Pushing it is the whole of it.

03 Oct 2026 - engine v286, A SUGGESTED PROMO CODE YOU CAN READ OFF A CARD (no database step, nothing to upload - pushing this one is the whole of it).

Why. You asked me to compare your promo codes with the printed-label codes munchies used to run and see if there was anything worth taking. There was one thing, and it is small: a code is read by eye **twice** - you type it when you make it, and the customer types it off the printed card. The pairs people get wrong doing that are **0 and O**, and **1 and I or L**. Your codes were allowed to contain any of them.

What you see now. On **More -> Promo codes**, the **New code** card has a **Suggest one** button beside the code box. Press it and the box fills with a code such as PCX68 - five characters drawn from an alphabet with no **0**, **O**, **1**, **I** or **L** in it, so there is nothing on the card to misread.

Nothing is chosen for you. The box is still yours to type in, and the button only fills it - type over it, or ignore it and type your own code as you always have. **Every code you already have keeps working exactly as it did**, including ones with those characters in them: a code that is already printed cannot be renamed, and none of them were touched.

It will not hand you a name you already use. Before suggesting, it checks every code on the list and never offers one that is spoken for - two codes sharing a name is the one thing the shop could not recover from, because it would take the wrong amount off.

The button leaves the room to the box. The code box asks for a sensible minimum width; if the screen is wide

enough they sit on one line, and if it is not the button drops below and the box takes the whole width. Either way the box is never the part that gets squeezed - measured at a 375-pixel phone screen the box is 181 pixels beside the button, and on anything narrower it is 208 or more.

Where this came from. Munchies' printed-label codes carried the same idea, and their comment names the reason exactly: *"a customer or the owner may type the code by hand off a printed label, and those pairs are the ones people get wrong."* This is the only piece of theirs worth taking; the rest of theirs is a printed-label pipeline for shops, which you do not have. **No database step and nothing to upload.**

03 Oct 2026 - engine v285, THE SHOPPING LIST CAN BE CORRECTED AT THE SHOP, AND AN INGREDIENT KEEPS ITS PRICES (no database step, nothing to upload - pushing this one is the whole of it).

Why. In your words: *"PO, say it is created an base on the po we go shopping, same supplier price change and we decide to buy more, i would like to change the price and the qty, i need the PO to be amendable, ingredient price journaled, ingredient price updated accordingly. So an ingredient need a journals."* A saved shopping list was frozen the moment it was saved. So when you got to the supplier and the price had moved, or you decided to take more, there was nowhere to put it - and the app's idea of what that ingredient costs stayed wrong until you remembered to go and correct it on another screen.

The list is now amendable - until you tap Bought. On a saved list you now have an **Amend** press. It opens the list with two boxes per line: **how many packs** you took, and **what each pack cost**. Change either and the line's total and the list's total follow as you type. You can also **Remove** a line you didn't take, **add a line** for something you picked up that wasn't on the list, and - if the shelf already covered something - press **buy some** on it to decide to stock up anyway.

Once you tap Bought, the list is what happened. That is deliberate. Bought has already put the packs on your shelf and asked you what you paid, so correcting the list afterwards would mean unpicking both. Amend is offered only while the list is still un-bought; the correction belongs at the shop, which is where you are when you discover it.

Your ingredient follows the price you actually paid. Change a price on the list, save, and the app writes that price onto the ingredient - so your recipes, your product costs, and your next shopping list all use it, with nothing else to set. The app tells you it did it rather than doing it quietly.

And every price move is now recorded. Open an ingredient and, once its price has moved at least once, it carries a **Journal** press. It lists **only the moments the price moved, newest last**, each row saying what it moved to and what it was before - and with **the price you are on now at the top**, because a list of changes on its own never says where you ended up. Buying the same thing again at the same price is stock, not news, so it is not listed. The journal prints and shares like every other book in the app, including as a PDF.

Why a journal matters here. Your Profit screen reads Cost of sales from your recipes and the ingredient prices as they stand **today** - so changing a price moves months that have already closed. That has always been true and the screen says so; what was missing was any record of **when** a price moved, and this is it.

A fault found and fixed while building this. The **"What did you pay?"** box you get after tapping Bought was **never pre-filled** with your list's total. It had been reading the total from the wrong place since the box was built, so it always opened blank and its sentence never named the figure. It is now filled in from the list's own total, which is also what makes amending worth doing.

What did not change. Nothing that was already saved was rewritten, no figure moved on its own, and the money is still recorded exactly where and when it always was - at Bought, from the account you choose in that box. A list's own totals still pre-fill that box. **No database step and nothing to upload.**

03 Oct 2026 - engine v284, A LONG COURIER LINK NO LONGER RUNS OFF THE CARD (no database step, nothing to upload - pushing this one is the whole of it).

Why. You sent me a customer's tracking card and said: *"the lalamove link is very long and go out of bound."* Once your Lalamove account went live, a real booking came back with a real share link - one unbroken run of

about 150 characters, with no spaces anywhere in it. Your customer's card draws that link as a **rounded button**, and a button cannot break a word it cannot find a space in. So the link stayed one enormous line and grew straight out of the card and off the screen.

What you see now. The link stays **inside** the card. It wraps onto as many lines as it needs, and the button grows downwards instead of sideways. Nothing else about the card moved: the wording, the order of the lines, the money and the progress line are all exactly as they were, and the link still opens the courier's own page in a new tab.

A short link is untouched. Measured on the card at a 375-pixel phone screen: a short link such as https://track.jt.com.my/A3F9C2 is still exactly the **38 pixels** tall it has always been, because the button's own floor and its centring are unchanged. Only a link long enough to need a second line behaves differently.

Measured before and after, on your own order. Your order's link is 131 characters. Before: it ran past the card's right edge and off the screen. After: it wraps to five lines, its right edge sits at 344 against the card's 361, and the page no longer scrolls sideways - the page's own scroll width equals the phone's width. **Your backoffice card was never affected**: the booked-trip card on your own screen already carried this rule, so both of your screens now stop a long link the same way.

Where else that link is drawn - checked, not assumed. You asked me to check everywhere it lives, so I did. The link is stored once, on the order, and reaches your customer in three places: the **customer's track card** (the one that was broken - fixed here); the **shipped WhatsApp message**, which carries it as plain text that WhatsApp wraps in the chat itself, so it was never at risk; and your **own booked-trip card**, which was already right. The delivery-run screen only mentions it in a passing message, and the orders list draws it as a number to read out rather than a link, so neither of those can overflow.

No database step, and nothing to upload. This is one style rule in the shop's stylesheet - no SQL, no Edge Function, no key, and your Lalamove account and your data are untouched. Pushing it is the whole of it.

03 Oct 2026 - engine v283, SHARE SENDS THE JOURNAL AS A PDF FILE (no database step, nothing to upload - pushing this one is the whole of it).

Why. You pressed Share on a journal and the share sheet opened **without WhatsApp in it**. Nothing was broken - a phone is handed a journal as ordinary text, and WhatsApp does not offer itself for a bare block of text. It does offer itself for a **file**. You said to make it share by PDF, and that is what this is.

What you see now. The **Share** button does the same thing it always did, and what comes out of it is different: the journal arrives in the chat as a **PDF document** - Cash journal.pdf, Profit and loss.pdf - which you can open, keep, forward and print again. It is listed for Mail, Notes and Save to Files exactly as before, and WhatsApp is on that list now.

Worked out with your phone, in this order. The app asks the phone what it can take rather than assuming:

- The phone can share a file - you get the **PDF**, which is what happens on yours. The phone says it can share, but not a file - you get the journal as **text**, so it still leaves the app rather than nothing happening. The phone has no share sheet at all - the **PDF is saved to the phone** and a short message says so. You send it from WhatsApp yourself. Even saving is impossible - the journal is **copied** instead, with a message saying so.

If you open the share sheet and change your mind, closing it does nothing else. Nothing is copied and nothing is saved behind your back, at any of those four steps. A cancel is your decision, not a fault.

What is on the document. Your shop name as a letterhead, the journal's name and the stretch it covers, every row and every total in the same words and the same order as the screen, the note that explains the figures, and along the bottom From More -> Profit · printed 3 Oct 2026. Money that went out reads -RM 12.00. A long journal runs onto a second page, which says at the top what it continues. A4, black on white, 20 mm margins. It is about 3 KB - small enough for any chat.

Print has not changed. The Print button still opens the phone's own print sheet, and **Save as PDF** is still a choice inside it.

What did not change. No figure moved and no rule about your money changed. The screen is identical. The

document is a fourth rendering of the same journal the screen, the paper and the sent message already shared, so the figure in the file is the figure on the screen, to the cent.

03 Oct 2026 - engine v282, EVERY JOURNAL CAN LEAVE THE SCREEN: PRINT IT OR SEND IT (no database step, nothing to upload - pushing this one is the whole of it).

What you asked for. You were reading a journal and said: *"those journals in profits and other journals should be printable and able to be shared"*. Until now a journal existed only inside the pop-up you were scrolling. You could read it, and there was no way to hand it to anyone or keep it.

What you see now. Every journal in the app wears the same two buttons, in the same order, under its last line:

- **Print** - the phone's own print sheet opens with the journal on it. That is also where **Save as PDF** lives, so a journal can leave as a file. **Share** - the phone's own share sheet opens with the journal as plain text, ready to go straight into WhatsApp, Mail or Notes. A phone with no share sheet - or a share that fails for any other reason - **copies the journal instead** and says so in a short message. You paste it wherever you want it. The fallback carries exactly the same text the share sheet would have been handed. **If you open the share sheet and change your mind**, closing it does nothing else. The journal is not copied behind your back. A cancel is your decision, not a fault. *Superseded the next day by v283: the phone is handed a PDF file now, not plain text, because WhatsApp was not appearing in the share sheet.*

Where the two buttons are. Every book in the app:

- **More -> Profit** - the Sales journal, the Cost of sales journal, and every spending category's journal, including Total expenses. **More -> Profit** - the **Profit and loss** statement itself. You chose that the whole statement gets the pair and not only the journals: Sales down to Net profit, with its two notes underneath. A journal on its own is half a document, and this is the other half. **More -> Money** - one method's book, opened either from its row on the screen or from **Books**. In Books, only the book you have opened can print; the list underneath never does.

What the printed sheet says. Your shop name as a letterhead, then what the journal is and the stretch it covers; every row and every total, in the same words and the same order as the screen; a line in the journal's own words saying what the figures mean and where they come from; and along the bottom, From More -> Profit · printed 3 Oct 2026. The statement's sheet carries the Cost of sales note as well, because anyone reading Gross profit needs it. It is plain black on white.

One rule this build rests on. The screen, the paper and the sent message are three renderings of **one** description of the journal. They cannot disagree about a row or a total: whatever figure is on the screen is on the paper to the cent, and the same figure is in the message.

On the phone. Pressing Print opens the phone's real print dialog - the same one any other app uses - and **Save as PDF** is a choice inside it. Choosing a printer there is the only way to see the true paper; the sheet itself is built by the app.

What did not change. No figure moved, and no rule about your money changed. Every pop-up looks as it did; the two new buttons sit under the last line of each.

03 Oct 2026 - engine v281, THE STATEMENT SAYS WHAT KIND OF COST IT IS SHOWING (no database step, nothing to upload - pushing this one is the whole of it).

08 Oct 2026 - THE CLAIM BELOW WAS TRUE WHEN IT WAS WRITTEN AND IS NOT ANY MORE. Read it with v380. The note on the screen used to say the cost was read fresh, so editing a recipe or a price moved past months too; **from v380 the cost is frozen onto each order as you take it, so a closed month no longer moves.** The note has been rewritten to say what it does now, and to admit plainly that an order taken before v380 was pinned at the cost it was already showing rather than at what it really cost that day. Everything else in this entry still stands: it is a recipe cost, not money you spent, and the Money screen is still where the cash is.

Why. You asked whether Cost of sales is really your recipe's ingredient cost - the theoretical figure - rather than money you had spent. It is, and the screen did not say so where the figure is. It only said a part of it in the small print underneath, and nothing at all about the part that matters most: that the cost is read from your recipe and

your ingredient prices *as they stand today*.

What you see now, on More -> Profit. A short line sits under **Gross profit**, inside the statement, in the same quiet grey as the other notes:

- **"Cost of sales is built from the recipe and the ingredient prices you have recorded, read as they stand today - so editing a recipe or a price moves past months too. It is not what you actually spent. Gross profit is therefore a guide to your pricing, not your bank balance - the Money screen is where the cash is."**

This is the wording as it stood in v281. Since v380 the same line reads that the cost is frozen onto each order as you take it, so a later change does not rewrite a sale already made, and that orders taken before v380 were pinned at the cost they were already showing.

Two things that follow from it, now said on the screen rather than left for you to work out:

- **A month that has already closed can still move.** If you fix a recipe, or update what an ingredient costs, the Cost of sales of a past month changes with it, and so does that month's gross profit. The sold price never moves - that is frozen on the order - but the cost side is read fresh every time you look. **Gross profit and Net profit are a plan, not a bank balance.** What you really spent is cash, and the Money screen is where you check it. A gap between the two is normal and not a fault.

The small print at the bottom changed too, so the two do not repeat each other. It now leads with the cash half - a pack bought today is cash on the Money screen and stock on the shelf, and only becomes cost of sales as the bread made from it is sold.

What did not change. No figure moved. The Sales journal, the Cost of sales journal, the month arrows and every spending line behave exactly as they did in v280. This is words on the screen, nothing else.

02 Oct 2026 - engine v280, THE WHOLE STATEMENT OPENS: SALES AND COST OF SALES TOO (no database step, nothing to upload - pushing this one is the whole of it).

What you asked for. You were in More -> Profit, reading the statement, and asked: *"at the profit section, can the sales and cost of sales be clickable to reveal its journal"*. The spending lines had opened on a tap since 17 September - these two did not.

What you see now, on More -> Profit. Tap **Sales** and its journal opens; tap **Cost of sales** and that one opens. Both read the same orders - one from the side of what the customer paid, the other from the side of what your recipes say those loaves cost to bake - so the two can never disagree about which orders the month held.

- **Each row is one order line**: the day it is delivered, what was sold and how many (Focaccia × 2), and the customer's own name, with the figure on the right. In the Sales journal that is the amount sold; in the Cost of sales journal it is what those loaves cost to make. **A product you have since renamed or deleted still reads as the loaf that was sold** - the name frozen on the order, not the live product's. The journal is a record of what happened, so nothing in it is re-described later. **The rows land on the figure you tapped.** The Sales journal totals to the Sales line, the Cost of sales journal to the Cost of sales line, exactly - the same guarantee the spending journals have kept since September. **A line your recipe prices at nothing is marked** "no recipe cost" rather than sitting there as a plain RM 0.00, and the footer counts them. That is a profit reading too high, and it is the one thing in these journals worth chasing. **A month with nothing in it opens and says so** - "Nothing was sold in October 2026" - rather than a line that looks live and does nothing when tapped, which is the fault you reported in September. **Sales counts on the day the order is DELIVERED**, not the day it was ordered, which is the same day the day headers, the weekly numbers and the customer's own calendar use. The journal says so under the total, because a sale counted on a day other than the one you expected is the one thing here that could look wrong and not be.

One thing that stays a figure, not a door. Gross profit and Net profit are worked out from the lines above them, so they have no rows of their own to open. They stay as they are.

02 Oct 2026 - engine v279, THE PRINTED CARD: PRINT IT, AND THE OFFER IS FIXED (no database step, nothing to upload - pushing this one is the whole of it).

What you asked for. Of the eleven steps, step 6 - Print it - was the last one with nothing behind it, and it is the step that puts a code into a customer's hand. You asked for the printed card next.

What you see now, on More -> Promo codes.

A public code's row has a "Print it" button. It sits first, ahead of Edit and Delete, and it only appears where a card makes sense:

- on a **public** code, because a card is a public thing and a personal code is never advertised; not on one that has **ended**, because there is nothing left to hand out; not on one that is **already printed** - that code wears its "Printed - fixed" chip instead, and the chip is the answer.

A **paused** code can still be printed. Pausing is a break, and printing a card for a code you are about to switch back on is exactly the right thing to do.

Pressing it asks first, and the question is honest about what it is. It says this is the point of no return: from here the amount, who it is for, the smallest basket and the name all go on saying what they say, because the card in the customer's hand cannot be amended. It also tells you what stays yours - the end date, later and never earlier, and nothing else - and, when the code has one, it names the ceiling the code will stop itself at.

It refuses a code with no cost ceiling, and says why. This is the one new rule, and it follows from the paper. A card carries no number and no end date, so the ceiling is the only thing left bounding what the card can cost you: without one, a launch that takes off has nothing to stop it and the paper cannot be recalled. The refusal names step 4, on that same code's own row, and it applies to every kind of offer - ringgit off, a percentage, and free delivery alike.

A code with no name is refused too, because a card that prints no code is a card the shop cannot accept.

What the card itself looks like. A new tab opens on a print-sized page holding **four identical cards on one sheet of A4**, with a hairline guide to cut along. Each card carries your bakery name and tagline, the offer in the customer's own words, the smallest basket if there is one, your own sentence for the shop when you wrote one, the rules that go on being true ("First order only", "One per customer", "Not with the bring-a-friend welcome discount"), the code in large letters, and a square to scan.

The square points at your shop with the code already filled in. A customer who scans it lands on the shop with the code already in the box - they type nothing. The same address is printed underneath in words, so a card that ever pointed somewhere wrong would be wrong visibly, on the paper, before anyone was handed it. Every square this version draws was scanned back and read correctly before it was called done.

What the card deliberately does NOT carry: any end date, any count, any ceiling figure. Those three are promises paper cannot keep - a date is one you may have to move, and "the first fifty orders" is one the card cannot count - so they are left off rather than printed and broken later. The rules it does carry are the ones that will still be true for the card's whole life.

Opened by itself, the page never prints a blank sheet. Opened with no code, or with a code that is not on that device, it says so plainly, tells you nothing has been changed, and offers no Print button at all. A blank sheet is the one outcome that could reach a customer, so it is the one outcome the page refuses to produce.

Undoing it. Printing is the point of no return, so nothing is printed until you press Print it in the app, and cancelling the question changes nothing. Once a code is printed you can still **raise** its ceiling, **extend** its end date, and **end** it - the same three freedoms as before. What you cannot do is change what it gives, who it is for, or the smallest basket, because a customer is holding a card that says so.

How this was checked. 2,497 automated checks pass, including the card against a frozen code field by field, the two promises it must never print, and every square against a reference encoder - and each new check was then broken on purpose to prove it can fail. The card was also opened and driven on screen: it draws, the square scans into the shop, and the empty and unknown cases say so.

02 Oct 2026 - engine v278, A CODE'S LIFE: PAUSE, END, AND WHAT PRINTING FIXES (no database step, nothing to upload - pushing this one is the whole of it).

What you asked for. The Promo codes screen had a half-finished piece. The eleven management steps were drawn as a plain grey box that did not match the rest of the app, and the Pause and End controls for a code's life were not wired to anything. You chose **"Finish it properly"** - the real screen, the buttons working, the rule about printed codes, and a test for each.

What you see now, on More -> Promo codes.

Every row says what its code's life is. A paused or ended code used to look exactly like a live one - the only tell was that it had quietly stopped being offered. Now the name carries a chip:

- nothing at all on a **live** code, because a row that is working needs no label; **Paused** on a paused one; **Ended** on an ended one; **Printed - fixed** on a code that is already on a card.

Pause and End sit on the row, under the offer. They are different things on purpose:

- **Pause** is a break. It stops the shop offering the code and accepting it, it can be switched back on at any time, and it changes nothing about the orders that already carry it. **End** is final. New uses stop and it cannot be switched back on. The confirmation says exactly that, and points you at Pause if a break is what you wanted.

Both tell you what they are about to do before they do it, including how many orders already hold the promise - "**2 orders carry it already and keep what they were promised - nothing already promised changes**" - and, when ending, how much the code has given away in all. An ended code keeps every order it already had: ending stops new uses, it never rewrites history.

A printed code is fixed. Once a code is on a card in someone's hand, the offer on that card has to go on being true. So a printed code refuses any change to what it gives, who it is for, the smallest basket, what it cannot sit beside, or its name. Two things stay open, and only one way round, because being generous with someone holding a card cannot hurt them and taking something back can:

- the **end date may be moved later**, or dropped altogether. It may never be pulled earlier, and a code that never ran out cannot be given a date now. the **ceiling may be raised**, or removed. It may never be lowered.

The Edit button on such a row reads **"Update the end date and ceiling"** rather than "Update code", so the label never promises a change the code will refuse. When a change is refused, the row says which part of the card it would have broken.

A printed code cannot be deleted. A card in someone's hand would simply stop working, with nothing to explain why, and the row would be gone so you could not even see that was what happened. The tap is not dead - it says why, and names **End** instead, which stops new uses and leaves the card honest.

The eleven steps fold. The life of a promotion, in the order you would do it, is now a proper list under a fold line - and the fold line is not a label, it is the answer you came for, for example "**1 live · 1 paused · 1 ended · RM 20.00 given away**". The sixth step, printing, is the only one that cannot be undone, so it is the only row shaded, and it says **point of no return** in words rather than by colour alone. Two of the eleven have nothing to press yet - there is no test code for step 5, and the printed card for step 6 is the next piece of work - and a note under the list says so, rather than leaving you to look for a printer that is not there.

One fault found and fixed before you saw it. The Pause, Resume and End presses were first built against the row's own working copy of a code rather than the stored one, so the confirmation opened, the press looked like it had worked, and the code never actually changed. It was found by driving the real buttons rather than by reading the code, and every press is now measured landing on the stored code. It never reached your phone.

Three things worth knowing.

- **Nothing sets a code to printed yet.** This version builds the rule and the screen that respects it; the printed card itself - which is what freezes a code - is the next piece of work. Until then the rule protects a code that arrives frozen by a sync or by an import. **Pause and End never touch an order.** A code written onto an order when the shop sent it stays on that order and keeps showing what you owe. Ending stops the shop accepting the code; it does not re-price anything. **The shop is unchanged.** Every change in this version is in the backoffice. Nothing a customer sees moves, so there is no new upload and no database step.

02 Oct 2026 - engine v277, THE MONEY LINED UP (no database step, nothing to upload - pushing this one is the whole of it).

What you asked for. The day after v276 you opened the receipt and said: "the format still not as clear as a receipt, the money have to align up." Asked which one looked wrong, you named the **customer's tracking page**, the **order list** ("non at all in the order list"), and the **customer's WhatsApp message**. Asked what lining up meant, you picked **"Line up in one column"**.

Two screens changed. What you see now.

The customer's tracking page. The money is no longer one run-on sentence ("Items total: RM 30.00 · Courier charge: RM 8.00 · Total: RM 38.00"). Each figure is now on its own line, in the same right-hand column, with the Total in full ink under a dashed hairline:

```
Items total          RM 30.00
Courier charge        RM 8.00
- - - - - - - - - - - - - - -
Total                RM 38.00
```

The digits are tabular, so the decimal points sit on that line too. A code the order carried appears on the same column, and so does the reason when a code paid nothing.

The order list. Every order row in your Orders screen now ends with its own total, on a line of its own across the foot of the row, at the same right edge as every other row:

```
Focaccia   storefront                    ×2   TNG   Courier RM8.00 · customer
Mei Ling · 012-345 6789
Order total                                                    RM 38.00
```

Before this, a row said "×2" and stopped - the one screen you work from all day was the one place an order's money never reached.

Three things worth knowing.

- **It is one figure, worked out once.** The row, the receipt in the Edit / Note pop-ups and the customer's message all read the same sum. A row pricing an order at RM38 while the pop-up priced it at RM35 is exactly the sort of disagreement this change exists to prevent. **An order nobody has priced says nothing.** It does not show "RM 0.00" - an unpriced order is not an order worth nothing. **The WhatsApp message is deliberately unchanged.** You were shown both ways - the plain lines with a bold Total, and the same figures inside WhatsApp's grey code block (which lines the figures up but loses the bold) - and you chose the plain lines. So the message a customer receives today is byte-for-byte the message they received yesterday.

Also left alone on purpose: the + New order card's single "Items total:" line (one figure has nothing to line up against), the shop's own cart bar (one figure), the packing label (no money on it), and the Money screen and Profit statement (already in this same column).

02 Oct 2026 - engine v276, AN ORDER'S MONEY AS A RECEIPT (no database step, nothing to upload - pushing this one is the whole of it).

What you asked for. "build the next promo screen, What is the promo screen for? Can the showing of promo be more streight forward, clearer, like putting them in an accounting format, clearly shown the working, how they add up." So the promo showing was rebuilt the way the Profit statement already reads.

What you see now, on both screens that show an order's money (the Edit pop-up and the Note / tracking card), in place of the one run-on sentence that used to be there:

```
Items total                RM 16.00
Courier charge              RM 8.00
Promo FRESH10               RM 0.00
- - - - - - - - - - - - - - - - -
Total                      RM 24.00
```

and under it, in small grey type, why the code paid nothing: **"FRESH10 needs a basket of RM 100.00 - this one was RM 16.00, so the code gave nothing. The rule is a guide, not a gate: you can still take something**

off by hand."

Four things this changed.

- **A code that gave nothing is now a LINE ON THE RECEIPT, reading RM 0.00** - not left out. Before, an order that carried a code which paid nothing looked identical to an order with no code at all, so you could not tell from the money whether a customer had tried a code. Now you can. **The reason sits under the figures, in your own terms.** The basket the code wanted, the basket this order actually was, and then a plain statement that the rule is a **guide** - it is your bakery, and you can still take something off by hand. **The customer's WhatsApp message says the same thing in one quiet line**: Code FRESH10 not applied: basket below RM 100.00. Until now that message said nothing at all about the code, which you told me reads as a code that was **forgotten** rather than one that never applied. An order with no code is unchanged, line for line. **Nothing about a missed code touches the money.** The line explains; it never deducts. v275's fix stands exactly as it was, and the customer's Total is identical to what it was before this version.

What is deliberately NOT here. There is still no box on an order to record a discount you grant by hand, and the customer's own message states the code did not apply **before** you have decided whether to honour it. Both are flagged rather than guessed at.

Under the hood. The receipt rows and the customer's message lines are built by one function each in admin/js/courier.js (receiptRows, moneyLines), both reading a single parts object, so a screen and a message can never quote two different figures for one order. 13 new tests; the suite is 2,440 green. The customer-visible change is the one line in the message; everything else is inside your app.

02 Oct 2026 - engine v275, A DISCOUNT NOBODY EARNED (no database step, nothing to upload - pushing this one is the whole of it).

What you reported. "the arithmatic is not rigght, the fresh10 promo code discount 10 for order of 100, but my order only 16, it deduct 10 and customer have to pay 6 only. if thats the case bakery will broke." You were right, and it was the worst kind of wrong: nothing crashed, nothing looked odd, and the app quietly took money off an order that had never earned it.

What actually happened. A code has two separate things about it, and only one of them was being asked. The **shop's code box** has always asked the first one - "is this basket big enough for this code at all?" - and it refuses a code below its smallest basket, telling the customer how much more they need. But the **arithmetic that takes the money off** only ever asked the second one: "what is this offer worth on these goods?" That is a question about the offer, not about the sale. So once a code was on an order, the smallest basket was never looked at again - and a RM10 code sitting on a RM16 order took its full RM10 off, leaving RM6 to pay.

The fix. The smallest basket is now asked **wherever the money is worked out**, which is one place: it is the same single figure the four customer messages and your own screens all read, so they cannot disagree with each other. A code on an order whose basket never reached its smallest basket now gives **nothing**, and the order reads exactly as an order with no code at all - no line, no total change, no stray "RM 0.00". Two more things ride on the same rule:

- **The tally on your Promo codes screen.** What a code has "given away" is counted the same way, so an order that carried a code it never qualified for counts **RM0 given** and cannot eat a ceiling it never touched. It still counts as a **use** - the code did ride on the order, and that is what "used" means. **The shop's own line.** A code the page states but does not act on (one per customer, or a first order only - the two it can only guess at) no longer promises the discount when the basket is too small for it. The customer gets the line they can act on instead: "Add RM84 more to use it." The order still goes through and the code still rides on it, exactly as before.

What has NOT changed, and it is the important half. A code you have since **paused, ended or used up** still comes off the order it was actually placed on. Ending a code is a decision about the next order and never about one you have already promised, and this version does not touch that. The distinction is exact: the **dates** are a fact about the code's life and are never re-judged; the **basket** is a fact about that one order, and an order's own basket never changes either. A code that gives RM10 on RM100 gave nothing on a RM16 order on the day it was placed, and it gives nothing for it now.

What this is not. It is not a way to stop you giving a discount by hand. Nothing here blocks an order, hides a sale, or refuses a code you want to honour - it only stops the app doing arithmetic on a discount that was never earned. If you want a customer to have money off an order that did not qualify, you still take it off yourself, exactly as you always have.

The engine is **v275**; the automatic test suite is now **2,427** checks, all passing. What a customer can see: store/app.js. Everything else is admin/ - your own screens - plus the version number itself.

02 Oct 2026 - engine v274, THE SHOP'S OWN SENTENCE IN A HAND YOU PICKED (no database step, nothing to upload - pushing this one is the whole of it).

What you asked for. The handwriting on the shop page. In v273 you asked for a font on the messages, and the honest answer there was that WhatsApp carries no fonts at all - so only your opening line could lean over. The shop page is a different matter: it is our own page, so a real hand can be put on it and it will look the same on every phone. You picked which hand by looking at five real versions drawn inside the shop's own strip, and you chose the **chalk hand** - rounder and fatter, like chalk on a board.

What this version does. The amber strip at the top of the shop has two lines. The first is the offer and the code - 'RM10 off when you spend RM40, use code FRESH10'. That line is untouched: same lettering, same size, still first, because it is the part a customer has to act on. The **second line** - the words you wrote yourself about the code - is now set in the chalk hand, and a little larger than before, so it reads as your own note under the offer rather than a second announcement.

One limit, said plainly. The hand is a Latin face. It carries **no Chinese characters at all**, and no other hand we looked at does either. So a sentence you write in Chinese keeps exactly the plain lettering it has always had - the words are the same and nothing moves - it simply cannot wear the hand. English and Bahasa Malaysia wear it. This is said here rather than left for you to find.

Nothing is fetched from anywhere else. The font travels inside the shop, in the shop's own folder, with its licence beside it. So there is no request to a font company when a customer opens the page, and the letters are the same on an iPhone, an Android or a laptop. It is trimmed to the Latin letters only, which is why it is about 45 KB rather than the roughly 390 KB the whole family weighs.

Two corrections to the v273 entry below, which was written before you had picked. It says your shop's sentence is "still in the serif it has always used" - it was never a serif; it was the plain system lettering, and it still is under anything in Chinese. And it calls the hand you picked a "felt-tip" - you picked the **chalk** hand, which is what this version now sets. The rest of that entry stands.

The engine is **v274**; the automatic test suite is now **2,418** checks, all passing. What a customer can see: store/app.css, store/index.html and the new store/fonts/ folder. Everything else is the version number itself.

02 Oct 2026 - engine v273, YOUR MESSAGES IN A VOICE YOU CHOSE (no database step, nothing to upload - pushing this one is the whole of it).

What you asked for. "Can you make font selection for message from baker to be customisable, example italic", and then, when I showed you what is actually possible: "i like to make it more casual with the right font from the kitchen". Here is the honest answer, and then the half of it you can have today.

One limit, said plainly. A font **cannot** be chosen in a WhatsApp message. WhatsApp carries no fonts at all - the letters you send are always drawn by the phone of whoever is reading them, and there is nothing you or I can put in the message to change that. Its only four marks are bold, italics, strikethrough and monospace. So "a different font in the message" is not a thing that exists; the closest thing that does is **making your opening line lean over.**

What this version adds: Message style, on More -> Settings. A card with one choice, and it applies to **all four** messages you send a customer - the confirmation, the payment reminder, the "on its way" message and the "ready / pickup" message:

- **Plain - exactly what you send today.** The default, and it stays the default until you change it. Every message is word for word, letter for letter, the message that went out before this card existed. **The greeting**

leans over. Only the first line moves, and only into italics. *Hi Mei Ling! Your order is confirmed.* Everything below it - the order code, the items, the money and the bold Total - is untouched.

Why one switch for all four, and not one per message. Four separate switches are four ways for your messages to end up opening differently from each other, which is exactly the kind of thing a customer notices and you would never have meant. One choice, four messages, no way for them to disagree.

The bold Total is not affected. It already carries its own marks, and the ones this adds sit on the first line only - so the figure the customer pays is still the boldest thing in the message.

It travels to your other phone. The choice is saved with your settings, and a phone that has never been opened to this card takes your choice from the cloud rather than overwriting it. Choosing **Plain** again is a real answer too, and it travels as one - it does not go silent and let the old choice come back.

What is NOT in this version. The shop page's own sentence, under the today line, is still in the serif it has always used. You picked a **felt-tip hand** for it, and for that to look the same on every phone - not only on iPhones - the shop has to carry the font file itself. It is a free font and it is about 25 KB, and **only you can fetch it**: it is one download on your laptop, and I will walk you through it and show you the real sample before anything is built. That is the next version, and it waits on that file.

The engine is **v273**; the automatic test suite is now **2,412** checks, all passing. admin/ only - nothing a customer sees, and no SQL.

02 Oct 2026 - engine v272, THE CODE COMES OFF THE TOTAL (ONE SMALL DATABASE STEP FIRST - run it before you push, the order matters; nothing to upload otherwise).

The thing you reported. "pushed. The whatsapp message still withhout the promo discount." You were right again, and this time it was by design that had gone stale: since the code first arrived, it rode on the order and every discount was taken off by **you**, by hand, in WhatsApp. You have now said the app should do that arithmetic itself, and this is that.

The code now comes off the Total, in all four places the customer reads it - the confirmation, the payment reminder, the "on its way" message, and their own track page. It appears as its own line, named by the code, between the workings and the total, so the figure can still be added up by the person reading it:

Items total: RM 30.00 · Promo FRESH10: -RM 10.00 · Total: RM 20.00

And on your own screen. The two figures you read on an order move with it: the **"The customer owes"** line in the Note / tracking window, and the **Order total** in the Edit window. The code is named in the sentence there too, so you can see which figure came from where. The **"Still to collect"** figure on your Money screen follows, because it was always worked out from the same number.

What does not change. An order with no code reads word for word as it did before - every line, every blank line, every total. A code that has since been paused, ended or used up **still comes off the order it was actually placed on**: ending a code is a decision about future orders, and it never reaches back and re-prices one you have already promised. A code whose terms come to nothing on this order prints no line at all rather than *-RM 0.00*. A discount larger than the order leaves you asking for RM 0.00, never a negative amount. A percentage comes off the goods and not off the courier's charge, which is money passing through you rather than yours to give away. A free-delivery code waives the delivery charge - and on an order you are collecting yourself it has nothing to waive, so it says nothing.

THE DATABASE STEP - do this first, before you push. In your Supabase SQL editor, run the file **supabase/promo_track.sql** (Dashboard, SQL, New query, paste, Run). It adds two columns, promo_code and promo_rm, to the customer tracking table. **Order matters here.** Your app publishes a whole tracking row in one call, and if those columns do not exist yet, that call is refused as a whole and **the customer's tracking page stops updating for every order** - not only the orders that carried a code. This is the same trap the courier charge tables set, which is why they say the same thing. Run the SQL once, then push. It is safe to run twice.

One thing that has NOT changed, so it does not surprise you. Your sales and profit figures still count the **full price** of what you sold; the discount comes out of the money you actually collect and not out of the goods' own

value. So on a discounted order, the amount you collected and the amount your books call a sale will differ by the discount. That is a real question about where a discount should sit in your books, and it is yours to decide - nothing is hidden, and no figure you read is wrong, they are answering two different questions.

The engine is **v272**; the automatic test suite is now **2,401** checks, all passing. admin/, store/ and store-lang.js.

02 Oct 2026 - engine v271, A CODE THAT KNOWS WHEN TO STOP (no database step, nothing to upload - pushing this one is the whole of it, and your app must reach the cloud once for the shop to be told).

The thing a printed code cannot do. A card has no number on it and no end date on it - that was decided deliberately, because a date is a promise a card cannot keep and a count is a promise it may not be able to honour. Which left one real question: how does a code with a card stop being a good idea? Until today, it did not. Now you say so yourself, on the code screen.

The new box: "Stop after giving away (RM)". Fill it in and the code stops the moment it has given away that much money. It sits beside the order limit you already had, and **whichever runs out first is the end of it** - so you can say "the first 5 orders, or RM50, whichever comes first", which is usually what you actually mean. Leave it empty and the code has no limit at all, exactly as before. It works for every kind of code, including a percentage, which otherwise has no natural end.

Where the numbers come from matters, so here it is. They are **counted from your own orders**, freshly, every single time - not kept in a little tally on the code. That is why the figure cannot drift: the moment an order reaches your app, or you delete one, the count is right again on its own. Nothing the shop writes down is ever trusted, because the shop is a public page and cannot be.

A basket counts once. A customer who orders three things in one go used the code **once**, and it gave away one RM10 - not three. Your app counts the basket, not the items, so a three-item order can never eat three of your "first 5 orders". (A code you have since deleted stops being counted; the orders that carry it keep it and keep reading correctly, exactly as deleting always meant.)

Your screen now tells you which limit ran out. Each code shows how far through its limits it is - *3 of 9 orders used · RM30.00 of RM30.00 given away* - and a code that has finished says so in an amber banner across its card: **Fully claimed - it has given away its RM30.00. The shop has stopped offering it.**, or **...all 2 orders used...** when it was the count. That distinction is yours alone. The customer never sees it: they get the one sentence that already existed, "that code has been fully claimed", with no figure and no reason, because neither belongs on their screen.

What the shop does differently, without a single line of the shop changing. Now that a code carries real counts, the shop can finally act on its own "fully claimed" answer: a code that has given away everything you allowed stops being advertised at the top of the shop and stops being accepted at the box. This is the one change a customer could feel, and it only happens once your app has published to the cloud - so after you push, **open the app once on your phone with a connection** and the shop will be told.

One honest limitation, written down rather than glossed over. An order remembers the code's **name**, not the code's terms - so if you change what a code gives, every past order is recounted at the new terms. That is the figure on your screen, and it is the reason a printed code is meant to be frozen rather than left editable. Nothing about an order's own promise changes; it is only the tally that follows the code as it stands today.

Nothing to set up in the database, nothing to upload anywhere. admin/ only, and no customer-facing screen is redrawn - the shop simply acts on better information. The engine is **v271**; the automatic test suite is now **2,374** checks, all passing.

02 Oct 2026 - engine v270, THE CODE THAT WAS THERE ALL ALONG, AND A PROMO SCREEN WITH SOMETHING TO SAY (no database step, nothing to upload - pushing this one is the whole of it).

The thing you reported. "i dont see the promo code fresh10 send over to app together with the order." You were right, and thank you for telling me. The code had been travelling with the order since v269 - the shop stamps it on, your app saves it - but **no screen ever showed it**. The one place you look for an order was the one place it did not appear. A code that arrives and cannot be read is a code that did not arrive.

What you will see now. An order that was placed with a code wears a small amber **FRESH10** tag beside its order number, on every row you read an order on: the delivery day's list, the unread inbox at the top, the search results, and the **Edit order** window. It shows the code itself and not the word "promo", because you take the money off by hand - so which code it was is what tells you whether it is RM10 or RM5. If the code is lower-case in some old record, it is shown the one way it is always matched, so you never see two spellings of one code. An order placed without a code looks exactly as it always did.

The shop's line now says the whole offer. When you are running a public code, the line at the top of the shop used to name the offer and the code. It now also says the smallest basket it works on and, when you have set one, the day it runs out - for example *Today: RM10.00 off on RM30.00 and above, until 31 October - use code FRESH10*. A customer can see what the offer is worth before they fill a basket, instead of finding out at the box.

And you can add your own sentence to it. On the code screen there is now a box for your own words about the code, in English, Chinese and Bahasa Malaysia, with a **Translate** button that fills the two other languages for you and never overwrites anything you typed yourself. Whatever you write appears under the shop's own line, in the customer's language - a short line about the code, in your voice. Two things about it: it is **extra, never instead**, because the shop's own line is what tells the customer the code exists and what to type; and if you leave the Chinese or Malay box empty, the shop shows the English you wrote rather than showing nothing.

The code screen is now a full one. A code is no longer only "what it gives". Six things about a code are real, all of them optional and all of them defaulting to the widest, simplest answer: **who it is for** (anyone, or a first order), **when it runs** (from a day, until a day), **the smallest basket** it works on, **how often it can be used** (as often as they like, once per customer, or only the first few orders), **what it cannot be used with** (nothing in particular, or not together with the bring-a-friend credit), and **who can see it** (everybody, or just the person you gave it to). Leave any of them alone and the code behaves exactly as it did before.

The shop now says exactly why it will not take a code. Before this, the shop knew only two answers - the code is good, or it has never heard of it. It now has its own plain sentence for nine different situations: we do not know that code; that code is on hold; **that code ended on 30 September** (it names the day, so it does not read as a glitch); that code has not started yet, it starts on a day; that code has been fully claimed; not quite enough yet, add a ringgit figure more to use it; and the two the shop states rather than refuses.

About those last two. The shop is a public page with no sign-in, so it cannot know who the customer is. That means **once per customer** and **first order only** can never be more than a good guess, made from what that phone remembers - and a new phone or a cleared browser guesses wrong. So the shop does not use them to refuse anybody. It says the code is for a first order and that you will confirm it when you take the order, and then it takes the code anyway. Your rule stands: a rule on the website must never block or hide a sale you take by hand. The real check is yours, in your app, where the true history is.

One refusal that only the app can give. If you type a number of orders into a code and that number is zero or blank, the app now refuses it by name. Zero would otherwise be read as "no limit at all" and quietly make that code the one code with no limit - the opposite of what you asked for.

What did not change. The total on the shop page still never moves - a code still does not change what a customer is charged, and you still take every discount off by hand in WhatsApp. An order placed without a code is posted, sent and shown exactly as it has always been: the message such an order sends is byte for byte the message it has always sent. No existing order, product, price, customer or setting was touched, there is no database step, and nothing needs uploading.

02 Oct 2026 - engine v269, PROMO CODES, THE FIRST SLICE (no database step, nothing to upload - pushing this one is the whole of it).

What is new for you. A new screen, **Promo codes**, at the bottom of the **More** menu. You make a code there: the code itself, what it gives, and whether the shop may show it to everybody or you are giving it to one person. Three kinds of offer are in this first version - a ringgit amount off, a percentage off, or free delivery.

What is new for your customers. The shop page has a **Have a code?** box. A customer types the code and presses **Use it**. If it is a code you made, the page says what it gives and tells them you will take it off when you

confirm the order. If it is not, the page says it does not know that code. And when you are running a code the shop may show, the top of the page carries a line naming it - for example *Today: RM10.00 off, use code FRESH10* - so a customer who was told about the code knows exactly what to type.

The most important thing about it. The total on the shop page NEVER moves. A code does not change what the customer is charged or what the page adds up. The customer types the code, the code arrives on the order, and you take the money off by hand in WhatsApp exactly as you already do for the bring-a-friend credit. That is deliberate: the code is a note that travels with the order, not a machine that edits your prices.

What comes with the order. An order that used a code carries that code into your app, so you can see at a glance which order had which code on it. An order that used no code, or one the page refused, is posted exactly as orders have always been posted - nothing is added to it. A code you made on one phone also reaches your other phone, so a card you hand out works whichever phone is in your hand.

One thing worth knowing about "personal" codes. A code marked personal is never shown on the shop page and never advertised. It still works when it is typed, which is the whole point of giving it to one person. But the shop's own data is readable by anyone who looks at the page, so personal means *never shown* and not *secret*. The real limits on a code will be the ceiling and your own hand-marking, which arrive in the next slices.

What is not in this version yet. This first slice is one code working all the way through. The finer rules - when a code starts and ends, the smallest basket it works on, how many times it may be used, who it is for, and the ringgit ceiling - are in the next slices, along with pausing, ending and printing the card. None of them is needed for a code to work today.

What did not change. Not one existing order, product, price, customer or setting was touched, and the message an order without a code sends is byte for byte the message it has always sent.

02 Oct 2026 - three of the reference documents were losing the end of their lines (no new engine; your phones did not change). Three of the printed references - the **Courier APIs** brief, the **money and profit guide** and the **ntfy and Resend** page - had lines that ran off the right-hand edge of the page and were cut off there. Worst of the three was the courier brief's own title, which printed 9.5pt past the edge of the paper; the money guide lost the end of two of its money cells, and one line inside an ntfy and Resend box lost its tail as well.

Nothing was wrong with what the pages said. Every one of those words was already in the file and is in the file now - the fault was in how the page was laid out. A long line was being drawn as a single strip that never wrapped, so whatever passed the right-hand margin was simply never printed, and a sentence you were reading ended early with nothing to show that anything was missing. Long lines now wrap onto the next line, which is what they should always have done.

What did not change. Not one word of any of the three documents was rewritten, reordered or dropped - this was a repair to the page layout only. Nothing you press in the app changed, no setting moved, and there is no database step. Each document still runs to the same number of pages as before (7, 7 and 4), and every page of all three was read back and measured to confirm that no line now reaches past the margin.

01 Oct 2026 - engine v268, FIVE PLACES THAT DISAGREED WITH EACH OTHER, AND THE BUTTON THAT CLAIMED TOO MUCH (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "pushed. when i schange the courier delivery to self pickup, the courier chages tag still there, just wondering are details well taken care of?" - and then the audit you asked for: "pls run whole system audit for the well being of the system later".

What was happening. v267 fixed the tag you could see. Reading the whole system after it turned up four more places where one screen had been told the truth about a self-collect order and another had not - and one button that made a claim it could not know. They are all the same fault underneath: **a fact is written on the order, and only some of the screens that read it ask the companion question that goes with it.**

Five things, fixed together.

1. The customer's track card was still being sent the courier's half of a self-collect order. A driver's name,

a plate, a phone number, "on the way", and the waybill number - on an order the customer is coming to fetch. Switching to **Self collect** keeps the booked trip, so the card was reading it. It now publishes none of it: the card says **Self collect** and nothing about a courier.

2. The trip card counted a parked charge. The card that shows what a courier trip is costing you against what the customer pays for it added up the charge on **every** order in the trip - including a charge parked on one you have since switched to self collect. It now counts only the orders actually going by courier, and says plainly when there is no charge on the order, so the whole trip reads as your own cost rather than as money nobody is paying.

3. "Paid by the customer" recorded a name, not a payment. Picking **Cash** or **TNG transfer** on the order wrote down how they paid and stopped there - so the row could wear a **TNG** tag while the day's till, the Money screen and the customer's card all still counted that order as owing. It now records the payment itself, the same three things the **Paid · Cash** / **Paid · TNG** buttons write, and **Not recorded** is the way back - the one control that can undo a payment marked by mistake, since those two buttons disappear once the money is in. It also says so, under the box, in words.

4. Print label vanished the moment the order was packed. It was offered on **Baked** and nowhere else, so a label that tore while the bag was being filled could not be printed again on an order still in your kitchen. It is now on every stage from **Baked** to the end, and still on none before it - there is no bag to kit yet.

5. Send confirmation turned the order green before the message was sent. Pressing it opens WhatsApp. It does not press Send in WhatsApp - and the Confirmed step went green anyway, on your row and on the customer's own card, for a message still sitting unsent in the box. You chose the honest version: the button now only drafts the message, and a small **I have sent it** appears beside it. **That** is what turns Confirmed green, on your screen and on theirs. One extra tap on each confirmation, in exchange for never marking something sent that was not. Moving the order on and back onto Confirmed starts that step again, so a second confirmation is never blocked.

Two pieces of paper that were out of date. The guide's step for sending a confirmation now describes the new press, and the page about the Google setup no longer says an API is still to be enabled - it has been live since 28 September, and the page now reads as a record rather than a to-do list.

What did not change. The charge box on an order that is going by courier, the parked-charge notes v267 added, the words on any message, and everything about how you take an order. Nothing here needs a key, a wallet or a database step.

01 Oct 2026 - engine v267, A COURIER CHARGE GOES QUIET WHILE THE ORDER IS A SELF COLLECT (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "when i schange the courier delivery to self pickup, the courier chages tag still there, just wondering are details well taken care of?" - and then, a moment later, "confirmation message still include courier charges".

What was happening. You were right on both counts, and they were the same fault. Switching an order from **Courier delivery** to **Self collect** never deleted the courier charge, on purpose - you might switch back, and an app should not ask you to type a fee over again. But nothing on the reading side ever asked whether the order was still going by courier. So a self-collect order went on wearing a **Courier RM 8.00 · customer** tag on its row, went on adding that RM8 into the customer's total in the **confirmation message you send**, in **every WhatsApp message** and on the customer's **track card**, and went on being counted in **still to collect** on the Money screen.

What changed on screen. An order that is a **Self collect** now carries nothing for a courier to charge for: no tag on the row, and no courier charge in the customer's total anywhere - the confirmation, the messages, the track card and the Money screen all read the items and nothing else.

The charge is parked, not thrown away. It is still written on the order, so switching **Fulfillment** back to **Courier delivery** brings the whole thing back exactly as it was, tag and all, with nothing to type again. Parking is not hiding either: both courier charge cards - the one under **Edit** and the one under **Note / tracking** - say in words that a charge of RM 8.00 is recorded on this order from when it was a courier delivery, that it is not added to the customer's total and not on any message while the order is a self collect, and that **switch Fulfillment back to**

Courier delivery is how it goes back to work. The charge box stays right there under that sentence, so clearing it is still the same gesture it always was - set **who paid the courier** back to **Not recorded** and save.

One thing that is deliberately left standing. If the charge was one **you** paid, the **Delivery & fuel** row it put on your books is still there while the order is a self collect. That is your cost, not the customer's, and nothing about a fulfilment switch makes it untrue - so the parked note says so, and names the Money screen as the place to take it out if the courier was never actually paid.

What did not change. The charge box on an order that is going by courier, the words it says, the COD tick, and the guard that refuses a save when there is an amount in the box with nobody down as the payer. That guard now asks on a self-collect order too, because it can only ever speak when the box is on screen with an amount in it - so it never refuses a save over something you cannot see, and no save can answer "Order updated" for a write it did not make.

01 Oct 2026 - engine v266, A BUTTON NOW SAYS WHAT HAPPENS IF YOU PRESS IT - AND WHAT HAPPENS IF YOU DON'T (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "when i see the button, i might self have to ask, i dont know what will happen or what will happen if i din press that button, these create confusion"

What was happening. Two places in the app put a choice in front of you without saying what the choice cost. On the **shop page**, the line above the pin buttons read "Optional: drop a pin where the courier should stop" - which tells a customer the pin does not matter, without telling them what they get if they do it or lose if they skip it. And in **your own pin window**, the **Coordinates, if you have them** box with its white **Use these numbers** button stood open under every map, on the one window in the app whose whole job is to be read fast at a door.

What changed on screen.

1. The shop's pin line now answers both halves. It reads: skip this and the driver goes to the address you typed - that is fine for most houses; pin it only if the address alone will not find your door, a condo block or a guard house; and either way put the block and unit number in the address above. A customer who does nothing now knows exactly what they are getting. Written in all three languages.

2. The coordinates box in your pin window comes out only when it is the answer. When the map cannot load, the box appears by itself, right under the sentence that names it - so "the map is not available right now, type the coordinates below instead" is finally pointing at something that is really there. On every other day the box and its white button are out of the way, and the window is three lines shorter.

3. It waits behind one press, not gone. On a working map a small button reads **Have a Google Maps link?** - press it and the box comes out. It is there because a Google Maps link is the most accurate point a customer ever sends you, and hiding the box outright would leave a working day as the one day you had nowhere to put one.

4. The box's own sentence says what the press does. It reads: press this and the pin moves to those numbers, a Google Maps link works too, and nothing is kept until you press **Use this spot** below. So the white button is now a button you can predict before you press it.

The limit, said plainly. The box appears when the map itself fails to load. A map whose tiles fail but whose map still works is not detected - the box stays behind its press there - and on that screen the map is still usable enough to drag a pin or tap an empty one, which is why it is not treated as a failure.

01 Oct 2026 - engine v265, THE CONSIGNMENT BOX IS A PARCEL'S, SO IT IS DRAWN WITH THE PARCEL (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "after i get price from lalamove, click use this fee, it closes the window, but user might confuse as at that page there is stick fields that user have not fill in, maybe that field or button should not be there? this create confuse"

What was happening. On the **+ New order** card the courier half drew the **Courier tracking number (optional)** box on every courier order, whatever kind of courier it was. A consignment number is the thing a carrier gives you for a **parcel**, and a Lalamove trip is not a parcel - so on a trip that box was empty, it was not about anything

on the card, and it stood beside a delivery price that was already finished. That is what read as an order with something still to fill in.

What changed on screen. The consignment box now belongs to the parcel, so it is drawn with the parcel and nowhere else. Take a Lalamove price and the card under it is the address, the pin, the courier charge, the price and the delivery note - no empty box waiting to be filled in. Name a parcel carrier and the consignment box appears **under that carrier**, which is where a consignment number comes from in the first place.

The order of the courier half on the card is now the address with its pin, then the parcel carrier, then its consignment box, then the courier charge, then the price. Only the last two positions are new: the parcel used to sit under the tracking box, and the tracking box used to sit on its own, above everything to do with a parcel.

Nothing became unreachable. The parcel carrier picker is still on the card, in the same place, saying **Not a parcel - nothing recorded** until you pick one - it is the press that brings the consignment box out. A number you had already typed is never hidden from you either: a box that holds something is always drawn, so the card can never keep a consignment number you cannot see or edit.

A trip's own number is untouched. When you book a real Lalamove trip from the order, the live share link is written into that same slot by the booking - the box is a parcel's, and the slot is still one slot, exactly as it was.

Swapping carrier does not move you. Because the consignment box comes and goes with the carrier list, the card now has to decide, each time you pick a carrier, whether the box should be standing. It only touches it when the answer actually changes - so if you are halfway through typing a consignment number and you change your mind about who is carrying it, the box you are typing in stays put and stays focused, rather than being thrown away and drawn again under your finger.

The Edit window and the Note / tracking box keep the layout they had. Their consignment box still stands above their parcel carrier, and the hint there still says so ("the consignment number goes in the tracking box above"). The two screens were allowed to differ before and they are allowed to differ now; only the + New order card changed.

What did NOT change. The delivery note, the courier charge and who bore it, the delivery address with its pin, **Get a delivery price** and **Use this fee**, the parcel carrier list, the hand-over press, and the rule that nothing at all is written until you press **Place Order**.

Nothing to run in Supabase, and one push. This is app code only, exactly like v264.

01 Oct 2026 - engine v264, A TAP ON THE MAP DOES NOT MOVE THE PIN (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "click on the map should not move the pin, only dragging the pin will"

What changed on screen. A tap on the map no longer moves the pin. Only dragging the pin itself moves it. This is true of both maps you use to place a pin: the pin map in the order form (the one behind **Move this pin** and **Done moving**), and the pin window that opens when you set a doorstep or the bakery's pickup pin.

Why it matters on a phone. That map is a 200-pixel strip inside a card you scroll. Before this, any accidental touch on it moved the pin to wherever your finger landed **and** pulled the map onto that spot - two things at once, from one touch you did not mean. The pin window even said so in its own words all along: *"Look the address up, then drag the pin to the exact door."* Dragging is what it promised, and dragging is now the only thing that moves it.

The one place a tap still works, and why it is not an exception. A map with **nothing on it yet** has no pin to move and nothing to drag - so there, a tap places the first pin. That is the state the pin window opens in when an order has no point at all, and taking that ability away would leave you with a map that does nothing. Once a pin exists, a tap does nothing at all.

What did NOT change. Looking the address up, the list of other matches, pasting a Google Maps link or a pair of numbers, pressing **Use this spot**, and the **Move this pin** button all work exactly as before. So does zooming - the plus and minus, the two-finger pinch and the double-tap.

What was deliberately left alone. The shop's own map, the one your customer uses at checkout, is untouched. For a customer the map is often the only way to place their pin at all (they have no coordinate box and no list of matches), and a map a customer can only aim at by dragging a small marker with a thumb is the one thing that would make them give up. If you want that map changed too, say so - it is a one-line change.

01 Oct 2026 - engine v263, A DAY YOU DO NOT DELIVER IS DARKER STILL (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "can make the quite day more solid?" - asked straight after seeing v262.

What changed on screen. A day you do not deliver is now a solid, dark day. v262 had already taken the fade off it, but it was still faint - **3.50:1** against the card. It is now **6.71:1**, very nearly twice as strong again. The colour is exactly **halfway between the app's quiet grey and its ink**, so it reads as a real day rather than a whisper, while still looking nothing like a day you deliver.

The three steps are further apart now, in the same order. The past is still the faintest thing on the calendar, a day you do not deliver is still the middle step, and a day you deliver is still the darkest. The middle step has simply moved a long way down towards the bottom. A day you deliver is still **twice as dark** as a day you do not - **13.45:1 against 6.71:1** - so the two can never be mistaken for each other.

What did NOT change. The past grey itself is untouched: still the shop's own faint grey, exactly as you asked for it to stay. A day you deliver is exactly as dark as it was. So are the **green pill** on a delivery day, the **green tint** on a day a product sells, the **brown** day you are looking at, and today's halo.

One thing worth knowing about, deliberately left alone. On a product's own availability calendar there is a second, rarer kind of quiet: a day you have **marked as selling** but which the bakery does not deliver. That one is kept faded on purpose - a green tint turned right down - because it means "this mark cannot become an order", which is something to notice rather than something to read easily. It is untouched by this version. Say the word if you would rather it read stronger too.

Why 6.71 and not darker. The same rule draws the product availability calendar, where it is the only thing separating **a day you can sell on** from **a day you cannot**. Pushed much further, those two would stop being different pictures. Measured on that calendar: a day you can sell still reads **13.45:1** against a day you cannot at **6.71:1**, which is a clear step and not a blur.

Nothing else moved. No number, price, booking count or order changed. The database was not touched, and there is nothing to upload - this is the app's own picture of a calendar.

01 Oct 2026 - engine v262, A DAY YOU DO NOT DELIVER IS NOW A SOLID DAY, AND THE PAST IS REALLY ONE GREY THIS TIME (no database step, nothing to upload - pushing this one is the whole of it).

What you said. "I still feel that the greyed and the non grey contrast is not big?" - and then, when offered the choice, **make the non-grey more solid.**

You were right, and v261 had not finished its own job. v261 put the shop's grey in the right place, but it left the older **half-transparent** rule standing on a day you do not deliver. That half-transparency is applied to whatever colour the day ends up with, so the past was still being drawn in **two shades** - and worse, a past day you do not deliver came out **fainter than a future day you do not deliver**, which is backwards. Measured on the real Orders calendar: a quiet day sat at **1.98:1** against the card and a past quiet day at **1.22:1**, so the two were almost the same picture.

What changed on screen. A day you do not deliver is no longer faded - it is a solid, plainly readable day that is simply quieter than a day you do deliver. On the real grid that took it from **1.98:1 to 3.50:1**, very nearly twice as strong. And because nothing fades the calendar cell any more, **the past is now genuinely one shade**: a past day you deliver and a past day you do not are the identical colour on screen, which is what v261 was meant to give you. Three clear steps now, everywhere: the past faintest, a day you do not deliver solid in the middle, a day you deliver darkest of all.

What did NOT change. The shop's past grey itself is untouched - you said you liked it, so it stays exactly as

faint as the shop's is. Anything on a past day that is telling you something still keeps its own colour, so on a day that has gone the **green pill** on a delivery day stays green and so does the **green tint** on a day a product sells. The day you are looking at still turns brown, on every calendar. Today's halo is untouched.

Why this is the right lever and not the other one. Between making the past darker and making everything else stronger, the second is the one that keeps the shop's calm look you asked for while still giving the past something solid to stand against.

Nothing else moved. No number, price, booking count or order changed. The database was not touched, and there is nothing to upload - this is the app's own pictures of its calendars.

01 Oct 2026 - engine v261, EVERY CALENDAR NOW GREYS OUT THE PAST THE SAME WAY YOUR SHOP DOES (no database step, nothing to upload - pushing this one is the whole of it).

What you asked. "Is all calendar passed day grey out same thru all calendar? I like the one at store."

It was not. On the Orders calendar a day that had gone by came out in **two different greys side by side in the same row** - one shade for the days you deliver, and a slightly different one for the days you do not. On the other calendars there was only one grey, but it was the app's own faded grey, not the shop's. The shop's is the one you picked.

What changed on screen. A day already gone is now drawn in the shop's own quiet grey, **flat and one shade, whether or not you deliver that day.** That covers the Orders calendar at the top of the Orders screen, the calendar inside **Delivery Dates**, the small calendar you open when you pick a delivery day, and the days you mark on a product's availability - they all now use the one colour, so a past day looks the same wherever you meet it. Nothing ahead of today is touched - a day still to come looks exactly as it did.

What deliberately did NOT go grey. Anything on a past day that is telling you something keeps its own colour, because greying it out would hide a fact you are looking at. So on a day that has gone by, the **green pill** on one of your delivery days stays green, and so does the **green tint** on a day a product sells. Greying those would have made a past delivery day look like a day you do not deliver, which is the opposite of useful.

The day you are looking at still turns brown. That has not changed anywhere - so on the Orders calendar you can still tell at a glance which day is open in front of you, and it stays plainly different from the days that have simply gone by.

Why it is safe to be this quiet. The store calendar's grey is deliberately faint but legible, so a day that has gone still reads as a day rather than a hole in the week. That reasoning is the shop's, and it now holds here too.

Nothing else moved. No number, price, booking count or order changed. The database was not touched, and there is nothing to upload - this is the app's own pictures of its calendars.

01 Oct 2026 - engine v260, YOUR MAILS NOW COME FROM YOUR BAKERY, NOT FROM THE APP (two functions must be uploaded - pushing alone does not change this one).

What a customer saw. The two mails this system sends - the wish-list mail you send out, and a customer's suggestion coming back to you from the shop page - each arrived with a person's name at the top of the inbox list. That name read **"BakeAdmin wishes"**. **BakeAdmin is the app's own internal name.** It is a word that appears nowhere on your homepage, nowhere in your shop, and on no label you print. It was the first thing a customer read from you, before they opened anything.

What they see now. Jien Luv 2 Bake - your own name, taken from your homepage title, which is where your brand is written down. The address the mail is actually sent from has not moved: it is the same verified address as before, and only the name in front of it changed.

One thing to check, because it can silently win. If the project has a saved setting called RESEND_FROM, that setting is used **instead of** the name in the code, and this change would make no difference to what a customer sees. Run supabase secrets list in **Terminal** to see whether it is set. If it is there, tell me, and I will give you the one command to set it to your name. If it is not listed, nothing further is needed.

Why pushing this one is not enough. These two mails are sent by the two functions that live in Supabase, not

by the app on your phone. Uploading the code - a push - leaves Supabase running the old copy. **Both functions must be uploaded separately**, from the repository folder, in **Terminal**: wish-mail and shop-feedback. The changelog cannot do that part for you; the commands are with the version notes.

No database step. Nothing was added or changed in the database, and no existing setting moved.

01 Oct 2026 - engine v259, THE EDIT CARD STOPS CLOSING ITSELF WHEN YOU LOOK AN ADDRESS UP AGAIN (no database step - pushing this one is the whole of it).

Your words were "when i edit order, go in and press look this address up again, it exit the page and go into the order list and without allowing me to get a price from lalamove". That is exactly what it did - and it was not the fault the four reports before it sent me chasing.

The whole of it, in one sentence. The Edit card holds two things that can save: the **Save changes** button at the bottom, and the pin work, which saves by itself the moment it finds a door - because a real vehicle on a real road must never be thrown away by closing a card. The line doing the pin's saving was written to reach for the app's own save - and both of them were called save, so it reached for the **Save changes** button instead. Pressing **Look this address up again** therefore found the address, wrote it, and then pressed Save for you: your order was saved, the card closed, and you were dropped back on the Orders list.

You can still see it in what the app said at the time - "Order updated". That is the Save button's own message. The pin's press never had a message of its own to give you, which is how one word gave the whole thing away.

What happens now. The press finds the address, writes it onto the order, says what it found on the card's own line, and **leaves the card open and exactly where it was, with the pin moved. Save changes** is still yours to press when you are ready - and **Get a delivery price** is still on the card to press, which is the thing you could not get to before.

Nothing else about the Edit card changes. The delivery day, the customer, the items, the prices, the charge and the parcel are all exactly as they were.

There is no database step and nothing to upload - pushing this one is the whole of it.

01 Oct 2026 - engine v258, THE APP NOW TELLS YOU WHEN YOUR PHONE IS RUNNING AN OLD BUILD (no database step - pushing this one is the whole of it).

Your words were "still the same problem engin no 257", which was the fifth time you reported the same press - and this time it sent me looking somewhere else. The v257 fix is on your website; I checked that from outside, and the file your phone downloads does contain it. So the fix was not missing. Something was stopping it reaching your phone.

What was happening. Your website hands out every file of the app with a ten-minute cache, and there is no build number written into any of the addresses. So a phone that had the app open - or that opened it within those ten minutes - was given the copy it already had: the old one. Nothing on the screen said so. And the **Engine v257** pill on the More screen could not say so either, because that pill reads a small file of its own rather than the one carrying the fix. It would have read v257 perfectly happily on a phone still running v254.

Two changes, so that this cannot be quiet again.

The app's own files are now confirmed with the website every time it opens. The part of the app that runs in the background used to ask the network for each file - but "the network" was allowed to answer out of those ten minutes of cache. It now asks the website to confirm that each file is still the current one before using its own copy. An unchanged file costs one very small check rather than a re-download, so opening the app does not get slower; a changed file is fetched whole. There is nothing here for you to press.

And a strip now says so when your phone is behind. Where the site is serving a newer build than the one on your phone, an amber strip appears just above the bottom menu - **New version ready** - this phone has v257, the site has v258 - with an **Update now** button beside it. Press it when you are ready and the app reloads on the new build. Its words say "save what you are writing, then update" for a reason: a half-written order does not survive a reload.

If pressing Update now does not take - which can happen inside those first ten minutes - the strip does not go on offering a button that did nothing. It changes its words, says the update still has not taken, and tells you to close the app completely and open it again.

What this means for you, from now on. When you report something, look at the bottom of the screen first. If that strip is not there, your phone is on the newest build and the fault is in the code. If it is there, press Update now and try again before you tell me. Either way you are no longer guessing, and neither am I - which is worth more than the fix itself.

Also put right while here: the change history had v256 printed above v257, so the two newest entries read in the wrong order. They are back the right way round.

30 Sep 2026 - engine v257, TAPPING "RESET THE PIN" NO LONGER FOLDS THE ORDER YOU ARE WRITING (no database step - pushing this one is the whole of it).

This is the one you have been reporting. You press **Look this address up again** on the order's own door block, the app asks you to confirm because there is already a pin to replace, you tap **Reset the pin** - and you land back on the Orders list with the half-written order gone. The look-up itself then carries on and writes the order.

The cause was the rule that folds the + New order card when you press anywhere else. That rule is right for a press on the page. It was wrong here, because the confirmation the card itself opened is drawn on a layer that sits **beside** the page rather than inside it - so the app counted your tap on **Reset the pin** as a tap on the page behind, folded the card you were writing, and then did the look-up anyway. A press on the confirmation, on a pop-up, or on the lock screen no longer folds the card. A press on the page still does, exactly as before.

Why the last two versions missed it, in plain words. v255 and v256 both changed the pin window's **Look it up**; your words named the button on the order's own door block, which is a different button on a different window. And the check that cleared that door block measured where each row sat on the screen - it never asked whether the card was still open, and it never pressed the case where the confirmation appears, which only happens when there is already a pin to replace.

30 Sep 2026 - engine v256, THE MATCHES IN THE PIN WINDOW NO LONGER PUSH THE CARD AT ALL (no database step - pushing this one is the whole of it).

Your words were "it is still the same, press look up this address again make it exit the page we are working in, why?", and the honest answer is that v255 made it worse. v255 tried to hold the map still, and while the list of matches takes up room on the card, holding the map still can only be done by scrolling the whole window by the height of that list - 269 pixels on a phone's width. That scroll carried the address box and the **Look it up** button you were pressing clean off the top of the card: the button you were holding went from 248 pixels down the screen to 21 pixels **above** the card's own top edge. Instead of the map jumping, the window jumped.

The list of matches now floats over the map rather than taking up room on the card. It drops from under the button you pressed, sits on top of the map, and is capped at about four rows with its own scroll. Because it takes up no space, nothing above it can be pushed anywhere - not the button, not the address box, not the window's own scroll.

And the list puts itself away the moment you choose from it. That is the one thing that had to change for it to float: a list left sitting on top of the map would hide the very thing you want next, which is the map, to check the pin landed on the right door. Pick a match and the list closes, the pin moves to it, and the map is there.

Measured on a phone's width with four matches: the button moved 0 pixels and the window scrolled 0 pixels. The map slides down by the height of the one line that appears under the button (38 pixels) and nothing else moves - including, now, the button itself. It holds the same at two matches as at four, when the look-up finds only one place, and when the look-up fails.

Separately measured, because your words name that button too: "Look this address up again" never moved anything. That is the same press on the order's own door block, and it was measured on both kinds of

window - the + New order card and the Edit order pop-up - and on both kinds of answer, a look-up that finds your house and one that fails. In every case the row stayed on the same line and the card stayed open. The window that misbehaved was the pin window's **Look it up**, and that is the one this version fixes.

What did not change: the first match still lands on the map by itself, the list is still only how you say "not that one", the line under the button still says when the answer is the road rather than your house, and no wording, price or order of questions is touched.

30 Sep 2026 - engine v255, LOOKING AN ADDRESS UP NO LONGER THROWS THE PIN CARD AROUND (no database step - pushing this one is the whole of it).

Your question was "when i say look this address up, why the interface jump out of the page?", and the answer was in the shape of that card. Pressing **Look it up** in the pin window asks the address service for other matches, and the list of them is drawn *under the button*, above the map. So the list appearing shoved everything beneath it down the window: on a phone's width the map was thrown 269 pixels down the card, and the **Use this spot** button went with it. The button you had just pressed did not move at all, which is exactly why it read as the page jumping rather than the button.

The map now stays where it is while the matches appear. The window is held by the map - the thing you are checking the pin against - so when the list arrives the window scrolls to keep the map on the same line of your screen, and the matches are revealed above it, where you are already looking. Measured on a phone's width, the map and the **Use this spot** button both stayed still to the pixel, and three presses in a row left them still. The same holds when the look-up finds only one place, and when you tap one of the matches.

This is the same rule that stopped the + New order card jumping at v254, and it now lives in one place. Both screens use one copy of it, so any later screen that needs the same protection inherits something already tested rather than a second hand-written version. Nothing else on either screen moved: the wording, the prices, the order things are asked in and the address look-up itself are all unchanged.

30 Sep 2026 - engine v254, THE + NEW ORDER CARD HOLDS STILL UNDER YOUR THUMB, AND ITS BIG BUTTON NOW SAYS PLACE ORDER (no database step - pushing this one is the whole of it).

The button inside the card says + Place Order. It used to say + Add order, and that is the only thing that changed: the card's own title is still **+ New order**, the + New order button on the Home screen is untouched, and the Edit order pop-up's **Save changes** is untouched. The two sentences that send you to the button now say "then press Place Order", so they no longer point at a button that is not there.

The card no longer jumps out from under your finger when you reset the pin. Your words were "once i click reset pin the screen jump". Resetting the pin makes the card change its own height in up to eleven ways at once - the answer line appears, the words re-wrap, the map is shown or hidden, the price list is rebuilt - and nothing was putting the page back where you were reading. The card is now held by the row of buttons under your thumb: whatever the block does to its own height, that row stays on the same line of your screen. Measured at a phone's width, the row sat still to the pixel while the page scrolled 15 pixels to absorb the change. The same press used to push everything below it by 210 pixels when the press was the one that made a map appear.

A map no longer keeps running behind a card you have moved on from. Every time the card is rebuilt - you change Fulfilment, you tap a day in the card's calendar, a pop-up refreshes, a sync pull arrives - the old map was left alive, holding its tiles and its screen-size listener until some later resize happened to notice. On a phone that notice is the keyboard opening, which is exactly what happens while you are typing in this form, so the maps piled up. Each new map now clears away the ones whose box has left the page, at the moment it is built. Nothing on screen changes.

A map can no longer paint over a question you are being asked. You put this as "sometimes pop up like half at back layer", and it was real. With a live map in the card, the map was drawn over the bottom half of the confirm dialog: the sentence cut off mid-sentence, both buttons hidden, and a tap in that band landing on the map's + / - zoom instead of the button. The map's own layers climb to 1000, and the box it sat in was not containing them, so they were competing with every layer this app draws. The map is now sealed inside its own box, in all four places the app draws one.

Nothing else moved. The door block's wording, the address look-up, the prices and the order the card asks things in are all unchanged.

30 Sep 2026 - engine v253, THE SHOP'S FEEDBACK MAIL NOW READS IN THE SAME ORDER AS YOUR WISH-LIST MAIL (no database step, but the shop-feedback function must be redeployed).

The heading block moved to the top, and the customer's words now come last. Engine v252 put the words first and the heading underneath a short rule. You asked for the wish-list mail to be the reference, and the wish-list mail does the opposite: it opens with what the mail is, then who it is about and when, and puts the content underneath. So the feedback mail now does exactly that, with no rule in between and no closing line:

New feedback for the shop page (Engine v253).

Project: jienluv2bake.com.my/store/

Sent: 2026-09-30 17:31

Written in English.

the words the customer typed

The language is written out as a word. Engine v252 printed the short code it was handed, so the line read Written in en. The shop asks its questions in three languages, so the line is worth having, but "en" is not a word anybody reads. It now says **Written in English.**, **Written in Malay.** or **Written in Chinese.**, and it is left out entirely when the page did not say which language the customer was reading.

The subject line and the Project line are unchanged from v252: the subject is still Shop feedback · Engine v<n> · <date>, the project is still the live address the customer was reading, and the time is still on your own clock in Penang. Only the order of the body and the wording of the language line have moved.

One step on your machine. The part of this that builds the email lives on Supabase, not on GitHub, so pushing is not enough on its own. Engine v252's version of it was never actually running there - the mail you received after v252 was still the build from before, which is why it carried none of the new heading. After pushing this version, redeploy it once:

supabase functions deploy shop-feedback --project-ref hzpyblqygnntixkijeem

30 Sep 2026 - engine v252, THE SHOP'S FEEDBACK MAIL SAYS WHICH SHOP AND WHICH BUILD (no database step, but the shop-feedback function must be redeployed).

A customer's message now arrives with the same kind of heading your wish-list mail has. Before this version the email was the customer's words and nothing else, so a message saying "the cart is confusing" could have come from either of your shops and from any build of the page. Now every one of these emails carries, under the words and a short rule:

- **Project** - the address the customer was actually reading, for example jienluv2bake.com.my/store. It is read off the live page rather than typed in anywhere, so your two shops can never be confused for one another, and the second half names which page of that shop the words were written on. **Sent** - the date and time on **your** clock, Penang time, so "16:52" means the same thing on the email, on your phone and in the backoffice. It follows the same shape as the wish-mail mail: 2026-09-30 16:52. **Engine** - the build the customer was looking at, both in the subject line and in the line above the project, so you can tell whether a complaint about the shop came from a phone that is still running yesterday's copy. **Written in** - the language, when the customer was reading the shop in Malay or Chinese.

The subject line names the build too, so the mailbox listing alone tells you what you need: Shop feedback · Engine v252 · 2026-09-30. Everything else about the mail is unchanged - the customer's words still come first, the heading still sits below them, and the mail still goes to the developer addresses you set in Settings.

And the shop itself now says which build it is running. At the very foot of the shop page, under "Website by" and the developer's WhatsApp and email links, there is a new small line:

Engine v252

It is the same number your app shows on More, and the two are read from one file rather than kept as two

copies that could drift apart. That is the point of it: when a customer tells you something looks wrong, you can compare the number at the foot of their shop page with the number on your own More screen and know in one glance whether they are seeing the same page you are. It is the smallest type on the page and the same muted grey as the credit above it, in the developer's own corner, where it is there when you go looking for it and out of the way of everything a customer came for.

One step on your machine. The part of this that builds the email lives on Supabase, so it must be redeployed once, exactly as before:

supabase functions deploy shop-feedback --project-ref hzpyblqygnntixkijeem

Nothing else to set up - no new secret, no new database table, no new key. The email reuses the same Resend account your wish list already sends with, and the same published settings row the shop page reads for its "Website by" line.

30 Sep 2026 - engine v251, THE SHOP'S QUESTION IN YOUR OWN WORDS (no database step, no redeploy, one push).

The line at the foot of the shop now reads exactly as you wrote it. "Webmaster: Like the User Interface? Tell me & I will improve it!" - your words, your punctuation, with "Tell me" back where you put it.

It is a measurement, not a rewrite, and this version takes the measurement in both directions. A phone gives that box about 323px of room, and "the User Interface" spelled out costs about 150px of that on its own. Your sentence runs to about 377px, so on a phone it takes **two lines** - and the guard built for exactly this opens the box to show the whole question rather than cutting the end off it. Nothing is clipped on any phone: at 320px the box simply opens a little further, and the page never scrolls sideways.

The Malay is your own sentence and holds one line - as long as it keeps the shape you wrote it in. "Suka User Interface ini? Komen & Saya akan perbaiki!" measures about 322px against a phone's 323px, which is a fit to within a pixel. It carries no "Webmaster:" prefix for that reason: the prefix would add another 76px and push it onto a second line. The English and Chinese still say who is asking.

The Chinese needed no change at all - it already used the everyday word for a screen's interface rather than a shorthand, which is why it reads the same as before.

30 Sep 2026 - engine v250, THE SHOP'S QUESTION SPELLS OUT "USER INTERFACE" (no database step, no redeploy, one push).

Shorthand is out. You looked at the built page and said it plainly - "UI should be user interface, ui is not a laymen term" - so the line at the foot of the shop read **"Webmaster: Like the user interface? I'll improve it"**, with the Malay and Chinese equivalents.

And it stayed one line on a phone, which is what made this a measurement rather than a rewrite. A phone gives that box about 323px of room, and "user interface" spelled out takes about 150px of it on its own. The full sentence you first wrote runs to about 373px - wider than the box - so its tail would have been cut off. Everything you asked for was kept: who is asking, the question in full words, and the promise that it gets improved. Only "Let me know" went, because the question mark and a box waiting to be typed in already say it.

The Chinese needed no change - it already used the everyday word for a screen's interface rather than a shorthand. The Malay used "UI" too and spelled the term out as "antara muka", for exactly the same reason the English does.

30 Sep 2026 - engine v249, THE BOX AT THE FOOT OF THE SHOP SAYS WHO IS ASKING (no database step, no redeploy, one push).

The one quiet line at the foot of the shop page now introduces itself. It reads **"Webmaster: Like this UI? Tell me, I'll make it better"** - the customer types over that sentence exactly as before, and the whole of it still goes out as their own words when they send, with nothing to press but Enter.

It is one line on a phone, and it is never a line and a half. The sentence was measured against the width of the box at phone size rather than guessed at, which is why it says "this UI" and "Tell me" rather than spelling

everything out: the longer wording you first wrote runs wider than the box on a phone, and its tail would have been **cut off mid-sentence** - the question would have read as far as "...Let me know, and I" and stopped. What you asked for is kept in full, in the shorter form, and all three languages carry it.

And if a phone is narrow enough that the question does wrap anyway, the box now opens itself to show the whole question instead of hiding the rest. That is the one thing this version can do that a shorter sentence alone could not: on any phone, in any of the three languages, the customer always sees the entire question rather than its first line.

The Chinese and the Malay say the same thing as the English, in the same few words: the Chinese names the person asking as the site's webmaster and asks whether the customer likes this page, and the Malay uses "Webmaster" as it is, since that is the word a Malaysian website uses too.

30 Sep 2026 - engine v248, NOBODY HAS TO PRESS ANYTHING, AND NOBODY LOSES WHAT THEY WROTE (no database step, no redeploy - the same one function, one push).

Two more things you asked for on the little box at the foot of the shop.

Closing the page is now a send. A customer who writes a sentence and then closes the shop, or taps a link out of it, has had their say - the words go as they leave, without them pressing Enter at all. **Nothing goes out while they are still on the page**, so a sentence somebody is midway through writing is never mailed off behind their back. Pressing Enter is still a send, and still shows your reply straight away.

And nothing they wrote is lost on the way. If they step away from the box with words in it, or the page is closed before anything could be sent, the sentence is **kept on their own device** and is waiting in the box the next time they open the shop - it goes only once it has been sent. If a send does not go through, their words are put back, so a failure nobody was around to read is never a sentence lost. And **with no internet, nothing is sent at all**: the words simply wait on their device for the next visit instead of going down with the page.

The box itself is unchanged and stays as quiet as it was: one line, no Send button, and one small grey line that says how to send.

And the Bahasa Malaysia has been rewritten to read the way Malaysians actually write. Several lines on the shop were correct textbook Malay but nobody writes that way on a Malaysian website - "pembangun" for your developer, "telah" where anyone here would say "sudah", "kemas kini baharu", "Menghantar..." where a shop would say "Sedang dihantar...", and a reply that read like a letter from an office rather than from you. The words are the same promise, in the voice a Malaysian customer expects: **"Idea anda sudah kami terima! Update terbaru akan datang tak lama lagi."**

The Chinese is now Malaysian Chinese, which is what your customers read. Malaysian Chinese read simplified characters exactly as mainland China does, so the characters stay as they are - what changes is the words. A few lines had been written the way a shopping app from China writes them: the items count used the mainland measure word rather than the one a Malaysian uses, the basket was called a shopping bag, two lines said "coming soon" and "enjoy" in the flat official way, and the two arrows on the delivery calendar were in **traditional** characters, which is simply the wrong script for your customers. The reply now says, in Chinese, that their idea has been received and will be taken good care of, with new updates coming very soon - not that it "will be processed". Taiwanese Chinese was not used: it would mean changing every character to traditional, and your customers do not read that.

30 Sep 2026 - engine v247, THE SHOP FRONT NOW ASKS THE CUSTOMER WHAT THEY WOULD CHANGE (no database step, one new function to deploy once, one push).

You asked for a way for a customer to tell your developer what they would improve about the shop page, in their own words, with a reply straight away.

A single line now sits at the foot of the shop, under the "Website by" line. It is printed with your own question - **"Like this UI? Comment, and I will make it better."** The customer types over that sentence and presses **Enter**. There is no Send button, because you asked for none: Enter is the send. The line is one line tall and opens further only when the words no longer fit, so at rest the whole thing is a single quiet row at the foot of

the page. Once there are words in the box a small grey line under it says **"Press Enter to send"**, so the one key that does something is never something they have to discover by accident.

The moment the send lands, the box is replaced by your reply - **"Your idea is well taken care of. New updates soon!"** - so they see the answer without doing anything else. They are never asked for an email address, and nothing is created on your side for them.

The words come to you as an ordinary email, at the developer address you already set in **Settings -> Website & developer** - the same address the shop's "Website by" line shows. It runs on the **same email service your wish list already uses**, so there is no new account, no new key and no new DNS record. Each message says which page it was written on, in which of the three languages, and when it arrived.

There is no box until you have set that address, because the box has nowhere to send to without it. That is on purpose: a box that quietly collects sentences nobody will ever read is worse than no box at all.

A send that did not go through says so, plainly. If the email could not be sent the customer is told, in words, to try again or to use the WhatsApp link that is already drawn just above the box - and **their words stay in the box**, exactly as typed, so nothing they wrote is lost. A thank-you drawn over a message that never left would stop them trying again and nobody would ever find out, so that is the one thing this never does.

It speaks all three languages. The question, the line that says how to send and the reply are written in English, Mandarin and Bahasa Malaysia, and a customer who switches language keeps whatever they had half-typed and whatever reply they had already been given.

What you have to do, once. This needs one small new function deployed in Supabase the same way the wish-list function was - the file is supabase/functions/shop-feedback. Until it is deployed the box still appears and still says plainly that the send did not go through. Nothing else is required: no database step, no new key, no new setting - it uses the developer email you have already saved.

30 Sep 2026 - engine v246, THE NOTES ON THE ITEMS ARE NOW UNDERLINED ON THE COMPACT LABEL TOO (no database step, no redeploy, one push).

You spotted this one: on the Compact label a note like **no nuts** was printing as ordinary words, while the same note stood out on the Full and Mailing labels. It was the last place a customer's own words went unmarked, and it was easy to miss - Compact puts every item on ONE line, so the note sat in the middle of that line with nothing to pick it out.

Each noted item's words are underlined on that joined line now, on the item they belong to, so a note on the second of four items is underlined on the second item and not somewhere else on the line. **If two items carry the same note, each one gets its own underline** - the line is not allowed to mark one item twice and leave the other plain.

The line itself has not changed at all. It reads exactly as it did - the items in order, the notes in brackets beside the item they belong to - and a label for an order with no notes on it is byte-for-byte what it was. Nothing else about any label moved: the Full, Mailing and Name-only styles print exactly as they did, and the delivery note on the Compact label is untouched by this.

30 Sep 2026 - engine v245, THE DELIVERY NOTE IS UNDERLINED TOO, AND IT NOW PRINTS ON THE COMPACT LABEL (no database step, no redeploy, one push).

You made the point that the **delivery note** is not a courier thing - it applies to a self-collect order just as much, so it should stand out and be printed wherever it belongs.

It is underlined now wherever it is drawn. The order row carries it as Aunty Bee · 012-345 6789 · deliver after 3pm, with the note itself drawn under, exactly as a note on a single item already was. The same on the printed label: the Note: line keeps its label plain and underlines the customer's words.

The Compact label prints it too. Compact used to join every item onto one line and leave the delivery note off entirely - so a note about the doorstep, the gate code or a collect time disappeared the moment you picked the denser label, on a self-collect order as much as a courier one. It now prints in the same place the Full label puts it: after the items, before the courier address.

Everything else is unchanged. The Mailing and Full labels read exactly as they did, the note box is still offered on both self-collect and courier orders, and an order with no note on it looks and prints exactly as before. Name-only stays a bag tag - code and name, no fields to read.

30 Sep 2026 - engine v244, A NOTE A CUSTOMER LEAVES ON ONE ITEM IS NOW UNDERLINED, SO IT IS HARD TO MISS (no database step, no redeploy, one push).

You asked where the words a customer leaves actually reach you. The answer turned out to be uneven: the **delivery note** on the whole order pings your phone the moment the order lands, but the **note on a single item** - "no nuts", "write Happy Birthday" - only ever lived inside the app, sitting quietly in grey beside its item where it is easy to skim past. You decided that is enough for it to be **underlined**, so it reads as something you were meant to notice.

On the Orders list, an item that carries the customer's own words now has those words underlined - Focaccia 800g ×2 (no nuts) with the brackets drawn under. It is the same line in the same place, just marked. Nothing moves, nothing is added, and an order with no notes on it looks exactly as it did.

On the printed label, the note beside its item is underlined the same way, on both the Full sheet and the Mailing sheet - so what you hold in your hand while you kit an order is as easy to read as the screen.

The line itself is unchanged. Only the way it is drawn changed, so the row, the label and the search still read the very same words - the underline cannot quietly disagree with what the line says.

30 Sep 2026 - engine v243, THE DELIVERY-DAY CALENDAR IN THE ORDER SCREENS NOW FOLLOWS TODAY, THE SAME AS YOUR SHOP (no database step, no redeploy, one push).

Your words: you wanted the calendar in the + New order card to behave like the one on your shop.

It does now, and it is the **same calendar in all three places it appears** - the + New order card, the strip at the top of Orders, and the day picker in the Edit-order pop-up. One behaviour in three places, so a day sits in the same square wherever you meet it.

What was wrong with it. It was a **month** - the whole of September, then the whole of October - exactly the calendar v231 replaced on your shop, for the same reason. On the 30th, most of what it drew was days already gone; a delivery day on the 29th could not be seen at all; and both arrows were drawn even at the ends, so a greyed button sat there inviting a press that would do nothing.

From this version it is five whole weeks anchored on today. The first row is the week just gone, the second row is this week with today ringed so you can see where "now" is, and the three rows under it are the weeks ahead, which is where your delivery days actually live. Today is always in the second row. **Every square is a real date now** - no blanks, no padding, so a delivery day at the start of next month is visible and openable from the moment you publish it.

Your delivery days are unchanged on it. Still circled, still carrying how booked each one is under the number, still FULL in red where a day is at capacity, and a day already gone still dimmed but **still open** - you backfill and review old days. A day you do not deliver still answers a tap, and still names where it gets added, so a holiday falling on a day you do not deliver is never the one square with nothing to say.

The arrows move one whole week, and are gone where there is nowhere to go. A dead control reads as a bug, so at the ends of the days you have set there is simply **no arrow** - nothing greyed out, nothing to press that would do nothing. The window still stops where your shop's stops for the same list, so the two calendars agree about where the far end is. A new week arrives **sliding up** going forward and **coming down** going back - the direction you are travelling in - and only when you pressed an arrow: an ordinary redraw never replays it.

The marks you made on Delivery Dates are on it too - the same single-day boxes and the same see-through bands, in their own colours and their own depths.

It opens on the day you are working on, not on today. The Edit-order pop-up opens on the day that order is actually on, so moving an order a season out does not leave you looking at a week with nothing to do with it. And picking a day that is already on screen **does not move the window**, so the grid never slides out from under your finger.

One thing I deliberately left alone. v231 ended with "nothing in the backoffice calendars changed; this is the shop's calendar only" - that was your instruction then, and this version is you asking for the opposite, for this one calendar. The **look** of the past days in the backoffice is untouched, and so are the other admin calendars: the free order-date field, Delivery Dates, Profit and Deliveries. What changed here is this one calendar's behaviour.

30 Sep 2026 - engine v242, A CUSTOMER WHO ALREADY HAS A COURIER BOOKING IS NEVER QUIETLY PUT ON A SECOND VAN (no database step, no redeploy, one push).

Your report: if a customer's order already has a courier booked, the delivery run should not tick that order, and it needs to say so, so that a delivery is not booked twice.

Here is what was happening. A booked trip is recorded on the order's own row, and the delivery run did not look at it when it built the day. Every customer on the day opened already ticked, the Tick them all press ticked them all, and a booked customer looked exactly like any other. The only thing that refused was the app's own sentence at the Book press - so you could tick a customer a driver was already on the way to, price the run, and press Book, and the refusal arrived afterwards as a message rather than as anything you could see while you were choosing.

From this version, a customer whose trip is already running opens OFF the run, with a line under their name saying so. The line names the courier and its own word for where the trip has got to, and it says plainly that ticking them and booking the run would send a second vehicle to the same door. The head of the list counts them too - it will read something like "Who is on the run - 3 of 5, 1 already booked", so the number you tick and the number of rows in front of you never disagree.

It is not a gate. The tick is still live and still yours. If you tick a booked customer by hand the app honours it, and the refusal comes at the price instead, in the same words as before. That is deliberate: the tick is a decision you make, and the app's job is to make sure you can see what you are deciding.

And the row offers the way out, rather than going grey. Under the warning there is one press: Call off the trip and add to this run. It asks you first, in the open, because calling a trip off is the half that cannot be undone - the driver stops being sent, and the customer's tracking box keeps its link but nothing will update it any more. If you say yes, the original booking is cancelled with the courier, the order's own row records that it was you who called it off and when, the warning disappears, and that customer is put on this run so their delivery can be consolidated with the rest of the day. If the courier refuses to call the trip off, its own reason is shown and nothing is written.

A trip that has already finished - delivered, or cancelled earlier - is not a barrier at all. Those customers open ticked like anyone else, with no warning, because a delivery that has to be done again still has to go on a run.

Orders going out by parcel carrier are untouched by this version. They were never on the run list in the first place, so there is no second van for them and nothing here to warn you about.

No database step, no redeploy, no new key, and no new setting to switch on.

30 Sep 2026 - engine v241, THE RESET NOW ANSWERS ON EVERY CARD, NOT ONLY A FRESH ONE (no database step, no redeploy, one push).

You pushed v240 and you were right again: the press worked, but the card could still stay silent. This version closes the last way it could.

Here is what was happening. The answer your press writes - the pin moved, or it found the same spot, or the address could not be reached - has two possible places to be said. One is the line under the buttons on the door card itself. The other is the line inside the delivery-price section. That section folds away, and its line folds away with it. v240 made the press work before the price section existed at all, which was the right fix - but the moment you opened **Get a delivery price** even once, the card handed every later answer to the folded-away section and stopped using the door card's own line.

So the sequence that caught you was this: open an order, press **Get a delivery price**, close it again, then reset the pin. The look-up really ran and the pin really moved, but the sentence saying so was written inside the

section you had folded away, and the door card said nothing at all. And where the look-up answers with the same spot as before, nothing moved on the map either - so the whole press looked exactly like a press that had never happened.

From this version the answer is always said on the line you can actually see. While the price section is open it is said there, exactly as it always was. While it is folded away - or if you have never opened it at all - it is said on the door card, right under the buttons. The two lines are never both used at once, so they cannot come to disagree about what your press did.

The same fault was on the pin you drag. Moving the pin by hand ends through the same sentence about the price section's numbers, and that sentence had the same problem: with the price section folded away, a drag you had just made reported into a section you could not see. It now says what it did on the door card in exactly the same way.

Nothing else changed. No database step, no redeploy, no new key, and no new setting to switch on.

30 Sep 2026 - engine v240, THE RESET BUTTON NOW ACTUALLY PRESSES (no database step, no redeploy, one push).

You pushed v239, saw the reset button at last, and pressed it - and nothing happened. That was exactly what your report said, and it was exactly right. The button was on the card, but its press was not connected to anything yet.

Here is what was happening. The reset press was wired up at the same moment as the delivery-price half of the card, and that half is only built the first time you press **Get a delivery price**. So on a card you had just opened, the button was drawn and waiting but its press led nowhere: no look-up, no message, no movement. The instant you happened to open the price fold once, the button started working - which is why it looked like an ordinary working button that simply ignored you.

From this version the reset press is connected the moment the card opens. It never waits on the delivery price again. You can open an order, see the pin, and reset it without touching the price button at all - which is the order you would naturally do it in anyway, because you reset a stale pin when you call the customer, not when you ask for a price.

And the press now always tells you what it did, on the card itself, right under the buttons. It says it is looking the address up, then it says one of the three things that can happen: the pin moved, or the look-up found the same spot and there was nothing to move, or the address could not be reached and the pin has been left exactly as it was. Those words used to appear only up in the price section, which is another reason a press on a fresh card looked like it had done nothing at all.

If a reset is about to replace a pin a person chose - the customer's own pin, or one you placed by hand - it still asks you first, exactly as v239 described.

Nothing else changed. No database step, no redeploy, no new key, and no new setting to switch on.

30 Sep 2026 - engine v239, THE RESET NOW SHOWS UP WHERE YOU ACTUALLY WORK (no database step, no redeploy, one push).

This is the fix for v238. You pushed it, opened an order, and the reset was still not there. You were right, and the reason is this: v238 offered the reset when the pin was the customer's own, and hid it when the pin was one **you** had placed by hand.

But placing the pin by hand is the only thing that card ever offered you. So a drag is what you did, and a drag is exactly what then hid the reset. Every customer whose pin you had ever corrected by hand still showed **Move this pin** and nothing else. v238 changed nothing at all for the doors you had touched yourself, which is why it looked like it had not arrived.

From this version the reset is offered wherever there is a pin to replace. If a pin exists on the order and an address is typed, the press is on the card - the customer's own pin, a pin you dragged by hand, a pin an older look-up wrote, or one a reset already wrote. The only case with no press is an order that has no pin yet, where there is simply nothing to replace; the price button looks one up on its own.

Where it would replace something a person chose, it still asks first. That is the customer's own pin, and it is also now your own hand-placed pin. The wording names which of the two it is about to replace. Over your own pin it says this is the door you placed on the map by hand, that a look-up may only find the road and can be a step back from a door you already had right, and that you can drag the pin again afterwards. Over the customer's pin it says, as v238 did, that this is their pin and you can switch back to it. Either way the confirming button says **Reset the pin**, and Cancel changes nothing at all.

Where it only replaces the app's own guess, it does not ask - a look-up's answer, or a reset of one. That is what this press has always done, so nothing new is put in your way there. The button says **Look this address up again** on those, and **Reset the pin from the address** on the two that will ask.

Nothing else changed. No database step, no redeploy, no new key, and no new setting to switch on.

30 Sep 2026 - engine v238, RESET THE PIN, AND LET THE MAP ZOOM (no database step, no redeploy, one push).

This one is yours, word for word: when you call a customer, their address may already have changed, and the only thing on offer was to move the pin. The map would not zoom out, and dragging a pin onto the right rooftop by thumb was slow and easy to get wrong. You asked why there was no way to reset the pin. There wasn't one. There is now, and the map zooms.

There is now a reset, on the pin card, where you can see it. When the pin the driver is sent to is the customer's own - the one they dropped on your shop page - the button under the map reads **Reset the pin from the address**. One press asks your address service for the address on the order again and puts the pin where that answer lands. That is the case you described: the pin was right when they dropped it, and they have moved since.

It asks before it replaces their pin. The customer's own pin is a fact they gave you, so a button that quietly overwrote it would be the pin moving on its own again - the fault several versions of this card have been spent ending. So the press asks first, in plain words: this is their pin, resetting replaces it with a fresh look-up of the address, and a look-up may only find the road. The confirming button says **Reset the pin**. Cancel changes nothing at all, not even a look-up.

And you can put their pin back. Once a reset has replaced their pin, the same card offers **Use the customer's pin instead**, one press, back to their exact point. The reset is not a one-way door.

A reset yields to a pin they drop afterwards. If that customer pins a new spot on your shop page after you reset, their new pin wins again straight away - the same rule as before. A reset replaces one specific stale pin, not every pin that customer will ever drop.

The map zooms the whole time now. The plus and minus buttons are drawn on the pin card, and pinch, double-tap and a desktop box zoom all work - before you press anything. Only moving the pin is still behind **Move this pin**: the map will not pan under a pin you are reading, and the pin will not take a drag, until you say so. That split is the point of this version. Looking and correcting used to be one action; they are two again.

Nothing else changed. Your prices, your delivery calendar, your messages, the order card, the packing slip and the label sheet are all exactly as they were. Adding an order is still the one pass v237 made it, and a pin you placed by your own hand is still not offered up for replacement.

On two phones. If your second phone has not been reloaded yet and is still running v237, it reads a reset as the older kind of pin and will price to the pin you replaced until it is reloaded. Reloading it fixes it, and there is nothing to run in Supabase for any of this.

30 Sep 2026 - engine v237, A COURIER ORDER IN ONE PASS, AND THE DAY ON ONE LINE (no database step, no redeploy, one push).

This one is yours, and it came from a complaint about your own way of working. Adding an order used to be two steps: fill in the New order card, press Add order, then find the order again and open the Edit window to finish it. You said that was redundant, that the flow was not smooth, and that it started in the wrong place.

A courier order is now one pass. The reason it was ever two steps is that four things lived only in the Edit

window and not on the card: the courier charge, the tracking number, the parcel carrier, and the Get a delivery price block. They are all on the card now. Choose Courier delivery and they unfold underneath it, in the order you need them: the address with the map pin, the charge, the tracking number, the parcel, and the price. Take the order, price the trip and record what it cost, all before you press Add order once.

The price is on the card. The booking is not. This is the way you asked for it. You can ask for a price and take the fee onto the order from the card, but Book this trip stays on the order itself, where it has always been. Booking a real vehicle against an order you have not finished taking would be the wrong thing to put under your thumb one press away.

The day is one line now. The card used to open on a whole month of calendar, even though you had already picked the day on the screen behind it. It now reads as one line - Delivering Fri, 2 Oct - with the calendar unfolding underneath only when you tap it. Nothing about how a day is chosen has changed: tapping one still switches the screen, so the products and the day's limits are right.

The items come first. After the day comes what the customer wants, then the customer. It is the order things are actually said to you on the phone: they tell you what they want, then who they are. Everything else on the card is unchanged and in the place it has always been.

Your own note box is now the delivery note. Your app and your shop both used to call the order's note box simply Note. Now that every item carries its own note, that box was left doing a different job, so it says what it does: it is the delivery note, for the gate code, the landmark and the time you should arrive. The box on your shop and the box in your app now show a grey hint saying exactly that, so a customer knows what belongs in it without you having to tell them.

Nothing is written until you press Add order. The card was, and still is, a draft. The charge, the tracking number and the parcel are written onto the order at the same moment the order itself is created, so an order you decide against leaves nothing behind it. And if you type a charge but do not say who paid it, the card refuses in words and changes nothing, exactly as the Edit window does.

Nothing else changed. The order's own field order in the Edit window, your prices, the delivery calendar, your WhatsApp messages, the packing slip and the label sheet are all exactly as they were.

30 Sep 2026 - engine v236, A NOTE ON EACH ITEM, SWITCHED ON PER PRODUCT (no database step, no redeploy, one push).

This one is yours, word for word: you wanted to switch a note box on or off on a product card, so a customer can leave a note on each thing they order - and switch it off on the products where a note makes no sense. And you asked for the same box in your own app, so a note you are told over the phone can be written down in the same place.

The switch lives on the product card. In your app, open a product and you will find a tick that reads **Ask the customer for a note on this item**, sitting beside the other selling switches. Tick it and the customer is offered a note on that item; leave it alone and they are not. It is off until you tick it, so nothing on your shop changes until you decide it should.

On the shop, the note is quiet. A customer who puts an item in the basket sees a small **Add a note** link under it. Nothing is opened on their screen until they tap it, so your menu still looks exactly as it looks today. Tap it and a box opens on that item, with a short example in grey: no nuts, or a name to write on the loaf. Whatever they type stays with that one item and rides on that item's line when the order comes in.

The note belongs to the item, not to the order. This is the whole point of it. If someone orders a focaccia and a sandwich and writes "no nuts" on the focaccia, that note comes through on the focaccia line only. The sandwich line is untouched, and the order's own note box at the bottom still works exactly as it did for anything that is about the whole order.

A note in your own app is offered on every line. Your New order card, and the Edit order window, both give every line its own small note box - whichever product it is, and whether or not you have switched the customer's box on. The switch decides what you ASK a customer for; it must never decide what you are allowed to write

down. A phone order of "no nuts on the focaccia" has to have somewhere to go.

You see the note where you work. On the orders list it appears in brackets beside the item it belongs to, so it sits next to the thing it is about rather than in a note line at the foot of the order. The same is true of the packing slip and the label sheet: it prints beside that item and never as the order's own note.

Nothing about a note blocks an order. An empty box and no box at all place exactly the same order - there is no such thing as a half-filled note to chase. Switching a product's note box off after the fact never hides or loses a note you have already collected, and clearing a note you wrote puts that line back exactly as it was.

Your WhatsApp messages do not change. The note stays in your app, which is what you asked for. Every message you send reads exactly as it reads today, and your customers' tracking page is untouched by this version.

Nothing to set up. There is no database step and nothing to redeploy. The switch travels to your shop with the rest of the product details, and a note travels with the order from one of your phones to the other on its own, the same way the rest of the order already does.

Nothing else changed. The order's own note box, the prices, the delivery calendar, the basket and the place-order bar are all exactly as they were.

29 Sep 2026 - engine v235, A BOOKED TRIP NOW SHOWS WHAT IT COST AGAINST WHAT YOU CHARGED (no database step, no redeploy, one push).

This one is yours, and it started as a worry rather than a request. A courier price you pick when you take an order can turn out to be higher or lower by the time the trip is actually booked, and nothing was watching that gap. You asked to see it rather than be protected from it: a real cost makes you aware, leaves room for a promotion later, and lets you decide not to collect the delivery at all.

So the booked-trip card now does the sum. Under the line that tells you when calling the trip off stops being free, the card compares what the trip actually cost against the courier charge on the order, and says which way the difference fell. If the trip cost more than you charged, it says so and tells you that much came out of your own pocket. If it cost less, it says that too, and that the difference stayed with you. Neither direction is treated as the bad one.

It shows both directions on purpose. A gap that only spoke up when you were short would be an alarm; you asked for a reading, so a surplus and a shortfall are worded the same way, in the same plain sentences.

A trip you decided not to charge for says so. If you have recorded no courier charge at all, the card tells you the whole cost of the trip is your own, rather than staying silent. That is the free-delivery case you named, and the one number worth seeing is the one you chose to give away.

A charge that matches the trip says nothing at all. When the two agree there is no line, because a difference of nothing is not a difference, and an "RM 0.00" row on every card would be noise you would learn to skip.

A delivered trip keeps the reading. The calling-off deadline disappears when a trip is over, because there is nothing left to call off. This line stays, because a cost is money that has already been spent and the whole point is to learn from it afterwards.

This is on your screen only. Nothing here goes near a customer. The price a customer was quoted does not move, their messages do not change, and their tracking page is untouched by this version.

Nothing else changed. The van you pick, the prices, the five-minute life of a price, the booking buttons, the charge box and the calling-off deadline are all exactly as they were.

29 Sep 2026 - engine v234, A BOOKED TRIP NOW TELLS YOU WHEN CALLING IT OFF STOPS BEING FREE (no database step, no redeploy, one push).

This one started as your own question: if you book a trip and then need to call it off, when does it start costing you money? The answer is a rule rather than a feeling. Lalamove lets you call off a scheduled pickup at no charge up to 45 minutes before it, and it may charge a fee after that. That rule was living in your head, so this version puts it on the page. Under the status line on a booked trip, the card now says **Free to call off until 30**

Sep, 10:15 am - Lalamove may charge a fee after that.

It turns over by itself when the window shuts. Leave the card open and the sentence changes on its own to **Lalamove's free calling-off window shut at 10:15 am, so a fee may apply from here.** You never have to work out which side of the deadline you are on, because the card says it for you.

A trip booked for as soon as possible is told apart honestly. On an immediate booking the free window runs from the moment a driver takes the job, and Lalamove never sends that moment back to the app. So rather than inventing a time that could be wrong, the card states the rule instead: **Booked for collection as soon as possible, so there is no pickup time to count back from. An immediate trip is free to call off only for a short while after a driver takes it - check with Lalamove before you count on it.**

Nothing is claimed that the app cannot stand behind. Every sentence says the fee may apply rather than that it will, because it is Lalamove's own terms that decide it and not this app. A trip that has finished, or one you have already called off, shows no deadline line at all, since there is nothing left to call off.

Nothing else about booking changed. The vans, the prices, the five-minute life of a price, the customer's tracking box, the charge box and the buttons on this card are all exactly as they were. This is one line added to a card you already had.

29 Sep 2026 - engine v233, YOUR ORDER PAGE NOW HAS A WAY BACK TO YOUR HOMEPAGE (no database step, no redeploy, one push).

Your order page had no route to your homepage at all. Someone who lands on it straight from an Instagram bio or a shared link sees the bakery name and the menu, and if they want the reviews, the gallery or the story, there was nowhere to tap. From this version there is a small **Our homepage** link in the orange banner at the top, marked with a little house, on its own line under your tagline - and it takes them to **jienluv2bake.com.my**.

It sits on its own line on purpose. The banner's top row already holds the "Made to order" line and the EN / Mandarin / BM buttons, and on a phone the "Made to order" line already wraps onto two or three lines - a third thing squeezed in there would have crowded all three. Under the tagline it belongs to the banner, reads clearly in white on the orange, and is a full-size tap target rather than a thin line of text.

It is translated like everything else on the page. A customer reading in Chinese or in Bahasa Malaysia sees it in their own language, and it changes with the rest of the page the moment they tap a language, with nothing re-loaded.

Nothing else moved. The menu, the delivery calendar, the basket, the order form, the track-your-order box and the place-order bar are all exactly as they were. Your homepage still opens your order page in a new tab, so a customer who came that way keeps both pages open.

28 Sep 2026 - engine v232, A TRANSLATED LINE YOU DELETE STAYS DELETED, AND THE GREY HINT IS SHORT ENOUGH TO READ (no database step, no redeploy, one push).

This one comes straight from you, and you described it exactly: you delete the Chinese or the Bahasa Malaysia wording on a product, and **the machine translation comes back** - greyed, sitting in the box you had just emptied, with the little sentence about filling in the English, and the right-pointing arrow over it. Then the arrow went away and left the greyed words behind with no way to accept them.

What was actually happening. The greyed words and the arrow are supposed to be **one offer**, made together: the arrow is drawn on the words, and the words are only there because the arrow can take them. They were being written in two different places, so they could disagree - and when you emptied a line, the words stayed and the arrow did not.

From this version a line you have made yours is yours. If you type your own wording into a translated line, or you **delete** what was there, that line goes **quiet**: nothing greyed comes back into it, and no arrow is put over it. It stays as you left it when you save.

The circular-arrow button is still your way back. A line you have emptied shows the **circular-arrow button** on its right instead of the arrow, so a translation is always one tap away if you change your mind - the line is never a dead end. The same circular-arrow button re-translates a line the machine filled for you.

A line you have not touched is unchanged. Empty, never typed into, with English above it - it still shows the suggested words greyed with the **->** to take them, exactly as before. What is on offer, and how to take it, are now decided in the one place, so they can never disagree again.

And the grey hint is short enough to read now. It used to say **"e.g. .....if blank, it will be filled with English"** after the suggested words, and that is too long for a one-line box: on your phone it ran past the right edge and was **cut off mid-word**, right where the arrow sits. It now shows just the suggested words with the plain **"e.g."** in front, the same shape as every other suggested box in the app. The promise it was making is still made, in full, in the card's own line above: leave a line blank and the English shows instead.

A line you emptied says what blank means. Instead of the machine's words lingering in it, an emptied line now carries a small grey note - **"Left blank - English shows."** - so it is clear that leaving it empty is a real choice and not a mistake, and that your customers will simply see the English.

Nothing else moved. The translations still come from the same free service, still only ever appear as grey suggestions, are still never saved over your own words, and the shop still falls back to the English wherever a translation is missing. Nothing here touches your shop, your orders or your prices.

28 Sep 2026 - engine v231, THE PICK-A-DELIVERY-DAY CALENDAR IS NOW FIVE WEEKS THAT FOLLOW TODAY (no database step, no redeploy, one push).

The calendar on your shop was a **month** - the whole of September, then the whole of October. That sounds tidy and it is not: on the 28th, almost everything on it is days you can no longer sell, and the days just before the 1st and just after the last day of the month were **blank squares**, so a delivery day at the start of next month was invisible until the month turned over.

From this version it is **five whole weeks anchored on today**, and it moves with the date.

The first row is the week just gone, faded the way it has been faded, and faded **harder than before** now that there is a week of it - it is there as a sense of where you are in the month, not as something to read. **The second row is this week**, with today ringed so the customer can see where "now" is. The three rows under that are the weeks ahead, which is where the delivery days actually live.

Every square is a real date now. No blanks, no padding - the days before the 1st and after the last of the month are simply part of the row they fall in, so a delivery day at the very start of next month is visible and bookable from the moment you publish it.

Your delivery days are unchanged on it - still circled green with the day you have open, and still struck out where a day is full. The line under the calendar still names the day they are getting.

Paging, and the motion you asked about. If you publish a delivery day further out than the window reaches, **arrows appear either side of the title**. Tapping the one on the right moves the window **one whole week later** and the dates arrive **sliding up**; tapping the one on the left moves it back and they arrive **coming down** - the direction you are travelling in, so a tap reads as a tap. The arrows are **not there today**, and that is on purpose: all five of your published dates sit inside the one window, so there is nowhere to go and **nothing dead is drawn on your shop**. They will appear the first time you publish a date roughly three weeks out.

Your shop's rules are untouched. The 18:00-the-day-before closing, the sold-out days, the day the customer has chosen, the product counts, the refresh that keeps a page open in someone's hand up to date - all exactly as they were. **Nothing in the backoffice calendars changed**; this is the shop's calendar only.

28 Sep 2026 - engine v230, THE DELIVERY ADDRESS BOX IS NOW THE FULL WIDTH OF THE FORM (no database step, no redeploy, one push).

Last version made that box **tall**; you told me it was still cramped, and you were right - it was **narrow**, and I had not seen it. From this version the delivery address box is **as wide as the form**, so a whole address sits on the line it belongs on instead of wrapping three words at a time.

Why it was narrow. The order form lays its fields out in **two columns**, and the address was sitting in **one** of them - with the space beside it standing **empty**. The address is the longest thing you type into that form, so it now takes **both columns**, and that wasted gap is gone. On a phone it went from about **137 pixels wide to**

about 315 - the same width as the shop's own address box, which was never the problem.

It is in both order forms - the **+ New order** card and the **Edit** pop-up - so the two behave alike, and **no other field on the form moved**.

The map card's address box is unchanged by this one. The box labelled **"The address you have"** was already full width - v229 gave it the height it needed, and nothing about its width was wrong.

Nothing else changed. The suggestions still appear as you type, the address still fills itself in when you pick a customer, and saving an order works exactly as it did. There is **no database step and no redeploy** - it is one push.

28 Sep 2026 - engine v229, THE DELIVERY ADDRESS BOX HAS ROOM FOR A WHOLE ADDRESS (no database step, no redeploy, one push).

A delivery address is four or five lines - the name, the street, the area, the postcode and the state. The box you typed it into was **one line tall**, so most of what you had written was hidden and you were reading your own address back through a letterbox. From this version that box is **tall enough to show the whole address at once**.

It is in both order forms - the **+ New order** card and the **Edit** pop-up - so the two behave alike. The box is also **draggable at its bottom edge** if you want it taller still.

The map card's address box got the same treatment. The one labelled **"The address you have"**, where you type the address before looking it up, is now multi-line as well - it holds a whole address too, and it had the same problem.

Nothing else changed. The suggestions still appear as you type, the address still fills itself in when you pick a customer, and saving an order works exactly as it did. There is **no database step and no redeploy** - it is one push.

28 Sep 2026 - engine v228, THE ADDRESS BOX SUGGESTS REAL ADDRESSES AS YOU TYPE (one Google setting to switch on, one redeploy - the app itself is one push).

You said last version's address helper only helped people who had ordered before, and that what you actually wanted was help **keying in an address for someone new**. That is what this version is.

From now, in the **+ New order** card and in the **Edit** pop-up, a **short list of real addresses appears under the delivery address box as you type**. Type the first few words of the street, and the list offers the addresses that match - each one a **whole address with its postcode and area**, not the half-remembered line you would have pecked in yourself. Tap the one you want and it fills the box. Then the app's own pin lookup can find the door from there, as it already does.

It never blocks you and it never changes what you have typed. Your words go into the box the moment you type them, whether or not the list appears, and nothing is saved until you press **Add order** or **Save**. If the list does not appear - no signal in the kitchen, or the lookup simply found nothing - you type the address by hand exactly as you do today, and the order saves just the same. **No suggestion can ever hold up a sale.** The list is the only thing that is ever missing.

It is only the delivery address, and only for you. The shop's address box is public and stays exactly as it is; this is for the orders you take by hand.

What you have to switch on, once. This one genuinely needs a Google setting: **Places API (New)** has to be turned on for your Google project, and ticked onto the key you already have, and the **courier** function has to be redeployed. **Until you do that, nothing is broken** - the list simply never appears and the address box behaves exactly as it does today. The step-by-step is in the Google Cloud setup guide, and I will walk you through it on your Mac.

What it costs: nothing at your size. The first **10,000 lookups a month are free**, and an order uses a handful. There is deliberately **no daily limit** on this one - you asked for it to keep working rather than stop partway through a day, so a runaway month would be a few ringgit rather than nothing. Only your signed-in phone can

reach it, too, so no one else can spend it.

Worth knowing: it sends your customer's half-typed address to Google as you type it. That is a step further than the door lookup, which sends a complete address once. It is the price of a suggestion list, and it is the same Google account the door lookup already uses - but you should know it rather than find out later.

28 Sep 2026 - engine v227, THE CUSTOMER'S ADDRESS FILLS ITSELF IN WITH THEIR NAME (no database step, no redeploy, one push).

When you type a customer's name into an order form, the app already offers you the people you have served before, each with their number and what they usually buy. Tap the right one and their **name and number fill in**. From this version their **delivery address fills in with them** - the address on the last order you actually delivered to them.

It is a default, never an overwrite. The address fills only a box that is still empty; it does not touch one that already says something. So an address you typed yourself, or the one already on an order you are editing, stays exactly as it is. A pick can never quietly move a delivery to a different door. And if a customer has only ever collected, or has no address on any of their orders, the box is simply left alone - nothing is written into it.

Both order forms do it - the **+ New order** card and the **Edit** pop-up - so the two behave alike rather than one of them surprising you.

It never asks the internet for anything. The addresses are the ones already in your app, from your own orders, so there is nothing to pay, no key to switch on and no Google account involved. It draws on the orders you have saved, so a customer who has never ordered before has nothing to offer yet, and it gets more useful as your customer list grows.

28 Sep 2026 - engine v226, THE SECOND KIND OF COURIER: PARCELS YOU POST YOURSELF, RECORDED IN THE APP (no database step, no redeploy, no key to set - one push).

You asked whether Lalamove is one kind of courier and SPX, Ninja Van, Line Clear and EasyParcel are another kind, and whether we could just have **two kinds**. Yes, and this is the second one.

There are exactly two kinds. Lalamove is *kind one* - on demand, point to point: it prices a vehicle for a journey, a driver accepts, and you get a name, a plate and a live share link. J&T, Ninja Van, Line Clear, Pos Laju and SPX are *kind two* - parcel carriers: they price a box by weight and by where it is going, you hand the box over, and it arrives next day or later. **EasyParcel is not a third kind** - it is one account and one API standing in front of most of kind two. So there are two, exactly as you said, and kind two is what this version adds.

What is new: More -> Parcel couriers. A list of your own, the same shape as your suppliers. You add a carrier by name, and there is **one press that adds the usual five** - J&T Express, Ninja Van, Line Clear, Pos Laju and SPX Express - so a fresh phone fills itself in. You can rename or delete any of them at any time.

On a courier order you now record the parcel. Open an order, and under the fulfilment picker you choose **which carrier took it**, type the **consignment number**, and press **Handed to the carrier** when you have actually handed it over. Press **Undo - not handed over yet** and that moment is taken back, while the carrier stays on the order. The cost goes into the **courier charge box the order already had**, with the same question of who paid it, so nothing about your books changes shape.

The order row shows one quiet line - for example *J&T Express · JT998877665 · handed over* - and only once a parcel is recorded. An order that is not a parcel looks exactly as it did before.

The customer is told. Their track card gains one line naming the carrier, and once you mark the parcel handed over it reads **Collected** - the same word a Lalamove pickup publishes. That carrier line is drawn **only on an order with no driver**, so every existing Lalamove card is left exactly as it was.

A product can now be ticked "Can travel as a parcel", and that tick blocks nothing. A courier order always offers the carrier picker. If a line is not marked, the app simply **names it** - "Not marked as able to travel as a parcel: Fresh Focaccia" - and you can still send it that way. Nothing is hidden, removed or turned off, because no rule in this app may block a sale you take by hand. The tick only feeds that one warning, and it is where a parcel's weight and box size will sit when the time comes.

Two things you should know about how it behaves. A parcel and a booked Lalamove trip are **mutually exclusive on one order** - whichever exists, the other's controls are replaced by a sentence naming it, and neither is deleted. And an order recorded as a parcel **can no longer be swept into a delivery run**, because it is not going on your van.

No price is shown anywhere, and that is deliberate. A parcel's cost moves with the zone and the box, and the way a box is measured into a price is not the same at every carrier - J&T and Line Clear divide by 6000, Ninja Van by 5000, which is about 20 per cent more on the identical box. One stored price per carrier would quietly lie to you, so there is no rate field at all. **There is also no key, no wallet, no account and no database step in this version** - you keep booking the parcel exactly as you do now, on the carrier's own site or over the counter, and the app remembers the rest.

Be prepared for this to look like almost nothing in your own shop, and that is correct. Every product here is fresh focaccia, and a parcel network is not where fresh bread goes. This is an engine feature and it is for **Munchies**, which is the shop whose products are the kind that survive a night in a van. You told me the plan: the parcel work is done here, in Bakeadmin, and Munchies syncs to it. When the booking side does come, it comes through EasyParcel with **Munchies' own account and its own wallet**, kept separate from this shop's - which is what you decided, and it is the right call.

One more thing, carried over from last version. A product that is on the shop can always be *hidden* - see v225 below - so nothing here changes how you take a product down.

Two faults were found and repaired on the way, and neither had ever reached you. In the app's own picker, the code that was supposed to run when you choose a new value had been quietly dead for about a hundred versions. **Nothing you could see was wrong**, because nothing in the app had yet depended on it - the new parcel card is the first thing that does, and a control that arrives one reopen late would have read as missing rather than as conditional. It is repaired once, in the shared picker, so it is now correct for the parcel card and for anything added later. The second was in the new card itself: what you had typed was being thrown away whenever the form redrew, and that is fixed before it ever shipped.

28 Sep 2026 - engine v225, A PRODUCT THAT IS ON THE SHOP CAN ALWAYS BE HIDDEN (no database step, no redeploy - one push).

You said *after publish a product, only delete is allow, it should not be*. You were right, and this is fixed. A product that is **on the shop** now always carries a *Hide* button. Before, the row decided for you: a product that already had orders, or that another product uses to make a set, offered *Hide*, but a **freshly published product with no orders yet offered only Delete**. So the one product you had just put up was the one you could take down only by destroying it.

Hiding never touches your recipe. Hide simply takes the product off the customer's menu and leaves everything of yours intact - the recipe, the ingredients and amounts, the photo, the price, the category, and any order history it has. *Unhide* puts it straight back. Nothing is deleted, and no database step is needed for any of this.

Delete is still there, beside Hide, on a product that is clean and unused - one with no orders and not used inside a set. That is the same rule as before, just no longer the only button on the row. So on a live product you now read *Edit · Hide · Delete*, and on one with history you read *Edit · Hide* (Delete never appears where it would break an order or a set).

Why this helps: the mistake this removes is a real one - publish a product, look at the card, and the only way off the shop is the button that also deletes the recipe. Now taking a product down is one tap and completely safe, and Delete stays for the case where you really do want it gone.

28 Sep 2026 - engine v224, THE PICTURE WINDOW IS SQUARE AND YOUR PHOTO IS CROPPED TO FILL IT, AND A PRODUCT ROW SHOWS ONLY WHAT A CUSTOMER SEES (no database step, no redeploy - one push).

You asked me for the window to be *square, and photo upload can be auto crop to fit the window*. Both are done. **The window is a square**, and **your photo is cut to a square as it is saved**, so it fills the window edge to edge

with no cream strips down the sides. The middle of your picture is the part that is kept, and the app shrinks it for you at the same time, so you never have to resize or measure anything. What you pick is exactly what shows.

This changes one thing from last version (v223), where your photo was shown whole. In v223 the window kept its shape and your photo kept its own, with the window's soft cream colour filling the spare space. You have now asked for the other thing, and this is it. Everything else about v223 stands: the window is still **one standard size in all three places** - your shop's card, your app's Products list, and the editor where you choose the photo.

What that means: whatever the square cuts off is gone. A very tall or a very wide photo loses its outer edges. If a crop has taken off a part you wanted, choosing the photo again is how you change it - the app keeps only the small square copy, so a picture cannot be un-cropped afterwards.

The pictures you already saved still work, and are now trimmed to the square. They were saved in whatever shape they were, so on your shop and in your list they show trimmed to the middle of the square window. Choosing the photo again gives you one cut for this window properly - that is the step you told me you would take.

A product with no picture is unchanged. It simply has no window, and its name and price read across the whole card exactly as before.

One thing to know, carried over from v223: a wider window leaves the words beside it a little less room. On a product that has both a picture and a long name, the name and price wrap onto one or two more lines than they used to. If that crowds the words more than you like once you see it, say so and I will make the window a little smaller - it is one number and a quick change.

The other half of this version: your product row is now only what a customer sees. You said *we dont need ingredient and cost price for products in app that shown on. On the shop*. So both came off **every** row in your app's Products list - the rows on the shop, the drafts, and the hidden ones alike.

What a row shows now is the name, the unit, an optional daily limit, the price you sell at, and the category. The two things that went are **the cost per unit** - the small figure that used to print beside your sell price, worked out from the recipe's ingredients - and **the ingredient line**, the list of what goes into the product with the amounts. Both are gone from the row.

Nothing was deleted, and nothing about your recipes changed. Both are still on each product's **Edit** screen exactly where they were, where you build the recipe and where the cost is worked out. The row was only ever meant to be a view of what you sell; the recipe is yours, and it stays yours on the Edit screen.

Why this helps: the row is shorter, so more products fit on the screen at once, and the one figure left on it - the price a customer pays - is the one you are actually checking when you scan the list. If you want to see a cost, tap **Edit**.

Nothing to run. No new table, no Supabase change, no redeploy of anything. This reaches your shop and your app the next time you push, and you will see **Engine v224** on your **More** screen.

28 Sep 2026 - engine v223, ONE STANDARD PICTURE WINDOW, BIGGER AND THE SAME IN BOTH PLACES (no database step, no redeploy - one push).

You asked me for *a standard photo window that is big enough for effectively shows a photo of product for APP side and Store side*. That is exactly what this is: one picture window, one size, used in every place you meet a product picture.

The window is the same size in all three places: your shop's card, your app's Products list, and the editor where you choose the photo. Before this version each place had its own size - 96 wide on the shop and following the card's own height, and only 64 in the app - so the same picture showed three different sizes, and none of them big enough to make out the product. Now it is one window everywhere, and it is bigger than before.

Your photo still goes in whole. Nothing about last version changes. A tall photo and a wide photo both still show in full, with the window's soft cream colour in the spare space beside them. The window is the fixed thing; your photo keeps its own shape.

A product with no picture is unchanged. It simply has no window, and its name and price read across the whole card exactly as before.

One thing to know: a wider window leaves the words beside it a little less room. On a product that has both a picture and a long name, the name and price now wrap onto one or two more lines than they did. If that crowds the words more than you like once you see it, say so and I will make the window a little smaller - it is one number and a quick change.

To use the new window at its best, choose the photo again. The pictures already saved were cut to the old shape back when they were saved, so re-choosing the photo is the step that gives you one made for this window. You told me you would reload the product photo, and this is the version it belongs to.

Nothing to run. No new table, no Supabase change, no redeploy of anything. This reaches your shop and your app the next time you push, and you will see **Engine v223** on your **More** screen.

28 Sep 2026 - engine v222, YOUR PHOTO KEEPS ITS OWN SHAPE AND NOTHING IS CUT OFF (no database step, no redeploy - one push).

You said the photo should be *restored to its original wide tall ratio*, and then you told me what to do about it: *fix tall and wide*. So this version fixes both. Your photo is no longer cut to fit a shape. It now goes on the card exactly as the shape you took it in, and the picture panel stays the same size and holds the whole of it.

A tall photo and a wide photo both work now. Take a photo of a tall stack of focaccia and the whole stack shows, top to bottom. Take a wide photo across a tray and the whole tray shows, side to side. Nothing is trimmed off the top, the bottom or the sides, whichever way round your photo is.

The picture panel is the thing that stays the same, not the photo. The panel beside a product is still the same size it was, so every product's card on your shop still lines up with the next. If your photo is not exactly the panel's shape, the panel's own soft cream colour shows in the spare space beside it. That is the panel, not a mistake, and it is the trade for never cutting your photo again.

This is the same in your own app. The picture beside a product in your Products list, and the preview in the product's editor when you choose a photo, both show the whole photo the same way.

A photo you have already chosen is still the old cut. I have to be straight with you about this one: the cutting happened when the photo was saved, so the pictures already in the app are the old cropped version and I cannot put back what was trimmed off. They will show whole on the card now, which is an improvement, but they are still the old shape. To get a picture that uses this properly, open that product in the app, choose its photo again, and the new one is kept in its own shape.

Nothing to run. No new table, no Supabase change, no redeploy of anything. This reaches your shop and your app the next time you push, and you will see **Engine v222** on your **More** screen.

28 Sep 2026 - engine v221, THE PICTURE IS NOW BIGGER AND RUNS THE WHOLE HEIGHT OF THE CARD (no database step, no redeploy - one push).

Last version I made the picture a tall shape. You told me you meant *bigger - take the fullest height as possible*. You are right, and last version only changed the shape; this one changes the size.

The picture now runs the whole height of the product's card. It starts at the top of the card and carries on down to the bottom, past the price and the description and alongside the + and - buttons, so the strip beside it is never left empty. On your shop the picture is half again as wide as it was, and the full height of the card rather than a small square in the corner. It is the biggest the picture can be before your product names would start breaking onto a third line, which is the point you asked me to stop at.

Your own Products list is the same. In the app, the small picture beside a product is now bigger and stretches from the product's name down past everything written about it. A product that has a description therefore gets a taller picture than a bare one - the picture follows the row rather than being a fixed stamp, so nothing is left empty and nothing is cut off.

A new photo you choose is now stored bigger, so it stays sharp at this size. The file the app keeps is

about 10 KB instead of about 6 KB. That is still tiny - your whole app still travels in one small package to the cloud and to every customer's phone - but it means the picture will not look soft now that it is drawn so much larger. Choose a product's photo again in its editor and you get the bigger version automatically.

The pictures you already have are fine, and nothing is stretched. An old square picture is fitted into the new tall box by taking a slice from the middle, so it fills the height without being pulled out of shape. It will look right straight away. If you want a picture that uses the new size at its best, choose that product's photo again - a tall photo of the tray or the plate is what suits this shape.

That middle slice is gone as of v222 above: the picture is no longer cut to shape at all, and the panel holds the whole of it instead.

Nothing to run. No new table, no Supabase change, no redeploy of anything. This reaches your shop and your app the next time you push, and you will see **Engine v221** on your **More** screen.

28 Sep 2026 - engine v220, THE PICTURE IS NOW TALLER THAN IT IS WIDE, AND YOUR OWN LIST ARRANGES ITSELF LIKE YOUR SHOP (no database step, no redeploy - one push).

Your request: *the thumbnail photo should be taller* - and, when I asked how far it should reach, *taller for store and app*. Both places, then: the picture a customer sees on your shop, and the small picture beside the product in your own Products list.

The picture is now a tall shape rather than a square - four parts wide to five parts tall. Last version you chose a square, and a square crop of a tray or a plate cuts off the top and bottom of the food, which is usually the part worth seeing. The app crops your photo to the tall shape by itself, taken from the middle, exactly as it did before - you still choose a photo and the app does the measuring. The same crop now sits at the left of the customer's row, in your own Products list, and in the little preview in the product editor, so the picture you pick is the picture you get in all three places, and no part of it is squashed out of shape.

Your existing product pictures are not lost, and they are not stretched. The photos you have already chosen are squares, and the new box is taller; the shop shows a square photo inside the tall box by taking a slice from the middle and filling the rest of the height, so it looks right and nothing is distorted. If you want a picture that uses the full height, simply choose that product's photo again in its editor and the app stores a genuinely taller one. Nothing to do if you are happy with what is there.

Your Products screen now arranges itself the way your shop does. Your request: *in app the product should arrange like it should at shop*. This is the change you will notice most. The three lists you already have - **On the shop**, **Draft** and **Hidden** - all stay, because those are about whether a product is on sale, not about what it is. Inside *each* of the three, your products now group under your own headings, in your own order, in exactly the order a customer scrolling your shop would meet them. A product with no heading yet sits last under a plain heading called **More items**, the same word your shop uses, so what you see in the app is what a customer sees.

And a product can be moved inside its heading, by a handle. Your request: *and it should be able to swap position by a handle in same category*. Beside each product's row there is a small handle; press it and drag the row up or down, and the products in that heading re-arrange behind your finger. You asked for a handle rather than up-and-down buttons or number boxes, so a handle is what it is. The handle appears only where there is something to move - a heading holding a single product shows none, because there is nowhere for it to go - and a drag can never pull a product out of its heading into another one, which is what you asked for. The position you set is stored, not just drawn: it travels to your other phone and it is the order your shop lists in, so dragging here is the same thing as arranging your shop.

One thing that was quietly broken, found while building this, and fixed. Adding a heading, renaming one, re-ordering one or deleting one did not reach your shop on its own. The shop only caught up the next time something else published - so a heading you renamed could sit under its old name on the customer page until you happened to edit a product. Every one of those four actions now publishes on the spot. It was there in v219 and you would have met it as a heading that "would not change"; it does now, at once and on its own.

Nothing to run. No database step and no redeploy: this is the app's own screens and the shop page, and one

push is the whole of it. **Engine v220** on your **More** screen once this is pushed.

---

28 Sep 2026 - engine v219, YOUR SHOP NOW FOLLOWS YOUR OWN HEADINGS, AND EVERY PRODUCT CAN WEAR A PICTURE (no database step, no redeploy - one push). Product categories and thumbnails.

Your request: *I need a category to be setup for each products, and category be manage like how we manage ingredient. Now our product listed in store are random, when there is category existed, the list will follow category, and in a category sort my name, or other criteria. Need a thumbnail photo for each product as well, with photo size adjusted accordingly.* - and, a moment later, *one product can be with more than one category*.

There is a new screen, and it starts empty on purpose. More -> Categories is where your headings live, built the same way your ingredients are: press **New category**, give it a name, and it appears in the list. Nothing is in it until you put it there, because your shop should say what *you* sell, not what a program guessed for you.

Your shop now lists by your headings instead of in whatever order the products happened to be stored in. That is the whole point of this version. Under each heading sit the products you filed there, and the order they appear in is the order *you* set - not alphabetical, not by price, not by anything the app decided. Your own arrangement is the arrangement.

A product with no heading yet is never hidden - it is shown last. A customer scrolling to the bottom of your shop finds a plain heading called **More items** with those products under it. So if you add a loaf this evening and file it tomorrow, it is on sale tonight; it simply sits at the end until you say where it belongs.

A product can be in more than one heading, and you choose which one it is listed under - by the order you tick. In the product editor your categories appear as a list of tick boxes, and the box you tick *first* is the heading the product is drawn under. Tick a second one and the product does not vanish from the first, and it is not drawn twice either: **one product appears in exactly one place on your shop**, under its first tick. The other ticks are kept for you, as the note that says "this one is also a snack", and the product's own row in the app spells them out so you can see them at a glance. You asked for this to be made simple rather than switchable, and that is what it is: *1st category, 2nd category, 3rd category; the first one groups it.*

A heading can hold headings under it, as deep as you like. Your bakery's headings are flat - Food, Drink, Snack, Appetiser - and they can stay that way. The engine underneath is shared with the Munchies shop you are planning, where you said the shape is **For Dog, For Cat** at the top and then **Pork, Duck, Fruits & Vegetables** underneath, so that is built in now and will need no second version later.

Every product can have a picture, and it is a square. Choose a photo from your phone in the product editor and the app crops it to a tidy square by itself, taken from the middle of the picture - you do not have to resize or measure anything. Your shop shows it at the left of the product's row, and your own Products list shows it too, so you can see at a glance which products have a picture and which are still blank.

Two things to know about pictures, said plainly rather than found out later. First, **the homepage is untouched** - you asked to leave it alone, and it still carries the photographs already written into it, so a picture set here reaches your shop and your own product list but not the front page. Second, **a picture costs weight twice**: it travels inside the copy of your records that every backup and every export carries, and it is fetched again by every customer who opens your shop. That is why it is kept small - a square about 160 pixels across, in the compact photo format - which looks right on a phone and keeps your shop quick to open. A picture that is far too large, or in a format the shop cannot show, is quietly left out rather than published broken.

Moving a heading, or a product inside one, is a drag. Each row carries a small handle at its left; press that handle and drag the row up or down, and the list re-arranges behind your finger. You asked for exactly this - a drag rather than up-and-down arrows or number boxes - so that is what it is. It is the first drag in the app and there was nothing here to copy from, so it is the part most worth a second look on your own phone.

Deleting a heading is guarded, the way your ingredients and suppliers are. A heading that still has headings under it will not delete. A heading with products filed in it will not delete, and it tells you how many. An empty one simply goes. There is no undo, and the question it asks is the last thing before it happens.

Two things you would have spotted, and they are fixed rather than left to puzzle over. The row for a heading shows how many products your *shop actually draws* under it. If you have ticked a product into a later heading, the shop has moved it there, so the row now also says *"1 more ticked here, listed elsewhere"* instead of quietly counting it and leaving you to wonder why the number does not match your customer page. And a product's own row in the app now *names* the other headings you ticked instead of saying "1 more ticked", because what you want to know is *which*, not how many.

Nothing to run. No database step and no redeploy: this is the app's own screens and the shop page, and one push is the whole of it. **Engine v219** on your **More** screen once this is pushed.

---

27 Sep 2026 - engine v218, ONE CUSTOMER'S DOOR IS NOT ANOTHER CUSTOMER'S BUSINESS (no database step, no redeploy - one push). The delivery run's tracking link.

Your report: *Can the courier link be not shared as one, as when the customer track it, they become aware of the other drop off point*.

You were right, and there is no way to split it. The courier hands back *one* link for the *whole trip* - its own documents describe that link as a thing for sharing delivery information with a third party, and there is no per-drop-off link at all. So a customer cannot be given live tracking of their own box without also being handed every other doorstep on the van. The only way to stop them seeing each other is to stop giving that link to customers on a run, and that is what this version does.

A run of two or more doorsteps no longer puts the courier's link on the customers' orders. It does not reach their track page and it does not go in their WhatsApp message, so there is nothing for them to open that would show them anybody else. **Everything else they had, they still have:** the app's own tracking card keeps saying how far the delivery has got - finding a driver, on the way, collected, delivered - and it still carries the driver's name, the plate to look for at the gate, and the press that rings him. All of that comes from your own app rather than from the courier's page, so none of it can leak a second address.

What changes for them is the live map. Without the courier's link there is no moving dot for a customer to watch, and their page is as fresh as the last time you press **Check the trip** - which is exactly how it already works for any order that has no link in its tracking box. That is the price of not showing them each other, and it is worth saying out loud rather than discovering.

A single customer's own courier order is untouched - you chose this, and it is deliberate. When you book one order on its own, that trip is your bakery to that one door and nobody else is in it, so its link gives nothing away and the customer keeps the live tracking exactly as before. The same is true if you ever run a day with one doorstep on it. **The rule is: the courier's link reaches a customer only when the trip carries nothing but that customer's own doorstep.**

The card tells you which one you are about to do, before the money is spent. The booking question on the Delivery screen now says which of the two it is: on a run it says the courier's link is deliberately kept off the customers' orders and why; on a single doorstep it says the customer's tracking box takes the link. And the note afterwards agrees with it instead of claiming everyone now shares one link. **Your own card still shows you the link** on every trip, run or not, because that is the one you use to check the trip and to call it off.

One more line promised the old thing, and it was found by driving a real run rather than by reading the code. Under the prices, the Delivery run screen's own note still read that the trip had *one share link that goes on every customer's own track card and message* - true until this version, and the exact opposite of what this version does. Nothing was asserting that line, so it had survived the first pass of this work; pricing a two-customer run in a browser is what caught it saying one thing while the booking did another. It now says what actually happens, and a test names that line by its own words, so the screen and the booking cannot drift apart again.

Nothing to run. No database step and no redeploy: this is the app's own screens, and one push is the whole of it. **Engine v218** on your **More** screen once this is pushed.

27 Sep 2026 - engine v217, A BUTTON THAT COULD GO QUIETLY DEAD CANNOT ANY MORE (no database step, no redeploy - one push). The courier price and booking presses.

Your report: *the get price from lalamove not responding*.

What I found, and what I could not. I drove both screens here and could not make the press go dead - the **Delivery** run's **[Price this run with Lalamove]** answered with a real price, and the order card's press answered too (it said the courier had refused, which is what this machine's test login gets). So this is not a version I can tell you *was* your fault - it is the one shape of fault that leaves no trace on the screen, and I found it by reading the code rather than by watching it happen.

The shape of it. Every press in those two screens carried its own little flag - it raised the flag before it asked Lalamove anything, and lowered it again on each way out that it knew about. Every throw that happened in between left the flag up. And a raised flag makes the button return **every** later press instantly, with **nothing said at all**, on a button still drawn grey. That is a dead control: it looks pressable, it is pressable, and it will never do anything again for the life of that card. It is also exactly what *not responding* looks like from your side. The app's own rule for this is written in the courier code: an error thrown at a button in a card, while you are standing in a kitchen, is *a dead screen with no words on it* - and the flag was the one place that rule was not kept.

What happens now, every time, on all eight presses. The flag is raised and lowered by one shared piece of code, and the lowering happens in its **last** step rather than on a list of happy endings - so **there is no way out of a press that leaves it up**. The button always comes back. And a throw is **said**, in a sentence, in the same place the price would have been: *The price could not be asked for, and nothing has been priced - [what went wrong].* Nothing is ever swallowed, and no press is ever swallowed again. This covers **Get a price from Lalamove**, **Get a price / Price this run**, **Book this trip**, **Book this run**, **Check**, **Call it off**, and the one-press address lookup - on the order card **and** on the Delivery run.

The one sentence that is deliberately not a promise. A booking that throws after the request left but before the reply came back cannot tell whether the trip was booked. That one says *The booking could not be finished - check the trip in Lalamove before pressing again, in case it went through.* It does **not** say nothing was booked, because it cannot know that, and telling you the money did not move when it might have is worse than telling you to look. An **unbooked** trip still says so plainly, as before.

Nothing to run. No database step and no redeploy: this is the app's own screens, and one push is the whole of it. **Engine v217** on your **More** screen once this is pushed.

---

27 Sep 2026 - engine v216, THE CHARGE BOX OPENS ON THE CUSTOMER, AND A GREYED BUTTON FINALLY SAYS WHY (no database step, no redeploy - one push). Two days after v215.

Two things you asked for, in your own words: *Can you default the who paid courier to The customer paid it?* and *now the greyed out book button*.

1. Who paid the courier now opens on The customer paid it. Open an order's **Note / tracking** card - or its **Edit** form - and the question **Who paid the courier** is already answered: **The customer paid it**. Type the courier's fee and it is theirs: the line under the box moves straight away, *The customer owes RM 38.00 - items total RM 30.00 + courier charge RM 8.00*, and **Save** writes it in one press. Before this, that same amount went in with **nobody named as the payer**, which is what v215 had to refuse - so this version removes the step you were being asked to take every time.

It is a default, and never an overwrite. An order that *already* records a payer opens showing the answer that order carries - press **Save** without touching it and nothing about it changes. **Not recorded** is still in the list, still yours to choose, and still how a charge is **deleted** (clearing the payer takes the amount with it, the rule you set on 19 September 2026). And the v215 refusal is still there for the one press that needs it: take the payer back to **Not recorded** with an amount in the box and press **Save**, and the app refuses in words rather than dropping it.

The delivery run is deliberately left alone. On the **Delivery** screen the same question looks identical, but there it is read **before** the amounts and decides **which amounts they are**: settling on *The customer paid it* makes the

app price every doorstep on its own and write a charge onto **every** customer's bill. That is a decision about money, so on that screen it stays a choice **you** make - the run still opens on **Not recorded**, and says so.

2. A greyed-out [Book this trip] now says why. You reported the button drawn grey with nothing to explain it. Each price row used to work out its own, shorter list of reasons - a price that had run out of time, no price id, fewer than two doors - so a price the courier's own reader refused for any **other** reason was drawn inert and **said nothing at all**. The reason is now asked of the one place that really refuses the booking, and printed once under the prices, in the courier's own words. The commonest one, and the one you are most likely to meet: *That price did not come back with the courier's own handle for each door, so it cannot be booked. Ask for a fresh price - and if this keeps happening, the app is reading the courier's reply wrongly, which is worth fixing rather than working around.* The same fix is on the **Delivery** run's **[Book this run]**, and the reason is drawn at the moment the row appears rather than a second later.

Nothing to run. No database step and no redeploy: this is the app's own screens, and one push is the whole of it. **Engine v216** on your **More** screen once this is pushed.

---

27 Sep 2026 - engine v215, THE COURIER CHARGE SAVES (no database step, no redeploy - one push). The order's charge box.

Your report: *the selected courier charges cannot save*. It was true, and this version fixes it.

What was happening. You pressed **Get a delivery price**, then **Use this fee**, and the courier's price appeared in the **Courier charge** box on the order. You pressed **Save**. The app said *Order updated* - and the charge was gone. The amount had been written into the box with **nobody named as the payer**, and a charge is the amount *and* who bore it. With no payer, the app read the whole thing as no charge at all, quietly dropped it, and still told you it had saved. A save that says it wrote something it did not write is the worst kind of fault, because there is nothing on screen to tell you to look again.

What happens now instead of that. Press **Save** with an amount in the box and no payer chosen, and the app **refuses, in words**: it tells you the charge is RM 12.50 but nobody is down as the payer, so it would not be saved, and it gives you the two ways out - choose who paid the courier under the box, or clear the amount. **Nothing is written, and your card is left exactly as you left it** - the amount still in the box, the note and the tracking number untouched. Choose a payer and press **Save** again, and it saves.

One thing this version is careful not to break. Clearing the payer is **your own way of deleting a charge** - the rule you set on 19 September 2026. So the refusal only applies to an amount that arrived while the card was open (from the courier's **Use this fee**, or typed by you). An order that already has a charge, whose payer you set back to *Not recorded* and then save, is still **deleted**, exactly as before. Two older tests pin that, and this version had to be narrowed until they passed again.

Both doors, one answer. The same refusal guards the **Edit** form's **Save changes**, in the same words - so a charge cannot be lost through one door after being protected at the other.

And the question you asked, answered. Booking a trip is not what saves the charge. The courier's charge saves on its own, with no booking at all - that was the whole of the bug. **[Get a delivery price]** asks the courier what it would charge and costs nothing. **[Use this fee]** copies that price into the charge box so you can save it. **[Book this trip]** is the one press that **sends a real driver to the bakery and spends real money** - press it when you are ready for the driver to come, and it is also what puts the trip's share link into the customer's tracking box. The charge is yours to record whenever you like; the booking is the dispatch.

Booking in advance: the app cannot do it yet. The time box on the card changes **the price only** - it does not schedule the trip, and the card now says so on itself. A booking is a driver sent now, so do not press **Book this trip** until you want the driver to arrive. When a booked trip can be scheduled for a later time, that becomes its own setting, and the card will say so then.

Nothing to run. No database step and no redeploy: this is the app's own screens, and one push is the whole of it. **Engine v215** on your **More** screen once this is pushed.

27 Sep 2026 - engine v214, THE ADDRESS IT FOUND IS WRITTEN IN FOR THEM (no database step, no redeploy - one push). The shop page.

Your words, once the lookup started landing on the right house: *the address is very accurate, it can go into the delivery address instead of customer type full*. So now it does.

What a customer sees. They type their address on the shop page. The list of doors appears underneath, as it has since v202. They tap the one that is their house - and the **Delivery address** box is filled in with the complete address the lookup found, in place of the part they had typed. Where they wrote *23 Jalan Seang Tek, Penang*, the box now reads the map service's own answer, house number, town, postcode and all. The pin still moves to that door, the map still comes to it, and everything else on the page is exactly as it was.

The house number is what decides it, and nothing else. The box is only written when the row they tapped **holds every number they typed** - so the address being written in is provably an answer about the door they meant, rather than a nearby one. If the row only reached the road, or a different town, or has one of their numbers but not the other, **their own words are left exactly as they are and only the pin moves.** That is deliberate, and it is the whole of the safety here.

Why this was refused for so long, and why that was right. Back at v205 you saw a tapped suggestion put its own words on your screen, and your words for it were that it *will contaminate the customer keyin address*. You were right: the free map services answer a Malaysian address with the **road**, and they throw the house number away. Writing a street and a town into the box would have replaced the customer's one complete address with a worse one. **What changed is not the rule but the service** - since v212 the lookup asks Google first, and Google's answer is a complete address, often with details the customer left out. So the rule you first refused is now safe, and the test for it is exact: the house number.

One consequence worth knowing. A customer who types an address with **no house number in it** - *Taman Sri Nibong, Penang* - can never have their words written over, because there is no number for the row to have found and so no way to tell a row that matched from one that guessed. Their words stand and the pin moves.

And the one thing that must not change, does not. The address on the order and the pin's name are still **one string** - the v205 promise. Whatever ends up in the box is what names the pin, so your screen still cannot show you two place names and have them disagree.

Nothing else moved. A customer who never taps a suggestion is unaffected. Editing the address afterwards still takes the pin, and still says so. **Use my location** and **Pin on the map** are untouched. The bakery side is untouched. No address the lookup could not match is written anywhere.

Nothing to run. No database step and no redeploy: this is the shop page's own code, and one push is the whole of it. **Engine v214** on your **More** screen once this is pushed.

---

27 Sep 2026 - engine v213, ASK THAT ADDRESS UP AGAIN (no database step, no redeploy - one push). The door card.

This is the missing half of v212, and without it you would have set up your Google key and seen nothing change.

The thing that was wrong. Once the app has worked out where a customer's door is, it keeps that answer and never asks again. Ask for a price a month later and it reads the saved point back and prices that. So every customer you pinned *before* v212 is holding a road-level point that nothing will ever revisit - and the better lookup v212 switched on only ever runs for an address the app has never looked up at all. Your key would have been set, working, and invisible.

So there is now one more press on the door card: "Look this address up again". It sits next to **Move this pin**, under the map. Press it and the app asks that address up a fresh, exactly the way it did the first time, and moves the dot to the answer it gets now.

It is a press, and never something the app does on its own. A pin that moves under your finger while you are reading the card is the exact fault v209 was written to stop, and this must not bring it back. Nothing here runs by

itself.

It appears only where the app itself put the door. One press is offered where a lookup decided that point, and where the door was saved before the app began keeping a note of how it was made - because before that, a lookup was the only thing that ever wrote one. **It is not offered over a pin you placed yourself**, by dragging or by picking on the map: that is your correction, not a guess, and the app will not offer to hand your work back to a service. And **it is not offered over the customer's own pin** from the shop page: they were standing at their door when they dropped it, and the app never looked that one up in the first place.

Every press tells you what it found, and there are three answers. It moved: the dot is on the new point, and the line says *The door moved - the lookup answers this address with a different point now.* It did not move: the line says it found the same spot, and names the spot it is standing on. It could not be asked: the reason is printed in the service's own words, followed by *The door has been left as it was.* A press that says nothing at all would read as a broken button, so there is no such press.

Prices follow the same rule they always have. If the press moves the door, the prices on screen were quoted for the old point and they go, with the line telling you to ask again. If the press finds the same spot, nothing moved, so the prices stand - a re-lookup that finds the same road does not make you pay for a quote twice.

And the road warning from v211 keeps its own rule. If the fresh lookup still cannot find the house number, the door keeps the warning line and the press says so in the same breath - *The door moved, and it is still only the road, not number 23.* A point that has been asked about twice is not on that account a better point.

Nothing to run. No database step and no redeploy: this is all inside the app, and one push is the whole of it. **Engine v213** on your **More** screen once this is pushed.

---

27 Sep 2026 - engine v212, THE PIN CAN NOW LAND ON THE HOUSE (no database step - but this one needs two commands from you before it does anything).

Your words: *the pin dont pin to the right house.* You were right, and v211 only put a warning on it. This version fixes the cause. It also comes with a small job for you, so please read the last part of this note.

Why the pin could not reach the house, and why no amount of work in the app would have fixed it. The two map services the app asks are free and need no account. Both of them are built from the same worldwide map, and that map holds Malaysian **roads**. It holds very few house numbers. So when you typed **23 Jalan Seang Tek**, the service was not being difficult - the number 23 is simply not written down in the file it reads. The best honest answer it could give was the middle of Jalan Seang Tek. Nothing in the app could turn that into your customer's gate.

So the app now asks a service that does have the house numbers. Google keeps house-level positions in Malaysia. When a key for it is set up, the app asks Google *first*, and the two free services stay behind it as the fallback. That is the whole of the change: the answer you get is now the house rather than the road.

The map you see and drag does not change at all. It is the same map, on the same screen, with the same pin, exactly as it is today. What is being replaced is not the map - it is the quiet lookup, the thing that decides where the dot lands when you type an address. You chose to leave the map alone, and it is left alone.

This is asleep until you do your part. With no key set, the app behaves exactly as it did at v211 - nothing better, nothing worse, and nothing broken. Setting the key is what switches the better lookup on. You can do it today or in a month; nothing in between is half-on.

Your part, in a terminal, about ten minutes, once. Never type a key into our chat, and never put one in a file. There are four small steps.

Make the key. At console.cloud.google.com, make a Google Cloud account and turn on the **Geocoding API**.

Restrict it. Create an API key, and then restrict it so a leaked key can do nothing else: under **API restrictions**, choose **Restrict key** and tick only **Geocoding API**. Google will ask for a card on file. The first 10,000 lookups each month are free, and a bakery this size will not come close to that.

Cap it. Set a **quota cap** on the key - a daily limit you are comfortable with. This is the important one. Past the

cap Google simply stops answering, and the app falls back to the free map services, so the pin goes to the road instead of the house. It cannot turn into a bill you did not agree to.

Set it and build it. In the terminal, in this folder, run the three lines at the bottom of this note - one to give the app the key, and two to build the two functions that use it.

What is behind the key, so the "card on file" sentence is not doing all the work. The key lives on the server, in the same place your Lalamove key lives. It is never put on the shop page, never in a file the browser downloads, and no phone ever holds it. A customer's phone only ever receives the finished answer. And the shop's own lookup has a rate limit in front of it already, so one person cannot spend your allowance in a loop.

If Google ever refuses - a spent cap, a mistyped key - the lookup does not stop and does not break. It falls back to the free services and you get a road-level pin with v211's warning line under it, which is exactly the behaviour you have today. Failing quietly back to what already worked was the whole design.

One genuine bug fixed on the way. One of the free services does sometimes send a house number, and the app was throwing it away - labelling a result that HAD found number 23 as **Jalan Seang Tek**, the road's name. So a correct answer looked like the road, and v211's own warning fired on it, telling you the pin was on the road when it was on the house. The number is kept now. This one is fixed whether or not you ever set up the key.

The commands (the = in the first one is required; without it the tool refuses):

supabase secrets set GOOGLE_GEOCODING_KEY=<the key> --project-ref hzpyblqygnntixkijeem supabase functions deploy shop-geocode --project-ref hzpyblqygnntixkijeem supabase functions deploy courier --project-ref hzpyblqygnntixkijeem

That is three lines, and the last two are two separate deploys on purpose: the lookup lives in both the shop's function and your app's function, and they are built separately, so doing only one would leave the other still asking the free services.

No database step. Engine v212 on your **More** screen once this is pushed.

---

27 Sep 2026 - engine v211, A LOOKUP THAT CAN ONLY FIND THE ROAD NOW SAYS SO (no database step, no redeploy - one push). The doorstep lookup, and the lines that report it.

Your words: *make it say this is the road not the house.* It does now.

Here is the thing you were seeing. A Malaysian address leads with its house number. The free map services the app asks - the ones that need no key and no bill - very often can only answer down to the street. So when you typed **23 Jalan Seang Tek**, the answer came back as **Seang Tek Road, George Town, 10400**: your street, named as your street. The pin landed on the road. And the app reported that answer in exactly the same voice it uses when it has found the actual house, so there was no way for you to tell the two apart.

From now on the line says which one it is. Where you type an address with a number and the geocoder's own answer does not contain that number, the line adds, in plain words: *The lookup found the road, not number 23 - drag the pin to the door.* Where the answer *does* contain your number, the house was found and the line stays quiet - no needless worrying.

It shows in all three places that report a lookup. The line under the pin window's **Look it up** press, whether it found one match or several. The card's own line under **The door the driver is sent to**, right where you are already reading where the dot is. And each row of the **Price this run** list, so a run you price in one press tells you the same thing about the same door.

And it is remembered, so it keeps being said. The number the lookup missed is written down with the door itself, not just printed once - so the caveat is still there on the next repaint, and on the next phone that syncs, and it survives the one-press switch from one door to the other. A door you place by hand or drag yourself carries no caveat at all, because no lookup wrote it.

It stays quiet where it should. An address you type with no number in it gets no such line - there is no house to have missed. And a postcode is not mistaken for a house number, so typing 10400 does not make the app look for a house called 10400. The comparison is by whole number, never by a piece of one: the postcode

10400 can never stand in for number 23, and 123 is a different house from 23.

Nothing is blocked, and nothing was gated. The point is still placed, the price is still asked for, the trip can still be booked - only the words changed, exactly as a warning should work. And when the lookup can only reach the road, the box for **Coordinates, if you have them** is right there: paste a pair of numbers or a Google Maps link and the pin goes exactly where you put it.

Nothing to run, nothing to set up. No database step and no deploy command. **Engine v211** on your **More** screen once this is pushed.

---

27 Sep 2026 - engine v210, THE WINDOW THAT OPENS WHEN THE MAP CANNOT (no database step, no redeploy - one push). The pin window on the courier card.

Small, but it is the one screen you fall back to when a map will not load - so it mattered.

On a courier order there is a press under **The door the driver is sent to**. When a map is on the card, that press unlocks the pin so you can drag it. When there is no map at all - its tiles never arrived, or there is no point to drag yet - the press opens a small window instead, where you type an address, look it up, and place the pin by its numbers. That window is the only way left to pin a door on a phone that cannot hold a map, and it had two things wrong.

It opened standing on **the door you keep for that customer**, even where the door in force was the customer's own pin from the shop page. And its sentence promised to look the address up - the one thing a customer who pinned their own door has already answered. So on exactly the orders it was most likely to be used on, the window stood on the wrong point under the wrong words.

There was one more wrong word on it, found while measuring this in a browser rather than in a test. Opened standing on a door, the line above the Keep still read *No spot chosen yet* until the window's own map happened to arrive - so on a phone whose map never comes, the line and the button beside it disagreed about the same point.

It now opens on the door that is actually in force, and says which one that is - from the first moment, map or no map. Where the door is the customer's own pin, the window says so; where their written address is all you have, it offers the lookup exactly as before. Keeping what you see writes whichever of the two you were standing on, and - like every press you make - it is recorded as your own hand, so it stays put.

The delivery run had the same shape of mistake, and it is gone. The window reached from a run's **Put it on the map** carried two sentences, but only one could ever be shown: that press is drawn only when an order has no door at all, and a customer who left a pin has a door. The unreachable sentence is removed; what you read there is unchanged.

Nothing to run, nothing to set up. No database step and no deploy command. **Engine v210** on your **More** screen once this is pushed.

---

27 Sep 2026 - engine v209, THE PIN THE CUSTOMER DROPPED IS NOW THE DOOR - ALWAYS (no database step, no redeploy - one push). The dot on an order, and the door a driver is sent to.

You told me where to look, and then what was wrong: *my app, the order; the dot is in the wrong place; a different town or area.* And when I asked you which door to trust when both exist, you said: *their own pin - always.* That is what this version is.

Here is what was happening. Every order could carry two doors: the pin the customer tapped on the map when they placed the order, and a door your app had kept for that customer - filled in, on an earlier version, by looking their written address up on a map.

The door your app had kept was the one in charge, so the customer's own pin was ignored. And a map lookup for a Malaysian house number can only answer down to the street - or, worse, the wrong town altogether. That is why your dot sat somewhere else.

From now on, when a customer tapped their own door, that pin IS the door. It is the point a delivery price is

asked for and the point a driver is booked to. You were standing nowhere near their house; they were standing at their own door when they tapped. Your app's own lookup is what no longer wins.

One exception, so you are never trapped. A pin *you* placed or dragged yourself still wins, and stays exactly where you put it. If you had already corrected a doorstep by hand, this version will not move it back.

And the card tells you which door it is using. Under **The door the driver is sent to**, the line now reads either "*<the address> - the door you keep for <name>*" or "*<the address> - <name>'s own pin from the shop page. This is the door the driver is sent to.*" When the two doors disagree, one press underneath switches to the other one - whichever of the two is not in use - so you can always overrule this with one tap.

This also applies to the delivery run and to the price. Pressing **Get a delivery price**, or **Price this run with ...**, prices the door on the screen, and a booked trip sends the driver there. Nothing to set up: an order whose customer left no pin simply behaves as it always has.

Nothing to run, nothing to set up. No database step and no deploy command. The engine on your **More** screen reads **Engine v209** once this is pushed.

---

27 Sep 2026 - engine v208, THE DOT NO LONGER SLIDES OFF THE CUSTOMER'S OWN DOOR WHEN YOU ASK FOR A PRICE (no database step, no redeploy - one push). The pin on the courier card.

Your report, the fourth time: the pin still wrong. When you told me where, it was **the order in your app**, and what was wrong was **the dot is in the wrong place** - not its name. This is that.

The courier card opens with the dot on the customer's *own* pin - the spot they tapped on the map when they placed the order. Pressing **Get a delivery price** then looked the written address up and moved the dot onto whatever the map's lookup found for it. But a lookup, for a Malaysian house number, can only ever answer down to the street - it has no house numbers. So the dot slid off the customer's own door and out onto the road, and the price that came back was a price for the road.

A customer who pinned their own door has already answered where the door is. Their pin is now the door, and asking for a price prices the door on screen and spends no lookup at all. Only a customer who left no pin of their own has their address looked up - and then, as at v207, the point is the map's and the words are the address on the order.

The same is fixed on the delivery run. Pressing **Price this run with ...** no longer looks up every customer in the list: each customer who left a pin of their own is priced at their own pin, and only the others are looked up. On that screen too, a looked-up door is now named with the order's address instead of the map's row - a fix v207 had missed on this one screen.

Nothing to run, nothing to set up. No database step and no deploy command. The engine on your **More** screen reads **Engine v208** once this is pushed.

---

27 Sep 2026 - engine v207, THE DOOR IS NOW CALLED BY ITS ADDRESS, NOT BY HALF AN ADDRESS FROM THE MAP (no database step, no redeploy - one push). The pin on the courier card.

Your report, three times: the pin still wrong. Read on the card, the one door was getting *two names*, and they did not agree:

12 Jalan Bunga, 10450 Penang - the door you keep for Mei Ling: Taman Sri Nibong, George Town.

The first is the address you and the customer both use. The second was the name the map's lookup had left on that door the first time it was looked up - and a lookup's answer is only a row off the map: a street and a town, no house number. So one door was being called by its address and by half of a different one.

A door is now named with the address it was looked up for. When your app looks the address up on its way to a delivery price, it keeps the *point* the map found and the *words* from the order - your address - instead of the map's row. The same is true when you press "Look it up" in the pin window and keep the spot: the pin moves to the map's point, and the door is named with the address you have. Where an order carries no address at all, the words the door already has still stand - it is never left as two bare numbers.

Old orders are fixed too, without touching a single record of yours: because the card now says the door with the address on the order, a door that was saved with a map's row for a name reads correctly from the moment this version is on your phone.

Nothing to run, nothing to set up. No database step and no deploy command. The engine on your **More** screen reads **Engine v207** once this is pushed.

---

27 Sep 2026 - engine v206, THE MAILING ADDRESS NO LONGER ASKS TO BE TYPED ON EVERY PHONE (no database step, no redeploy - one push). One sentence on More > Settings was telling you something that stopped being true at v200.

The card under **Mailing labels (courier)** used to end with "Type it on each phone you print labels from." Since v200 the mailing address is one of the settings that travel, so it is typed once and your other phones have it the next time they open - which is what the manual has said all along. The card now says so: *Saved with your settings, so it travels to your other phones - type it once.*

Nothing else changed, and nothing you have typed is touched. This is a sentence, not a behaviour: the address has been travelling since v200 and still does.

Nothing to run, nothing to set up. No database step and no deploy command. The engine on your app's More screen will read 206 once the phone has the new build.

27 Sep 2026 - engine v205, THE PIN IS NAMED WITH THE CUSTOMER'S OWN ADDRESS (no database step, no redeploy - one push). Version 204 gave a dropped pin the wording of the suggestion row the customer tapped. That was the wrong half: those rows are pieces of an address, not addresses. From now the pin is named with the address the customer *typed themselves*, and the row's wording is used for nothing except the row on their own screen.

What you said, and it settled it. "the customer know their address well, when i tap the address the address is not a complete one, if it is plaste into the address line, it will contaminate the customer keyin address." You are right, and it is worth saying exactly why: a suggestion row reads like "Taman Sri Nibong, George Town" - *a street and a town, and no house number*. It is a good enough answer to "roughly where is this" and a hopeless answer to "which door". Putting it on the order was giving a fragment the standing of a complete address.

What now happens.

- **The pin travels named with the customer's own typed address**, word for word as they wrote it. So the address on the order and the name of the pin are the same string, and there is nothing left that can say a different place. **The suggestion row's wording never leaves the shop page.** It is printed in the list they tap and nowhere else - not in the pin, not on the order, not on your screen. **A tap on a suggestion still never writes into their address box.** Their words stay exactly as they typed them. A suggestion moves the pin; it does not edit anybody's address. **On your card the address is printed once, not twice.** Because the pin now carries the same words as the address, the card's line about their own pin adds nothing when the two match - and it still says the thing that matters, that this is *their* pin and not yet the door the driver is sent to. A pin from an older order, which does carry a name of its own, still shows that name, so you never lose the ability to see where a point was meant to be. **A pin with no address to name it goes as the bare point it is.** A customer who pins without writing anything has nothing to name the spot with, and the order stays exactly what it has always been.

One consequence worth knowing. The pin's words are the customer's, taken at the moment they ordered, so they do not follow a later edit you make to that order's address. The pin stays where they drew it and keeps saying what they called it, which is the honest thing for it to do.

Nothing to run, nothing to set up. No database step and no deploy command: the change is in the shop page and the app itself, and it arrives with the push. The engine on your app's More screen will read 205 once the phone has the new build.

27 Sep 2026 - engine v204, THE PIN AND THE ADDRESS CAN NO LONGER DISAGREE (no database step, no redeploy - one push). A customer's dropped pin used to arrive as two bare numbers with

nothing tying it to the address written beside it, so the two could say different places and your screen had no way to show you. That is fixed, and the one sequence that actually produced it - taking a door from the suggestion list and then editing the address - now takes the pin back with the address instead of leaving it behind.

What you were seeing. "When the pin arrive at backoffice, it did not tally": the order carried an address that said one place and a pin that sat somewhere else entirely. Both were shown, and they contradicted each other.

Two things were wrong, and only one of them was visible.

The first: **the words the customer tapped were being thrown away.** When they tap one of the address suggestions, the row they tapped says which door it is - "Taman Sri Nibong, George Town" - and that wording was dropped one line before the order was sent. What reached your app was "5.33250, 100.30204" and nothing else. A pin of two numbers cannot be checked against an address by eye, so the app had no way to tell you the pin and the address were in different places. A pin taken from the list now travels with the words it was found for, so the two can be read side by side.

The second, and this is the one that produced the order you saw: **the pin was never told that the address had changed.** A customer types "Taman Sri Nibong, Penang", taps the suggestion, and then edits the box to "Bayan Lepas, Penang". Measured on the live page, that gave an order with Bayan Lepas written on it and a pin **5.8 km away** at Taman Sri Nibong. Two answers to one question, written at two different moments, with nothing joining them. Nothing on your screen could show it, because the pin had no words to compare.

What now happens.

- **Changing the address takes back a pin that came from the suggestion list**, and says so in a line above the list in the customer's own language: the words they tapped no longer go with what is in the box, so tap a suggestion above or set the pin again. The list they were reading stays where it is until its replacement arrives - it is not yanked out from under their thumb. **A pin they placed with their own hand is left alone.** Dragged on the map, or taken from "Use my location": that is a door they chose themselves and not an answer to what they typed, and no edit to the address box touches it. **A suggestion row tapped a moment too late is refused out loud.** The list is deliberately left up while they keep typing, so for the length of the lookup's own pause a row drawn for the old address is still tappable. That tap is now refused with its own sentence, and the rows it came from are taken away with it. A tap that does nothing at all is the dead-control fault this shop has a standing rule about, so it always says why.

One more thing, on the customer's confirmation card. When no product has a change/cancel window, the card has one line fewer - and that missing line was printing as the word **null** between two sentences. It prints as nothing now. This was reachable on your live storefront: none of your products sets a window, so it was on every order's receipt. Your cancellation wording in More -> Settings is a different thing and is untouched.

Nothing to run, nothing to set up. No database step and no deploy command: the change is in the shop page and the app itself, and it arrives with the push. The engine on your app's More screen will read 204 once the phone has the new build.

26 Sep 2026 - engine v203, AN ADDRESS WITH A UNIT NUMBER IS NOW FOUND (no database step, no redeploy - one push). Some shop addresses the lookup used to give up on now bring the map to the door. There is nothing for you to run and nothing to set up: this is a change to the shop's own page, and it is live the moment you push.

What was going wrong, exactly. The lookup matches the words a customer typed against the words the map service holds, and a Malaysian address carries two kinds of word that service does not have.

The first is the part that says WHICH door: "No 5", "Blk A", "Lot 1234", "Tkt 3", "Mukim 12" - unit numbers, block letters, lot numbers, floors and sub-districts. The map service holds STREETS and BUILDINGS, and those are neither. Worse than being no help, they crowd out the street that IS in the map, so a perfectly good address came back empty. This was measured on real Penang addresses rather than reasoned about, and every one of these four pairs repeated every single time:

Taman Sri Nibong, 11900 Bayan Lepas, Penang found 1 door 12-3-4 Blk A, Taman Sri Nibong, 11900 Bayan Lepas found nothing

Pangsapuri Sri Indah, Penang found 5 doors Blk 12-3-4, Pangsapuri Sri Indah, Penang found nothing

Lorong Seri Nibong 3, 11900 Bayan Lepas found 4 doors No 5 Lorong Seri Nibong 3, 11900 Bayan Lepas found nothing

Jalan Teluk Kumbar, Penang found 2 doors Lot 1234, Mukim 12, Jalan Teluk Kumbar, Penang found nothing

The second is shorthand. "Rd" for Road, "Jln" for Jalan, "Tmn" for Taman. The map service holds one spelling and not the other, so the same street written the short way came back empty too.

What now happens. When the lookup finds nothing, it asks once more - with the unit number and the block letter taken out and the shorthand spelled out - and only when it found nothing. Checked against the live map services, those four addresses now come back with one, five, four and two doors instead of an empty list.

The customer's own words are always asked first, and this is the part that matters. The exact address a customer types is the first thing the page sends, every time. So an address that already worked is not slowed by so much as a moment and cannot come back with a different door. The forgiving wording is only ever spent on a question that has already come back with nothing, which is also why taking a word out can never do harm: it is only ever tried where the full address has already failed.

What it costs. A lookup that finds nothing now takes about a second longer before it settles, because it is asking a second question. An address that works is not affected at all. It also means a customer who types an address the map truly does not hold waits a moment longer to be told so - which is the honest price of the four addresses above now working.

No deploy and no database step. Unlike v202, nothing here lives on your Supabase. The change is entirely in the shop's page, so if you pushed v202 and ran its deploy command, you are already done and this simply arrives with the push.

26 Sep 2026 - engine v202, THE SHOP NOW COMES TO THE ADDRESS THE CUSTOMER TYPES (no database step - but ONE deploy command, and the feature does nothing until you run it - see the bottom of this entry). A customer on the shop page used to type their address into a plain box and then press "Pin on the map", which opened the map on the whole of Penang. They then had to drag and zoom around the island looking for their own street. You asked whether the map could bring them to their address instead, the way a ride app does. It now does.

What they see. As they type their address, a short list appears under the box with the doors they might mean - the street on one line, the town and postcode with it. They tap the right one, and the map opens already standing on that door, close enough to pick out the house, **with the pin already on it**. Then they drag or tap to fine-tune it, exactly as before, and press Keep this spot. The list holds at most five doors, only in Malaysia, and only doors it has a real point for. Once they have tapped one, the line above the list steps aside - the screen must not go on telling them to do something they have just done. The doors themselves stay where they are, so changing their mind is one more tap rather than typing the street again.

Nothing is asked until they have actually typed an address, and nothing is asked mid-word. The page waits until they pause, and it only bothers with a question worth asking - a word or two, or a postcode on its own, gets no list. It also never blocks anything: if nothing matches, or if the service having a bad day, one short sentence appears in their own language, and the map button is exactly where it was. A customer who ignores the list completely can still pin by hand the way they always could. That has not changed at all.

Their typed address now leaves the phone one more time - as your bakery, not as them. This was your own decision and it is worth stating. The address goes to your Supabase, which asks the map service. So the map service sees your bakery's name and your server, not the customer and not their phone. That is the same posture your own app already takes when it looks an address up, and it is the reason the shop could not simply ask the map service from the customer's own browser. Nothing is stored anywhere: the service is asked a question and gives an answer.

One thing to do before this works, and it is the only step. The shop page is a plain website, so the lookup has to live somewhere that can hold your bakery's name - a small function on your Supabase. Until you put it there, the list says the lookup is unavailable and the customer pins by hand, which is exactly what they did before this version. Push alone changes the shop's screen; it does not make the list work. The command is repeated at the bottom of this entry, and it is one line.

Nothing about the pin's standing has changed. A pin the customer drops - from the list or by hand - is still a SUGGESTION and never a fact. Nothing is priced from it, nothing is booked from it, and no driver is ever told it. It is still offered to you on the order's delivery-price card and on the delivery run, and it still only becomes their doorstep when you press **Use the customer's pin**. A suggestion is a suggestion however the customer arrived at it.

And the door you keep for them still wins. A lookup the shop did has no say in what you have saved against a customer. Where you already keep a door for somebody, that is the door the driver is sent to, and their new suggestion is offered beside it as it always was.

The one step, written out. In the Terminal, from the folder that holds the app:

supabase functions deploy shop-geocode --project-ref hzpyblqygnntixkijeem

The --project-ref part matters - without it the tool asks you to pick a project from a list, and that list also holds an unrelated one. There is nothing to configure and no key to paste: this function holds no secrets and can reach nothing but the two public map services. Run the line once and the list works from then on.

One thing was wrong in this version, and it is fixed. No second deploy. The first build of this feature still said "the address lookup isn't available right now" after the function was deployed and working. The reason was small and entirely in the shop's own page: the lookup was sending its key twice, and one of those two ways is a header the function was not allowed to accept. A browser refuses that before the request ever leaves, which is why the screen said the service was down when it was in fact answering perfectly. The fix is one line in the shop's page and nothing on the server, so the deploy above is still the only command - if you have already run it, you are done, and the lookup now works. It was found by asking the function directly, not by reading the code: the same call with the extra header failed and without it returned four real Penang streets.

26 Sep 2026 - engine v201, THE COURIER BOX NOW SHOWS THE DOOR (no database step, no redeploy - one push). Open an order's Note / tracking box and, for an order that goes by courier, the top of it now tells you which door the driver would be sent to: the customer, the address, and a small map with the pin on it. You can look at it without pressing anything, and you can move it if it is not the door. That is the whole of it - and it is the box you asked about, the one that used to send you to Edit for the address, which meant the one place you book a trip was the one place you could not check where the van was going.

Why this was worth doing, in one line. A courier is not given an address, it is given a point. The pin is the stop the van drives to. A pin that landed on the wrong estate is a delivery you pay for and then chase, and until now the only place you could see it was a map screen that threw away whatever you had half-typed in this box. So the door moved to the top of the box you were already in.

It is a picture until you say otherwise. The map opens read-only, which is what you asked for: a drag does nothing, a tap does nothing, and pinching does nothing. Press **Move this pin** and that same map wakes up - drag the pin, and the moment you let go it is saved against that customer, with nothing you had typed beside it thrown away. A second press (**Done moving**) puts it back to a map you cannot nudge by accident while reaching past it. If you do not want to drag a small pin with a thumb, a tap anywhere on the unlocked map drops it there instead.

If the customer pinned their own door on the shop page, that is what you see, and the box says so. It reads as their own pin, not as the door the driver is sent to - because it is not one until you make it one. Press Move and it becomes the door you keep for them, exactly as it always worked, and from then on the line reads as the door you keep for them rather than as their suggestion. Nothing they dropped reaches a driver unaccepted.

Moving the pin clears the prices that were on the screen, and tells you why. A price belongs to the door it was asked for, so a price quoted for the old one would have been a lie in the box you book from. The box says, in the line where prices appear, that the door moved and to ask again for this spot. It does not quietly re-ask on its own - asking is eight requests, and that has always stayed your tap.

And it still works when the map does not. On a phone with one bar of signal, or a day when the map service itself is having trouble, you get the words instead of the picture: the door named as before, the pin's own coordinates so you can still check them, and a line saying the map is not available right now. There is always a way to put a doorstep on the map even then, because the pin screen can take coordinates as well as an address.

An order with nothing pinned shows its address and offers to pin one. No map, because a map with no pin on it is a picture of nothing, and it would spend a fifth of the box saying so. And an order that is collected rather than delivered gets no door block at all - there is no door to check.

One thing to know about the Edit form. It gains the same block, so the customer end has one kind of control rather than two. There it sits directly under the address row rather than at the very top, because the address on that screen is a draft you are part-way through typing and is not the address the driver is sent to - so the check you can trust sits beside the real one rather than floating above a box you have not saved yet.

Nothing here can stop you taking a delivery, and that is deliberate. An order with no pin, or a pin you have never looked at, prices and books exactly as it did before. This is a way to look and a way to fix; it is not a step you have to pass.

The map stays where you leave it. If you drag the map across to look at the next street, nothing on the card will pull it back. And once you move the pin, it lands where you let go of it and stays there, rather than sliding back to the middle of the box. This was wrong in the first build of this version and was caught by dragging the map on a real render and watching the box snap back the moment anything redrew. It is the same rule the planner already follows: a redraw must never move you. The one thing that *does* move the map is the address lookup finding a door the map has not shown yet, because that is a door you would otherwise not be looking at.

Seven new tests, and ten faults put back to prove they work. The faults included the block being built inside the price fold instead of at the top of the box (where it would be invisible until you asked for a price - the exact fault you reported), a dragged pin never being saved, a dragged pin leaving the old door's prices on the screen, the map not following the pin the panel looks up on its way to a price, the block being drawn for collect orders, a missing map being left as a blank space with no word about why, the map opening already draggable, a locked map swallowing the card's scroll on a phone, the map never being measured once the card had settled, and every redraw pulling the map back to the pin. All ten were caught by the tests that name them, and each was injected, seen to fail, and removed again with the files proved byte-identical. The whole suite is 1711 tests.

26 Sep 2026 - engine v200, THE SETTINGS THAT STAYED BEHIND NOW TRAVEL (no database step, no redeploy - one push). Five things you set on one phone never reached the other: the two lists on Money - what an expense was for, and the ways you get paid - your mailing address, and the planner's people, meaning the names you give the workers and which of them the day calls. Set them on your phone, open the app on the other, and the built-in names were back, the address box was empty, and your workers were numbers again. Nothing you set was ever lost; it saved properly on the phone that set it and had no way across. All five now travel, so whichever phone you pick up has them.

Here is what was actually wrong, and it is not what it looks like. Your settings go up to the cloud as one block, and the app only puts a thing into that block if it knows to look for it. Five of them were never on that list. So they were saved on the phone, left out of everything sent up, and the other phone - never told - kept whatever it already had. It was never a failed upload and never anything you did; the app simply was not carrying them.

Set any of the five on either phone and the other phone has it the next time it opens. The two Money lists travel as the whole list, not blended together - editing them replaces them, exactly the way the rest of Settings behaves, so the two phones cannot end up with a half-and-half version of your chart. Your mailing address

reaches the other phone as the FROM block on your labels. The planner's names are the same names on every plan you make, because the app keeps them by the worker, not by the plan - name somebody once and every plan reads as people.

Emptying one is an answer and now travels as one. If you clear your chart back to nothing, that means "back to the built-in names", and that meaning now reaches the other phone instead of the old list coming straight back at you. The same goes for an address you clear, and for a worker's name you delete. Before, a list you had cleared and a list you had never touched looked identical on the way up, so the phone that still had your old list would have handed it back - the app could not tell "she cleared this" from "she has no opinion here". It can now.

And it is safe to open the app on a phone you have not set up. A phone that has never touched one of these can never delete it from the other phone: staying quiet about a setting is not an instruction to remove it. If the phone that has never seen your chart opens first, it receives yours rather than pushing its own built-in one over you. This is the same rule that stopped a newly set-up phone from emptying your saved days, and it now covers these five as well.

Still on the phone it is on, on purpose, and worth knowing. Your sign-in details and your app lock never travel - those belong to the phone in your hand. The short list of occasion names Home remembers stays put too, since it is a convenience for the phone you name them on; if you would rather it followed you as well, say so and it is a small change.

Twenty-one new tests, and ten faults put back to prove they work. The faults included each list being left out of what goes up, each one being made deletable by a phone that never set it, a cleared address going out as silence instead of as an empty one, and a name list arriving as something that would have stopped the planner drawing. All ten were caught by the tests that name them; a test that stays green while the thing it tests is broken is worth nothing, so each one was injected, seen to fail, and removed. The whole suite is 1704 tests.

25 Sep 2026 - engine v199, THE TOTAL NOW SHOWS ITS WORKING (no database step, no redeploy - one push). On the WhatsApp message at every stage of an order, and on the Track page your customer opens, the money is set out as a small addition rather than one figure on its own: what they ordered comes to this, plus the courier's charge if there is one, and it all reaches the total - with the total set apart by a blank line and shown bold, so the eye lands on the number that matters. Every message that names money now names it the same way, and every track card says the same figures in the same order as the message.

There were two different things wrong, and only one of them was the emphasis you asked about. On a message, the total was simply the next line down from the working, in the same weight as everything else, so it read as one more line of the paragraph rather than the answer. And on an order that carried no courier charge at all there was no working to read in the first place: the "Items:" line is followed by what the customer ordered - two loaves, one sandwich - and not by what it costs. So a customer with a plain self-collect order had a total arrive from nowhere and nothing to check it against. That is the one you spotted, and it was the more serious of the two.

What every message now reads, from the confirmation onwards. The goods, then what they come to, then the courier's charge if the customer is bearing one, then the total. The charge keeps the two wordings it has always had - the plain charge, and the charge the courier collects at the door, which says so and stays out of the total because the customer hands it over in person. A self-collect order with no charge skips straight from what the goods come to down to the total, so the two figures sit one above the other and agree.

One of the three messages was different from the other two, and it is now the same. The message you send when a courier order is on its way used to leave the total out completely unless the order carried a charge - so the one message a customer gets when their order is in a van was the one message that never said what they had paid, while the payment reminder for the same order always quoted it. It now carries the same lines as the rest. Nothing about where the order is, or the tracking number, has changed.

The Track page draws the same three lines, in the same order. Before, it named the courier's charge and then ran the goods and the total together onto one line joined by a dash, so the charge appeared with no clue

how it related to the figure after it. It now lists what they ordered, what that comes to, the charge if there is one, and the total underneath - the same reading order as the message, so the two cannot be read as two different sums. It also used to write the charge as "RM8.00" with no space, unlike every other figure on the site; it now matches.

Bold is done the way each place can do it. WhatsApp has no bold setting in a typed message, so the total is wrapped in a pair of asterisks, which is the one way to make a line stand out there. The Track page is a web page, so it simply styles the line instead - you will never see an asterisk on the shop side, because that would be the WhatsApp message leaking onto the page. Both were worth saying, because they are the same intention carried out two different ways.

The figures are worked out from what the app already knows, so no new column was needed. The total is the one figure the app publishes to the customer's page; the courier's charge is published beside it, and it is inside that total unless the courier is collecting it at the door. So the Track page takes the charge back off to show what the goods came to - which is why a doorstep charge does not move that figure at all, since it was never added to it. The page also reads the currency off the published total rather than assuming ringgit, so a price written in another currency cannot end up with its working-out written in RM.

Sixteen faults put back, all sixteen caught. The faults included the total quietly losing its asterisks, the blank line going missing above it, the working disappearing altogether so the old "total from nowhere" came back, and - the one most worth having - the Track page taking the courier's charge off an order the courier is collecting at the door, which would have shown a customer a subtotal of RM22 on an order they are being asked RM30 for. Also among them: the shipped message going back to naming a total only when there was a charge, and one of the translations losing the placeholder that carries its figure, which would print "Items total:" with no amount beside it. Every file was restored byte-identically and proved by sha256 rather than by counting lines. The suite is 1670 passing with none failing.

No database step, and no redeploy. This one is all in the app itself, so the push is the whole of it - nothing to run in Supabase, and the courier function is not touched. The v196 and v198 lookup changes still need the redeploy they were waiting for; this one does not add to that. After pushing, close and reopen the app on your phone.

25 Sep 2026 - engine v198, THE ADDRESS LOOKUP OFFERS ITS OTHER MATCHES (no database step; the courier function must be redeployed). When you press "Look it up" on the pin card, the card now lists the other places the lookup found, underneath the button, so the right door is a press rather than a drag. The first match still lands on the map by itself and nothing else about the flow changes - the list is there for one thing only, saying "not that one".

Every answer but the first was being thrown away, on every lookup. The address services answer with a handful of places, nearest-sounding first, and the app used to keep only the first one and drop the rest on the floor. So when the top answer was the wrong end of the street, the pin landed there, and the only way to fix it was to notice and drag it. The other answers were arriving all along.

What you do now, and it is one press more than nothing. Type or paste the address, press "Look it up", and the pin lands on the best match exactly as it always did. If the lookup found more than one place, the line under the button says so - "Found: 12 Jalan Bunga - and 3 more below" - and the other places are listed under it, each showing the street on one line and the town and postcode underneath. Press the one you want and the pin goes there.

The row the pin is on wears a tick, and the tick follows the pin rather than your press. It is not a note that you tapped a row; it is read back from where the pin actually is. So if you tap a row and then drag the pin, or paste a pair of coordinates, the tick moves with it and stops claiming a street the pin is not on. The list stays open after you press a row, so you can try one, watch the map move, and try another before you settle.

One match draws no list. A single answer is not a choice, and the line under the button already names it, so nothing is drawn under it. This is also what keeps the card behaving exactly as it does today until the courier function is redeployed.

A new lookup clears the old rows first, and a lookup that finds nothing takes them away with it. The previous answer is never left sitting under a new question.

Why the same free map service, and not Google. You asked about Google's address service, and I read its terms before building anything. It requires the results to be shown on a Google map and it forbids keeping them - and this app keeps your customers' doors for good, which is exactly what the pin card is for. The two cannot be made to agree by any amount of code, so the answer was to ask the service already in use for all five of its answers instead of one. It is built from the same open map data, it covers Malaysia, and it needs no account, no key and no card.

The suggestions arrive when you press, and not while you type, and that is deliberate. This is not the typing-ahead suggestion box you see on shop websites. A suggestion per keystroke would put several requests per word onto a free service that asks to be used lightly, and its own terms say heavy use may be throttled or refused. One press is one request - the same number of requests the app made before this change.

A quiet detail, put where a test can see it. Long addresses are split into the two lines of a row, and how to split them is decided in the app's own rules file rather than inside the card, so the split that cannot be got wrong is the one that gets tested. The honest limit: an address from the second service can read as a bare house number on the top line, with the street underneath. The coordinates are unaffected, and the list is a correction rather than the thing that places the pin.

Ten faults put back, all ten caught. The faults included pressing a row and the pin landing on the first match anyway, a single match drawing a list of one, and the old rows being left behind under a new lookup. Every file was restored byte-identically and proved by sha256 rather than by counting lines, because a one-letter difference has hidden behind a matching length here before. The suite is 1658 passing with none failing.

This one needs the courier function redeployed. The list is assembled on your server, so pushing the app's files is not enough by itself - the deploy step is what puts it live, and until then the card behaves exactly as it does today. One redeploy settles this and the v196 lookup fix together. After that, close and reopen the app on your phone.

No database step. Nothing stored is added, changed or moved.

25 Sep 2026 - engine v197, THE CUSTOMER CAN PIN THEIR OWN DOOR (no database step). Your customer, on the shop page, can now drop a pin on the door the courier should stop at, and you are offered it where you already look at a doorstep: on the order's delivery-price section, and on the delivery run under that customer's row. Their pin is a SUGGESTION and never a fact - nothing is priced, booked or sent to a driver from a pin you have not taken up yourself. An order from a customer who pins nothing is posted exactly as it always was, byte for byte.

The two ways a customer pins, because one is not enough. A customer standing at their own door presses "Use my location" and the phone hands the pin over. A customer ordering from somewhere else - the office, the bus, their mother's house - presses "Pin on the map" and drags a pin to their door, which is how everybody already drops a pin in a ride app. Where they order from has nothing to do with where the focaccia goes, which is the whole reason both exist.

Your three answers are the feature. Offered every time, even when you already keep a door for that customer: the offer is not silenced by your own door, only the words change - "pinned a different spot this time, the doorstep you keep for them is untouched until you take this one", and the press says "Use the customer's pin instead". A courier customer who pins nothing sends as it always has. And it appears in both places, the order card and the delivery run, drawn the same way in both so the same press means the same thing.

Taking it up is the only thing that makes it a door. Pressing "Use the customer's pin" writes it with the same single writer your own map pin already uses, so a pin you accepted and a pin you placed by hand are the same record on the same customer. That is also what makes it stop being offered - no "dismissed" flag is stored anywhere, because accepting is what ends it. And because a delivery run's row is one big label you tap to tick a customer, the offer sits beside the row rather than inside it: a press in there would have ticked the customer off the run instead of pinning their door.

Where to pin, for a block or a condo. The shop now says it in words, in all three languages: drop the pin where the courier can stop - the guard house or your block's entrance - and put the block and unit number in the address box above. A pin is two numbers and cannot hold a unit number, and a pin dropped inside a tower is exactly where a phone's own position goes vague. Lalamove's own Malaysian guidance says the same thing: pin the lobby, the main entrance or a landmark a driver can actually stop at, and sort a wrong pin out by talking to the driver. That is also why the typed address still travels with every order, with or without a pin.

A vague position is kept and said, never refused. If a phone reports a position that is only accurate to more than about 150 metres, the customer is told the number and pointed at the map to move it - rather than being blocked from ordering. Your own rule, applied here: the shop never stands between you and a sale, and you reconfirm every pin anyway.

Every way the location press can fail has its own sentence - a browser with no location at all, a permission refused, a position unavailable, a wait that timed out, and a map that could not load. None of them blocks the order.

The map is only fetched when somebody presses for it. A customer standing at their own front door never loads it and never pays for it, and the location press works with no map at all.

The suite is 1649 passing with none failing, and nothing of yours was written. The new coverage drives the real customer controls: a stubbed phone position, a press, and the posted order read back. One finding worth naming was caught by these tests rather than by you - a half-written pin would otherwise have become a real point on the equator, because a missing latitude reads as zero to the computer. No database step and no redeploy: the pin rides inside the order the shop already posts.

One honest limit. Your phone's own permission sheet is something only a real phone can show, so "Use my location" is proved by tests rather than by me driving it. Please try it once on your own phone - stand somewhere you know, press it, and check the words that come back.

25 Sep 2026 - engine v196, THE ADDRESS LOOKUP ASKS A SECOND SERVICE (no database step; the courier function must be redeployed). You pressed "Look it up" on the pickup-pin card and read "The address lookup service did not answer. Put the pin on the map instead." - and then read it again, a second time. The lookup now asks a second address service first, and that one answers: typing an address fills the pin in again. Nothing about your day, your chart, your orders or your saved data was touched, and putting the pin on the map by hand still works exactly as it did - it is still the thing the pin is really for.

What you read, and what it proved. That sentence is written in one place, and it appears only when the address service has answered and refused. So the news was not that your phone could not reach it, and not that the service was down: it was that the service looked at the question and turned it away. Three other sentences cover the unreachable, the too-slow and the not-found cases, and you would have read one of those instead of this one.

The service itself is healthy, and I checked rather than assumed. Asked for the same kind of request your app makes, it answered with a correct Penang result - the right street, the right town, the right postcode, first time. So the address was not the trouble, and neither was the question or the way the question is written.

Why it turned the app away. That service asks every caller to say who it is, and it turns away callers it cannot account for. The app does say who it is, and says it the way that service asks - but the request leaves through a shared machine, and whether the name survives the trip is not something the app can see. What settled it was not the wording but the repetition: an ordinary address, refused identically twice, while the same service happily answers for a different caller (me). That is a house rule being applied rather than a fault to retry, so the answer was not to ask harder but to ask somebody else as well.

The fix: a second service is asked first. The new one is built from the same open map data, covers Malaysia, needs no account and no key, and asks nothing of the caller beyond the question itself. The original service is kept and is now the second ask rather than the only one. Whichever answers first wins, so in ordinary use only one service ever sees a customer's address - and an address still never leaves your own server for the browser,

which is how this was built and how it stays.

The old service was kept rather than thrown out. It is the better-resourced index of the two, and keeping it costs nothing: if whatever turned it away is ever lifted, it is already in the queue. Losing its first place is the only thing it lost.

The four ways a lookup can fail are now told apart. Not found, refused, too slow, and unreachable each get their own words, so the sentence you read describes the problem you actually have - a house in a new estate that the map simply does not carry is not the same news as a service that refused, and neither should wear the other's words.

A quiet trap in the new service, caught by a test rather than by you. The new service writes a point's two numbers in the opposite order to every other point in this app. Read in the order it writes them, a Penang street becomes a latitude of 100 degrees, which is not a place at all - and a pin that silently never appears is a fault nobody notices until a driver is sent somewhere wrong. A test now names that order explicitly.

Eleven new tests, six faults put back, all six caught. The faults included the original fault itself - a refusal from the first service ending the whole lookup - and every one of them failed a test that named it. Every file was restored byte-identically and proved by sha256 rather than by counting lines, because a one-letter difference has hidden behind a matching length here before. The suite is 1618 passing with none failing.

This one needs the courier function redeployed. Unlike the app's own screens, this fix lives on your server, so pushing the files is not enough by itself - the deploy step is what puts it live. After that, close and reopen the app on your phone before trying the address again.

No database step. Nothing stored is added, changed or moved.

25 Sep 2026 - engine v195, THE PIN CARD DRAWS ITSELF (no database step). The screen that puts your bakery's door on the map was not drawing a card. It was printing the words "[object HTMLParagraphElement],[object HTMLDivElement],[object HTMLDivElement]..." - which is what you saw on your phone, and it is what the card had become. Every field, button and map it was supposed to hold had been built and then thrown away, so there was nothing on it to press and no way to save the pin. It draws now. Nothing about your day, your chart, your orders or your saved data was touched, and this was the same day the courier card above it began reading "Lalamove: sandbox", which is that connection working.

What you saw, and what it was. You opened Settings, pressed "Put the pickup pin on the map", and instead of a card you got a line of bracketed text. Those brackets are how JavaScript says "here is an object" when something asks it to write an object out as words. Each one was one part of the card that never reached the screen: a line of explanation, the box for your address, the map itself, the note under the map, the box for coordinates, the line reading where the pin sits, and the row with the press that keeps it. Seven parts, all of them lost to one line of code.

Why it happened, in one sentence. The app builds a pop-up's contents as a list, and then hands that list to a browser call that only accepts bare items, not a list - so the browser turned the whole list into a single piece of text instead of drawing the seven things inside it. This is a subtle trap rather than carelessness: the helper used everywhere else in the app happily takes a list, so the two look as though they should behave the same, and only one of them does.

This was the only card of the thirty-eight that did it. Every other pop-up in the app - the order edit, the product editor, the money journals, the day's availability - hands over one piece at a time, which is why none of them has ever shown you this. The pin card is the one that hands over a list, and it is the one that broke.

And it got that far because that screen had no test of any kind. Not a weak test, not an outdated test - none. It is the newest screen in the app and it was the only one nothing checked. The file that keeps this class of mistake from ever reaching a screen already existed, and it already covered three other screens that had printed stray words at you; it simply did not know about this one. It does now.

Where the fix sits, and it is deliberately not the pin card. The pin card could have been rewritten to hand its

parts over one at a time, which would have cured the symptom and left the trap in place for the next card. Instead the pop-up helper itself was taught to accept a list as well as a single piece, so the two builders in the app now behave the same way - and a test asserts both halves: that a card handing over a list gets the list drawn, and that the thirty-seven cards handing over one piece each still draw exactly as they did.

Three faults put back, all three caught, every file restored byte-identically. The original fault was reinstated, and the test that names it failed; the list-check was removed, and cards across the app failed, proving that guard is what keeps the other thirty-seven working; and the map's own failure message was silenced, and the pin card's test failed. Restoration was proved by sha256 of each file, not by counting lines, because a one-letter difference has hidden behind a matching length before. The suite is 1607 passing with none failing.

No database step. Nothing stored is added or changed, and there is nothing to run in Supabase. This is a screen-drawing fault, so it is cured by the app itself being updated - push it, then close and reopen the app on your phone before you try the pin again.

25 Sep 2026 - engine v194, THE COURIER STOPS BLAMING THE BUILD (no database step; the courier function was redeployed). The sentence that appears on the Settings card when the courier cannot be asked was false, and it pointed you at the wrong place. It said the build has no courier called Lalamove, when the truth was that no Lalamove key had been added to the server yet. Those are two different problems that until now shared one sentence, and they no longer do. Nothing about your day, your chart, your orders or your saved data was touched.

What you read, and why it was wrong. Having just deployed the courier function, you opened More -> Settings and read this: "Lalamove could not be asked: This build has no courier called lalamove." Read plainly, that says the app is missing the courier - so the hunt begins in the build, in the deploy, in the code. There is no fault in any of them, and there never was. The server knew the real reason and had already written it to its own log, and then handed you a sentence that hid it.

Where the sentence came from. One small function answers for the courier, and it builds its configuration through a single door. Two different things could shut that door: a request naming a courier this build does not carry, which is genuinely a fault in the app - and a Lalamove key that has not been put on the server, which is one secret's worth of setup. Both produced the same nothing, and one sentence was written for both. It was the sentence for the fault, so a missing key was reported to you as a broken build.

What it says now, and it is the exact line you will read. "Lalamove could not be asked: Lalamove's api key and secret have not been added to the server yet." The courier's name is still there, the thing that is missing is named, and the word build is gone from it, because the build is not the problem.

Two problems, two sentences, and a test that keeps them apart. The sentence about the key comes from the courier's own file, beside the other two sentences that courier can be refused by, so a second courier one day brings its own words rather than inheriting Lalamove's. And because the function that chooses between them cannot be loaded by the test suite at all - it begins with a line only Deno understands - a test now reads those two branches out of the file and requires each to keep its own words. Put the old single sentence back into either one and the suite goes red.

A guard of your own caught this change while it was being made. The first attempt wrote the missing key's name into a log line inside the dispatcher, and the app's existing rule - only the registry and the courier's own file may name the courier in live code - refused it. The guard was right and the line was changed to suit it, rather than the guard being widened to suit the line. That guard exists because a courier's name written by hand into a screen is the defect that stops a second courier ever being added cleanly.

One more thing corrected in the same file, and it cost real time today. The setup note at the top of the courier function told us to run supabase secrets set LALAMOVE_KEY <key>. That command is refused: the tool requires the equals sign, LALAMOVE_KEY=<key>, and answers "Invalid secret pair ... Must be NAME=VALUE." The note also now says to name the project when deploying, because the tool otherwise offers a list in which a wrong project can be chosen by a stray keystroke, and it records that the sandbox switch lives in the top right corner of Lalamove's own page.

No database step. Nothing stored is added or changed, and there is nothing to run in Supabase. The courier function itself was redeployed, which is the only thing that had to change for this to take effect. Suite is 1604 passing with none failing, and five faults were put back in one at a time - the collapse of the two sentences among them - each watched failing a test that names it and then restored byte-identically.

25 Sep 2026 - engine v193, THE WORKERS' LINE BECOMES A TRUE TIMELINE (no database step). The seven changes you asked for to the person's window on the Production line, and one question you asked about its countdown. A coach is now as long as the minutes your hand is needed on it, the three people's lines share one clock so they can be read against each other, the line starts at the left and the clock sweeps in to the middle before the line itself moves, a coach turns from green to light red when its moment has passed and eases rather than snapping, and the line is dragged with the right mouse button or with two fingers on a phone. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "1. The should start left hand side, not centred, this request I did in previous chat, but never implemented. 2. Make the coach width relative to its duration, and others batch duration, their labour requirement mark by hand needed, not batch arrival. 3. So the 3 person's train head, should be timed and position relatively to each other, when time start, 3 train started together. 4. The wait time a person interested is not the batch arriving, but their hand needed. 5. The coach should have 2 color only. Green and light red, the coach change to red only when their hand needed time is up, the change of color should be fade in and out, not sudden so that the user dont feel the screen is jumping here and there." And then: "and the drag, should be by right mouse button hold down". One more, asked after the first pass: "can the time show under their names, accurate to 5m 55s?"

2. A coach is as long as your hand is needed, and nothing else. Until this release every coach was the same width, and each person's line had its own private map from minutes to pixels, so a twelve-minute job and a forty-minute one looked identical. The width is now the minutes your hand is needed, multiplied by the day's own scale - the same figure the modules' window above it is already drawn from. Measured live at the Detail scale: a hand-needed stretch of 6 minutes drew 43.2 pixels, one of 2 minutes drew 14.4 pixels, one of 18 minutes drew 129.6, and every gap between them was an exact multiple of the same figure. A one-minute fold is therefore a hairline, which is your own answer for it.

3. Three people, one clock. Because every line is now drawn from that one scale, a minute of the morning is the same point on all three, and the three trains start together. Measured live through one whole walk: all three people's rulers stood at the identical pixel at every moment - 0 at the start of the day, then 121 pixels, then 452 - and all three strips had slid by the identical amount. Before this release each line carried its own map, so the three could not be compared at all, which is the fault you named.

4. The line starts at the left, and the clock sweeps to the middle. This is your point 1, the one you had asked for in an earlier conversation and which had never been built. At the start of the day the day's own left edge is glued to the left edge of the window and the clock stands at the left, so the whole queue ahead of you can be seen without dragging anything. Measured live on a walk: with the clock 0.7 minutes in, the ruler stood at 5 pixels and the line had not moved at all. Once the clock reaches the middle it stops there and the line begins to move instead. Measured live at the same scale: the ruler pinned at 121 pixels - the middle of the 241-pixel window - while the line slid on to 181 pixels. Your own words for the reason were that at the start a person would otherwise see only one train head, with the rest hidden until the line was dragged.

5. Two colours, and they ease rather than snap. Green means nothing is outstanding: the moment has not come yet, or you have already taken it. Light red means the moment your hand is needed has passed and nobody has taken it - the same minute the bell rings on. There are exactly two, read off the drawn coaches: green as rgb(231, 242, 232) and light red as rgb(251, 233, 228). The coach carries a 0.45-second fade on its colour, so it eases from one to the other instead of jumping, which was your worry about the screen moving under you. A tap on a coach takes it and fades it back to green, and where a coach is wide enough to hold its face it also carries a tick.

6. The coach is the hand-needed section, labelled. Your points 2 and 4 are one thing seen twice, and it was

already true in the model underneath: the section a coach is drawn from IS the stretch your hand is needed, taken from the same hands-window the day's own plan is built on, and never from the moment a batch arrives. What this release changes is that you can now see it, because the length of the coach is that stretch.

7. A coach too small for its words shows no words, not smaller words. Your answer was that if there is no room it should simply not be shown, because the same job is already named under the person's name. Measured live: coaches of 14.4 and 28.8 pixels showed no face at all, while one of 43.2 pixels showed its icon, its name, its clock and its batch. The line under each person's name still names what is coming, so nothing is lost. And a hairline coach is still a 36-pixel tap target, so even a very short job can be taken by hand.

8. The drag is the right mouse button, or two fingers on a phone. Your correction, taken literally. On a computer the line is dragged by holding the right mouse button, and a left press no longer drags at all - which is the point of it, because a left press is also the tap that says "I am on it". On a phone the gesture is two fingers, as you asked. Measured live on the drawn board: a left press with a hundred-pixel sideways move left the line exactly where it was, to the pixel; the same move with the right button carried it a hundred pixels; two fingers carried it forty. A drag also leaves behind a click, and that click is not read as a tap on the work - measured live, the press that ended a two-finger drag took nothing, while a still press on the same coach took it.

9. Scale, on the board's own row. You asked for it and there it is: the control row is three groups now, with the Scale step as the third, laid out like the planner's own so that the two screens behave alike. It moves through the app's six scales and stops at both ends rather than running past them. Measured live: stepping from the Standard scale to the Close one drew every coach half again as wide and the line from 2304 pixels to 3456, and pressing on at the top of the range left the button looking switched off and changed nothing at all.

10. The countdown reads to the second. Your question. The line under each person's name now gives seconds as well as minutes - measured live as "Next 22m 7s", "Next 160m 7s" and "Next 918m 7s" - and because the same instant is also printed on the link beside each coach, both are written by one shared piece of arithmetic, so the screen cannot state one moment two ways.

11. Both windows are kept. You chose to keep the train AND the planner's own people's row underneath it, and both are on the screen. Measured live, the three windows stand in the order train, modules, people, with all three present at the phone width and at the desk width.

12. What is measured live, and at what widths. Everything above was read off the drawn screen at 375 by 812 and at 1280 by 900, including the whole sweep of a walk: the ruler's own pixel, the line's own slide, a coach's drawn width against its minutes, the three rulers standing at the same pixel, the two colours and the fade read off the drawn element, the tap target of a hairline, and the Scale step's ends. The desk width was reached by widening the window with no reload at all, and the board re-measured itself within its own second - which is how it will behave on your own desk.

13. Nothing of yours is rewritten, and one honest note about that claim. Your stored data was read before anything was measured, and every write the measurement made was captured and never allowed to land - seven of them across the whole pass - so your phone's storage holds exactly the three keys it held before and not one byte of it changed. The note: the first copy I made of your data to compare against carried a typo that I typed, a single letter in the name of a stored list, so the file-by-file comparison could not honestly be made from that first copy. The copy was taken again from the screen itself and verified, and it does match your stored data byte for byte. But this is stated plainly rather than overclaimed: the strongest form of the claim here rests on the reading taken before the measurement began and on every write having been intercepted.

14. Every rule that stands on it was proved load-bearing. Seven faults were put back into the real files, one at a time; each was watched failing a test that names it, and every file was put back byte-identically afterwards. Among them: the moment the ruler is meant to stop at the middle, the one scale that every line must share, the minute at which a coach is allowed to turn red, and the left press being allowed to drag again. The suite is 1602 passing with none failing.

15. No database step. Not one stored field is added or changed. The train is drawn from the day's own plan and the day's own scale, both of which the app already saves, so there is nothing to run in Supabase.

CHANGELOG and the guide both carry all of this, both PDFs rebuilt and read back with PyMuPDF, and the app's own More screen reads Engine v193.

25 Sep 2026 - engine v192, THE SAVING STAYS WITH YOU, AND A CUSTOMER IS CHARGED THE ORIGINAL PRICE (no database step). A correction to v191, in your own words and taken as the rule. When a customer bears the courier charge, they are charged what their own doorstep costs on its own - never a share of the one-trip fee - so the money that consolidating saves stays with you rather than quietly becoming their discount. When you bear it, the run's real fee is your own cost and is apportioned across the orders as before. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "the benefit of consolidated charges, should go to merchant, not the customer. And if the courier charges were reveal to them, it will shown as the original cost." Two sentences and one rule. v191 split the run's one fee evenly across the orders, which was right for your own books and wrong for a customer's: a customer who was quoted RM 16.25 to their own house was being asked for RM 11.00 - the saving from the shared vehicle had been handed to them without you deciding to. The rule now: a customer who pays carries their own doorstep's standalone price; the difference between that and the trip's real fee is yours. Nothing else about v191 changed.

2. Where the money lands, and which half of the screen changed. One function decides the amounts, and it now has three answers rather than two. If the customer bears the charge, each order carries its OWN doorstep's standalone cost, exactly as Lalamove quoted it - and if the number of costs does not match the number of doorsteps, the answer is no charge at all rather than a guess, because a mismatched list means the list is about a different run. If YOU bear it, the trip's one fee is apportioned across the orders as v191 did, and the customer is charged none of it - the amount then lands on your books as a "Delivery & fuel" cost and never reaches a customer's total. If nobody has answered who paid, the answer is no charge, which is the rule the charge model has carried since v127.

3. Where those standalone costs come from, and what they cost you. They are real quotations from Lalamove, one per doorstep, asked when you press Book. There is no bulk pricing call, so two doorsteps is two requests - and Lalamove allows two requests a second, so the app spaces them and waits for every one before anything is booked. Measured live on the wire: the requests went out 602 milliseconds apart, and the confirmation appeared 633 milliseconds after the press - so the question genuinely waits rather than showing you numbers it does not have yet. If you have ALREADY pressed Compare with sending them separately, those costs are reused rather than asked for and paid for a second time; that is the same set of prices, so it is the same answer, and the screen says where they came from.

4. What the screen now says, measured on the drawn page rather than reasoned about. The line under the price list no longer talks about an even split. With the customer bearing it, it reads: "Each customer is charged what their own doorstep costs on its own, never a share of the one-trip fee - so the saving from going together stays with you. Those 2 costs come from Lalamove when you book, and the confirmation shows each of them before anything is asked for." The confirmation then does the arithmetic in front of you: "Each customer is charged what their own doorstep costs on its own, never a share of the trip: RM 13.50 - RM 16.25 - RM 29.75 in all, which is RM 7.75 more than the RM 22.00 the trip costs you - that difference stays with you." The RM 7.75 is the consolidation saving, said as a number and said as whose it is. When the same run LOSES you money - which happens when the doorsteps are close together and the trip fee is high - the sentence says so in those words rather than polishing the arithmetic into a saving that is not there: on a run built for the purpose whose standalone costs came to RM 12.00 against a RM 22.00 trip, it read "RM 12.00 in all ... RM 10.00 SHORT of the RM 22.00 ... loses you money".

5. Your own half of it, unchanged in the arithmetic and clearer in the words. With "I paid it" chosen, the line under the prices reads: "You are bearing it, so the run's own fee is your cost and the customer is charged none of it: RM 22.00 over 2 orders - RM 11.00 - RM 11.00." The confirmation says the same thing as a fact about your books: "The run's fee goes on your books as one cost of RM 22.00, over 2 orders - RM 11.00 - RM 11.00 - and the customer is charged none of it." Measured live after confirming that run: two "Delivery & fuel" rows of RM

11.00 each, the orders carrying RM 11.00 each and paid-by-you, and the customers' own track cards carrying no courier charge at all - their totals read RM 45.00, which is their items and nothing else. That is the half where the even split is correct, and it is the only half where it is.

6. The odd cents, and why the last order carries them. When a fee does not divide evenly, the apportioning puts the remainder on the LAST order rather than spreading a fraction of a cent. Measured on a fee of RM 22.00 across three orders: RM 7.33, RM 7.33 and RM 7.34, which add up to the RM 22.00 you actually paid. This matters only on your own books - it is the arithmetic that stops the three rows summing to RM 21.99 or RM 22.02 - and it is unchanged from how the charge model has always rounded.

7. What is measured live, and at what widths. Every number in sections 4 and 5 was read off the drawn screen at 375 by 812, with the world built for the purpose and taken back out afterwards. The Book press measures 110 by 36 pixels, the row 315 wide with no sideways scrolling, and the line under the price list 90 pixels tall, which fits without clipping at that width. At 1280 by 900 the two presses measure 110 by 36 and 270 by 36.

8. Nothing of yours is rewritten. The only press this release changes is the one that decides what each order's charge box is to hold, and it writes the same three fields it has written since v188 - the amount, who paid it, and whether it is collected on delivery. Measured across a whole pass - a price asked, a run priced separately, a run booked with the customer bearing it, and the same run booked again with you bearing it - the only fields that moved were the charge on those three orders and the trip itself. Your saved days and your settings were not touched.

One limit stated plainly rather than glossed: the copy of the app's stored data taken before the measurement was lost when the browser reloaded partway through, so the usual byte-for-byte comparison could not be made this time. The bakery the measurement ran in was an empty one built for the purpose and it is empty again, with nothing of yours in it at any point - but the comparison itself was not performed, and that is a weaker claim than the one every other release in this history makes.

9. One thing this release deliberately did NOT change, and you should know it. If you book a run with the payer question left unanswered, the orders on it lose any charge you had set on them by hand. That is not new - it is how v191 behaved, because the charge is written as a set of three keys and an unanswered payer is an empty answer. It is left exactly as it was rather than quietly changed under this release's heading, and it is named here so you can decide whether you want it to behave differently.

10. Every rule that stands on it was proved load-bearing. Ten faults were put back in the real source, each watched failing a test that NAMES it and then restored byte-identically. Among them: a wrong-length list of costs accepted as a charge; a customer given the even split instead of their own doorstep; your own share repeated whole on every order instead of apportioned; the first customer's cost written onto every order; the costs re-asked for even when a comparison was already on screen; and the losing run's shortfall smoothed into a saving. One of the ten is worth singling out because only ONE test catches it: taking the spacing between the two standalone requests away fails the timing assertion and nothing else, so the rate limit is guarded by the only thing that can see it. Two tests were added to the pure charge function for the answers that had no test - a list of the wrong length, and the no-payer case. The suite is 1596 passing with none failing.

11. No database step. Not one stored field is added or changed. The charge lands in the fields the order already carried, the trip in the fields v189 already writes, and an order record syncs whole - so all of it reaches your other phone with no sync change and no migration.

CHANGELOG and the guide both carry all of this, both PDFs rebuilt and read back with PyMuPDF, and the app's own More screen reads Engine v192.

25 Sep 2026 - engine v191, ONE TRIP, SEVERAL DOORSTEPS, AND WHAT IT SAVES (no database step). The fourth and last step of the courier work you asked for, and the one that saves you money. The app now groups a day's courier orders into ONE trip, prices it against the same doorsteps sent one at a time, and shows you the difference as two numbers rather than as a claim. Booking it spends one vehicle on one journey and gives every customer on it the same share link. What your shop promises

those customers becomes a WINDOW that you type, and only when you type one. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "i want to do Lalamove, API to manage courier, make it ready for other courier as well. Focus now for laalamove. And check our backoffice readiness. From my understanding lalamove support last day deliver consolidation to save on cost." v188 was the price, v189 was the booking, v190 was the customer watching the trip. This is the fourth and last, and it is the one your own hunch was about. Your consolidation belief was confirmed in v188: one trip carries two to sixteen drop-offs, charged as one base fare plus a fee for every stop after the first, which outside the Klang Valley is RM 1 on a motorcycle, RM 2 on a car, RM 5 on a van or four-by-four and RM 10 on a lorry. This release is that mechanism with a screen around it.

2. One trip, and the arithmetic you can read rather than take on trust. The run screen prices the whole run twice. Once as the one multi-stop trip it is about to become, and once as the same doorsteps sent one at a time, and then it writes the difference out as a sentence. Measured live on a two-customer run on a car: "One trip RM 22.00 - one at a time RM 44.00 - this one trip saves RM 22.00 against sending them one at a time. Each of those was asked now, for the same collection time, so they are this journey split up rather than a guess at it." That last sentence matters more than it looks: the comparison prices are REAL quotations asked at the same moment for the same collection time, not an estimate the app did on its own. The saving is therefore a number you can act on rather than a claim off a leaflet. It also costs you waiting time rather than money - there is no bulk pricing call, so a run of eight drops is still eight price requests - and the screen says so rather than making the second pass look free.

3. Your four decisions, and where each one landed. You were asked four things before any of this was built, because a single trip cannot let each customer choose their own hour, and the answers are built rather than described. **The window** - "A window I type when I book the run" - so the run screen has two boxes, opening and closing, that you fill in yourself, and nothing is promised to any customer until you do. **The fee** - "Split evenly into each order's box" - so the trip's one fee is divided across the customers on it and each order's own charge box carries its share, which is what lets the existing charge model, the books and the customer's own total keep running exactly as they run. **What fits** - "Show me the load and let me judge" - so the screen counts the doorsteps, counts the items and names them, and leaves the judgement to you rather than refusing the run on your behalf. **The trip link** - "The same trip link for everyone" - because there is one vehicle on one journey, so every customer on it gets the one share link and there is nothing to keep apart.

4. A customer's own doorstep, and the rule that makes consolidation honest. One customer who ordered three things is ONE doorstep, not three. The screen counts stops per customer, and so does the price it asks for: measured live, a two-line order for one customer was priced as one drop, and a two-customer run was priced as two. Without that rule a customer ordering twice would have been charged as two deliveries to the same house, which is exactly the overcharge consolidation exists to remove. The row for each customer says what is being sent and where it is going, and a customer whose doorstep has never been pinned carries a press to put it on the map instead of a price it could not honestly be given.

5. What your shop promises now, and only when you have said it. A delivery window you type reaches the customer through the slots that were already there rather than through anything new: it is folded into the delivery line on their track card, and into the shipped WhatsApp message and the reminder. Measured live: the card read "Sat, 26 Sep - Courier - 1 Jalan A, 2-5 pm", and both messages read "Delivery: Sat, 26 Sep, 2-5 pm - Courier delivery" with "Track your delivery:" above the share link. Leave the two boxes empty and every customer keeps the promise they already had - the day, and no hour - so a run booked in a hurry is not a run that promises something nobody chose. And the window is checked before it is priced: a window whose end is before its start is refused in words on the spot rather than being priced and then written onto customers' cards, which it would otherwise have been, because the formatter that draws a window renders an end-before-start one perfectly happily. That is a fault this release's own test sweep found and the test now holds.

6. Nothing is booked until you press the press that books it. The run screen prices freely and books once, and the two are separate presses with a question between them. Measured live: the question read "Book the Car with Lalamove for RM 22.00? It carries ONE trip with 2 doorsteps and 6 items. Customers will be told 2-5 pm...

(RM 11.00 - RM 11.00), marked paid by you." - so the number being spent, the vehicle being spent on, the load it has to carry, the promise about to be made and what each customer is about to be charged are all on the screen before anything is sent. Behind it there is exactly ONE booking call for the whole run, and the toast afterwards read "Run booked with Lalamove - 2 customers now share one trip and one link."

7. The money, and why a share and not a fee. One trip is one fee, so charging each customer the whole fee would take the cost of one van several times over. The trip's fee is divided across the customers on the run and the remainder lands on the last order rather than being lost, so the shares add up to the fee exactly. Measured live on a RM 13.50 motorcycle price for two customers, the split line read "Split evenly over the 2 orders ticked above, in that order: RM 6.75 - RM 6.75" - and on the booked car run the two orders were stamped RM 11.00 each against the RM 22.00 fee, with two RM 11.00 "Delivery & fuel" expenses on your books, one per order, paid by you. An odd cent is not invented and not dropped: on three customers the same RM 22.00 fee lands RM 7.33, RM 7.33 and RM 7.34, and there is a test that fails if the parts ever stop adding up to the whole. Answering the charge questions is optional and stays optional: a run booked with nothing said about the money records no charge at all rather than a charge of nothing, which is the difference between a customer who owes you nothing and a customer whose box says they owe you zero.

8. Three places where the honest answer is not the convenient one. A price describes one list of customers on one journey, so anything that moves either throws it away rather than leaving a stale number standing beside a new question: untick a customer after pricing and every Book press on the screen goes inert with the sentence "The list of customers has changed since this price was asked, so it describes a run you are no longer taking. Price the run again for the list above." Change the collection day and the prices go entirely. And a live trip on any order in the run stops the whole run being booked, with one sentence saying which order is already out - the same one-trip-at-a-time guard the single order card carries, at the level of the run.

9. Two faults found by reading the drawn screen rather than by reasoning. The first: a customer's tick row was drawn as a whole pressable line, but the live area was only the words inside it - the top and bottom nine pixels of a 49-pixel line landed on the row rather than the tick and ticked nobody, which is the exact opposite of what that row's own comment promises. Measured with probes before and after, and fixed by moving the padding from the row onto the label so the whole line is the tick's. The second: ticking ONE customer left the head line reading "Who is on the run - 2 of 2" over a list with one customer on it, while the tick-everyone press corrected it. Two controls that look alike behaving differently, and the count she reads is the one that went stale. The cause was that the single row's handler repainted the load, the prices and the charge boxes but not the list the head lives in, and only the bulk press rebuilt it. Fixed so a single tick corrects the count and the bulk press WITHOUT rebuilding the rows - measured, the row she pressed is still the same node afterwards and the list does not move under her finger. Both are named rather than glossed because no assertion had ever watched either one.

10. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Your stored data was copied before the presses that write, put back afterwards and compared byte-identically, with no backup key left behind. The one thing this release adds to an order is the window you typed and the trip and share of the fee you booked - the same fields the single-order booking already writes - and an order with no run on it is drawn exactly as it was.

11. Every rule that stands on it was proved load-bearing. Thirteen faults were put back in the real source, each watched failing a test that NAMES it, and each restored byte-identically. Among them: pricing one drop per order LINE rather than per customer, charging each customer the whole fee instead of a share, writing the window onto only the first line of a customer's group, stamping the trip on only the first line, a stale price still booking, an end-before-start window slipping past its refusal, the odd cent dropped so the shares no longer add up, a zero fee recorded as a charge, and the stale head from section 9. Two of the thirteen are worth singling out because the sweep is what found the holes rather than what confirmed them: the publishing gate that decides whether a window reaches a customer was asking "can this be read" rather than "is this a window at all", so an end-before-start window published happily - the test now compares the customer's card against the card of an order with no window, which is the only comparison that catches it. And the zero fee was initially written as a

literal zero rather than left absent, and nothing failed, because no test had ever booked a run with the money questions left unanswered. A test was written for it and the fault fails by name. The suite is 1587 passing with none failing.

12. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. A run's window is folded into the delivery line the customer's card already carried, the trip and the share of the fee land in fields the single-order booking already writes, and an order record syncs whole - so all of it reaches your other phone with no sync change and no migration.

13. Still to come, and only if you want it. The public callback is the one piece left unbuilt from the courier work. It would let a customer's page update by itself within seconds of a parcel being picked up, without anybody pressing anything. It is also the one part of this work that would be reachable by anybody on the internet rather than only by your own signed-in app, so it stays your decision and nothing about it is promised here.

CHANGELOG and the guide both carry all of this, both PDFs rebuilt and read back with PyMuPDF, and the app's own More screen reads Engine v191.

25 Sep 2026 - engine v190, THE DELIVERY'S PROGRESS AND ITS DRIVER REACH THE CUSTOMER, AND AN ORDER MOVES ITSELF (one database step). The third step of the courier work you asked for. A booked trip's own progress now shows on the customer's track card in the customer's own language, with the driver's name, the number plate and a press to ring them. And when the courier says the parcel is on the vehicle, the order moves itself to Collected / Shipped - a note appears on the row naming who moved it and when, and one press puts it back. Your order row is the first thing in this app that has ever changed itself, so it is built to be SEEN and UNDONE rather than to be trusted. There is one database step, and it must be run BEFORE this build is deployed or every customer's tracking page stops updating silently. Not one module, batch, start time, cycle or saved day of yours was rewritten.

1. What you asked, in your own words. "i want to do Lalamove, API to manage courier, make it ready for other courier as well. Focus now for laalamove. And check our backoffice readiness. From my understanding lalamove support last day deliver consolidation to save on cost." v188 was the price, v189 was the booking, and this is the third of the four stages you chose. The seam and your consolidation belief are unchanged, and the fourth stage - the delivery run that saves you money - is still to come.

2. What the customer sees, and it is three lines rather than a tracking number. A customer who opens their track page on a courier order now reads where the parcel has got to ("Delivery: The driver is on the way"), who is bringing it ("Driver: Ah Meng · PMM 1234"), and a press that rings them ("Call the driver"). Each line draws only when there is something to put on it: an order with no trip on it is drawn exactly as it was before any of this existed, and a trip booked a moment ago carries no driver yet because the courier has not matched one - so the card has two lines and not a line reading "Driver: -". That is deliberate. A card that named an empty driver would be telling your customer to wait for somebody who does not exist yet.

3. The customer's page is never taught a courier's vocabulary. This is the seam, said as a rule a customer can be caught by. What the backoffice publishes about a trip is one of a handful of NEUTRAL words - finding, on_the_way, collected, delivered, stopped, nodriver - and the customer's own words for those live in the shop's own language file, in all three languages. The courier's status string (ASSIGNING_DRIVER, PICKED_UP) never leaves the backoffice. A phase the shop's page has not been taught draws NOTHING rather than the raw word, so a second courier arriving with a status list nobody has heard of cannot make your customers' page read wrongly. Six new words were added to all three languages, and a test holds every one of them and refuses a label that has lost its placeholder - the broken-label class this shop has been caught by before.

4. The driver is a SECOND call, and for most of the day it refuses - which is not an error. The trip record itself carries only a driver id, and the driver's own record - the name, the plate, the number - sits behind its own endpoint. That endpoint answers nothing until about an hour before the pickup, so a check you make in the morning comes back with the trip and no driver at all. That refusal says nothing about the trip, which is perfectly healthy, so it is folded in as an absence and never reaches your screen as a fault. Failing a working check over a name that is not available yet would be breaking a feature over a detail you had not asked for. The same rule

runs in the other direction: a check that learns nothing this time keeps the driver it already had, because treating an empty answer as an erasure would take the driver's name and number off the customer's card at the exact moment they became useful. The card in the backoffice says all of this under the booking button, in words.

5. The order moves itself, and the row says so where you will actually see it. When a check reads the trip as collected, the order moves to Collected / Shipped and a note appears on the row in your own list: "Lalamove says it collected at 10:35 am - this row moved itself. Put it back if that is not right." Written on the row rather than said in a message, because a message is gone in seconds and this is a change to an order's own money that you may not look at until the evening. The note names WHO moved it and WHEN, so the row can never read as something you did and cannot remember doing. The message at the moment it happens carries both facts in one line - what the courier said and what this app did about it - because two messages would let you read the first and miss the second, and the second is the one that changed an order.

6. A check made at the end of the day still moves it. A parcel cannot have arrived without being collected first, so a trip that already reads delivered moves the order exactly as a trip reading collected does. Without that, an order whose collection you were never told about would sit on Baked for ever and the app would look like it had missed the trip entirely - and the check you actually make is often the evening one, after the run.

7. The Undo, and the two things it deliberately does not do. The row gains an Undo beside its other presses. It puts the stage back and it puts the money flag back, and it is careful about both. An order that had no payment flag at all goes back to having none, rather than to a "false" this Undo invented - because a flag that was never there and a flag saying she has not been paid are two different things on this app's own rules. And the stock rule is deliberately NOT run on the way back. That is not an oversight: the stock rule reads a row's current stage, which at that moment is the last one, so it cannot tell "put this back" from "step this out of Baked" - and stepping OUT of Baked puts ingredients back on your shelf. Charging one order's ingredients twice because an API spoke is exactly the kind of thing this release must not do. The move itself changed no stock (only stepping into or out of Baked does), so there is nothing to undo. This was measured, not reasoned: the whole shelf count was read before and after.

8. An Undo has to survive you looking again. The guard that stops the rule moving a row twice is the stamp it writes when it moves it - NOT the trip's phase. That difference is the whole thing. If the guard were the phase, then looking at the trip again would move the row straight back, and an Undo that a look undoes has not undone anything. So after an Undo the row is once again a row the rule has not moved, and the trip being collected does not move it a second time. Measured live: after an Undo, a fresh check left the row exactly where she had put it back.

9. A row you have already marked Collected / Shipped is left alone. There is nothing to move, so nothing moves. A trip booked on an order you had already marked as gone needs no help from anybody, and the row stays yours.

10. The callback you chose to leave until after this release, and what that costs. The plan for this stage included a public callback - a small function the courier could call the moment a parcel is picked up, so the customer's card would update by itself within seconds. You chose to leave that until after this release, and the reason is on the record: Lalamove's webhook specification could not be read from here, and a public endpoint built against a guessed shape is exactly how a security hole gets shipped - it is the one part of this work that is reachable by anybody on the internet rather than only by your own signed-in app. So the honest description of where things stand is this: the customer's page is as fresh as the last time a phone asked the courier, and pressing Check is what asks. Her own order row catches up the same way. Nothing is lost, and nothing is claimed that is not true - the callback needs no database change and can be added whenever you want it.

11. One database step, and the order it must be run in. supabase/courier_job.sql adds five columns to the order_tracking table: the courier's name, the neutral phase, the driver's name, the plate and the number. It is safe to re-run and each column is added only if missing. RUN IT BEFORE THIS BUILD IS DEPLOYED. The backoffice publishes the whole tracking row in one call and PostgREST rejects the entire call if it names a column that is not there - and the publish is deliberately silent about its own failures, so it would not look like a failure: every customer's tracking page would simply stop updating, for every order, with nothing on any screen

saying why. This is the same trap the courier charge migration documented, and the SQL file carries the warning in its first paragraph.

12. The shop's own page had to learn five new columns, and forgetting one fails silently. The customer's card asks PostgREST for a named list of columns, and it returns ONLY what is named there. A column the backoffice publishes and this list forgets is a line the card can never draw - and it fails with no error anywhere: the row arrives complete, the card is simply missing a section, and nothing says why. So the five names were added to the request, and a test now reads the request the page actually sends rather than trusting that it was edited.

13. A fault found by measuring the drawn card, which no test in this repository could have seen. The two things on the customer's card that matter most at the door - ringing the driver and opening the courier's own tracking page - were built from padding alone and came out 30 pixels tall, while every other thing a customer taps on that page (the delivery day, the collect / courier pair, Track, Place order) is 37 to 48. The two smallest targets on the page were the two the customer most needs at the door. Found by measuring the drawn card at 375 pixels, not by reading the rule. They now take the page's own floor and centre their word in it, re-measured at 38 pixels. And the reason no assertion caught it is worth the sentence: every stand-in screen in this repository answers a box with the number it was handed, so a link 30 pixels tall and a link 48 pixels tall measure exactly the same. The guard therefore reads the stylesheet's own declared numbers instead, and it holds both halves at once - the floor and the box that makes the floor mean anything.

14. Every rule that stands on this was proved load-bearing. Sixteen faults were put back in all, each watched failing a test that NAMES it and then restored byte-identically. Among them: a check that moves an order every time it looks; an Undo that restores an invented "false" where there was no flag at all; a stage change that forgets to tell the customer's own card; an order already marked paid being un-paid by a move past it; a late check that skips the collection it was never told about; a courier's raw status word published to the customer; a plate labelled as a driver; a number that cannot be dialled turned into a button anyway; a driver read as a bare string; and the customer's card dropping one of the five columns from its request. Two of the sixteen are worth naming because a test could not have caught them before a person read the screen: the dispatcher's own import list, where a name called but not imported is a 500 that only appears when that action is used, and the two links on the customer's card being 30 pixels tall, which is section 13. One guard turned out to be unreachable by pressing anything at all - the transition rule in the courier panel and a second guard inside the order rule both answer a repeated check, so the inner one can never be reached from the screen. Rather than leave a guard nothing can prove, the rule was exported and a test written for the one state that isolates it: the state where YOU have moved the stage back by hand after the rule moved it, which is a state you really can reach. The suite is 1545 passing with none failing.

15. Where the secret lives, and where it does not. Unchanged and worth restating, because this release adds a second call that spends nothing but reaches the courier's own records: the signing of every request happens on Supabase, in a function only your own signed-in app can call. The key and the secret are Supabase secrets you set yourself; they never enter this repository, this app's screen, or any message. The driver's number is the one thing this release publishes on purpose, and it does so because you asked for it and because the person at the door is a stranger your customer has to meet - but it is published only for a trip that has one, and only as a number that can actually be dialled.

16. Nothing of yours is rewritten. Measured on the running app, on a phone-sized screen, with the whole state read out of storage before the pass and read out again afterwards and compared field by field. The pass was the full one: a trip checked, the order moving itself, the Undo pressed, and then a SECOND check made afterwards. The result, in full. After the Undo, the order's stage and its payment flag were byte-for-byte what they had been - the Undo really does put the row back and not merely draw it back. The only fields that changed at all, across the entire pass, were the trip's own: what the courier last said, when it said it, and the driver it named. Nothing else in the whole stored state moved - not one ingredient, not one order's items, not one setting. The three saved days still read 24, 4 and 24 pans, exactly as they did before. And the phone's storage holds exactly the same three keys it held before the pass, with no backup key and no scratch key left behind.

17. What is still to come. The delivery run - one trip carrying several drop-offs at once, which is the stage that saves you the money your own hunch was about. It is v191, and it opens with a question to you rather than with code.

25 Sep 2026 - engine v189, BOOKING THE TRIP (no database step). The second step of the courier work you asked for: the price you chose becomes a real van. One press books the trip with Lalamove, the share link it sends back goes into the tracking box your orders already have, and one press checks where the driver has got to or calls the trip off. Two real defects were found by reading the drawn screen rather than the code - the job card had been printing the courier's own word ASSIGNING_DRIVER instead of "Finding a driver", and the price panel was printing the word "null" under its last row - and both are fixed, guarded and proved. The customer-facing half had no test at all and has three now. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "i want to do Lalamove, API to manage courier, make it ready for other courier as well. Focus now for laalamove. And check our backoffice readiness. From my understanding lalamove support last day deliver consolidation to save on cost." Three separate things. v188 built the quote, which was the first. This is the second: booking. The seam and your consolidation belief were both answered in v188 and nothing about them changed here.

2. What a booking does, and what it deliberately does not. The trip comes back with its own share link, and that link is written into the tracking box this app has carried since v97 - the box that is already printed on the customer's card and in the shipped WhatsApp message. That reuse is why a booking needs no database step at all: an order row syncs whole, so a trip booked on one phone reaches your other phone with no migration and no new column. A share link IS the courier's reference for this delivery, so putting it where a tracking number used to go is honest rather than a shortcut. One copy touch goes with it: when that box holds a web address the message says "Track your delivery" instead of "Tracking number", because it is a page the customer opens rather than digits they read out.

3. The driver's name and number plate cannot be shown at booking - and the plan promised one. This is the one place this release does not deliver what was planned, and it is said here rather than glided over. Lalamove has no tracking number and no driver at the moment a trip is booked; the driver is assigned afterwards and Lalamove does not release the name and plate until about an hour before pickup. So booking shows the price, the vehicle and the share link, and the driver arrives later, if you want that screen. There is nothing to fix; the promise was simply wrong when it was written.

4. The two defects, both found by reading the drawn screen rather than the code. The first: the job card printed "ASSIGNING_DRIVER" where the app has a translation table with a test on it. The table was correct and tested; what was missing was the wire from the card to it, because the card asks the courier "what is this status in words" and the courier object did not answer that question. The fallback printed the courier's own enum instead of complaining, which is what a quiet fallback does. It reads "Finding a driver" now, and the whole interface has a guard so a second courier cannot ship without the answer. The second: the price panel printed the word "null" under its last price row, because two optional lines were handed to a browser function that turns a null into the word. It was invisible to every test the app had, because the ordinary stand-in screen quietly drops nulls - only the strict one keeps them, and it now draws this panel too.

5. Your three decisions, carried forward and unchanged. "Locate it, let me fix it" for addresses, so a courier is handed a point and never words on a card. "All four stages", so this is the second of four and the remaining two follow as v190 and v191. "Build now, keys later" - the app still carries no key and no secret, and says so in plain words rather than failing in a way that looks like a bug.

6. Measured live on the drawn screen, at 375 by 812 and at 1280 by 900. The job card reads "Lalamove is on this order", then "Car · RM 14.00", then "Booked 1:51 am", then the two presses, then "Status: Finding a driver - read just now", then the customer's own link. Pressing Check the trip against a reply of PICKED_UP put up "Trip status: Collected" and the card then read "Status: Collected - read just now". Pressing Cancel trip opened the app's own confirmation, which names the consequence and the money: "Call off this Lalamove trip? The driver stops being sent, and the customer's tracking box keeps the link but nothing will update it. This cannot be

undone from here - you would have to book again, at a fresh price." Confirming it put up "Trip called off" and the card then read "Status: Called off from here 9:28 am". The presses measure 115 by 36, 99 by 36 and 110 by 36, this app's own tap-target floor, and the confirmation's two presses 87 by 46 and 100 by 46.

7. Where the honest answer is not the convenient one. Calling a trip off sends the courier a request that answers with nothing at all - so the card does NOT claim the courier said "Cancelled". It records that YOU called it off, with the moment, in the app's own words: "Called off from here". The same reason runs the other way on a finished trip: the heading changes from "Lalamove is on this order" to "A Lalamove trip on this order", the Cancel press is withdrawn rather than offered and then refused, and the Book presses come back to life so you can rebook a trip you cancelled. And a checked status is merged into the trip's record rather than replacing it, so looking at a trip never forgets the price it was booked at.

8. The money guard, measured. A trip that is still running blocks a second one. Measured live with a live trip on the order: all three Book presses drawn disabled, and the panel's own sentence under them - "This order is already on a trip." That is the guard that stops you paying for a second van to the same door.

9. The customer-facing half had no test, and that gap was found by writing the faults rather than by reading the code. v189 puts a trip's share link into the tracking box, which makes the "is this a link or a number" branch customer-facing on every courier order. The number half was covered; the link half, which is exactly what a booking newly produces, was not - on either side. Three tests were written: the storefront now proves a link is drawn tappable and a javascript: value is not; and the shipped message now proves it routes the value through the one label-maker rather than labelling a URL a tracking number.

10. Nothing of yours is rewritten. No module, batch, start time, cycle or saved day. Nothing was written to Supabase, and no real network call left the browser for the whole measurement: the courier was answered by a stand-in. The pane's own stored data was compared field by field afterwards and put back - the only fields the measurement had changed were the trip's own state and the app's two customer edit-stamps, and the fake trip, the fake share link and the fake session token are all removed. The pane holds exactly the three storage keys it held before, with no backup key left behind.

11. Every rule that stands on it was proved load-bearing. Twenty-four faults were put back in all, each watched failing a test that NAMES it and then restored byte-identically. Among them: the request envelope the API needs; a reply not being unwrapped; the stop ceiling checked AFTER the doors; a doorstep matched by position instead of by where it is; a finished trip still counting as live; a phone number losing its plus; a trip number dropped into a URL unescaped; the vehicle and time being re-sent on a booking, which would be a second and quieter way to choose a vehicle; a GET carrying an empty body; and a cancel sending one. Four of them are named rather than counted, because a test could not have caught them and one was found only by a fault: the printed "null" and the printed enum, both read off the drawn panel; the storefront's link branch; and the missing "done" flag on a job written before that flag existed, which every test passed until a fault exposed it. The suite is 1519 passing with none failing.

12. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. A booked trip is kept on the order row, which syncs whole.

25 Sep 2026 - engine v188, A DELIVERY PRICE BEFORE YOU PROMISE ONE (no database step). The first step of the courier work you asked for: on an order you are delivering, one press asks Lalamove what the trip would cost on every vehicle it runs, shows the prices side by side with the distance and the quotation's own five-minute life counting down, and the one you choose fills the charge box that was already there. Nothing is booked, nothing customer-facing changes, and not one module, batch, start time, cycle or saved day of yours was rewritten. There is no database step.

1. What you asked, in your own words. "i want to do Lalamove, API to manage courier, make it ready for other courier as well. Focus now for laalamove. And check our backoffice readiness. From my understanding lalamove support last day deliver consolidation to save on cost." Those are three separate things, and this release is the first of them. The other two are answered below and built on top of this.

2. Your consolidation belief is correct, and here is the mechanism. Lalamove does not consolidate by

holding parcels back for a later trip; it consolidates by letting ONE trip carry several drop-offs, two to sixteen of them, priced as one base fare plus a fee for each extra stop. Outside the Klang Valley, which is Penang, that fee is RM 1 on a motorcycle, RM 2 on a car, RM 5 on a four-by-four or van and RM 10 on a lorry, so eight cakes to eight houses on one motorcycle is one base fare plus RM 7 against eight base fares. Three things go with it and all three are said rather than hidden. The vehicle has to physically hold the load, so one motorcycle is not eight focaccia and the run screen will say so. There is no bulk pricing call, so a run of eight drops is still eight price requests, which costs you waiting time and not money. And route optimisation is described two different ways in Lalamove's own documents, so the stop order is yours to arrange and the screen will never claim it found you an optimum. Building that screen is the last stage of this work and the only one that changes what your shop promises, so you will be asked before it is built.

3. The seam, which is what "ready for other courier as well" had to mean - and a hole in it that the audit found. A claim in a comment is not readiness, so readiness here is a file that holds the one interface every courier must answer, a file that is the only place any courier's own name and rules live, and one channel that both use and that never throws at you. Behind it all is a **fake courier written for the tests, which quotes without the name Lalamove anywhere near it** - and that fake is the proof that a second courier is a new file and a line in the registry rather than a rewrite.

And the audit you asked for, walked against the built thing rather than against the code meant to satisfy it, found the claim was not yet true. It found the same defect twice, once on each side of the seam. On the app's side, two screens and the one shared file named the courier in their own words: the Settings card printed "Courier (Lalamove)" and sent the literal word to the server itself, and the shared channel defaulted the courier to Lalamove when nobody said otherwise. A default is that file quietly choosing a courier, which is the one thing the seam exists to prevent - the day a second courier arrives, a default left behind would send its calls to the first one and nothing on screen would say so. On the server's side, the dispatcher wrote the courier's name into a refusal you would read: "Lalamove did not name any vehicles for this market". That is the identical fault, and it is the one that matters more, because a refusal is exactly the sentence a second courier would have to have found and edited by hand. Both are fixed with no visible change: the card asks the registry for the name, the channel refuses to name one, and the refusal now comes from the provider's own label. The words on your screen are character for character what they were. And the rule has a test it never had, which is the real lesson here and the same one v187 taught. It now covers both halves: inside the app the word is allowed in two files only, the provider and the registry that imports it, and the walk refuses any third file that names it; on the server the walk does the same, and the dispatcher gets a tighter rule of its own because its environment variables ARE named after the courier and it therefore had to be on the allowed list - it may name the courier only in those variables and in the provider's own file path, and in nothing that could be shown to you.

4. Your three decisions, and where each one landed. You chose "Locate it, let me fix it" for addresses, so a courier is never handed words on a card: it is handed a point. The address already on the order is looked up once, the answer is kept against that customer so a second order from the same number costs no lookup at all, and where the lookup misses a map opens and you tap the exact door. The bakery's own door is pinned once, from the same place you asked for it. You chose "All four stages", so this is the first of four and the other three follow as v189, v190 and v191. You chose "Build now, keys later", and that is exactly where this stands - see paragraph 9.

5. What the screen does, on your own order, measured live at two widths. The price panel opens inside the card that already holds the charge box, folded away until you ask for it, and that is deliberate: this app has one pop-up layer, so a second card would destroy the charge box, the note you were part-way through and the Save button, and the fee would land in a box nothing could save. Measured at 375 by 812 with the test courier answering, the panel drew three price rows - Motorcycle RM 8.50 counting its own life down, Car RM 14.00, Van RM 0.00 - then the line that says which vehicle could not be priced and why, then "3 prices from Lalamove, just now, for collection 2026-09-30 at 10:00. Each one dies on its own clock." The four presses at the foot of the rows measure 99 by 36 pixels, which is this app's own tap-target floor, and the line naming the two ends of the trip reads "From Jienluv2bake, George Town to 12 Jalan Bunga, George Town". At 1280 by 900 the same card drew

all three rows, the failure line and the live status line with nothing clipped.

6. Taking a price is one press, and it goes through the money path that already worked. Pressing Use this fee on the RM 14.00 car put 14 in the charge box and said so, and the panel folded itself away so the number it had just written is what you are looking at with the payer question under it. Setting the payer to "The customer paid it" then read back "The customer owes RM 58.00 - items total RM 44.00 + courier charge RM 14.00". Nothing new was invented for the money: the fee lands through the same model that has carried the courier charge since it was built, so the amount, who paid it and COD all keep running the way they already ran.

7. Three places where the honest answer is not the convenient one. First, a quotation DIES - Lalamove's last five minutes - so every price carries its own clock, a dead price greys itself out and says "ask again" rather than sitting there looking live. Second, a courier can reply with a total of zero, and a zero is not a charge: measured live on the RM 0.00 van, the press left the charge box exactly as it was and said "Lalamove priced this trip at RM 0.00 - that is not a charge, so nothing was put in the box. Ask again, or type the amount," and the panel stayed open with all three rows still there. A toast claiming a number had landed in a box the box had refused is the screen lying to you, and this is the one place in the app where a lie about money is easiest to ship by accident. Third, a price that arrives after you have closed the card is not written into it: measured live with a deliberately slow reply, closing the card the moment after asking left the pop-up closed, no charge box and your stored data unchanged, character for character the same length before and after.

8. Where the secret lives, and where it does not. The signing of every request happens on Supabase, in a function only your own signed-in app can call, and the key and secret are Supabase secrets that you set yourself and that never enter this repository, this app's screen or this conversation. The browser half is built to the same shape as the mail function that already works: it never throws, and every refusal arrives as a sentence you can act on rather than a code. Two you will meet: a missing pin reads as a pin to be placed rather than a failure, and an address Lalamove cannot find reads as an address to put on the map. Until you open the partner account and set those secrets the app carries the courier and no keys, and it says so in plain words instead of failing in a way that looks like a bug.

9. What is measured and what is not, said plainly. The panel, the two ends of the trip, the five minute clock, the money, the refusal and the closed-card case were all measured live, in a real browser, on the drawn screen, at 375 by 812 and at 1280 by 900 - against a stand-in courier that returns Lalamove's own documented replies. What is NOT measured is Lalamove itself, because there is no account yet: the signing and the reading of its real replies are tested against its published sample answers in the suite, and three details are flagged at the head of the provider file where the first real call will show them - whether a stop's point is sent under one field name or another, the exact language tag this market expects, and whether the distance comes back with a unit or as a bare number. Each is one line in one file, and each will be settled in the first minute of the sandbox, not guessed at now.

10. Nothing of yours is rewritten, and every rule that stands on this was proved load-bearing. Measured live: your stored data was backed up before the single press that writes anything, put back afterwards and compared, and your phone's own storage holds exactly the two keys it held before, with no backup key left behind. Eleven faults were put back in all, each watched failing a test that names it and then restored byte-identically - among them the zero being accepted as a charge, the closed card being written into anyway, the panel's clock failing to clear itself when the card it belonged to was gone, and all four seam leaks named in paragraph 3. Two of the eleven are worth naming because a test could not have caught them, and both are said rather than glossed. The clock one: the panel stops its own clock when a fee is accepted and folds away, so only the case where the fee is REFUSED and the panel stays open exercises the rule at all, and the test that catches it is the one written for the refused zero. And the server's own hand-written refusal: when that fault was put back, the first version of the new rule did NOT catch it, because the rule read only ordinary quoted text and the sentence was written as a template literal - so a real leak sat in a blind spot of the very guard written to catch it. The rule was widened to read those too, and the fault then fails by name. The suite is 1460 passing with none failing. **No database step** - a price is transient and an accepted one lands in fields that already exist, so there is nothing to run in Supabase.

24 Sep 2026 - engine v187, EVERY PERSON'S WINDOW IS THE HEIGHT OF ITS OWN PEOPLE (no database step). One fault, found by measuring the drawn board against your own eleven points and then fixed: the workers' windows on the Production page carried a ceiling, so on a phone only the first person stood on screen and the other three sat inside a box that scrolled. Not one module, batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked. You told me the Production page deferred from your instruction, and then: "fix it". So this release is your eleven points read back against the board that was actually drawn, not against the board that was intended, and it carries the one place they did not agree.

2. The fault, in one sentence. Your point 8 says "The person's window be taller" - the *window*, not the row. v186 made the row taller and left the window's own ceiling exactly where it was.

3. Why a ceiling was there at all, and why it was wrong here. The window a person's rows sit in had a ceiling of min(20vh, 190px) - 190 pixels on your phone. That number was written for the Scenario planner's own people window, where the rows are 35 pixels tall and five of them is 190. The Production page's rows are 72 pixels, because a taller row is what you asked for. Four of them need 406. The ceiling did not know the rows had grown.

4. What that looked like on your phone, measured. At 375 by 812 the window was 162 pixels tall while holding 406 pixels of people: one person on screen, and persons 2, 3 and 4 behind a scrollbar that was not visible and that you had no reason to look for. Read off the drawn board: 0 of 4 rows fully visible. Your point 2, "Create windows for each person", was answered with one window and three hidden behind it.

5. The fix. One line of the stylesheet. The board's own person window is now as tall as its people and carries no ceiling of its own; the ceiling stays exactly where it was written, on the planner's window, which still has the five 35-pixel rows it was written for. Nothing else about either window changed, and nothing inside either of them was restyled.

6. Measured after, at the same two widths. At 375 by 812: the window 407 pixels tall, its inside 406 and its content 406 - equal, so nothing scrolls - all 4 of 4 rows fully visible, the coaches still 37 pixels wide with four filling the line, four red rulers still standing at 107 and the one clock face still at 215 reading 11:27 pm. At 1280 by 900: the window 401 pixels, no scroll, 4 of 4 visible, the coaches 155 and the line 738. The modules' window above and the planner's borrowed window below are the heights they were.

7. Nothing of yours is rewritten. The whole of this release is one stylesheet line and one test, and neither of them writes anything: a ceiling is a drawing rule, and no module, batch, start time, cycle or saved day is read by it, let alone changed. Measured live on a four-person day built for the purpose, the app's own stored data was put back from a copy taken beforehand and compared byte for byte the same, with no backup key left behind.

8. And the rule has a test now, which it never did. This is the part worth reading, because the reason the fault shipped is not the fault itself. No test could ever have caught a window that was too short. The stand-in screen the tests draw on answers every box with the same measurement - nothing high and nothing of content - so a window with three persons hidden inside it looks exactly like one that fits. A stand-in that says every box is zero high cannot see a box that is too short, and a guard that measured that screen would have passed just as happily with the bug in place. So the guard reads the stylesheet's own declared numbers instead, and it holds both halves of the rule at once: the board's window declares no ceiling, and the planner's still declares its own. Two faults were put back and each was watched failing that test by name - the ceiling returned to the board's window, and the ceiling taken away from the planner's - so the test cannot be passed by deleting the cap everywhere and leaving the screen these windows were borrowed from unbounded. Each was restored byte-identically. The suite is 1343 passing with none failing.

9. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. A ceiling lives in the stylesheet and in the test, and your plan, your days and your own settings are exactly as they were.

24 Sep 2026 - engine v186, THE MORNING CAN BE WALKED FROM THE WORKERS' OWN WINDOW (no database step). Eleven things you asked for on the Production page, built as one: the lines drag left and right and everything on them moves together, Start the day now / Stop / the bell / Back to now sit where the workers read them, and the walk's own pace is a drop-down of Slow, Mid and Fast. Not one module,

batch, start time, cycle or saved day of yours was rewritten, and there is no database step.

1. What you asked, in your own words. "let me tell your what i want for production page. 1. Keep the Next Up. 2. Create windows for each person, the detail able to be drag left or right. 3. Put the person windows parallel, close to each other, the 2nd person and subsequent person have the centered red ruler, but the 2nd and other don have to show the clock face. 4. Replicated for person 2, 3, etc. 5. I need walk the day, Button 'START THE DAY NOW', 'STOP' , 'The N Called', these button are for the person's windows. 6. The scenario dual window dont join to the person's window. 7. Name of the scenario should be on Top of page. 8. The person's window be taller, so that coach size better fit wordings. 9. Each of the person line has a clock line at centre of line, when we drag to the right, the train move to right and the clock and red ruler move relatively. 10. When i click outside the person window, the clock back to center. 11. For simulation, i want some button to do this, left buttton make the clock back to current, Right button, when clicked, dropdown a choice list, slow, Mid, fast. This button is to make the chart move faster for simulation purpose."

2. The Next Up stays. Your point 1, and nothing about it moved. The card that says what is next is still there, still first, still telling you the job and its clock.

3. The scenario's name, on top of the page. Your point 7: the heading card now carries the name of the day you are looking at, so a page of numbers cannot be mistaken for another day's.

4. A taller person's line, so the words fit. Your point 8. Each of the workers' lines is now *72 pixels* tall, against the 56 it was and the planner's own 35 - measured at both 375 pixels wide and 1280, and every coach measures 52 of those 72. Nothing on a coach face is clipped at either width: the face reads its whole derived word - at 375 pixels the words on your own day read *mixer, fold, Wash, Load, Unload, Dimple, oven, Cool* - and not one word is cut off.

5. The two windows stay your borrowed pair, and they do not join the lines. Your point 6. The modules window and the planner's people window are drawn below the lines by the very code that draws them in Scenario planning, and they are **not** tied to the train: the trains do not scroll sideways at all, so no drag on a worker's line can move the two windows above it. What v185 put back - the two borrowed windows panning together on both screens - is still true and still theirs; it is the train that stands apart, and that is your point 6 said as a rule.

6. A red ruler on every line, but only one clock face. Your points 3 and 4. Every person's line carries its own red ruler standing at the middle of that line - at 375 pixels wide the ruler sits at 107 and at 1280 the line's own centre is 694 - and the face is drawn **once**, on the first line, reading the clock at the line's own centre in the pane. The second person and every person after carry the ruler and no face, exactly as you asked. Measured at 375 pixels: the face stands at 215 in the pane and each row's ruler at 107 in its own track, and the two differ by exactly the 108-pixel name column between them - so they are one line, drawn in two coordinate systems.

7. The drag, and everything on the line moves together. Your point 9, in your own words: "when we drag to the right, the train move to right and the clock and red ruler move relatively." A sideways drag on a person's line moves the strip, that line's red ruler and the clock face by the same number of pixels - measured live at 375 pixels wide, an 80-pixel drag took the strip from -437 to -357, the ruler from 107 to 187 and the face from 215 to 295, *all three by exactly 80*. Nothing about it is saved, so it dies with the screen.

8. Letting go of the line, or pressing off it, puts the clock back at the centre. Your point 10, and your point 11's left button. A press anywhere off the line re-centres it, and the *Back to now* press above the lines does the same thing on purpose - both put the ruler, the face and the trains back where the day puts them. It is never saved, so it cannot move a module, a batch or a start time.

9. Walking the day, without touching the day. Your point 5, in your own words: "I need walk the day, Button 'START THE DAY NOW', 'STOP' , 'The N Called', these button are for the person's windows." So the presses sit above the lines, inside the workers' own window:

- **Start the day now** begins the walk at today's own clock time and takes the screen to the day's first minute. **Stop** ends the walk and puts the lines straight back on the real clock. **The bell** tells you how many jobs would be called and throws the calling on and off. **Back to now** re-centres the line whether or not a walk is running.

Measured live on a 375-pixel screen with the walk started: the line went to the day's first minute and the row read *Now: mixer*, the link's count was 30m, and *12 coaches stood red*, falling away as the clock ate into them. Nothing about the walk is written down anywhere, so a day you saved is exactly the day you saved no matter how long the walk runs.

10. The walk's own pace, and it is only the walk's. Your point 11, in your own words: "Right button, when clicked, dropdown a choice list, slow, Mid, fast. This button is to make the chart move faster for simulation purpose." The drop-down carries *Slow 1x, Mid 10x and Fast 60x*. Measured live at Fast: three seconds of walking took the clock face from 10:58 pm to 11:01 pm and brought a link's countdown down from 30m to 27m with it. Choosing a speed changes nothing in your plan - it is the walk's speed and it lives nowhere but the walk. **Stop** then put the face and the line straight back to 10:58 pm, turned the press off and brought Start back, saying "Stopped - the line is back on the real clock."

11. The countdown comes down by itself. Nothing has to be tapped for a worker to watch their next job get nearer: measured on the board, one minute of the clock brought a row's own count from *Next 20m to Next 19m* with nobody touching the screen. And that minute rebuilt nothing - the app's own rule stands, because a repaint throws away a worker's scroll, so a beat writes the numbers on the line and never rebuilds the line around them.

12. One more thing was mended underneath. A tap on a coach had to be proved to reach your other phone and the store, not just the screen. Reading the tests showed that the tick's write - and the Clear press's write - were being read off the screen's own copy of the day rather than off the saved day, so a tap that moved the picture and skipped its save would have looked exactly like a working one. Both are read off the saved day now, and both fail by name if the write ever goes missing again. This is the app's own standing rule: a stand-in must be as unforgiving as the real thing, because a forgiving one hides a bug instead of catching it.

13. Where a tick lives, and why there. A tick is written in a key of its own beside your plan, never inside it - so **clearing the board can never touch a saved day, and deleting a day can never wipe the board.** *Clear the board* asks first, in the app's own words, before it takes a single green coach back.

14. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Measured live, across a whole pass of this page - the board drawn, the day started and stopped, the speed changed three times, the line dragged and put back, a coach ticked and then the board cleared - the one and only thing that changed in the app's stored data was that new key of the tick's own, sitting beside the plan: an empty board written as a spoken empty so your other phone cannot read a clear as ignorance, and the plan itself untouched. The app's own version reads *Engine v186*.

15. Every rule that stands on it was proved load-bearing. Eleven faults were put back in all, each watched failing a test that names it, and each restored byte-identically. One of them escaped on the first pass and is named rather than glossed: with the beat no longer repainting a row's next-job line, nothing failed - because no test had ever watched a countdown come down on a beat alone, with nobody touching the screen, which is the half a worker at a bench actually reads. A test was written for it, and the fault then fails by name, reading that a minute of the clock moved a row's own countdown from 20 minutes to "Next 20m". The suite is *1343 passing* with none failing.

16. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The tick is written under a settings key the app already guards, the walk's speed is a number in the walk and nowhere else, and nothing new is written anywhere.

24 Sep 2026 - engine v185, THE COACH KEEPS ITS WORD, AND THE BOARD'S PAGE IS PUT IN YOUR ORDER (no database step). Two things: every coach on the Production line's train now carries a word at every width, and the board's page reads top to bottom the way you asked - the person's train first, then the two windows borrowed from the Scenario planner, kept original. Nothing of yours is rewritten, and there is no database step.

1. What you asked, in your own words. On the coach face: "Phone coach" reads as "Two lines: icon + name, then time", and on the broken label, "Yes, fix it." Then, on the page itself: "let me emphasise, the production page start with the person's windows the tarin inside, after the N person, we have the 2windows that we bring in

from scenario planning, we keep it original." Asked what should sit below the train rows, you chose both planner windows - the module window AND the people window exactly as Scenario planning draws them.

2. The coach keeps its word. v184's refinement - "if the coach box is too smalll to house the full words, then just show meaningful, hoover mouse on show tooltip, with more details" - had been built as *drop the name and keep the icon*, and reading the drawn board on a phone showed what that cost: at 375 pixels every coach was nameless. So a narrow coach now gives up its type size and its side padding instead - 10 pixels to 9, and 3 to 2 - and never the word. The stylesheet rule that hid the name outright is gone, and a test now refuses to let it come back.

3. And the word keeps its punctuation off. The word on a coach is derived from the module's own name, and it was keeping whatever punctuation sat against it: "Wash, oil and fill" was showing as *Wash,* with the comma welded on. It reads *Wash* now. That one was found by reading the drawn board rather than by reasoning - the same way the last one was found - and it is exactly the kind of broken label your refinement refuses.

4. The board's page, in your order. The Production page now reads top to bottom: the person's train rows first, then the modules window and the people window, drawn by the same code that draws them in Scenario planning. They are borrowed, not re-drawn - a second renderer of one day is the thing this app refuses - so the bars, the rows and the counts in both windows are the planner's own.

5. What that order costs, said plainly rather than hidden. Two things follow, and neither is dressed up. The people now appear twice on that one page: once as train rows at the top, and once inside the planner's own people window below. And because that borrowed people window sits on the planner's minute axis, it pans together with the modules window again on both screens, while only the train stays fixed - v184 had taken the pan off it because the train had replaced it, and you have put it back. The train itself still never scrolls sideways, which is what keeps its centre the centre.

6. A seam, so the train is not mistaken for the map. A hairline rule now runs above the modules window on the board, so a worker can see where the coaches end and the day's own map begins. It is a seam on the wrapper only - nothing inside either borrowed window is restyled, which is what "keep it original" means.

7. One more thing had to be mended underneath. The app remembers where you had scrolled in each window and puts it back after it redraws - your standing rule that a redraw must not move what you are looking at. It was finding the FIRST window of each kind; with one page now holding two people-shaped windows, the second one was being left behind. It keeps every one of them now, each by its own numbers.

8. Measured, on the drawn board at 375 x 812, signed out of the cloud. The three windows are in your order on the page - the train 57 pixels tall at the top, the modules window 493 under it, the planner's people window 35 at the foot - with 31 coaches 37 pixels wide, all four to the visible line, the clock standing at the line's own centre, and every coach face reading a whole word: *mixer, fold, Wash, Load, Unload, Dimple, oven, Cool*. Not one clipped word and not one nameless coach anywhere on the board.

9. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Your three saved days still read 24, 4 and 24 pans, and there is nothing to run in Supabase - not one stored field is added or changed.

10. Every rule that stands on it was proved load-bearing. The coach's word, the punctuation it keeps off, and the seam above the modules window were each put back as a fault, each watched failing a test that names it, and each restored byte-identically. Two of the three were caught by name; the seam was caught by the page-order test, which now refuses a board whose train and day's map run together. The suite is 1341 passing with none failing.

24 Sep 2026 - engine v184, THE TRAIN ON THE PRODUCTION LINE, AND TWO CLOSER STOPS ON THE DIAL (no database step). The board's workers' window is no longer a strip of time - it is a train of work. One coach is one job, joined in the order they happen, with the clock standing at the centre of the line, a countdown on every link between two coaches, and one tap on a coach turning it green. Your own words for it: "For the production, the effectiveness need to be improve, for each of person line, make his series of work join up like a train, coaches represent work, the current time at the center sharing with all person, showing next work in minutes on top, something like next station coundown for subway. Person

should click on their work to turn it green indicating acknowledgement." Nothing of yours is rewritten, and there is no database step.

1. What you asked, in your own words. Six points, and then one refinement while it was being built: (1) "i want to have 2 step more closer scale"; (2) the Production line should join each person's work up "like a train, coaches represent work, the current time at the center sharing with all person, showing next work in minutes on top, something like next station coundown for subway", and "Person should click on their work to turn it green indicating acknowledgement"; (3) "The coach can pass the current timeline, but stay red, click it turn green"; (4) "Each title of the line show their name, make each line distinctively seperated"; (5) its purpose - "to tell person, what they have done, and what is coming dynamiccally, on the work that have not done, every person can see it, and can land. helping hand if their process can tolerate an interuption"; and (6) "Ask me question if not sure" - which is what settled the six design answers below. The refinement: "if the coach box is too smalll to house the full words, then just show meaningful, hoover mouse on show tooltip, with more details."

2. Two closer stops on the dial. The Scenario planner's Scale control went from four stops to six. The two that were there stay where they were and only one name changed: the old Closest is now called Closer, and the new top stop is Detail. Three lists have to be the same length or the ruler loses its step and the chart breaks, so a test now says so out loud rather than trusting it.

3. The answers you gave, which the build follows. Asked five rounds of questions, you chose: the clock is "Pinned to the centre"; one coach is "One job they do"; "Anyone, on any line" may tick; clearing is "A Clear the board button only"; green means "I'm on it (seen)"; a coach goes red "When the job is due"; coaches are "All the same width" with numbers "Only on the links"; ticks "Travels to your other phone"; the two windows are kept with the person's line taller; and the architecture is "The train is its own strip". And on the clock itself: "lock the clock relative to the train, when drag, clock line move, when click outside, the clock centred on the line, make the coach not too big, in a full shown line should be able to visualize 4 coach, we can also show the coming time countdown on the coach link."

4. What the person's line is now. A taller 56-pixel row per person, separated from the next by a rule and a rail in that person's own colour - your point 4, "make each line distinctively seperated". The name cell carries who it is and what is next ("Next 26m", "Now: rests"). Alongside it, one coach for every job that person has, all the same width, in the order the jobs happen. Each coach carries its icon, a word that means it, and the clock time. That word is DERIVED and never cut: "Into the proofer" shows as "proofer", "Dimple and top" as "Dimple", "The oven swap and the bake" as "oven" - because a face showing the word "The" would be showing a word and saying nothing. The full module name, the two clocks, the length, the batch and the state in words are all on the coach's tooltip, which is your refinement read at the only moment it applies.

5. The clock, the countdown, and the colours. The clock is one line at the centre of the line, shared by every person, and its label is the clock time - never the word "now", so it can never be mistaken for the planner's own now-line. Every wait between two jobs carries its countdown in minutes on the link, "like next station countdown for subway", and it counts down live. A coach turns red the moment its job is due and STAYS red after the clock has gone past it - your point 3 exactly. One tap turns it green and names the job back to you; the same tap takes it off again. A green coach travels to your other phone, and a Clear the board press appears beside the calling switch, asks first, and clears every coach at once.

6. Dragging reads the day without moving it. A sideways drag on the workers' window moves the clock line only and leaves every train exactly where it was, because a drag is somebody asking "how long have I got" and not a change to the day. The label then reads the minute the line is standing on and the line is drawn in the reading colour. Letting go, or pressing anywhere off a row, puts it back at the centre. It is never saved, so it dies with the screen.

7. One wording fault found by measuring, and fixed. Live at 4:30 am on your own day, Person 3's one job had run 4:00 to 4:12, so their row read "All done" - beside a coach that was standing RED and untaken. That is the board telling a worker they have finished something nobody has touched, and talking them out of the very tap your point 3 asks for. A row with work still untaken now counts it instead ("1 not taken"), and gives the count up the moment the last coach goes green.

8. One real bug, found only by measuring in the browser. The widths of a train's coaches were being assigned to the strip in a way that looks right and writes nothing at all, so every coach fell through to the stylesheet's own fallback size. Measured at 375 pixels wide: 34-pixel coaches with SIX of them filling the visible line, where the plan promises four. It is fixed, and the test that watches it was made stricter FIRST so the fault was caught by a named failing assertion before the fix landed - the same lesson as every earlier version where a forgiving stand-in let a real fault through.

9. Measured on your own day, at 375 x 812 and at 1280 x 900, signed out of the cloud. At 375: the line the coaches stand on is 213 pixels, and a coach is 37, a link 18, a side stub 4 - four coaches and three links and two stubs come to 210, so exactly four coaches fill the line, which is the figure you gave. The clock stands at 215 pixels, the centre of the line. At 1280 x 900: a 723-pixel line, 151-pixel coaches, 30-pixel links, and again exactly four to the line, with the clock at the centre at 470. Wei's row read "Now: rests" with her five coaches at 4:00, 4:13, 4:30, 5:01 and 5:32, the first three red and the one at 4:30 ringed as the job running now; the links counted 31m and 62m, and five minutes later the same links read 26m and 57m with the clock label moving to 4:35 am. One tap on Person 4's coach turned it from red to green, said "Oil the pans and weigh the dough out at 4:00 am - marked as taken.", wrote the tick, and turned that row's line from "1 not taken" to "All done". A drag from the centre to 150 pixels read 4:13 am in the reading colour with every train's position unchanged to the pixel; letting go put the clock back at 215 and the label back to 4:30 am. The Clear press asked first, wrote the empty map rather than deleting it, turned every coach back to red, and made itself quiet again.

10. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. After all of the above - the taps, the clear, the drag, the clock moving, the window resized - your stored data was put back from a copy taken beforehand and compared: byte for byte the same 13,888 characters, your three saved days still reading 24, 4 and 24 pans, and only the app's own two keys left in the phone's storage.

11. Every rule that stands on it was proved load-bearing, not assumed. Sixteen faults were put back in all, each watched failing a test that names it, then restored byte-identically. Among them: with a coach sized as a plain quarter of the line instead of out of the budget it fails reading that four coaches no longer fill the visible line; with a wait between two jobs given no width it fails the round trip that reads every minute of a row back off its own strip; with the due test given an upper bound it fails reading that a coach which has passed the clock is no longer red; with the tap no longer deleting the key it fails reading that the last tick off left something behind; with the centre fixed instead of measured it fails the placement of every row; with the pan guard removed it fails the test that says the board's window cannot pan while the planner's still can; and - the two worth naming, because they escaped the first sweep - with the tap's save or the clear's save skipped, nothing failed at all until the tests read the STORE and not just the object the app holds in memory. A tick that moved on screen and never reached the phone's saved file looks perfect until the phone is opened again, and no test could see it. Both now fail by name. The suite is 1339 passing with none failing.

12. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The ticks live in the settings row your phones already sync, beside the calling switch, and that row learned one thing in this version: an emptied board is now SPOKEN rather than passed over in silence, so the other phone takes your clear instead of reading it as ignorance and putting every tick back.

CHANGELOG and the guide both carry all of this, both PDFs rebuilt and read back with PyMuPDF, and the app's own More screen reads Engine v184.

24 Sep 2026 - engine v183, REASSIGNMENT HAPPENS BY ITSELF, AND THE PERSON'S CARD SAYS SO (no database step). One sentence on one card, and nothing else in the app changed. You went looking for a reassign button on a person's card and there was not one, so the card now says why there is nothing to press: the day hands out its own work, so changing that person's hours or their ticks changes their job list on the spot, and a job that no longer fits them is given to somebody else. Your own words that settled it: "if you know that reassignment is automatic, then forget about my request for reassignment button, just need to mention work reassignment is automatic with worker card detail change." Nothing of yours is rewritten, and there is no database step.

1. What you asked, in your own words. "where is the reassign button?" - and then, having answered it

yourself: "you can put the reassign work button inside person's card", and then the correction that ended it: "if you know that reassignment is automatic, then forget about my request for reassignment button, just need to mention work reassignment is automatic with worker card detail change".

2. Why there was nothing to find. Reassignment has been automatic since v177, and it is automatic on that very card. The day gives out every job nobody has been named for, and the only two things it reads about a person are the two facts you type on their card: the hours they are here, and the modules they are ticked for. Change either and the day re-picks hands immediately - which is why the card never needed a button, and why a button there would have been a second way to say the same thing.

3. What the card says now. One sentence, in place of a control: the day gives out this work by itself, so changing the hours above or the ticks changes these lines on the spot, and a job that no longer fits is handed to somebody else. And what you placed by hand stays yours - a person you put on a module yourself, or a stretch you moved off somebody's row, is not moved by this card. That last half matters as much as the first: the card only ever moves the work the day was already choosing between.

4. The rule itself, written down the way you stated it. "in a module finding his person, module will look at person's with the least skill and fit its module. if fit then assign him first." So a module asks the person who FITS it - trained for that module, and inside the hours they are here - and between two people who both fit it asks the one trained for the fewest modules first, with the person who can work anything at all asked last of all, exactly as you confirmed it. Ties go to the lower person number. It is a preference and never a refusal: a job nobody trained is free for still goes to somebody, and the row says why.

5. Measured, on your own day, at 375 x 812, signed out of the cloud. Your card for Wei opened reading "11 min of work, in 3 places, and nothing collides", "Here all day", and her five jobs: 4:00 am Dimple and top, 4:13 am The oven swap and the bake, and the three stretch and folds at 4:30, 5:01 and 5:32. The new sentence was on that card, and the card carried no press about the work at all - the only button on it was its own close. Then one tick, "Mixing the dough in the tub", and four of Wei's five jobs left her on the spot: her list fell from five lines to one, and Person 3's rose from one to five. The day still stood at the same four people, so nothing invented a hand and nothing was lost - the work was handed on, which is the whole claim. Taking the tick off put all five lines back exactly as they were.

6. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day, and no word of your plan. Measured live: after working the card, ticking a module and taking the tick off, your stored data was put back from a copy taken beforehand and compared - byte for byte the same 13,888 characters, with your three saved days still reading 24, 4 and 24 pans, and only the app's own two keys left in the phone's storage. One trace is named rather than glossed, because the same shape was named in v180: taking the last tick off leaves twelve characters - an empty skills entry - where there was nothing before, exactly as your stored data already carries for merges. It reads as no training at all, it changes nothing about the day, and a tick you never press never writes it.

7. The sentence was proved true, not just printed. A card that said reassignment was automatic and did not do it would be worse than saying nothing, so the test reads the job list off the card itself after one tick. Three faults were put back and each was watched failing a test that names it: with the sentence deleted, the test fails reading that the card says nothing about the reassign button; with the training made unable to keep a job off anybody, it fails at the day's own end, reading that the day went on giving the free job to somebody trained for something else; and with the card no longer re-reading itself on a tick, it fails at the card's end, reading that the card went on drawing the day it was built from. The claim can break at either end and the two are told apart on purpose. The suite is 1313 passing with none failing.

8. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The card's sentence reads the same two facts v177 already stores in the scenario, and nothing new is written anywhere.

24 Sep 2026 - engine v182, THE SHADE IS ONLY WHAT YOU SET, ONE PRESS MOVES ONE BATCH, AND THE PRODUCTION-LINE NUMBERS ARE PROTECTED (no database step). Three things in one

version, all of them inside the Scenario planner except the last. First: the shaded band the app used to work out for you is gone for good. Her rule, in her own words: "Wei should not have any shade because there is no work time set for him, shade should just follow what i set, not other consideration." So the only shade a person's row ever carries is the two hours you typed on their card, and a person whose hours nobody has typed carries none - no band, no "Working ..." line, nothing. Second: "delta t on one batch of the module dont change the module batches, it should not be." It did - one press on the first batch of a module moved every batch of it - and it does not now. One press moves the batch under your finger and nothing else, and Back onto the line takes off that one batch's hold and leaves any other batch's hold standing. Third, carried over from what v181 measured and left unprotected: the numbers on More -> Production line are now guarded exactly like your saved days, so a newly set-up phone can no longer push its own defaults over the numbers you typed.

1. What you reported, in your own words. "build v182. And the working time shade is wrong", and then, while this was being built: "delta t on one batch of the module dont change the module batches, it should not be." Both were measured before anything was changed, and both were real.

2. The shade, and why it goes rather than gets quieter. v179 drew a band on each person's row worked out from their own day - their first job to their last, the gaps between included - and v180 made that band stand down on a row where you had typed hours. Both were answers to a question you had already answered yourself, and v180 left the app holding two opinions about one working day. The measurement that settled it was on your own day: the computed band was drawn 93 minutes wide, 4:00 am to 5:33 am, on a row whose five jobs come to eleven minutes, beside a card that already read "11 min of work". It was measuring the stretch and calling it the work. So it is deleted - from the day, from the row, from the card, and from the drawing itself - and what is left is the one thing you actually said: the hours you type. Typed hours draw a band; no typed hours draw nothing. The card still says "Here all day" for a person with no hours set, because that is the real answer for every person on every day you have ever built, and a blank where a fact should be reads as a fault.

3. The delta t, and the batch it belongs to. A press on a batch is a hold - the batch runs later than the line puts it, and it comes with the module above it when that module moves. Since v154 the first batch of a module had been read as the module's own offset, so holding it moved every batch of that module with it: on a module running four batches you pressed one bar and four moved. Batch three, pressed the same way, moved only itself - which is the answer you asked for, so every batch now behaves like batch three. Moving a whole module is still one press away and it is the honest one: the module's own start time on the module card, from which every other batch is spaced by the module's own pace. The words changed with the rule, so nothing on the screen claims a move the app no longer makes: the toast now reads "Batch 3 held back 5 minutes from where the line puts it", the tip on a held bar "held back 5 min", and the way back "Batch 1 back at 4:00 am - exactly where the line puts it."

4. The Production-line numbers, which v181 measured and left open. v181 closed the same family of fault for your saved days - silence must never delete - and named this one as not fixed: the numbers on More -> Production line have no "empty" of their own, so a brand-new phone could still push its defaults over what the cloud had. It is fixed here, and it needed no new stored field. The signature of a phone with no opinion about the line is a plan that still reads exactly the stock numbers; that signature is now asked of both sides, so a phone that has never touched the line takes the cloud's numbers on its first pull and carries them in its first push, a phone whose numbers are your own keeps them against a stock cloud and queues one publish to put them back, and a plan you deliberately set to the stock numbers is the one case that cannot be told apart from silence - by design, and said plainly rather than hidden.

5. Measured on your own phone, at 375 x 812, signed out of the cloud. Your own day first: four person rows, not one computed band anywhere in the app, not one typed band, and the word "Working" appearing nowhere on any screen; Wei's card read "17 min of work, in 3 places", "Here all day", and nothing else about her hours. Then hours typed onto Wei's card, 5:00 am and 9:00 am: stored as 60 and 300 minutes from your 4:00 am start, and drawn as one band 96 pixels along the row and 384 wide at your own 1.6 pixels a minute - 60 minutes and 240 minutes - with the computed band still absent and the card reading "Wei is set to be here 5:00 am -> 9:00

am, but Dimple and top runs to 4:06 am - a job that cannot fit inside those hours." Then the delta, measured by giving the oven four batches for the length of the test: its four bars sat at 0, 139, 278 and 418 pixels. Pressing + 5 min on batch 1 moved that bar 0 to 8 pixels - five minutes at 1.6 - its card read 4:00 am -> 4:15 am becoming 4:05 am -> 4:20 am, its tag read B1 delta t=+5, and batches 2, 3 and 4 did not move a pixel, at 139, 278 and 418 exactly as they were. Pressing + 5 min on batch 3 moved batch 3 alone, 278 to 286. Pressing Back onto the line on batch 1 put it back to 4:00 am with a plain B1 tag and said "Batch 1 back at 4:00 am - exactly where the line puts it", and batch 3's hold was still standing.

6. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. The copy of your stored data taken before the first press was put back afterwards and compared: byte for byte the same 13,888 characters, with the four batches used for the measurement gone and the oven back to the one batch you had it at, and with the two hours typed onto Wei's card taken back off again. Your three saved days are exactly the ones you had - 24, 4 and 24 pans. The only keys left in the phone's storage are the app's own two, and nothing here blocks a sale or reads an order.

7. Every rule that stands on it was proved load-bearing, not assumed. Eleven faults were put back in all, each watched failing by a test that names it, then restored byte-identically. Among them: with the computed band put back into the day's own model it fails reading that a person's row carries nothing but the hours she typed; with it put back into the row's drawing it fails that same test from the other end; with the rule put back into the styling it fails the third time; and with it put back onto the worker's board it fails reading that the board shows the hours she typed and nothing the app worked out for her. With the first batch's press carrying the rest of the module again, three named tests fail - including the one on the last module's own bar. Of the production guard: dropping the key from the guarded list fails four tests; making every plan count as somebody's opinion fails five; dropping the cloud-side guard fails two; and reading a plan that has no field at all as an opinion fails the test that says a cloud with no plan and a phone that has one is put back. One fault escaped on the first pass and is worth naming: restoring the old "Back onto the line" spread was NOT caught, because every test pressed the way back on a batch that happened to be the only one held - the tests had pinned the press and left the undo unpinned. A test was written that holds two batches of one module and takes one back, and the fault now fails by name. The suite is 1312 passing with none failing.

8. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The production guard reads a signature off the numbers your phone already syncs, and the batch hold has always lived in the delta your plan already carries.

24 Sep 2026 - engine v181, A NEW PHONE NO LONGER EMPTIES THE OTHER ONE (no database step). You reported that signing in on a new phone showed none of your saved scenarios. It was real, and it was the app rather than the phone. The cause is one word: silence. When the new phone pulled from the cloud it had nothing to say about your saved days, and the app read that silence as a decision to delete them - so the next thing it wrote back to the cloud left them out, and the cloud lost them. Three rules now stand on one sentence: silence means this phone has no opinion, and silence must never delete. A newly set-up phone now receives your saved days and the day you built, and its first push carries them back up instead of dropping them; and a phone that still holds something the cloud has lost puts it back by itself, with no press from you. A deletion you actually made is spoken out loud, so it is honoured and never refilled.

1. What you reported, in your own words. "i try log in with the new phone, the save scenario is not loaded? why". The "why" is answered below. This version is the answer to it happening again.

2. Why it happened. Everything you have set up in this app - your saved scenarios, the day you built, your to-do list, your wish list, the developer line - travels to your other phone in one single row called "settings". There is only ever one of those rows, on every phone, from the first day you install the app. A phone that has just signed in therefore already has a settings row of its own to publish, even though it has never seen the cloud copy, and the app decided which side to keep by the clock: whichever phone noticed last won. A new phone wins that comparison simply by being newer, and when it won, it wrote its own empty settings up to the cloud. The saved days were not emptied on purpose and were not refused - they were simply not there to write. A

push replaces the whole row, so a key the payload leaves out is a key the cloud loses. That is the fault: silencing a thing is not the same as deleting it, and the app could not tell the two apart.

3. The three rules, and they are one sentence between them. Silence means this phone has no opinion, and silence must never delete. First: if you empty a shelf on one phone, that deletion is now spoken out loud - the app says "the saved days are now an empty list" instead of saying nothing at all, so the other phone cannot mistake your decision for ignorance. Second: when a phone's own settings win, every one of those five keys it has no opinion about is taken from the cloud instead of being left out - taken into what it is about to publish, so its push cannot delete them, and taken into its own app, so a phone that has never had your days receives them on that same pull. Third: if the cloud has lost one of those keys and this phone still holds it, this phone queues one publish to put it back - no press from you, and never for a key you emptied on purpose.

4. What a newly set-up phone does now. Measured. Against a stand-in cloud holding your own saved days, the new phone receives all three of them and the day you had built, and the settings it then sends up still carries those same days - so the push that used to empty the shelf now leaves it whole. Measured again from the other end: with the cloud's settings row missing your saved days and your own phone still holding them, one pull from that phone queues exactly one publish carrying them, and a second pull queues nothing, because there is nothing left missing. And a shelf you really did empty is not put back - the empty answer is spoken, read as an answer, and left alone.

5. What this version protects, and what it does not yet. Read this part. This version protects your saved scenarios, the day you have built, your to-do list, your wish list and the developer line. It does not protect the numbers on More -> Production line. Those are a different kind of value, with no "empty" of their own - a brand-new phone's numbers can still replace what the cloud has, and that was measured, not guessed: the cloud's production plan was pushed over as a brand-new phone's defaults. It is the same family of fault one key over, it is not fixed here, and it is not dressed up as fixed. Until it is, check More -> Production line on both phones and set the numbers on the phone you trust; and if a number has already been replaced, the daily cloud backups are the way back, not a re-typing.

6. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day was changed by this version, and nothing here blocks a sale or reads an order. Measured live on your own phone: after unlocking the app, signing out of the cloud and opening the Scenario planner, your stored data was put back from a copy taken beforehand and compared - byte for byte the same 13,888 characters. Your three saved days read 24, 4 and 24 pans. That last figure is worth one line, because two earlier notes in this history said "24, 12 and 24": 12 is what the third day is set to aim at, 4 is what that day as built actually makes, and 4 is what the app has always computed and still computes. The number in those two notes was the mistake, not the app, and the earlier entries are left standing rather than quietly tidied.

7. Every rule that stands on it was proved load-bearing, not assumed. Eleven faults were put back in all, and each was watched failing by a test that names it, then restored byte-identically. Taking the shelf the phone is silent about out of what it publishes, so the push deletes it again, fails reading that a phone that never had her saved days receives them. Taking the same value out of her own app instead, so the phone stays empty, fails the same test from the other end. Letting the phone overwrite its own days with the cloud's, and refusing the cloud's empty answer when she has deleted a shelf on purpose, each fail the test that says a phone with its own days keeps them and a shelf she emptied stays emptied. Dropping the repair of a short cloud row fails reading that a phone that still has her days puts them back. Making that repair fire on every pull fails reading that a cloud row with nothing missing queues no publish. Dropping the built day from the guarded keys fails reading that a phone that never built a day receives the one she built. Removing the spoken deletion fails reading that emptying her shelf is said out loud, and making every phone speak - including one that never held the key - fails reading that a missing field is not announced as an empty shelf. Two of the eleven were faults in the first draft of the fix and were found by this sweep rather than by reasoning: one change was unfalsifiable as written, because the cloud's value was copied into her own app before the comparison the fault was supposed to break, and one guard turned out to be protecting a real case - a hand-edited or corrupted file - which had no test until the sweep asked for one. The suite is 1309 passing with none failing.

8. No database step. Not one stored field is added or changed on the server, so there is nothing to run in Supabase. The three rules are read by the sync code your app already runs, and they act on the same five keys that already travel between your phones.

23 Sep 2026 - engine v180, THE HOURS YOU TYPED TAKE THE SHADE OFF, AND THE LAST BAR OF THE DAY CAN BE HELD (no database step). On a person whose working hours you have filled in on their card, the shaded band the app worked out for you is no longer drawn. Your two clock boxes are the answer for that person, and a window you typed with a second, computed window drawn inside it was two answers to one question - which is what you saw and reported. Nothing else about the shade changes: on every person whose hours you have NOT filled in, it is drawn exactly as it was, and on every person it is still named in words - on their row, in the tip and on their card. The bar at the very end of your last module now carries the same readings and the same pair of presses as every other bar, so it can be held like any other; the card it opens is still the day's own card, with the whole day read off it underneath.

1. What you reported, in your own words. "dont shade if the person have indicated work time", and then, so there was no doubt which screen you meant, "and those person with work time specify, the shade is wrongly indicated". Both true, and neither was your phone.

2. What was wrong. The shade added in v179 is the stretch of the day a person's own jobs actually occupy, worked out from the day. It was drawn on every person, including the ones you had already filled the two hour boxes in for. So on those rows two bands lay on top of each other: the hours you typed, and a second, computed band sitting inside it. Measured on a row with 5:00 am to 9:00 am typed in, at 1.6 pixels a minute: the typed band ran 96 pixels along the row and 384 wide, and the computed one was drawn on top of it, starting later and ending earlier wherever the work did not fill the hours. Two different lengths for one person's working day, and no way to tell which one the row meant.

3. What it does now, on your rule. The moment you have said the hours for somebody, that person is not shaded at all. Your typed hours are the only band on their row, and they are read as the answer. Where you have NOT filled the boxes in - which is every person on every day you have ever built, including all of today's - nothing changes at all: the shade is drawn from their own first job to their own last exactly as it was in v179, because there is no typed answer for it to argue with. The rule is one line, held in the model rather than on the screen, so the row, the worker's board and the person's card cannot come to different conclusions.

4. And the words stay, deliberately. Taking the drawing away must not take the fact away. The working stretch is still named in words everywhere it was: on the person's row ("Working 4:00 am -> 5:33 am"), in the tip, and on their card, where it sits beside the hours you typed so you can still read the two against each other. That is the half of this release you may not have asked for, and it is deliberate: you said do not SHADE it, not do not SAY it. If you would rather the words went too on a person whose hours you have typed, that is a one-line change - say the word.

5. And the other thing you reported, which I had explained away. "the last module batch pop up, still dont mark his delta?" - the very last bar of your day. Tapping it opened the day's own card and nothing else: it carried no batch number, no delta t, no cycle reading and no way back, while every other bar on the chart carried all four. In v179 I answered that this was by design, because since v152 that bar is where the day is worked backwards from, so the card it opens is the day's own card. You reported it again, so I measured it rather than explaining it again - and it was real. A bar you can see and tap but cannot hold is missing the one thing every other bar on your chart has, and the day's own card does not need that to be true.

6. What that bar does now. The card underneath is unchanged: it is still the day's own card, still headed the end of your first batch, still naming the day's first batch, the day's finish and every module on it. What is new sits on top of it - the four readings every other bar carries (which batch it is, the delta t, the cycle, the pans), the same pair of presses, + 5 minutes and - 5 minutes, and Back onto the line. A hold there is a hold on that whole module, exactly as a hold on batch 1 of any other module is. Measured on your own day at 4:00 am: five minutes on the last bar moved that module five minutes later. Its row read "starts 4:05 am", the tag on its bar read B1 delta t=+5, and its card read Batch 1, delta t = +5 min, 4:05 am to 4:17 am, where it had been 4:00 am to 4:12

am. Not one module above it moved a minute. Nor was any of them given more room: every one of their latest starts still read 4:00 am, which is where they already were, and the day card's own line - that every module of your line is already as late as it can go - was as true after the hold as before it. The day's own finish did not move either: 6:03 am both times, because the day finishes with whichever module ends last, and on your day that is The rests and the stretch and folds, which the packing never delays. The pans a day did not change at all: 6 before and 6 after. Pressing Back onto the line put the batch back to 4:00 am and the tag back to B1 with no delta t, and the round trip left your stored data byte for byte the same.

7. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Nothing here is stored at all: the shade is worked out from the day each time the screen is drawn, and the only thing that decides whether it is drawn is whether you have filled the hour boxes in. Measured live, after typing hours into a person's card, taking them out again, and opening and shutting the cards: every field of the day read exactly as it had before, with one trace left behind, and it is named here rather than glossed over. The first time an hour is typed, the day grows an empty shifts key, and taking the hours back out leaves the key standing at nothing. It is twelve characters, it reads as no hours at all - as does an absent key, which is the one thing it can never mean anything else - and it is the same shape your own stored data already carries for merges. So the day is not byte-identical to itself after a round trip; it is identical in every field, plus those twelve characters of nothing. The copy taken beforehand was then put back and compared: 13,888 characters, byte for byte, and your three saved days read 24, 12 and 24 pans.

8. Every rule that stands on it was proved load-bearing, not assumed. Ten faults were put back in all, and each was watched failing by a test that names it, then restored byte-identically. With the rule dropped so everybody is shaded again it fails reading that a person whose hours you typed is still wearing a shade; with the shade dropped everywhere it fails the other way, reading that a day with no hours typed anywhere lost its shade; with the rule inverted so only a typed person is shaded it fails both; with the row drawn from the work again, past the rule, it fails reading that the two bands are back; with the words silenced along with the drawing it fails reading that the fact left the row; and with the shade made to answer to the whole day rather than to the person it fails reading that one person's hours took the shade off somebody else's row. Of the last bar: with its card taken back to the day's card alone it fails reading that the card no longer says which batch it is; with the day's own reading dropped off it, the other way; and with the two pairs of presses stopped saying which one moves the whole day it fails reading that the card no longer names the pair. The suite is 1298 passing with none failing.

9. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. You do not need to type your hours again on the other phone: the shading is decided by the hours you already had.

23 Sep 2026 - engine v179, WORKING HOURS AND THE DELTA T (no database step). Each person's line is now shaded across the hours they are actually working - drawn from their own jobs, from their first of the day to their last, with nothing for you to type. And the delta t is back across the whole chart: a press on batch 1 of a module that keeps its own time is a hold again, with its own way back. Your widest scale gives up the delta t before it gives up the batch number, on your own rule.

1. What you asked, in your own words. "can you shade their working hours in their line?" - and, asked which of two ways you meant it, you chose the one that needs nothing typed. And then, separately: "i ask for delta time, that function is not worker across the chart. And the ruturn to original button is missing?" Both were real, and neither was your phone.

2. The shade, drawn from the day itself. Each person's row carries a wash from their own first job of the day to their own last, the gaps between jobs included - somebody who is back in an hour has not gone home. It is worked out from the day rather than typed, so it is drawn on every day you have ever built, including the days with no hours set on anybody. Where you HAVE typed the two hours on a person's card, those stay exactly as they were: the typed hours are the outer band and the working shape sits inside them.

3. Measured live, and one thing to expect on a row. On your own day at 1.6 pixels a minute, Wei's shade was drawn at 0 and 149 pixels wide, which is 93 minutes of a 4:00 am day - 4:00 am to 5:33 am - against her five jobs. A job shorter than four pixels is still drawn four wide so you can see it and tap it with a thumb, so on a row whose last job is one minute the last bar reaches a couple of pixels past the shade. Measured live, Wei's last job

is the fold, one minute, drawn at 147 for 4 pixels and ending at 151, where the shade - which reads the minutes themselves - ends at 149. The shade is the true length of the work and the bar is the true size of your thumb, and the difference is only ever that.

4. The delta t, and why it had gone. A press on a batch used to be a hold measured from where the line puts that batch - a delta t. Since v154 that had been narrowed to batches below batch 1, and to batch 1 only of a module set to follow the one above. Every module of your own day is set to "Its own time", which is the one case the narrowing left out, so batch 1 of every module wrote an absolute start time instead: no hold was ever written, no delta t was drawn anywhere on your chart, and the card's own way back - which exists only while there is a hold to take off - disappeared with it.

5. It is back, and the way back is called "Back onto the line". That is its name from v151, and it can be renamed in one line if you meant the other wording. Measured live on your own day at 1.6 pixels a minute, pressing + 5 min on batch 1 of Dimple and top: the card read Batch 1 and delta t = +5 min with 4:05 am to 4:11 am, the bar's own tag read B1 delta t=+5, the note read "Dimple and top held 5 minutes later than the time you gave it - every batch of it moved with that.", and the three modules below it did not move a minute. Pressing Back onto the line put the batch back to 4:00 am and said so: "Every batch of Dimple and top is back at the time you gave it." Going there and back left your stored data byte for byte identical at 13,888 characters, because the hold is kept in the delta itself and none of your stored start times was rewritten.

6. The badge now knows when to keep quiet. Your own rule: "if the scale is too wide to show batch no. and delta t then forgo delta t". So the batch number - the thing that tells one batch from another - is never given up, and the delta t is printed only above the widest stop. Measured live at the widest reading, 72 pixels an hour: the tag read exactly B1 while the bar kept its own tint, which takes no room and still answers which of the batches is held, and the tip on that tag still named the hold in full.

7. Why the last module's batch 1 is different - the question you asked while this was being built. It behaves differently on purpose, and it has since v152: tapping batch 1 at your LAST module does not open an ordinary batch card at all. That batch is the moment the whole day hangs from - it is where the day is worked backwards from - so its card is the day's own card, headed "the end of your first batch". Measured live, tapping batch 1 of Cutting and packing read "Cutting and packing - the end of your first batch", named the day's first batch ending at 4:12 am and the day finishing at 6:03 am, and listed all eight modules. On the FIRST module it is the other way round: batch 1 there has no module above it to be held back from, so a press is simply that module's own start time. Everything in between is the ordinary batch card with the delta t and its way back.

8. And where "Put my start times back" lives, since you looked for it. It is on that day card and it is a one-shot undo: it appears only once you have pressed "Pull them back to their latest start", because there is nothing to take back until something has been pressed. Measured live on your own day, every one of the eight modules was already ticked as late as the line allows, so the card said so - "Every module of your line is already as late as it can go - there is nothing to pull back." - and offered neither press, which is why you could not find it there.

9. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day. Measured live, after working through the batch card, both presses and the widest scale, your stored settings were put back from a copy taken beforehand and compared: byte for byte the same 13,888 characters, and your three saved days still read 24, 12 and 24 pans.

10. Every rule that stands on it was proved load-bearing, not assumed. Thirteen faults were put back in all, and each was watched failing by a test that names it, then restored byte-identically. The shade: with it drawn from the whole day instead of from that person's own jobs it fails reading whose work it covers; with the bars painted before the shade it fails reading that the bars no longer cover the wash; with the shade given a stacking order it fails reading that the hours climb over the module titles, which is the fault this screen already had once. The delta t: with the hold narrowed back to the v154 rule it fails reading that a press on batch 1 wrote a start time instead of a hold; with the module's other batches left behind it fails reading that the hold did not take them with it; with the widest scale still printing the delta t it fails reading that the batch number was given up for it; and one fault in the tests' own stand-in screen was found and fixed on the way. The suite is 1294 passing with none

failing.

11. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. Both facts live inside the scenario, which already syncs whole, so they travel to your other phone with the rest of your plan.

23 Sep 2026 - engine v178, REASSIGNMENT (no database step). "Share them out" now really hands the whole day back. It used to clear every module and leave the stretches you had placed by hand exactly where they were - and a hand-placed stretch is read first, ahead of the module's own person, so the press said one thing and the chart drew another. Both full-day presses now take those stretches back, and they say how many they took back, because taking back a decision of yours is a change you are told about.

What you asked, in your own words. "how to rerun job assignment after we set the more details?" - and then, once that question had been measured and answered, "build v178, Name it as REASSIGNMENT". The answer to your question is the People box, and measuring it is what found the fault below.

1. The fault: the press promised a day it did not draw. The People box holds three answers, and two of them redraw the whole day's hands at once. "Share them out" cleared every module - but a stretch you had placed by hand is kept on the module as its own entry (that is what makes a hand-over move one stretch and not a whole row), and it is read BEFORE the module's own person. So the day kept your hand-placed stretch standing while the press said it had handed everything back. Measured on your own day: after the press the toast said the day had been shared out, and the box's own closed words still read "Your own" - and your Dimple and top was still sitting on Wei, where you had put it.

2. What it does now. Both full-day presses - "Share them out" and "One a module" - take every hand-placed stretch back along with everything else. The stretch is not deleted quietly either: the press names how many it took back, and a day with none adds nothing to the sentence, because a count of zero is not news.

3. Measured live on your own day, starting at 4:00 am at 1.6 pixels a minute, with your one hand-placed stretch on Wei. Pressing "Share them out": the box went from "Your own" to "Sharing them out", your hand-placed entry was gone from the record, and the toast read "Sharing them out - every job reassigned to as few hands as can cover the day. 1 stretch you had placed by hand goes back on the day's own arrangement." Dimple and top moved off Wei's row onto the row the day picks for itself, and every other job on the day was untouched. Pressing "One a module" on the same day: six people, one to the six modules that hold hands, the box reading "One to a module", and the same sentence naming the one stretch it took back - and the stretch no longer overrides the person that press had just given that job.

4. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario was changed by building this, and nothing here blocks a sale or reads an order. The only writes are the ones your own press makes. Measured live: after both presses and a full pass of opens and closes, your stored data was put back from a copy taken beforehand and compared - byte for byte the same, 13,888 characters - and your three saved days still read "No fridge, 1 person" 24 pans, "My sister proposal 21/9/2026" 12 pans and "One baker day" 24 pans.

5. Every rule that stands on it was proved load-bearing, not assumed. Eight faults were put back in all, and each was watched failing by a test that names it, then restored byte identically. With the stretch left standing, the test fails reading that the old arrangement is still on the chart; with the count left out of the sentence, it fails reading what the press said; with the count worked out after the day had already been cleared, it fails the same way; with "One a module" left holding the stretch, it fails reading that one job is short of the person the press gave it; with the empty entry written back as an empty map instead of no entry at all, it fails reading that nobody has one spelling; and with the count written always in the plural it fails on the first hand-back. The suite is 1287 passing with none failing.

6. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. Everything here is inside the Scenario planner, under More.

23 Sep 2026 - engine v177 (no database step). Each person now has two facts of their own, and both are

yours to state once: the hours they are here, drawn as a band along their row and a real bound on who the day may hand work to, and which modules they are trained for. A job nobody is named for goes to somebody who can work it, inside their hours, and between two of those the day asks the most specialised person first. Nothing is refused and nothing is hidden: a job that lands outside somebody's hours or training is shown and said in words.

What you asked, in your own words. "back to the scenario planning. 1. I need each person to have the start work time and end work time, 2. which module they. are dedicated for, and only taking over others job, as priority 1,2, or3." And then, while this was being built, the two sentences that settled the shape of it: "a person we should be able to specify their skill, by module, can be more then one, by module", and "module should select the one that specialised". Of the takeover order you said the thing that made it simple: "the person card supersed the module person assignment. maybe no priority rank is needed to be sinple". So there is no 1, 2, 3 rank anywhere: the card says what each person can do, and the day picks by itself.

1. Their hours, counted from your day's start and not from the clock. You asked for a start time and an end time, and told me how to read them: "the working hours is relative to the chart, not exact hours, end also relative to the chart". So the two boxes are clock times you can read at a glance, and what is stored is the minutes from the start of your own day - which means moving your day's start slides every person's hours with it and nothing is retyped. Measured live on your own day, which starts at 4:00 am: typing 05:00 and 09:00 stored 60 and 300, which is exactly five and nine o'clock less four.

2. An empty box is a real answer, and it is "here all day". Leave both empty and that person is here all day, which is what every person on every day you have ever built already is. The hint under the boxes says so in one line, and it also says which clock the times are read against, because two empty boxes with no explanation would read as two boxes you had not filled in. Typing the whole day by hand - 4:00 am to midnight - is stored as no entry at all, so "here all day" has one spelling rather than two that mean the same thing and could ever drift apart.

3. And a pair that comes out back to front is refused and SAID. An end earlier than the start is a mistake, and the app now puts both boxes back to the hours the day is actually using and says which way round they go. A box that took a number and quietly did something else with it is the kind of thing that reads as a fault for weeks, and this app does not do that.

4. What the hours are for. Your own aim, in your own words: "the person that have the most consistant work is the person handling oven, colling and packing, he only help do others when he has free slot, by one person not doing too many functions, it is easier to train that person as will. Probably can employ part-timers." So the hours and the training are what the day reads when it hands out a job nobody is named for - and they are a PREFERENCE, never a refusal. A job nobody qualified is free for still goes to a free pair of hands, and the row says why.

5. Which modules each person can work, and the specialist is asked first. The person card now carries a tick for every module that actually holds hands, so there is never a tick against a module switched off or against the proofer, which would be a tap that does nothing. Only six of your eight modules hold hands on your own day, and those six are exactly what the card offers. Tick nothing and that person can work anything, as they do today - the same absence rule this app already uses for automatic cycles and merged jobs. Tick one or more and the day only ever gives them those.

6. And between two people who can both do a job, the more specialised one is asked. Your own rule, in your own words: "say one person have all skill, person2 have have the skill, person3 have only 1 skill, module should select the one that specialised". Somebody trained for one module is asked before somebody trained for five, and somebody who can work anything is asked last of all - which is what leaves the specialist on their own job and the generalist on what is left over.

7. Measured live, on your own day. With everybody untrained and no hours set, your day packs to four people and 58 minutes of hand work, and the three stretch-and-fold slots all sit on Wei. Ticking Person 4 for "The rests and the stretch and folds" and nothing else moved all three folds off Wei and onto Person 4, even though Wei is

the lower number and was free at that minute - the specialist was asked first, exactly as you asked for. Wei's row fell from five jobs to two and the other rows were untouched. Unticking put the day back to where it was, to the job.

8. Their hours, drawn on their own row. Each person's hours are painted as a wash along their own row, under everything the day draws on it, so you can see at a glance when somebody is here without reading a number. It carries no stacking order of its own on purpose: a band given one would climb over the pinned module names, which is the fault this screen already had once. Measured live at a phone's width on your own day: a shift of 05:00 to 09:00 was drawn at 96 pixels along the row and 384 wide, which is 60 minutes and 240 minutes at your own 1.6 pixels a minute. It takes no taps at all, so it can never swallow one meant for a bar over it.

9. And when a job cannot fit somebody's hours, the card says so in words. This is the part that matters most to the way you work, because it is the same rule this app has always kept: a plan that does not fit is shown, never quietly rearranged. Measured live: with Wei set to be here 4:00 am to 4:05 am while the dimple runs to 4:06 am, the card read "Wei is set to be here 4:00 am to 4:05 am, but Dimple and top runs to 4:06 am - a job that cannot fit inside those hours." The same sentence appears on their row's tip and on the Production line board, all three printed from one funnel so they can never disagree.

10. And a job outside somebody's training is said in the same place. Somebody given a job they are not trained for is told so in the same words and the same place: the job is named, the times are named, and the note says it is said here rather than the job being quietly moved. The card also reads back what they are trained for - "Trained for 1 module: The oven swap and the bake" - so the answer you typed is the answer you can see.

11. On the Production line, the worker's own hours. The board a worker reads already says what they do next and how long they have; it now also says when they are here, and names any job that cannot fit inside those hours. Withholding that while printing their jobs would be perverse - when you are here is the first question of a working day. The board is still a READING and writes nothing: measured live, a tap on a person's row on the board left your stored settings character for character identical, and the board offers no hour boxes and no ticks.

12. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day was changed by this version. Measured live: after a full pass of presses on the planner - the two hour boxes typed into and cleared, ticks set and taken off, cards opened and shut - your stored data was put back from a copy taken beforehand and compared, byte for byte the same 13,888 characters, with the two new keys absent again because writing them is something only a press you make can do. Merely opening the screen adds nothing. Your shelf itself is untouched: No fridge, 1 person still reads 24 pans, My sister proposal 21/9/2026 still reads 4, and One baker day still reads 24. And the two chart windows kept their pan at 300 across two of those presses, so a card write no longer throws either window back to the start of the day.

13. Every rule that stands on it was proved load-bearing, not assumed. Eighteen faults were put back in all, and each was watched failing by a test that names it, then restored byte-identically. Two of the eighteen were faults in earlier drafts that the fault sweep found, and both were unfalsifiable as first written - a test that could not fail - which is the same lesson this file has carried before about a stand-in being more forgiving than the real thing. Among the rest: with the band given a stacking order the test fails reading that the hours climb over the module titles; with the band drawn last it fails reading that the bars no longer cover the wash; with the boxes storing a clock reading instead of minutes from the day's start it fails reading where the band was drawn; with an untick leaving an empty list behind it fails reading that unticking is an answer and not a missing entry; with the copy of a day not carrying the two keys it fails reading that the hours of the day you left are still on the day you opened.

14. No database step. Not one stored field is added or changed on the server, so there is nothing to run in Supabase. Both new facts live inside the scenario, which already syncs whole, so they travel to your other phone with the rest of your plan and need no new table and no new SQL.

23 Sep 2026 - engine v176 (no database step). The Production line is now a board a worker reads, not a

form you fill in: the day chart from your plan, a clock line showing where the real time falls in it, and a strip saying what each person does next and how long they have. The screen used to ask you for 21 numbers; it now asks for 12, and 8 of those are folded away.

What you asked, in your own words. "can we simplify the productioon portion, it ask for alot of key in" - and then the shape of what you wanted: "maybe one function for production line, 1. copy the chart into Prodcution line, but should not be editable, it become a dashboard for worker, give the worker good inform of what next for them, how long, and others, you can be creeative". And of the boxes on the screen you said the thing that made this simple to build: "must of the question are answered in scenario plan". They are, and that is where the board now reads them from.

1. The day chart is carried across, not drawn a second time. The board and the Scenario planner are the same drawing, produced by the same code, so the two screens can never disagree about a day - not its modules, not its batches, not its times, and not who is on what. What the board drops is everything that changes the day: the Scale buttons, the People box, the backwards chip, the undo chip, the Start the day now button, New module, and Save changes. Measured on your own day: the planner wears 8 buttons above the chart, and the board wears none of them - with the two cards shut the whole screen carries two buttons, and those two are the fold heads at the foot. What the board keeps is everything that only reads: you can pan the day with a right press, the clock follows your finger along the ruler, and the tips on the module names still open. A board that could not be panned would be a worse board on a long day.

2. A clock line, drawn with no press at all. A board nobody has touched yet still has to say what time it is, so the line is placed the moment the screen opens and kept in place once a second. It is drawn exactly where the real clock falls in your day: at 96 pixels an hour on your own day, a press of nothing at all puts the line at 156 pixels for your 4:00 am start, and the clock it reads out matches the ruler under it. Measured live on your own day: at 5:00 am the line stands at 252 pixels and reads 5:00 am - which is 156 plus exactly one hour of 96. At 4:30 am it stands at 204 and reads 4:30 am. Before your day begins it parks at your day's own start and says "day starts 4:00 am"; after your day has ended it parks where the day ended and says "day ended 6:03 am". It never wraps around to the other end of the ruler, which is the tempting wrong answer and would have put the line in the middle of a day that is not the one on the screen.

3. What is next, in clock times and not countdowns. Above the chart the board says, for each person, the job they are on and the one after it, with the clock time each starts and how many minutes it runs: on your own day at 5:00 am that reads "Wei - the rests and the stretch and folds, 5:01 am, 1 min, then the rests and the stretch and folds at 5:32 am". Clock times and not a countdown on purpose: a countdown means rewriting a box somebody is reading every single second, and this app has a rule about that - a redraw must never move what you are looking at. On a day that is not shared out to people, the strip falls back to one list of what the line does next.

4. And the strip answers AGAIN as the clock runs, without touching the chart. This one was found while measuring, not reported: the strip was drawn once when the screen opened, so a worker opening the board at 4 am would still be reading "next: mix at 4:01" at eight in the morning. It now looks again whenever the clock has changed the answer, and only then - inside the working day it redraws at the exact minute a job starts or ends, and never in between. The chart itself is not redrawn, which is the planner's own rule and the reason a worker's scroll is never lost. Measured live across a job boundary: at 9:00 am the strip names Cutting and packing, at 9:10 am it names the oven swap, and the number of times the strip has been rebuilt does not change between them.

5. Three things about a day the clock has already gone past. The first was a real fault, found by opening the board in the evening. The strip had one fallback sentence for two different empty lists, so a day that finished at 6:03 am was being described as "Nothing on this day needs hands" - which is a different thing entirely, and the opposite of true. A day whose jobs have all been and gone now says so: "The day is finished - every job on it has been and gone." The second is the offer to make a sound. The board rings people one minute before each job, and it rings by working out from the clock which jobs are due - but on a day that is already over there is nothing left due, so switching calls on would ring for nobody. That offer is now withdrawn on a finished day rather

than left promising a sound it cannot make. Measured live at 7:16 pm on your own day: no Start calling button, and the line under the strip reads "The day is finished". The third is the other half of that, and it matters more: calls you turned on earlier can ALWAYS still be turned off. Stop calling stays on a finished day, so a sound started before cannot be left ringing with no way to stop it.

6. What left the screen, and where the numbers went. Twenty-one boxes became twelve. Nine of them are now read off the day you built under More -> Scenario planner - how many pairs of hands, the pans a tub makes, oiling and weighing out, the dimple, the packing, the oven and the bake-and-swap - so those questions are asked once and answered on both screens. The twelve left are the ones genuinely yours: four behind the first fold (how long you will bake for, the pans you own, what your proofer holds, and the pace between batches you would like), and eight behind the second fold (the rests, the fold, the two proofs, the cooling, the clock you want the first batch at the oven, and the minutes of soft tolerance around a start). Both folds start shut and open in place, so opening one never redraws the chart above it. The line card still answers while it is shut, in one line: measured on your own day, "44 pans today · the proofer is the wall that holds it back", worked out from the day rather than typed anywhere.

7. And the timetable card is folded, which is where it belongs. The backwards timetable - the last moment each stage may start - was left on this screen on your word ("you decide, that is your baby") because it is worker information: it is the one thing on the screen that says when the dough has to be in the tub. But it is a page of advice, not a glance, and it was sitting out on the board in front of a worker who came to read what is next. It now lives behind its own fold head. Measured at a phone's width with both cards as they come: the timetable card is 64 pixels tall, its head and nothing else, where it was 943 while the timetable sat outside it. That one change alone takes the whole screen from 2005 pixels to 1126.

8. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved day changed by this version, and nothing here blocks a sale or reads an order. The board is a READING of your day, and every number on it is computed rather than stored - the box you type a number into writes the one number you typed and nothing else. Measured live, going through a full pass of presses on the board: a bar tap, a module tap, a right press, and both cards opened and shut twice each, with your stored data compared before and after. Your data reads 13,888 characters, and it is byte for byte the same characters it read before the first press.

9. Every rule that stands on it was proved load-bearing, not assumed. Nine faults were put back in all, and each was watched failing by name. The strip: with the finished-day sentence taken off, the test fails reading that a finished day falls back to the sentence for a day with no hands on it; with the calling offer always drawn, it fails reading that a finished day offers a sound it cannot make; with the foot dropped whenever the day is over rather than only when calls are off, it fails reading that a sound already started cannot be turned off; with the strip never rebuilt, it fails reading that the strip still names a job the clock has passed; and with the live minute taken off the strip's own signature, it fails reading that the parked sentence no longer names the hour it is standing in. Two more put a null where the parked sentence goes, and each was caught by the word "null" appearing on the screen. The fold: with the timetable put back outside its fold, the test fails reading that a shut card is already showing the clock times its caret says are behind it; and with the test's own reader taken back to one that walks straight through a hidden box, the same test fails - the reader had been reading text the browser never paints, which is the same lesson this test file already carried about a stand-in being more forgiving than the real thing. Every one was restored byte-identically afterwards. The suite is 1266 passing with none failing.

10. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase. The day the board draws is the one already synced, and the four numbers behind the first fold were already synced with the rest of your production plan.

CHANGELOG and the guide both carry all of this, the app's own More screen reads Engine v176, and the Production line names itself as the board it now is.

23 Sep 2026 - engine v175 (no database step). Two things you reported: the scenario windows no longer jump back to the start when you press something, and the right-button drag now works while a card is open on the chart. Both inside More -> Scenario planner.

What you reported, in your own words. "Few problem of screen jump here and there. One obvious one is the scenario windows, when i click, window reset. And right button drug dont work". Both were real, and neither was your phone.

1. Why a press threw the windows back to the start. When you press something that changes the day - a batch card's + 5 min, the People box, the scale, opening a saved day - the screen is drawn again from scratch. That has always been how it works, and it is what keeps the whole panel agreeing with itself. But a box drawn again is a NEW box, and a new box starts its own scroll at nothing. So every one of those presses put both windows of the day chart back at the far left and the top, and would have put the list of saved days back at its first row with them. Measured live at a phone's width with the modules' window panned 300 pixels: one press of + 5 min left both windows reading 0.

2. What it does now, and where you were is put back only after the measuring is done. Where each of the three sliding boxes was - the two windows and the list of days - is read before the screen is drawn again, and written back afterwards. The order is the whole of it: the list of days is measured to work out its height, and a row's foot moves when the list under it slides, so measuring a list that had already been slid back down would come out short by exactly the amount it moved and would cut your fourth day in half. Nothing is slid until the measuring is finished, and the measuring all happens before anything is drawn on the screen, so you never see the day sitting at its start either. Measured live on your own day: panned 300, one press of + 5 min, and the window reads 300 on the day drawn after it - and both windows are put back by their own numbers, never one number shared, because they pan together and a redraw is not the place to find out they had drifted.

3. Why the right-button drag did nothing with a card open. A card is drawn on a full-screen sheet that sits over the whole chart, so with a card open your right press landed on the sheet, not on the day behind it - the day you were actually looking at could not be moved at all. Measured live at a phone's width with a batch card open: a right press 200 pixels into the chart, dragged 60 to the left, left the window reading 200 exactly where it started. The sheet carried no gesture of its own, so nothing was being taken from anything else.

4. What it does now. The sheet answers the same right press the two windows do, and which window moves is worked out from where your hand is - the upper window for a press on the upper half of the chart, the lower one for the lower. Your own card is left alone: a right press on the card, or on any of its buttons, moves nothing, so the day cannot slide about while you are reading the card that describes it. The left button is untouched and still opens and works cards. Measured live with a batch card open: a right press 60 pixels to the left moved the modules' window from 200 to 260, the same press on the people's window moved that one from 260 to 310, the two windows came back into step at 310 on the next frame, and the same right press on the card itself moved nothing.

5. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before. Your stored data was put back from a copy taken before the measurements and compared: byte for byte the same, and your three saved days still read 24, 4 and 24 pans.

6. And the rule that stands on it was proved rather than assumed. Nine faults were put back in all and each was watched failing by name, then restored byte-identically: with the putting-back left out entirely the test fails reading that the press threw the windows back to the start; with only the modules' window put back it fails the same way; with the two windows collapsed into one number it fails; with the putting-back done before the measuring it fails reading that the shelf was capped short of where the fourth day ends; with the class name matched dot and all it fails (this is the fault that was in the first draft of the fix, and the test caught it); with the sheet never wired it fails reading that a right press over an open card did not move the day; with the sheet always panning the upper window it fails on the lower press; with your own card not spared it fails reading that a right press on the card dragged the day behind it; and with the left button let through it fails the older drag test, which says the left press is the one that opens cards. One more thing was found and fixed on the way: the tests' own stand-in screen did not move a box when a box it sits inside was scrolled, which is the opposite of what a window does, and that gap had been hiding whether the measuring happened before or after the sliding. The suite is 1220 passing with none failing.

7. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v174 (no database step). Two things you asked for: the modules' window no longer lets the day's writing come up between the module titles when you pan it right, and the Your scenarios card now shows four days at a time with the rest reached by sliding the list. Both inside More -> Scenario planner.

What you asked, in your own words. "the windows when panning right, some of the chart writing shown in between module titles" - and "now i had 8 scenario, we have to make the secenario just shown 4, the rest shown by slider". You were right about both, and neither was your phone.

1. The writing between the titles, and why panning right brought it there. The module titles on the left of each window are pinned: they stay put while the day slides under them, which is the whole reason you can read which row is which at any hour. What was wrong was the ORDER they were stacked in. Every B-number, every band of your hands and every line down the bars was drawn at or above the title's own level, and when two things sit at the same level the one drawn later wins - and the day is always drawn after the titles. So nothing had to go wrong for a number to appear over a name; the day simply took its turn on top of it. Measured live on your own day at a phone's width, with the chart panned 25 pixels: nine B-numbers and six bands of your hands were the topmost thing at their own position inside the 156-pixel title column.

2. What it does now. The three things that can be stacked are written down once, on the chart itself, as three levels: the titles lowest, the clock strip above them, and the readings you take highest. The clock strip has to outrank the titles because the titles scroll under it, and a reading has to outrank the clock strip because the reading's own top line lands on the strip. Everything the DAY draws - its marks on your bars, the B-numbers and the names on them, and a module's own note - is deliberately under the titles, so no pan position can bring one of them up over a name. Measured live on your own day across six pan positions, 1,400 points sampled inside the title column: not one of them reads as the day, while the same eight B-numbers and five bands are still drawn exactly where they were, under the titles where they belong.

3. And the clock you take is refused where the day is not in front of you. The same fault had a second half. Pressing the clock strip over the title column could still read a minute and draw the hairline and its number across the titles - because the strip's own ruler scrolls with the day, so once you have panned, the strip's left edge has travelled behind the titles and the arithmetic no longer knows where the titles end. It now measures where the day actually starts rather than trusting the strip's edge. Measured live with the day panned 25 pixels: pressing at 140 pixels across - inside the titles, which end at 183 - reads nothing and draws nothing, while the same press 17 pixels further on reads 4:26 am as it always did. On a day that is not panned nothing about this changes at all.

4. The shelf now shows four days and slides to the rest. On your eighth saved day the card grew a row per day, and a shelf of eight days plus the buttons under it ran past the bottom of your phone. It is four days at a time now, with the list itself sliding up and down for the rest, and a line under it that says so: measured live with eight days, "Showing the first 4 of your 8 days. Slide the list itself up for the rest." The headings stay put as the list slides, so a row you have scrolled to is still under the heading that names it. A card with four days or fewer is exactly the card it was before - no scroller, no height, no note.

The height is the part worth being careful about, and it is measured rather than assumed. A long day's name wraps to a second line at a phone's width, so the fourth row is not always as tall as the first three, and a shelf capped by arithmetic would have cut the bottom off the fourth day. Measured live on the eight days: three rows of 59 pixels, then a fourth of 73 because its name took three lines, and the shelf stops at exactly 290 pixels - the fourth row's own foot. The fifth row begins at 290, so nothing of it shows. That is the reading this was built from: capped by arithmetic it would have been 236 pixels and the fourth day would have been cut in half.

5. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before. Your stored data was compared before and after every measurement: 13,888 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans. The eight-day card used for that measurement was put on a

copy of your data and taken off again, with your own file restored byte for byte and nothing left behind.

6. Every rule that stands on it was proved load-bearing, not assumed. Eight faults were put back in, one at a time, and each was watched failing by name, then restored byte-identically. The ladder: with a B-number put back above the titles the test fails reading that panning right puts the day's writing between the module titles; with the title cell taken off its own level it fails reading that the levels can no longer be moved in one place. The reading: with the guard taken out it fails reading that a reading was taken with the pointer over the module titles. The shelf, four ways: capping it off the last day instead of the fourth fails reading where the fourth day ends; writing no height at all fails the same way; measuring it as a multiple of one row instead of reading the fourth row's foot fails reading 40 pixels where the fourth day ends at 220; and taking the slide away fails reading that the days past the fourth are unreachable. The suite is 1217 passing with none failing.

7. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v173 (no database step). Two faults in the "Minutes from one batch to the next" box, both on a day that runs more than one batch: on a module that runs only one, the box did nothing at all and said nothing about why; and on a module whose batch two had been nudged closer by hand, a smaller pace could send that batch in front of batch one, and the day then drew both on the same minute. All inside More -> Scenario planner.

What you reported, in your own words. "similar problem happen. 1st module, time betwenn each batch, set to 21, but no effect, or the chart dont workout" - and then, when asked which of the two you were seeing, "both batch start the same time". You were reading a 12-pan day, and the first module of that day is "Mixing by hand in the tub (set the minutes)": one batch, and the same pace it has always had. Both halves of your report were real, and neither was your phone.

1. On a module that runs one batch, the box does nothing - and now says so. A pace is the minutes between one batch and the next. A module that runs a single batch in the day has no second batch for a pace to sit between, so whatever you type there cannot move anything. That is what your first module is: measured on your 12-pan day, "Mixing by hand in the tub" runs one batch of four pans at a pace of 20 minutes, so a 21 typed into it had nothing to act on and the chart was already telling the truth. The fault was that the box said nothing at all about that - it looked exactly like a box that was broken. It now explains itself in your own words: measured live on that very module, the line under the box reads "This module runs one batch in the day, so there is no second batch for a pace to sit between yet. How many batches it runs in the day is the box above that sets it - make that two and these are the minutes between them." Make it two and the line goes.

2. And on a module that does run more than one, a smaller pace could put batch two in front of batch one. This is the "both batch start the same time" half, and it is the more serious of the two.

A saved day remembers the exact minute each batch starts, one module at a time. A module that carries that list is re-spaced at the pace you type, so a batch you have nudged by hand keeps the minutes you nudged it by rather than being pulled back onto the rhythm. That is right for a batch you nudged LATER - and exactly wrong for one you nudged EARLIER, because its deviation from the rhythm is a negative number, and adding a negative to the smaller rhythm lands it in front of the batch above it.

Measured on your own first module, in the shape that does it. A nudge of batch two by hand with the box still reading 87 leaves the module carrying the times 4:00 am and 4:21 am - and that batch is then 66 minutes earlier than the rhythm, so its deviation is -66. Typing 21 sent it to minute -45: three quarters of an hour before your day begins. The day cannot draw a batch that starts before it does, so both batches were drawn on minute 0 - both at the same time, exactly your words. This is also where the "no effect" reading of the chart comes from on a two-batch module: the box reads 21, and the two bars sit on top of each other and will not come apart.

3. What it does now, and it is your own app's rule rather than a new one. Two rules, and both were already the app's own. A hold is only ever a HOLD - every other press on this screen already obeys that, so a batch can be held back off the rhythm and can never be pulled in front of it. And no batch is left at or behind the batch above it: where a hold has become impossible in the smaller rhythm, the batch takes the rhythm rather than a minute that is not its own. Measured live on that same module, through the same card: typing 21 leaves the two

batches at 4:00 am and 4:21 am, drawn 21 minutes apart with the box reading 21, and the second batch is no longer able to overtake the first.

4. And a third fault, found while proving the second: the Auto box did not save. This one you had not reported and would have met eventually. "Minutes from one batch to the next" carries an empty box with the word Auto in it, meaning each batch starts the moment the one before it ends. Emptying the box moved the flag and nothing else.

The zero that means Auto was being written through the same writer every other box on this screen uses; that writer refuses any number below the box's own minimum; and this box's minimum is 1 - so the refusal came back BEFORE the save and before the repaint. The choice was made in memory, nothing on the screen moved, and the next reload threw it away. Measured on your own mix with two batches: emptying the box left both batches 30 minutes apart and the screen did not move at all; the card reopened with its empty Auto box drawn over the very same unchanged chart; and a reload forgot the choice altogether. The zero and the save are now written directly, so Auto is an answer the app keeps: measured live, emptying the box redrew the two batches 20 minutes apart - the module's own cycle, which is what Auto means - and a reload kept it.

5. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before. Your stored data was compared before and after every measurement: 13,888 characters, byte for byte the same, and your three saved days still reading 24, 12 and 24 pans.

6. Every rule that stands on it was proved load-bearing, not assumed. Six faults were put back in all and each was watched failing by name, then restored byte-identically. The hold rule: with the clamp taken out the test fails reading that a batch pulled closer than its line carried a hold it does not have. The forward sweep: with the sweep removed the test fails reading that a batch was left in front of the batch above it. The one-batch note: with it gated off the test fails reading that a pace box with nothing to space does not say so. The Auto half, three ways: with the 0 written through the writer again the test fails reading that Auto did not write the 0 that means Auto; with the save and the repaint left out it fails reading that the Auto press saved and repainted nothing; and with the list not deleted it fails reading that Auto left a list of times behind, so the cycles no longer decide. One fault in the tests' own stand-in screen was found and fixed on the way - it had no firstChild, so the "did the screen repaint" assertion was comparing nothing with nothing and could never fail - which is the same lesson that file already carries about a stand-in being more forgiving than the real browser. The suite is 1215 passing with none failing.

7. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v172 (no database step). Four things: handing one stretch of work to somebody else no longer moves the people's rows, their names or their jobs; the "Minutes from one batch to the next" box reaches your batches again, and says so when your oven is what is holding them apart; the Your scenarios card is a table you can run your eye down; and the question a card asks before deleting something is drawn over the card that asked it instead of behind it, with the row of buttons at the foot of a card given the shape it never had. All inside More -> Scenario planner.

What you reported, in your own words. First, and the largest thing in this version: "i have 4 persons, 1st Jien, 2nd Wei, 3rd and 4th. I saw Jien is heavy loaded, so i click one of Jien session, and select to switch that session to Wei, where he is free. BUt what happen is Jien disapper, and Jien name chage to Wei, and the originally Wei sessions disappeared". Then: "im now on my 36loaf scenario, pls check why im lowering the minutes frrom one batch to the next, cannot reflex even i reduce it to 13, look like something prevent it from loweriing below 25". And, separately: "can you make the Your Scenario, reorganised it to a better visual clarity, details in table, action button below?" And then, while that was being built: "please check the process of delete a module and the process of deleting a scenario, it dont look right".

1. Handing one stretch of work to somebody else moved everybody on the chart. This was your report, and it was the largest thing in this version.

All three of the things you saw were real, they were one fault, and none of it was your doing.

A day that is sharing them out has nobody written on it. The rows you see are worked out fresh every time the chart is drawn: the first job nobody is on goes to person 1, the next to person 2, and a new row appears only when everybody is busy. That is why a shared-out day reads as perfectly sensible people - and it is also why its row numbers mean nothing at all. They are not yours and they are not on the modules.

Your names, though, are kept by number, and every day you have shares them. So the moment one stretch was given to Wei, a row for Wei existed before the free jobs were handed out - the packing began at Wei's row instead of at the first row, every invented number moved up by one, and the row that had been person 1 became person 2. Jien was person 1, so Jien was gone; the row he had been on now wore Wei's name; and a fifth row appeared at the top with what was left over. Exactly the three things you saw.

2. What that press does now, and what it was measured doing. Before a stretch changes hands, the arrangement already on the screen is written down: each job is given the person the chart is already showing at it - one name for a module whose batches all sit on one row, one name per line for a module worked as lines, and a note on the single stretch where a module's batches really are split between two people. A stretch you had already placed by hand, including one you had handed back to whoever is free, is left exactly as it is. From that moment the numbers on the rows are facts rather than the packing's invention, and your names go on meaning what you see.

Measured live on your own day at a phone's width, going through your exact press. The card read "Move this slot off Jien" and its button read "Move it to Wei". Before the press the chart drew four rows - Jien with five jobs on him, Wei with one, and two more below them. After it, the same four rows were in the same order with the same names: Jien had lost the one stretch you moved and nothing else, Wei had gained that one stretch and nothing else, and the two rows below were untouched. The People box stopped saying the day was sharing them out and read "One to a module", which is what the card had told you it would do before you pressed it. And the chart held there: the page was read again from scratch and drew the same four rows, the same names and the same jobs - where before this fix that very re-read was what moved them.

3. And one sentence on that card was describing a day that was not in front of you. Found by measuring the fix on your own day rather than on a made-up one. The card carried a note explaining that the press settles the whole day, and it said so in the words "This day is sharing them out" - on a day whose People box plainly read "Your own", because one stretch on it had been placed by hand. Two claims, about a day that was making neither. The note may now say only what is true of the day on the screen: the part that says the press settles the whole day is always said, and the part about the box stopping saying it is sharing them out is said only when the box is saying it.

4. The box that would not move, found and fixed. Two different things were true at once, and only one of them was a fault.

The fault: a saved day keeps, for each module, a list of the times its batches start. A list is read as the whole truth about where every batch sits, so a module that carries one ignores the box entirely - the box could be set to 13, or to 87, or to anything, and the batches would not move. Measured on the engine, not reasoned about: a module with the list 0, 25, 50, 75 and the box at 13 still ran 25 minutes apart; set to 87 it still ran 25 minutes apart; with no list at all, 13 ran 13.

How the list got there: the press that moves one batch's time wrote the module's whole list back onto it, and so did the day card's "work the day backwards", which does it for every module in the day at once. So one press of one button used to freeze the pace box of the entire line - which is why it looked like something was preventing it, rather than like a box that had quietly stopped working.

5. What changed. Three things, all small. A press that moves a single batch no longer writes a list onto a module that never had one - batch one of a module IS the module's start time, so writing that one number is the same answer the pace already gives, and leaving no list behind is what keeps the pace alive. A module that already carries a list has it re-spaced to the pace you type, with every batch keeping what you did to it by hand: each batch moves by the minutes the rhythm moved, so a batch you deliberately held off a collision is still held off by the same minutes after. (23 Sep 2026, added at v173: this re-spacing was unbounded in v172, and a

batch that had been nudged EARLIER by hand carried a negative deviation that the smaller pace then added on, putting that batch in front of the one above it - the day drew both on the same minute. See the v173 entry above: a hold is now only ever a hold, and no batch is left at or behind the batch above it.) And the box now carries a live note under it, so a pace the oven cannot reach says so where you typed it instead of doing nothing.

6. The other half of it, which is not a fault and now says so. A module that takes one batch at a time cannot start two of them closer together than one batch lasts - that is your oven, not the app. Measured on your own "Mixing the dough in the tub": its cycle is 20 minutes, so typing 13 still runs the batches 20 minutes apart, and the note under the box says exactly that, names the 20 minutes, and names both ways to run closer: give the module a second production line under "How many production line you have", or switch on "Allow multiple production line" if those minutes are the dough's time and not the machine's. Measured live on your own frozen day: typing 13 turned that module's stored list from 1, 88, 175, 262 into 1, 14, 27, 40, and the day then ran 1, 21, 41, 61 - 20 minutes apart, its own cycle. Typing 40 ran the day 1, 41, 81, 121, and the note went back to being a plain explanation.

7. The Your scenarios card is now a table. One row for each saved day: the day's name, the pans a day it turns out, and how many modules it has, with the pencil at the end of its own row and Save changes and Save a copy below the table rather than mixed into it. The day the planner is currently plotting is filled in and carries a line down its left edge, and every name starts at the same place so the days read as a column.

8. And one measurement that made the table worth having. At a phone's width the card has 315 pixels to spend, and two headings were taking 170 of them - "PANS A DAY" on one line claimed 91 pixels to sit over a two-digit number and "MODULES" claimed 78 to sit over a one-digit one - which left the name 92 pixels. Every name you have is wider than 92, so every row wrapped, each to a different number of lines, and the card read as three ragged paragraphs. The heading is written over two lines now instead, which hands its slack back to the name: measured at 375 by 812 the name goes from 92 pixels to 138, all three rows come back to the same 59-pixel height instead of 59, 72 and 59, and every name starts at the same x. The mark that said which day was open used to be a dot typed in front of its name, which pushed that one name 17 pixels right of the other two; it is a line down the row's left edge now, which costs the row nothing.

9. Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before - a row opens that day, and the pencil beside it opens the rename card without opening the day as well. Your stored data was compared before and after: 13,888 characters, byte for byte the same. The one press in this version that does write to your day is the hand-over above, and it writes only what the chart was already showing - the same people, on the same rows, at the same jobs. Its own round trip was measured on your day, starting from your day as it stood and putting it back byte for byte afterwards: 13,888 characters before it and the same 13,888 after it. Your three saved days are untouched, and each still asks for the pans a day it always asked for - 24, 12 and 24 - and still carries the same 9, 8 and 8 modules. One reading in the table is worth naming, because it is not the same number: "Pans a day" in the table is what each day's line turns out, which is exactly the number the app already puts in a day's own summary line. For two of your days that is the 24 you asked for; for "My sister proposal 21/9/2026" the one tub passes 4 against the 12 you asked it for, and the shelf says 4 for the same reason its summary line already does - so the two screens cannot disagree.

10. The delete button opened a question you could not see. This was the fault you found, and it was not a wording problem or a small one: the question was being drawn underneath the card that asked it.

Both of these deletes are asked from inside a card - Delete this module from the module card, Delete this scenario from the card the pencil opens - and the card stays on the screen behind the question, which is deliberate, so you can see which module you are about to lose. What was wrong is the order the two are drawn in: the layer the question goes on sat at 60 and the layer of a card sat at 65, so the card was painted over the question every time.

Measured on your own day at a phone's width, not reasoned about. With the question open, the topmost thing at the question's own Delete button was an input belonging to the module card underneath, and the topmost thing at the question's own centre was a label from that same card. So the sentence could not be read, neither button

could be pressed, and the only thing you saw was the screen dimming a second time with nothing appearing on it. Tapping anywhere did what the card underneath does instead, and the question sat there unanswered until something else happened to close the card, at which point it surfaced on its own - a delete box appearing when you had stopped asking for one. Both flows were the same fault, and both are measured fixed: at the question's Delete button the topmost thing is now the Delete button, at Cancel it is Cancel, and at the centre of the box it is the box.

The question's layer now sits at 72, above every card and above a toast, and only the app-password lock outranks it. Nothing about the flow changed otherwise: Cancel leaves the module and leaves the card open where it was, and Delete takes the module out and closes the card with it - both driven and watched on a live day, a module going from eight on the line to seven with the day answering without it, and a saved scenario going from three to two with the shelf redrawn and no name left standing on the card. And the question now names the module in quotes, as every other delete in the app names what it is about. Unquoted, a module called "The proofer again" read as one sentence with the name lost inside it: "Delete The proofer again?".

11. And the row of buttons at the foot of a card had no shape at all. The same report's other half, and a smaller thing, but the same kind of thing: a class used in 26 places across 12 files that never had a rule of its own.

Without one, its buttons flowed as plain inline blocks - flush against each other and hard against the left edge of the card. Measured on the module card at 375 by 812: "Delete this module" and "Done" were 0 pixels apart, so a thumb reaching for Done was touching the button that deletes the module, with 88 pixels of dead card to their right. The scenario card was the same shape the other way up - its Delete stacked directly above its Save at exactly the same left edge - and both read as two buttons run together rather than as a row of choices.

The row is now one rule, written once: the card's own edge, 8 pixels between the buttons, and a line's worth of air above them. wrap is part of it rather than a detail, because the module's three buttons need 437 pixels and have 303 - a row too wide for its card has to go onto a second line, not off the card's edge. It is the shape the label printer's own sheet already carried in full and the shape three other cards had each patched in by hand, a little differently every time; those three patches are gone, because one answer in one place is what stops a fourth being added. Measured after, on the same card: Duplicate this module on a line of its own ending at the card's edge, then Delete this module from 94 to 264 and Done from 272 to 345 - 8 pixels apart, Done sitting where a thumb expects it at the card's own right edge.

12. Every rule that stands on it was proved load-bearing, not assumed. Nineteen faults were put back in all, and each was watched failing by its own name before being restored byte-identically.

For the hand-over: with the arrangement left unwritten the test fails reading that the day has stopped being able to say what it is, which is the fault put back; with a person written onto a module without the crew that carries it, the same test fails with your own three symptoms - five rows where there were four, renumbered 2, 3, 4 and 5; with the People box allowed to read a crew of zeroes as an arrangement, the hand-back test fails reading that a day which has not moved no longer says it is sharing them out; and with the card's note gated off, it fails reading that the card does not say what the press settles.

For the pace: with the re-spacing taken out the test fails reading "the stored list did not follow the pace she typed"; with the note about the oven taken out, "a pace the machine cannot reach is not explained"; with every batch pulled back onto the rhythm instead of keeping its own offset, "a batch held off by hand was pulled back onto the rhythm"; and with the press left writing a list again, "Mixing the dough in the tub came out of the press carrying a list of times, so its pace box is dead".

For the question: with the layer put back under the card it fails reading "under the card that asked it"; with it put above the app-password lock, "drawn over the app-password lock"; with the module's name left unquoted, "does not name the module in quotes"; with the destructive button put first, "answered by the safe button and then the destructive one"; with Cancel made to run the delete, "Cancel deleted the module"; and with the card left open behind the delete, "the card stayed open after the module went".

For the foot row: with the rule declared a second time it fails reading "declared 2 times"; with the buttons put

back to inline blocks, "flow as inline blocks"; with the 8-pixel gap taken away, "sits flush against"; with the row back at the left edge, "not laid out against the card's own edge"; with wrap taken away, "cannot wrap"; with one card patching its own row by hand again, "lays its own foot row out by hand"; and with the foot's buttons reordered, "does not carry its three buttons". The two faults found while measuring the note on your own day are put back the same way: with the note always claiming the day is sharing them out it fails reading that the card claims a box that is not saying it, and with the note taken off the card altogether it fails reading that the card does not say what the press settles. The test suite is 1215 passing with none failing.

13. No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v171 (no database step). Three things you asked for while looking at the day chart: a right press and hold moves a window's day about without opening anything, the clock balloon rides two inches clear of your pointer, and the batch card is brief. All inside More -> Scenario planner.

What you asked, in your own words. First: "And can i drag the module window up/down, left/right by right click and hold? Dont let this action open up the card". Then, while that was being built: "put the clock balon 2 inches higher than. cursor". And: "the batch pop up, make it as brief as possible".

1. A right press and hold now takes hold of a window's day and moves it, both ways at once. Press the right button anywhere on either chart window, keep it down and move: the day comes with your hand, sideways and up and down together, exactly as if you had taken hold of the paper and pulled it. The window under your hand moves and the other one comes with it sideways, so the two can never disagree about a minute - that part was already there and is untouched. Measured live at a phone's width (375 × 812): a right press on the modules' window and a 60-pixel drag moved its day from 0 to 60, and the people's window to the same 60. The window marks itself while your hand is down and unmarks the moment you let go.

And it opens nothing, which was the other half of your sentence. The gesture is the RIGHT button, and a browser fires no ordinary click at all for a right press, so a press for panning can never be read as a press for opening. The ruler's own press, which drags the clock reading, now ignores the right button too - one press, one gesture. Proven live, not assumed: a right press on a bar dragged the day and left the card shut, and a left click on that very same bar opened it. Your own hand is on the safe side of it either way - a finger cannot set the right button at all, so nothing here is reachable from your phone and the way you touch the chart is exactly as it was. The browser's own menu, which a right press would otherwise put on the screen in the middle of the gesture, is stopped over the chart.

2. The clock balloon rides two inches clear of your pointer. It used to sit 8 pixels above your finger, which on a phone is under it. It is lifted two CSS inches now - 192 pixels, and named in the unit you used so the number is one this can be held to rather than a feeling. Above your pointer wherever the window has the room, and where it does not, the same two inches BELOW it - which is the side that matters, because the clock strip you press with a finger is at the top of the modules' window, and a balloon that came to rest at the window's top would be lying on your hand. Measured live at 375 × 812: the full 192 pixels of clearance in both directions, and never nearer than 175 pixels anywhere on the chart. It is still held inside the window you are looking through rather than to the day, so a tall day cannot carry it off the screen.

3. The batch card is brief, and its own reading is no longer cut off. What left the card: the paragraph under each pair explaining what a move means in that module's own terms, and the second line that repeated the two times the line above it already reads out. What is on it now: one line of four readings - which batch, how it sits against the line, the two times, and what it costs your hands - then the two pairs, five minutes and one minute, and the way back where there is one. Measured at a phone's width, the card came down from 338-366 pixels tall to 265-288. The explanations are still true and are still there, on the press itself as its accessible name, and every press still answers in words when it lands.

And one fault found while measuring that card, fixed here. Four readings on one line ran off the card's right edge - "Batch 2 on the line 8:51 am -> 8:5" - with the end time, the clock you opened the card for, cut off. A reading is set never to be squeezed, on purpose, so a time can never be shrunk to nothing; a line one word too wide therefore overflows rather than shrinks. The card's readings now WRAP to a second line instead, which

only ever happens where the line would otherwise be clipped. The same fault was possible on every card that shares that line, and it is fixed on all of them by the one change.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before. Your stored data was compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

Three rules now stand on it, and every one was proved rather than assumed. One says a right press and hold drags a window both ways and that a right press over a bar opens nothing, while a left press on the same bar still does. One measures the balloon against the palest bar and pins the two inches of clearance, above or below, with the reason the window's own top is not good enough written into the test. One pins that a card's readings wrap and that they still refuse to be squeezed - the two settings asserted together, so neither can be changed alone. Each was proved by putting the fault back and watching it fail by name, then restored byte for byte: no downward fallback for the balloon read "a balloon with no room above her pointer sits -50px below it, not the two inches she asked for"; no upward preference read "the balloon sits 232px above her pointer, not the two inches she asked for"; and removing the wrap read "a card's readings are back on one line, so a row too wide for the card is clipped at its edge again". The test suite is 1209 passing with none failing.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

CHANGELOG and the guide both carry all of this. Both PDFs rebuilt and read back with PyMuPDF, the app's own More screen reads Engine v171, and your phone is untouched by nothing here.

23 Sep 2026 - engine v170 (no database step). The frame around every bar is drawn quietly now, so a block of work still reads as a block without every bar being a box drawn round it. Nothing else moved. All inside More -> Scenario planner.

What you reported, in your own words. "the border look like over emphasized". You were right, and it was my own v169 line being too loud. Nothing about the day, the bars or your data was wrong - the frame was simply drawn heavier than it needed to be.

What was too much. The frame went in at a line one and a half pixels wide, drawn at full strength in the bar's own solid colour. Round a pale bar that is not an edge, it is a box drawn on: every bar on the chart was outlined that way, so the day read as a page of boxes rather than as a day with edges on it.

What changed, and it is only the frame's own two numbers. The line is one pixel wide rather than one and a half, and the whole frame is laid down at 55 parts in 100 rather than at full strength. The colour is untouched and is still the bar's own tone - the colour your hands already wear inside that bar, and, for a person, the colour of their card's stripe - so the frame still says at a glance which module a bar belongs to. Measured live on your own day at a phone's width: the frame reads the bar's own colour, a 1-pixel line with the layer at 0.55, on a batch bar, on a lane bar that shares its row, and on a person's time slot alike.

Your B-numbers and your people's names are untouched, and no line crosses a label, exactly as at v169: the frame is still drawn below them.

One number softens every frame together. The strength is declared once, on the chart's own wrapper beside the marker's transparency, so the modules' window and the people's window and all sixteen tones are all softened by the one edit and cannot be tuned apart. Lower it towards 0 for a whisper, raise it towards 1 for v169's line back again - it is one number either way, so say the word if this is still not the weight you want.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before, and the frame still takes no taps at all. Your stored numbers were compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

And the rule that stands on it was extended rather than left standing on the old numbers. The frame test now carries your report in its own name, pins the 1-pixel line, pins the layer's strength and refuses a strength of 0 or 1 - so a later edit cannot quietly put the shout back - and refuses a second declaration of it. Each was proved

by putting the fault back and watching it fail by name: the 1.5-pixel line back made it read "the frame is no longer a 1px line"; full strength back made it read "--frame is 1: one at 1 is v169's shout back again"; and the strength declared twice made it read "the frame's strength is declared in more than one place, so the two windows can disagree". Each was restored exactly as it was afterwards. The tests are 1,201 passing with none failing.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v169 (no database step). Every batch bar and every person's time slot now wears a frame in its own colour, so a block of work reads as a block and not just as a patch of colour on the paper. And the clock's lines have gone to the back of the day while every bar became see-through, so the ruler reaches inside them the way you asked instead of sitting on top of them. All inside More -> Scenario planner.

What you asked, in your own words. First: "Can the batch and person time slot having a more highligted frame?" Then, while that was being built: "its scale should be at backgrond, the marker should be ontop, with some degree of transparency". Those two together are this release, and the second one turns round the way the ruler was drawn at v165 and v166.

Every marker now wears a frame, and it is drawn in its own colour. A batch bar and a person's time slot are the very same shape in the app - one rule frames both, so a batch and a slot cannot end up looking like two different things. The frame is the bar's own solid colour: the same colour the hands inside that bar already wear, and, for a person, the same colour as the stripe down the side of their card. That matters more than it sounds - a frame in a colour of its own would be one more thing to keep in step, and it would drift the first time a tone changed. Nothing about the bar's shape moved for this: a batch bar is still 16 pixels tall and a person's slot still 11, measured, and the frame is a line laid inside the bar's edge rather than a border added round it. A border was tried and thrown away - it would have squeezed the inside of the bar, which is where your hands and the cycle shades are placed, and an inner shadow was no better because those same shades are painted over it.

The ruler is at the back and the markers are over it, as you asked. Since v165 the lines were drawn as a layer of their own on top of everything. They are now the paper's own ruling again - the surface the bars are drawn on - and every bar is a wash over it rather than a solid block. The bar is 72 parts in 100 opaque, which is the "some degree of transparency" you asked for: enough that the clock reads through it, not so much that the day looks faded.

What that cost, honestly, and what was done about it. A line seen through a bar loses about a quarter of its strength, because the same wash covers both the line and the paper beside it. Left alone, the ruler would have gone faint again - and faint is the fault you reported twice. So both lines were drawn heavier: the hour line at 60 parts in 100 of full strength where it was 34, and the finer grid at 45 where it was 20. Measured through one of your batch bars, the hour line now differs from the bar beside it by 29 parts in 255 and the grid by 22, against the 20 that your own v166 rule demands. The honest other side of that: on the bare paper between the bars the ruling now reads about twice as strong as it did at v168. That is the price of drawing it behind the markers, and it was paid deliberately.

And no line crosses a label. Your B-numbers and the name of the person on a stretch of their day are laid above the frame, so thickening the frame cannot put a line through a number you have to read.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Every tap still opens exactly what it opened before - a batch its Batch card, a cycle its own card, a person's stretch the hand-over card; the frame takes no taps at all. Your stored numbers were compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

Three rules now stand on it, and every one was proved rather than assumed. One says the ruling is the paper's own background and that no second copy is drawn over the markers - two grids, one in the gaps and one across a bar, is how a reader ends up seeing two different grids. It also pins the transparency to one number declared once, so the modules window and the people's window can never be tuned apart. One says every batch bar and every person's slot wears a frame in the tone its own hands or card already wear, and that no bar

grew a border or a shadow of its own. One is the arithmetic above: measured, not trusted, that both lines still differ from a marker by at least 20 parts in 255. Each was proved by putting the fault back and watching it fail by name, then restored exactly as it was. Five faults were put back in all, and one of them is the point of the whole release: with the grid's ink back at v166's strength, the test fails reading "differs from the marker beside it by only 9.8 parts in 255, which is invisible on a phone - the fault she reported twice at v165 and v166". The tests are 1,201 passing with none failing.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

Softened at v170: the frame described above went in at a line one and a half pixels wide and at full strength, and you reported it as "the border look like over emphasized". It is now a one-pixel line laid down at 55 parts in 100, in the same colour as before. Read the v170 entry above for the rest.

23 Sep 2026 - engine v168 (no database step). The extra slider between your two chart windows is gone, so each window's own sideways bar is the pan. The modules window is taller. The "People at once" row is out of the people's window, and an overlap is still shown in red. All inside More -> Scenario planner.

What you asked, in your own words. First: "since the both windows have their own slider, additional slider is redundent. Remove that" - so the slider that sat between the two windows is removed. What you asked for when the two windows were built still stands and is untouched: the two windows stay in step, so any minute is in the same place on both, whichever one you drag. Then: "can you increase the modulle window further?" - the modules window is taller. And: "im thinking of remove the people at once?", settled by "you can just show when there is overlapping highligt in red box, like previously have" - the row is out, and the overlap is still shown in red.

The extra slider is gone, and the two windows still move as one. Measured on your day at a phone's width, 375 by 812: nothing sits between the two windows any more but their own two edges - the chart is the modules' window and the people's window, in that order, with nothing in between. Dragging one window's day sideways moved the other to the same pixel, and the window under your finger took no write back at all, so a drag is never fought by the window that is following it. On the two windows being able to reach the same place: they could not quite, and it was one pixel. The ruler's last hour sits at the very end of the day, and its line was being drawn one pixel past it, which made the modules window 2,461 pixels wide against the people's 2,460. Both windows now read 2,460 and both pan to 2,139. That last hour's name is also written to the left of its own line now, where a ruler with an end on it has always put it, and it reads in full inside the day - measured, all 25 of the ruler's names sit inside the day, the last of them ending 3 pixels short of its end.

The modules window is taller, and here is exactly what that buys. It takes the larger half of the chart's ceiling now: 64 of the 84 parts where it had 50 of the 80, and 590 pixels where it had 470. On your phone at 812 tall this changes nothing you will see - your day's modules already fitted the window exactly (403 pixels of modules in a window that was already wide enough for them), so the taller window is room there rather than a difference on your screen. Where it shows is a shorter screen: at 667 tall the old window held 333 pixels and made your day scroll, while the new one holds 427, so your whole day sits in front of you with nothing to scroll. It is also the room the next two module rows will need as you add them.

The "People at once" row is out of the people's window. You were right that the day already says it without that row. An overlap is still shown in red, in the places that were already saying it: the red outline on the colliding bar in the people's window, the colliding stretch of that person's own row, the red line in their card, and the collisions written out in words under the diagram - "two jobs at once", or how many. Nothing about an overlap is hidden by the row going, and it is one line to bring back when there is a second pair of hands to weigh up.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. No bar changed colour, and every tap still opens exactly what it opened before - a batch its Batch card, a cycle its own card, a person's stretch the hand-over card. Your stored data was compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

Four rules now stand on it, each proved load-bearing rather than assumed. The ones from v167 still hold,

with the slider's part of them narrowed to what is left: the two windows are adjacent with nothing between them, and a drag on either moves the other while the window under the finger takes no write back. Three are new. One asserts that the "People at once" row leaves nothing on the chart while a collision is still marked in red. One asserts the ruler's last name is drawn inside the day, and that the ruler's own track clips, so the two windows cannot disagree by even a single pixel. One asserts the modules window holds the larger half of the chart's ceiling, 64 of the 84. Every one was proved by putting the fault back and watching it fail by name, then restored byte-identically.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v167 (no database step). The day chart is now two windows sharing one horizontal slider - the process above, the people below, one ruler's lines down both - and the blank space under your people's rows is gone. This release also carries two smaller faults you reported while it was being built.

What you asked for, in your own words. First: "the window for person stay over size", which you confirmed means the empty space under the people's rows. Measured at a phone's width, 375 by 812: the chart was one window, fixed at 650 pixels tall, the day's own content was 473, and your people's rows ended at 457 - so 177 pixels of empty window sat under them. Then your proposal, which was better than a tweak: "instead of consider then one window, why not create 2 windows, and let the ruler sync in the 2 windows. As the business grows, the persons will grows, processs might not, we need to have a better way to manage, what you say?" And on how the two should sit: "can the 2 window share the horizontal slider, place between the 2 windows, make the 2 windows as close as possible", with the slider going straight into the people's bars underneath it.

What is on the screen now. Two windows, one above the other, with a single slider lying between them and touching both - measured gap of 0 pixels either side. The upper window holds the clock and every module, batch and cycle; it scrolls up and down on its own when you add more modules. The lower window holds your people's rows and nothing else, and it scrolls on its own too, so as your people grow the second window grows and the first one does not have to. The ruler's lines reach down both. The slider between them is the only horizontal control: drag it, and both windows move together; drag either window's own day sideways, and the other follows and the slider follows with it. Measured on your day at 375 by 812: the slider reads a travel of 2,139 pixels, a drag on one window moved the other to the same pixel, and the window under your finger received no write back at all - the echo is dropped rather than sent, which is what would make a phone stutter mid-drag. On a day short enough to fit, the slider dims and cannot be dragged, because a control you can grab that does nothing reads as a fault.

The blank space. It is gone. The chart is no longer forced to a fixed 650 pixels; it is now exactly what its two windows need. Measured on your day: upper window 403, slider 20, lower window 70, totalling 493, with your people's last row sitting exactly on the window's floor - 0 pixels of empty window under it, where there were 177. The height ceiling from v157 survives as a ceiling: the chart is never taller than it was, it is only allowed to be shorter when your day is shorter.

The first of the two smaller faults: a deleted saved day left its name on your card. In your words: "why the deleted scenario stil listed?" You were right, and it was reproduced rather than explained. Deleting "My sister proposal 21/9/2026" left the card still offering "+ My sister proposal 21/9/2026 · Mix by hand, no chiller", dressed exactly like a day still saved, while your storage held only two. The cause was that the card offered each ready-made day by name whenever that name was not in your list, so deleting one put the name straight back. Now the card offers one plain row, "+ Add a ready-made day", which names nothing; opening it lists whichever ready-made days you do not already have, and one tap puts that day on your own list, where it is yours to change like any other. Proven live: after deleting that day the card named it nowhere, and the chooser handed back exactly that one day.

The second: on a computer, a person's name card was cut. Hovering a person's name on the chart opens a card of what they do that day. In the new lower window that card was taller than the window it opened into - measured 180 pixels of card inside a 69-pixel window, 96 pixels of it below the window's own bottom, with no scroll position that could show it whole. It is now laid on the screen rather than inside the window, so it reads in

full wherever the name is, and it is pulled back from the screen's edges rather than running off them. The "People at once" row's card, which sits in the same window, is placed by the very same function, so the two rows cannot behave differently. Measured on a computer at 1280 by 900: both cards fully on the screen, with 8 pixels to spare at the foot.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. No bar changed colour, and every tap still opens exactly what it opened before - a batch its Batch card, a cycle its own card, a person's stretch the hand-over card. Your stored data was compared before and after: 13,561 characters, byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

Three rules now stand on it, each proved load-bearing rather than assumed. One asserts the two windows are adjacent with the slider between them, and that a drag on either moves the other and the slider while the window under the finger takes no write back. One asserts that a ready-made day you have deleted leaves no name standing on your card, with the offer row itself excluded from the check. One asserts that both a person's card and the tally row's card are placed beside their name and pulled back from the screen's edges. Every one was proved by putting the fault back and watching it fail by name, then restored byte-identically.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

Added at v168: the single slider shared by the two windows, described above, has since been removed - each window now carries only its own sideways bar, and the two still move as one. Read the v168 note at the top of this file for what is on the screen now.

23 Sep 2026 - engine v166 (no database step). The ruler's lines are now drawn strongly enough to be seen on a phone - including across your batches, your cycles and your people's occupied slots.

What you reported. In your words: "At ver165, you still fail to reveal the ruler at person's time slot". You were right, and this release is that same fault finished rather than a new idea.

What v165 got right, and what it got wrong. v165 did draw the ruler over the bars - that part worked, and it is untouched. What it got wrong was the strength of the ink. The grid line was drawn at about one part in six of full strength, and over one of your slots it came out as a colour that differs from the slot by about ten parts in 255 - a hairline the eye cannot find on a phone in normal light. So the ruler was drawn, and it was invisible, which on a screen is the same as not being there. That is why you came back with the same words a second time.

What changed, and it is only the ink. Both lines are drawn in a strength measured against the bar they cross rather than against the cream paper between the bars - the hour line at about a third of full strength and the finer grid at about a fifth. The lines themselves are in the same places as before, at the same intervals for the scale you are reading at, and the ruler at the top of the chart still draws none of them over its own ticks. Nothing moved; the day simply reads as one ruled sheet now, so the edge of a marker has a line beside it wherever you look.

And a line never crosses a label. Your B-numbers on the bars, and the name of the person on a stretch of their day, are now laid over the ruling rather than under it. A gridline drawn through a B2 is a number you have to guess at, and the whole point of the ruling is to make the day easier to read, not harder.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. No bar changed colour, and every tap is untouched - the ruling takes no taps at all. Your stored numbers were compared before and after and are byte for byte the same, and your three saved days still read 24, 12 and 24 pans.

About the height of the planner window, which you also raised. That is not in this release. You are right that it is a real thing, and what you proposed - two windows, one for the process and one for the people, sharing one ruler - is the honest fix rather than another tweak. It is written up as its own piece of work, so it is decided on its own rather than half-done here.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

Changed at v169: the ruler's lines are no longer drawn over the bars, as described above. The ruling is the

paper's own background again and every bar is a see-through wash over it, so the clock reads through a marker rather than across it. The goal is the same one named here - the clock must hold inside a slot you are working - reached the other way round. Read the v169 note at the top of this file.

23 Sep 2026 - engine v165 (no database step). The ruler's lines are now drawn over your batches, your cycles and your people's occupied slots. Before this the lines stopped dead at the edge of every bar, so inside a slot you were working there was no ruler at all.

What you asked. In your words: "i requested a ruler to draw into person's line time slot, maybe you cant visual my complaint". Asked how the lines should reach that slot, you said: "Just like the ruler draw over the batches and cycles". You were right that I had not pictured it, so I went and measured instead of answering.

What was actually happening, and it was worse than "the lines are only on some rows". Your person's row did carry the ruler's lines - it has since v160. But every marker, batch and cycle is a solid coloured bar, and a solid bar sits ON TOP of the lines behind it. So the ruler vanished wherever a bar was. Measured on your own day: of one module row's ruler lines, 5 fell inside your batches; of the slots on one person's row, 14 fell inside them. Every one of the 14 was covered. Pointing at a ruler line inside one of your slots, the topmost thing on that pixel was the slot itself. That is why the ruler looked present on an empty stretch of the day and absent wherever you were actually working.

And it was the same on the module rows. So the lines were not reaching over your batches and cycles either - what you were seeing was them showing in the gaps between bars, with your eye joining them up. This release makes that true for real rather than only in the gaps.

What changed, and it is one thing. The ruler is now a layer of its own, drawn OVER the bars - over the batches, over the cycles and over your people's occupied slots alike. The lines themselves are the same lines as before: the same hour line, the same step for the scale you are on, at the same positions. Nothing about the ruling moved; it is only no longer hidden. The whole day now reads as one ruled sheet, and a slot is ruled exactly like the empty day beside it.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. No bar changed colour, and every tap on a bar is untouched - a batch, a cycle and a person's own stretch are still tapped and still open the same cards. Your three saved days were read before and after: One baker day, No fridge 1 person and My sister proposal 21/9/2026 each still on the target they had, stored numbers untouched.

Two rules were proved load-bearing rather than assumed. One test now stands on that there is exactly one ruling - if the lines were ever drawn both under the bars and over them, the translucent grid would read darker in the gaps than across a bar. Another stands on that the ruling is drawn at a height above the bars, that it lets every tap through, and that the ruler's own strip at the top turns it off so it cannot put lines between the ticks you read the clock off. Each was proved by putting the fault back: removing the ruling's height made the first fail, and putting a second ruling back on the track made the other fail.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

Changed at v169: the lines described above are no longer drawn over the bars as a layer of their own. The ruling is the paper's own background again and every bar is a see-through wash over it, so the clock reads through a marker rather than across it. What has not changed is why: the lines must reach into a slot you are working, and the ruler's own strip at the top still draws none of them over its ticks. Read the v169 note at the top of this file.

23 Sep 2026 - engine v164 (no database step). A number you type into the cycles of a module now stays where you put it. Before this, typing into a second box put the first box back to what it was when the card opened, while the first box still showed the number you had typed.

What you asked. In your words: "why some parameter that i save can change by itself", and then, when I had guessed wrong about where: "it was on stretch and fold module, the cycle1,2 , loading, unloading. It is not like you explain". You were right and my first answer was not. I stopped explaining and went and did it on your own fold module until I could see the number move on its own.

What was really happening, and it is exactly where you said. Inside a module's card there is a list of cycles, and each cycle has its own boxes: the name, Minutes, Load, Unload. The card does not redraw itself while you type, and that is on purpose - it was your own instruction at v142, because a redraw would throw the box out from under your finger. But the code was keeping a copy of the whole cycle list from the moment the card opened, and every box wrote that old copy back. So the first box you used was still carrying the numbers as they were when the card opened, and the moment you touched a second box, the first box was put back to those old numbers. The screen kept showing what you had typed, the day behind it kept the old value, and the two disagreed. That is what "changing by itself" was.

Measured on your own day, so it is not a theory. Your fold module runs four cycles. Cycle 1's Minutes was 31; typing 40 into it stuck. Then touching cycle 2's Unload put cycle 1's Minutes straight back to 31, while the box under your finger still read 40. The same thing happened on one row: Load 5 stuck, then typing 44 into the Minutes beside it put Load back to 0 while its box still read 5. Renaming a cycle stuck only because nothing else was waiting to be undone.

The fix, and it is small. A box now changes only its own value, on the list as it stands at the moment you touch it, instead of writing back the copy from when the card opened. Nothing else in the card can be dragged back with it. The comment beside the code names your words and the exact numbers above, so the next person to touch it cannot put the copy back.

Nothing of yours is rewritten. Not one module, batch, start time, cycle or saved scenario changed, and nothing here blocks a sale or reads an order. Your three saved days were read before and after and compared byte for byte: One baker day, No fridge 1 person and My sister proposal 21/9/2026 each still make what they made, their stored numbers untouched.

A test now stands on it. It opens your fold module's card, types into cycle 1's Minutes and then cycle 2's Unload, and asserts the first number survives the second box - the exact fault. It also covers two boxes on the same row, and a cycle renamed after a number was typed. Proved load-bearing: with the fix taken back out it fails, reading "the second box she touched put cycle 1's Minutes back to where it was when the card opened".

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

23 Sep 2026 - engine v163 (no database step). The clock above your people's rows has been taken back out. The clock is drawn once again, at the top of the chart where it has always been, and the ruler's lines still run down through your people's rows exactly as they did.

What you asked. In your words: "i dont want the clock, i just want the ruler to draw into person row". That was said after v162 had put a second clock above your people's rows, so this release takes that second clock back out. Nothing else about the chart moves.

The ruler's lines were never the thing missing, and they are still there. The lines that run down your people's rows came in v160 and they are untouched: every row of the day carries the faint grid, so a marker's edge still lands on a line you can follow up the chart to the top. That is the part you asked for, and it is the part that stays.

What comes back. The chart window is about 35 pixels taller again, because the block pinned at the foot of the chart no longer carries a clock row inside it. On your own day at a phone's width the foot block goes back from 105 pixels to 70.

One rule is kept, and it is worth naming. There is now a test whose whole job is that the clock is drawn once and that no clock is ever drawn above the people's rows again. It was written because the second clock came from a shared builder, so putting it back is one line - and a thing that can come back by accident should have something standing on it.

What has not moved. Not one module, batch, start time, cycle or saved scenario of yours changed, and nothing here blocks a sale or reads an order. Your three saved days are the same three days, each on the target it had, their stored numbers byte for byte the same.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

The operations guide (section 23) carries the same change, and the tests are 1188 passing with none failing.

23 Sep 2026 - engine v162 (no database step). The times along the clock are written a second time, standing directly above your people's rows, so a marker at the foot of the day can be read without tracing it 368 pixels up to the top.

What you asked. In your words: "why cant you draw the ruler into person's lines?" I went and measured it rather than answering from memory, because the answer turned out to be "half of it already is". The ruler's *lines* have reached your people's rows since v160 - every row of the day carries the faint grid, so a marker's edge already lands on a line you can follow up the chart. What had not reached them is the *clock*: the times written along the ruler, which lived only at the top of the panel. On your own day the top of the chart stood 368 pixels above your people's rows - the whole height of your eight module rows - so reading the clock and reading a marker meant looking in two places far apart. Asked which of the two ways to close that gap you wanted, you chose this one: "yes, draw the clock above the people's rows".

The clock is now drawn twice. Once at the top of the panel where it has always been, and once more standing directly above the people's rows, inside the block that is already pinned to the foot of the chart - so it rides the same pin as the rows it was asked for, and the day scrolling can never leave it behind. The foot one sits on the first person's row with no gap at all between them, measured on your own day.

The two clocks cannot disagree, because they are the same clock. They are built by one and the same function, from the same minute, with the same step and the same labels - so the top of the chart and the foot of it can never say two different things about where a minute is. Measured on your day at your usual reading: 97 ticks each, every label the same, and the 8:00 am tick at exactly the same pixel on both, whether the day is scrolled or not.

And the foot clock is read exactly as the top one is. Point at it, or drag your finger along it, and the same hairline follows with the same reading in the same words and to the same minute, snapped the same way. It wears a crosshair under a computer's pointer, as the top one does. Two clocks that look alike now behave alike; there is no second way to work the foot one.

What it cost, measured rather than promised. About 35 pixels of the chart window. On your own day at a phone's width the block at the foot grew from 70 pixels to 105, and your day still fits inside the window with nothing to scroll - so those 35 pixels came out of empty space under your bars, not out of your day. At a wide screen the window is 720 tall holding 705 of content, and the day still fits.

One fault was found by testing on the real app, not by a test. The code that finds a clock's own track asked the wrong kind of list for the answer, and a real browser's list does not have that question to ask - so the screen failed to open at all with "This screen couldn't load", while every test stayed green, because the test's stand-in list is more forgiving than a real one. That is the standing lesson about stand-ins, and it is written into the code beside the line so the next person to touch it cannot repeat it. Fixed, and the tests now cover the wiring that hid it.

What has not moved. Not one module, batch, start time, cycle or saved scenario of yours changed, and nothing here blocks a sale or reads an order. Your three saved days were compared before and after: One baker day makes 24 pans still, No fridge, 1 person makes 24 pans still, and My sister proposal 21/9/2026 makes 4 pans still, each on the target it had, their stored numbers byte for byte the same.

What is not in it, so you are not left looking for it. The other way to reach the markers - moving the clock times around rather than drawing a second one - stays available if you ever want it instead. And the per-batch person, letting one batch go to somebody else while its neighbours stay, is still a release of its own because it needs a stored field and a one-time step in Supabase.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

The operations guide (section 23) carries the same change, and the tests are 1188 passing with none failing.

23 Sep 2026 - engine v161 (no database step). Your people's rows are the app's own white now, and a tap on one of their markers moves THAT STRETCH and nothing else - the batch it belongs to is

untouched.

What you asked, and it was two things. In your words: "I want to have white background for person's line", and "The reassign job to next person is not whole day, it is that slot only". Asked what one tap should take with it, you told me: "we dont change the batch. Say a labour slot belongs to person1, clicking that slot, will offer to swap it to others, this basically to balance work load".

Your people's rows take the app's own white. They wore the cream band that every soft strip in the app wears, which reads as a highlight sitting under the markers rather than as the plain paper the day is read against. They are now the same white as the panel behind them and as every module's name cell, so the foot of the chart is the same paper as the rest of it. What still marks the block as the foot of the day is untouched: its own line above it, its own shadow, and the hairline between one person's row and the next.

A tap now moves one stretch, not the whole job. Until this release, tapping a marker handed over the whole module on that line - every batch of it. Your fold loop folds three times in one batch, so a tap anywhere near it took three of your markers with it. That is not what you asked for and it is not what balancing a day needs. What changes hands now is the single stretch your finger is on: the one batch, and which part of that batch. Its neighbours stay exactly where they were, and the batch itself is never touched.

The card says so before you press it, with the number. If the batch you tapped has more than one stretch on that person's row, the card counts them and says only the one you tapped moves. You are told what is about to happen rather than finding out from the chart afterwards.

And what you can hand a stretch to. The list is the other people who have a row on this day, named as you have named them - plus "whoever is free". That last one is not a person; it is what the day does when nobody is named, and without it a stretch on a module nobody was put on could be handed to somebody and could never be handed back. Pick it and the press under the list changes to say so, before you press it.

Your module's own card keeps up. A hand-over is now one stretch of a module rather than the whole of it, so "Who is at this module" is the module's default and no longer the whole answer. The card now carries a note naming what has been handed on, and to whom, with the count - so a card reading Person 2 while a marker of that module sits on Person 1's row cannot happen. On a day where you have handed nothing over, there is no note at all.

One thing the earlier note got wrong, said plainly. v160's entry said that moving the person on one batch of a module would need a new stored field on the module and a one-time step in Supabase. The stored field was indeed needed - it is there now - and the Supabase step was not. Your planner travels inside the app's own settings record, so nothing had to be run.

Nothing of yours is rewritten. The only thing that changes is the one stretch you hand over with your own press. Your day was read before and after: One baker day still makes 24 pans, and its one person's row still carries the same 32 stretches across the same clock. No module, no batch, no start time and no cycle moves by being read, and there is nothing to run in Supabase.

Measured on the real app rather than assumed. A stretch was tapped on your own day and handed over: the toast named the module and the clock, the marker left one person's row and arrived on the other, and the module's other three batches stayed where they were. Tapping in a gap still opens that person's own card, exactly as it always did.

Measured rather than promised. The suite is 1187 passing with none failing. The two tests that stood on the module-level move are now three that stand on the single stretch - that only the one key is written and the module's own person, its crew and its times are untouched; that a second hand-over keeps the first; that a stretch handed back to whoever is free survives as its own answer; and that the chart puts exactly the tapped stretch on the other row. The app's own More screen reads Engine v161.

23 Sep 2026 - engine v160 (no database step). In the Scenario planner, the ruler's lines now carry on down through the people's markers, and a tap on a stretch of somebody's day hands that job to another of your people.

What you asked, and it was two things. In your words: "can the ruler extend down to peoples marker area?", and "can the personX marker be click to change it job to personY, by a drop down person selector". Asked which way you wanted the ruler to reach the markers, you chose **a faint grid down every row** - the clock itself stays where it has always been, at the top. Asked what the tap should do, you told me: "click on the person's occupied time slot, a drop down list, list the other people available".

The ruler's lines reach the markers. Every row of the day has always drawn a faint line on the hour, so a marker could be traced up to an hour and no closer. That grid now carries on down through the module rows and through the people's rows, so the edge of a marker always lands on a line you can follow up the chart to the clock. On your day the lines are 36 pixels apart at the widest reading, 24 at the standard one, and 12 and 16 at the two close ones - so the grid you read a marker against is the ruler's own step everywhere the ruler can be read at all.

Why the grid has a step of its own. At the closest reading a single minute is 3.2 pixels wide. A line every 3.2 pixels is not a grid, it is a wash of grey. So the grid takes the smallest step that is both at least the ruler's own - which is what makes every grid line also a tick on the ruler, so the two can never disagree about where a minute is - and at least twelve pixels of your day. The clock strip at the top keeps its own ticks and is untouched; a module you have switched off keeps the grid rather than losing the only thing its dimmed bars are read against.

It costs nothing on the screen. The grid is painted, not drawn one line at a time, so your day is exactly as fast to open and to scroll as it was - which matters, because the ruler alone draws one tick for every minute at the closest reading.

Tap a stretch of someone's day and hand it over. Tap a person's marker on the timeline and a small card opens naming the job it is about - the module, its line where it has one, and the clock it runs from and to, in the same words the marker says under your pointer. Under it is a list of the other people on this day, and one press - "Move it to ...", named with whoever you picked - hands that job over. The two rows swap it the moment you press, and anything that now collides is drawn red, exactly as before.

What the card tells you before you press it. A module carries one person for all of its batches, so if the module you tapped runs four batches on that line, all four move together - and the card says so, with the number, rather than letting you find out afterwards. The list offers only people who have a row on this chart, because somebody who is not standing anywhere on your day is not an answer the day can give. If you have only one person on the whole day, the card tells you that instead of opening an empty list.

The card goes with the job. The moment you press, the job leaves that person's row and the card closes with it, so what is left on the screen is the day as it now stands. A tap on the marker again - now on the other person's row - opens a fresh card naming whoever holds it. A card left standing would have gone on offering to move a job off somebody who no longer has it, and the press under it would have done nothing.

The tap is read by the minute under your finger, not by the marker. A person's marker is 11 pixels tall in a 34-pixel row, and at the widest reading one minute of work is 1.2 pixels wide - so a fingertip lands beside a job far more often than on it. The tap takes the whole height of the row, and it takes the nearest stretch of work within about six pixels, which is what makes the narrow end of the scale usable at all. A tap where that person is not working is not swallowed: it still opens the person's own card, which is what every tap on a person's row did before.

Nothing of yours is rewritten. The only thing that changes is the one job you hand over with your own press. No module, no batch, no start time, no cycle and no saved scenario moves by being read, and there is nothing to run in Supabase - the whole of this release is drawn from the planner you already have on your phone.

What is not in it, so you are not left looking for it. Moving the person on one batch of a module, leaving its other batches where they are, cannot be done without a new stored field on the module and a one-time step in Supabase. It is a release of its own whenever you want it. And a clock repeated above the people's rows was the other way to reach the markers - you chose the grid, and the other way stays available.

Measured rather than promised. The grid's spacing at each of the four readings, its reach and the card's own list are all covered by tests, as is the card closing the moment the job moves: the suite is 1186 passing with none

failing, six tests new for this release and one of those standing on the card going with the job. The app's own More screen reads Engine v160.

23 Sep 2026 - engine v159 (no database step). Every card that opens over a screen can now be pushed aside by its own title bar, and the Scenario planner's chart window is one fixed height that no longer changes as you add people.

What you asked. In your words: "whenni click on batch to adjust the start time the pop up windows should be allow to drag" - and then, while I was working on it: "the chart window became shorter and shorter as no. of people grows. the window height should be fix".

What you settled, in three answers. You asked about the batch card and told me it should be **every** card, not only that one. You told me the position should **always reset** - a card opens where cards have always opened, and nothing is remembered. And you asked for **a faint grip line** on the title bar so it reads as something to hold rather than something to discover.

The batch card sat over the very bars it moves. The card that adjusts a batch's start time opens over the day it is about, and on a phone it covers the bars you are deciding between. That is what the drag is for: push it down out of the way, keep working, and it stays where you put it until you close it.

It is every card, not only the batch one. All of the cards that open over a screen share one title bar, so all of them now move the same way: the batch clock, a module, a person, an order, a customer, a product, a supplier, a unit and the rest. A card that looked the same but refused to move is a difference you would only find by trying it twice.

A card always opens where cards have always opened. The drag is a way to push a card aside while you work - it is not a setting and it is not remembered. Open the next card and it is where every card has always been.

The grip is what says so. There is a short faint bar near the top of every card's title bar, and the pointer over that bar becomes a hand. Measured, the bar is 34 pixels wide.

A card is held inside the screen, and that is deliberate. A card can be pushed until its own edge comes close to the edge of the screen, and no further. This matters more than it sounds. A card that could be dragged past the edge would hand the screen behind it something to scroll that it never had - and you could then scroll your own close button off the top, where no finger reaches it. Holding **the whole card** inside the screen means no button on it can ever be parked out of reach. On the phone a card is already nearly the full width, so what you get is vertical travel, which is the direction you want.

Measured rather than promised. On your One baker day the batch card for the mixing module is 337 pixels tall. Dragged down by 300, it moved by exactly 300 and its whole box stayed on the screen, and the screen behind it gained nothing to scroll. Dragged by 900 - far further than a finger can reach - it moved by 8 pixels and stopped, because 8 pixels is all the room it had. A module card, 784 pixels tall, moves 8 pixels; a person card, dragged 2000 pixels across and down, stopped with its whole box still inside the window; a wide card on a phone, 347 pixels wide in a 375-pixel screen, moves 8 pixels sideways, which is all the room it has. A drag on the card's close button is a press on that button and nothing else: measured, the card does not move and the cross still closes it.

The buttons keep working after a drag. Pressing the plus-five-minutes button on a batch card that has been dragged aside moves the batch and leaves the card where you put it - the card redraws its contents without going back to the top.

A drag on the day itself still scrolls the day. The title bar asks to be handled by hand; nothing else on the screen does. Dragging the chart - or the page - moves the page exactly as it did before, and the page does not move while you are dragging a card.

The chart window is one fixed height now. This is the second thing you asked for. The panel showing the day used to be a rubber band: its height was whatever the day's rows added up to, held between a floor and a cap. Measured at 375 pixels with one person it was 473 pixels tall and with a taller people list it was 648 - so the

window grew as you added people, and once it had grown to its cap the pinned people list at its foot carried on growing into the days above it. At the same phone size your eight-module day left eight rows of modules visible with one person and four with a crowd. The panel is now **one height** - 80 percent of your screen, up to 760 pixels - and it does not change whatever you add. On your phone it is 649 pixels at 375 by 812 and 720 pixels at 1280 by 900. Your own eight-module day draws 473 pixels inside it, so **all eight module rows are visible at once**, and even your tallest people arrangement - seven person rows, which draws 648 pixels - still fits in the window with nothing to scroll.

Nothing of yours was rewritten. Your three saved days were compared byte for byte before and after this release: One baker day, No fridge 1 person and My sister proposal 21/9/2026 still stand on the targets they had, and not one module, batch, start time, cycle or saved scenario changed. Nothing here blocks a sale and nothing here reads an order.

No database step. The whole of this is drawn from the scenario already on your phone, so there is nothing to run in Supabase.

22 Sep 2026 - engine v158 (no database step). The Scenario planner's people rows come down to one line each and hold still, everything about a person moves into a tip on their own name, the "PEOPLE" heading is gone, and the clock ruler now reaches a single minute where you are zoomed in the closest.

What you asked. In your words: "can i have all the peoples cards noted in tool tips as well?", "i dont need the PEOPLE title", and "make the ruler resolution to 1min". Every one of those is below. Where a thing left the screen you told me where it should go instead: every person's own notes into the tip on their own name, and the ruler stepped so a minute is drawn where a minute can be drawn.

The people rows are one line each, and their height can no longer move. A person's row used to print its own notes as text on the row - how much work they are carrying, how many places they have to be, and a red line when two of their jobs collide. That made it the only part of the chart taller than the bars beside it, and the only part that changed height as you worked: measured at 375 pixels, a row was 56 pixels on a clean day and 70 pixels when someone had a collision, so the whole strip measured 91 pixels one moment and 105 the next. **The strip sits pinned at the foot of the window**, which is what you asked for at v154 - and a pinned block that grows and shrinks as you scroll is a block that moves the very thing you are reading it against. The notes have left the row, so a person's row is **35 pixels** now, **on a clean day and on a day with a collision alike**, and the strip measures a flat 70 pixels either way. Those 35 pixels are what the row is: the same name cell the modules wear, and nothing else.

Everything about a person is in the tip on their own name. Point at a person's name on a computer and their own box opens beside it, carrying all of it: who they are (the whole label, including " (with 2, 3)", which the row shortens), how much work they are carrying, how many places they have to be, the red collision line in the words it already used, and **What they do today** - their jobs, each with its clock, up to six of them and then a count of the rest. It is the same list their own card carries, from the same builder, so the tip and the card cannot drift apart and say different things. The tip opens from the **title** alone, exactly as the module tips do since v156: pointing anywhere else on the row does nothing, and the tint that used to follow the pointer across the whole row is gone. That was your own v156 instruction - "only the note should shot as a too tip" - now applied to the last rows on the chart that had not had it.

Told rather than hidden: the tip is bottom-anchored, because the last row sits on the bottom edge. A module tip is centred on its row, which is what keeps a first or a last module's box off the edge of the window. The people strip is different - its last row **is** the bottom edge - so a centred box would be cut in half. A person's tip now sits on its row's own bottom line and grows **upward** into the day. Measured at the window's tightest, on a person with a collision and six jobs, the tallest possible tip is 196 pixels tall and it clears the clock strip above it with room to spare.

The widest-person trap, measured rather than promised. A person's tip is taller than a module's, so its lines have to be allowed to fold - and a folding box inside a 156-pixel name cell collapses to the width of its widest single word: measured live at 66 pixels wide and 923 pixels tall, one word to a line. The box is now told to be as

wide as its own content, and the rule carries a note saying why, because the fault looks like a styling detail and is not. Measured after: 272 by 180 pixels, sitting on its row's bottom line, inside the window.

The "PEOPLE" heading is gone. It named what every row under it already names - each one wears a person glyph - so the strip is now its rows with nothing over them. The strip keeps its line and its shadow, so it still reads as the foot of the day.

The tally row's three facts are in its own tip, and on a phone they are unreachable. The **People at once** row is the one you asked for: the most hands the day needs at once, the busiest stretch by name, and the minutes paid for twice. Those three facts are printed nowhere else in the app, and that row has no tap of its own - so once they are in a tip, a phone has no gesture that can open them. **You chose this, with the consequence in front of you** ("tip only, no card"), and it is written here rather than left to be found: on your phone those three facts are now unreadable. On a computer nothing is lost. Giving that row its own card is one line whenever you want it, and this note is the place that says so.

The collision is not lost by the row giving up its red line. Under the day there is still a note for every collision, in full: who it is, **two jobs at once** with the minute it happens, both jobs named on lines of their own, and the sentence saying what to do about it - move one along the day, or combine with another person and accept it. The red outline on the bar itself is unchanged. So the row's red line leaving costs nothing you were reading.

The ruler now reaches a minute where a minute can be drawn. The day is drawn at four scales, from 1.2 to 3.2 pixels to a minute, and a minute hairline is physically undrawable at the two wide ones - so the tick step follows the scale, which is what you chose ("a minute where it can be drawn"). **Wide** ticks every 30 minutes, 36 pixels apart, which is the ruler you had. **Standard** ticks every 15 minutes, 24 pixels apart. **Close** ticks every 5 minutes, 12 pixels apart. And **Closest** ticks every **single minute**, 3.2 pixels apart. The hour is still solid and labelled at every one of the four, and the half hour is still the dashed tick it always was - the one you judge a 15-minute pass by. Everything finer is a new thin hairline at reduced weight, so a ruler at Standard does not read as a ruler made of dashes. The count is honest: at Closest your 24-hour window draws 1441 ticks, and a full redraw takes 2.8 milliseconds against 2.3 at Wide - half a millisecond for the lot, so the ruler is drawn as elements rather than as a painted pattern.

And the cursor reads the exact minute. The hairline that follows your pointer down the day used to snap to the five minutes a batch's buttons step by. That was stricter than the screen it reads: a batch can be nudged by **one** minute, so the cursor could not name a time the chart was already able to set. It reads to the minute now - measured, a pointer 101 minutes into the day reads **5:41 am** where it used to read 5:40 am - and it still refuses to name a time outside your day, and still hangs its label to the left at the far end of the evening.

What has not moved, measured rather than promised. Your three saved days were compared before and after: **One baker day** still reads 24 pans across 8 modules, **No fridge, 1 person** still reads 24 pans, and **My sister proposal 21/9/2026** still reads 4 pans, each with the target it had. Their stored numbers were compared byte for byte and are untouched - the whole of this release is drawn from the scenario already on your phone. Nothing here blocks a sale and nothing here reads an order.

No database step. Not one stored field is added or changed, so there is nothing to run in Supabase.

CHANGELOG and the guide (section 23) both carry all of this, both PDFs rebuilt and read back, the app's own More screen reads Engine v158, and the tests are 1174 passing with none failing - eight of them new: that no person row prints its own notes and the tip carries them all, that the tip and the person's own card agree word for word, that a collision changes nothing about a row's shape, that the tip is a sibling of the name and not inside it, that the two tip rules both survive as file text, that the tally's three facts are in its tip and off its row, that the "PEOPLE" heading is gone, and that the ruler's step follows the scale at all four stops.

22 Sep 2026 - engine v157 (no database step). The Scenario planner's list of controls is much shorter and the modules window is taller: the module tags have moved into the tip, the scale is two presses around the stop it is on, the people are one box, and "The line" has left the row for the module's own card.

What you asked. In your words: "i want to make the modules window taller. The module tag, can you put them

in the tip tips? Make the scale with+ -, People: make it drop down option, The line: this name is very unclear,, we just need the button, WORK BACKWARD. The idea is to make this list shorter and make room for a bigger modules windows." Every one of those six is below, and where a thing left the screen you told me where it should go instead: every tag into the tip, the line's own settings into the module's card, and the two paragraphs above the chart removed.

Two paragraphs left the day. The two sentences that used to stand between the "The day" heading and the first bar are gone: the one teaching how to read the chart, and the one naming the moment your day hangs from. At 375 pixels they measured 330 pixels between them, 210 and 120, and they were the whole of what you scrolled past before you reached a bar. With them gone the first bar sits 278 pixels below the heading where it sat 707. Nothing is lost by it: how to read the chart is section 23 of the operations guide, and the button that moves the day now says its own name in the row under the chart, so the move is still findable without a sentence pointing at it.

The modules window is taller, and now it fits. The window has a floor of 500 pixels and a cap of 760 where before it had only a cap of 560: the floor is what makes it a window rather than a box that resizes itself every time you edit a module, and the raised cap is what lets a long day use the room the rest of this release gave back. On your seeded day at 375 pixels the window is 616 pixels tall where it was 536, and the day's own drawing is 616 - so before this release the window hid 80 pixels of your day and now it hides none of it.

The tags are off the row and into the tip. Every tag a row used to wear is gone from the row: how many of a module you have, how many of its lots are in it at once, whether it waits above, whether it runs itself, and whether you have asked for more passes than a day holds. Point at the module's own title and the tags are the first line of the box, joined with a dot, with the two notes under them. Measured at 1280 wide on your day, the box is 42.6 pixels for the two notes alone and 58.9 with the tag line, and all nine boxes measured inside the panel left and right.

The tip is now taller than the row it belongs to, and here is why that is all right. Three lines of tip is 58.9 pixels against a 45-pixel row, so the box no longer fits inside its own row the way v156's did. What keeps it inside the panel is that it is centred on its row: half of it goes up into the clock strip's row above your first module, and half of it goes down into the rows below your last one. Every tip was measured inside the panel at both ends of the day, and nothing on the chart moves while one is open.

The scale is two presses. The four stops are kept exactly as they were, with the same pixels behind them: Wide, Standard, Close, Closest. What changed is that they are no longer four buttons. There is a round minus and a round plus, and the stop you are standing on is the word between them, so you can see where the dial is without opening anything. At either end of the dial the press that would go past the end is off rather than doing nothing when you tap it.

The people are one box. Three chips became one drop-down. What the box says when it is closed is the arrangement in force right now - Sharing them out, One a module, One to a line, Your own, or Combined when you have combined two people yourself - so the box names what it has done rather than hiding it. The four lines inside it are the arrangement in force and the two things it can do to your day, One a module (a person for each job, one job each) and Share them out (as few hands as can cover the day), and Combine two people on the end.

"The line" left the row, and its name with it. You said the name was unclear and that its place is the module's own card, and it is there now. The control row reads Scale, People, Walk the day and nothing else. On the module's card, under the words **Waiting on the module above**, is the same sentence the old card opened with - which of your modules wait on the one above, by name, and which keep the times you placed - and the same two presses it always offered, Chain every module and Take the waiting off. One module's own three answers, As the one above finishes, Never before the one above finishes and Its own time, are untouched and still on that same card, which is where they have been since v154.

The button alone. You said we only need the button, so the label that used to stand beside it is gone. It read "Your day backwards · first batch out 9:09 am", which at 375 pixels wrapped to two lines: 54 pixels of a phone

screen spent explaining a button that already says what it does. What is left is the button, **Work the day backwards**, in the row under the chart.

What the row came back to. At 375 pixels your control row is 222 pixels tall where it was 305, so 83 pixels came back to the chart. Together with the two paragraphs that left and the taller window, the run from the "The day" heading to the first bar is 278 pixels where it was 707, and the day itself is not scrolled to at all: the window is 616 and the drawing is 616.

Nothing on a phone is left without a way to read it. A tip opens where there is a pointer that can hover at things, and a phone has not got one - measured at 375 pixels with a touch phone emulated, no tip opens at all. Everything a tip carries, the tag line and the two notes, is written out word for word at the top of the module's own card, in the row's own tag colours, and that card is what a tap on the row has always opened. There is nothing new to learn and nothing new to find.

What has not moved, measured rather than promised. Your three saved days were compared before and after this release. One baker day still reads 24 pans across 8 modules, its first batch out at 9:09 am and its day finishing at 1:30 pm. No fridge, 1 person still reads 24 pans. My sister proposal 21/9/2026 still reads 4 pans. Not one module, batch, start time, cycle or saved scenario of yours changed in this release. Nothing here blocks a sale and nothing here reads an order.

No database step. The whole of this is drawn from the scenario already on your phone, so there is nothing to run in Supabase.

The app's own More screen reads Engine v157. The tests are 1166 passing with none failing, eight of them new: that no tag is left on any module row and that every tag the tip carries is the same tag the module's card carries word for word, that a module you have two of is carried the same way, that the taller tip is still held inside the panel it is drawn in, that the scale is exactly two presses around the name of the stop it is on and that each press is off at its end of the dial, that the people are one drop-down whose closed text is the arrangement in force, that the line's own group has left the control row for the module's card, that the window is taller and wider-capped and the two paragraphs above it are gone, and that the day-backwards group is the button and nothing else.

22 Sep 2026 - engine v156 (no database step). The module's notes now open as a tip over the day when you point at its title, and the row keeps its height while you read them.

What you asked. In your words: "i prefer not to expend the module with mouse hoover. Remain the height. When mouse hoover to the module title, only the note should shot as a too tip." That is a correction to v155, and it is a fair one: v155 brought the notes back by making the row taller, which moves every row below it the moment you point at it, so the bar you were reading is no longer where you left it. Reading a number should never move the thing you were reading it against.

The tip, and where it comes from. Point at a module's own title and its two notes open in a small box beside the name: the time the module starts and how many batches it runs, then what a batch costs in minutes and in your own hands. It is drawn over the day rather than in it, so nothing on the chart moves while it is open - not the row it belongs to, and not one row under it. Point away and it is gone.

Measured rather than promised. On your One baker day, with the tip open on the mixing module, the row reads 46 pixels before and 46 pixels after, the row underneath starts on the same pixel it did, and the day's own total is 527 pixels either way. The box is 43 pixels tall, which is inside the row it belongs to - that is deliberate, because it means a tip on the first or the last module of a long day cannot be cut off by the edge of the panel. The whole box measured inside the panel on the first module and on the last one.

Only on a computer, and only on the title. The tip opens where there is a pointer that can hover at things, which a phone has not got: on your phone pointing is a finger and a finger is a tap. There the notes are written out word for word at the top of the module's own card, which is what a tap on the row already opens, and that is unchanged from v155. Pointing at the row anywhere other than the title does nothing, because you asked for the note and not for the row to react.

What has not moved. Everything v155 did to the row is kept: a row is still exactly as tall as the bars it draws, the name still takes the row's first line with the badges on the line beneath, the name column is still 156 pixels, and

your whole day still fits the window. Your three saved days were compared byte for byte before and after: One baker day still reads 24 pans across 8 modules with its first batch out at 9:09 am and its day finishing at 1:30 pm, No fridge, 1 person still reads 24 pans, and My sister proposal 21/9/2026 still reads 4 pans. Nothing here blocks a sale and nothing here reads an order.

No database step. The tip is drawn from the scenario already on your phone, so there is nothing to run in Supabase.

This change history and the operations guide (section 23) both carry all of this, both PDFs rebuilt, the app's own More screen reads Engine v156, and the tests are 1158 passing with none failing - one of them new: the tip is off the row, laid out absolutely so the row cannot grow, centred on the row it belongs to, opened by the title alone, and its two lines are held to a height that fits inside a row.

22 Sep 2026 - engine v155 (no database step). Every module row is now as tall as its own bars, so your whole day fits the window with no scrolling; the two note lines come off the row and come back when you point at it on a computer, and they are always written out on the module's own card.

What you asked, and you were right. In your words: "The window for scrolling module became too small, why not reduce the height of each module, the notes can just shown up upon mouse hoover." You were right that the rows were too tall, and the cause was not the bars. On your One baker day the eight module rows measured 912 pixels of drawing inside a window about 536 pixels tall, so only about four of them could be seen at once and the rest had to be scrolled to. The bars themselves need only 45 pixels; it was the words beside them that made the rows tall. The name column is 138 pixels wide, and a name like The rests and the stretch and folds wraps to seven lines in it, with each of the two note lines wrapping to two more.

Every module row is now the height of its bars. A row measures 46 pixels on its own day, and the whole of it - the clock strip at the top, your eight modules and your people at the foot - now measures about 502 pixels inside the 527 the window gives it, so the day is read in one look and nothing is scrolled to. Every number in it is the number it was.

The two notes leave the row, and two things bring them back. On a computer, pointing at a row brings both notes back and tints the row, and the row grows taller to hold them rather than covering what is beside it, so a note is never cut in half. On your phone there is no pointing, so a reveal that needed one would be a reveal your phone could never reach; the notes are written out instead at the top of the module's own card, which is what a tap on the row already opens. Nothing new to learn and nothing new to find. On your mixing module the card now opens with what the row used to whisper: the time it starts and how many batches, then what a batch costs in minutes and in your own hands.

The name takes the row, and the badges drop below it. This was the real damage of the old row, and it is worth naming. Your longest module name is 273 pixels wide, and it was sharing its line with a small badge, which left just 35 pixels of the name to read - about two words of it. The name now has the row's first line to itself and the badges sit on the line beneath, and four of your eight module names now read whole where one did before. The name column is 18 pixels wider as well, from 138 to 156, which is the difference between one and three more of your names reading whole. A name still too long for the row is shortened with three dots, and the whole of it is on the module's card and under the pointer on a computer.

A fault found while building it, worth naming. A module drawn as two production line and set to wait on the module above it could not be drawn at all: the drawing asked for the waiting badge's words in a place where they did not exist, and that is not a blank space, it is an error that stops the screen. Two taps on your own module card would have found it. None of your saved modules is drawn as two production line, which is the only reason you never met it. It is fixed, and a test now stands on it so it cannot come back.

What has not moved, measured rather than promised. Your three saved days were compared byte for byte before and after this release. One baker day still reads 24 pans across 8 modules, its first batch out at 9:09 am and its day finishing at 1:30 pm. No fridge, 1 person still reads 24 pans. My sister proposal 21/9/2026 still reads 4 pans. Not one module, batch, start time, cycle or saved scenario of yours changed, and the change is in how a row is drawn, not in what any row says. Nothing here blocks a sale and nothing here reads an order.

No database step. The row is drawn from the scenario already on your phone, so there is nothing to run in Supabase.

This change history and the operations guide (section 23) both carry all of this, both PDFs rebuilt, the app's own More screen reads Engine v155, and the tests are 1157 passing with none failing - three of them new: one on the notes being carried by the row and printed by the card, one on the name taking its own line ahead of its badges, and one on the module drawn as two production line that waits on the module above.

22 Sep 2026 - the documents, not the app (no engine version, no database step). The guides and this change history print their own italics now, so the stray asterisks are gone.

Every one of these documents is built from a markup in which one asterisk on each side of a word means "print this in italics". The builder stopped understanding the single asterisk at some point and reads only the pair, so a word meant to lean quietly came out wearing its own asterisks instead: *'why no button for work backward'*, *Cutting and packing*, *status New*. Not one sentence read wrongly, which is why it went unnoticed for so long. There were 116 of them in this change history, 44 in the operations guide, and 9 more across the money and profit guide, the alert guide and the safety guide. All of them now use the pair of underscores the builder does understand, so they print as the italics they were always meant to be.

The asterisks that were meant to be read stay exactly as they are: the one marking a required field on the order form, and the fill-in lines on the weekly note sheet. Those are marks, not markup.

Nothing in the app changed. More still reads Engine v154, every screen is exactly where it was, and no test moved. All five documents were rebuilt and read back to check: this change history prints no asterisk at all now, and the only two left in the operations guide are the two that are meant to be there.

22 Sep 2026 - engine v154 (no database step). A module now says how it takes its start - three answers where there was one switch - so the packing begins the minute the cooling ends; every batch's own card offers the move again; and the people's rows are held below the modules so you can read a slot against them.

What you asked. Your three reports, in your own words: "the cut and packing batch din follow the earlier batch end", and then your own correction, "cutting and packing sit below cooling down, so cutting and packing batch start should follow cooling down batch end". "The delta t disappeared, before this we have it. I want each batch start time to be adjustable, like the 1st module. Just need to show delta on the batch 1st offset only, then following module of that step dont have to show the delta because it follow the previous module tightly." And "I want to freeze the persons card, so that by scrolling thru modules i can see exactly where that slot of that person tie up to and searching for opportunity to move some batch start time to reduce the number of person needed." You also set the rule the first of them had to be built by: "the behaviour has to base on configuration, not a hard wired".

1. **What was actually wrong, and it was words against behaviour.** The old switch was called **Waits for the module above**, and its hint promised a follow - *2nd module 1st cycle should follow 1st module 1st cycle completion*. What the day really did with it was a **floor**: a batch could be pushed later than the module above, never pulled back onto it. So a packing step sitting 8 minutes after the cooling ended could never come back onto it, however the switch was set. Your report was right and it was the most useful kind: the sentence on the card and the arithmetic under it disagreed.

2. **A module now says how it takes its start, in one of three answers.** Open any module and the card asks **How this module takes its start**, with three choices: **As the one above finishes** (batch 1 of this module starts the minute the module above finishes its own batch 1, batch 2 on batch 2 - a tight follow with no gap); **Never before the one above finishes** (a batch here can never start before the module above has finished that same batch, but a later time you set yourself stands - this is the old switch, kept exactly as it behaved); and **Its own time** (the module above has no say over this one). Each one carries a sentence saying what it means before you tap it.

3. **The choice is on the module, not on a module's name.** This was your rule and it is the way it is built: nothing in the day knows a module by what it is called. Set **As the one above finishes** on your packing and it

lands on whatever really sits above it - cooling, or anything else you put there, in whatever order you keep your modules. There is a test that renames every module of a day and confirms the same answer, so this cannot quietly become a rule about cooling.

4. **Nothing of yours was re-laid by being read.** A saved module keeps the behaviour it already had: a module that was waiting on the one above it opens on **Never before the one above finishes**, and a module that stood on its own opens on **Its own time**. A **new** module arrives on **Never before the one above finishes** too, because **As the one above finishes** on a module you have not yet told what feeds it would move your day the moment you added it. Your three days were compared before and after: 24 pans on One baker day, 12 on the seeded line, 4 on your sister's, and not one start time of yours changed.

5. **Every batch's card offers the move again - this is your second report.** A batch below batch 1 of a module set to wait used to be **refused** the two button pairs, with a sentence telling you to go and change the module instead. Every batch at every module now offers **Five minutes at a time** and **One minute at a time**, and the last paragraph of this note says plainly what a move means in each place so the two cannot disagree.

6. **The offset is read out on batch 1, and only on batch 1 - exactly as you asked for it.** On batch 1 of a module that is not the first, the card reads how this module takes its start, in words: **follows the module above**, **never before the module above**, or **on the line**. The batches below batch 1 do not repeat that, because they ride the module above at this module's own pace - there is no offset of the module's to read on them. A batch you have **held** on purpose is the exception, and it is a deliberate one: it reads **delta t = +5 min** on its own card, because a batch that is off the line must never look like one that is on it.

7. **A move below the first module is a hold, and a hold can only ever put a batch later.** That is not a limitation of the buttons: a batch cannot start before the dough it is made of exists. What makes it safe rather than a wall is that a hold is an **offset measured from where the line puts the batch**, not a time - so move the module above it, or change how many batches it runs, and the held batch comes with it, keeping its place in the queue. The card of a held batch carries **Back onto the line**, which takes the hold off, and taking a hold all the way off now says so - *Cutting and packing is back on the line* - where it used to say *held 0 minutes behind*, which is a sentence about nothing happening.

8. **A press on batch 1 of a module that follows moves the whole module.** A module set to **As the one above finishes** has no start time of its own to write: it begins where the module above ends, whatever a stored start says. So a press there writes the module's own hold instead, and **every** batch of it comes with it - which also keeps a press on batch 1 meaning the same thing on every module, rather than being a button that appears to do nothing on one of them.

9. **The people are held below the modules - this is your third report.** The module rows scroll inside the day's own panel and the **People** rows and the **People at once** total are pinned at its foot, so scrolling down through eight modules never takes the person you are reading away from the bar you are reading it against. This is the move the whole report was for: hold a slot in view, look for the batch you can move to empty it, then go and move that batch. The red now-line of the walk through still runs over the people's rows, because that line is the day happening rather than a reading of it.

10. **One fault found while testing, and it is worth writing down.** The test that stands in for the browser kept an attribute a real browser also exposes as **dataset** - and the chart reads which batch you tapped off that. So every simulated tap had been reading as **batch 1**: a whole class of taps the tests could not see, which is exactly why "the delta t disappeared" could be true on your phone while a test suite said all was well. The stand-in is unforgiving now, the way the browser is. Two tests written against it failed immediately and both were real: the wording of a hold taken off, and the day's own card opening where a batch's card should.

11. **No database step, and it is on every phone.** Nothing to run in Supabase. The planner travels with the app's own settings row, so this new per-module answer goes up with everything else the planner holds and comes down onto your other devices the same way - the choice you make on your phone is the choice your tablet reads.

12. **What has not moved.** One baker day still reads 24 pans across 8 modules with its four stored rhythm times

untouched; the seeded line still makes 12 with the chiller as the wall; your sister's line still makes 4. v152's backwards press and v153's button on the row are unchanged, and neither this release nor those reads a single order or blocks a single sale.

The CHANGELOG and the guide (sections 23 and 26) both carry all of this, both PDFs are rebuilt, and the app's own More screen reads Engine v154. The tests are 1154 passing with none failing - 14 of them new: 8 on the model (the three answers and how a saved module is read through them, the packing landing on the cooling's end, the follow being a choice on the module and not a rule about a name, a switched-off module stepped over, a module above with fewer batches answering for the batches it has, a hold riding the module above and never pulling a batch earlier, and her three days reading the same times) and 6 on the screen (the three pills on the module card, the packing landing on the oven, the pairs on every batch, the offset read on batch 1 alone, a press on batch 1 of a following module moving the whole module with its way back, and the people's rows held below the modules). No database step - nothing to run in Supabase.

22 Sep 2026 - engine v153 (no database step). The day's own row works your day backwards - one press, and it is a button that acts instead of one that only tells you where to tap.

What you asked. Your words: *'why no button for work backward, this is to reeposition the batches latest start time'*. You were right, and it was about a control that only talked: v152 put the move behind a tap on one bar and left the row carrying a chip that named the moment and then told you which bar to tap. A thing shaped like a button whose whole effect is to send you somewhere else reads as a dead control, and that is what it was.

1. **The button does the work now.** The row under the chart reads **Your day backwards · first batch out 9:09 am**, and the one pressable thing beside it is **Work the day backwards**. One press and the day is worked back on the spot - no card in the way, the same seven modules pulled to their latest starts, and the same message saying what moved and what it cost your hands.

2. **It is the card's own press, not a second one.** Same writer, same words, same snapshot, same way back - so the two doors cannot drift apart. Press from the row, then open the card on batch 1 of *Cutting and packing*, and the card is showing you the day the row just made, with the way back on it.

3. **Pressed on a day already worked back, it says so and writes nothing.** This was your own choice and it is the rule the button lives by: a button that comes and goes reads as a fault, so it stays in the row on every day. On a tight day the press answers *Every module of your line is already as late as it can go - there is nothing left to pull back.*

4. **The way back is beside the button that moved the day.** While there is something to put back, the row carries **Put my start times back** next to **Work the day backwards** - the pattern Stop beside Start already uses. The card keeps its own copy, so the undo is either one tap away or three, and never missing. It is still one-shot and still lasts only while the screen is open: leave the screen and the day you pressed is the day you have.

5. **A line with no chain says so, rather than drawing a dead button again.** A day of a single module is not a chain - its first batch is simply that module's own start time - so the button is drawn, carries no moment in its label, and says in words that there is no module above the last one to work back through. The fault you found is the one thing this release refuses to repeat.

6. **The sentence above the chart now names the button.** It used to tell you to tap batch 1 of *Cutting and packing* and stop there, which was the only door v152 had. It now names both: the button in the row does it in one press, and tapping batch 1 opens the same move with the whole chain laid out before you press it, which is the one to read first.

7. **One fault found while building it, and it was in this release's own new code.** The undo works off a snapshot of the day kept in the screen's own memory and never stored, and v152's card took it from the right place. The row's new press first took it from the app's saved state instead, which put that snapshot into your data - and a reload would then have offered you a way back to a day you had already left. It was found by measuring the stored data after a press rather than by reading the code, it is fixed, and there is now a test that fails if the snapshot ever reaches your data again. Nothing of yours was changed by it.

8. **What has not moved.** Your three shelf days were compared byte-for-byte before and after: *No fridge, 1 person* still reads 24 pans, *My sister proposal 21/9/2026* still reads 4, and *One baker day* still reads 24 pans with the first batch out at 9:09 am and the day finishing at 1:30 pm. A press followed by the way back leaves the stored data byte-for-byte identical to what it was - measured, not promised - and the snapshot is gone when the screen is. Nothing here blocks a sale and nothing here reads an order.

CHANGELOG and the guide (section 23) both carry this, both PDFs rebuilt, the app's own More screen reads Engine v153, and the tests are 1140 passing with none failing - five of them new: the row press moving the day exactly as the card does, a finished day keeping its button and saying it is done, the way back sitting beside the button that moved the day with the card offering it too, a line with no chain saying so instead of drawing a dead button, and the one that pins the snapshot out of your saved data. No database step - nothing to run in Supabase.

22 Sep 2026 - engine v152 (no database step). The planner can work your day backwards - the day hangs from your first batch, and every module gets its latest start.

What you asked. Your words: *'how to make the calculate backward works?'* The backward count already existed on More -> Production line (section 24), where it counts back from the minute the first pans must be at the oven. The Scenario planner had none: it could only ever push a day later, never pull it earlier. So the planner has one now, on the planner's own day.

The moment your day hangs from is your own last module, not the oven. The anchor is the end of your first batch at the last switched-on module of your line - whatever that module happens to be. Add, remove or reorder a module and the anchor moves with your line. On your own One baker day that module is Cutting and packing, so the day hangs from 9:09 am.

Where it is, and how you find it. Tap the B1 over the first bar of that last module - the card you already get by tapping a batch - and it is now the day's own card. It reads the module, the end of your first batch and your day's finish, then every module with the time it sits at and the time it could sit at. On your line it reads: Mixing the dough in the tub 4:01 am to 4:43 am, The rests and the stretch and folds 4:21 to 5:03 am, Oil the pans and weigh the dough out 6:24 to 7:06 am, Into the proofer 6:39 to 7:21 am, Dimple and top 7:24 to 8:06 am, The proofer again 7:30 to 8:12 am, The oven swap and the bake 8:00 to 8:42 am, and Cutting and packing 8:57 am with a tick. A tick means that module is already as late as the line allows. So the control is findable rather than hunted for, the day card also carries a chip reading First batch out 9:09 am with the sentence saying which bar to tap.

The rule is measured off your own day, not subtracted from the anchor. This is the one place I departed from what I first planned, because the arithmetic was wrong on your own numbers. Counting backwards is not the sum of the cycle minutes: on your 24-pan line the eight modules' minutes add to 853 while the chain you actually run spans 1107 minutes, because your modules do not sit end to end - the fold loop and the proofer are wider than the few minutes a batch takes in them. Take 853 off your anchor and the mix is handed a LATER start than the one it already has, which is the opposite of what you asked for. So the calculation measures the room between each pair of modules as your stored times really stand, and pulls each module back by that room. On your line it gives the mix 4:43 am, and the 42 minutes it gains are the slack your day really has.

One press, and a second press does nothing. Pull them back to their latest start moves the seven modules that can still come later and leaves every other one exactly where it is. A module that is switched off is stepped over, exactly as the day already steps over one. A press never moves a module EARLIER than it already stands, and pressing it again finds nothing left to take.

It is a real change, and the way back is on the same card. Said plainly rather than quietly: this rewrites the start times of your switched-on modules and saves them into your scenario. It is not a preview. So while you are on that screen the card carries a second button, Put my start times back, which puts every module the press moved exactly where it was and then goes. It is one-shot, and it lasts only while the screen is open, because it is never stored - the same reason the walk through's own clock is not. Leave the screen and the day you pressed is the day you have.

Below that are the same two pairs of buttons you already know, and here they move the whole day. Five minutes at a time and One minute at a time, each with an earlier and a later button. Press one and every module moves together by that amount, so the day keeps its shape: on your line all eight modules move five minutes and every gap between them is the number it was. That is the move for reading what an earlier or a later start does to the whole day without re-drawing it.

A tighter day can need another pair of hands, and the app says so instead of hiding it. This is the honest cost of taking the slack out. Squeeze the day to its shortest span and two of your jobs can land in the same minute, so a press can raise the people the day needs - on your One baker day from one person to two. Nothing blocks and nothing is refused: the toast reports it in words, the People area below shows the collision in red, and the button under it takes it back.

The clock you point with now follows your finger, and the clock strip stays where you can read it. Your four asks of that day chart, all in this release. The reading follows the pointer DOWN the day as well as across it - it used to sit at the top of the chart, which is the right place only while the top of the chart is on screen, so pointing at a bar four rows down meant reading a clock a screen away. It rides 8px clear of your finger rather than under it, because a reading under your own finger hides the bar you are holding it against - and near the top of the chart, where there is no room above, it drops below the pointer instead.

Lifted at v171: the 8 pixels above your finger described here are two CSS inches now, because 8 pixels is under a fingertip rather than clear of it. The rule itself is unchanged and this entry still describes it: above your pointer where the window has the room, and below it where it does not. Read the v171 entry above for the rest.

The chart is a panel of its own now: the clock strip is pinned at the top of that panel and the rows scroll under it, so the hours stay beside the bars you have scrolled to. The panel is capped, so it never grows past what you can see at once.

And one fault found while building it. The red now-line of the walk through was being drawn in the chart's name column - 138px, about an hour and a half, to the LEFT of the minute it was naming, because its position was measured when the chart was built and a chart that is not yet on the page has no position at all. Measured off the clock strip's own column instead, it now stands on the minute it names. Nothing about the walk through changed except that the line is where it says.

What has not moved. Your three shelf days were compared byte for byte before and after: No fridge, 1 person still reads 24 pans, My sister proposal 21/9/2026 still 4 pans, and One baker day still reads 24 pans across 8 modules with the first batch out at 9:09 am, the day finishing at 1:30 pm and the stored rhythm of 4:01, 5:28, 6:55 and 8:22 am untouched. Not one module, batch, start time, cycle or saved scenario changed, because nothing moves except by your own press. Nothing here blocks a sale and nothing here reads an order.

CHANGELOG and the guide (section 23, cross-referenced to section 24) both carry all of this, both PDFs rebuilt, and the tests are 1136 passing with none failing - nine of them new: four on the calculation (the exact latest start of every module on One baker day, that it is measured off the day rather than subtracted from the anchor, that a switched-off module is stepped over, and that a press never moves a module earlier nor does anything a second time) and five on the card and the drawing (the card opening on batch 1 at the last module and nowhere else, one press leaving the day's own numbers alone, the way back putting every start time exactly where it was, the step pairs moving the whole day with its shape kept, and the chip signposting the moment the day hangs from). No database step - nothing to run in Supabase.

22 Sep 2026 - engine v151 (no database step). The module card answers to you, every batch is numbered, and the day calls your people.

What you asked for, in your own words. Eleven things in one message, all of them inside More -> Scenario planner: the scenario's name first; the module card asking only what you alone can answer; a B number over every batch; moving a batch at any module the same way you move one at the first; naming Person 1 and Person 2; a marker that reads as a person; and the day calling you a minute before each job. It all ships together, as you asked. One of the eleven was not a wish but a fault: a person's card could not be opened at all before this, which is why naming Person 1 and Person 2 had no way in.

The scenario's name comes first. The card at the top of the planner now reads What this scenario is called, then My day starts at, then Pans a day you want from it. Going to that screen to rename a scenario used to mean reading past two numbers you did not come for.

The module card no longer asks you two things you do not think in. Which job of your line is this has gone. How you place each batch's start time has gone as well, and you said an alternative is coming - so the batches' stored times are kept exactly where they are, and nothing reads differently for it. What taking those two blocks out costs you is written at the end of this note, honestly.

A new module now arrives ready to use. One cycle of 20 minutes with all 20 of them as your hands at the start, Pans in one batch 1, waits for the module above switched on, allow multiple production line switched on. Before this a new module arrived at 15 minutes with both switches off, and you had to type five things before it was a module at all.

Batches follow the one above, unless you say otherwise. Minutes from one batch to the next and How many batches it runs in the day both carry an Auto. Auto is written as an empty box with the word in it, so you can always tell a number you set from a number the day worked out. The batch count is a box on the first module - that is where your day's number comes from - and on any module that has a number of its own. A module on Auto says its number out loud and names what it is following, and beside it is one button, Give it its own number, which keeps the count it is running now so taking the wheel never moves your day a minute. Auto spacing is not new behaviour: the day has always read a module with no pace of its own as one batch following the one before it end to end. It has a name now.

How many production line do you have - the number is unchanged, the words are yours. You said it yourself: "that 2 refer to 2 production line, if the meaning is for that purpose, a better remark needed for that. no functional changes needed." So the behaviour is exactly what it was, and the remark is rewritten around production line instead of the word lines. It also says plainly what happens when Allow multiple production line is on: then the module is one production line with one person, and this number is not in force. And Let its batches overlap is called Allow multiple production line now, everywhere on the card and in the note under the day.

Duplicate this module. At the foot of the module card, beside Delete, there is now Duplicate this module. The copy lands directly under the one you copied and carries its cycles, its numbers and its person exactly - a copy is a copy and none of the new-module defaults reach it. You asked for this because you can add modules as you wish now, and two ovens of yours are two of the same thing.

Every batch is numbered, on every module and every production line. A small B1, B2, B3 above each bar, in a band the row grew for it, on every module including one drawn as two production line. The number carries the batch it names, so tapping B3 opens batch 3's own clock - a number you use to tell one batch from another is no use if it is not also a way to pick one.

Move a batch at any module, not only at the first. Tapping a bar or its number opens that batch's clock, and the card now says the batch number and its delta t beside each other. On the first module a move is the batch's own start time, which is what it always was. On every later module a move is a delta t - a hold measured from where the production line puts that batch - so the batch rides the chain: move the module above it and this batch comes with it, keeping its offset. A delta t can only ever hold a batch back, because a batch cannot start before the dough it is made of exists, and the tag over the bar shows delta t = +5 whenever there is one. When your batch is held back on purpose the card also offers Back onto the line, which takes the hold off, so a nudge is never a one-way door.

Your people have names, and the card that was unreachable opens. This is the fault you reported: "in the person card, now person card is not accessible". The person's row had no handler of any kind, so no gesture anywhere opened their card. Tapping a person's row below the day now opens it: a box for what you call them, a switch for whether the day calls them, and the jobs they are on today. The name is kept in the app's settings rather than in the scenario, so you type it once and every scenario uses it - person numbers start again at 1 in each scenario, so the name you give person 1 shows wherever person 1 is working, and the card says so. The card's own heading and its "Call them a minute before their next job" line follow the name as you type it, so the

card never goes on calling them Person 1 while the row behind it already says their name. The box itself is never rebuilt while you are typing in it, which is what keeps your cursor where you left it.

A person's marker is theirs now, not the module's. The People rows used to be tinted by the module each stretch of work came from, so one person's row was eight colours and read as nothing in particular. Each person now has their own colour, with their name on the stretch where it is wide enough - so a row reads as one worker's day and two rows are two people. A collision still wears its red outline over the top.

The day walks, and it calls you. A Start the day now button in the Walk the day row, with Stop beside it while it runs, and a red now-line that follows the real clock down the chart. The call comes one minute before each job a person owns - one minute and not the minute of it, because your fold is a one-minute job and telling you to fold as you should already be folding is telling you too late. A call card appears over everything: their name, their own colour, the job, and the clock time it is for, with one OK. OK takes it away. If you do not tap it, the next call simply replaces it rather than queuing behind it, which is your own rule. Each person's call is a different note, so two people called in the same minute are two sounds. The card and the guide both say the limit plainly: it runs only while this screen is open and awake, and a phone that locks or a browser that closes stops it. It is a call for the day you are walking through, not an alarm that survives the app being shut.

The call names the cycle only when you have named it. A module's own name can be a mouthful ("The rests and the stretch and folds"), and the thing you can actually go and do is the cycle inside it, so the call carries the cycle's name as a line of its own when that cycle has one. An unnamed cycle gives no line at all. That last part is a fault this release nearly shipped: the line was drawn through the same number-formatting helper the rest of the chart uses, and on a cycle with no name that helper answers "0", so the card went out printing a bare 0 above the clock. There is a test on it now, and it is written so it fails if the line ever comes back.

Your day start is never rewritten by it. Starting the day sets the run's own minute zero to the minute you press, and nothing is saved - so a 4:01 am plan started at 9 am calls at 9 am, and your scenario's own day start is exactly where you left it.

What this release takes away, said plainly rather than quietly. Three things are gone with the two blocks removed from the module card. The job list is gone, so the way to CORRECT a module whose name the Production line does not recognise is gone with it - a module still crosses over by what you call it, exactly as before, and every module of yours already crosses. The per-batch start boxes are gone from the card, so the one control that put every batch back on its own even spacing, Space them evenly again, is gone too; tapping a bar gives you the same five-minute and one-minute moves, but not that reset. Say the word and any of the three comes back in one line.

What has not moved. Your two saved lines and the seeded starting line were compared byte-for-byte before and after this release: No fridge, 1 person still reads 24 pans, My sister proposal 21/9/2026 still reads 12, and the seeded starting line still reads 12 with the chiller as the wall. Not one module, batch, start time, cycle or saved scenario of yours changed, because every new default is for a module you have not filled in yet - a stored number or a stored start is yours and is never rewritten, which is the same rule the cycles migration used. Nothing here blocks a sale and nothing here reads an order.

CHANGELOG and the guide (section 26) both carry all of this, both PDFs rebuilt, and the tests are 1127 passing with none failing - fourteen of them new: five on the model (a new module's own defaults; Auto following the one above while a number of your own is never overruled; a delta t riding the chain and never pulling a batch in front; reading a line twice still a no-op; all three seeded lines pinned batch time by batch time; and the call landing one minute before the job) and nine on the drawing and the card (a batch number on every bar including two production line, a copy carrying its source's cycles, the person's name on the row, the card's own heading following it, a named cycle on the call and no line at all for an unnamed one, and the four halves of the live clock). No database step - nothing to run in Supabase.

22 Sep 2026 - engine v150 (no database step). The word is cycle - and a batch of four now reads as four.

What you told me, in your own words: "there is not steps term in this project, i want the cycle shown in

batch, and the cycle labour shown visually". You were right on the first point, and the app was the thing saying it. I measured the day chart before changing anything: nothing had disappeared. On your One baker day the fold is one batch of four cycles and its bar draws four shaded bands with your three folds on them, and the oven is one batch of two cycles with your two-minute swap on the second - 24 cycle bands and 32 stretches of your hands across the day. The feature was built and it was drawing. What was wrong was the word on it.

Why your own two saved lines look plain. Both the line you call No fridge, 1 person and My sister proposal 21/9/2026 have exactly one cycle in every module, because both were saved before cycles existed and were given one cycle each so the bar's shape would not move. One cycle is drawn as one plain bar. You settled that yourself, and the app agrees with you: each batch has at least one cycle, and a batch with one cycle is that cycle. So nothing of yours was rewritten - both lines are correct as they stand, and not one module, batch, start time or saved scenario was touched.

The cycles box no longer calls a cycle a step. It used to say "Each cycle is one step of this module", which is exactly the sentence that made you doubt the feature was there. It now says "Each cycle is one piece of this module's work", in the order you work them. The hint over the day chart says the same thing in the same word: the paler bands inside a bar are the separate cycles of that batch. Three more places where the app said step while it meant in sync now say what they mean - batch counts that differ from the module before them, and the line under the day naming any that no longer match. The batch's own card used to say "the step the day is read in"; it now says the amount the day is read in, which is what it is.

The Production line says stage, not step. That screen walks your bake day in order, and it called each part of the day a step. It now calls each of them a stage. The word job is left alone there on purpose: on that screen job already means a hand-job - one of the six jobs you do - and the proofer is not one of them, which the screen says in its own sentence. Calling the proofer a job would have put two contradicting sentences on one screen.

A batch of four now really does read as four. This is the one fault in the drawing, and it was a real one. The shades the cycles are drawn in said "a batch of four reads as four steps", but the ladder only had three shades and started again at the fourth - so on a four-cycle batch the first cycle and the fourth wore the same palest band, and the palest of the three was so faint over the pastel bars that it barely showed at all. There are four shades now, spaced evenly, and the palest has been lifted off the floor so it still reads as a wash. A fifth cycle and beyond reuse the darkest rather than falling back to the palest, so a long batch never ends on the faintest band.

What has not moved. Not one minute, position or labour band on the day chart. Your own saved line still reads 24 pans across 9 modules, the seeded starting line still 12 pans with the chiller as the wall, and your sister's still 4 pans in 8 modules. Nothing here blocks a sale and nothing here reads an order.

The tests are 1113 passing with none failing - three of them new, and one of them fails if a four-cycle batch is ever drawn with two cycles sharing a shade. No database step - nothing to run in Supabase.

22 Sep 2026 - engine v149 (no database step). The way down - taking batches off, so a day that makes more than you asked for can come back to your number.

What you found, in your own words: "This scenario makes 36 pans, and you want 24. this does not agrees?" You were right, and it was two faults in one card, not one. Your One baker day made 24 pans. You typed 36, pressed **Use these numbers**, and the ladder did what a ladder does - it raised all eight modules from 4 batches to 6 and the day made 36 pans, which is what you had just asked it for. Then you put 24 back in the box. The card underneath said "This scenario already makes your 24 pans" while the line directly above it said the day makes 36. Both sentences were drawn from the same card, one of them had to be wrong, and the ladder that could have settled it only ever went one way.

The climb only adds. Every rung of it raises something: more batches in the day, or a second machine. Not one rung has ever taken a batch away, so there was no way on that screen to bring 36 back to 24 - no amount of pressing would have done it. That half simply did not exist, and no wording change could have covered for its absence.

There is now a way down. When the day makes more than the number you asked for, the card is headed **The**

way down and carries one rung: every module the day is waiting on, and the batches to take off each of them. On the day you were looking at it reads: all eight modules holding the day at 36 pans, 6 -> 4 batches in the day at 6 pans a batch, the same change in each of them - that takes the day from 36 to 24. The same **Use these numbers** button applies it, on the heading and at the foot of the card.

Why it is one move and not a ladder. The day is the **least** any module turns out, so every module sitting on that least is holding the day there. Bring one of them down alone and the day does not move at all, because one of the others takes its place as the limit. They have to come down together, and then the day lands on your number in a single move. That is why this half is a settling move rather than a rung-by-rung climb.

The batch times you have set by hand are left exactly where you put them. This is the one place the way down deliberately differs from the climb. Raising a count makes a new rhythm, so the climb re-spaces a module's batch times from its start. **Lowering** a count only takes batches off the **end** of the day - so if you have dragged a batch to a time that suits your morning, that time stays. Nothing on this move rewrites a minute you chose.

When a batch is bigger than the number you asked for, it says so. Six pans a batch and one batch a day, and you ask for three: no count of batches gets under one batch, so there is no move to offer and no button to press. The card says it plainly instead - one batch is 6 pans, so this line can never come below 6 - because a card with nothing on it and no explanation reads as broken.

The two numbers can no longer contradict each other. The branch that reads "already makes your 24 pans" is now only reached when the day really does make 24. When the day makes more, you get the way down and the honest reading of the gap: "This scenario makes 36 pans, and you want 24. That is 12 pans more than you asked for."

What has not moved. Not one module, batch time, cycle or saved scenario. Your own saved line still reads 24 pans across 9 modules, the seeded starting line still 12 pans with the chiller as the wall, and your sister's still 4 pans in 8 modules. Nothing here reads an order and nothing here blocks a sale.

Engine 149. Tests are 1110 passing with none failing, eight of them new - four on the move itself (the 36 -> 24 settling move on all eight modules, a day already at your number offering nothing, a dragged batch time surviving the move untouched, and a batch too big to come down past) and four on the card (its heading and button, the two numbers never contradicting each other, the button actually bringing the day down, and the plain answer when a batch is too big). No database step - nothing to run in Supabase.

22 Sep 2026 - engine v148 (no database step). A cycle's labour drawn where you work it, and the ladder's button found.

A cycle's labour is now drawn where you work it, and wide enough to see. This is your own sentence from 22 September: "cycle should be drawn with its labour shown at the appropriate timing." It was being drawn at the right minute already - the fold at the end of its rest, the oven swap after its bake - but it was drawn too small to find. A one-minute fold is under two pixels at the scale the day is read in, and where it fell it landed exactly on the seam between two rests, so what you saw was the seam and not your hands. Two changes, and neither of them moves a minute:

- **The width has a floor.** A stretch of your hands inside a bar is now never drawn thinner than 4 pixels, which is the same floor the person row below already gives the same stretch of your day. So the bar and the person attending it now agree about what a one-minute job looks like - before this, the fold was 1 pixel on the bar and 4 pixels on the person row, which is two drawings of one fact disagreeing with each other. **Each end of it wears a fine light edge,** so a band that lands on the seam between two cycles still reads as a band rather than as the join.

Only the width is floored. **The position is the model's own minute**, so the picture still says WHEN: the fold sits at the end of its rest and the oven swap sits after the bake. Hover or long-press a bar and it now names the cycle and the exact minute, so the true length of a floored band is never lost - it reads, for example, "1 min of you at 4:51 am".

The button that applies the ladder is now on the ladder's own heading. You told me you could not find **Use**

these numbers. You were right, and the reason was arithmetic: it sat at the very foot of The climb card, and a day of nine modules makes a nine-rung ladder, so on a phone the button was about four screens below the heading it belongs to. It is now on the heading as well, beside the words The climb, and it is still at the foot of the card. Both press exactly the same thing.

When there is nothing to apply, there is no button - and the card says so. If your scenario already makes the number you asked for, the card reads "This scenario already makes your 24 pans" and offers nothing to press, because there is nothing for it to change. That is the other half of why the button could not be found: on a scenario that already reaches your number it does not exist at all. Raise the number you want above what the line makes, and the ladder appears with its button on the heading.

What has not moved. Not one module, batch time, cycle or saved scenario. The day still reads the same numbers it read in v147: One baker day is still 24 pans in 8 modules, one person, 3 h 52 min of hands, 4:01 am to 1:30 pm. Nothing here reads an order and nothing here blocks a sale.

Engine 148. Tests are 1102 passing with none failing, five of them new - they pin the fold's drawn minute, the bar and the person row agreeing about a one-minute job, the oven swap being drawn after the bake, and the apply button being on the heading. No database step - nothing to run in Supabase.

22 Sep 2026 - engine v147 (no database step). Module, batch, cycle - the three words on the Scenario planner, and the model that moved with them.

What you asked. In your own words: *"let us realign our terms use in the production line"* - a module instead of a brick, batches inside a module, and cycles inside a batch, with a batch's time moved by buttons instead of a drag. The words and the model under them were changed together on purpose: renaming first and building the cycles later would have left the word **cycle** meaning two different things while you were learning it.

The three words, and what each one is. A **module** is the equipment plus the hands tending it - what this app has called a brick until now. A **batch** is one lot: the six pans of one oven load, one tub of dough. A module can run several batches in a day, and how many is yours to set. A **cycle** is one step of that module's work, and it is the cycle that owns the minutes and the labour. Your own sentence settled it: *"batch is not a process but cycle is."* So a batch carries no minutes of its own - the time it holds is the sum of the cycles inside it. A batch with a single cycle is drawn on the day as one plain bar, exactly the way a module has always been drawn.

Labour now sits at its real end of a cycle: its load, or its unload, or both. Open a module and under **Cycles in one batch** every cycle has a name box, **Minutes**, **Load** and **Unload**. Load is your hands at the start of that cycle, unload is your hands at the end, and the day draws them where they really happen. A cycle cannot start until the cycle above it ends, so a cycle is never offered a start time of its own - its position is set by its minutes.

Two of your eight steps were in the wrong place, and both are corrected. Your 1-minute fold was drawn two rests early; it is now the **unload of each 30-minute rest**, so the rests run back to back inside the batch - 30 minutes, fold, 30 minutes, fold, 30 minutes, fold, 30 minutes - and the fold lands at the end of the rest, where your hands really are. And your **oven swap is a cycle of its own after the 13-minute bake** rather than labour before it, so the first swap now reads 8:13 am, thirteen minutes after the pans went in. Nothing else about the eight steps moved: the tub is still 4:01 am and the first six pans are still at the oven at 8:00 am, the minute the whole day is counted backwards from.

Moving a batch is buttons now, and the drag is gone. Tap a batch on the day chart and a small card opens with that batch's clock and its own two pairs of buttons: **Five minutes at a time** and **One minute at a time**, each with an earlier and a later button, plus the batch's own times written out. Five is the step the day is read in and the step the time line reads out; one is the nudge for lifting a batch off a collision. This replaced dragging a bar completely, and it was your call: a bar being dragged across the day and the hairline of the time cursor were fighting over the same finger. Nothing about a press does arithmetic of its own - a pressed time and a typed time give the same answer - and when the day's own rule puts the batch somewhere else the card says which rule it was. A cycle drawn inside a bar is not a separate control: tap it and you get that batch's card, because a cycle's start is its batch's start plus the cycles above it.

A batch that waits on the module above offers no buttons, and says why. Where a module has **Waits for**

the module above switched on, a batch's time is not its own - batch 3 of that module IS batch 3 of the module above - so there is nothing to step and nothing to mislead you. The card says exactly that, in as many words, and names the switch that would free the time. The row itself carries a **waits above** tag and the words **waits on** and the module it is waiting for, so a chained batch is never silent about why it is where it is.

Batch counts follow the one before, and the day says when one does not. Editing **How many batches it runs in the day** carries down to the modules after it that were running the same number, and stops dead where a module has a number of its own - it is never quietly overwritten. A standing line under the day then names each module whose batch count differs from the one before it, both numbers written out, with a plain sentence that it is a note and not a rule: the day runs on the numbers you have put in. On the seeded starting line, for instance, it reads **Wash, oil and fill - 6 batches, The fold loop - 4**. Nothing is locked by it and nothing is refused.

The tile that said Cycle time now says Minutes a pan. Same quantity it always was - the minutes one pan takes off the slowest module, with that module named under it - said in the words this app already uses elsewhere, because cycle now means a step inside a batch. No arithmetic moved with the name.

A module can have as many lines as the day needs, and lots go to the free one. How many of these do you have is that - a mixer, a sink, two ovens. At the even spacing the app works out, the lots take their turns down the lines exactly as before; where you have moved a start, a lot leaving one module now goes to whichever line is actually free rather than the one whose turn it is. No figure on an untouched day changes because of it.

One baker day was re-derived, and it is an hour later than the last version of it said. Correcting where your fold and your swap really happen costs the day a work slot, so the pace one pair of hands can hold is **87 minutes** between batches, not 81: the tubs now go in at 4:01, 5:28, 6:55 and 8:22 am, the packing starts at **8:57 am**, the last batch is packed at 1:18 pm and the day ends at **1:30 pm**, which is 9 h 29 min. The 81 no longer holds, and the reason is checkable rather than a matter of opinion: at 81 the second rest's fold lands on minute 82, the very minute the second tub's dough has to be mixed, so the app puts a second person on it. It still reads 24 pans a day in 8 modules, one person, and 3 h 52 min of hands, and nothing collides. Guide section 25 carries the whole story, including what the sweep of every other pace found.

What has not moved. Your own saved line still reads 24 pans across 9 modules and exactly the day you asked for; the seeded starting line still reads 12 pans with the chiller as the wall; your sister's still 4 pans in 8 modules. Not one saved module, batch time or scenario was altered - the words you read and the way you move a time changed, and the days themselves did not.

22 Sep 2026 - engine v146 (no database step). One baker day - your own bake day set on the bricks, with a latest start on every batch.

What you asked. In your own words: *"i need you to create one scenario and save it as One baker day. Set the bricks for me, with latest start time each brick batch."*

Where it is. More -> Scenario planner. Under **Your scenarios** there is one new row, **+ One baker day**, reading **One pair of hands, set for you**. One tap puts it on your shelf and opens it. From then on it is an ordinary saved scenario, exactly like your own line and your sister's proposal: it stays in your list, it opens when you tap it, and the row that offered it is gone because you already have it. Delete it the way you delete any scenario, and it stays deleted - it is never dropped back in on you.

The eight bricks are your eight steps, in your own order. Mixing the dough in the tub; the rests and the stretch and folds; oiling the pans and weighing the dough out; into the proofer; dimple and top; the proofer again; the oven swap and the bake; cutting and packing. Nothing invented, nothing left out. The fold and the two proofer bricks carry no line step against them, because the Production line has no field for them - and the screen names them rather than writing their minutes onto a field that means something else.

Every brick's first cycle is a latest start, which is what you asked for. Open the day and the eight bricks read:

- Mixing the dough in the tub - **4:01 am** (20 minutes, all of it yours) The rests and the stretch and folds - **4:21**

am (123 minutes, 3 of them your hands) Oil the pans and weigh the dough out - **6:24 am** (15 minutes of you) Into the proofer - **6:39 am** (45 minutes, no hands) Dimple and top - **7:24 am** (6 minutes of you) The proofer again - **7:30 am** (30 minutes, no hands) The oven swap and the bake - **8:00 am**, the anchor: the minute the first 6 pans have to be at the oven (2 minutes of you) Cutting and packing - **8:27 am** (12 minutes of you)

The four batches sit **81 minutes apart**, so the second tub starts at 5:22 am, the third at 6:43 am and the fourth at 8:04 am. The last batch of the day is packed at 12:30 pm and the day is done at 12:42 pm.

One of the eight is not where your chain puts it, and that is the whole finding of the scenario. Your own chain would start the cutting and packing at **8:15 am**, fifteen minutes after the oven. Alone, you are still folding the last tub at 8:27, so **8:27 is the first minute that job needs nobody else**. It is not a preference: start it a single minute earlier, at any minute from 8:26 backwards, and the day asks for a **second pair of hands**. The window you can move it in without buying anybody is **8:27 to 8:33 am**, and the seeded day sits at the start of it.

Why the batches are 81 minutes apart and not 40.5. This is your own proofer figure, from the same morning: a batch is in the cabinet 81 minutes - 45 in, 6 out to be dimpled, 30 back in - so the cabinet could take a new batch every **40.5 minutes**, about 8.9 pans an hour. But 40.5 assumes somebody is free to carry a batch out and the next one in the moment the cabinet empties, and with one pair of hands there is no such minute: the hand-work of a single batch is 58 minutes by itself. Rather than guess, every rhythm from 40.5 upwards was swept through the app's own arithmetic, and **81 minutes is the pace one pair of hands can hold**: every hand-window in the day then tiles exactly, each one beginning where the last one ends, and the proofer is never asked to hold two batches at once. It still works at 80.5 minutes; at 80, or at 40.5, the very same day reads **2 people**. So the two figures you have are both right and they answer different questions: the card's 40.5 is what the cabinet allows when you are not alone, and this scenario's 81 is what one baker can actually feed it.

What the day reads. 24 pans a day in 8 bricks - four batches of the six pans you bake. **One person**, **3 h 52 min of hands** across the whole day, and a run from **4:01 am to 12:42 pm**, which is 8 h 41 min. There are **no red notes** under the day, because nothing collides: one pair of hands, no job doubled up. The brick setting the pace is the **mixing**, and seven others pass the same 24 pans - so the card says they are one limit between them, and relieving only one of them will not move the day.

It loads onto the Production line as your own numbers. Under **Start from a scenario** at More -> Production line, tapping **One baker day** shows you everything it would change before anything moves: minutes to mix one tub 20, oiling the pans and weighing the dough out 15, dimple and top 6, take 6 pans out and put 6 in 2, minutes the pans are in the oven 13, pans in one load 6, cutting and packing 12, the target 24 pans, and 1 person. The fold and the two proofer bricks are listed by name, because that screen has no field for them. Nothing is written until you press the button.

Every number is yours to move. The day is seeded, not frozen: drag a bar, open a brick and type a number, change how many pans a day you want, or move the time your day starts, and the whole thing answers again the way every other scenario does.

What has not moved. Your own saved line still reads **24 pans across 9 bricks** and exactly the day you asked for, the seeded starting line still reads **12 pans with the chiller as the wall**, and your sister's still **4 pans in 8 bricks**. One baker day arrives as a row you choose: not one brick, start time, merge or saved scenario of yours changed, and nothing here blocks a sale.

CHANGELOG and the guide (section 25) both carry all of this, both PDFs rebuilt, and the tests are 1090 passing with none failing. No database step - nothing to run in Supabase.

22 Sep 2026 - engine v145 (no database step). Your real bake day on the Production line - and the same day read backwards from the oven, so no dough is ever mixed too early.

What you asked. In your own words: *"work backwards, from end of process, the previous process whould have a latest start time, by going this way, we prevent preparing dough too early and prevent dough from over fermented. But we can fine tune each process to start earlier to better manage manpower utilisation."* You had corrected your own step order that morning, after checking it with your sister: *"i think my self confused, i double check with my sister and now understood the right flow. we cannot mix the dough too early, it will over*

fermented." And you had already worked your ceiling out yourself and asked about it: *"6pan sit in it for 81min, proofer able to hold 12pans, so effectively the poofer can produde 6 pan/40.5min? So the proofer became the bottle neck?"* Yes - 40.5 minutes exactly. This release gets that figure onto the screen out of your own numbers rather than out of a guess.

The flow on the Production line was not your day. It read mix, then the chiller, then washing and oiling the pans, then weighing the dough out, then the oven. Against your day that put the chiller where you have none, transposed the oiling and the weighing out, and - the worst of it - had no fold in it at all. The map now runs your order, top to bottom: mix the dough in the tub; the rests and the stretch and folds; oil the pans and weigh the dough out; into the proofer; dimple and top; proofer again; the oven; then cutting and packing. The proofer is drawn twice because your day uses it twice, with the dimple in between.

The wall was a chiller you have not bought. The screen watched a chiller of 12 trays over 5 hours and made that your ceiling - 12 pans a day - which is a cap on a cabinet that is still only a what-if on your planner, kept switched off. The station you actually own is the proofer, and it is the honest ceiling: a batch sits in the cabinet for the whole 81 minutes (45, then the dimple, then 30), so 12 pans means one batch every 40.5 minutes and no faster. Change **Pans your proofer holds at once** to 18 and the screen answers 27 minutes instead, and says which of your limits is the wall now. It never tells you to buy one: at one pair of hands a bigger cabinet buys you nothing, and the screen says so.

Your day, read backwards. A new card beside the flow, **Your day, backwards**, gives every step the last moment it may start, worked back from the oven - the end of the process, which is where you asked to start from. For the first six pans to be standing at the oven at 8:00 am, the dough goes into the tub at 4:01 am:

- Mix the dough in the tub - 4:01 am The rests and the stretch and folds - 4:21 am Oil the pans and weigh the dough out - 6:24 am Into the proofer - 6:39 am Dimple and top - 7:24 am Proofer again - 7:30 am The oven, bake and swap - 8:00 am

That is **239 minutes** from the tub to the oven and **254** from the tub to the pans coming out, which is your own arithmetic and nothing else. One batch is **46 minutes of your hands**: 20 mixing, 3 folding, 15 oiling and weighing, 6 dimpling, 2 swapping.

The clock is a field, not a rule. The minute the first batch must be standing at the oven is yours to set - it arrives at 8:00 am - and every start time below is worked back from it. Move it and the whole day moves with it.

Five minutes early is allowed, and nothing is locked. Every step also carries a five minute band above it: a step may begin up to five minutes early to suit your hands, which is the room you asked for to manage the people. It is advice drawn on the screen and it is never a gate - no time here stops you baking, hides an order, or overrules your own eye. You watch the dough; the timetable watches the clock.

Five numbers changed what they were asked, and each one is named here rather than changed quietly. Your saved numbers move across by themselves the first time your phone opens this version, so you will see these the moment you load it - and every one of them can be typed back in one tap if your own timing says otherwise:

- **Minutes to mix one tub of dough**: was 6, now 20. The old box asked how long it took you to load a mix; this one asks how long the whole mix takes, which you gave as twenty minutes for one tub. **Pans one tub of dough makes**: was 28, now 6. The old box asked what your mixer's bowl holds. Your cycle is one tub of dough into the pans of one oven load - six - and the mixing and folding are spread over that tub, so this is the number the chain needs. **Oil the pans and weigh the dough out, for 6 pans**: was two jobs under two names (your washing and oiling read 20, and the weighing out was never timed), now one job at 15. Your own count of the hand-work has it once, which is the proof they were one job: with the two merged, the app's 46 minutes a batch is your 46 to the minute. **Minutes to dimple and top 6 pans**: was 8, now 6 - your one minute a pan. The 8 was measured with the topping in it. **Minutes to take 6 pans out and put 6 in**: was 4, now 2 - *"2min covers out and in"*, as you put it. One turn of the oven is the baking and the swapping together, 15 minutes as before.

Everything else of yours is untouched: your two pairs of hands, your 5 hours, your 24-pan day, your 12 pans, 6 pans a bake, 15 minutes a turn, the 30-minute cooling - and your 12 minutes of cutting and packing, which you

had timed and which is still yours.

A step with no minutes on it is still named, never counted as free. Cutting and packing is the one job that arrives with an empty box - every other step of the day comes seeded with your own figure - so on a fresh phone the flow says **not timed yet** and the backwards card says in a sentence that it is not counted in the 254 minutes, rather than quietly treating your time at the table as free. You had already timed yours at 12 minutes for 6 pans, so on your phone that box holds your 12 and the card's heading reads **After the bake** rather than **Still to time**. A heading left reading **Still to time** over a figure you had given would read as the screen not having noticed it. Either way the step is drawn on the backwards card as sitting **after the bake, outside the 254 minutes**, because those minutes happen while you are waiting on the next batch. Nothing was reset.

One thing on the join between the two screens reads differently, and it is a correction. When you tap a saved scenario under **Start from a scenario** on the Production line, the list of numbers it would change is now the corrected list: **oiling the pans and weighing the dough out** is one row where it used to be two, because it is one job under two names, and a **retard** brick is named as something the scenario said but the line has no field for rather than quietly writing a tray count onto a field that no longer exists. The bricks themselves are untouched, and nothing is written until you press the button. What did **not** ship is a button on the Scenario planner to push a planned day's own start times onto the bricks: your bricks are timed from a day that starts at 3 pm, and the bake day here is a morning that has to be ready by 8 am, so writing those times in would have moved every saved brick. The timetable is drawn, and it is not pushed anywhere yet.

What has not moved. The Scenario planner, brick for brick. Your own line still reads 24 pans across 9 bricks and your sister's still 4 pans in 8 bricks, and both are stored exactly as they were - not a brick, a start time or a merge has changed. The chiller bricks stay where they are, switched off, because the fridge is still worth planning for.

CHANGELOG and the guide (section 24) both carry all of this, both PDFs rebuilt, and the tests are 1084 passing with none failing. No database step - nothing to run in Supabase.

21 Sep 2026 - engine v144 (no database step). The People row counts the places one person is in, a line of a brick counting as a place - and a collision is written so it can be read.

What you asked. In your own words: *"fix the hats count and the collision sentence. And what is the concept of hat, i dont understand, why name them by numbers?"*. Both were real faults, and the question was fair - "hat" was my word, and it was a poor one. It is gone from the app and from this book; what it meant is said plainly below.

What a "hat" was. It was my shorthand for **the places one person has to be in the day**. You put it better yourself: *"one person would not need to wear too many hats"* - one worker asked to stand at six different jobs in one day is being handed six hats. That is the thing the People row is there to count, so that is what it now says: **in 6 places**.

The count was wrong. Each person's row read **at 6 bricks** - the number of different bricks, not the number of places - so a worker on **line 1 and line 2 of the same fold** counted as being in **one** place. That is exactly the doubling-up the number is supposed to show you, and it was the one thing the row said nothing about. A line of a brick is a place of its own now, so that same worker reads **in 2 places**. On your own line the row has gone from **at 6 bricks** to **in 7 places**, because your fold is drawn as two lines. The number only appears when it is more than one, exactly as before.

The collision sentence was unreadable. It was one long sentence - *Person 1 is at X and Y at the same time, ...* - and your bricks have long names: *Mixing by hand in the tub (set the minutes)* is forty characters, twice in one sentence, with the word *and* between them lost in the middle. It is now three parts on four separate lines. Line one, who and when: *Person 1 - two jobs at once, 4:00 pm to 4:05 pm*. Lines two and three, each job on a line of its own with a bullet in front of it: *Mixing by hand in the tub (set the minutes), line 1* and *Mixing by hand in the tub (set the minutes), line 2*. Line four, what to do about it: *Move one of them along the day, or combine with another person and accept the collision.* Two lines of one brick still say which line they are, because without that they read identically. And if a day has more than three of these, the last block **counts the rest** instead of quietly

stopping - a list that stopped at four read as though the day had four problems when it had nine.

Why the people are numbers. Because a number is what the app stores: who is on a brick is a number, combining two people is joining two numbers, and the same person on three bricks is the same number in three places. That is what makes "put person 1 on both lines" work in one tap. It also has nothing to do with you - **Person 1, 2, 3** are places in the day, not names of people, and the app never shows them to a customer. Naming them (Mei, Kak, Adik) is a small change on top and I have not built it - say the word if it would help you plan.

What has not moved. The day itself. Not one minute, one pan or one figure. Your own line still reads **24 pans across 9 bricks**, the starting line still **12 pans with the chiller as the wall**, your sister's still **4 pans in 8 bricks**. Only the words on one row and the shape of one red note changed.

CHANGELOG and the guide (section 23) both carry all of this, both PDFs rebuilt, and the tests are 1062 passing with none failing. No database step - nothing to run in Supabase.

21 Sep 2026 - engine v143 (no database step). A time cursor on the day chart: point at the chart and a line follows your pointer down the whole day, reading out the clock, so you can line two bricks up against each other.

What you asked. In your own words: *"i need a time x axis cursor on the time chart"*.

Where it is. The day chart in **More -> Scenario planner**. One change, in one screen.

How to use it. Run the pointer across the chart and a fine line follows it, down through every row - the ruler, every brick, every person, the total. A little brown label at the top of the line reads the clock at that point: *5:05 pm*. Move the pointer and the reading moves with it; move the pointer off the chart and it goes away.

On a phone. A finger cannot hover, so on a phone the **clock strip is the handle** - the one band along the top of the chart where the hours are written. Put your finger on that strip and slide it left or right and the line follows, reading out the time. Everywhere else on the chart your finger still scrolls the day the way it always did, so you have not lost the swipe.

Why it only counts in fives. The line reads the clock to the same five minutes a bar snaps to when you drag it. That is deliberate: the cursor can never name a time the chart could not also set, so you are never shown a time you cannot ask for.

At the ends of the day. Point before the day starts or past the end of it and the line goes away rather than naming a time that is not in your day. At the far end of the evening the reading hangs to the left of the line instead of the right, so it is always written out in full and never half cut off.

Two bricks lined up. This is what the line is for. Point at where a brick starts, read the time off the label, then point at where the next one starts and read that - no counting squares. It reads the chart and it writes nothing: no brick, no time and no number changes because you pointed at something.

Nothing you have has moved. Not one figure. Your own line still reads 24 pans across 9 bricks, the starting line still 12 pans with the chiller as the wall, and your sister's still 4 pans in 8 bricks.

CHANGELOG and the guide (section 23) both carry all of this, both PDFs rebuilt, and the tests are 1061 passing with none failing. No database step - nothing to run in Supabase.

21 Sep 2026 - engine v142 (no database step). Typing a number into a long card no longer throws the card - the box keeps your cursor and the card stays where you were reading.

What you asked. In your own words: *"why the display jump and it is very confusing"*. It was a real fault, and this is it fixed.

What was actually happening. The brick's card is a long card - fifteen boxes, a box per line, a row for every cycle - and **every single number you typed into it was rebuilding the whole card underneath your finger**. That is the jump you saw: the box you were typing in was thrown away and a brand new one put in its place, so the cursor went with it and the keypad shut. On a phone that reads as one digit per tap, on a card that had moved somewhere else. It was worst on the three boxes that change how tall the card is - typing a 6 into *How*

many of these do you have grew the card by a line's worth of boxes on every keystroke.

What it does now. Typing a number writes the number and **leaves the card alone**. The box keeps your cursor, the keypad stays open, and the card stays exactly where you had scrolled to - checked box by box: the brick's own numbers, the people box, and every line box under *How many of these do you have*. Three boxes still change the card's shape, because their answer really does add a row or move a time - **How many times it runs** (one more row per cycle), **How many of these you have** (one more line box), and **Minutes in, when its first pass starts** (cycle 1's own time). Those three now redraw **once, when you leave the box**, instead of once per digit.

And it is not only the brick's card. The jumping was in the shared pop-up that every screen in the app opens, so the same fix covers all of them: any card long enough to scroll keeps its place while you work down it, and a card you open fresh still starts at the top where it should.

Nothing you have is affected. This changes only when the screen redraws, never what it redraws: every figure, every scenario and every saved line reads exactly as it did in v141. Nothing new to run in the database.

21 Sep 2026 - engine v141 (no database step). A brick you have two of is now drawn as two lines, and each line carries its own person - so one worker stands at one job instead of wearing every hat in the day.

What you asked. In your own words: *"i want is just better drawing.. Since now one brick having 2 lines, each line can have its own person, and few line can have a combine person sharing to different lines. This will make training of worker easier, because one person would not need to wear to many hats"*. So this is the **drawing**, and **who is on each line** - nothing about how the day adds up has moved.

A line, and why a brick you have two of is two of them. Everything on this screen was already worked out as if a brick you have two of were two streams taking the lots in turn - lot one and lot three on one, lot two and lot four on the other. The screen was drawing those two streams flattened into a single row. Now each one is a row of its own: the brick's own name, badges and minutes on the first line, and a row per further line underneath, tied together by an accent down the left and a dotted seam between them.

Who is on each line is yours to set. Open the brick and, once *How many of these do you have* says 2 or more, you get a box per line - **Line 1 - who is on it**, **Line 2 - who is on it**. 0 means whoever is free. Put the same number on two lines and that is one person covering both, which is exactly what you asked for; if those two lines really do need them in the same minute, their row in the People list goes red and names the minute and which line collided with which. The line boxes and the brick's own *Who is at this brick* are the same number, so the two can never disagree.

The People list now counts the places one person has to be. Each person's row reads how many different places that one person is in - and every job on it names its line ("Stretch and fold in the tub, line 2"). That number is the whole point of this release: it is the one that tells you when one person is being asked to be in too many places at once.

*(This entry was first written as "counts hats", as **at 3 bricks**. Both the word and the count were wrong, and **v144** fixes them: the row now reads **in 3 places**, and a line of a brick counts as a place. See the v144 entry at the top.)*

Your starting point in one tap. One a line (it was *One a brick*) gives every line its own person, so a brick you have two of takes two people rather than one. It says *One a line* only when a brick in the scenario really is drawn as lines; otherwise it reads exactly as it always did. Its name and the message it leaves say which of the two it just did.

What has not moved. With no lines drawn and nobody set - which is every brick you have - the day is computed exactly as before, and every figure you have already been given reads identically: your own scenario still 24 pans and exactly the day, the seeded line still 12 pans with the chiller as its wall, your sister's still 4 pans in 8 bricks. Two of a brick still buys cycles out of the pans you planned; it does not bake pans you did not plan.

One thing it deliberately does not do. With **Let its cycles overlap** switched on, a brick's lots no longer take

turns, so there are no lines to draw and the brick goes back to one row with one person. The people you set per line are **kept**, not thrown away, and the brick's card says so - switch overlapping back off and your lines come back, each with the person you gave it.

Nothing to run in the database. This is a screen and a drawing; no SQL, no new table.

21 Sep 2026 - engine v140 (no database step). A brick can now be told that its cycles may overlap - one switch, with the rule for when to use it written beside it - and the two lots are drawn as two bars.

What you asked. In your own words: *"i also want to allow each brick cycle to overlap, this will improve the visualisation"*, and then, as I was building it: *"allow overlap button and the overlapping criteria"*. So it is a **button** in the brick's card, **Let its cycles overlap**, and the criteria sit under it in plain words rather than in a rule you would have to work out.

What a brick did before, and still does unless you switch this on. A brick held its own cycles apart: lot two could not begin until lot one was out of it. That is the right rule when the dough is physically *in* the thing, and it is what stops you folding dough that is still in the mixer. Where you switch it on: when the minutes in a pass are the **dough's** time and not a machine's - dough resting between folds, a second bin on the go. Then every cycle sits exactly where you put it.

How you can tell it is on. The brick's row wears a **2 at once** tag naming how many lots are in it together, and the row grows taller so the overlapping lots are drawn as separate bars in their own lanes instead of painting over each other. A brick that is not overlapping looks exactly as it did - same bar, same height, same place.

Three things it deliberately does not do. It never lets a lot start before the dough exists: a brick that waits for the brick above still waits, so switching this on is not a way around the chain. It does not put pans in the day that you did not plan. And it does not quietly double you up - if two of your own cycles need the same person in the same minute, that person's row still goes red and names the minute.

And it is not the same as a second brick. How many of these do you have means one lot in per brick you have - with 2 of them, lot three may start while lot one is still folding, but lot four waits for lot two. **Let its cycles overlap** means every lot at once. The switch is the wider of the two.

One correction to v139, written down so it never looks like a bug. v139 said a brick whose passes come round faster than one pass takes "now has its passes held apart instead of drawn overlapping - that day could never be worked". That is still what happens **with the switch off**, which is how every brick arrives. With it on, those passes are drawn overlapping, in lanes, which is the whole point of the switch.

Nothing else in the app moved, and both of your scenarios read exactly as they did - the sister's line still 4 pans in 8 bricks. Tests are 1048 passing with none failing, and the guide (section 23) carries all of it.

21 Sep 2026 - engine v139 (no database step). The scenario planner now builds a real line: a brick can be two of them, every pass in a brick has its own start time you can drag or type, and a brick can be set to wait for the brick above it - lot 10 waits for lot 10.

What you asked. In your own list: *"1. allow me to set the same brick to 2 or more, if that module became bottle neck."*, *"2. in a brick, let say we might need 8 cycles, each cycle should be able to have its own start time, drag-able."*, *"3. 2nd brick 1st cycle should follow 1st brick 1st cycle completion, and so on."*, *"4. Each brick, cycle10 only starts after one brick earlier cycle 10 ends, etc."*, and *"5. objectives are: less man, less manhours thru better planning, shorter hours."* Then, as I was building the screen: *"from your layout, the stretch and fold should have 2nd , 3rd brick as the person who do that process and start another cycle without having to wait for thee 1st cycle to finish"* - which is what **How many of these do you have** does: a second fold starts the next lot without waiting for the first to finish.

What changed. Everything below is inside More -> Scenario planner. Nothing else in the app moved, and every figure you have already given reads exactly as it did.

Every pass in a brick is drawn on its own, and dragging one moves only that one. Before this version a brick's passes were a single rhythm - a gap and a count - and every pass sat where that rhythm put it. Now each pass is **its own bar on the timeline** and can be dragged by itself: cycle 3 moves and cycles 1, 2 and 4

stay where they were. The toast when you let go names the one you moved - *"Cycle 3 -> 4:15 pm"* - so there is no doubt which pass changed. For a small nudge, a number beats a careful finger: the brick's card now has **a row per cycle**, each with a box for the exact minute and the clock reading beside it, so a five-minute correction is typed rather than dragged. And **Space them evenly again** puts every pass back on the even rhythm from the first one, once you are done experimenting.

A brick can be two of them - or more. A new box in the brick's card, **How many of these do you have**, and the row wears a **2 of them** tag when you raise it. What a second one really buys you, and what it does not: it lets the brick hold **twice as many passes in a day**, and it runs at **twice the pans an hour** - but it does **not** change the pans you planned, because a second mixer standing idle does not bake anything. To use it you raise the passes. That is the honest reading of buying one: it removes a wait, it does not invent work. With **2 of them**, that brick's own passes are allowed to overlap, because two machines can run at once - which is the whole point of having two.

A brick can wait for the brick above it - and that is the chain you described. A switch in the brick's card called **Waits for the brick above**, and it does exactly your points 3 and 4: lot 1 of this brick cannot start until lot 1 of the brick above has finished, lot 2 waits for lot 2, and so on down the line. Turn it on for the bricks that really are one line and a slow pass upstream holds every later lot behind it - which is precisely why a second fold, or a second mixer, is worth buying. Two things worth knowing. **A chained bar cannot be dragged**, because where it sits is decided by the brick above: those bars are drawn faded and wear a **waits above** tag, so nothing on the screen looks like a control that does nothing. To move a chained lot you move the cycle it is waiting on, and the whole line moves with it - which is the lesson the screen is for. And **if two bricks have different numbers of passes**, the screen says so instead of guessing: passes past the end of the brick above follow its *last* pass, and the row names it. **The chain arrives switched off**, so nothing you have already measured moves until you turn it on yourself.

The line, in one tap - and your three objectives as three numbers. You named the goal as three things, so the screen now says all three in one line under the day: **2 people · 13 h 40 m of hands · a 16 h 20 m day**. **People** is the most hands needed at once (*less man*), **of hands** is every person-minute of the day added together (*less manhours*), and **the day** is first job to last (*shorter hours*). Drag a cycle, move a start, add a second brick, and all three answer again at once - so a change either improves all three or you can see which one it cost. Beside the timeline a chip reads **3 waits above**, or **Not chained**: tap it and it explains the rule, names which bricks are waiting, shows those same three numbers, and offers the whole thing in one tap - **Chain the whole line** or **Take the waiting off** - so you never set eight switches by hand.

The climb can now offer a second one, not only more passes. The ladder used to relieve the wall by raising that brick's passes, and stopped where a day could not hold more of them. That stopping point is exactly where a second brick is the real answer, so a rung can now read **2 of them** as well as **more passes**, with the same before-and-after pans figures. A rung that adds a second one says what it costs in as many words - **that is a second one to buy** - because that rung is money where the others are time. A brick that needs both is offered both, in order.

One arithmetic change worth writing down, so it never looks like a bug. A brick whose passes come round *faster* than one pass takes - a 5-minute gap on a 28-minute fold - was always a drawing of something that cannot happen: the same hands cannot start pass 2 while pass 1 is still on them. Those passes are now held apart to one at a time, when the brick is the only one of it and is not chained. None of your own bricks is built that way and every figure you have given reads exactly as it did, but if you ever type a gap shorter than a pass, the screen will hold the passes apart rather than draw a day you cannot work.

21 Sep 2026 - engine v138 (no database step). The scenario planner is yours to build in: bricks you make, move, give to a person and drag along the day, a ruler and a start time, scenarios you save and name - and the scenario you choose can now be loaded onto the Production line.

What you asked. In your own list: *"1st i cannot change the time scale, say i want a strat at 8am. And i need an time scale ruler"*, *"I need to move the sequence of the lego brick"*, *"I want to be able to create, edit and delete lego brick"*, *"i want to be able to same and name a scenario. Edit or delete"*, *"how to have person 2?, I would*

want one person for each module, And a Total person showing overlapping persons time slot if any", *"The lego brick should be able to restart? if choosen to, and how, can it be drag tomove the start time of each brick"*, and *"The fold brick should show 28min, and its person should show 2 minutes, when i adjust the brick until minumum total person overlap, that will be best scenario for manpower planning."* Then, as I was building it: *"can allow to save a scenario? scenario and production line has some function overlapping, the chooses scenario should be able to be loaded to production line"*.

What changed. Everything below is inside More -> Scenario planner, plus one card on More -> Production line. Nothing else in the app moved.

The day starts when you say, and there is a ruler to set it by. One box at the top sets **the time your day begins**, and every clock reading on the screen is counted from it - so *"say i want a strat at 8am"* is one field rather than ten start times to retype. A ruler now runs under the timeline: a tick every half hour with the o'clock ones named, so a bar can be read against the clock instead of guessed at. And the **time scale** is yours: four chips - **Wide**, **Standard**, **Close**, **Closest** - change how wide a minute is drawn, from the whole day on one phone screen to close enough to tell a 15-minute bake from the one before it. The hour lines on the timeline follow whichever you pick, so grid and ruler always agree.

A brick is now a thing you own, not a row you read. + New brick adds one, and it arrives **switched off** so an empty brick cannot become the wall and answer your day with a nonsense figure. Open any brick and you can rename it, change its picture, change every number in it, and - this is new - **delete it** or move it **up Earlier / down Later**, because the list order *is* the day's order: the timeline is read top to bottom and a tie for the wall goes to the earlier brick. So *"move the sequence of the lego brick"* is two buttons, and the day re-sorts around it.

One person for each module, and the Total person you described. Every brick can be given to a person: **0** means whoever is free, **1, 2, 3...** names someone. Two chips do it in one tap - **One a brick** gives each job its own person, **Share them** hands them all back to whoever is free. Under the timeline there is now **one row per person**, each showing the jobs that person is at and how much work they carry, and where two of their jobs overlap the bar wears a red outline with a sentence naming **both jobs and the minute they clash**. Nothing blocks a plan that cannot work - the whole point is that it is *shown*. The **Total person** row stacks person 1, person 2, person 3 and the rest on top of each other, which is what you asked for: *"with that i can see, which time period utilise more manpower"*. It reads out the most people needed at once, the busiest stretch by name, and how many minutes are being worked two at a time. Bring the bricks closer together and that row flattens - that is the manpower planning.

Combining two people, which was your own example. Combine two... asks two questions: the person who keeps the job, then whose bricks move to them. Person 1 then person 3 makes the row read **Person 1+3** - *"to generate person 1_3"*, in your words - and whatever then collides is exactly what that combination cannot pay for. Give one brick back in its own editor and the **1+3** label prunes itself the moment person 3 has a brick again, so a label can never claim somebody who is not really there.

A brick restarts, and a brick can be dragged. Same number, two ways to move it. Type the gap and the number of runs and the brick restarts itself later in the day - your fold, four times. Or take the bar on the timeline with your finger and **drag it**: the start moves in 5-minute steps, the row's times update as you drag, and letting go will not open the brick's card. Drag to move, tap to open. Your pans do not change when you drag; only your people do.

Your fold brick now reads the way you described it. It **holds for 28 minutes**, it is **2 minutes of you**, and it restarts 30 minutes after the last fold, four times. One honest note: if your phone's fold brick still reads **30** in *minutes one pass holds it*, that is the copy your phone saved from the first version of this screen. Type **28** into that box and it is right. I have not overwritten it for you, because your numbers are never changed by an update - that is a rule of this app, not an oversight.

Save a scenario, name it, open it, rename it, delete it. Save changes keeps the one you are in; **Save a copy...** makes a second named one and carries on in it - which is how you keep the line you have and the line you are thinking of buying side by side. **Your scenarios** lists everything saved: tap to open (the row you are in

wears a filled dot), or tap the **pencil** to rename or delete it. Deleting a saved scenario never touches the bricks you are working on.

Your sister's line is in there too, built from her own words. Next to your saved scenarios there is one more row - **+ My sister proposal 21/9/2026** - and one tap puts it on the shelf and opens it. It is the line she wrote out, and it has **no mixer and no chiller in it**: everything mixed by hand in one big tub, folded in that same tub for 2 minutes every 30 minutes three times, tipped into four pans, an hour in a proofing cabinet, then dimple and oil in the pan, half an hour of rest, the oven, and cooling and packing. As she described it that is **4 pans a day in 8 bricks**, and the thing stopping it is **mixing by hand in the tub** - not the oven. Bring the bricks closer together and the Total person row will show you where the three folds collide with filling the pans.

Two of its numbers are mine, not hers, and both say so on the brick. She never gave a mixing time - she mixes by hand, so there is no machine time to give - and I have seeded it with the length of your own mixer pass so the brick is not blank. The **Into the pans** brick is worked out from **your** measured 3 minutes a pan, which is 12 minutes for four. Everything else is her own numbers, and every number on every brick is editable. The tub mixes **4 pans in one go**, so if her tub is bigger, change the batch on that brick and the whole day re-reads itself.

One repair under this. Opening a saved scenario used to leave the box at the top of the screen still reading the scenario you were in *before* it - the name, the start time and the pans you want. It now repaints with the scenario you actually opened, so the top of the screen and the row you just tapped always agree.

And the chosen scenario loads onto the Production line. A new card at the top of More -> **Production line** - **Start from a scenario** - lists everything you have saved. Tapping one shows you **every number it would change before anything moves**: what it was, what it would become, and which numbers it would leave exactly as they are. So that the two screens can agree without retyping, every brick's card now carries one more line - **which step of your line is this** - where you say whether it is the mixer, the wash, weighing the dough out, the topping, the oven, the cooling or the retard. A brick's **minutes** and its **batch** are what cross over; a brick that is not a step of the line (your fold, loading the chiller) is named as left out rather than silently dropped. **How many pans you own** and **how many hours you will bake for** stay exactly as they are, because a scenario does not know them and the screen says so instead of filling them with a nought. Nothing is written until you press the button, and the button counts the numbers that would really move - so a scenario you have already loaded says *"nothing here changes"* rather than promising eleven.

21 Sep 2026 - engine v137 (no database step). More -> Scenario planner: build your line out of modules, watch the day as a timing diagram, and raise the pans you want until something stops you.

What you asked. "can i have a apps function that visualise the flow... something for me to optimise line balancing", and then, once we were talking about it: a screen you could **configure**, a **timing diagram**, and - this was the important one - *"not optimised for mobile, for a clearer interface"*, so you could "toase about all possibility to configure the production flows". You described the parts yourself: a module around the oven, a module around the mixer, a module around a fermentation, each with its **own cycle time**; the folds as a **loop**; ingredients going in "from time to time, not one shot"; and the fridge and the bins as **half-finished goods** sitting between steps. And you set the rule for the whole exercise: *"alll this are design, and require proofing, and before the proofing, the fugures have to be right first, at least logically."*

What changed. A new screen, **More -> Scenario planner**. It is the one screen in the app that is **wider than the rest** - up to 900px instead of 540 - because a day laid out across a timeline needs the room. Everything else about the app is exactly as it was.

A module is a Lego brick: a machine *and* the pair of hands that tends it, counted as one thing. Each module has its own cycle time - how long one pass holds it - how many pans that pass deals with, how many minutes of you it takes, and how often it repeats. Your fold loop is just a module that repeats more often, so a loop needs no special anything. Ten modules come seeded: the mixer, the fold loop, wash/oil/fill, load the chiller, retard overnight, unload, dimple and top, the oven, cool and pack - and the fridge, which arrives **switched off**.

Scenario 1 is the line without a fridge, and the fridge is a switch, not an absence. It sits in the module list

switched off and reads *"not in this scenario"*, so the screen answers without it. Switch it on and it appears. That is deliberate, because you said it plainly: *"if the output fulfilied my need now, then fridge would not be the 1st thing to buy for me"* - so the buying question and the design question had to be the same screen, one tap apart.

The day is a timing diagram. One row per module running across the day, with the clock along the top every two hours. A **solid block is you standing there**; a pale one is the module running without you - the retard and the bake have no solid block at all because they need nobody. Below the modules is a row **per person**, carrying the jobs that person is at. There is one person on your seeded day, doing 4 h 56 min of work. Swipe sideways; the day is wider than the phone.

The climb - this is the part you asked for. You described the method yourself: *"i will slowly increase the output needed per day until 1st limitation hit, tweaking the 1st limitation and expecting the output to grows, until the next limitation surface"*. So: type how many pans a day you want, and a ladder appears. Each rung names the one module that stops you at that number, exactly what to change about that module, what the day becomes, and which module becomes the wall next. **One tap puts the whole ladder into the modules.**

And the ladder has to be honest, which is where it earns its keep. Your seeded day makes **12 pans** and it is the chiller that stops it. Raise the target to 36 and the ladder climbs to **24** - then stops, and says why in plain words: *"A 12 h pass fits 2 times in a day, and it is already running that often. So this one cannot be run more often - it has to take more pans at once."* Twelve trays, twice a day, is twenty-four pans and no amount of wanting makes it more. **Thirty-six pans needs eighteen trays.** It is the same trays-to-a-day arithmetic from the Production line, now holding this screen to account - and it means this screen can talk you **out** of a purchase as well as into one.

Every figure is yours to move. Tap any row - or any brick in the "Not in this scenario" list - and one card opens with that module's numbers, each label naming its own unit: minutes one pass holds it, pans in one pass, minutes of you, minutes from one pass to the next, passes in the day, minutes in when the first pass starts, people this pass needs. Nothing here reads an order and nothing here blocks a sale; it is a planner you type into.

One more honest note. A day can only hold so many passes of something: a 12-hour retard fits twice in 24 hours, and the screen will not count a third however many you type. Your number is kept exactly as you typed it - the row just wears an *"a day's limit"* tag and counts the passes that are really there.

20 Sep 2026 - engine v136 (no database step). The Production line now opens with a flow map: your bake day top to bottom, in the order it actually happens.

What you asked. "can you do a flow map so that i can visualise better, make a intuitive one". The four bars answered *how fast* each part is, but nothing on the screen showed the **order** - that the dough is mixed first, rests overnight, and only then goes pan by pan through the oven and back round again. Numbers side by side don't tell you what happens next.

What changed. A new card called **The flow** now sits at the top of the screen, above the bars. It is the bake day written down the page, one step under the next with a numbered dot and a line running between them, so your eye follows it like a recipe:

Weighing in and loading the mixer -> Retard overnight in the chiller -> Wash, oil and fill -> Weighing the dough out into pans -> Dimple and top -> The oven swap and bake -> Cooling and packing -> The day.

Each step carries its own cost: the hand steps say **minutes a pan**, the chiller, the oven and the pans going round say **pans an hour**, and a step you have not timed says **"not timed yet"** rather than quietly reading as free. The last node states the day the line actually makes. The four bars you already use are still there, unchanged, directly below the map.

The step that sets your pace is the red one. The same red the bars use, so the two agree. Whichever step is the wall wears a **"the slow one"** tag and the one sentence that explains why - at your current numbers that is the chiller, holding 12 trays; fill in the three untimed steps and the mark moves to the wash, and the traffic goes to your hands instead.

One honest note about the hand steps. The six hand jobs are one shared pool - you do them in turn - so no

single one of them has a pace of its own. The map therefore shows each as **minutes a pan of work** and never as a rate, and when the wall is the hands the "slow one" mark goes on the heaviest job while the sentence beside it says the minutes are shared across your pairs of hands. Only the chiller, the oven and the pans are given a rate, because those are the only three that are not you.

20 Sep 2026 - engine v135 (no database step). The mixer box on the Production line now explains itself, and its number is right: 25 kg is about 28 pans, not 25.

What you asked. "Most dough in one mix, in pans - What is this?" Fair question, and the fault was mine. Your form asked for your mixer in **litres or kg** ("Most dough your mixer takes in one go - look for the bowl size on the mixer or in the manual"), and the Production line asks for it in **pans**, because every other number on that screen is in pans so it can be compared to the rest. I converted your answer without telling you, and never checked which unit you had meant. So the box read 25 and quietly meant "25 pans" - a figure you had never given.

What changed. You told me your mixer takes **25 kg** of dough. One pan takes about **900 g** of dough - that is what bakes down to the 800 g you sell, the rest being water that bakes off and oil that goes in the pan - so 25 kg is about **28 pans**. That is now the figure the box is seeded with, and the hint under it says the whole thing in one line, so the conversion is on screen rather than in my head: *"Your mixer's bowl size, counted in pans: one pan takes about 900 g of dough - that is what bakes down to your 800 g - so a 25 kg bowl is about 28 pans."*

One caveat, said plainly. A 25 kg mixer means the bowl holds 25 kg, but a wet focaccia dough climbs the hook and will overflow a bowl filled to the brim, so what you can actually mix in one go is probably less. The hint says so too - lower the number if you never fill it. You know your mixer better than the arithmetic does.

If your box still reads 25, type 28. The figure the app remembers for you is the one you last typed, so your own phone keeps showing 25 until you change it - the new seed only reaches a fresh phone. This is the one small thing to do by hand.

What it affects. Two things, and both are small at your scale: how many mixes a day takes (the day's pans divided by this), and - once you time "Minutes to weigh in and load one mix" - the mixing minutes divided across each pan. At 28 pans a day of 36 is two mixes.

20 Sep 2026 - engine v134 (no database step). Every time box on the Production line now says what it wants: minutes.

What you asked. "the field input unit is not shown?" You were right, and it was worse than a missing label - it was six missing labels. Every other box on that screen carries its unit in its own name ("Minutes per bake", "Pans per bake", "Hours you'll bake for"), so the pattern was already there and the time boxes did not follow it. You saw "Wash, oil and fill 6 pans" above a box holding 18 and had to guess whether that meant minutes or seconds.

What changed. All six hand-work boxes now lead with the unit, so the box and its number read as one thing:

- **Minutes to wash, oil and fill 6 pans Minutes to dimple and top 6 pans Minutes to take 6 out and put 6 in Minutes to weigh in and load one mix Minutes to weigh the dough out into 6 pans Minutes to cool and pack 6 pans**

The two card headings say it as well - the times you measured are "each for 6 pans, in minutes", and the rest of the kitchen work is "the hand-work around the bake, in minutes".

Nothing else moved. No number changed, no arithmetic changed, no figure on the screen reads differently. This is the labels only, so if you had already worked out what went in each box, nothing you typed has been affected.

Nothing to set up, and no database step. Open the app and the boxes say minutes.

20 Sep 2026 - engine v133 (no database step). The Production line now counts the whole bake day, not three jobs of it: weighing in and loading the mixer, weighing the dough out into pans and cooling and packing each have their own box to time, and any job you have not timed is named on the screen rather than quietly counted as free.

What you asked. "no time alocation for handling the mixer? and weighing the dough, no other process that need

manpower, or should still need some resources?" You were right, and it was a real hole in the screen. It was timing three jobs - wash oil and fill, dimple and top, and the oven swap - and treating those three as the whole day's work. The mixing, the weighing out and the packing were not in the arithmetic at all, so the day it promised you was a day nobody could actually have delivered. Worse, the plan I wrote for this screen had named cooling and packing; the screen I built dropped it.

What changed. "The rest of the kitchen work" is a new card on the Production line with three more boxes to time, the same shape as the three you already timed: weighing in and loading the mixer, weighing the dough out into pans, and cooling and packing. The list of jobs now reads all six, each with its minutes a pan or the words "not timed", and every pair of hands is placed against all six and not just three.

A job you have not timed is named, never assumed free. Leaving a box blank is allowed - you may already have that time inside your wash, oil and fill figure, and there is a line under the box saying so. But a blank box does not silently count as nothing: the screen lists the jobs still untimed and tells you plainly that the day above looks longer than it really is until you time them. A time with nowhere to go - minutes on the mixer but no batch size for it to spread over - is treated the same way, rather than being divided by nothing.

The mixer is timed per mix, not per pan. Its minutes are spread over the batch it makes, because that is the honest reading and it means a bigger mixer really does cost you less work for every pan. Twenty minutes of weighing in is one seventh of a pan's work in a 25-pan mixer and half that in a 50-pan one.

And the answer changes when you time them. This is the part worth knowing before you do. At your own numbers with the three new boxes filled in - 20 minutes a mix, 3 minutes to weigh out 6 pans, 12 minutes to cool and pack 6 - the whole day's hand-work is 8.3 minutes a pan, so one pair of hands does about 7.2 pans an hour instead of the 12 the screen used to claim. Five hours then makes **36 pans, not 60**. The wall moves too: it stops being the chiller and becomes your own hands. Every number on the screen is yours to correct, and the day figure moves the moment you type.

Nothing to set up, and no database step. Open the app and the three new boxes are there, blank and waiting to be timed once.

19 Sep 2026 - engine v132 (no database step). Your customer's own track page now keeps up with the order: anything you change about an order reaches their page when you save it, and an order you have just added has a page behind its link from the moment it exists.

What you asked. "add new order and edit order din sync?" They were, on your own phones - an order is one record, so an edit made on one phone is on the other within seconds. What was NOT in step was your customer's track page, and that was a real defect. Their page was only rewritten when the delivery day moved, the tracking number changed or the courier charge changed - so if you corrected the items, the price, the address or the customer's name and saved, their page went on showing the order as it used to be: the old items, and the old total, in the very message you had already sent them the link in. Adding an order wrote no page at all, so a link you sent straight away could answer "not found".

What changed. Every way of saving an order - adding one, editing one, the Note / tracking box, and deleting the courier's Delivery & fuel line on Money - now hands your customer's page the whole new version of the order, and the page itself decides whether anything it shows has moved. Where the app used to hold a short list of the things it believed that page showed, it now compares the entire row, so a figure on that page can never be left out of the decision again.

An added order has its page at once. Adding an order publishes your customer's page there and then, instead of waiting for the first time you happen to change its status, so the track link in your confirmation message always has something behind it.

A save that changes nothing writes nothing. The page compares itself with what you last published, so saving an order without changing anything they can see writes no page at all - and if a write is ever refused (a phone with no signal, a SQL script not yet run), the next save tries again rather than counting it as done.

Nothing to set up, and no database step. Open the app and it is already in step.

19 Sep 2026 - engine v131 (no database step). A Production line screen: your bake day drawn as a row of bars, the station that is holding you back named and made the only red thing on the screen, and a ranked list of the changes you could make with what each one actually buys you.

What you told me. "can i have a apps function that visualise the flow, bottle neck hightlight and something for me to optimise line balancing, so that i can have an optimise manpower usage?" Then, when I asked how it should help you decide: "Both - ranked levers and what-if", and the numbers should come from "A planner I type into". You asked for all four levers to be yours: **People**, **Pans and trays**, **The task times**, and **Days and hours**.

The line, drawn. Open **More -> Production line** and your day sits on one screen: one bar each for **your hands**, **the chiller**, **the oven** and **your pans**, every one measured in the same unit - **pans an hour** - so they can be compared honestly. The shortest bar is the one setting your pace. It wears a red wash, is named in a sentence just underneath, and it is the only red thing on the screen, so your eye lands on the one thing that matters. Each bar carries a plain line saying what that station is: "1 pair of hands · wash, oil, fill, top, oven swap", "12 trays of dough over 5 hours", "6 pans every 15 min".

What your own numbers already say. The three tasks you timed for every 6 pans - 18 minutes washing, oiling and filling, 8 dimpling and topping, 4 for the oven swap - are 30 minutes of work for 6 pans, which is **5 minutes of work for every pan**. One pair of hands therefore finishes **12 pans an hour**; your oven does **24**. And the chiller, holding your trays of dough overnight at one tray to one pan, turns 12 trays through a 5-hour day at only **2.4 pans an hour**. So the screen names **the chiller as the wall, not the oven** - your day comes to **12 pans**, 48 short of the 60 you want - and it tells you that **one more pair of hands buys nothing at all** while the trays are already full.

The third pair of hands buys nothing. That is the manpower answer, and the screen says it in words rather than leaving you to work it out: once the chiller is out of the way, the **second pair doubles the line** (12 pans an hour to 24) and the **third buys nothing**, because the oven stops at 24. Spending on help before fixing the wall is money standing around.

Where the hands go. With two or more pairs the screen splits them by how much work each job is, so two pairs reads "1 on the wash, oil and fill, and 1 on the dimple and top", with the oven swap taken by whoever is free. Every job finishing at the same moment is the whole trick to balancing a line, and it is the same arithmetic at two pairs of hands or at six.

What to change, ranked. Six moves are scored by the pans they actually buy on a day like yours and sorted best first - **one more pair of hands**, **six more trays**, **six more pans**, **a wash 10% faster**, **one more hour**, **one more pan per bake**. Each row says what the day becomes and which wall is left standing after it. At your numbers only **six more trays** gains anything (+6 pans, taking the day to 18) - and the other five keep their place and say so in as many words ("No help - the chiller still sets the pace"), so nothing on the screen ever looks like a button that does nothing.

A planner, not a gate. Nothing on this screen reads an order, nothing here blocks a sale, and nothing here asks you to keep anything in step with it - the numbers are yours to correct right on the screen, and the screen's answers move as you type. It is seeded from the kitchen form you filled in, so it says something true the first time it opens. The planner shares between your phones like the rest of your settings, so both of you plan from the same numbers.

One number worth settling once. Your form answered **8** trays where you have also told me **12**. The tray count is your day's hard ceiling, so it decides everything above - the field says so, and it is yours to set.

19 Sep 2026 - engine v130 (no database step). The courier charge can now be changed from Edit as well, and deleting the courier's line on the Money screen takes the charge off the order with it.

What you told me. "q2, yes make deleting it clear the charge too", and "edit a courier related now only at one button, we should make edit an order able to alter details for courier."

Deleting the courier's line now clears the charge. A charge you paid has two halves: the **Delivery & fuel** row in your books, and the charge on the order. The Money screen only shows you the row, so deleting it used to take that half alone - the order still wore the charge, and the next Save in the order's own courier box quietly

wrote the row straight back. Now the row and the order go together: the amount, who bore it and the COD mark all come off. The confirmation says so before you agree, naming the order, and the row is labelled **Courier (order #A3F9C2)** so you can see which order you are deleting it from. Deleting an ordinary expense still touches no order at all.

The courier charge is under Edit too. The amount, who paid the courier, how they settle it and how you paid are now asked in the Edit order pop-up as well as in the **Note / tracking** box - one shared set of controls, so the two forms cannot word it differently or save a different charge. Edit's own order total moves as you type the amount, so you can check the figure before you save rather than after.

Nothing was taken away. The Note / tracking box keeps all of it and is still the quick way in; Edit simply gained the same questions, so an order can be corrected end to end without opening a second box. Nothing changes on its own: a charge is only written when you save, and saving an unchanged form writes an unchanged charge.

19 Sep 2026 - engine v129 (no database step). Still to collect now counts the courier charge the customer pays with the order, so the row promises the same money their own message asks them for.

What you told me. "q1, it should reflex rm72." RM72 was RM64 of bread plus the RM8 the courier was charging, and the Money screen was showing RM64 - the items alone.

What it does now. When the customer bears the courier charge and pays it with the order, **Still to collect** counts what they will actually hand over: the items plus that charge. So the figure on your Money screen is the same RM72 their own message asks them for, and you are never left counting a charge you are about to be paid.

A COD charge stays out of it. When the courier collects the charge at the door, that money was never yours to collect - so it is not in this row either, exactly as it is not in the total you ask them for.

A charge you bore is not in it. Nothing changes there: the bread is what they owe you, and the postage is your own cost.

Your takings do not move. What has already come in - Cash in, TNG in, the net, every journal - still counts the items alone, because a charge the customer pays is pass-through: it arrives and leaves in the same breath, so it was never profit. Only **Still to collect** changed, because that is the one figure about money still on its way to you.

19 Sep 2026 - engine v128 (one database step - run supabase/courier_cod.sql BEFORE you deploy). A courier charge the customer bears can now be a Courier COD charge: the courier collects it at the door instead of it sitting in the total you ask for.

What you told me. "courier charges can be collect, that means customer pay courier upon collect." A charge the customer bears was always being added to the total you ask them for. That is right when they pay it with the order - but when the courier collects it at the door, asking for it by TNG as well means the same money is asked for twice, once by you and once by the courier.

What it does now. The Note / tracking box asks how they settle it. Under **Who paid the courier**, answering **The customer paid it** now opens a third box, **How they pay it**, with two choices:

- **With their order (in the total)** - the charge sits inside the total you ask for, exactly as it always has. Nothing you have already recorded changes, because this is where that box starts. **COD - the courier collects it on delivery** - the charge comes out of the total you ask for, and is named everywhere as payable to the courier when the order reaches them.

Nobody is asked for the same money twice. The confirmation, the payment reminder and the shipped message all quote the items alone as the total, and name the charge underneath as COD with the instruction spelled out - "Courier charge: RM 8.00 - COD, pay the courier when your order reaches you". The customer's track card says the same. COD is the word the couriers themselves use for a parcel the receiver pays for; it is also called a reverse charge, but COD is what your customers will understand.

What you see. The box tells you what the customer owes as you type, and it now reads differently for COD - "the customer owes RM 30.00 - items total RM 30.00, plus RM 8.00 collected by the courier on delivery" - so

what you are asking for and what the courier is asking for are never confused. The Edit pop-up's order total follows the same rule. And the order's row is tagged **Courier RM 8.00 · customer · COD**, so you can see at a glance which charges the courier is collecting rather than looking for that money in your tin.

Your books are untouched, either way. A charge the customer bears was already pass-through, and COD does not change that - it still never reaches your takings. Only a charge you bore yourself becomes a Delivery & fuel expense, exactly as before. (v129 above changes where that charge is counted under **Still to collect**.)

One database step. Run **supabase/courier_cod.sql** in the Supabase SQL editor BEFORE you deploy this build. The app publishes the whole tracking row in one call, so if that column does not exist yet the call is rejected as a whole and the customer's tracking page stops updating for EVERY order - not only the ones with a COD charge.

19 Sep 2026 - engine v127 (no database step this time). Deleting a courier charge really deletes it - the tag comes off the order and the charge leaves the customer's track card.

What you told me. "why i delete courier charges and the tag is not remove?" You were right, and it was a real fault, not a display quirk.

What was wrong. The charge is two things: the amount, and **who paid it**. Putting **Who paid the courier** back to **Not recorded** felt like deleting the charge - and it is the obvious way to do it, because that picker is what decides where the charge goes - but it only deleted the answer to *who*, and left the amount behind. The order's row kept a tag reading **Courier RM 8.00 · customer** for a charge nobody had assigned, and nothing you could do in that box would take it off. Worse, it named the customer as the payer when nobody had said so, so a row could claim they owed money they did not.

What it does now. A charge is the amount AND the payer. Put the payer back to **Not recorded** and both go: the tag comes off the row, the customer's total goes back to the bread alone, and the expense row goes out of your books. Clearing the amount still does the same, as before. And a charge with no payer recorded - one restored from an older backup, say - is tagged with nothing at all rather than being labelled as the customer's.

Clearing a charge now reaches the customer's track card too. It did not before: the card was only republished when a charge was added, so a customer could keep looking at a courier line you had just deleted. Now the card is republished whenever the charge changes, in either direction.

Where to see it. Orders -> an order with a courier charge -> **Note / tracking**, then set **Who paid the courier** back to **Not recorded** and Save. The tag goes from the row.

One thing left as it is, so you know. Deleting the **Delivery & fuel** expense from your Money screen does *not* take the charge off the order - that row is your own record of money going out, and the order is the record of what happened, so the two are deliberately separate. If you would rather that deleting the expense cleared the order's charge as well, tell me.

19 Sep 2026 - engine v126 (no database step this time). The messages that ask your customer for money now show how the total is reached, and the app shows you the same arithmetic.

What you told me. You pasted the confirmation for order #7E51BC - 4 × Focaccia 9"x12" 800g, self collect, RM8 courier, Total RM 72.00 - and said: "the message need to show the add up for rm72. And it should be the same for APP."

What was wrong. The RM72 itself was right. What was missing was any sign of where it came from. The line above it named the bread without a price, so the RM64 of foccacia and the RM8 courier charge were both invisible, and the total read as a figure picked out of the air. That is the worst place to leave it, because this is the message that asks them for the money.

What it does now. When the customer bears a courier charge, the messages no longer state the answer on its own - they show the sum, in the same order every time, with the total underneath:

- Items: Focaccia 9"x12" 800g x4 Items total: RM 64.00 Courier charge: RM 8.00 Total: RM 72.00

The app says the same words, in the same order. One piece of code works out those two figures and their

total, and the confirmation, the payment reminder, the shipped message and your own screens all quote it. Two surfaces cannot drift apart when there is only one sum.

On your side of the app. The **Note / tracking** box's "The customer owes..." line and the **Edit** box's order total both now read the breakdown as **items total ... + courier charge ...** - the same two words the customer sees, so the figure in front of you and the figure they were sent are provably the same arithmetic.

An order with no courier charge is untouched. Its messages come out byte for byte as they did before, and the tests hold that.

Where to see it. Orders -> an order with a charge the customer pays -> **Send confirmation**, **Remind for payment** or **Send shipped message**. Your side: that same order's **Note / tracking** and **Edit** boxes.

One thing left for you to decide - not changed. The Money screen's "still to collect" counts the items only: RM64 on this order, not the RM72 the customer will hand over. That is deliberate, because the line measures what stays in your purse and the RM8 arrives and leaves again in the same breath. If you would rather it showed the full amount you are about to be handed, tell me and I will change it.

19 Sep 2026 - engine v125 (no database step this time). The courier charge now appears in the two totals YOU read, and in the one customer message that was still missing it: the confirmation.

What you told me. Straight after deploying v124: "if customer were to paid courier, the total is not shown to me and to customer." You were right, and it was three places, not one.

What was wrong. v124 did put a charge the customer bears into the payment reminder, the shipped message and the track card, and named it in each. But the **confirmation** - the very first message, the one carrying the TNG QR that asks them for the money - still quoted the items alone. And on your side of the app, the **Note / tracking / courier** box never showed what the customer owed, and the **Edit** box's order total left the charge out. So both totals you read were lower than the figure the customer had actually been asked for. A customer paying what the confirmation said would have come up short, and you would have had no way to see why.

What it does now. The **confirmation** names the charge on its own line just above the total, exactly as the payment reminder already did - so the message that asks for the money and every message after it agree. The **shipped message**, which named the charge but never said what the order came to, now ends on the total as well. In your app, the **Note / tracking / courier** box tells you what the customer owes as you type the amount, and it moves the moment you say the customer paid it - showing the arithmetic, items plus courier, so the figure explains itself. The **Edit** box's order total counts the charge the same way.

A charge you bear still appears in none of the customer's totals and in none of the figures they are shown - that part has not changed, and it is still deliberate. Your own books are untouched by a charge the customer pays, exactly as v124 had them.

Where to see it. Orders -> open an order -> **Note / tracking** (the total sits under the question of who paid). **Edit** -> the Order total line. The customer's side: all four messages, and their track card.

19 Sep 2026 - engine v124 (one database step: run supabase/courier_fee.sql in Supabase BEFORE you deploy this build). The Note / tracking box now records the courier charge itself, and who bore it - and the two answers behave differently on purpose: a charge the customer pays is added to their total and named where they can see it, and a charge you pay comes off your profit under Delivery & fuel.

What was wrong. You looked at the box behind an order's **Note / tracking** button and said the **Paid by** picker there was for courier charges. It was not. It records how the CUSTOMER paid YOU - Cash or TNG - and it is what puts an order's money into the right book on your Money screen. It has been there since v101, from your own words about needing to reconcile cash against the TNG app. Your reading was the more useful one though, because the thing you were looking for did not exist anywhere in the app: there was no field for what the courier charged, and none for who paid it.

What it does now. That box carries the **Courier charge** - what it cost to send that order - and, under it, **Who paid the courier**.

When the customer pays it, the charge is added to what they owe. It is named on its own line above the total in

their payment reminder and their shipped message, and on their track card, so a total that runs a little more than the bread explains itself rather than looking wrong. Your books are not touched at all - and that is deliberate, not an oversight: in your purse that money arrives and leaves again in the same breath, so counting it as takings would only make a good day look better than it was.

When you pay it, a third box asks **How you paid the courier**, and the charge becomes an ordinary **Delivery & fuel** expense dated the day you paid it. It then shows up on your Money screen, in that account's own journal, and as a line on your profit statement - so it comes off your profit like any other cost, which is exactly what you asked for: off profit only if you paid it. Recording the same order twice updates that one expense rather than adding a second, and clearing the amount takes it away again.

Either way, the order's row is tagged with the charge and whose money it was - amber when it was yours, quiet grey when it was theirs, sitting beside the Cash / TNG tag.

The box you misread is renamed. The picker that was called **Paid by** now reads **Paid by the customer**, and it sits below the courier charge rather than above it, so the two questions cannot be confused again: one is what THEY paid YOU, the other is what YOU paid the courier.

Where to see it. Orders -> open an order -> **Note / tracking**. The customer's side: the payment reminder, the shipped message, and their track card.

Before you deploy. Run **supabase/courier_fee.sql** once in Supabase (Dashboard -> SQL -> New query -> Run). The app publishes each order's tracking row in a single call, so if that column is missing the call is refused whole and the customer's track page stops updating for every order - not only the ones with a charge. Run the SQL first, then deploy.

Nothing else moved - your prices, customers, delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v123 (no database setup needed). The product dropdown's middle section now groups sold-out and not-sold-that-day together, and a product the shop does not sell on a day is named by the days it IS sold instead of being given a count for a day it was never on.

What was wrong. The dropdown's middle section was called **Sold out** and read its count for the day you were adding to. But a product the shop does not sell on that day at all - a Saturday-only loaf on a Wednesday - was sitting in **On the shop**, and worse, could be labelled "sold out" from a count for a day it was never on the menu. Neither told you anything true.

What it does now. The middle section is **Unavailable**, and it holds both kinds together: a product that has **sold out** for the day you are adding to, and a product the shop **does not sell on that day**. The second kind is named by the days it IS sold - a Saturday-only loaf on a Wednesday reads "(only Sat & Sun)". The headings are now **On the shop**, **Unavailable**, **Taken down**.

Why they are one section, not two. They are the same thing to you: an active product you can still choose. You bake to a plan of your own, and as you put it - every day the kitchen may make 12 focaccia, ordering closes the day before at 6pm, and the 12 are not necessarily sold off, so some go to a walk-in or into the fridge. **The picker is a guide while you sell, never a gate**: nothing is blocked, and the order-by deadline you set per product is deliberately not counted, because that deadline stops a stranger ordering - it says nothing about what you may sell by hand.

Where to see it. Orders -> **+ New order** -> Items, and Orders -> open an order -> **Edit**.

Nothing else moved - your prices, customers, delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v122 (no database setup needed). The product dropdown now shows the three kinds inside the list itself: each kind sits under its own heading, and on a newer phone each section is tinted in its colour.

What changed. Adding items to an order - in **+ New order** and in **Edit** - the product list now comes in three labelled sections instead of one long run of names: **On the shop** first, then **Sold out**, then **Taken down**. The

headings appear on every phone, because they are part of the list itself.

The colour. On a phone with the very latest system - a new iPhone or a recent Android - each section is also tinted inside the open list: **green** for on the shop, **amber** for sold out, **grey** for taken down. The menu then reads as three blocks you can tell apart without reading a word.

On an older phone. The open list is drawn by your phone itself, and an older phone will not let a page colour it. You get the same three headings in the same order, and the closed box still wears the colour of whatever it holds - you lose the section tint and nothing else. Nothing is broken, and nothing has moved; only the colour is missing.

This replaces the v121 note below. v121 put the colour on the closed box alone, because the open list could not be reached. Newer phones can be reached, so the list carries the colour too.

Where to see it. Orders -> **+ New order** -> Items, and Orders -> open an order -> **Edit**. The list is otherwise unchanged: the same products, the same "(hidden)" marker, the same "12 left" count.

Nothing else moved - your prices, customers, delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v121 (no database setup needed). The product dropdown now shows you which kind of product each choice is, in colour, and lists them in the order you would look for them.

What changed. Adding items to an order - in **+ New order** and in **Edit** - the product list is now sorted into the three kinds you already know, in the order you would reach for them: what is **On the shop** first, then anything that has **sold out** for the day you are adding to, then whatever you have **taken down** (still marked "(hidden)"). A product still in Draft stays out of these lists entirely, as it always has.

The colour. The closed box wears the colour of whatever is picked - **green** for on the shop, **amber** for sold out, **grey** for taken down. On an order with several lines that means one glance down the list tells you which lines need a second look, without reading each one.

Why only the closed box. The open list is drawn by your phone itself, and iOS does not let a page colour the rows inside it. So the colour sits on the box you see while the list is shut, which is where you read it anyway.

Where to see it. Orders -> **+ New order** -> Items, and Orders -> open an order -> **Edit**. The list itself is unchanged: the same products, the same "(hidden)" marker, the same "12 left" count.

Nothing else moved - your prices, customers, delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v120 (no database setup needed). One customer is one row. The app now recognises a phone number by its digits, so the same person can no longer appear twice just because their number was written with a + in front of it - and the duplicate already in your app joins itself into one the first time you open it after updating.

What was wrong. A customer used to be identified by their number written in exactly the same way every time. So **+60123456789** and **60123456789** looked like two different people, and the Customers list showed the same customer twice. Your messages always went to the right person - only the list disagreed. Numbers are now read by their digits, which means **012-345 6789**, **+60 12-345 6789** and **60123456789** are one person, the same way "012345" already found "012-345" in the finder.

It tidies itself once. The first time you open the app after updating, it joins any customer who was split this way, carrying their dog's name and photo, what they like and avoid, and their note across to the one that remains. There is nothing to do - you will simply find one row where there were two. If you are the careful sort, take a cloud backup first from **More -> Backups**: that is the way back if anything looks wrong, and it costs one tap.

New: "Join with another customer". Open a customer's history and you will find the button under their number. Tap it, pick the duplicate from the list, and confirm. Reach for it when two rows are split for some other reason - a misspelt name with no number beside it, say - which the tidy above cannot guess at. It moves the other person's orders under the name and number of the customer you were looking at, so open the row you

want to KEEP. This one cannot be undone inside the app; your cloud backup is the way back, and the confirm box says so before it acts.

Two things worth knowing. Update both phones - a phone still on v119 can save a number the old way from its Profile card, and only the updated build does the tidy. And on the Profile card itself, the number you type is now stored in the one form every label, message and link already uses, so a "+" or a space can never split a customer again.

Nothing else changed - your prices, your delivery dates, the shop and the way messages are written are all exactly as they were.

18 Sep 2026 - engine v119 (no database setup needed). Start typing a customer's name in an order form and the people you have already served appear underneath - tap one and their name and WhatsApp number fill themselves in.

How it works. Type two letters - "aun" is plenty - and a short list appears right under the name box. Each row shows their number, how many orders they have placed, and the item they buy most, like **012-345 6789 · 5 orders · usually Focaccia**. Tap the right person and the name and the number both go in; the delivery day, the items and self-collect-or-courier stay yours to set. One letter on its own is not a search - it would offer you half your address book - and a name you have never served offers nothing at all. Typing a number without its dashes finds them too, so "012345" works the way "012-345" does.

Both order forms do it - the **+ New order** card and the **Edit** pop-up, the same way.

Two things worth knowing. It draws on the orders already in the app, so on a brand-new setup it has nothing to offer until your first order is saved. And on the **Edit** form, saving a changed name has always carried that name to every other order belonging to the person it was - that is the one-saved-name rule from v68 and it has not changed, but picking a name is now much easier than typing one, so it is easier to do by accident. If you ever pick the wrong person while editing an order and save it, tell me and I will add a confirmation that asks first.

Nothing else changed - the Customers screen, its finder, your prices and the shop are all exactly as they were.

18 Sep 2026 - engine v118 (one SQL step in Supabase; nothing to re-install). The order alert on your phone now arrives as a real notification instead of a wall of raw text. The ping that reaches you when a customer orders was showing the whole message as unreadable computer text - curly brackets, quote marks, the word "title" - instead of a notification. The cause was the way the message was handed to the ntfy service: the channel name has to travel inside the message for ntfy to lay it out as a title and a body. It was missing, so ntfy did the only thing left and printed everything as plain text.

What you see now. A notification titled **New order - Jien Luv 2 Bake**, with the delivery day written out in full (Sat 19 Sep 2026, not 2026-09-19), the total as money (RM 10.00), a bread icon, and a higher priority, so it is more likely to make a sound and pop up rather than wait quietly. A courier order also carries the address; a self-collect order no longer shows a delivery address it does not have.

A tidy-up came with it. Your channel name used to be typed inside the rule that sends the pings, and three leftover test rules from setting this up carried copies of it too. It now lives in one small table, and the three leftovers have been removed - so the name that guards your alerts exists in exactly two places: your database, and the ntfy app on your phones.

Nothing about ordering changed, and no phone needs re-subscribing.

17 Sep 2026 - engine v118 (no database setup needed). The Paid step now stays in its place on every order, wearing an X until the money is in - and turns into a green tick when it is. This replaces what v117 did, on your instruction.

Your reasoning, and you are right. The customer's track page is the design: they must see the same steps in the same places every time, and your own app has to read the same way. v117 left the Paid step **off** an order that skipped it - five dots instead of six - which meant the two lines did not line up, and a missing step is itself a thing to puzzle over.

What happens now. A regular who pays when she collects still never passes through Paid. But that step keeps

its place in the line, and it wears an **X** - amber, deliberate - instead of a tick, and it never turns green while the money is outstanding:

New tick · Confirmed tick · **Paid X** · Baked tick · Packed tick · Collected / Shipped dot

Press **Paid · Cash** or **Paid · TNG** when she hands the money over - at whatever stage the order has reached - and **the X becomes the tick**: the ordinary green tick, in the same spot, with the **Cash** or **TNG** tag beside the row. The order itself stays where it was: taking the money never drags it backwards.

Both lines change together. Your row's map and the **customer's own track page** draw the same six steps, so the X is on their page too, and it turns green there the moment you record the money. A step that is *deliberately gone past but not paid* now looks different from a step that has not been reached (grey) and from a step that is done (green) - which is exactly the distinction that was missing.

17 Sep 2026 - engine v117 (no database setup needed). A regular who pays when she collects no longer has a Paid step on her order - and the Paid · Cash / Paid · TNG buttons stay on until the money is actually recorded, wherever the order has got to. From your note about close customers paying by TNG or cash at pickup.

What was wrong. Move an order from Confirmed straight to Baked - no payment - and the row's map still showed **Paid with a green tick**. It was claiming money you had not taken. The app assumed "past Paid means paid", which is fine for her old orders but false for a regular who settles up at the counter.

What happens now.

- **A bypassed order has no Paid step.** Her route reads New -> Confirmed -> Baked -> Packed -> Collected / Shipped: **five dots instead of six**, and no tick for a payment that never happened. **The Paid · Cash / Paid · TNG buttons stay on** at every stage from Paid onwards - Paid, Baked, Packed, and Collected / Shipped - until the money is recorded. So the moment she hands over the cash at the counter, wherever the order has got to, you press the button right there: **the order stays where it is** (marking an order paid never drags it backwards) and the money is stamped with the day it landed, which is the day the Money screen counts it. **And the Paid step comes back once it is paid** - a green tick, in its proper place in the line, with the **Cash** or **TNG** tag beside the row. So the map ends up telling the truth about that order either way. The buttons disappear the moment the payment is recorded, and they never appear before the order reaches the Paid stage. **The customer's own track page follows the same rule** - the same five dots, no Paid tick for a payment that has not happened, and the step appears there too once it is recorded.

One rule behind it, worth knowing: moving an order into Paid **or any later stage** without having pressed a Paid button is taken as *this order owes money*. That is what your older orders are protected from - nothing already in the app changes, because only orders you move now are marked.

17 Sep 2026 - engine v116 (no database setup needed). A Day one form: your opening balance - the cash in the tin, the money on your phone and what is on your shelf - entered once, in one place. From you telling me *"we need to enter opening balance"*.

What an opening balance is here. Four answers about the morning your books begin: what is in the tin, what is on your phone, what is on your shelf, and who still owes you. Until now the app could take all four, but in three different places and one ingredient at a time, which is a poor way to start.

The form. More -> Money -> the new **Day one** line (under **Books**) -> **Set**. It asks for:

- **Cash in your tin** and **Money on your phone (TNG)** - two boxes, and the day your books begin, which takes any date but should be the day you are starting from. **What is on your shelf** - your ingredients, one box each, already in the unit you use for that ingredient (kg for the flour, g for the salt - no converting). Ingredients you have marked as *not something I buy*, like labour, are left out: they have no shelf to count. **Who still owes you** needs no box - those orders stay unpaid and show under **Still to collect**.

Nothing is half-saved. Every box is checked before anything is written: if one has something that is not a number, nothing saves at all until you fix it. And **a box you leave empty is left exactly as it is** - so the form is safe to reopen later just to correct one figure, and safe to run again without wiping what you already had.

What it writes. The tin and the phone become ordinary money-in rows dated that day (so from then on the Money screen's **Net** is what you should really hold, and the weekly check works from the first week), and the shelf becomes each ingredient's **On hand** - the figure your shopping lists subtract, so your first list buys only what you are genuinely short of.

What it deliberately does not do. Money you **owe** - a loan for the oven, a supplier you have not paid - has nowhere to go: the app keeps a loan as a way of paying, not as a debt, so record what it pays for as it happens. Your own money in shows under **Capital you put in** on Profit, never as income. And the tin must not be entered as a sale: a sale has a customer behind it.

Section 3 of the **Money and Profit** manual has been rewritten around the form - it is attached.

17 Sep 2026 - engine v115 (no database setup needed). The month arrows on Profit work both ways again. From you telling me *"the profit month can move earlier but cannot move later"*.

What was wrong. Open Profit - it starts on this month, where the **right** arrow is correctly switched off, because there are no numbers after today. Press **left** to step back to August, and **right** stayed switched off too, so you could not get back to September. You could walk backwards through the months but never forwards again: to return you had to leave the screen and open it again. The arrow's state was worked out once, when you opened the screen, and then reused on every step - so it kept answering for the month you started on rather than the month you were looking at.

Fixed. The arrows are worked out fresh on every step, so **left** and **right** now move both ways between the months you have, and **right** switches off again only when you are back on this month. Verified by stepping back and forward repeatedly. The same fault did not exist on the other calendars - Orders, Delivery dates and the date pickers all recompute their arrows each time they draw.

17 Sep 2026 - engine v114 (no database setup needed). Every spending line in Profit & Loss now opens - including the ones reading 0.00 - and the lines are twice as tall, so they are easy to tap. From you asking *"in profit the expenses is not clickable, is that a bug?"*

You were right, and it was my mistake. The lines that open their journal were the ones with money in them. A line showing **0.00** was deliberately left dead, because there is no journal to show - but I made it look **exactly** the same as a live one: same colour, same font, no hint. So if your month has most categories at 0.00, nearly every line you tapped did nothing, and "not clickable" was the only sensible conclusion. A line that looks alive and does nothing is worse than no line.

What it does now. Every spending line opens: *Packaging*, *Utilities*, *Delivery & fuel*, the categories you have added, and **Total expenses**. A line with money in it shows its journal as before. An empty one opens and **says so in your own words** - *"Nothing recorded under Rent in September 2026"*, In / Out / Net at zero, and a line telling you it will fill up on its own as you record spending under that category. The statement's own totals (**Sales**, **Cost of sales**, **Gross profit**, **Net profit**) stay figures rather than doors - you did not ask for those, and they are not spending.

Two more things found while fixing it.

- **The lines were only 17 pixels tall** - a small target for a finger, and easy to land in the gap between two lines and hit nothing. Tappable rows everywhere (the Profit statement, the Money screen's cash rows, the Books list) are now **36 pixels**, comfortably thumb-sized. **The Total expenses journal did not say which category a row belonged to.** It read *"2 Sep · boxes · Cash"* with no way to tell what "boxes" was for. It now names it - *"2 Sep · Packaging - boxes · Cash"* - while a single category's journal still reads short, since its title already says which category it is.

17 Sep 2026 - engine v113 (no database setup needed). Every way you pay now has a book you can always open - including a pocket that has been quiet - behind a new Books line on the Money screen. From you asking *"where can i find pocket journals"*.

What was wrong. The journal rows live on the money card, and that card only shows a pocket's row **when that pocket moved money in the stretch you are looking at**. The card opens on **Today**, so a pocket that paid for

something last week had no row today - and no way into its book at all. Cash and TNG never had the problem: their four rows (Cash in, Cash out, TNG in, TNG out) are always there. The pockets were the only books that could go missing.

The Books line. Under the two money cards there is now a second line beside *Categories & ways to pay*:

Books Every way you pay · 5 books, each opening into its own rows [Open]

Open lists **every** way you pay - Cash, TNG, Loan, Personal Pocket Kean, Personal Pocket Suan, and anything you add later - each with what moved by it in the stretch on screen, each opening into its own book **under the line you tapped**, so the list stays in front of you and you can step from one book to the next. The first method with anything in it opens ready; a quiet pocket says so plainly - *"Nothing moved this way in this stretch"* - with its In / Out / Net at zero.

Two things this keeps honest: a method you have since **renamed or deleted** still gets a line, because its money is still in the books and a figure you cannot open is a figure you have to take on faith; and every figure in the list comes off the Money screen's own rows, so the list and those totals cannot drift apart.

A small thing the screenshot caught while building this: the first version printed the word "null" under each closed line. Fixed, and there is now a test that fails if it ever comes back.

17 Sep 2026 - engine v112 (no database setup needed). You can pay a personal pocket back out of the till in one go, and the category list on the expense form is no longer cut off. From your question - *"can my own withdrawal payback to the cash register like Personal pocket Kean or Suan?"* - and your note about the categories.

Pay back a pocket. On **More -> Money**, beside **+ Put money in**, there is now **+ Pay back a pocket**. It is for the case where a pocket of yours - **Personal Pocket Kean**, **Personal Pocket Suan**, a loan, or one you add later - paid for something, so the money left you rather than the till; that pocket's line on the Money screen then reads **-RM 40.00**, which means *the till owes it 40*. Open the form and it comes up **already on the pocket that is owed**, with the amount filled in - Kean's RM 40 sitting in the box, and a line saying so. Say whether the money came out of **Cash** or **TNG**, check the day, and press **Pay back**.

It writes both halves at once, which is the whole point: the **till goes down** by that amount (it shows in Cash out and in the Net, and in the Cash journal with your note), and the **pocket's line comes back to zero**. One pocket can be paid back without touching another. Because it is your own money going back to you, the till's side is recorded as a withdrawal - so it never counts as a cost and your **profit does not move**. On the pocket's own list it reads *"Paid back by the till"*, not *"From my pocket"* - the two are different things and now say so. Paying it back in part is fine: pay RM 25 of the RM 40 and the pocket reads -15 after.

The category list was being cut off - fixed. Thank you for catching this. On the expense form, the row of categories ran off the side of the screen: it was laid out as one single line, so at phone width everything past about the sixth category - *Salary (you)*, *EPF / SOCSO*, *Marketing*, *Equipment & tools*, *Other*, *My own withdrawal*, and the **+ New category** chip - was off the screen and impossible to reach. Measured on a 375-pixel phone, the row was 744 pixels wide inside a 343-pixel box. The pills now **wrap onto as many lines as they need**, the same fix applies to the ways-to-pay row and to the new pay-back form, and the whole list is visible and tappable. Nothing about the categories themselves changed.

17 Sep 2026 - engine v111 (no database setup needed). Every figure on the books is now a door: a journal behind each line of spending in Profit & Loss, and a book for every way you pay - Cash, TNG, Loan, Personal Pocket Kean, Personal Pocket Suan, and anything you add later. From two notes close together: *"the expenses items in Profit & Loss should reveal its journals"*, then *"each CASH, TNG, LOAN, Personal Pocket Kean, Personal Pocket Suan, and others that might be added in future need a journal"*.

Profit & Loss - tap a spending line and see what made it up. On **More -> Profit**, the running-cost lines (Packaging, Utilities, Delivery & fuel...) are now tappable, and so is **Total expenses**. Tapping one opens its journal for the month on screen: every expense behind that figure, oldest first, each reading *"<day> · <what it was for> · <how it was paid>"* and ending on the total the statement itself shows. So *"Packaging RM -58.00"* opens two lines - the boxes paid from a pocket and the bags paid in cash. A line with nothing in it that month is

not tappable; there is nothing behind it. The header now says **Running costs · tap a line to see the spending behind it**.

Money - one line and one book per way of paying. The single *"Paid by loan / other"* row is gone. In its place, every way of paying that is not cash or TNG gets **its own line** with what moved by it, and its own journal behind it: **Paid by Loan**, **Paid by Personal Pocket Kean**, **Paid by Personal Pocket Suan**, **Paid by Bank OD** - whatever you have called them. Money of your own you put in that way reads as a plus on that line; what was paid out of it reads as a minus. Cash in, TNG in, Cash out, TNG out and the Net are exactly as they were, and none of these other ways is in the net - it never came out of your purse.

Why they appear by themselves. The lines are read off your own transactions, not off a fixed list, so a method you add tomorrow needs nothing from me: record one payment with it and it has a line and a book. And a method you have since **renamed** keeps a line for the name its older rows were written with - those rows are still in the books, and a figure you cannot open is a figure you have to take on faith. Each book ends on **In / Out / Net** for that method.

One small correction while here: the note at the foot of the Profit screen about your own hours now matches v110 - either pay yourself a **Salary (you)** expense (with EPF / SOCSO as their own category), or mark Labour as a not-bought ingredient and put the hours into the recipes. Count them one way, never both.

17 Sep 2026 - engine v110 (no database setup needed). An ingredient you never buy - your own labour, electricity, gas - can be marked as a cost, and it then never appears on a shopping list again. From your note: *"certain ingredient we dont purchase, in ingredient we can set that as a non purchase item, like labour and electricity"*.

What it does. Open **More -> Ingredients** and tap **Edit** on the ingredient (or fill in the New ingredient card at the top). Above the cost there is now a switch: **Not something I buy - it only ever costs. It still counts in a recipe's cost, but it never appears on a shopping list or a purchase order.** Turn it on for Labour, Electricity, Gas, your own time.

Where you see it.

- The ingredient's **card** stops showing an On hand and a Keep-at-least line, and says instead: *Not bought - a cost in your recipes, never on a shopping list*. Nothing else about the card changes, and the cost you typed is untouched. Every **shopping list** (the preview, the list you save, and an "orders changed" follow-up) leaves that ingredient out - and says so at the foot of the list: *Not on this list: Labour, Electricity - marked as not something you buy. Their cost still counts in the products that use them.* So it reads as deliberate, not as something the app forgot. The **recipe cost** is unchanged: a recipe using Labour at RM 8 an hour for 15 minutes still adds RM 2.00 to that product's cost, and that runs on into the Profit screen's cost of sales.

Switching it OFF restores the ingredient exactly as it was - the app removes the mark rather than storing a "No", so nothing that exists today changes unless you turn this on.

And a word on the books, since this touches Profit: if you mark Labour as not-for-purchase because it is your own unpaid time, it is already inside each product's cost of sales through the recipe - so do **not** also record it as a Salary expense, or the same hours get counted twice. The **Salary (you)** category is for when you actually pay yourself money out of the till.

16 Sep 2026 - engine v109 (no database setup needed). A journal for each way the money moves - the cash book, the TNG book - and a bug that was hiding in the TNG column. Two things, both from you asking how to see the journals.

A bug first, because it matters. Every order you marked **Paid · TNG** since v106 was landing in "Paid, no method" instead of the TNG column. The app was checking for the old way of writing it ("tng") while the paid buttons write the list's own label ("TNG") - so the TNG in figure read RM 0.00 with the money in your phone, and the same money showed as unaccounted for. Fixed, and there is now a test that would have caught it. If your TNG column has looked wrong since v106, that was it.

Tap a figure, see its journal. On the Money screen, **Cash in / TNG in / Cash out / TNG out** (and the loan line)

can be tapped. Each opens that method's book for the stretch: every order paid that way with its code and customer, everything you spent out of it with its note, anything of your own you put in - in date order, ending on In / Out / **Net**, and a line saying what that is meant to be (your purse, your phone, or money that never went near either).

16 Sep 2026 - engine v108 (no database setup needed). Where the two lists live, and how to change one. One fix and one addition, both about the categories and the ways to pay.

Why. v106 put them in a Settings card and you could not find it, and there was no way to change an entry once it existed - "dont put the setting separately, it should be at where it suppose to be."

They are on the Money screen now, as a line under the money lists reading "Categories & ways to pay · 11 categories · Cash, TNG, Loan" with an **Edit** beside it. Tap Edit and the lists open: every category with its kind (ingredients / running cost / your own money) and every way to pay, each line saying "edit".

Tap a line and it turns into the form: rename it, change what kind it is, or delete it - in place, with the rest of the list still in front of you. Renaming moves what you have already recorded with it, so a slip of the thumb does not split your history into "Packaging" and "Packing".

And the + chips sit where you are working: the Add an expense form and Put money in each end their pills with **+ New category** / **+ New way to pay**, which opens the same small form in place - the amount and the note you had already typed stay put.

No window inside a window. These forms expand in place instead of opening a pop-up of their own, the same rule the date field follows: the app has one pop-up layer, so a second one would wipe out the form underneath it.

16 Sep 2026 - engine v107 (no database setup needed). The item line on an order reads properly again. One fix, on the two forms where you type an order's items.

What was wrong. When the selling price joined the line (v101) it took most of the row. Measured on a 375px phone: the price box was 183 pixels wide and the product dropdown was squeezed to **2** - so the name of what you had picked was pushed onto a line of its own and unreadable, and the whole line looked broken.

What it does now. Each item line is two clean lines: **the product across the top**, full width so its name reads, and its controls under it - how many, the selling price on the right, and the X to remove. The same shape in the + New order card and in the Edit pop-up, since both are the same row.

16 Sep 2026 - engine v106 (no database setup needed). A note on every transaction, and two lists you shape yourself. Three small things, all on the money screens.

A note on each one. Add an expense now takes a note ("Mydin run, 2 boxes") - the Put money in form already had one - and both lists on the Money screen show it beside the amount, so a row reads as what it was, not just a figure.

The categories are yours. More -> Settings gains a **Categories & ways to pay** card. Every category shows what it means to the accounts (ingredients / running cost / your own money), each has a X to delete it, and you can add your own - "Baking class", "SSM licence" - choosing which of the three it is. Deleting one never touches what you have already recorded: those rows keep their label and still count, printed at the end of the statement's cost list rather than dropped.

And the ways to pay are yours too. The list is Cash, TNG and **Loan** - the third choice you asked for - and you can add your own, a **Bank OD** or a cheque, so a shopping run nobody would pay for out of the till can still be recorded honestly. The same list appears on the Bought prompt (where the receipt is still in your hand) and on both money forms.

One thing that follows from it. Money that never came out of your purse - a loan, an overdraft - no longer sits in the net. The Money screen leaves it out of the cash figures and names it on its own line, "Paid by loan / other", with a sentence saying why. Net now means exactly what it says: what should be in your purse and on your phone.

16 Sep 2026 - engine v105 (no database setup needed). The books: a profit and loss account, a month at

a time. One new screen, More -> Profit.

Why. Money told you what was in your purse. It could not tell you whether the baking made money - a RM250 flour-and-butter run makes one week look like a disaster while the bread it went into sells the week after.

What it shows. More -> Profit, month by month (arrows either side), reading down like a set of accounts: Sales, Cost of sales, Gross profit, the running costs one line each, Total expenses and Net profit - with the gross margin beside the order count. Then, kept apart from all of it, your own money: capital you put in, drawings you took out, and what is left in the business.

The one rule that makes it accounting rather than a cash total. Ingredient cost is what the baking cost to make, taken from your RECIPES as the bread sells - not the packs you bought. Buying and using sit in different places on purpose: a shopping run is money out on the Money screen and stock on your shelf, and it becomes a cost of sales here as the bread made from it goes out. Nothing is counted twice, and a big stock-up week stops looking like a loss.

Your own money is neither income nor a cost. What you put in is capital; what you take out is drawings (including the "My own withdrawal" category). Both move cash - Money counts them - and neither changes profit. If you pay yourself a proper salary, record it as **Salary (you)**, with **EPF / SOCSO** as its own category, and those DO count as running costs.

The category list is now a chart of accounts: Ingredients & shopping (stock), Packaging, Rent, Utilities, Delivery & fuel, Salary (you), EPF / SOCSO, Marketing, Equipment & tools, My own withdrawal, Other - the same list the Add an expense form offers, so every row lands on the right line of the statement.

Two things it will not pretend to know. Your own unpaid hours are not a cost - if your time is worth paying, pay yourself a salary and it becomes one. And a product whose recipe prices to nothing is counted as costing nothing; the screen says how many lines that was, so it is visible rather than flattering.

16 Sep 2026 - engine v104 (no database setup needed). Money you put in yourself, and taking it back out. One change, on More -> Money.

Why. Some spending happens before any order money arrives - a packet of flour paid from your own purse, a float of change for the day - and there was nowhere to put it. The Money screen only counted what customers paid.

Put money in. More -> Money now has a **Put money in** button beside Add an expense: how much, the day you put it in, cash or TNG, and a short line saying what it was for. It counts into the Cash in / TNG in rows, because your own money really is in your purse and those rows are what you check the purse against - and one line underneath says how much of the money in was yours: "of the money in, RM 100.00 was your own". The entry is listed in its own **Money in** card, so you can see where it came from, and take it off again.

Taking it out again. Add an expense has a new category, **My own withdrawal** - money you take back for yourself rather than a cost of baking. It leaves through the same money-out list as everything else, so the net, the day lines and your backups count it with no new machinery. Your own money in and out is shared between your phones, like every other list.

16 Sep 2026 - engine v103 (no database setup needed). Money out: what you spend, beside what you take in. One change, on the shopping side of the app.

Why. v101 gave you the money coming in. Spending had nowhere to go: a purchase order worked out what a shopping run should cost, and tapping Bought put the packs on your shelf, but nothing recorded what you actually paid.

A shopping run now asks. Tap **Bought** on a saved list and, once the packs are on your stock, a small box opens: **What did you pay?** - pre-filled with the list's own total, with **Cash** or **TNG** beside it. Accept the guess with one tap, or type what the receipt really said. **Skip the money** leaves the stock added and records nothing, exactly as the app behaved before.

Everything else goes in by hand. More -> Money now carries an **Add an expense** form for what a purchase order never sees - packaging, delivery and fuel, utilities, equipment, or a market top-up - each with the day you

paid it, so a receipt found in your bag later still counts on the right day.

The Money screen adds it up both ways. Cash in, TNG in, Cash out, TNG out and **Net** - what should be in your purse and on your phone for Today / This week / This month - with that stretch's spending listed underneath, each line showing what it was for and how you paid. What is still to collect stays out of the net: it is money owed, not money held. Expenses are shared between your phones, and counted in your backups, like every other list.

One fix along the way. A payment taken just after midnight was being counted as the day before, because the app read the date off the stamp in UTC. It counts on your own day now.

16 Sep 2026 - engine v102 (no database setup needed). Swipe across the calendar to pick a run of delivery dates. One small change, on More -> Delivery dates.

You asked whether the delivery-date calendar could work the way a product's availability calendar does. It now does: **drag across days and the whole run fills in under your finger**, then **Add selected** puts them on the calendar. It is the same gesture as dragging a sell period on a product. A single tap still picks one day (or puts it back), and a day that is already a delivery date is left alone by a drag - it comes off by tapping it, the way it went on.

That makes four ways in: tap dates, swipe a run, tap a weekday letter for every one of that day in the month, or **Generate the next dates** on your Settings pattern.

What I deliberately did NOT copy: the From / To boxes. On a product's availability card those two boxes are how a SEASON is said - a sell period genuinely has two ends. A delivery date is one day, not a stretch, so a From/To here would only ever have been a slower way to do what swiping now does.

16 Sep 2026 - engine v101 (no database setup needed). A price you can change on an order, and the money recorded as cash or TNG. Three things, all about money.

A price you can change on an order. Every item line on the + New order form and in the Edit pop-up now carries its selling price. It opens on what the product costs; type over it and THAT order is sold at your price - the confirmation, the payment and pickup reminders, the receipt, the customer's own track page and every money number in the app follow it. Your menu price is untouched, and a menu price you change next week never rewrites a sale already made. Blank keeps "whatever the product costs", which is how an unpriced product has always behaved. Both forms also show the items total as you go.

Cash or TNG, written down when the money lands. The Paid button is now **Paid · Cash** and **Paid · TNG** - one tap each. The order remembers which, and when, and shows it beside its status. The Note / tracking box (now **Note / tracking / payment**) carries a **Paid by** box as well, for an order you marked paid before you could tell, or one to correct later. None of this reaches the customer: their page just shows Paid.

Two places to check it against. The delivery day's own header on Orders now carries a short till - "Cash RM 95 · TNG RM 120 · 1 to collect". And More -> **Money** is a screen of its own, with **Today / This week / This month**: cash collected, TNG collected, paid but no method recorded, and what is still to collect with the number of orders. Collected money is counted by the day it LANDED, so a transfer that comes in today counts today even when the order delivers on Friday; what is still to collect is counted by delivery day, because that is the day you hand it over. An order paid before any of this existed shows under "Paid, no method" rather than being guessed at.

16 Sep 2026 - engine v100 (no database setup needed). Making delivery dates for a weekday is one tap again. One change, on More -> Delivery dates.

Why. You found the function missing: adding your delivery dates for, say, every Monday meant tapping each Monday by hand. Nothing had been removed - the app's only "do the dates for me" button lived on Home, inside the "No delivery dates yet" message. That message shows only while you have no upcoming dates, and pressing the button creates the very dates that hide it: it was a one-shot button that took itself off the page the moment it worked. So once you had dates, which is every day since your first week, it was never there again.

Three ways in now. On More -> Delivery dates: tap dates and **Add selected**, as before; or **tap a weekday**

letter - tap **M** and every Monday in the month shown is picked for you, tap it again to put them back, and the letter underlines solid once that whole weekday is picked; or press **Generate the next dates**, which is always there, follows the delivery days in your Settings, and adds the next six dates that are not on the calendar yet. Generate moves the calendar to the month the new dates start in, so you can see what it just did.

One wrong line fixed with it. Home's "no dates yet" message used to say "the next Mon/Wed/Fri dates" whatever your delivery days really were. It names your own days now (and so does the line under Generate).

15 Sep 2026 - engine v99 (no database setup needed). The customer can tap the "Next available" line and be given that day. One change, on the order page.

Why. A product you keep on the shop, on a day it cannot be ordered, tells the customer when it can be had - "Next available: Sat 19 Sep". Until now they had to go and find that date on the calendar themselves.

What it does now. The line is a control: tapping it takes that delivery day, so the product in front of them can be ordered straight away, and the page moves up to the calendar, where the chosen day is written out in words. It keeps its gentle pulse, and wears a small arrow so it reads as something to tap.

The one time it does not take the day. If the customer already has something in their basket, the line stays a plain label and tapping it answers instead: "Your basket is for Fri, 18 Sep. To order for another day, choose it on the calendar above." Taking a day then would silently move their whole order to a new date - and drop anything that does not fit there - which is not something a tap on a product line should ever do. Changing the day is what the calendar is for.

15 Sep 2026 - engine v98 (no database setup needed). The last status reads Collected / Shipped, and a placed order gains a Note / tracking button. Two small changes on the Orders screen.

One label for the last stage. v97 named it per order - Collected on one the customer fetches, Shipped on one you post. You asked for the pair itself instead, so the stage now reads **Collected / Shipped** everywhere it is named: the row's status list, the row's own journey map, the day's status filter, and the map on the customer's track page (in all three languages). Which message the row offers is still decided by the delivery method - a courier order shows **Send shipped message**, a self-collect one **Send pickup reminder**.

A way in for the two fields you reach for most. Beside Edit, every order now carries a **Note / tracking** button. It opens a small box with exactly two fields - the order's **Note** and the courier's **Courier tracking number** - and a Save. Send, and the note lands on the order, while the tracking number also goes onto the customer's track card and into the shipped message. Nothing else moved: **Edit** still opens the full form (the delivery day, the customer, the address, the items, the quantities) whenever you need it. The tracking box that sat on the row itself in v97 is gone - the button is the quicker, tidier way to the same field, and it works on any order rather than only a courier's.

15 Sep 2026 - engine v97 (one database line to run once). The last status is named for how the order leaves, and a posted order can carry its courier's tracking number. One change, across the order row, the WhatsApp messages and the customer's track page.

The last status is Collected or Shipped. The final step used to say "Delivered" for everything. It now reads **Collected** on an order the customer fetches from you and **Shipped** on one you post - the same step in the journey, named for what actually happened. That is true everywhere the status is named: the row's status list, the row's own journey map, the day's status filter, and the map the customer sees on their track page (in all three languages). Nothing about the step's behaviour changed, and every order already sitting at that step keeps its place - only the word is different.

A tracking number, typed where you post the parcel. A courier order shows a **Courier tracking number** box under its details from Packed onwards, and a **Send shipped message** button beside the status. Type the number the courier gave you, press the button, and WhatsApp opens with the message already drafted: the order number, what was sent, and the tracking number. The box is in the Edit pop-up too, for a number you need to fix or read back over the phone.

The customer sees it. The number is published with the order, so a posted order's track page shows it on its

own line under the delivery details - they can read it to the courier without messaging you. Self-collect orders show no such line. The shipped message is only offered on courier orders: a self-collect order keeps its **Send pickup reminder**, which is the message that fits that hand-over.

What you need to run once. The tracking number needs one column added to the tracking table - supabase/track_no.sql (a single ALTER, safe to re-run). Until that runs, everything else works; the number just reaches the message and not yet the customer's page.

15 Sep 2026 - engine v96 (no database setup needed). "Set day's availability" now shows how the day adds up. One change, on the pop-up you already use.

Why. v95 made the day's total count only the products you sell that day, but the pop-up that sets those numbers still showed nothing about where the total came from - and a product you do not sell that day sat in the list looking as if it were part of it.

What it does now. Under the product rows, the pop-up ends with **How the day adds up:** one line per product on sale that day with what it contributes, and a final line with the total - **the same number the order page uses for that day**. Type a + or a - and the lines and the total move as you type. Underneath, the pop-up names what it left out ("Not counted: Saturday loaf - not sold on this day, so no order can go on them here") and says what is already booked, so you can read how much the order page can still take. A product not sold that day keeps its row so you can still set its numbers; it simply does not enter the total.

One number, one place. That sum and the day's capacity are computed by the same function the shop's numbers come from, so what you read on screen and what your customers get cannot drift apart.

15 Sep 2026 - engine v95 (no database setup needed). A day's total now counts only the products you actually sell that day. One change, on your side of the app, with one knock-on for customers.

What it was doing. The "x/y" on a day - the chip beside the date on Orders, the same number on the day in the month calendar, and the dials on Home - was adding up the daily limit of every product on your menu, whether or not that product can be ordered on the day you were looking at. So a day with one product on it (a limit of 12, say) still read **1/42**: thirty of those units were limits belonging to products that could never take an order that day.

What it does now. Only the products on sale that day count. A Saturday-only loaf adds nothing to a Wednesday, and a product kept on the shop as **Unavailable** adds nothing to a day it is not sold on. The same day now reads **1/12** - what that day can really take. Products with no sell marks are unchanged (they sell every day, so they always count), and a day where nothing on sale has a daily limit falls back to the **default capacity** in Settings, exactly as before - it does not drop to 0, which would make the app call the day Sold out.

The knock-on worth knowing. This same number is what the app tells the order page about a day, so a day now turns FULL at the moment every unit you are actually selling that day is booked, rather than at some larger number padded out by products that were never on sale. Past every real slot being taken, the order page stops offering that day - which was always the intent, and now happens at the right moment rather than late.

Nothing else moved. Your over-capacity warning, the "N ordered · N left" line on a delivery date, the purchasing plan and each product's own "N left" stamp all read the same number they always did for products that are on sale.

15 Sep 2026 - engine v94 (no database setup needed). One sentence on the customer's greyed card, reworded. Wording only - nothing behaves differently.

On a product you keep on the shop, the line saying why it cannot be ordered today now reads **Only available on Mon.** where it read "Only sold on Mon." The old wording was the seller's voice on a customer's card, and "available" is what the rest of that same card already says ("Only available for delivery from 19 Sep"). The Chinese and Bahasa Malaysia say the same thing as before, in their own words.

15 Sep 2026 - engine v93 (no database setup needed). The track link now lands on the track card, lit. One change, on the customer's side of the shop.

Why. Your confirmation message ends with a Track your order link. Tapping it opened the order page at the very top, with the tracking card somewhere below - the customer had to scroll and hunt for the very thing they had

just tapped, on a page they have never seen before.

What it does now. The page comes up with the Track your order card already in view AND gently glowing, so it cannot be missed. The glow does not fade on a timer: it goes when their pointer or finger reaches the card, so it is still burning while they are looking around for it. That is the same rule your own app uses when you tap an order in the New orders box and it flashes the row - the flash is there until you arrive.

Nothing else changed. Typing a code into the box on the page behaves exactly as before (the card is already in front of them, so nothing moves), and a customer who simply opens the shop sees no glow at all.

15 Sep 2026 - engine v92 (no database setup needed). The New product card on Products folds away. One change, on one screen.

Why. The card that adds a product sat open at the top of the Products screen, and open it is a long form - it pushed your three lists (On the shop, Draft, Hidden) right off the bottom, so every look at what you sell began with a scroll past a form you were not filling in.

What it does now. The card arrives folded to a single line reading **+ New product**. Tap it and the form opens; tap the title again, or tap anywhere else on the screen, to fold it away. It stays open after you add a product - the toast still says it was saved as a draft - so you can type the next one straight away, and it starts folded again the next time you come to the screen. Nothing about adding a product changed: the form inside is exactly as it was, and it is still the same card that stays on the page while an Edit pop-up is open over it.

15 Sep 2026 - engine v91 (no database setup needed). The Orders calendar now answers a tap it cannot act on. One change, on one screen.

Tapping a day you do not deliver on. The month calendar at the top of Orders rings your delivery days in green with how booked each one is, and draws every other day of the month plain. A plain day never did anything before - tapping it was silent, which left you guessing at the reason. Now the calendar answers: a warm line appears under the grid naming the day and saying it is not a delivery day, and that delivery days are added in More -> Delivery dates. A marked day says its usual name above the day at the same time, so a holiday that is not a delivery date now tells you both what it is and why no order can go on it. Opening a real day, or tapping another plain one, replaces the line. Nothing can be ordered on a plain day, and nothing is created by the tap - the calendar simply says so now instead of staying quiet. A plain day already gone stays silent: there is nothing left to add to it.

Where it appears. The month calendar at the top of Orders, the day picker inside the + New order card, and the Delivery day picker in an order's Edit pop-up - every calendar in the app that offers only your delivery days.

Nothing else changed. Delivery dates are still made in exactly one place, More -> Delivery dates (or **Generate next delivery dates** on Home). Tapping a calendar elsewhere in the app does not create one, and a tap that could open a day before still opens it.

15 Sep 2026 - engine v90 (no database setup needed). A product can stay on the shop on a day it cannot be ordered, with the date a customer can next have it. One new switch, off unless you turn it on.

Keep it on the shop. A product's Availability card now opens with a switch at the very top, above the sell-day calendar. It is for the two or three hot items customers come back looking for - if it is not listed they wonder whether you still make it - but leave it on for everything and the menu fills up, so it is a decision you make one product at a time. The card's own title carries the state too, so you can read it while the card is still closed.

What the shop shows with it on. On a delivery day the product is not sold on, the card stays instead of vanishing: greyed, stamped **Unavailable** (a day you never sell it is not the same as a day it ran out), with the reason it already writes ("Only sold on Mon") and one more line - the next date it can be ordered. On a day it has sold out the card is exactly what it was, **Sold out** stamp and all, and gains that same line. The line reads **Next available: Sat 19 Sep**, and names how many are left that day when a daily limit is set; it pulses gently so it reads through the dimming, and sits still - simply lit - if your phone asks for reduced motion. Nothing on a greyed card can be ordered: the + and - do nothing, it never enters a basket, and a line already in a basket still leaves with a note when the day changes. The switch never overrides **Draft** or **Hidden** - those still take a product off the

shop entirely.

With the switch off, nothing at all changes. Off is the default and no product you already have carries the setting, so every product behaves exactly as it does today until you switch one on. One consequence worth knowing: on a day where every product is unavailable, a shop with kept items shows those greyed cards instead of the "nothing is on the menu" line.

14 Sep 2026 - engine v89 (no database setup needed). The Delivery Dates screen trades its list for the calendar: a ticked day comes off when you tap it again, and the past dates fold away into one group. Three small changes, all on one screen.

Tapping a ticked day takes it back off. A delivery date that was already on the calendar used to come off only from the list underneath it. Now tapping that green day takes it off on the spot - the same tap that put it there undoes it, so the calendar is the one place you both add and remove dates. If that day already has an order on it the app asks first ("2026-09-16 has 2 order(s) on it. Delete the date? The orders are kept in your delivery history."), because that is the one removal worth a second look; a day with nothing on it just goes, with a small "Delivery date removed" note. The date is all that is ever removed - the orders stay in your history either way.

The list of dates still to come is gone. The calendar already draws every delivery date as a green tick, and now it lets you take any one of them back, so the list below was repeating what was already in front of you. That is most of why this screen reads so much lighter. Nothing else moved: the green tick, the coloured holiday tags and the "N ordered · N left" line are all as they were.

The past dates are one group, folded away. Dates already gone keep their green tick but cannot be tapped, so nothing on a past day can be changed by a stray touch. They now sit together under a **Past dates (12)** heading that starts folded - tap the heading to open the list, tap it again to fold it back. Every past date is in there now: the old list showed only the ten most recent and left the rest out of reach. Each one still opens its orders or can be deleted exactly as before, with the same "Del" button and the same question when orders sit on it.

14 Sep 2026 - engine v88 (no database setup needed). The holiday tint sits on the day numbers at the Orders screen, and a mouse resting on a marked day names it. Two small things, both about the Orders screen and the bubble.

The tint now sits on the line of dates. On the Orders calendar each delivery day carries its booking count, and that count sat *under* the date number, which pushed the number 5.5px above the middle of its row. The occasion wash is centred on the row, so on this one screen the tint sat lower than the numbers it was meant to be washing - and the dates themselves did not line up with each other. The count now sits at the foot of the day, every date number sits on one line, and the tint runs through that line. Nothing else about the calendar moved: the green ring, the chosen day's fill, and the counts all read exactly as before.

On a computer the name comes up on hover. v87's bubble needed a tap. With a mouse, resting the pointer on a marked day now shows its name on its own, exactly as the customer's page does. This is deliberately only for a pointer that can hover: a touch screen keeps the tap, so a finger sweeping a run of days across the Availability or Delivery Dates calendars never drags a bubble along with it.

14 Sep 2026 - engine v87 (no database setup needed). Tap a marked day and it tells you its name - on every calendar in the app. v85 drew your marks on all five of the app's month calendars, but only as colour: a wash and a band, with nothing to say which holiday it was. The customer's own shop page has answered this since its calendar arrived - tap a marked day and a small dark bubble above it says the name - so the app now does the same thing, with the same look.

A tap names the day and still does that day's usual job. The bubble appears above the day ("Malaysia Day") and goes away as soon as you tap anywhere else. Nothing else about the tap changed: on the Orders calendar it still opens the day, on a product's Availability calendar it still marks or unmarks the sell day, on Delivery Dates it still starts the range you are about to mark. A day that sits inside two marks takes the name of the shorter one - the same rule the colours already follow, so the name and the colour always belong to the same mark. A day already gone is never named, because no mark is drawn on a past day either.

Two kinds of day could not be read at all before, and now can. On the Orders calendar a marked day you do

not deliver was a quiet number with nothing to tap, and on the Delivery Dates calendar a marked day that is *already* one of your delivery dates could not be tapped either - its tick is removed from the list below the calendar, never by tapping its square, so its tap was free. Both now say their name when tapped, and nothing else about them moved: a Delivery Dates square still never ticks or unt ticks a date.

The Order date boxes name their day above the date, not above a day of the month. Those two fields fold their calendar the moment a day is picked, so a bubble drawn on the day would vanish with it. The name sits above the date the box is showing instead, which also means an order recorded on a public holiday says so for as long as that date is on the field.

Nothing was added, moved or removed, and the shop is not touched: this is the same bubble the shop's calendar has been showing, arriving on your own calendars.

14 Sep 2026 - engine v86 (no database setup needed). On the Orders screen your marked days now sit on their days, at the same depth as everywhere else. v85 put your marks on every calendar in the app. On the Orders screen the wash did not line up with the days underneath it. Two things were behind that, and both came from the same fact: the days on that one screen are taller than the days anywhere else, because a delivery day there carries its booking count ("3/12") under its number.

A holiday band was stretched to fill its week row. Where every day is the app's usual height, a band came out the same 30px sheet a single-day box is, so the two shapes matched - which is how the rest of the app has always drawn them. On the Orders screen the rows are taller, so the band there came out deeper than a single-day box, and deeper in the weeks that held a delivery day than in the weeks that did not, so one holiday's own band changed depth from row to row inside the same month. **A band is now a sheet of one depth, centred on its row** - the same depth a single-day box has, on every calendar, whatever height the days around it happen to be.

A plain day kept the app's usual height inside a taller row, so it sat at the top of the row rather than on its middle. A single-day mark was therefore drawn 6px higher on a day you do not deliver than on a day you do: the same holiday, on a different line, decided by whether you deliver that day. **Every day of the Orders month is now drawn at one height**, delivering or not, so a mark sits on the same line as the day it covers, beside the booking counts.

One more thing fell out of that. Today's soft glow is drawn around a whole day on the app's shorter calendars, and around a taller day it would have become a box again - the very thing that glow replaced. **Today on the Orders screen now glows around its number**, the same round glow a delivery day's number carries, so today reads the same whether or not you deliver that day.

Nothing else changed. The shop's calendar and your Delivery Dates calendar look exactly as before (their days are all one height, so they never had this), and no mark was added, moved or removed: this is the drawing only.

14 Sep 2026 - engine v85 (no database setup needed). Your marked days are now drawn on every calendar in the app, not just the one you marked them on. Until now a holiday only appeared on the customer's calendar and on the Delivery Dates screen you marked it on. Every other calendar in the app drew a plain month: the month calendar at the top of the Orders screen, the little **Order date** calendars in the new-order form and in the Edit-order pop-up, the delivery-day picker inside that pop-up, and a product's **Availability** calendar where you mark the days it sells. So a holiday you were planning around was invisible on every screen you actually plan on.

Now all of them draw your marks, in the same two shapes and the same colours as the Delivery Dates calendar - a see-through band across a run of days, deeper the shorter the run, and a box for a single day. A mark on a day you do not deliver is drawn too, on a dimmed day, since that is the only way a holiday on a quiet day can be seen at all. Nothing about the marks themselves changed: they are still only ever something to see, and a mark still never adds or removes a delivery date or changes what a product sells. Two small pieces of the drawing were tidied while moving it: the green tint on a day a product sells is now a see-through wash rather than a flat fill, so a holiday band running behind it still shows through instead of the mark disappearing on exactly the days you sell (the tint itself looks the same as before); and a marked day that happens to be today now

keeps the soft glow that says "this is today", the way the shop's calendar has always drawn it.

14 Sep 2026 - engine v84 (no database setup needed). A day you marked is now the same gentle wash everywhere, whichever calendar you are looking at. A marked day was drawn one of two ways, decided by a single question: is this a day you deliver for? On a day you deliver for the mark went pale, so the green "you can order" ring stayed the first thing the eye read. On a day you do not, the mark was drawn at full colour, as a solid block with the number turned white. September's Malaysia Day sat on a delivery day and read as a quiet tint; October's World Animal Day fell past the delivery days published so far and read as a loud solid block - so the same kind of holiday looked like two different things, and a mark quietly changed its own appearance as the delivery window moved towards it.

Now a mark is only ever a see-through wash of its own colour, and its depth follows how long the mark runs: a one-day holiday the deepest, a long school break the palest. The customer's calendar and your own Delivery Dates calendar read that depth from the same place, so a day wears exactly the same mark on both sides of the shop and in every month. Nothing is ever a solid block of colour any more, which also means the date number and the green delivery pill always read on top of a mark, whether or not you deliver that day. Only the look changed: marking, removing, renaming and loading standard occasions all work exactly as before, and a mark still never adds or removes a delivery date.

14 Sep 2026 - engine v83 (no database setup needed). Marking a holiday now tells the shop straight away. A day you mark on the Delivery Dates calendar is drawn on the customer's own calendar (engine v79 to v81). But saving a mark only ever sent your delivery dates to the cloud - it never sent the customer's page anything - so the marks reached the shop only by accident, whenever you next happened to save a product or a setting. Mark the holiday, open the shop, and nothing had changed. Found while you were looking for a holiday on the shop and finding none.

Now the same save does both. Marking a day, removing a mark, loading a set from the standard list, or renaming or recolouring one all republish the customer's page within a couple of seconds, exactly the way editing a product does. Nothing else about a mark changed: it is still only ever something to see while planning, it still never adds or removes a delivery date, and a name you typed yourself still never leaves your phone.

One tap to catch up the marks you already have. The fix sends the marks from now on; the ones already on your calendar have never been sent at all, so open More -> Settings -> Storefront and tap **Publish now** once. After that the shop's calendar follows your marks by itself. Nothing to run in Supabase.

14 Sep 2026 - engine v82 (no database setup needed). A product's selling days are now marked on a little calendar inside the product's own screen - weekend, certain days, a season, or any mix of them. Until today a product could say only two things about *when*: a notice period, and one from-date and to-date for a season. You asked to set it freely instead, so the two date boxes have become a card you mark.

Open a product and tap the Availability card. It sits under Daily limit and stays folded until you tap its title, exactly like the product-text card you already know. Folded, its title tells you what is marked - "Every day", or "Sat & Sun", or "1-24 Dec 2026" - so the Products screen alone tells you a product's selling days without opening anything.

Inside, it is a month calendar and four gestures. Tap a weekday letter (M, T, W...) and every one of that weekday **in the month shown** is marked: tap S twice and the product sells on Saturdays and Sundays. Tap one day to mark just that day. Slide your finger across several days and the whole run is marked. Tapping or sliding over a day that is already marked takes it back. Everything you mark is added together - mark Saturday and Sunday, then a week in the middle, and the product sells on all of them.

Nothing carries over into the next month, which is the point of it: marks belong to the days you marked and no others, so a December-only set stays December-only and January starts with nothing marked. Page to another month with the arrows and mark that one too if you want to sell then. And a product whose card you have never opened sells on every delivery day exactly as before, so nothing on your live shop moves until you mark it.

A period that crosses a month or a year is one mark with two ends you can stretch. The card lists your

marks under the calendar; tap one and its **Starts** and **Ends** dates appear above it, and you can push either end as far as you like - a Sat & Sun run from 1 Dec to 4 Jan is one line, not two marks. Leave an end empty and the mark simply has no bound on that side ("from here on", or "up to here"). Type the two dates the wrong way round and they are put in order for you. A small x on a listed mark removes it whole.

On the order page, a day a product is not sold for is not a note - the product is simply not there. No card, nothing to want and not have. If a customer already had it in their basket and then switches to a day it is not sold for, it leaves the basket and a short line says why. The one rule that keeps its note is the notice period: a product needing 14 days IS sold that day, it only has to be ordered earlier, so it stays on the menu reading "Orders close 14 days before delivery - pick a later date" - the two are different questions, and only the second one hides anything.

Two things worth saying plainly. The old From/To season boxes are gone; a product that had them opens with that period already sitting in the card as one mark, so nothing you set before is lost. And a product whose selling days have all gone by keeps them, on purpose: your own rule is "if I have not indicated a selling date, it is not selling", so a mark you leave on says exactly that, and the way to put a product back on sale every day is to take its marks off with the x. No database setup, and delivery dates, capacity, value packs, ingredients, suppliers and everything else are untouched.

14 Sep 2026 - engine v81 (no database setup needed). A holiday you marked is now drawn on the shop's calendar exactly the way your own calendar draws it: a pale band across a run of days, a solid box for a single day. The marks were already reaching the shop; it was only their shape that differed, and a mark that looks different on the customer's page than on yours is a mark your customer has to learn twice.

On the shop, a marked day now wears the mark your own calendar gives it. The shop had been drawing each marked day as its own flat pale square, so a holiday running over a week read as a row of separate blocks where your own calendar shows one continuous shape. Now a holiday running over **several days** is a **pale see-through band** across those days - one unbroken rounded band per week row it crosses, in its own colour, and **deeper the shorter the run**, so a one-day holiday still stands out inside a long school break. A **single-day** holiday is a **solid box** in its full colour on its one day, a touch narrower than the band so a longer band still peeks out at its sides. The eight colours are your own calendar's eight, and the strength steps are the same three, so a marked day reads the same to your customer as it does to you.

The one place it gives way: a day you deliver for. Green is the colour the delivery ring uses, and a solid box would swallow that ring - so on a day the bakery delivers for, the day a customer can actually order, the box goes pale there and the **green ring stays the first thing the eye reads**. The mark is still there on that day, and the day is still named on a tap. This is the one guard on the shop's calendar, and it covers every mark colour rather than green alone, so no colour can ever hide the days that can be ordered for.

Nothing else moves. Tap a marked day - or, on a computer, rest the mouse pointer on it - and the same small bubble names it, e.g. "Malaysia Day"; tapping anywhere else puts the bubble away, and still nothing is listed under the grid. Today still breathes its soft glow. A run still stops at today rather than colouring days already gone, exactly as your own calendar leaves past days alone. And your own typed-in days still never reach the shop: only a day that is, by name and date, one of the built-in standard days can be published, exactly as before. No SQL, no schema change. The calendar the shop draws with is still a copy of your app's own, and a test now pins the two copies' mark rules together as well as their month grids, so a day can never again look like one thing here and something else there.

14 Sep 2026 - engine v80 (no database setup needed). Today glows instead of wearing a box, a holiday you marked is a soft tint you can tap to name, and the Orders screen's sideways date strip becomes the same month calendar the shop shows. Three small things, all pulling the same way: a calendar should read the same on both sides of the shop, and nothing drawn on it should look like something it is not.

On the shop, today no longer wears a little red rectangle. It always meant to be a ring around the date, but on a day you do not deliver the app had no circle to draw that ring on, so it traced the day's number as a box instead - which is why it looked like a red square on the quiet days and only looked right on the days you deliver

for. Today is now a **gentle glow** that breathes around the number, with no box at all, and it sits happily on a day you deliver for too: the **green ring** says you can order for this day, the soft brown glow says it is today. Under **prefers-reduced-motion** (a phone set to calm its animations) the glow holds steady instead of pulsing, so it is still there to find. The same fix reaches the app: the today marker on the Delivery Dates calendar, and inside every date picker, was a hard outline and is now the same soft glow.

A holiday you marked is now a soft wash, and its name comes out on a tap. The shop's calendar used to mark a standard day with a small coloured dot and then list the month's names in a "Holidays" line under the grid - which crowded the very line that line was there to keep clear, and made a holiday compete with the days a customer can actually order for. A marked day now wears a **pale wash** of its own colour over the whole day. The wash is deliberately neither a ring nor a filled dot, because a ring is what "you can order for this day" looks like, and the two must never be read for each other. Tap a tinted day - or, on a computer, rest the mouse pointer on it - and a small dark **bubble** above it names the day, e.g. "Malaysia Day". Tapping anywhere else puts the bubble away. The "Holidays" line under the grid is gone. Your own typed-in days still never reach the shop: only a day that is, by name and date, one of the built-in standard days can be published, exactly as before.

In the app the same holiday is named **beside the day** instead of in a list: open a delivery day and, if it carries a mark, its name sits under the date in the mark's own colour. Nothing at all is drawn on a day you have not marked.

The Orders screen's date strip is now a month calendar. The row of date chips at the top of Orders grew a little longer every week. It is now the same **month calendar** the customer sees - one month at a time, with left and right arrow buttons either side of the month's name. The days you deliver for are ringed in **green**, and under each one is **how booked it is**: the same "3/12" the Home dials show, or **FULL** when the day has no room left. A day whose orders have already closed shows its count in **red**; a day already gone is dimmed but still opens, because you backfill and look back at old days. Every other day of the month is plain and cannot be tapped. Tap a green day and the app's view moves to it - that day is filled in on the calendar and its orders appear below - and tapping a day in the list below moves the calendar with it, so the two always agree. The arrows reach only the months your delivery days fall in, so paging never strands you on an empty month.

The + New order card arrives folded, and leads with the calendar. It used to sit open on every delivery day, which is not what that screen is for day to day. It now starts **shut**, showing just its title, "+ New order" - tap the title to open it, and tap the title again (or tap anywhere else on the screen) to fold it away. Opened, it reads in the order you think in: **the month calendar first**, the same one the shop shows, already sitting on the day you are looking at, so tapping another green day moves the whole screen before you have typed anything; **then the customer's details** - name, WhatsApp number, self-collect or courier, a note; **then the items**, with "+ Add another item" and "+ Add order" at the end. It stays open after you add an order, so you can type the next one straight away, and coming back to the screen later starts it folded again.

The Edit pop-up now reads the same way. Its **Delivery day** control was a button you had to tap to reveal a calendar; the calendar is simply there now, already open on the order's current day. The pop-up then reads the day, then the customer's details, then the items - the same order as the card. Every day still to come is tappable, plus the order's own day even if it has passed, so an old order always shows where it is, and the soft notes re-read underneath the moment you pick a different day. Moving an order is unchanged in every way that matters: still one delivery day per order, and capacity, the confirmation and the payment reminder, the track page and the WhatsApp messages are all exactly as they were. The button-style day picker the app used in these two places has been deleted along with its tests, because nothing calls it any more.

14 Sep 2026 - engine v79 (no database setup needed). The customer picks a delivery day on a calendar, and so do you in the app. The shop used to show its delivery days as a row of date chips laid out sideways. Now, under **Pick a delivery day**, the customer sees a **month calendar** - one month at a time, with the arrow buttons either side of the month's name to look ahead or back. The days you are taking orders for are picked out in **green**; every other day of the month is plain and cannot be tapped. They tap a green day, and the day they chose is written out in words just underneath - **"Your delivery day: Wed, 16 Sep"**. That is the one line of text below the calendar, exactly as asked: no dates floating outside the grid, and still **one delivery day per order**.

The first open day is already chosen for them when the page opens, so ordering can never be blocked by forgetting to tap. Nothing else about taking orders changed: a day you have cut off is still not shown at all, and a day that is full still greys out with a red **Sold out** badge and cannot be tapped.

Only the standard days you have marked reach the customer. If you have loaded days from the built-in list onto your own Delivery Dates calendar, those days now wear a small coloured dot on the customer's calendar, and a short **Holidays** line under the grid names the month's ones (e.g. "16 Sep · Malaysia Day"). They follow the colour you gave them. A day **you typed yourself** - a birthday, a promo, a school break - never appears: only a day that is, by name and date, one of the built-in standard days can be published, so your own notes stay private to your phone. Mark nothing in the app and the shop's calendar simply carries no dots and no caption.

In the app, the same little calendar now stands wherever you name a date. Three places changed. In the **Edit** pop-up, **Delivery day** is no longer a drop-down list: it names the order's current day, and tapping it opens a small month calendar beneath it - every day still to come is tappable, plus the order's own day even if it has passed, so an old order always shows where it is. The **Order date** box in that pop-up is a calendar too, with a **Today** shortcut for an order you are typing in right now. And the **+ New order** card has gained a **Delivery day** picker at its top: tap another green day and the whole screen moves to that day, so the product list you are choosing from is always the one that is actually sellable for the day you picked. The soft notes that help you decide - the change/cancel window, a warning, a closed day, too little left - all still sit under the calendar, and none of them stop you.

The calendar expands in place under the button rather than opening in its own window, because the app's pop-ups share one layer and a calendar opened inside one would wipe out the Edit pop-up it was opened from; expanding in place also reads better on a phone. No database setup needed - nothing about how orders are stored, counted, bought for or messaged changed.

13 Sep 2026 - engine v78 (no database setup needed). The two confusing "Add" buttons on the Delivery Dates calendar now say what they do. On the Delivery Dates screen, **Mark an occasion** used to carry a small button reading **"+ Add occasion"**. The name did not tell you that it opens a ready-made list of Malaysia's holidays and fun days, all ticked and waiting - so you had to open it to find out. It now reads **"+ Load standard occasions"**, and the window it opens is titled the same.

Inside that window there were **two buttons both called "Add"** - the big one at the bottom that files the days you ticked, and a smaller one in the **My own day** strip just above it, which files a day you name yourself (a birthday, a one-off promo). Reaching for "Add" and hitting the wrong one gave you **"Type a name first"**, for a day you never meant to name. The smaller one is now spelled out as **"Add my own day"**, so the only plain **Add** in the window is the one that adds the ticked days.

Nothing else moved. Ticking a row and pressing the big **Add** still files exactly what you ticked and nothing you did not; **Untick all** / **Tick all** and the heading tick boxes work as before; the **My own day** box still takes any name and any date. No database setup needed.

13 Sep 2026 - engine v77 (no database setup needed). The greyed line in the translated-text card now says what happens if you leave it blank. Engine v76 turned the product screen's Mandarin / Bahasa Malaysia text into one folding card, and each empty line showed its translation on offer. That greyed line read "e.g. <the translation> - blank keeps English". It now reads **"e.g. <the translation>.....if blank, it will be filled with English"** - the same words you asked for, saying plainly what the customer gets rather than asking you to work it out. The line in the card's own explanation says the same. Deleting what is in a line brings the greyed line and its arrow straight back, exactly as before. Nothing else moved: the card still folds away when you tap outside, the -> still takes the words, the circular-arrow button still translates that line again.

13 Sep 2026 - engine v76 (no database setup needed). The translated text on a product is now one quiet card that folds away, filled in for you one line at a time. The **Product text for your customers** part of the New/Edit product screen - the same words in **Mandarin** and **Bahasa Malaysia** - was the busiest corner of the whole app. It was always open, so a long screen got longer before you needed any of it; eight lines each carried its own **"Translate this one"** button; two **"Fill all"** buttons sat above them; and a little **"auto"** tag marked the

machine-made lines, in a word you had to remember the meaning of.

Now the card starts **shut** and stays shut - it is the setup part of the screen, not the part you touch every day. **Tap its title to open it.** Opened, each line already has its translation worked out and waiting as the same pale grey suggestion you know from every other box, with the little **->** at its right edge: **tap the -> and the suggested words drop in** - exactly the gesture you already use. The moment you take them, the arrow is replaced by a small **circular-arrow button** - tap it and that one line is translated again on the spot, which is what the old "Translate this one" button did. **Type your own words over any line and the circular-arrow button goes away** for that line: it is yours, and it is never overwritten. **Leave a line blank** and that language keeps the English, the same as before. And when you are done, **tap anywhere outside the card and it folds away again.**

The **"Fill all Mandarin" / "Fill all Bahasa Malaysia"** buttons are gone, the eight **"Translate this one"** buttons are gone, and there is no **"auto"** tag to read any more - a line showing the circular-arrow button is machine text, and a line with no button is yours. Nothing about what your customers see has changed, and nothing needs setting up: a product is still translated for the shop the first time you save it and again when you Publish, exactly as before, so a product you never open this card for still reads in all three languages on the order page.

12 Sep 2026 - the TNG QR box no longer claims to be on the order page (no new engine; your phones did not change). The box in **Settings -> Storefront** holding your Touch 'n Go QR image link said it was **"shown on the customer's track page"**. It never was - your payment QR reaches the customer in the **WhatsApp confirmation** and the **payment reminder**, and nowhere else. The help text now says that, and the order page's own settings no longer carry the QR at all, so it cannot be drawn there by accident. Nothing about how you take payment changed, and nothing about how the app works changed.

12 Sep 2026 - engine v75 (no database setup needed). The glow on an order you have just jumped to now stays until you reach it. Tapping a row in the **New Orders** inbox (or a result in **Find an order**) takes you to that order's delivery day and marks the order with a green glow - but the glow used to fade after less than two seconds, which is no time at all when your eye is still travelling down a long day, and you were left hunting for a row that was no longer lit. **The glow now keeps pulsing until you get to the row.** It stops the moment your finger or the mouse arrives on that order, or when you tap it open - so it is waiting for you however long you take, and it never lingers as a puzzle later. (If a phone has **Reduce Motion** turned on in its accessibility settings, the same green outline holds steady instead of pulsing - it still stays until you reach the row.)

12 Sep 2026 - engine v74 (no database setup needed). The little arrow in an empty box now keeps working. Engine v73 added the small right-arrow that sits at the right edge of a box showing a greyed example: press the **-> key** on a keyboard, or tap that arrow, and the example drops in as real text. On a phone that tap worked the first time and could then go dead - once you had tapped into a box, tapping the arrow on the next one did nothing.

Two things were behind it, and both are fixed. The tap was measured against the **page**, and the phone's keyboard shifts the page under your finger as soon as any box is focused, so the arrow was no longer where it had been. And taking a suggestion **opened the keyboard itself**, so the page jumped after the very first tap. The tap is now measured against the **box itself**, taking a suggestion no longer opens the keyboard at all, and the strip you can tap is about half as wide again - so a thumb that lands a little wide of the arrow still counts. A tap on the greyed text stays an ordinary tap that just puts your cursor there. **The arrow looks exactly as it did**, and nothing else changed.

12 Sep 2026 - a new order now opens on the order itself (no new engine; nothing to set up). Tapping a row in the **New Orders** inbox took you to that order's delivery day - and stopped there. On a busy day, with a long list of orders, you then had to read down the list to find the very order you had just tapped, and it was easy to open the wrong one.

The app now goes the whole way: it opens the delivery day, **scrolls the order to the middle of the screen and flashes it** for a moment, exactly as the **Find an order** search box already did. If the list had been narrowed with the status filter so the order was hidden, that filter is cleared first, so the order is always there waiting for you.

Nothing else on the Orders screen changed.

12 Sep 2026 - engine v73 (no database setup needed). Five changes, all about the shop's promise to the customer and the baker's control over it.

Your rule is that an order is **not cancellable and money is not refunded, but it can be moved to another day**, and a customer may ask until a set number of days before delivery. That number now lives **per product**. Open a product in **Products** and, under its daily limit, a new box **"Changes or cancellations (days before delivery)"** - type 2, say, and that product's card on the shop reads **"Change or cancel up to 2 days before delivery."** It only ever *tells* the customer; it never blocks you, and you still move any order by hand. Leave it blank (or 0) and no note shows for that product. When one order holds several products with different numbers, the customer sees the **strictest** one - the largest, i.e. the earliest cut-off - so a mixed basket always gives one clear figure. The same sentence appears on the green **"Order received!"** card just after they place the order.

Second, **you can now move an order to another delivery day in the app.** Open an order with **Edit** and, in the pop-up, a new **"Delivery day"** box lists every day still to come, soonest first. Pick the new day and **Save changes**: the whole order moves as one - items, customer and the price it was sold at all stay put - and the customer's **"Track your order"** page follows to the new day straight away. Under the box are quiet notes to help you decide: the order's own change/cancel window, a reminder when the new day falls **inside** it, a note if that day's orders have already closed, and a note if that day does not have enough left for an item. None of them stop you - they only tell you - and the capacity check now counts the order on the **new** day, so you can see at once if that day would go over. This is the **only** way to change an order's day: never delete an order and re-type it, because deleting reads as a cancellation. A day left empty stays in your list - to remove it, use **Del** on the Deliveries screen as before.

Third, **a Policies box** for your own cancellation and refund wording. It sits in **More -> Settings -> Storefront**. Type it once in English, tap **Translate**, and the Mandarin and Bahasa Malaysia boxes are filled for you (edit either and it becomes yours). It appears on the shop **under Track your order**, in the customer's own language and with the line breaks you typed, and it stays hidden until you write something.

Fourth, and everywhere in the app: **an empty box showing a greyed suggestion can now be accepted with one tap - the arrow.** A small right-arrow sits at the right edge of such boxes; press the **right-arrow key** or tap that arrow and the suggestion drops in as real text you can edit - the same gesture as accepting a suggestion in an AI chat. It is offered only where the grey text is a genuine recommendation; the sign-in and cloud boxes never offer it, so a made-up key or password can never be accepted by accident.

No database setup and no new SQL - the new product box and the Policies wording travel with your existing shared data.

9 Sep 2026 - the wish list's automatic email went live (no new engine; the phones did not change). Adding a wish on **More -> Software wish list** now emails the **full wish list** - every wish, ticked or not, newest first, with the app version and the date - to the developer address(es) you set under Settings -> **Website & developer**, all on its own, whenever the phone is signed in and sharing data. Before today that automatic email needed a behind-the-scenes service that was not switched on yet, so only the **"Email the full wish list"** row sent it by hand; that row is still there as a backup any time.

10 Sep 2026 - the homepage opens faster (no new engine; the phones did not change). Two things about the public homepage were quietly costing every visitor time, and both are fixed.

The homepage's six photographs were stored *inside* the page itself, so the page was **249 KB before a single picture had even been asked for** - and every visitor downloaded all of it, on every visit, before they could see anything. The photos are now six separate files that the browser fetches and **keeps for next time**, which brings the page itself down to **22.7 KB**. The pictures are byte-for-byte the same ones, so the homepage looks exactly as it did, and each one now states its size up front so nothing jumps around as the page loads.

Second, the "What customers say" section used to fetch **every published review photo up front** - including the reviews a visitor never waited around to see. A photo is now fetched **as its review comes around**, one step ahead of time so it is ready before it slides in, and the reviews a visitor never reaches are never downloaded at

all. On a phone that means less data and a quicker page, and it changes nothing about how the section looks or behaves.

A review photo is also shrunk a little harder before it is stored - to 1000 pixels across instead of 1600. The card on the homepage shows a picture no more than 340 pixels high, so the smaller copy looks just as sharp while being about **a third of the size**, which is kinder to the customer's mobile data and keeps your stored photos light. This applies to photos left from now on; the reviews already on your page keep the picture they were stored with.

Your sign-off line - **Jien Luv 2 Bake | Small Batch. Big Love.** - has moved from the very bottom of the homepage footer, where it sat shoulder to shoulder with the developer credit, up to the **top of the footer**, above your contact details and set off by a thin line. It reads a touch larger now, and in whichever language the visitor has chosen, like the rest of the footer. The "Website by ..." credit stays where it was and now stands on its own.

10 Sep 2026 - the homepage's language buttons can no longer go dead (no new engine; the phones did not change). On some phones the homepage's **EN / Mandarin / BM** buttons could stop responding - the page would sit in English and tapping Mandarin did nothing - often after the page had been left alone for a while, or refreshed. The culprit was a leftover background worker from when the backoffice was served from the top of the website instead of from /admin/. (A background worker is the machinery the app installs so it can open without a connection.) That old worker had cached the whole website, and whenever something failed to download it answered with whatever it had - including handing the browser **a web page where it had asked for a script**. A page that receives a web page instead of its script never runs it, so the code that wires up the language buttons never ran; the English words you saw are written into the page itself, which is why the page still looked perfectly normal.

The old worker is now cleared away wherever it is still installed: a small cleanup file sits at the same address the old worker used, so the next visit replaces it, empties everything it had cached, and removes itself. The backoffice's own offline cache is deliberately left untouched, because the app needs it to open without a connection. The app's worker has also been taught, for good, never to answer a script or a picture with a web page - so the same fault cannot come back on any part of the site.

v72 - The shop's sold-out notes read in the customer's language (10 Sep 2026)

The shop page switches language in place, but two corners of it had stayed English no matter which language a visitor chose: the short note on a product that **cannot be ordered for the chosen day** (that it only opens from a certain date, only runs up to one, or that orders close so many days ahead - with the date written in the page's language), and the notes that appear above the menu when the page **has to change a basket** - an item that just sold out, or a quantity trimmed to what is actually left.

Both now read in **English, Mandarin or Bahasa Malaysia** along with everything else. The reason a product is closed is now handed to the page as a plain fact - which kind of rule, and the day or number of days behind it - and the page writes the sentence itself, so no English can slip into a Chinese or Malay page. An English visit is word for word what it was.

v71 - The shop's Mandarin / BM buttons answer instantly (10 Sep 2026)

On the shop page, tapping **EN / Mandarin / BM** used to reload the whole page. That re-downloaded the page and fetched the menu, the "left today" numbers and your saved storefront settings all over again, so each tap sat there for a moment before anything changed - and the slower the phone's connection, the longer the wait. A reload could also wipe a basket that was already half filled.

Now the shop switches **in place**, exactly like the homepage already did: the words change the moment you tap, and nothing is downloaded again. Whatever the customer has already done is left untouched - the items in their basket, the delivery day they picked, the name, WhatsApp number and address they typed - and if they are looking up an order, that card re-reads in the new language too. Nothing to set up.

v70 - An order keeps the price it was sold at (10 Sep 2026)

Until now an order only remembered **which product** was bought, not what it cost or what it was called. Everything - the amount on the customer's confirmation, the total on their own "Track your order" page, their lifetime spend on the customer list, the Home estimate - was worked out from **today's** menu. So the moment you re-priced or renamed a product, every past order silently moved with it: last month's order suddenly read RM22 instead of the RM15 the customer actually paid, and a renamed loaf rewrote what people had already bought.

Now an order is a record of a sale. When an order comes in - from the shop or typed in by hand - it keeps the **name and the price it was sold at**, and those are what every message, the tracking page and the customer's spend use. Rename or re-price a product today and yesterday's orders stay exactly as they were. A line you deliberately swap to a different product when you **Edit** an order does pick up that product's current price - that is a new sale - but a line you leave alone keeps its old price.

Existing orders (taken before this version) are stamped once with today's values, which is exactly what they are showing right now, so they stop drifting from here on. Nothing to set up; the phones pick it up with the next sync.

v69 - Either place updates the customer (10 Sep 2026)

Yesterday's fix put a customer's name and number back in one place - but it made the **customer card** the side that always won, so if the card and an order disagreed, the card's name was the one shown. That is no longer the rule: now it is simply **whichever you edited last**. Correct a customer's name with **Edit** on one of their orders and that is the name the list, the history pop-up and the "Hi {name}!" greeting use; open their card, type the name and save, and the card's is the one used. Either place updates the customer - neither is the boss. They normally read the same anyway, because each save writes through to the other side; this only decides the rare case where an old copy and a newer one are both sitting there. Nothing to set up; the phones pick it up with the next sync.

v68 - One name and one number per customer (10 Sep 2026)

A customer's name and WhatsApp number used to be held in two places at once - on their orders, and on their customer card - and the two did not always agree. If you opened a customer and typed their name on their card, the customer list could carry on showing the old one, because it read the name off the order. And renaming someone who has **no WhatsApp number** on file could lose them: the name is what ties a numberless customer to their history, so the new name slid off their orders and appeared nowhere at all.

Now the name and number you save on a customer are the ones that count, and fixing them once fixes them everywhere: saving their card writes the name and number onto **every order they have**, so the labels, the WhatsApp messages and the customer list all follow. It works the other way round too - correct a wrong number with **Edit** on an order and that customer's card updates as well, so the two can never drift apart again. Correcting a number carries that customer's bring-a-friend credits across with it, and renaming someone who has no number no longer loses them.

Leaving a box empty means "leave this as it is" - it never wipes a number your confirmations depend on. And if you renamed anyone before today and it never stuck, the app puts their orders right the first time this version opens. Engine v68, guide v68.

v67 - Naming a "No name" customer (10 Sep 2026)

A customer who reached you without a name - a storefront order that carried none - used to stay **"No name"** in your customer book even after you opened them and typed a name. Now it sticks: open them, tap **Add details**, type the name (and the dog's name, what they like, and so on) and save, and their row shows that name straight

away - in the list, the finder and the 'Hi {name}!' WhatsApp greeting. A tidy-up; nothing else about your customers changes. Engine v67, guide v67.

v66 - Auto-translate product text, plus Draft / Publish / Hidden states (9 Sep 2026)

Product **description, selling unit word** (the "RM15.00 / loaf" bit), **serving tip** and now the **name itself** read in English, Mandarin or Bahasa Malaysia automatically - a free translator fills the Mandarin/BM boxes for you the first time you save or publish a product. A small "auto" tag marks a machine translation; if one "sounds weird", open the product's edit screen, type over that box (it then becomes yours and is never overwritten), or tap a "Fill all Mandarin" / "Fill all Bahasa Malaysia" button or the circular-arrow button next to a single line to have it translated again now. Good translations are simply left alone, and you never hand-fill a whole product.

Products also gain a **third state**. A new product now starts as a **Draft** - fully built with its recipe, but **not on the shop and not orderable** until you **Publish** it. The Products screen shows **three separate lists**: **On the shop**, **Draft - not on the shop yet** (each with a Publish button) and **Hidden - taken down** (with an Unhide button). Hiding a product still keeps its history, recipe and purchase orders.

The **"Copy follow-up"** bring-a-friend message gains a small **EN / Mandarin / BM** choice, so the whole check-in you paste into WhatsApp - including the product's serving tip - can be in your customer's language. No SQL this version.

v65 - Reviews preview + "waiting to publish" cue (9 Sep 2026)

More -> **Reviews** now shows each review **exactly as customers see it** on the homepage - the photo, the stars, the message, then the reviewer's name and the language and date, in the same warm card - and they play **one at a time**, like the homepage's "What customers say" carousel. The reviews waiting for your **Publish** come first; once you have worked through them, the ones already on the homepage follow. It never slides on by itself while you are deciding - a swipe, the arrow buttons, or a dot under the card moves it on, and Publish / Take down / Delete moves you straight to the next review. A line at the very top of the screen says how many are waiting ("3 reviews waiting for you to publish"). The **Reviews** row on the **More** menu shows how many are waiting too, and the app's **Home** screen gains a "**N new review(s) to publish**" card that opens the Reviews screen in one tap - so a new review is never left waiting unseen. No SQL this version.

v64 - Review carousel, EN/Mandarin/BM site, developer contact + wish email (9 Sep 2026)

The homepage **"What customers say"** reviews now play as a **carousel** - one review fills the section at a time, then slides on every few seconds. Visitors can swipe or use the left / right arrows and the dots below to move through them, and it pauses while someone is reading. One review, or none, still shows simply.

The whole **homepage and storefront now read in English, Mandarin or Bahasa Malaysia** via an **EN Mandarin BM** switch near the top of each page (English is the default, and a visitor's choice is remembered next visit). Every product can carry its own **translated shop names** - Products -> Edit -> **Shop names** - so shoppers in Mandarin/BM see a proper translated name on the order page; a blank box keeps the English name, which always stays the real name in the app, labels and history.

The thank-you after sending a review is warmer and honest: it reassures the sender that their review and photo reached the baker and will appear once approved, and a critical review (3 stars or fewer) gets an extra line of thanks for honest feedback. A reviewer now attaches their photo **two ways**: **Take photo** (camera) or **Choose photo** (gallery).

Developer contact is now in the app: under More -> Settings -> **Website & developer**, type the developer's name and one or more email addresses. A small **Website by ...** credit links from the bottom of the homepage,

the storefront footer and More -> About (each address tappable). Add the developer's optional **WhatsApp number** in the same card and that link opens WhatsApp (Business) with a ready "Hi!" - message the developer straight from the page - with the email links staying underneath. And adding a **software wish** on the More screen emails the developer the **full wish list** - every wish, ticked or not, newest first, with the app version, the project address and the date - to every address you listed; an **Email the full wish list** row on the card can send it from your own mail app too. No SQL this version.

v63 - Customers tab, wish list, holidays fix (8 Sep 2026)

Customers is a real tab of its own, swapped with **Purchase Order** (now under **More -> Purchase Order**). The tab stays your automatic customer list - who ordered, how much, their favourite - and now every person can carry a **profile** (the **Edit** button in their history pop-up): their dog's name and a small photo (auto-shrunk to a thumb), what they like, what to avoid, and a note. Profiles save into a synced customers collection, so both phones know the same people, and they build up quietly for a future AI chat about your customers. A **finder** box at the top of the tab filters the list as you type - name, number, dog's name, likes, notes - and says how many match, just like "Find an order" on Orders. A saved profile shows as a paw or photo beside the name and a small line under the row.

On **Home**, the "Upcoming holidays" card is fixed: it used to vanish whenever nothing was marked on the Delivery calendar. It is now always there - up to three marks at a glance, a scroll for any more, and a friendly empty state that still opens the calendar so you can add days.

On **More**, a small green **Engine v63** pill (in the header card) shows which build this phone runs, a **Software wish list** sits at the bottom and behaves exactly like your to-do list (add, tick, reword, remove - a tick stays ticked and never touches the weekly routine), and **Full change history** opens this history as a PDF kept on your website. No SQL this version.

v62 - "Not sharing" strip + read-only View (8 Sep 2026)

When a phone is not on the shared cloud - shared data off, never set up, or signed out - a thin amber strip at the very top of every screen reads **" Not sharing right now"** with a one-line reason and a **Fix it** tap straight to Settings -> Shared data. It never locks the app, and it disappears on its own the moment sharing is back on. (Set up once per phone: no SQL - the shared-data setup from v53 does everything.)

Also, every backup copy in Settings -> **Backup & safety** now has a **View** button: a purely read-only look inside that copy - its orders grouped by delivery date (customer, product, quantity, status), the product price list and ingredient stock at that time. Looking at a June price confirms the detail without moving today: no rewind, nothing written, ever.

v61 - Cloud backups (8 Sep 2026)

Every time the app opens while signed in, it quietly saves a full **Daily** copy (first open each day), a **Weekly** copy (Mondays) and a **Monthly** copy (the 1st) to your own Supabase cloud, keeping the newest 7 daily / 4 weekly / 3 monthly. A "Back up to cloud now" button, and one "Before restore" copy saved before every restore, stay until you delete them. Each copy can be **Restored** (steps the phone - and the shared cloud - back to that copy; a "Before restore" copy is always saved first, so it is never one-way), **Downloaded** as jienluv2bake-backup-YYYY-MM-DD-KIND.json (which re-imports like an Export), or **Deleted**. Lives under More -> Settings -> "Backup & safety". The copies never contain the per-phone app-login email/password or the app password. Setup: run supabase/backups.sql once.

v60 - Homepage customer reviews (8 Sep 2026)

Customers can leave a **What customers say** review on the homepage - name, 1-5 stars, a message in English / Mandarin / Bahasa Malaysia, and an optional photo. Every review lands *unpublished* in More -> Reviews, where **Publish** shows it on the homepage, **Take down** hides it, and **Delete** removes it - nothing becomes public until you tap Publish. Reviews appear on the homepage only. Setup: run supabase/reviews.sql once.

v59 - Occasion import tidy-up (8 Sep 2026)

Added **World Animal Day** (4 Oct) to the fun/pet-day import, plus whole-group ticking - tick a heading box, or use **Untick all / Tick all**.

v58 - Occasion import (8 Sep 2026)

A one-tap **+ Add occasion** import of Malaysia's days plus fun days (pet, baking, kindness), and a **My own day** quick-add.

v57 - Private notes per product ingredient (8 Sep 2026)

Each ingredient line in a product can carry a short private description *for that product only* (e.g. which flour that product uses). Typed and seen on the Products screens only - never on the shop, a label, or the product cards.

v56 - "Keep at least" reserve (8 Sep 2026)

Ingredients can set a **Keep at least** level. The order list tops an ingredient back up to that level in whole packs when your stock drops more than 10% under it, and low items ride along even when today's bakes don't use them. The Ingredients screen shows a red low-stock strip. Guide synced to v56.

v55 - Time & length units (7 Sep 2026)

min/hr and **cm/m** are pre-loaded in Units (under More), like g/kg - so bake times and pan sizes can be measured and converted (1 hr = 60 min, 1 m = 100 cm). Baking stock is guarded so a time or length unit can never distort a gram count.

v54 - Ingredients on hand (7 Sep 2026)

Ingredients now have **On hand** stock. The PO buys only what you don't have (rows you already have drop out), "Bought" on a saved list adds its packs once, and marking an order **Baked** subtracts its ingredients (undo restores them).

This history covers the two-phone cloud era (v54+). Earlier versions (pre-v54) were never recorded version by version, so they are not listed here rather than invented. Each new version is added here as it ships.